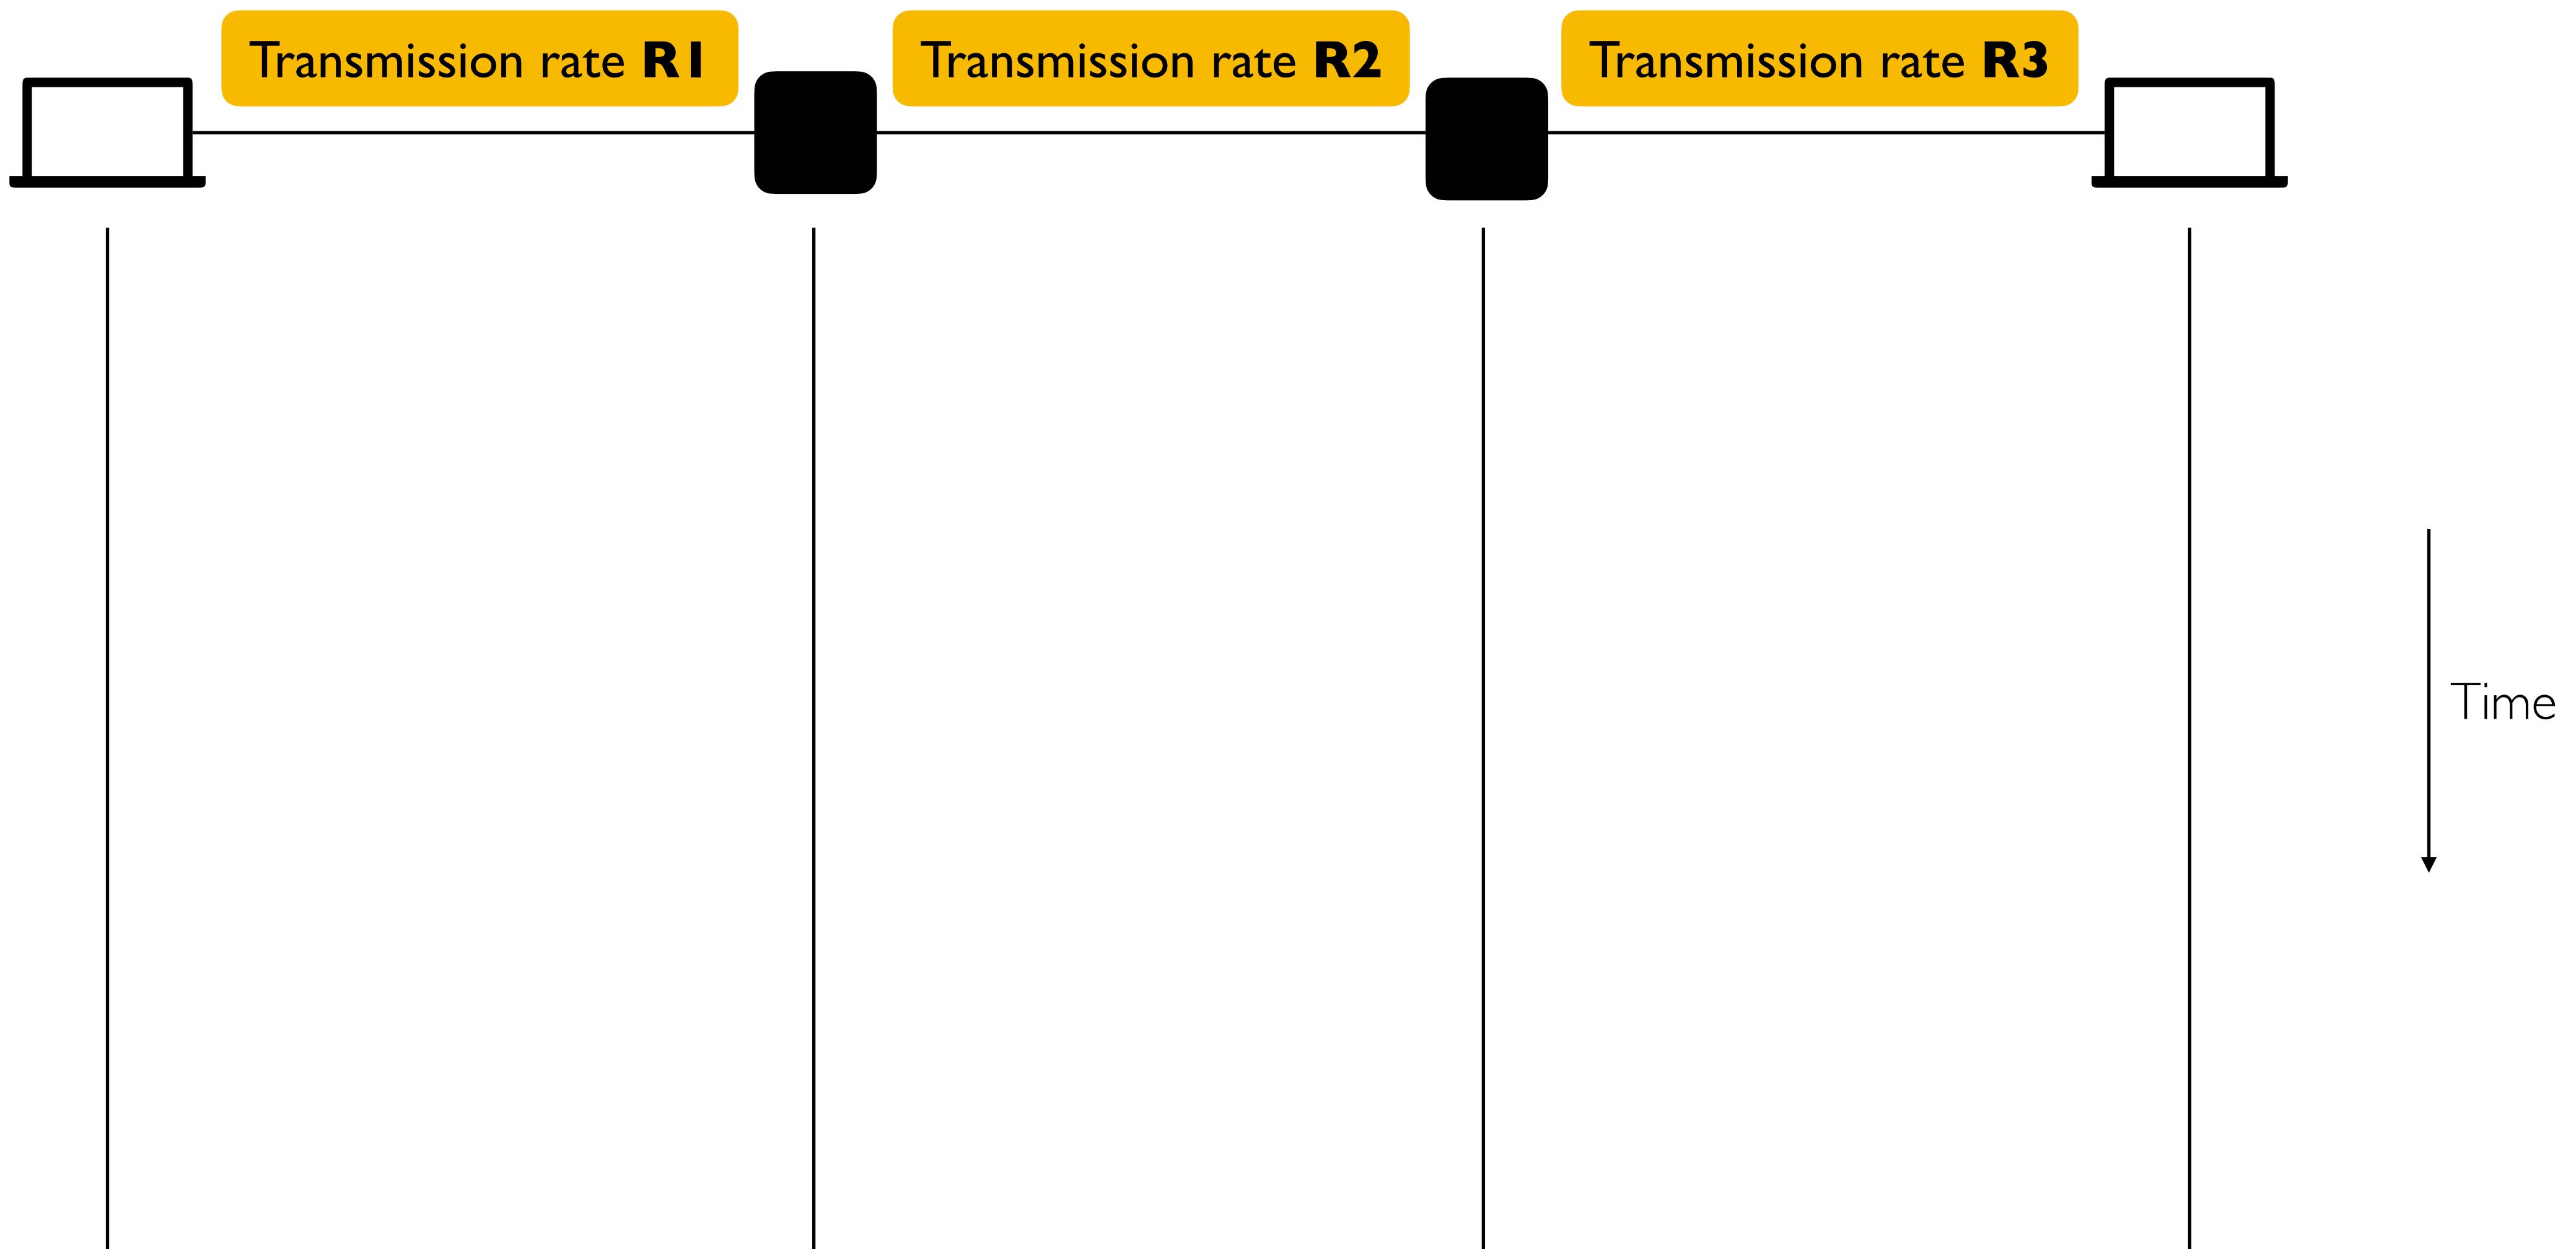


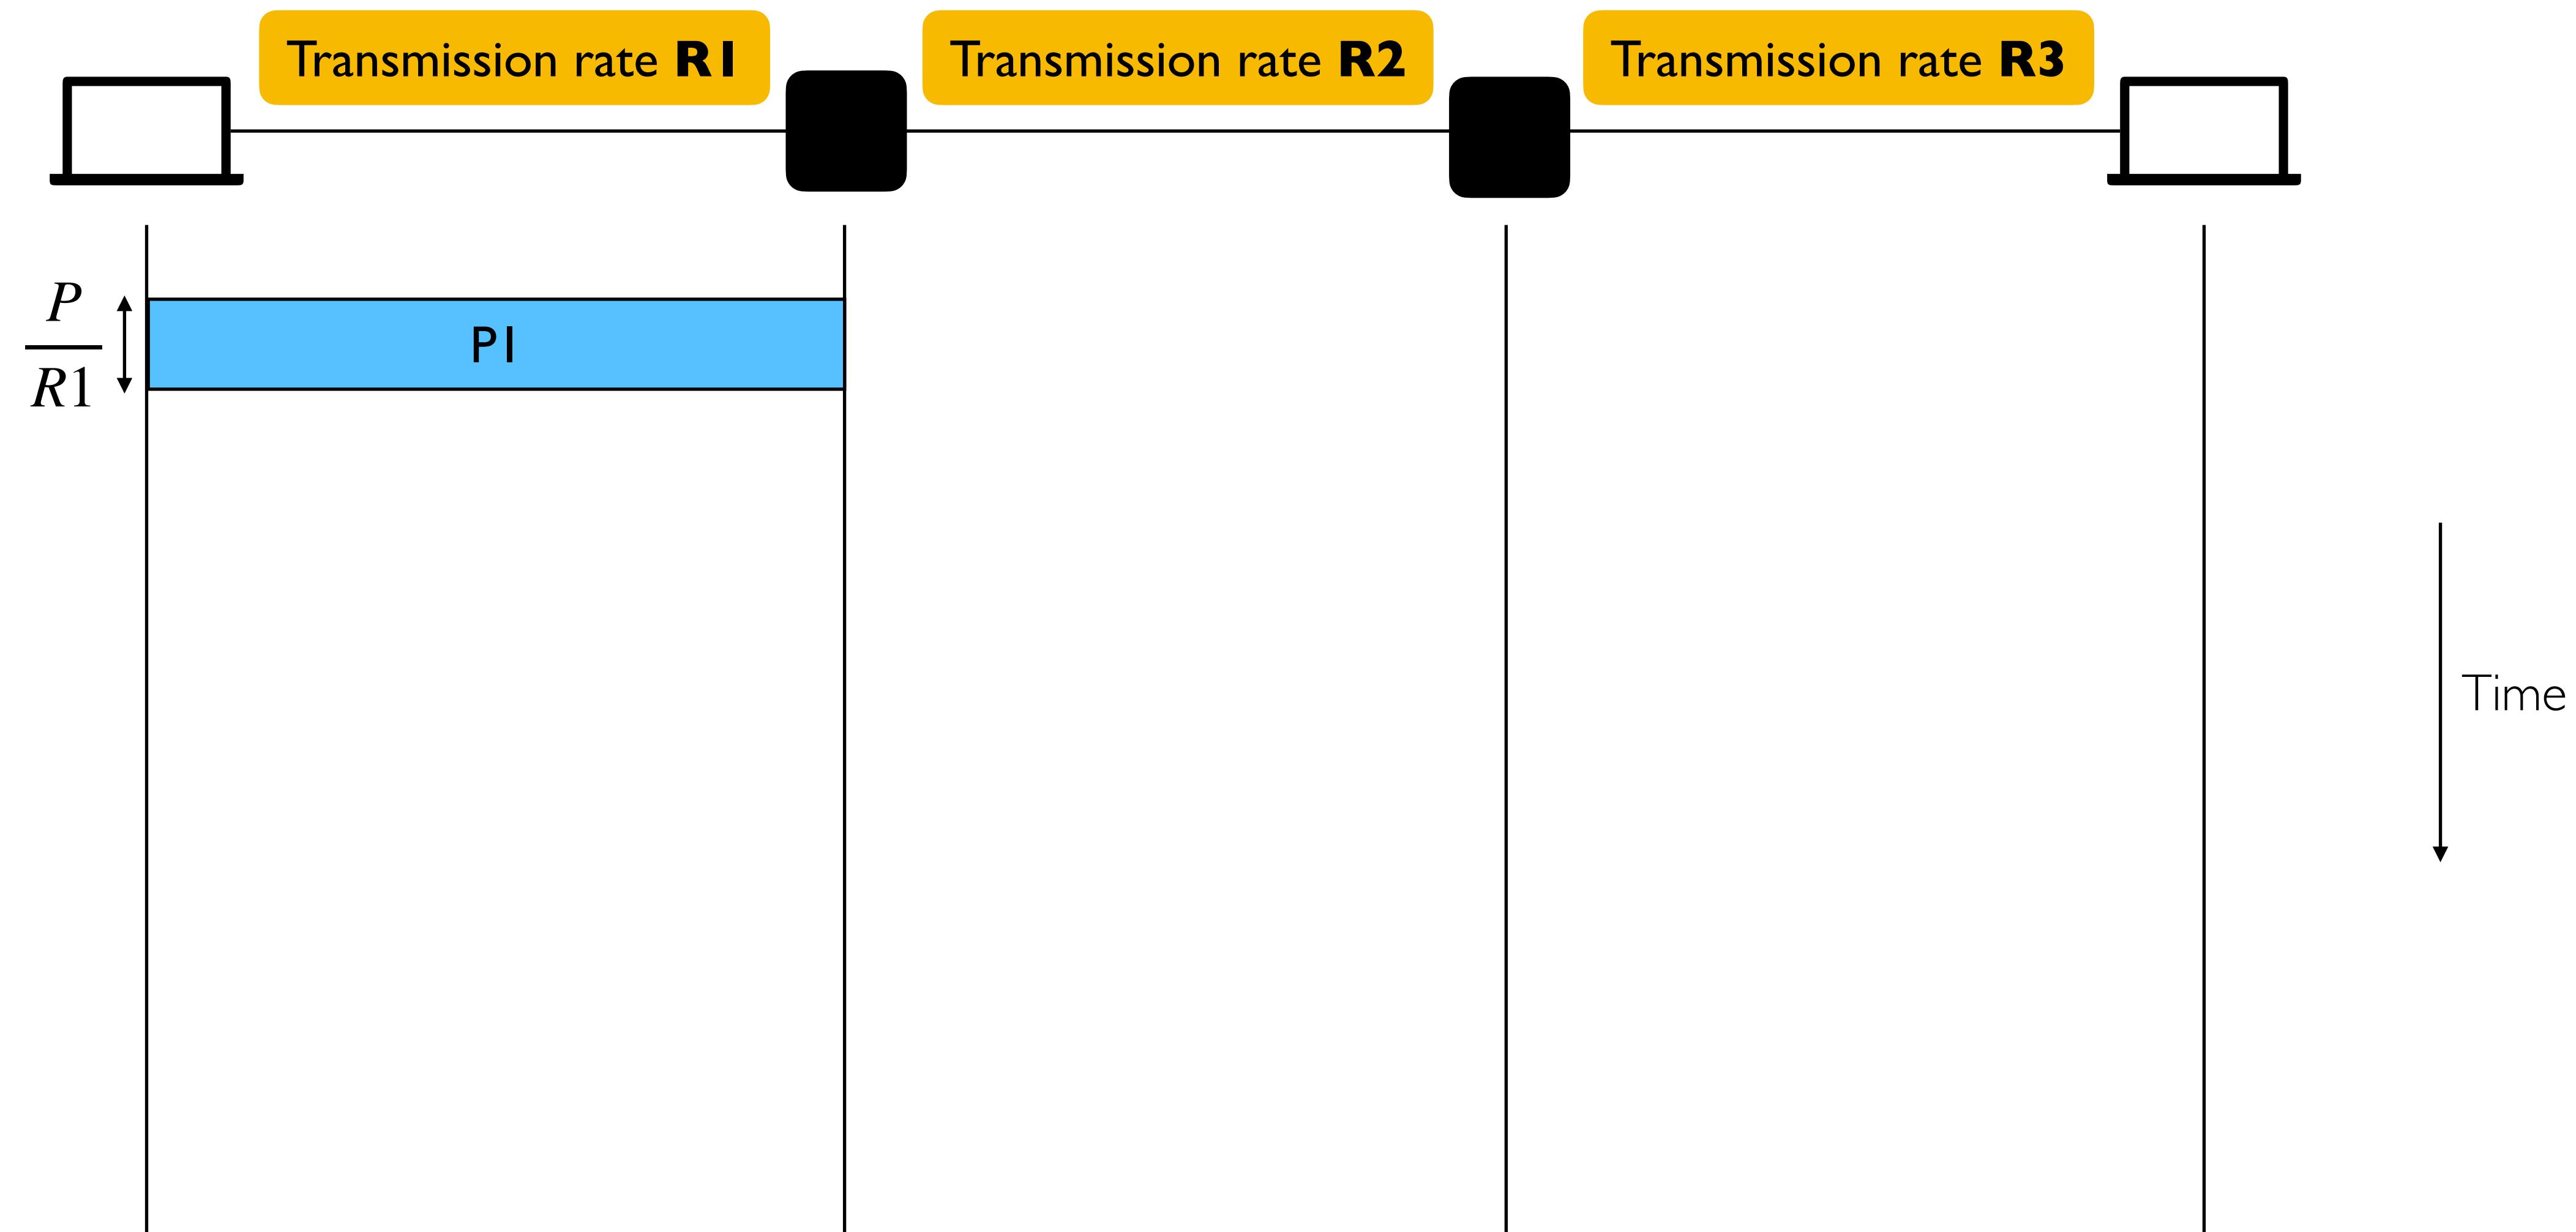
# How is communication organized? (Contd.)

CPSC 433/533, Spring 2021  
Anurag Khandelwal

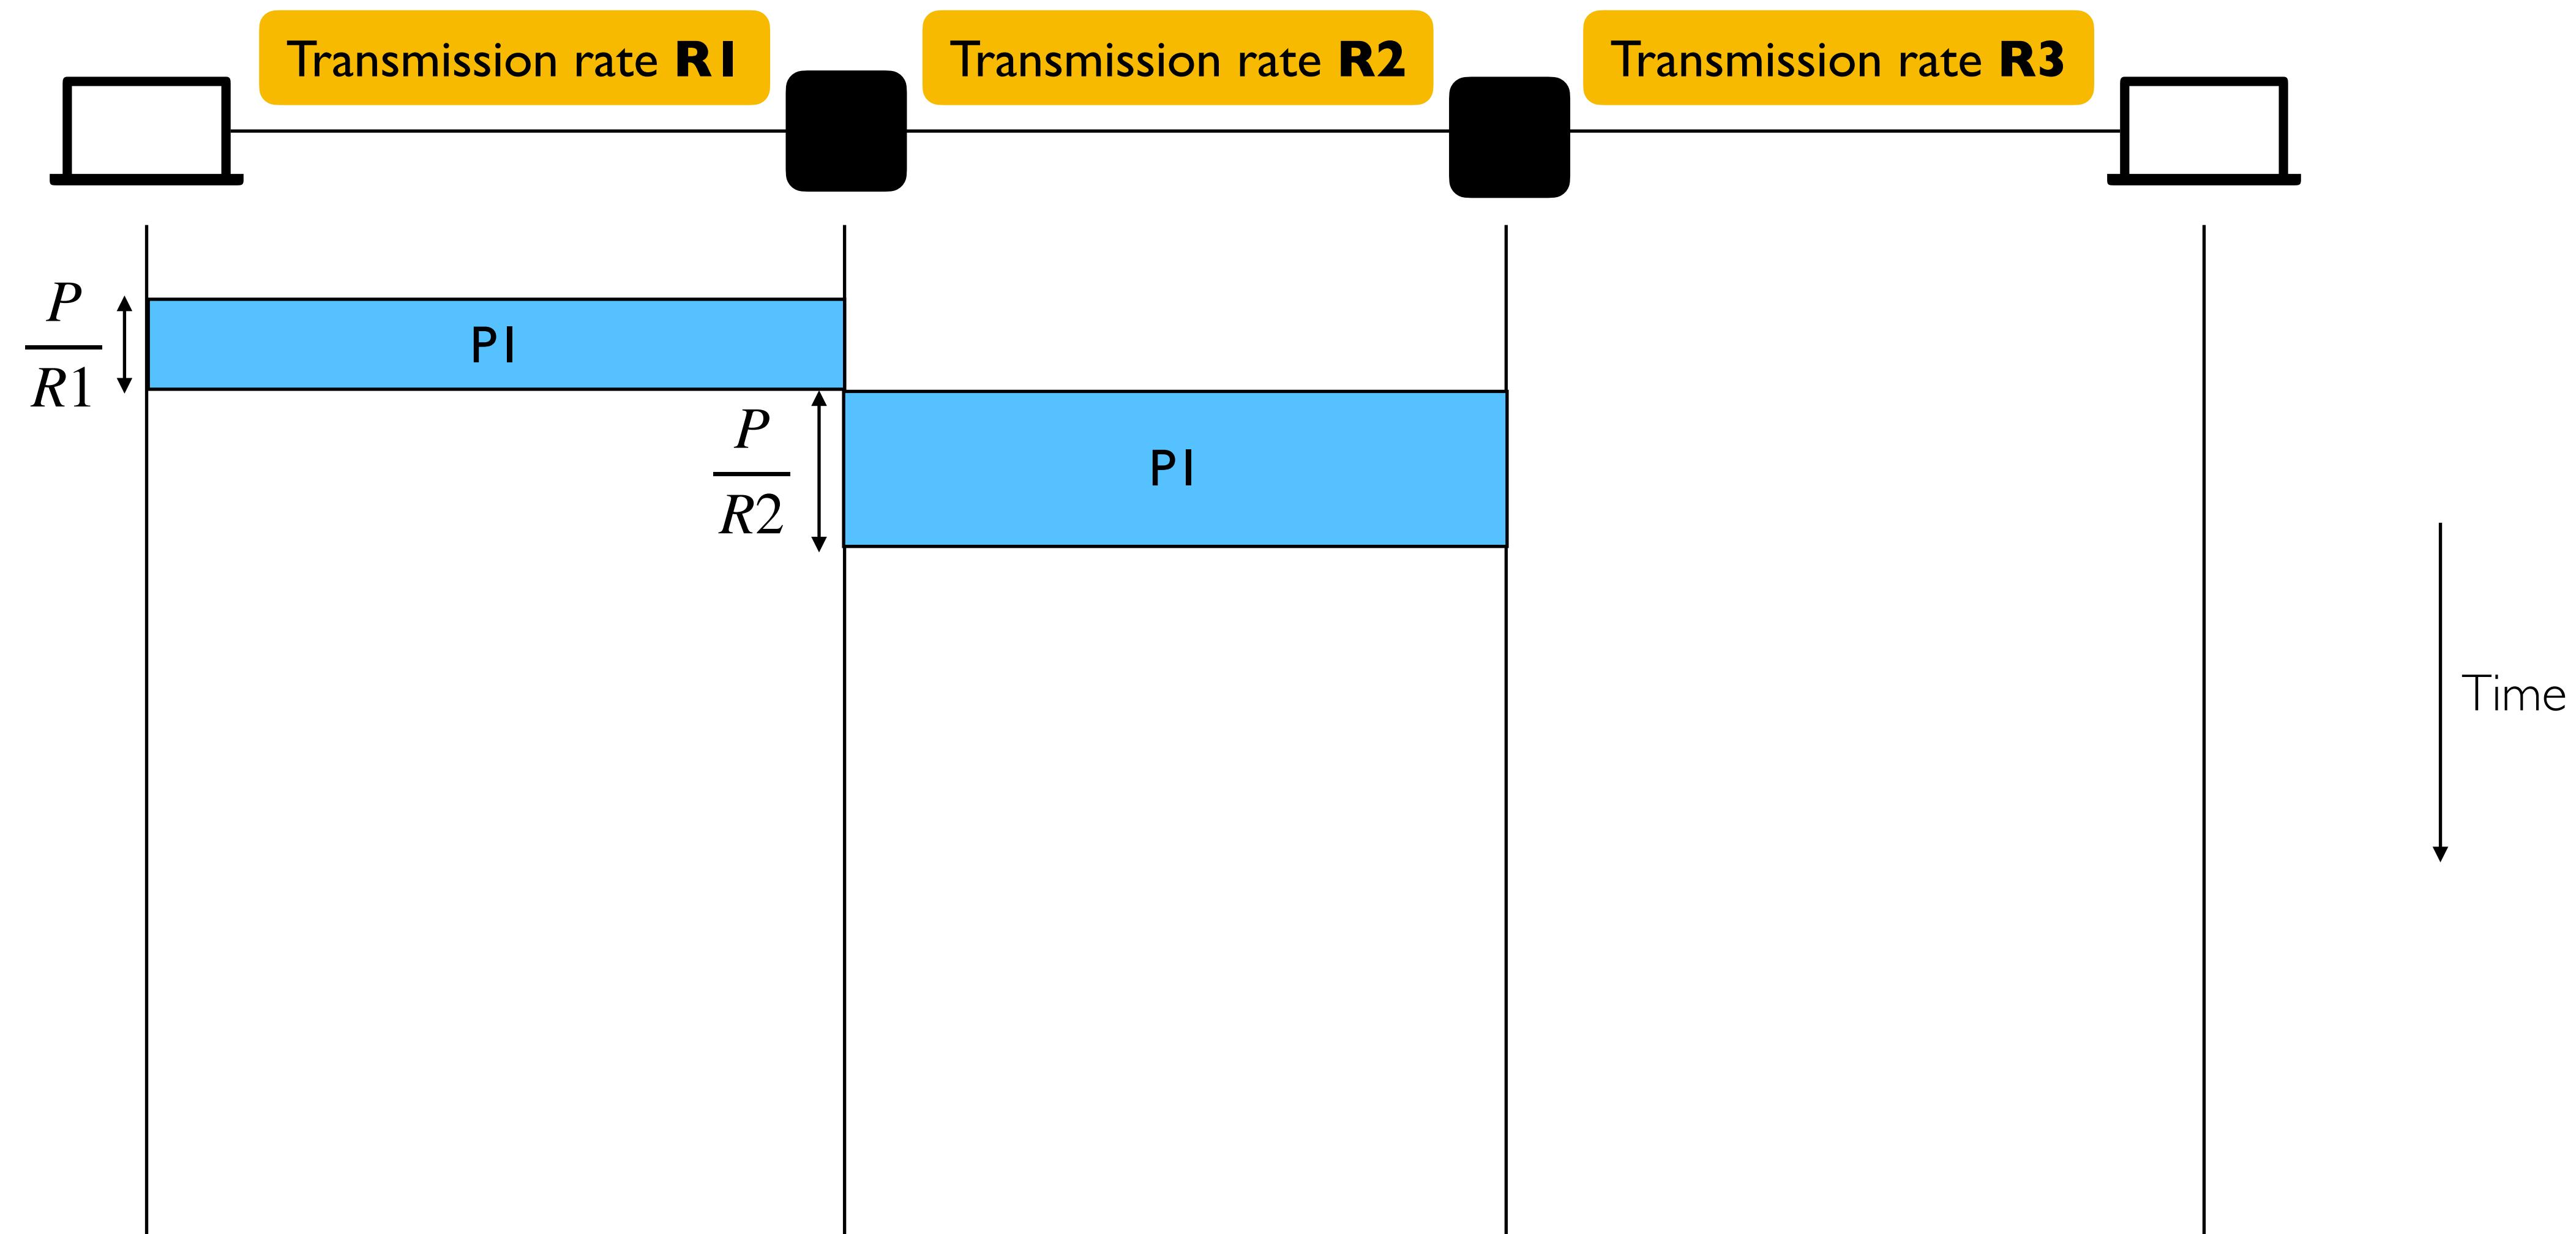
# But before that... Throughput Calculation



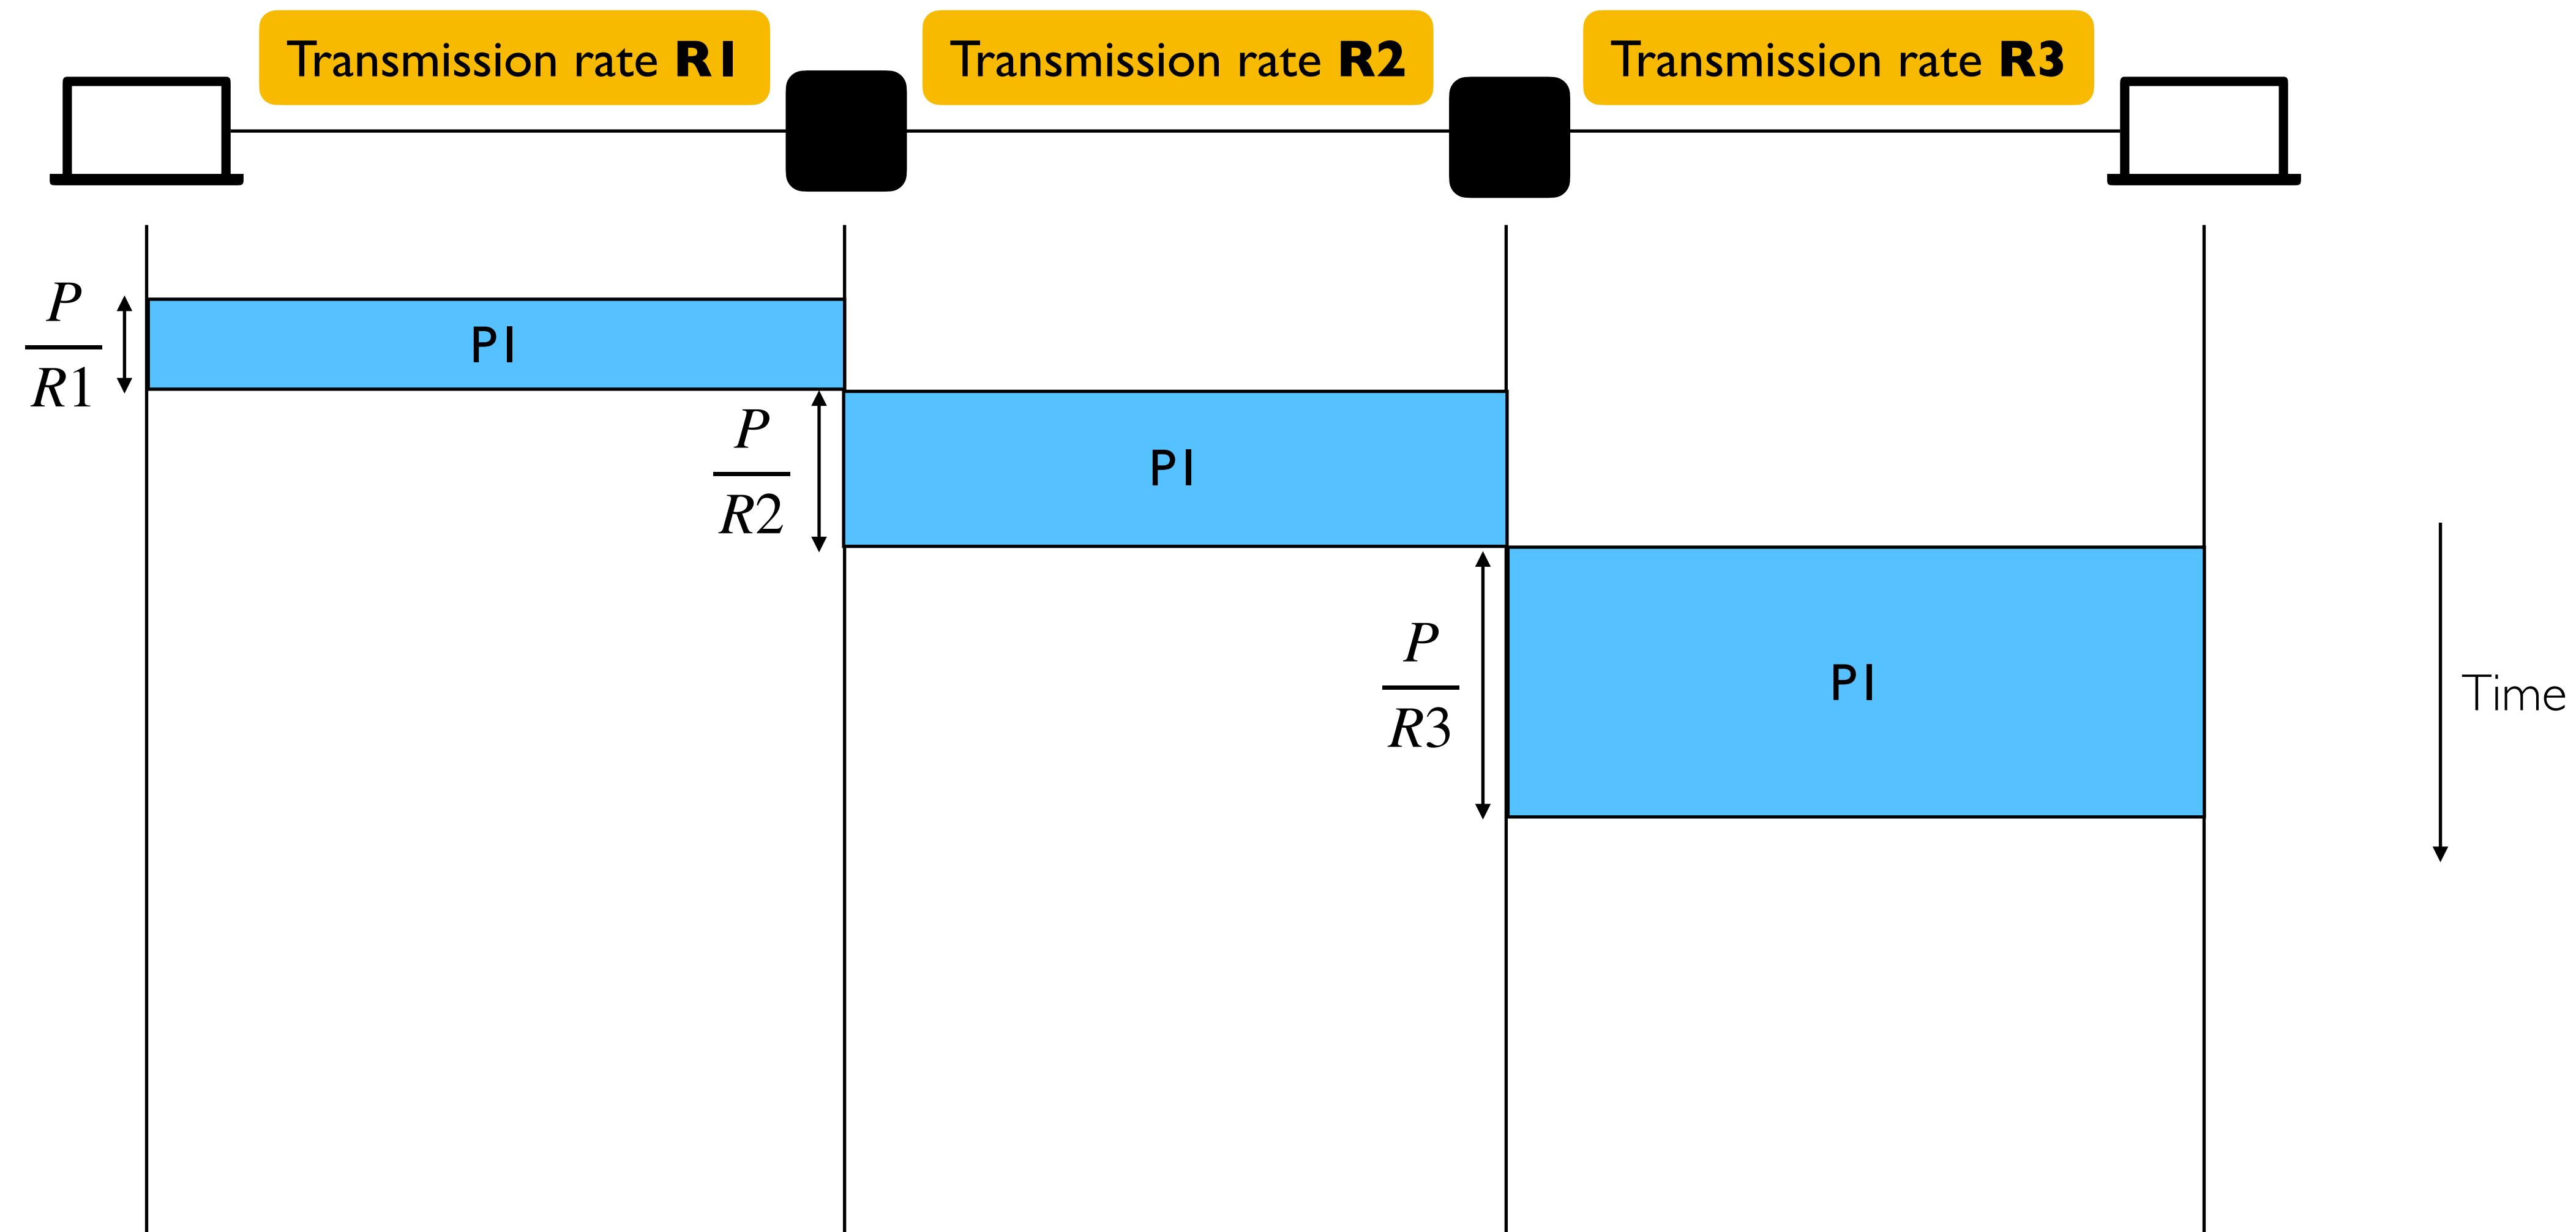
# But before that... Throughput Calculation



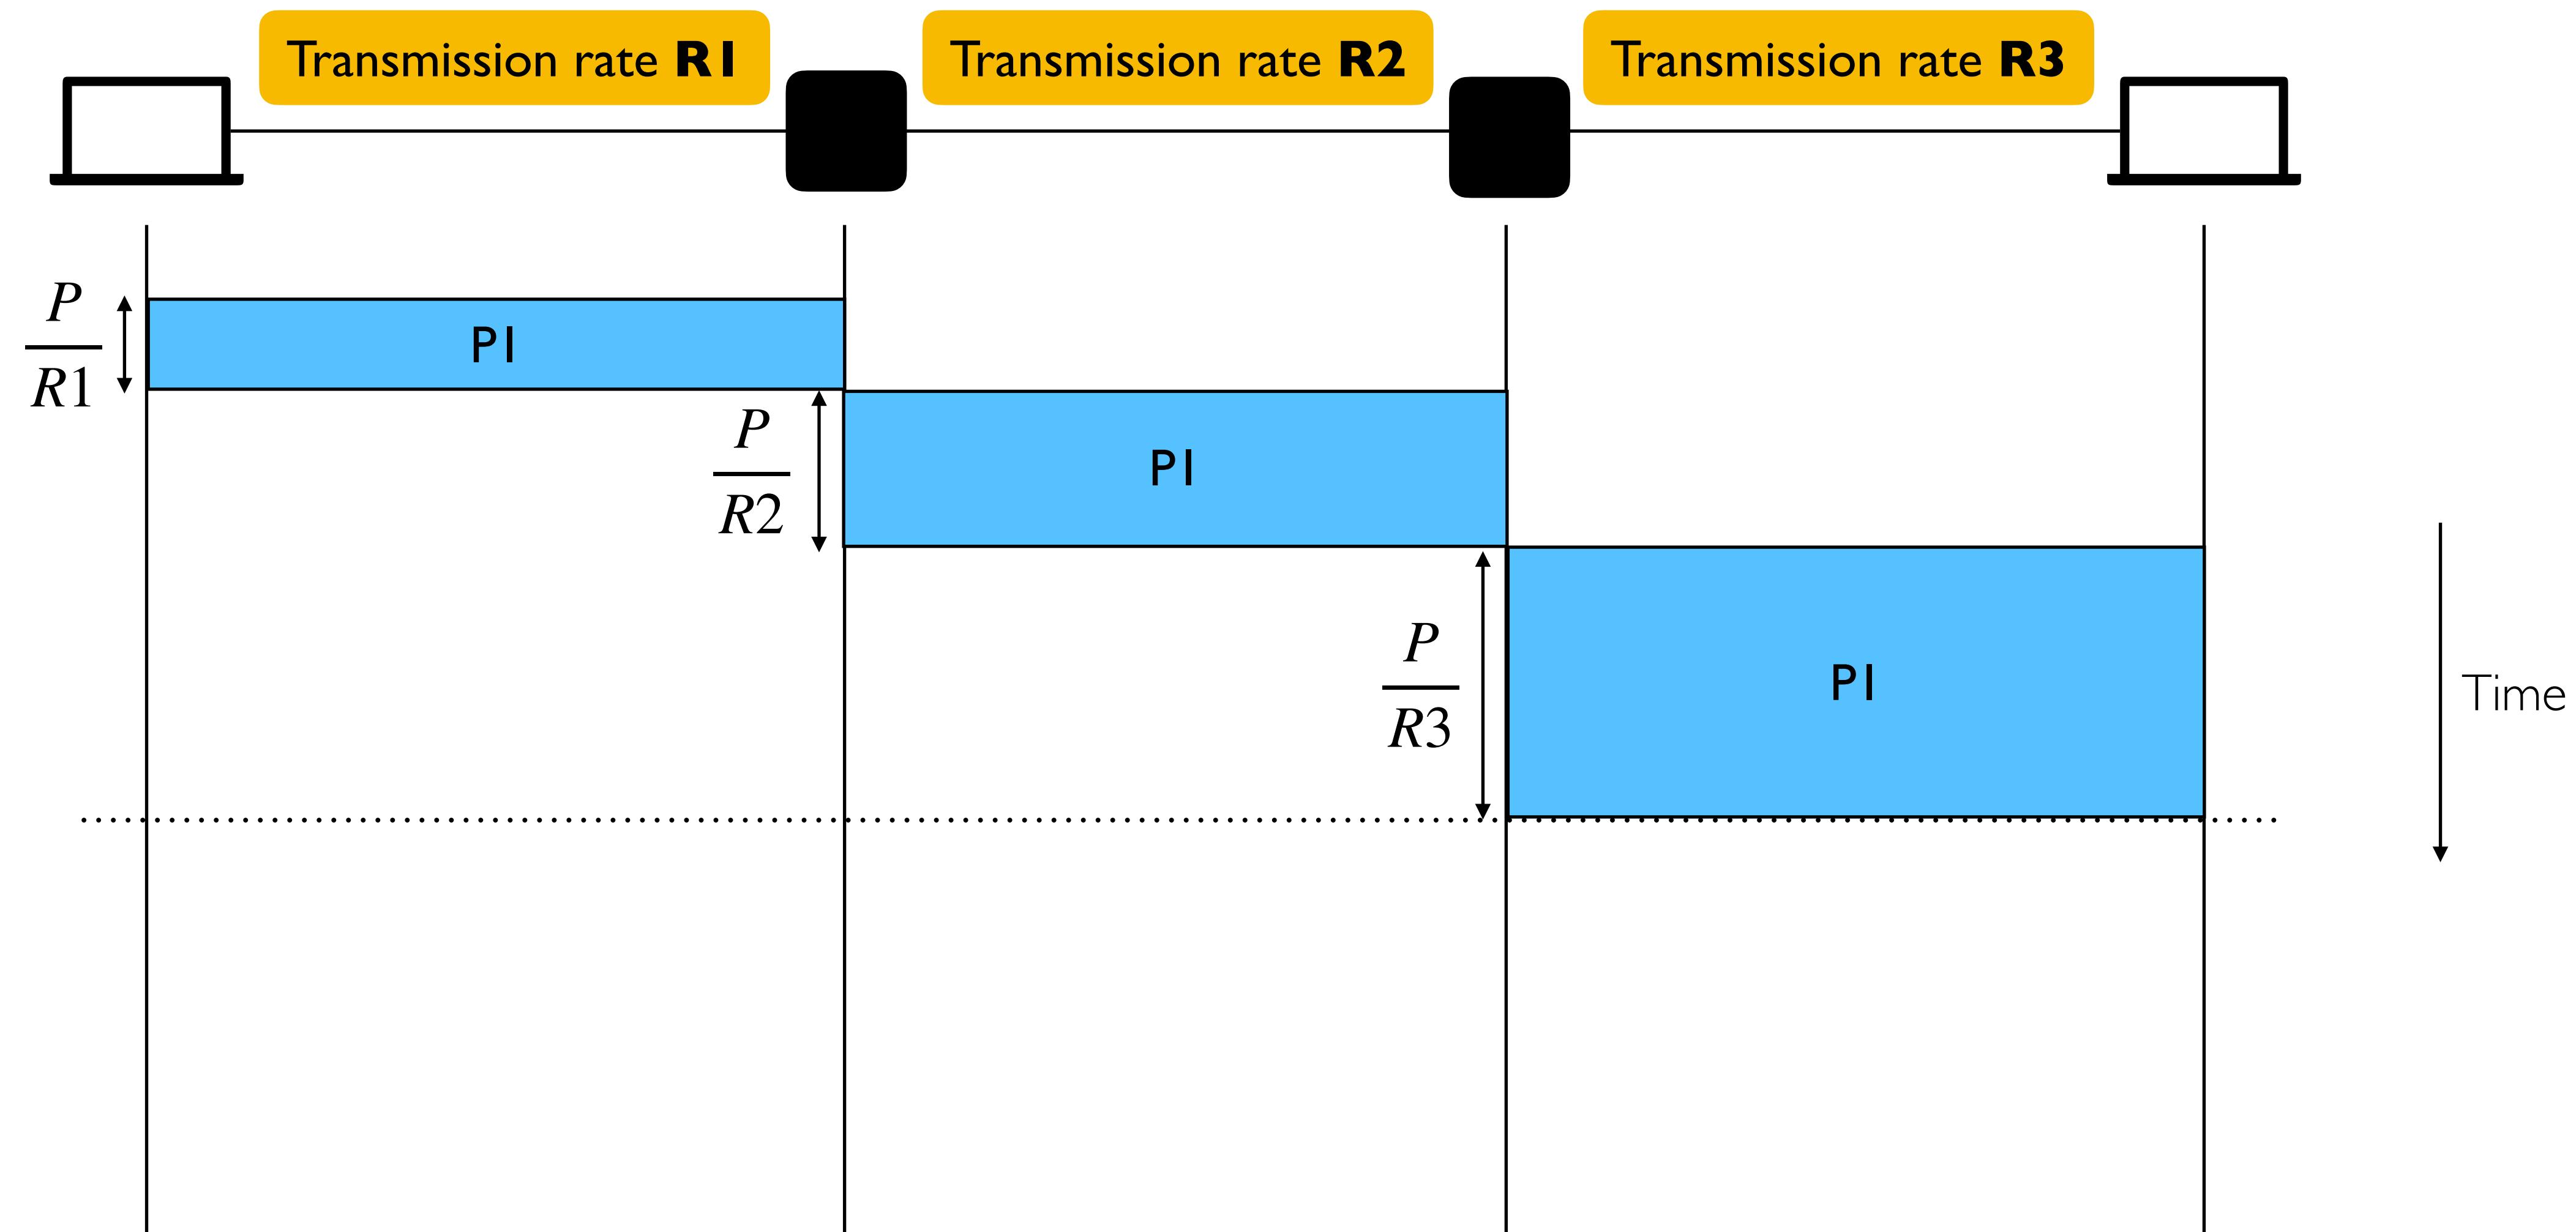
# But before that... Throughput Calculation



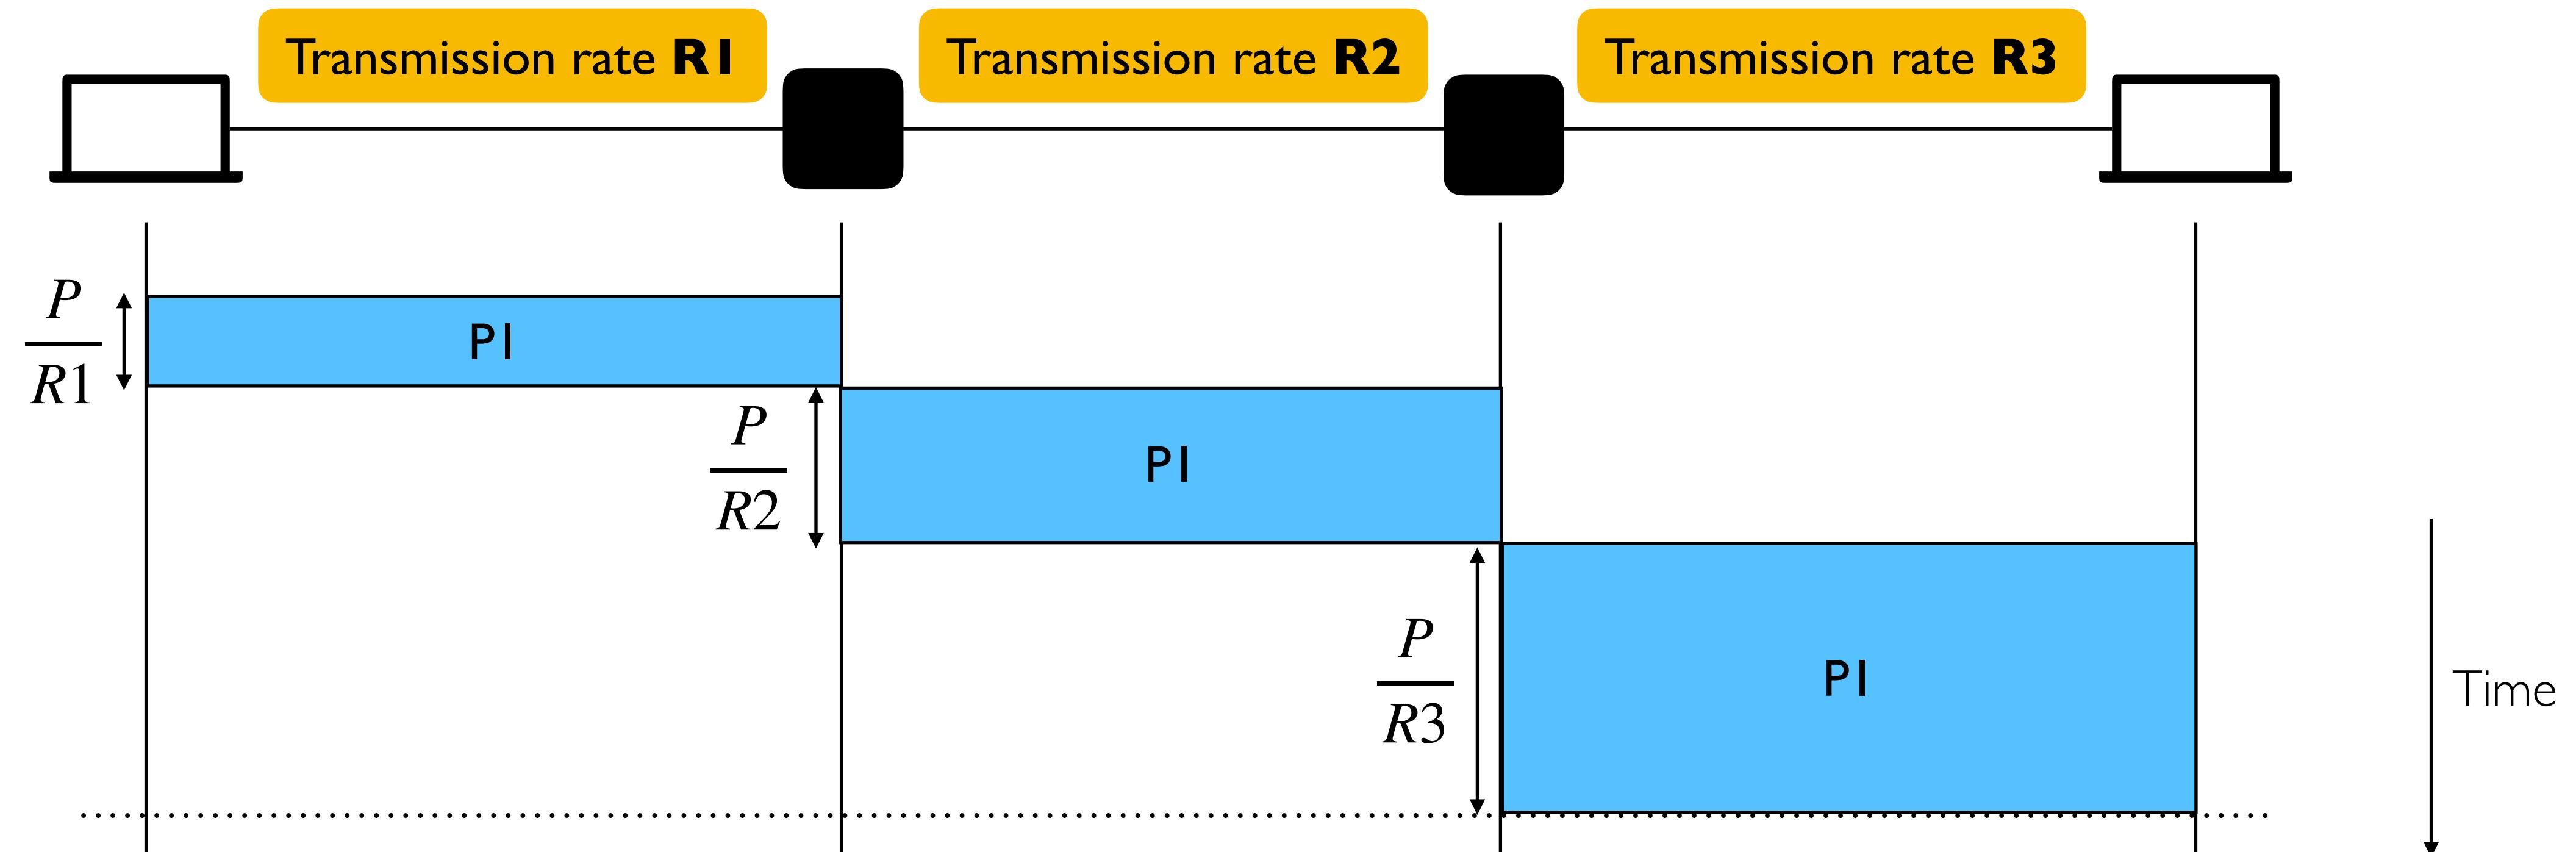
# But before that... Throughput Calculation



# But before that... Throughput Calculation

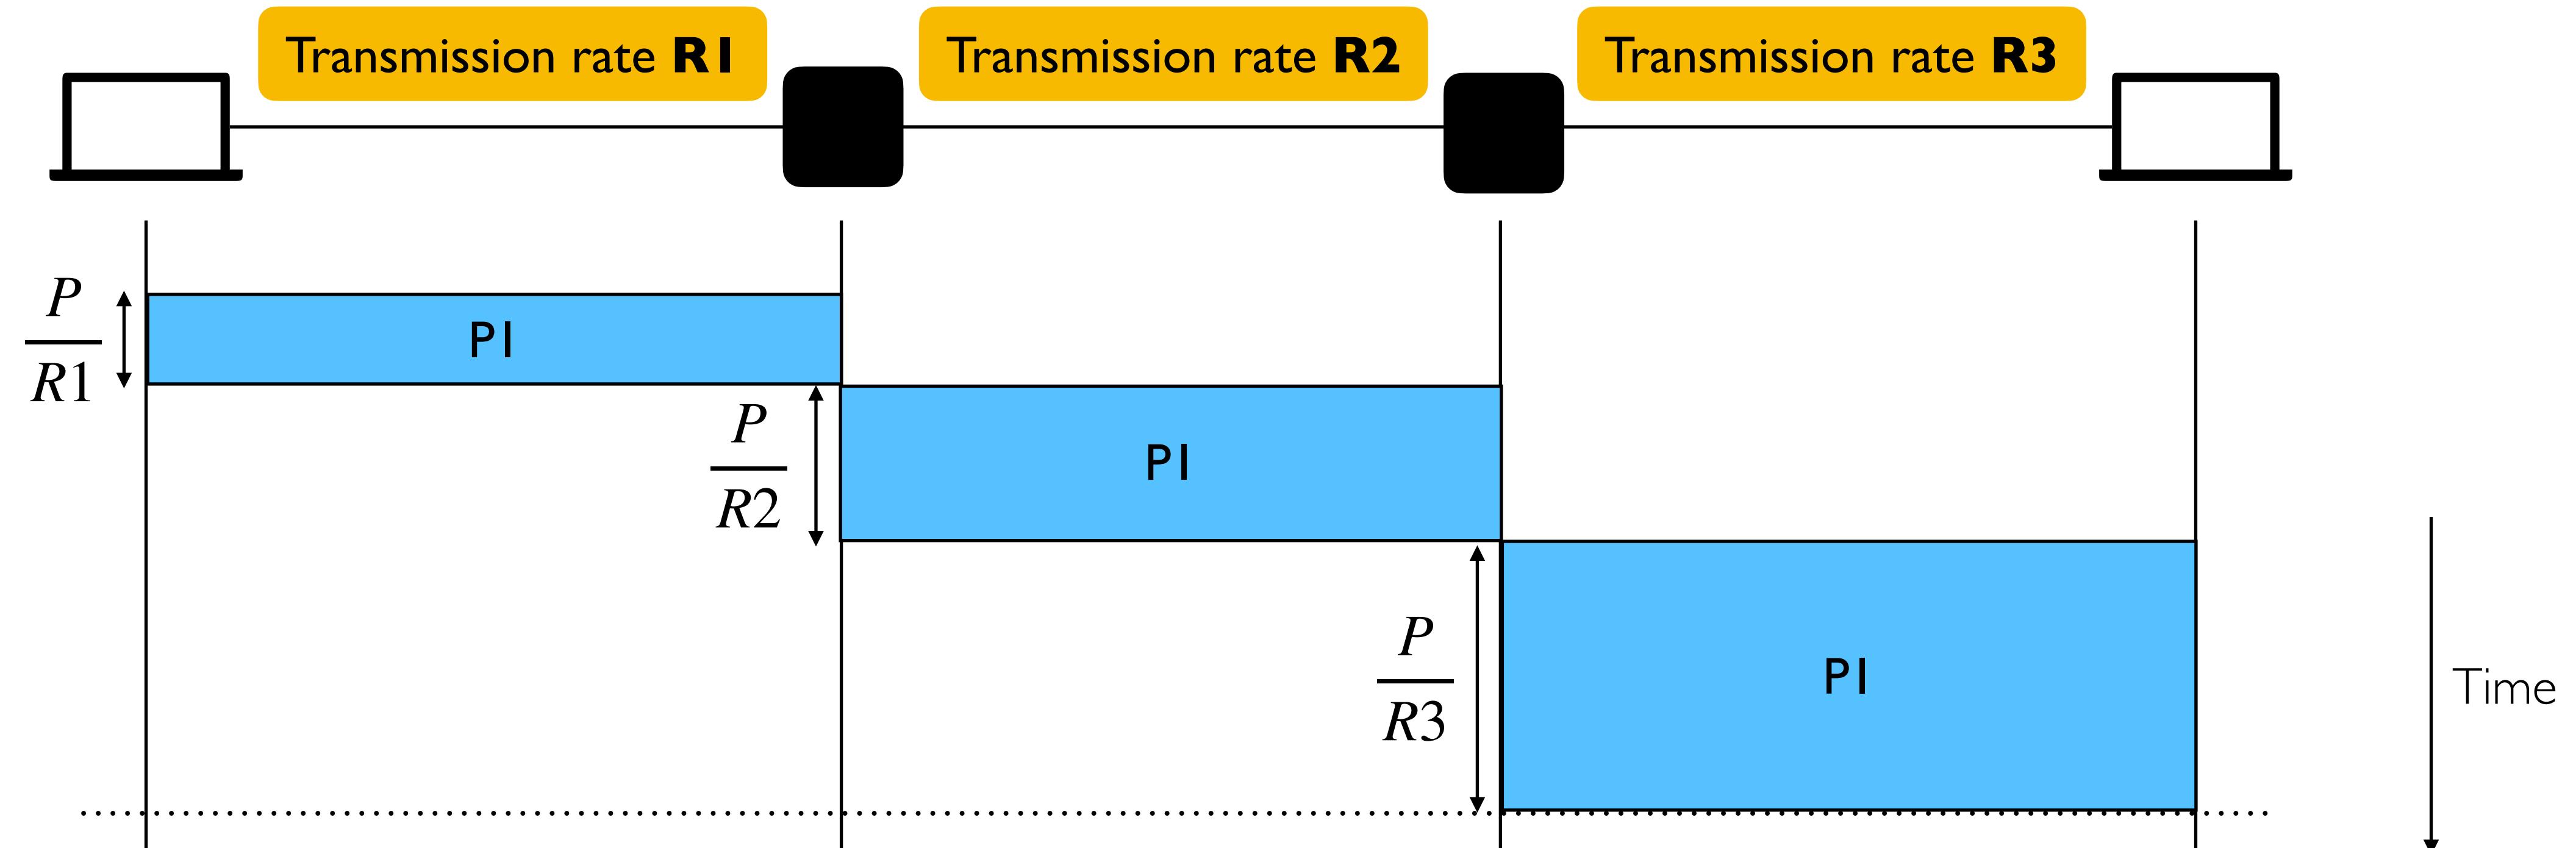


# But before that... Throughput Calculation



Ignoring propagation, queueing and processing delays, packet delay for a packet of  $P$  bits =  $\frac{P}{R_1} + \frac{P}{R_2} + \frac{P}{R_3}$

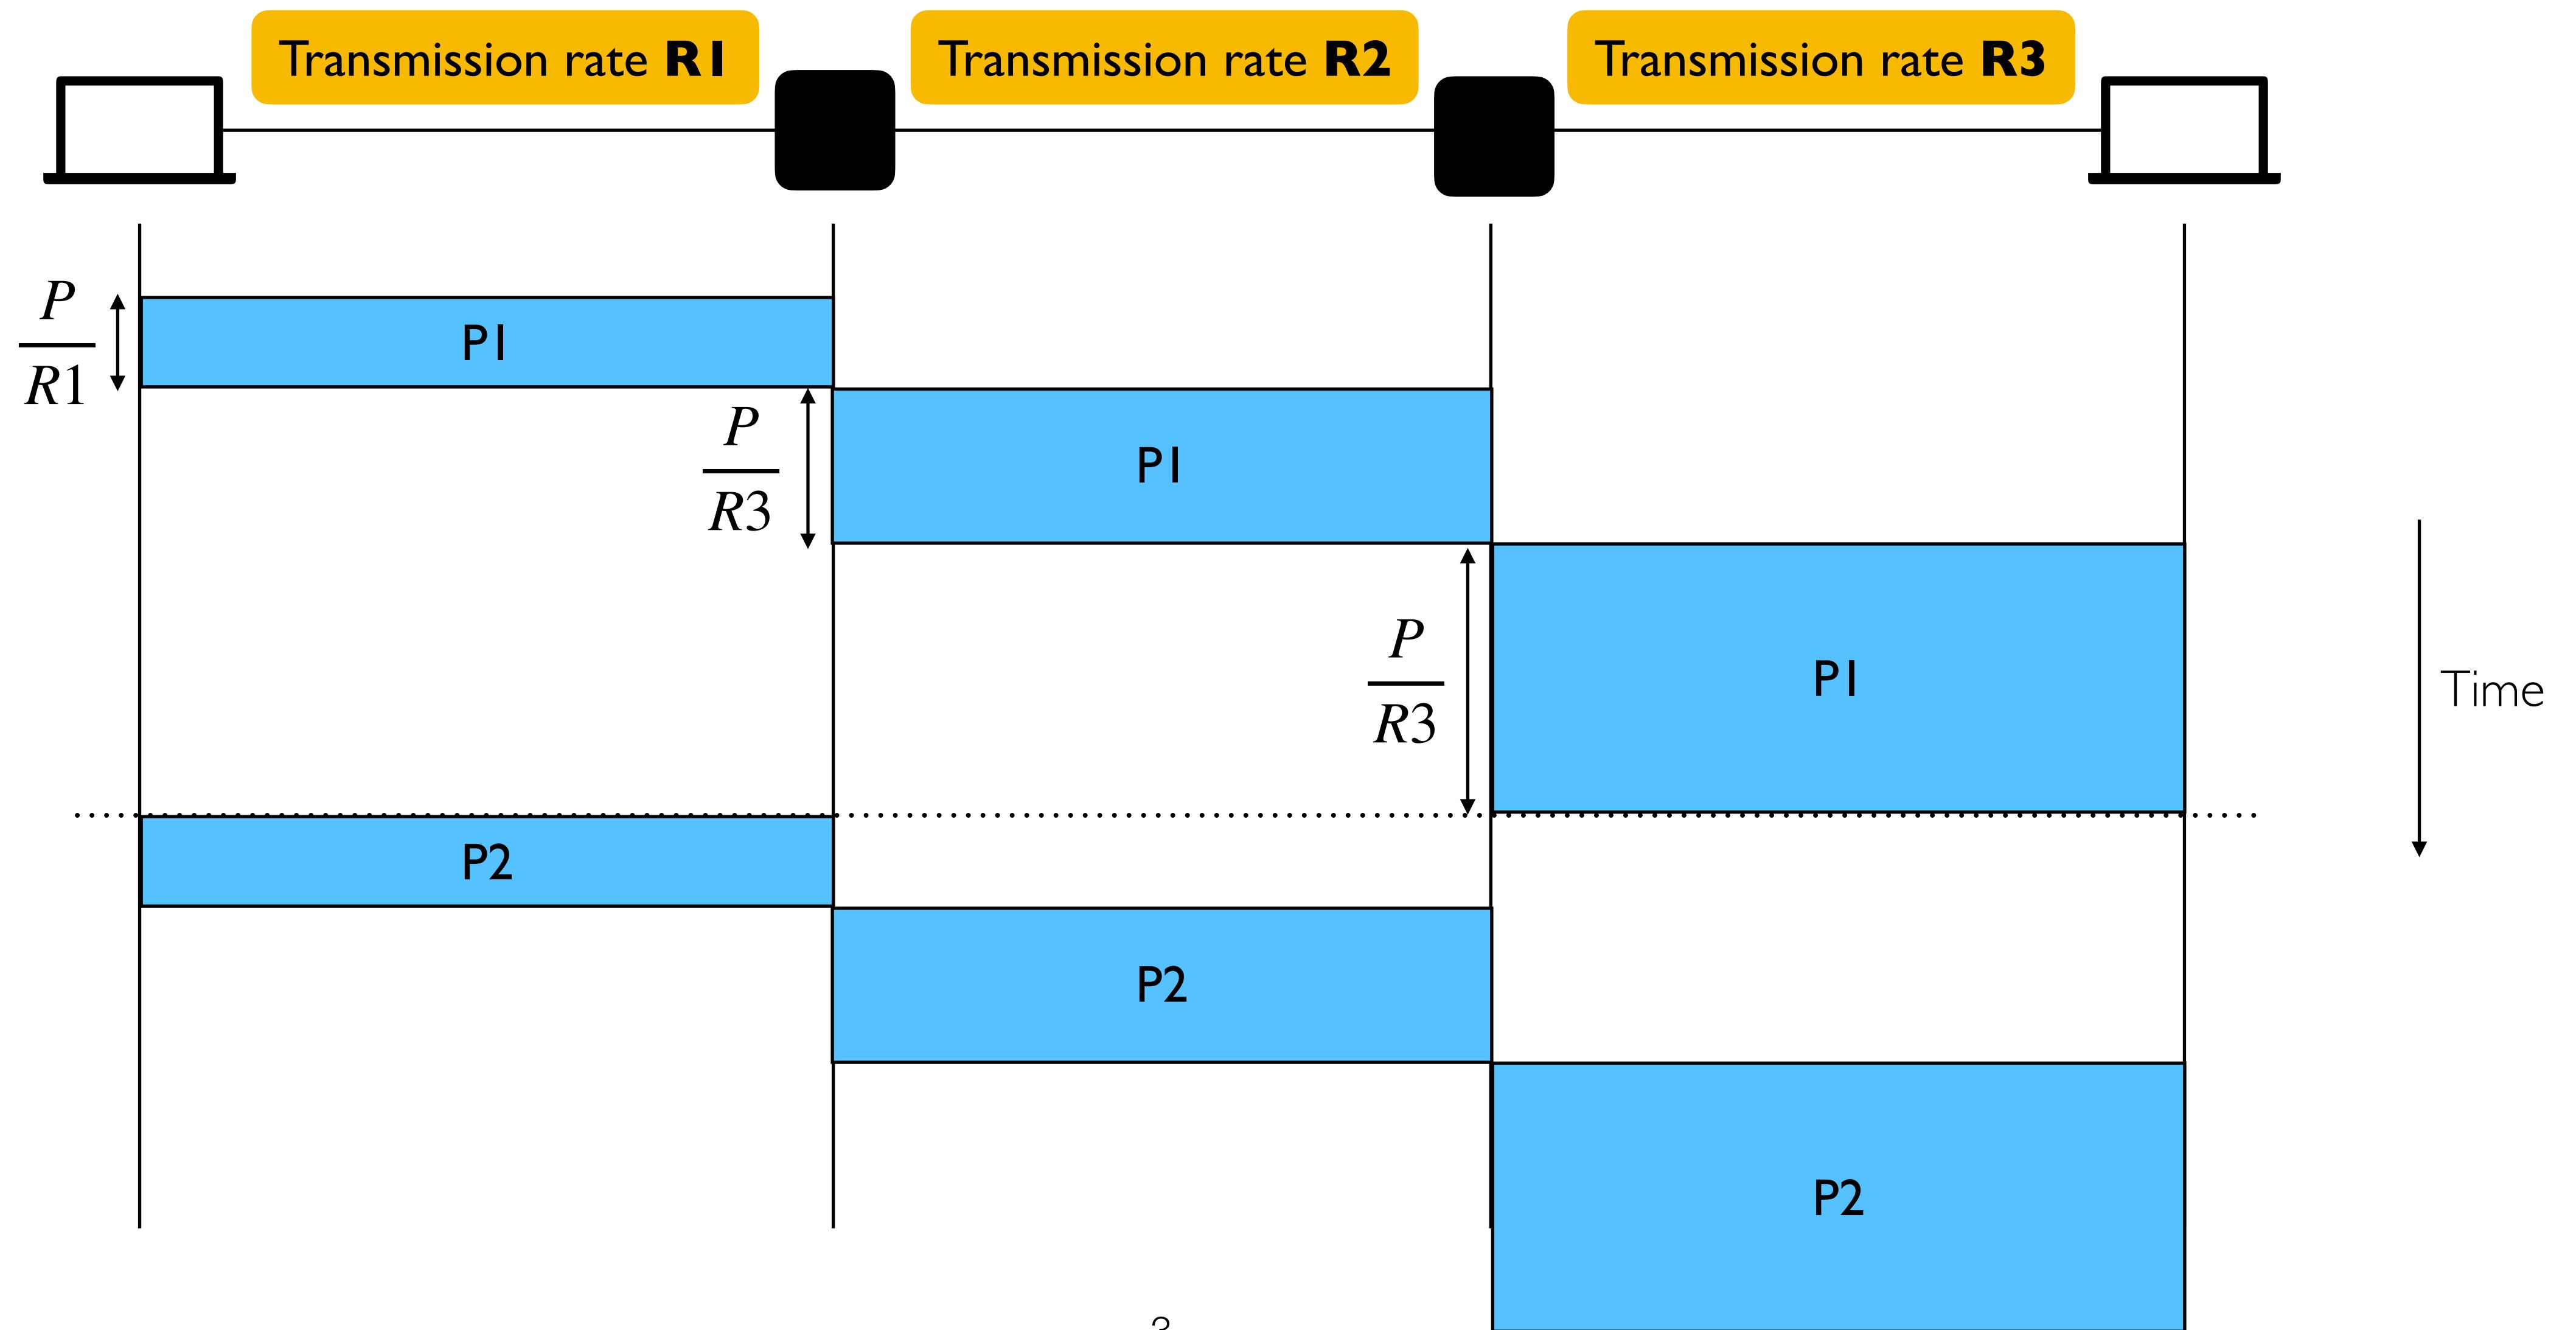
# But before that... Throughput Calculation



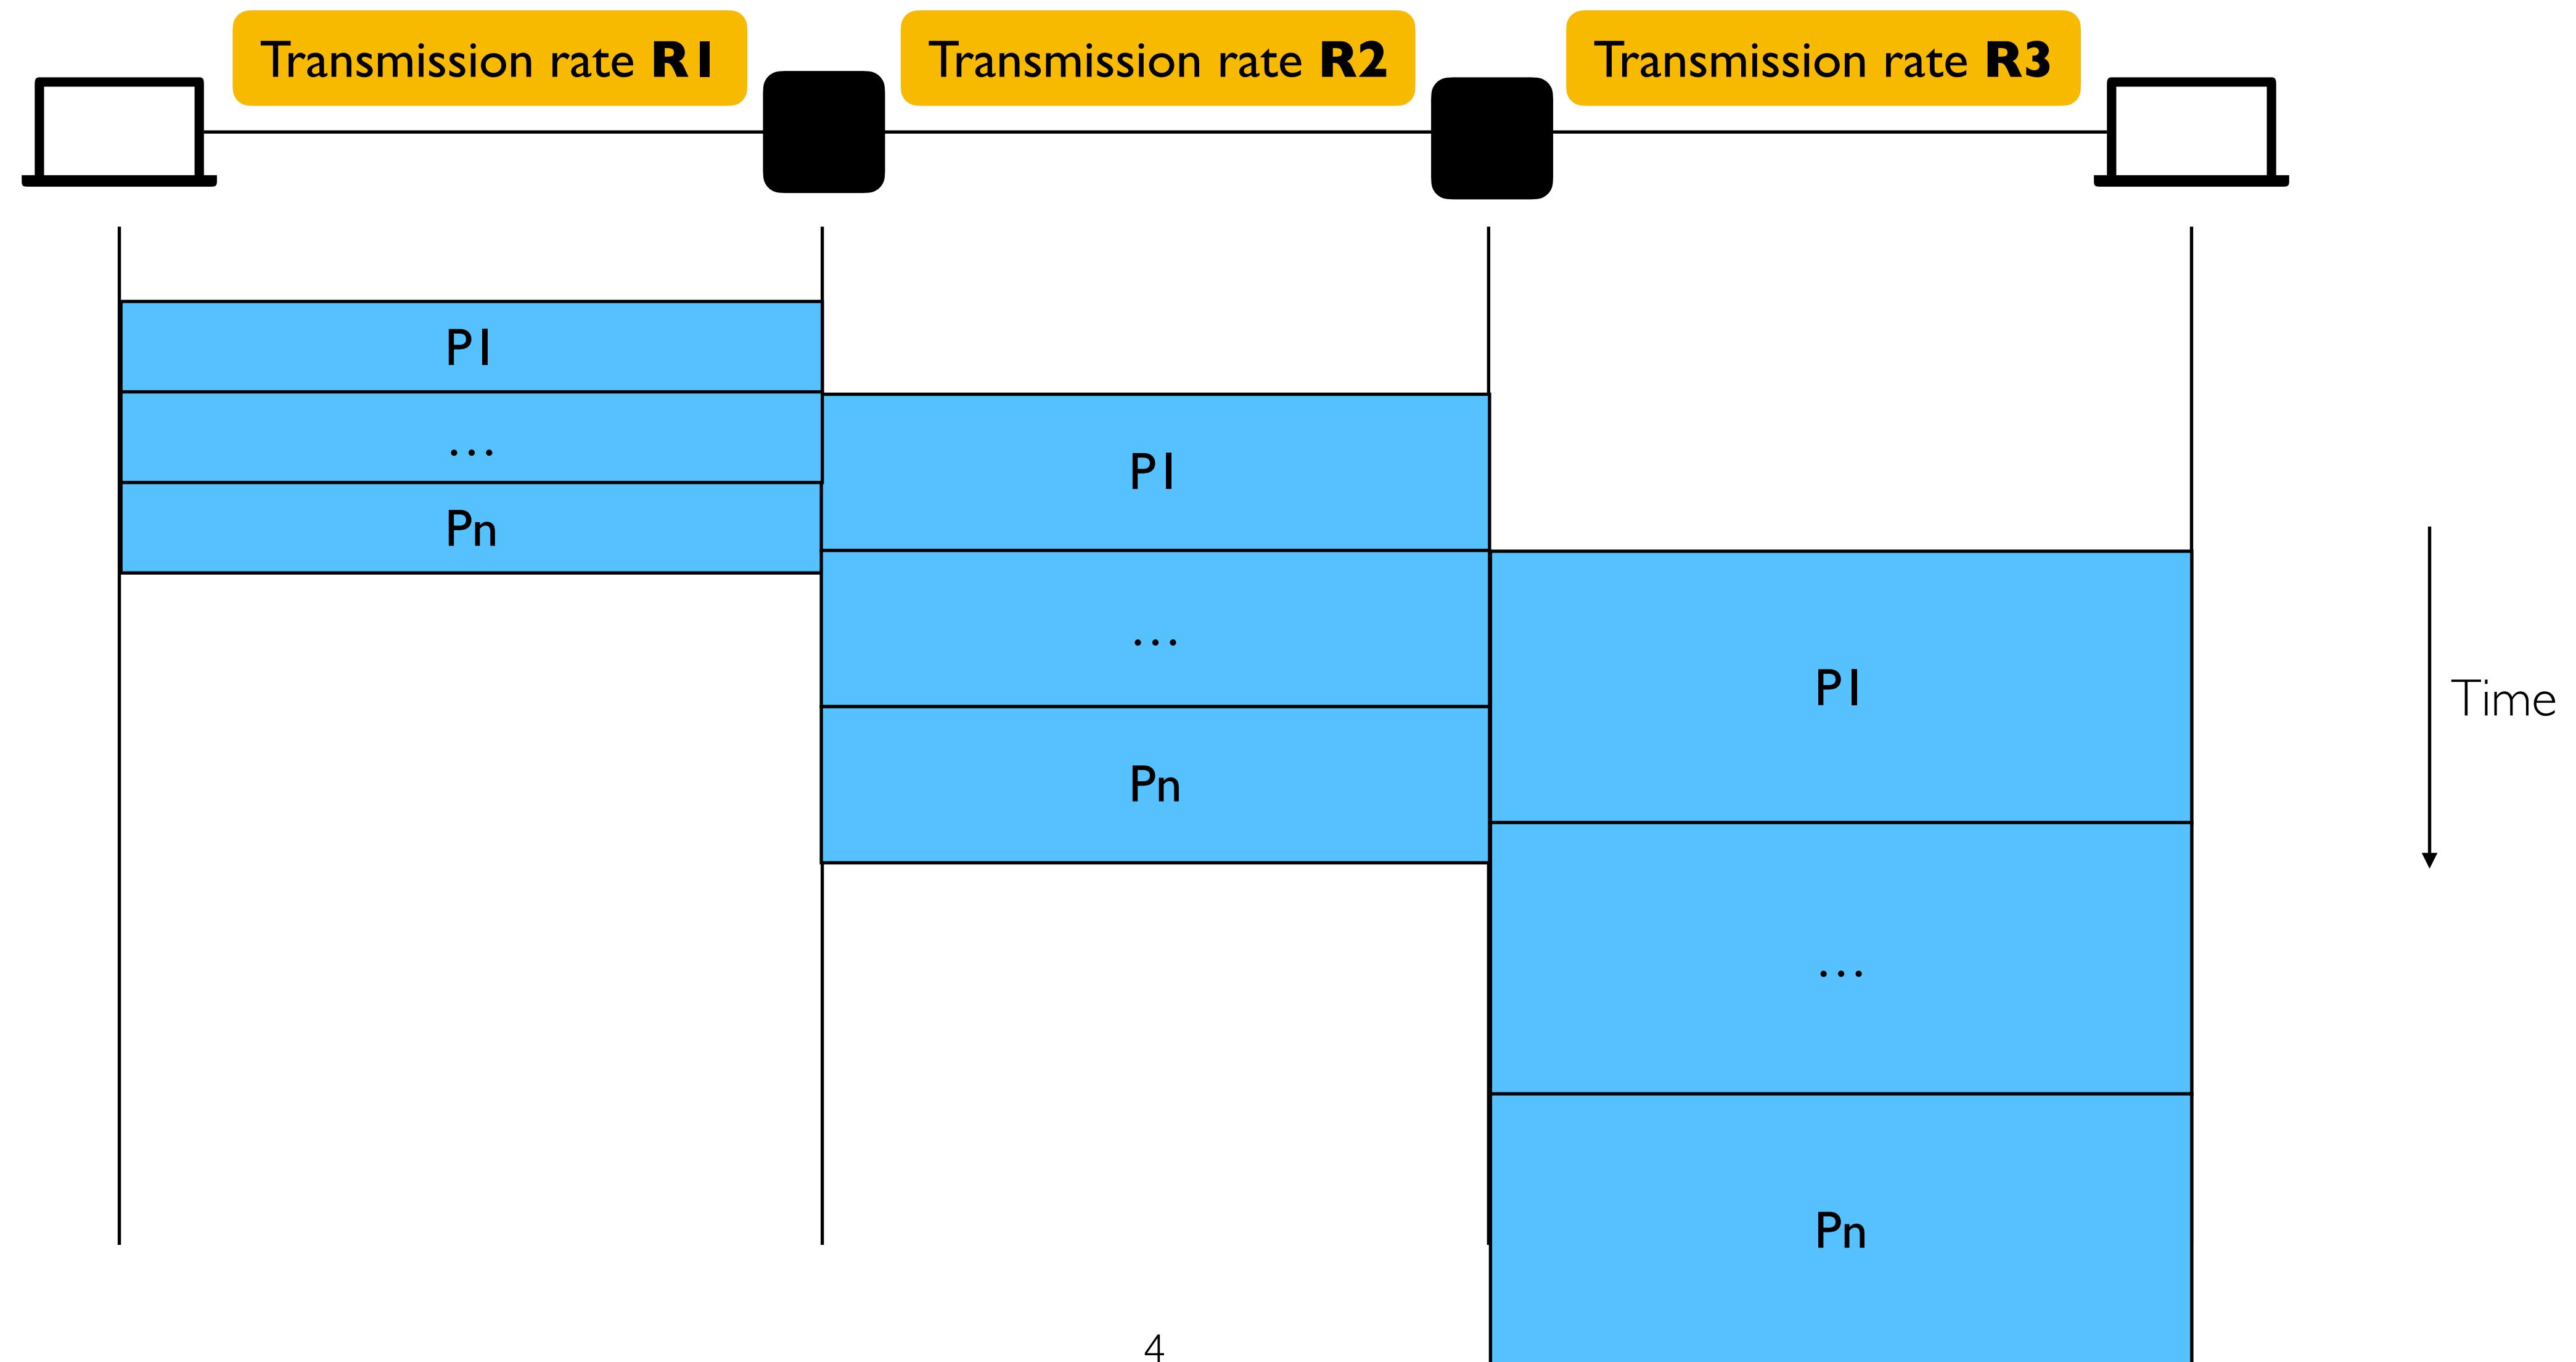
Ignoring propagation, queueing and processing delays, packet delay for a packet of  $P$  bits =  $\frac{P}{R1} + \frac{P}{R2} + \frac{P}{R3}$

Shouldn't the throughput of a file of  $F$  bits sent in  $P$  bit packets be  $\frac{F}{P} \cdot \frac{P}{\frac{P}{R1} + \frac{P}{R2} + \frac{P}{R3}} = \frac{F}{\frac{P}{R1} + \frac{P}{R2} + \frac{P}{R3}}$ ?

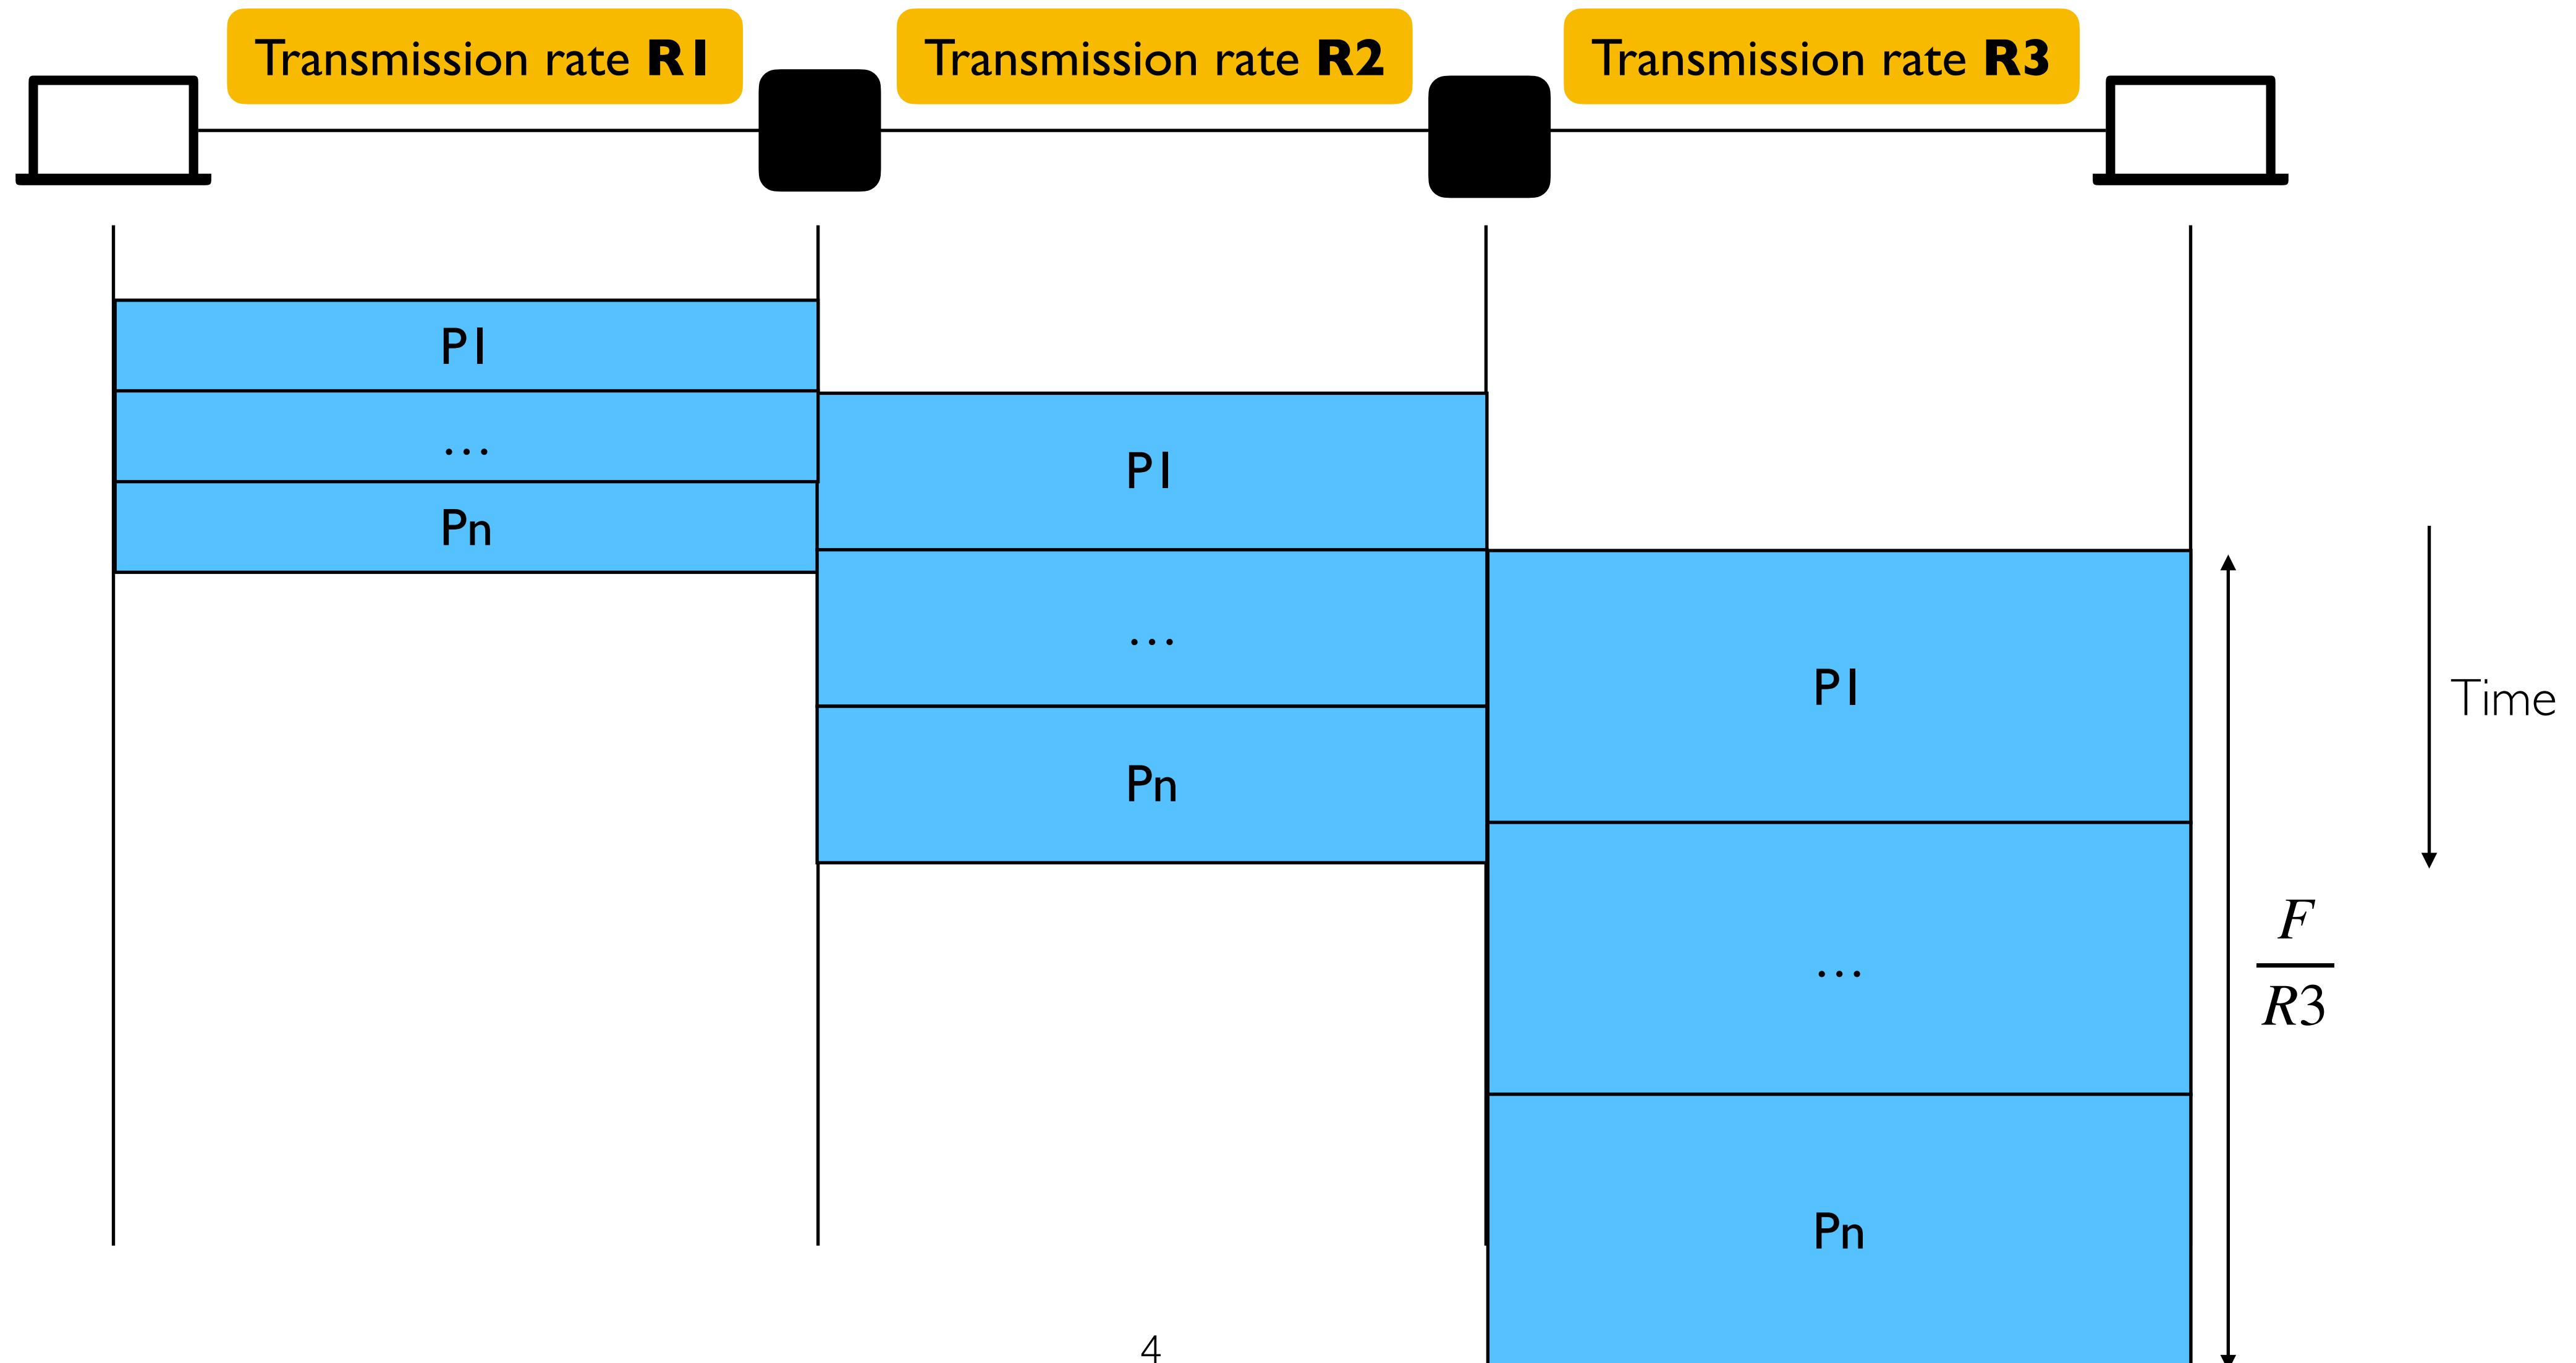
# Throughput Calculation: Packet by packet



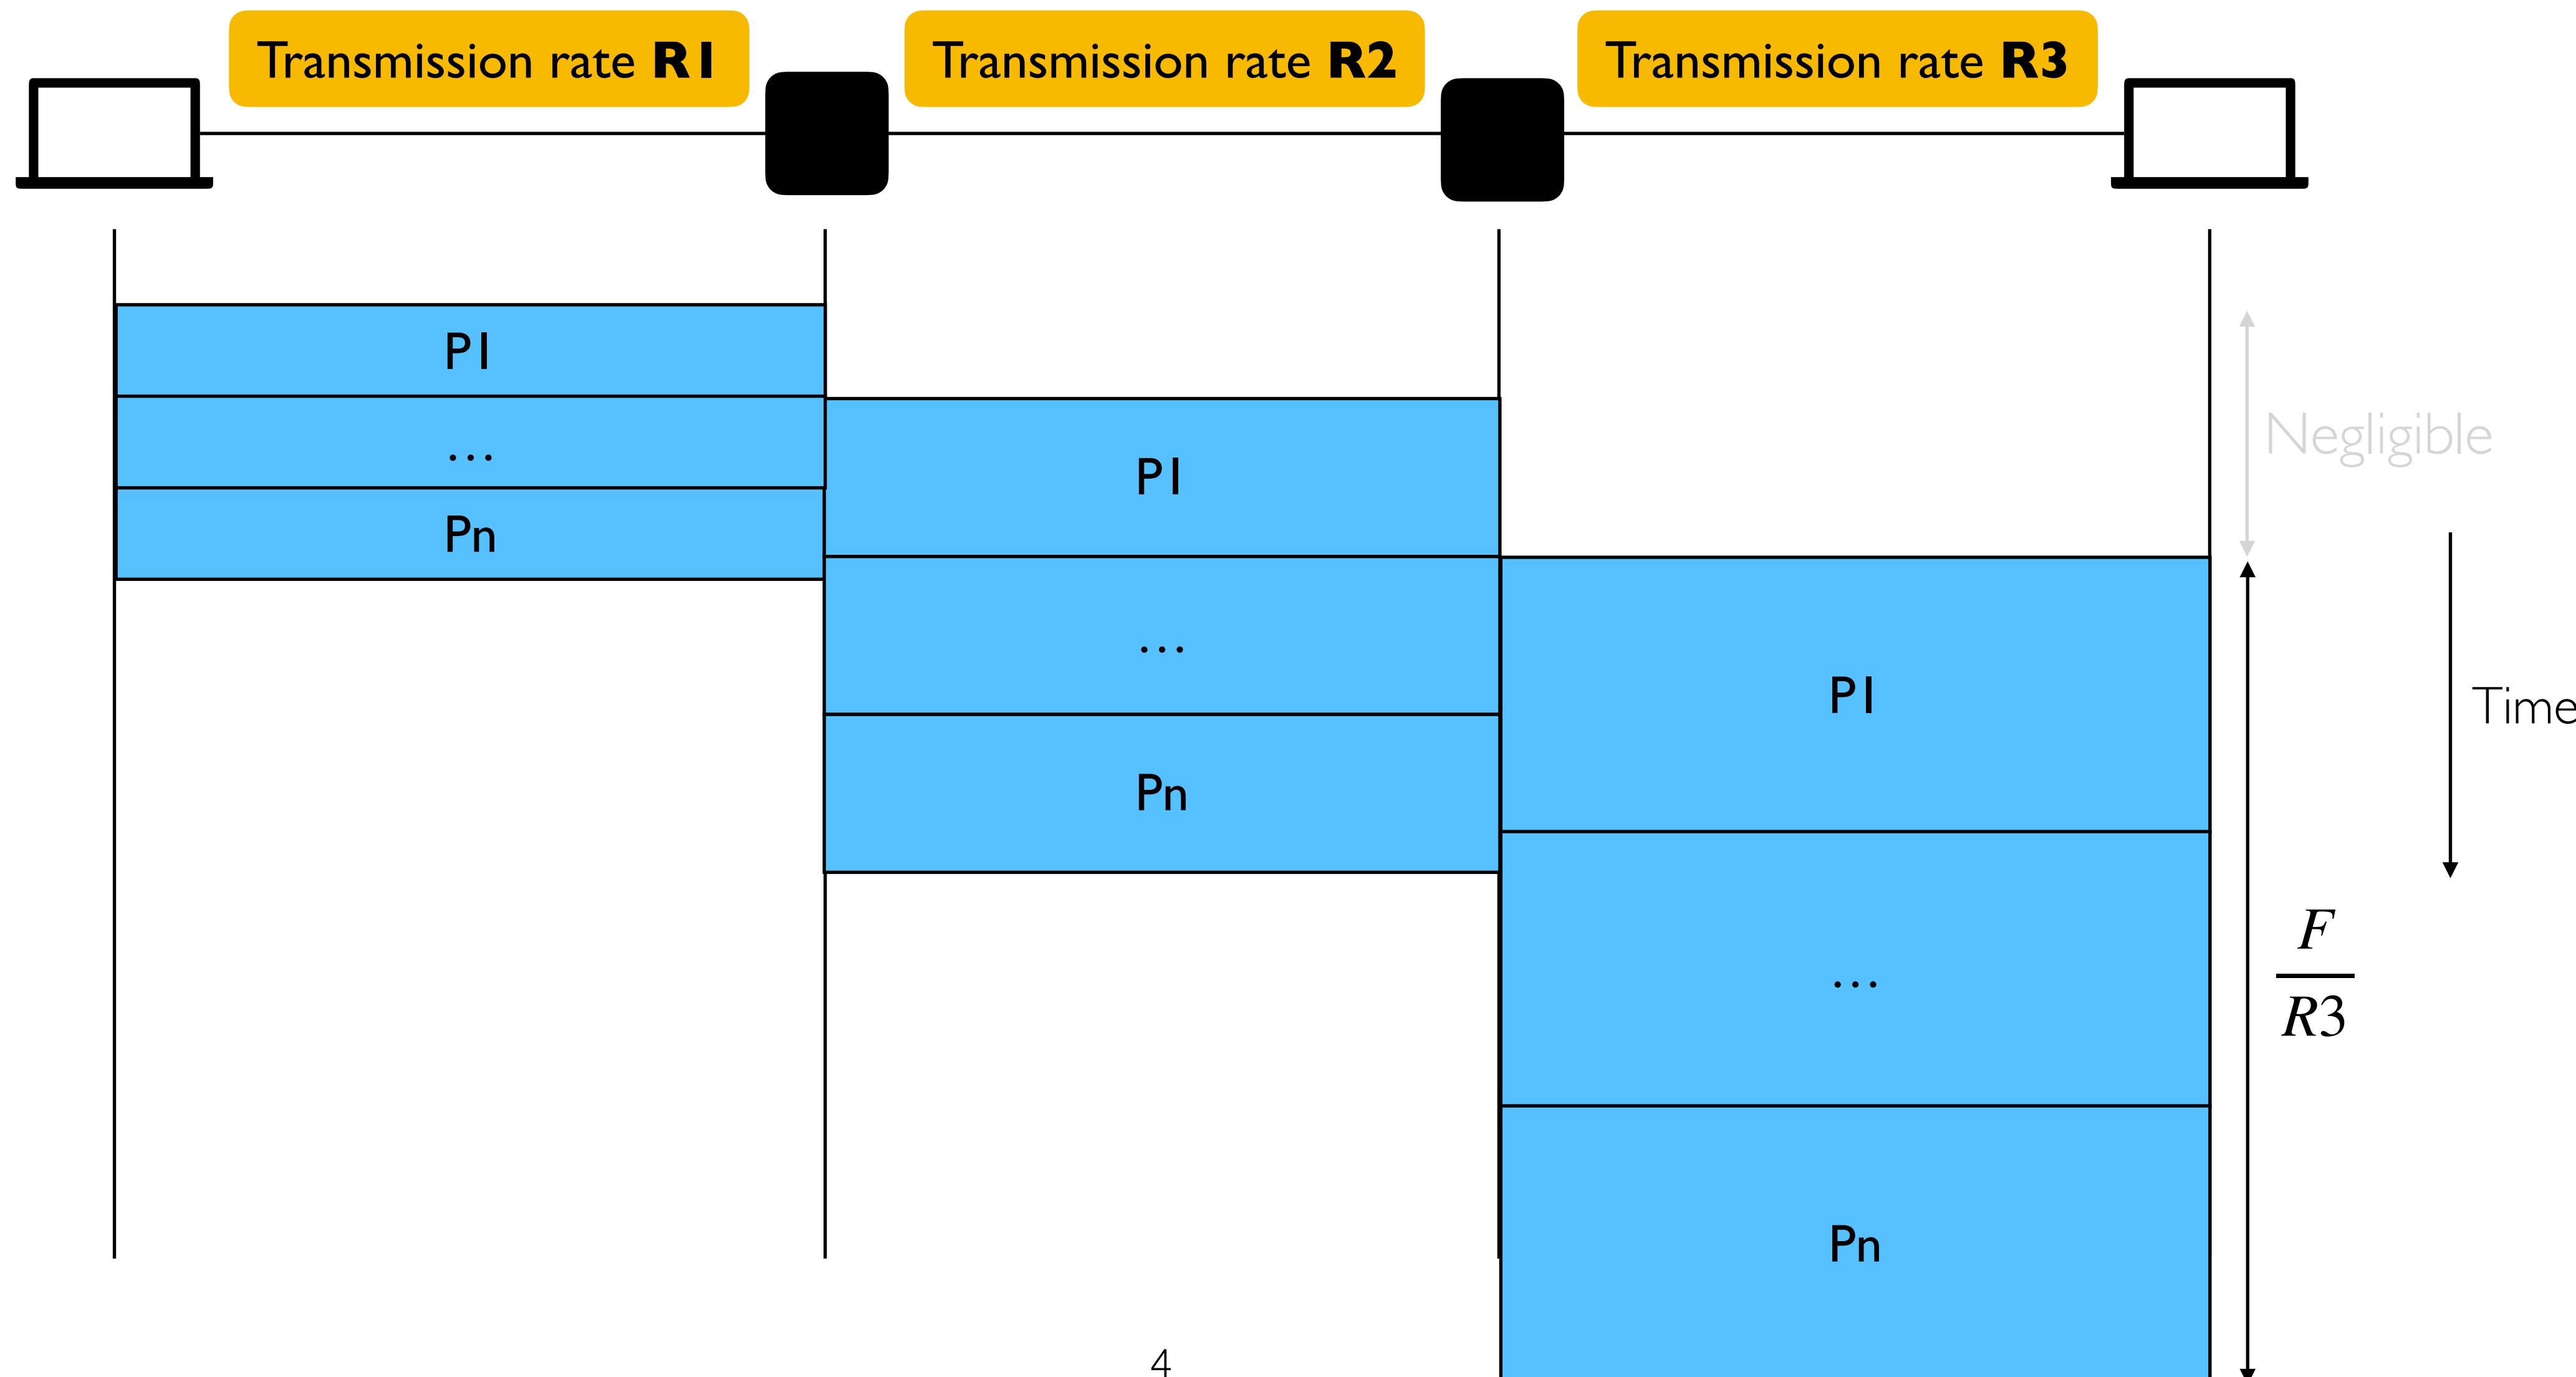
# Throughput Calculation: In Reality



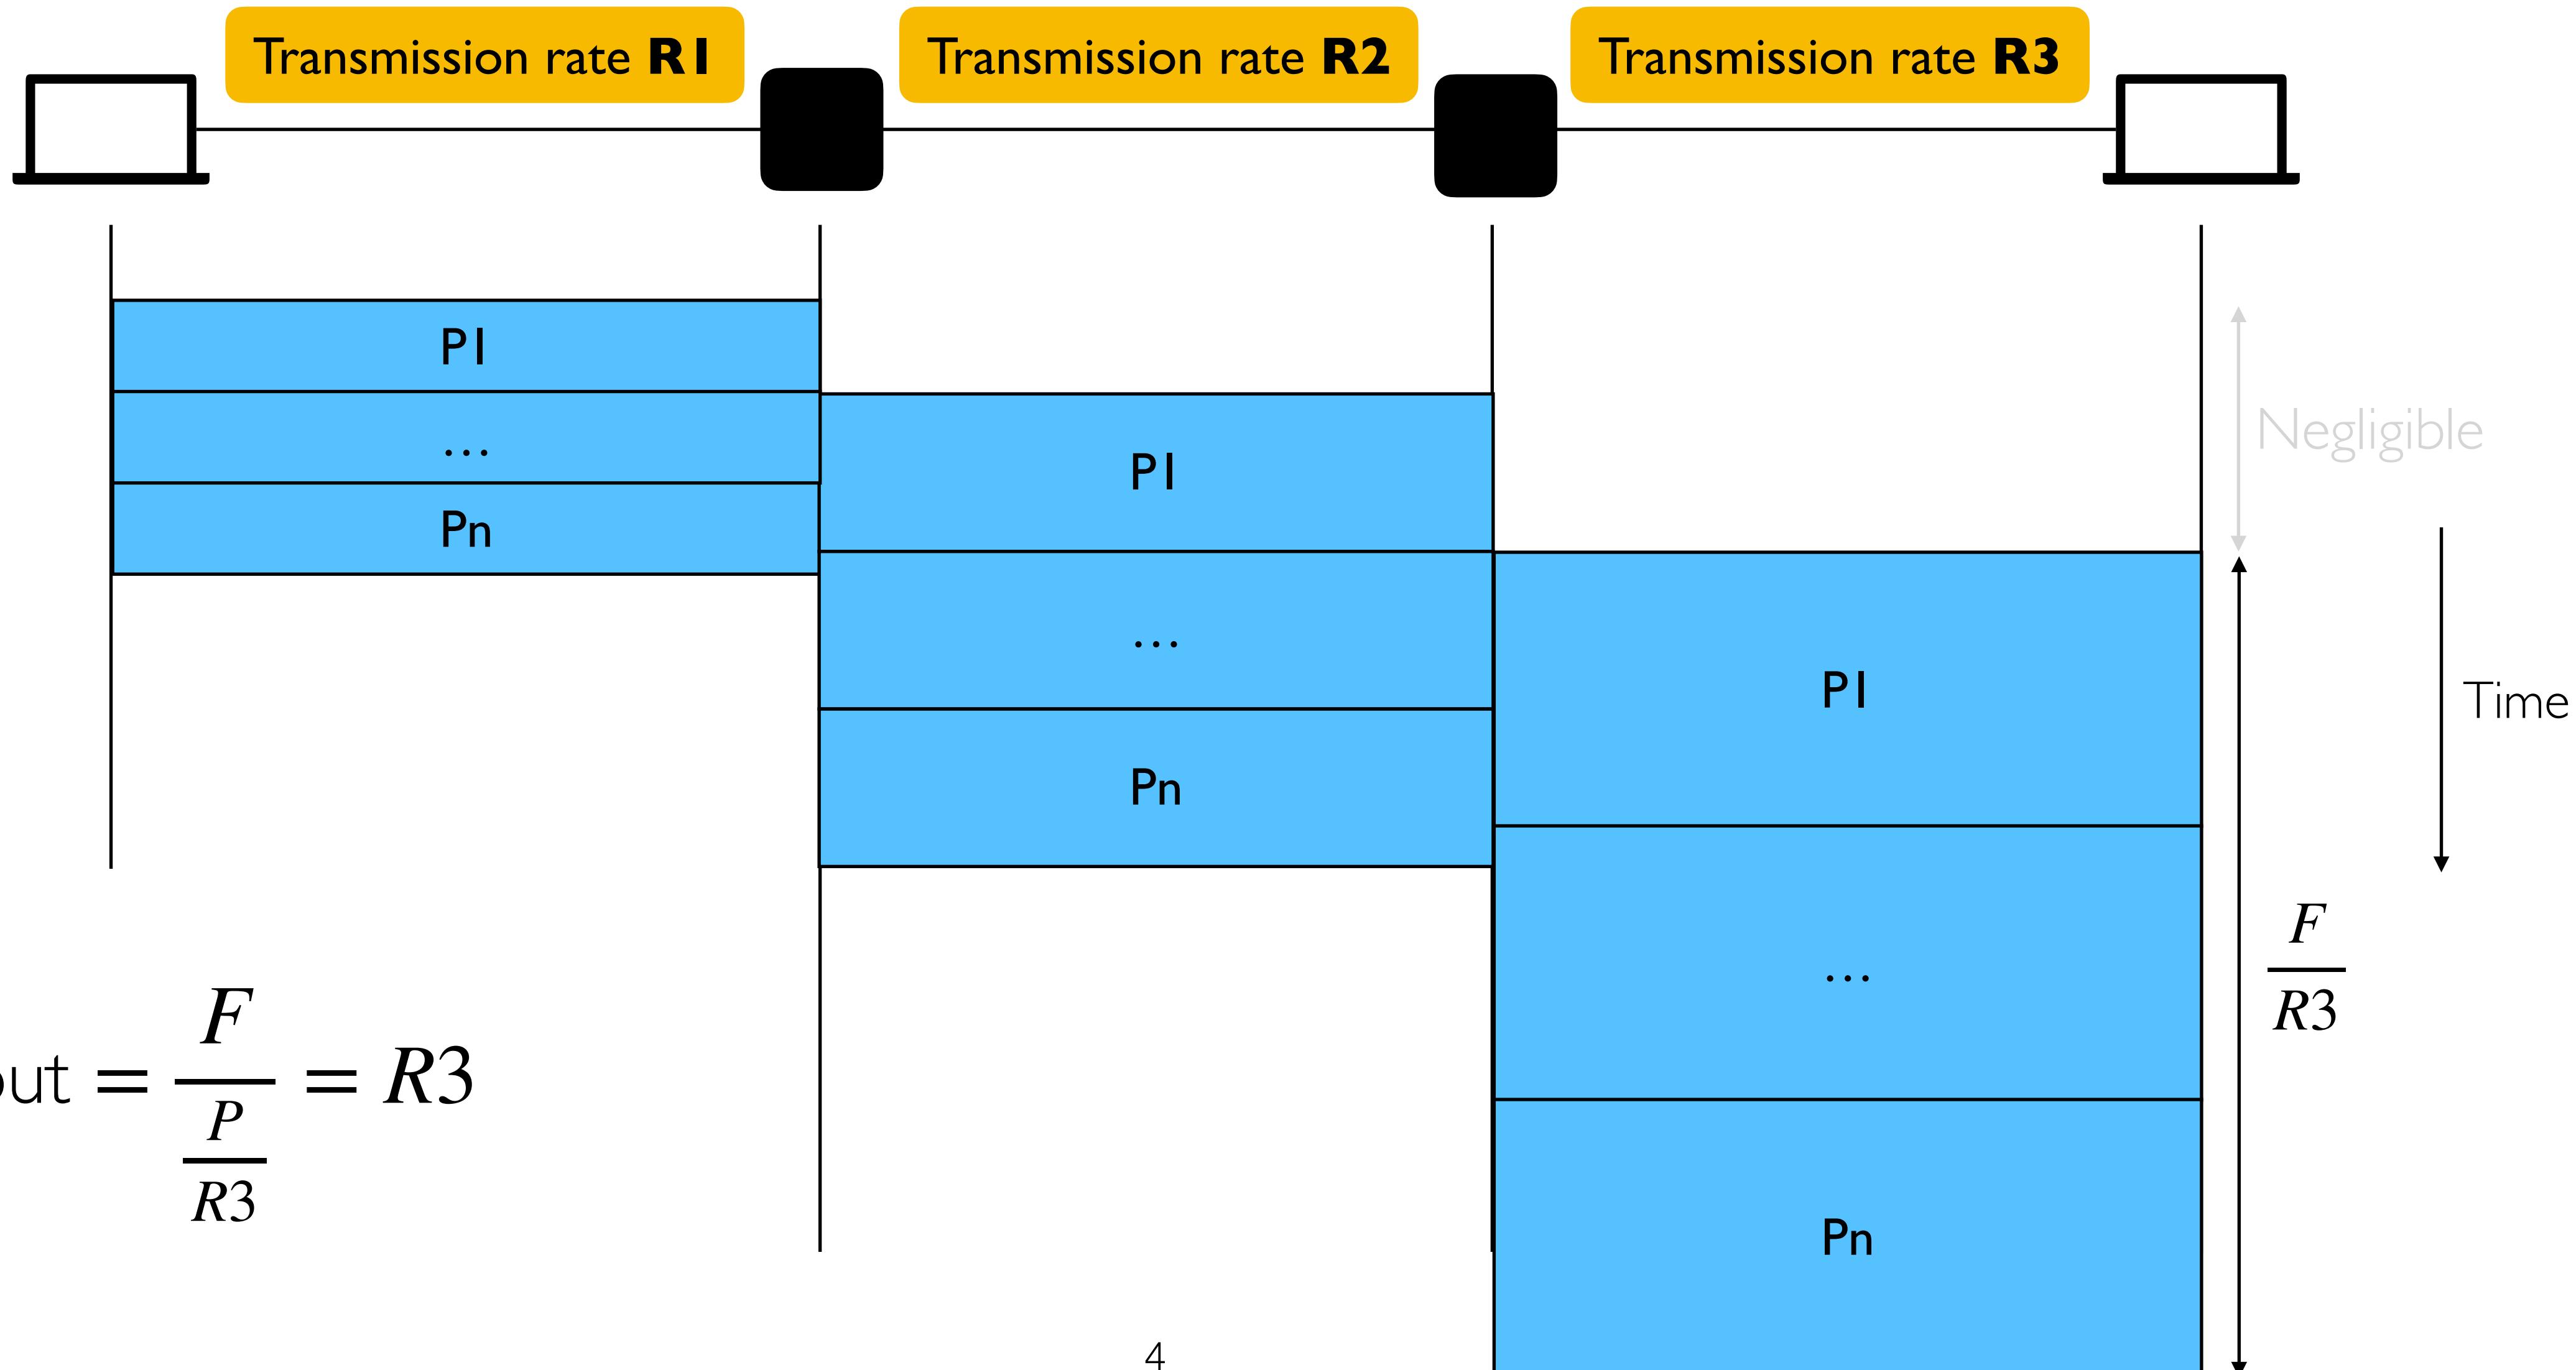
# Throughput Calculation: In Reality



# Throughput Calculation: In Reality



# Throughput Calculation: In Reality



# Three steps

- **Decompose** the problem into tasks
- **Organize** these tasks
- **Assign** tasks to entities (who does what)

# Three steps

- **Decomposition**
- **Organization**
- **Assignment**

# Three steps

- **Decomposition**
- **Organization**

Layering principle

- **Assignment**

# Three steps

- **Decomposition**
- **Organization**
- **Assignment**

# The Path through FedEx

FE = FedEx Envelope

Truck

Sorting Office

Airport

Sorting Office

Airport

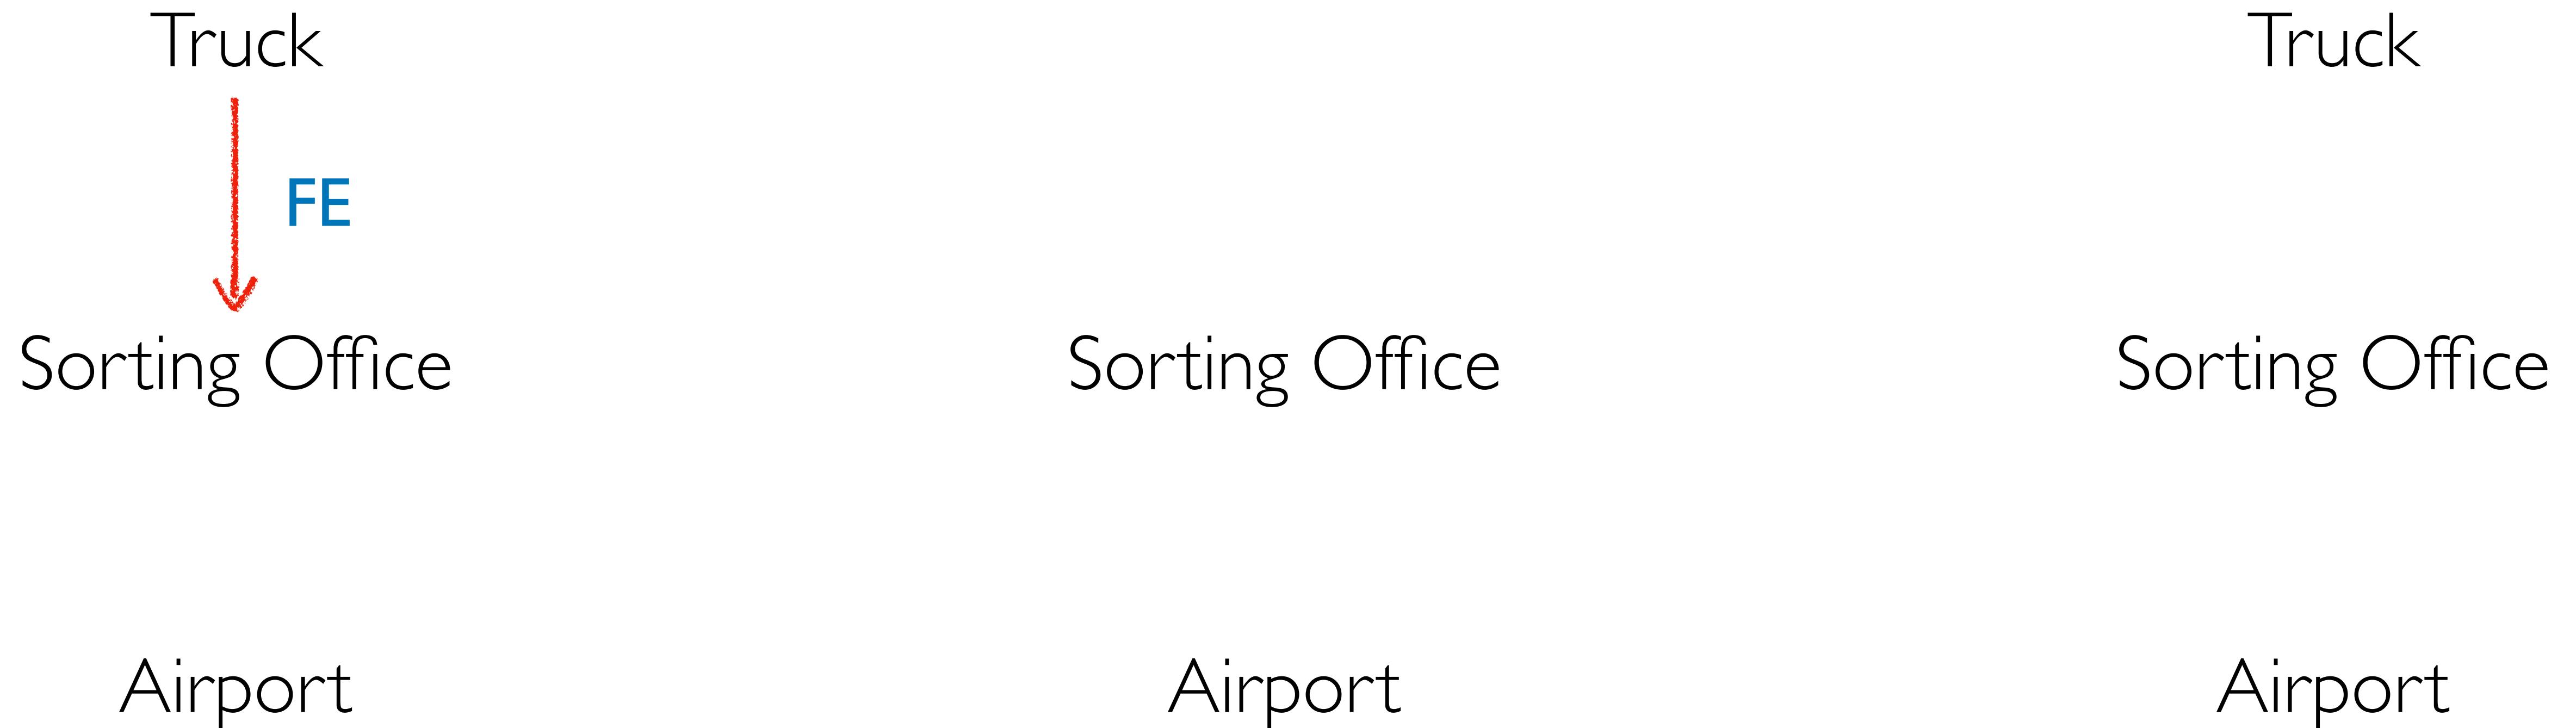
Truck

Sorting Office

Airport

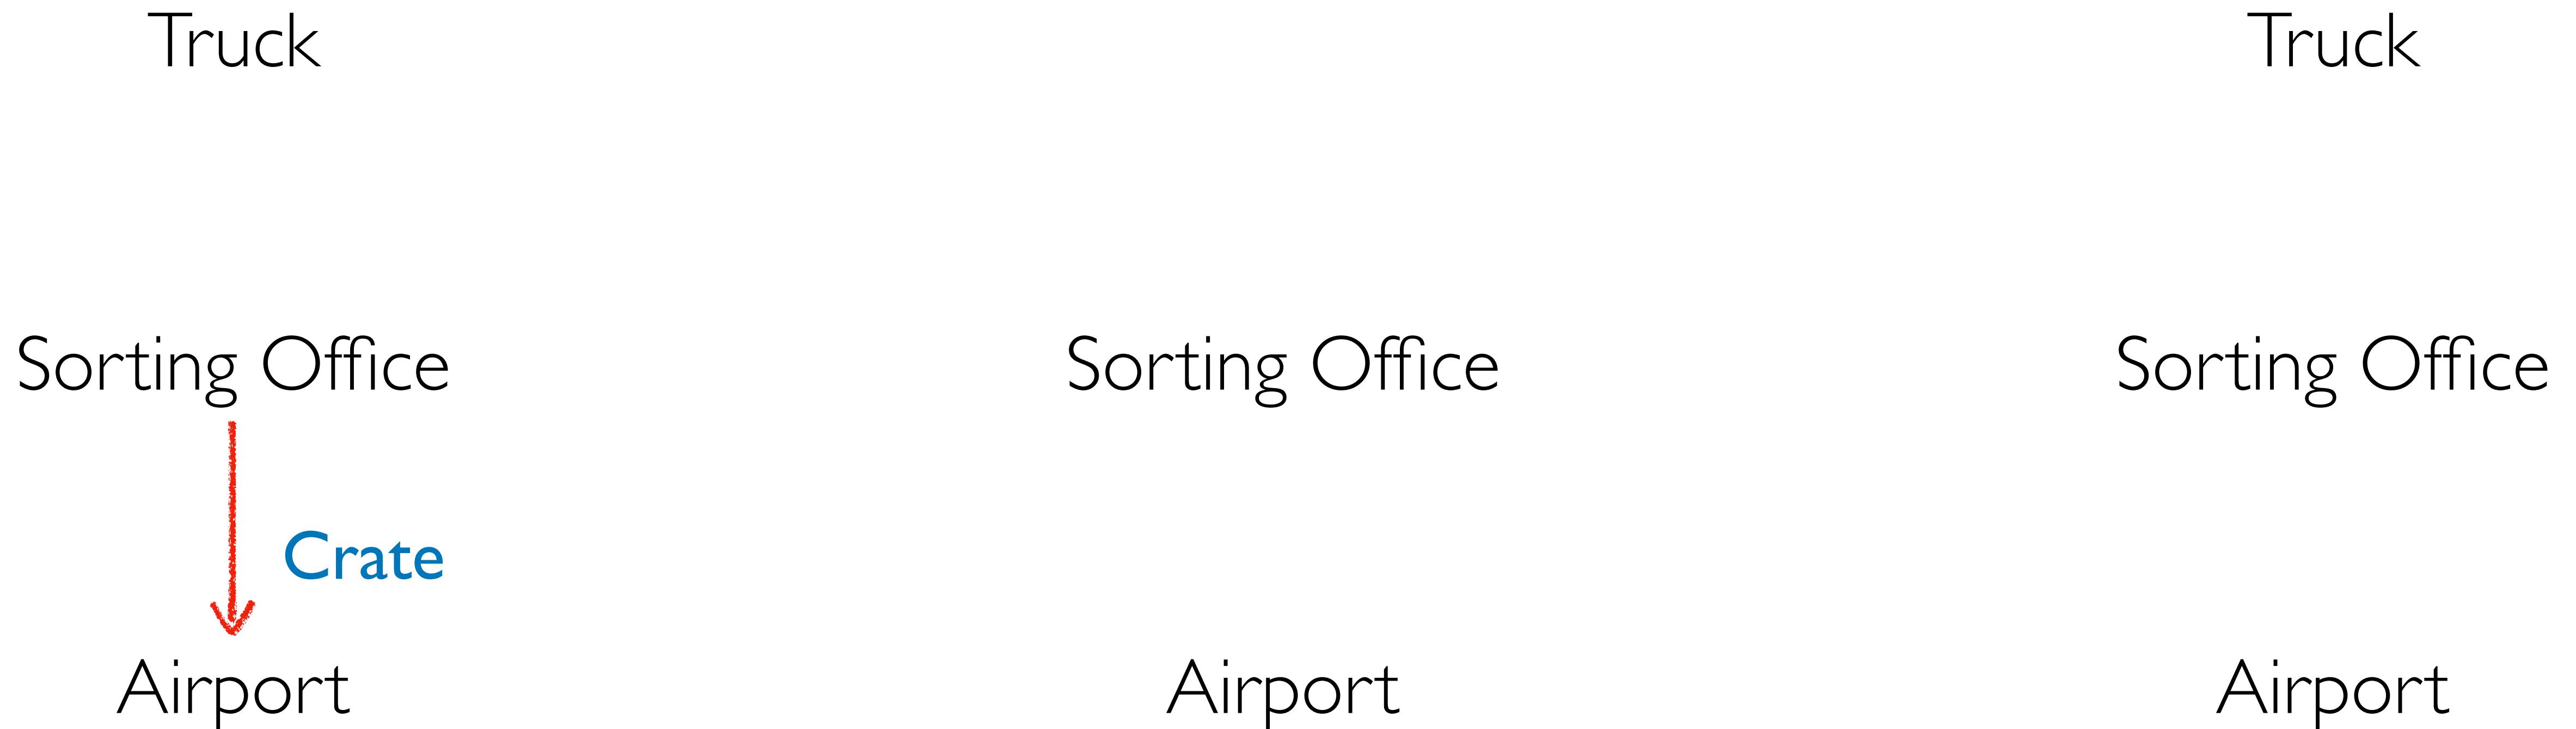
# The Path through FedEx

FE = FedEx Envelope



# The Path through FedEx

FE = FedEx Envelope



# The Path through FedEx

FE = FedEx Envelope

Truck

Sorting Office

Airport

Sorting Office

Airport

Truck

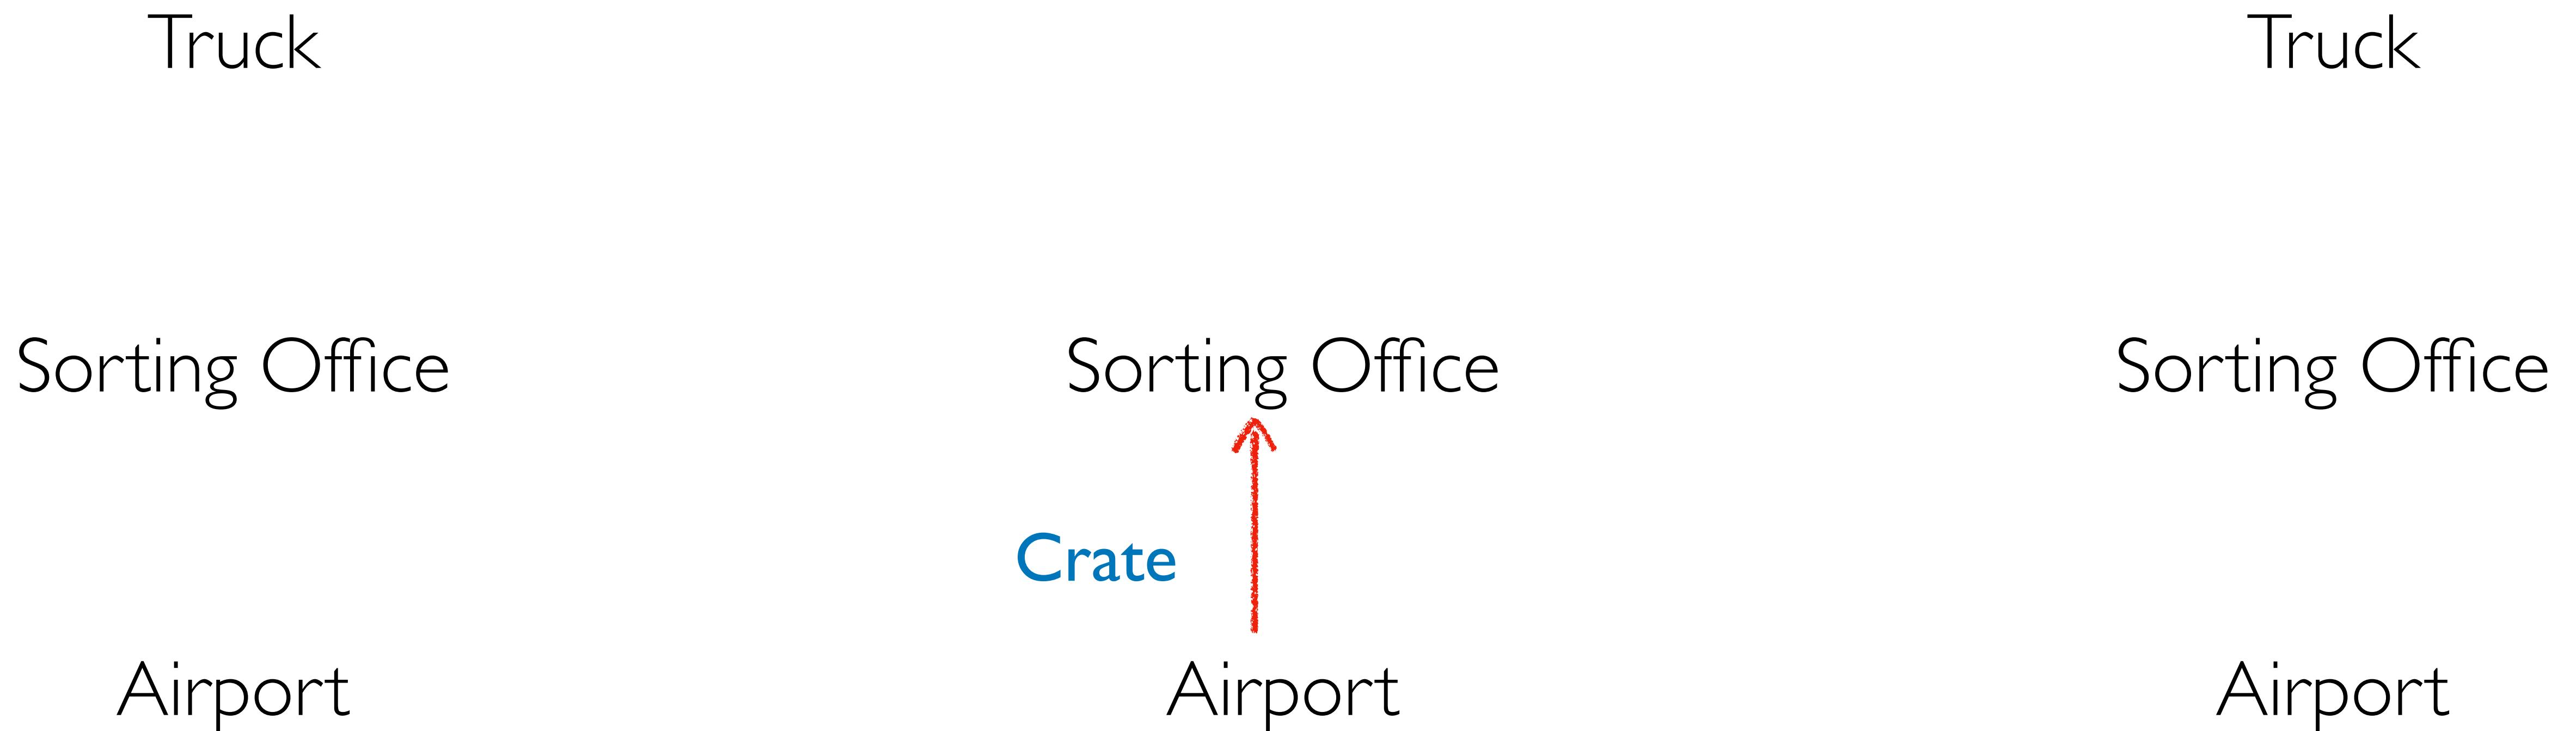
Sorting Office

Airport



# The Path through FedEx

FE = FedEx Envelope



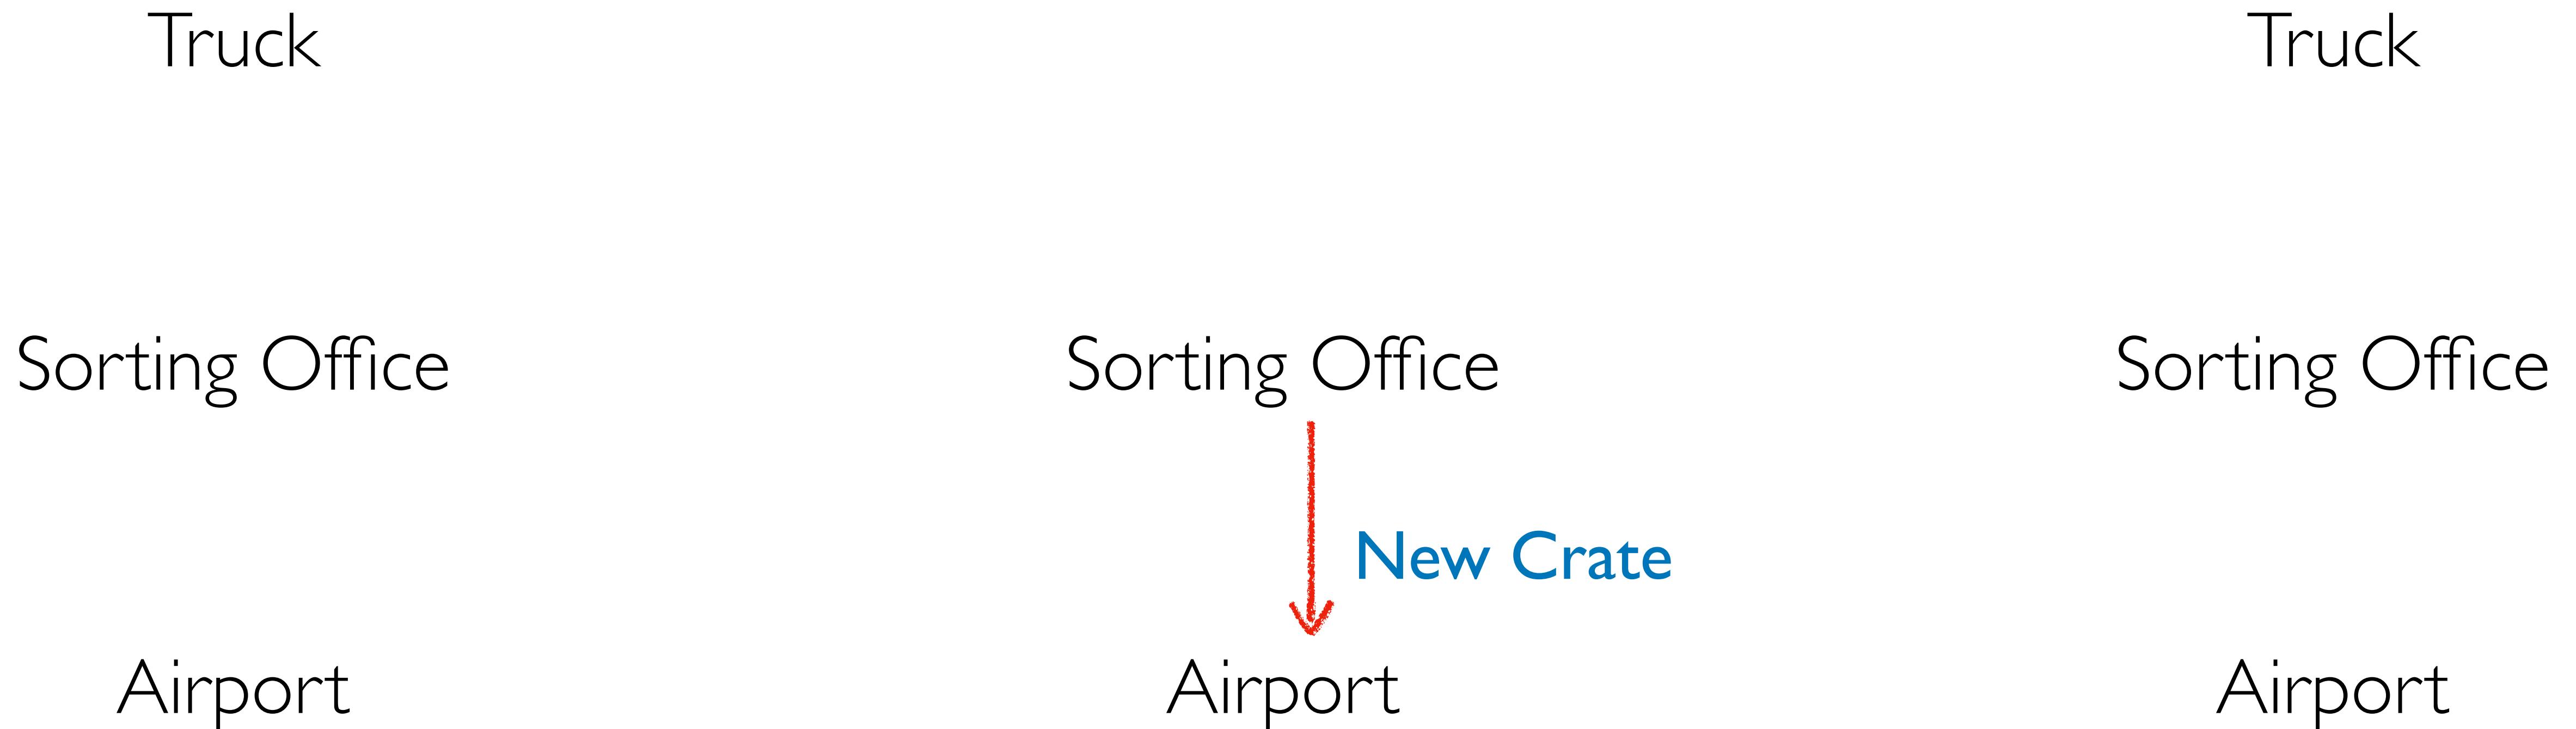
# The Path through FedEx

FE = FedEx Envelope



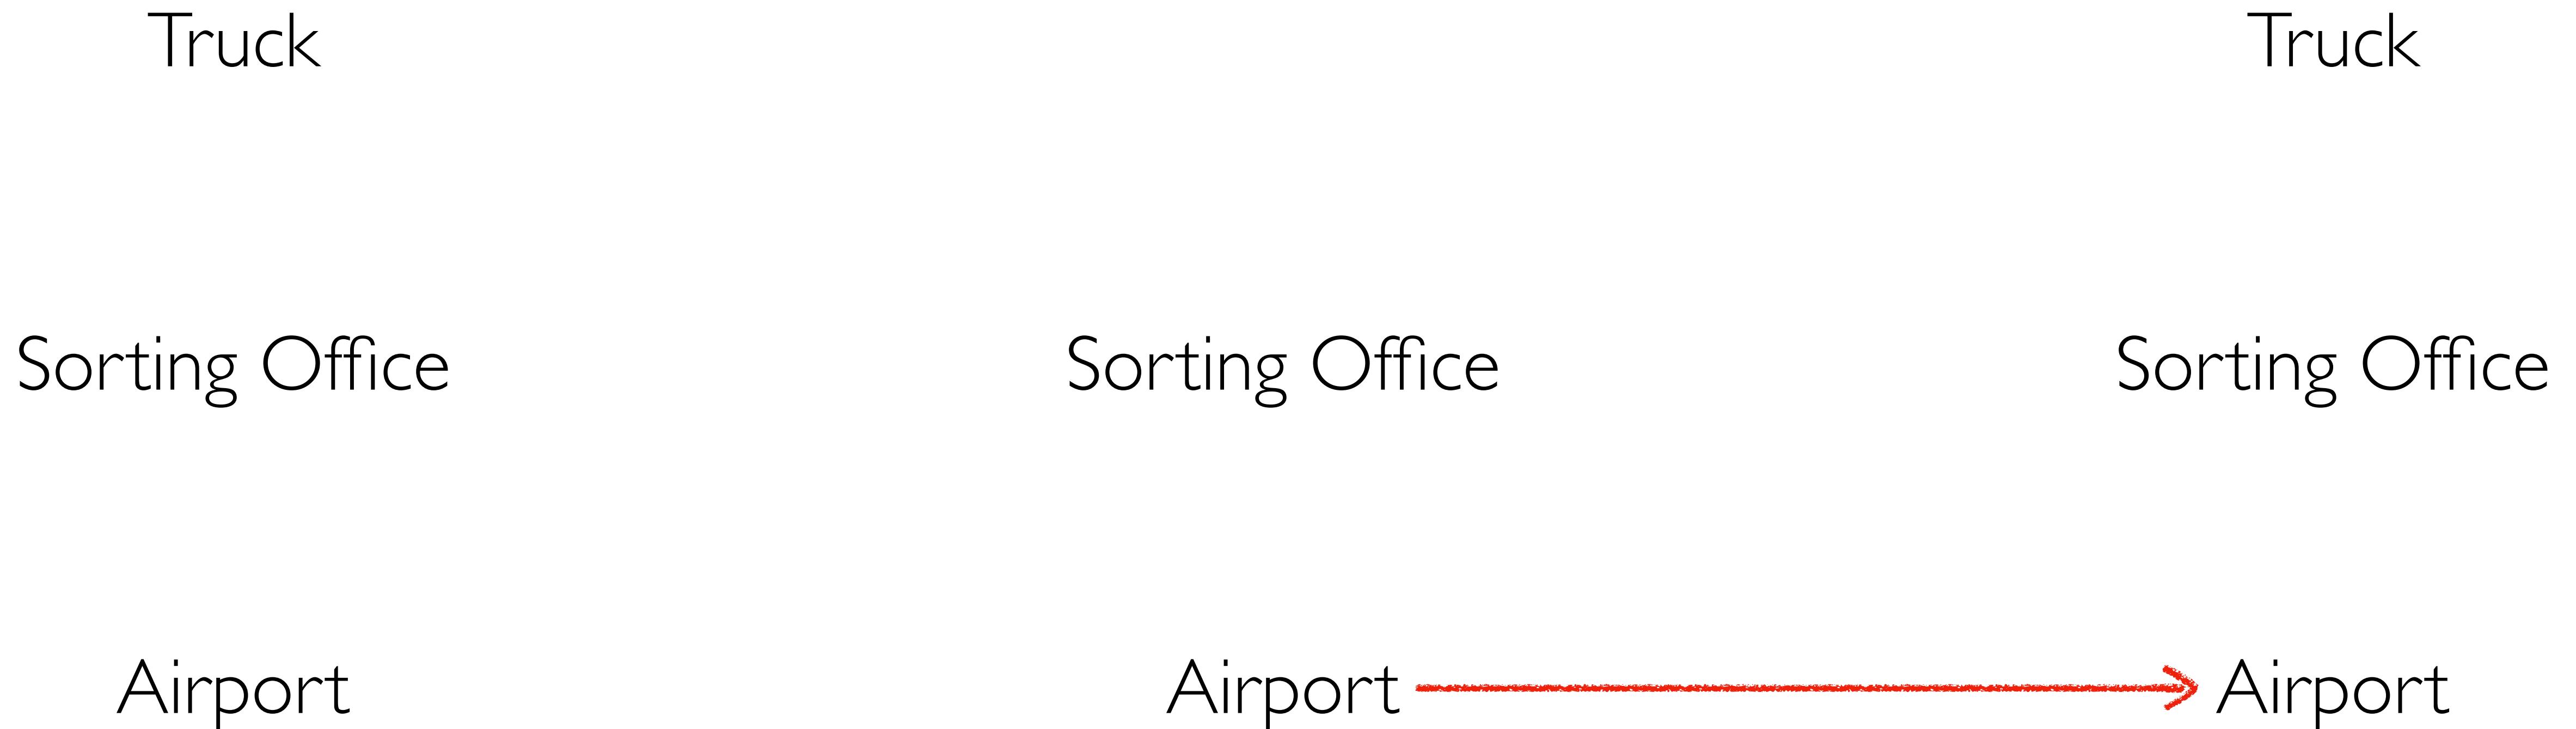
# The Path through FedEx

FE = FedEx Envelope



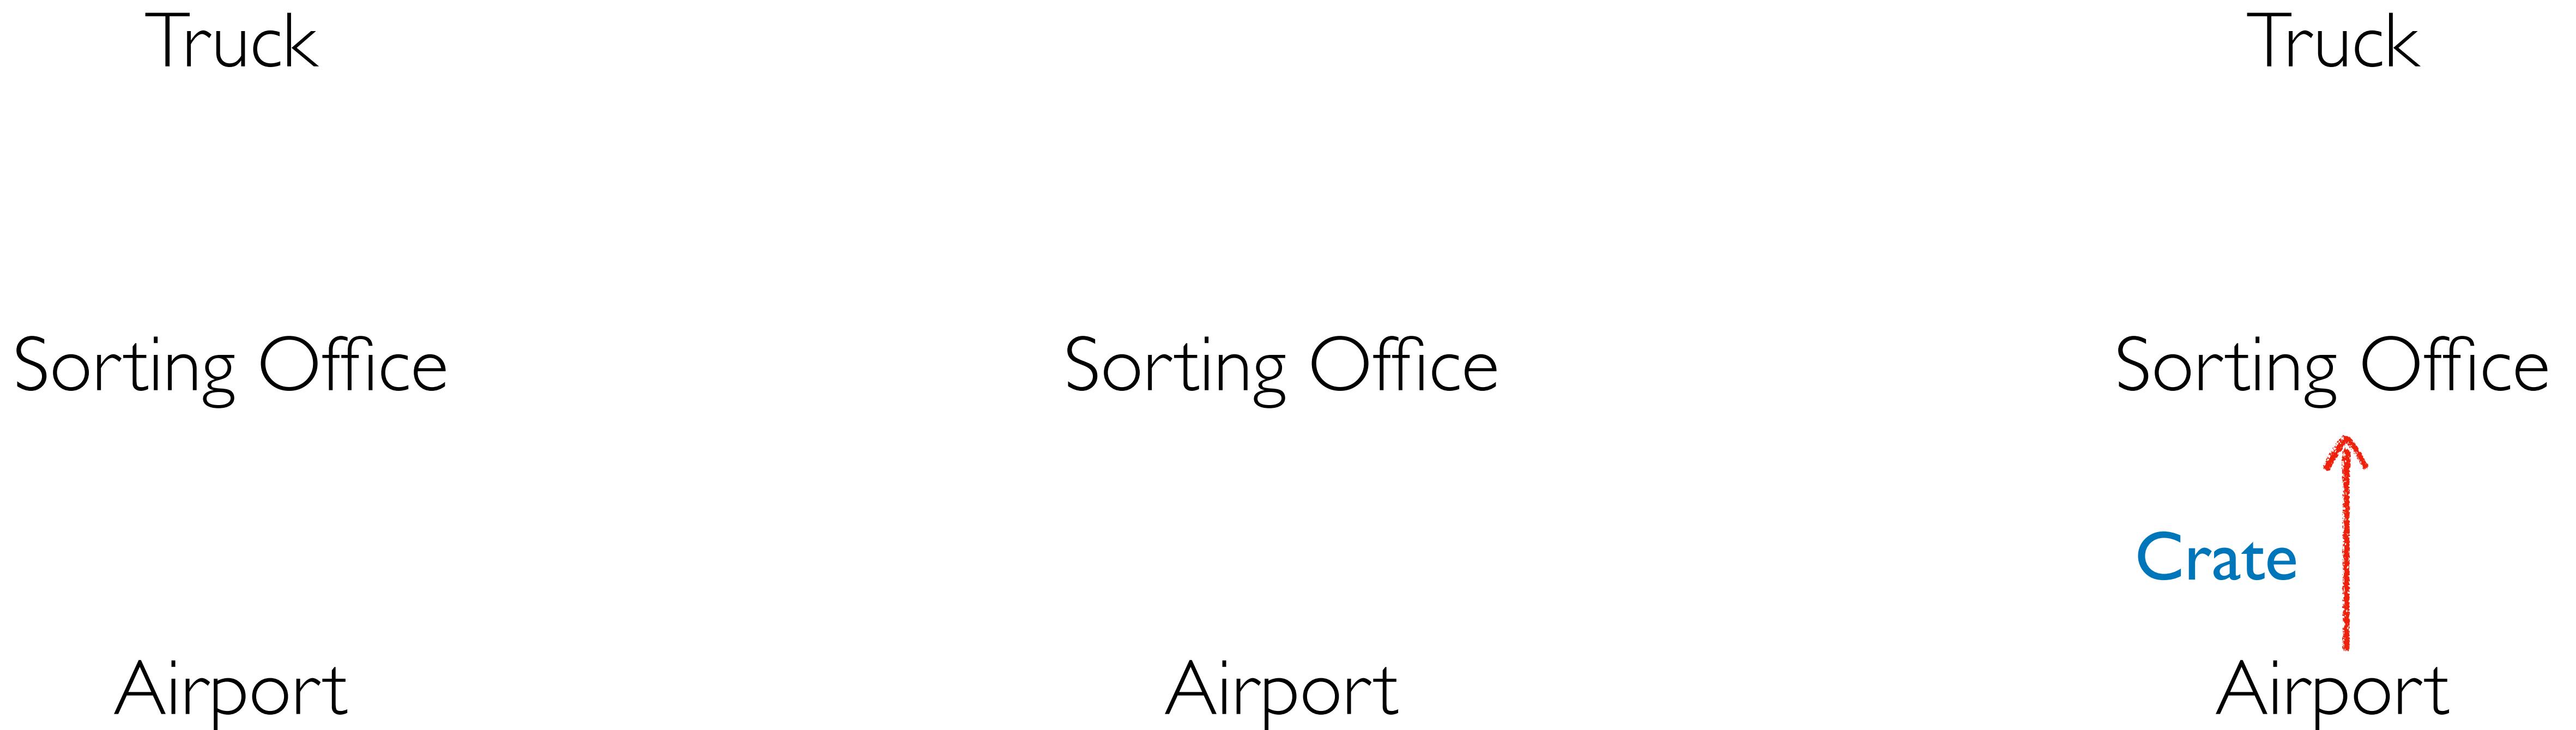
# The Path through FedEx

FE = FedEx Envelope



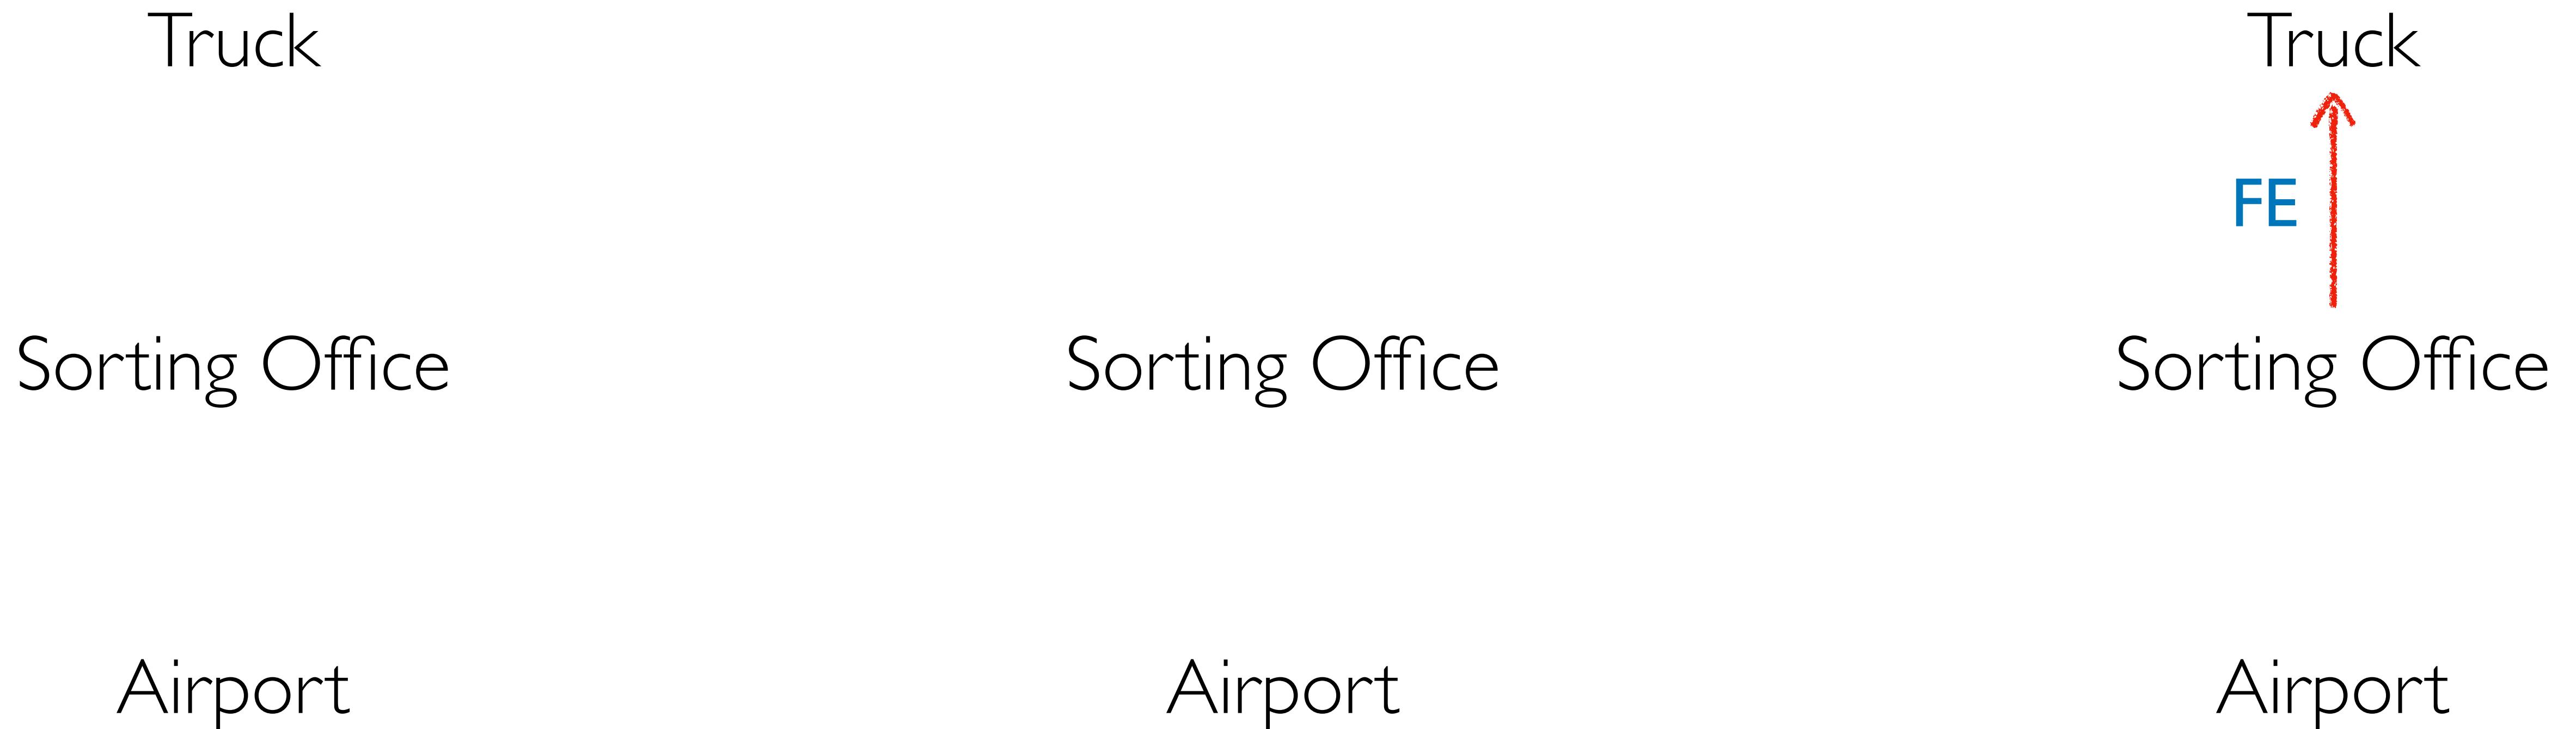
# The Path through FedEx

FE = FedEx Envelope



# The Path through FedEx

FE = FedEx Envelope



# The Path through FedEx

FE = FedEx Envelope

Truck

Sorting Office

Airport

Sorting Office

Airport

Truck

Sorting Office

Airport

# The Path through FedEx

FE = FedEx Envelope

Truck

Truck

Sorting Office

Sorting Office

Sorting Office

Airport

Airport

Airport

Deepest packaging (Envelope+FE+Crate) at Lowest level of transport

# The Path through FedEx

Higher “Stack”  
at Ends

Truck

Sorting Office

Airport

Deepest packaging (Envelope+FE+Crate) at Lowest level of transport

FE = FedEx Envelope

Sorting Office

Airport

Sorting Office

Airport

# The Path through FedEx

Higher “Stack”  
at Ends

Truck

Sorting Office

Airport

Deepest packaging (Envelope+FE+Crate) at Lowest level of transport

Partial “Stack”  
during Transit

Sorting Office

Airport

Truck

Sorting Office

Airport

FE = FedEx Envelope

# Path through Internet



# Path through Internet



What gets implemented where?

# What layers get implemented at the end-systems?

# What layers get implemented at the end-systems?

- Bits arrive on wire, must make it up to application

# What layers get implemented at the end-systems?

- Bits arrive on wire, must make it up to application
- Therefore, all layers must exist at host!

# What layers get implemented in the network?

# What layers get implemented in the network?

- Bits arrive on wire → physical layer (L1)

# What layers get implemented in the network?

- Bits arrive on wire → physical layer (L1)
- Packets must be delivered across link and local networks → datalink layer (L2)

# What layers get implemented in the network?

- Bits arrive on wire → physical layer (L1)
- Packets must be delivered across link and local networks → datalink layer (L2)
- Packets must be delivered b/w networks for global delivery → network layer (L3)

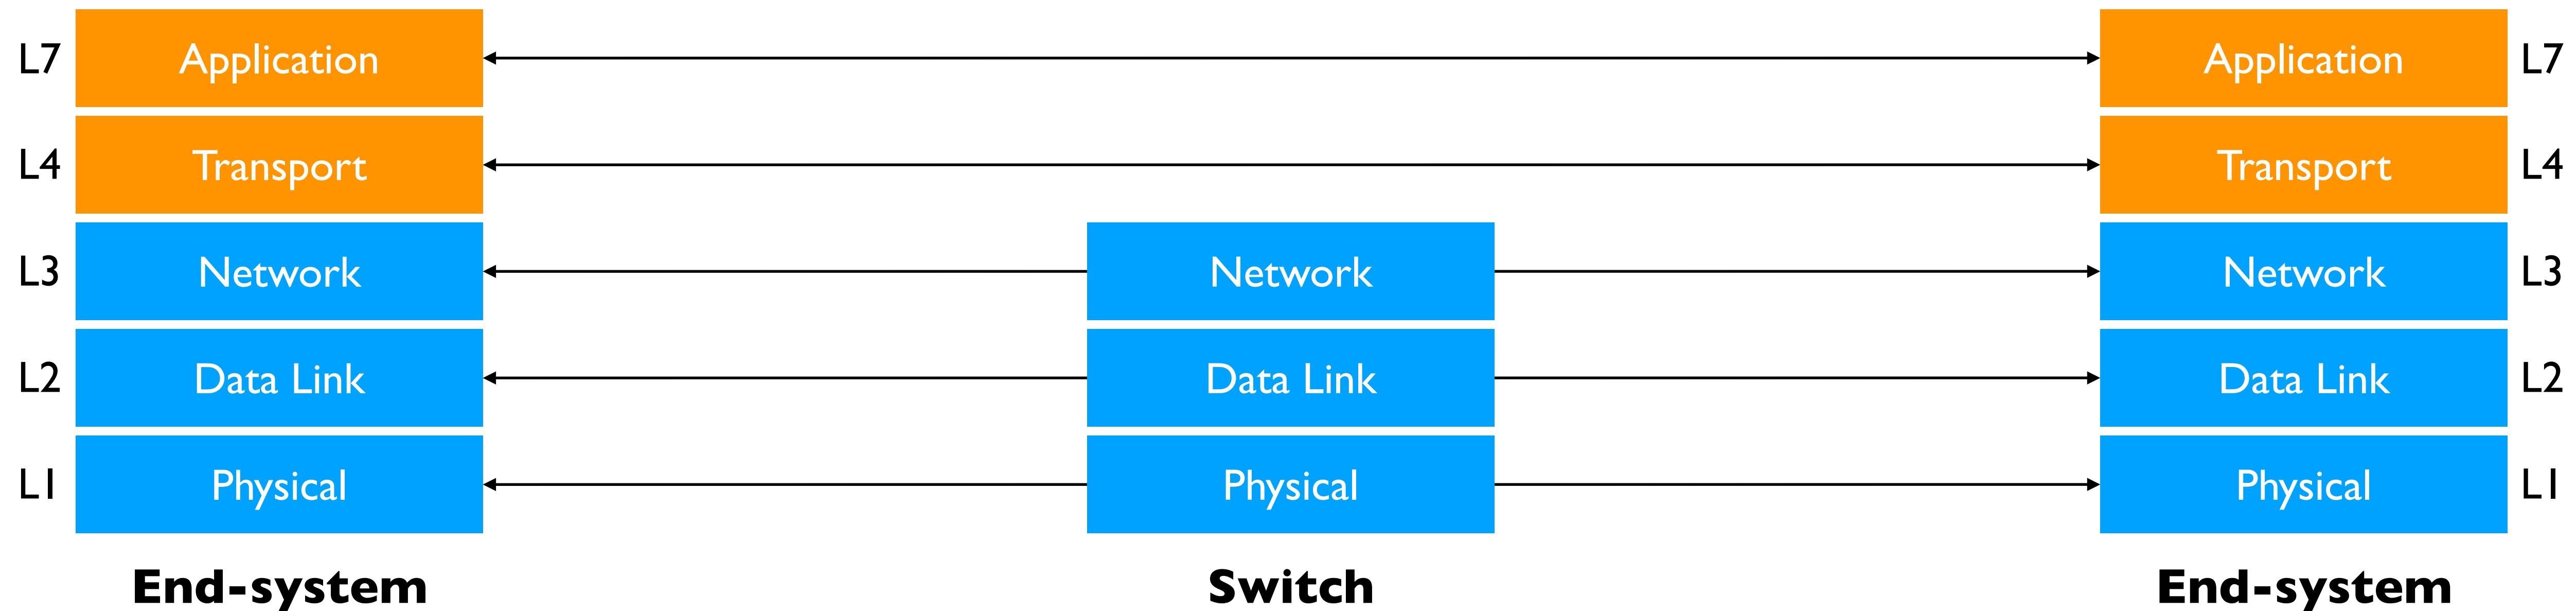
# What layers get implemented in the network?

- Bits arrive on wire → physical layer (L1)
- Packets must be delivered across link and local networks → datalink layer (L2)
- Packets must be delivered b/w networks for global delivery → network layer (L3)
- The network does not support reliable delivery

# What layers get implemented in the network?

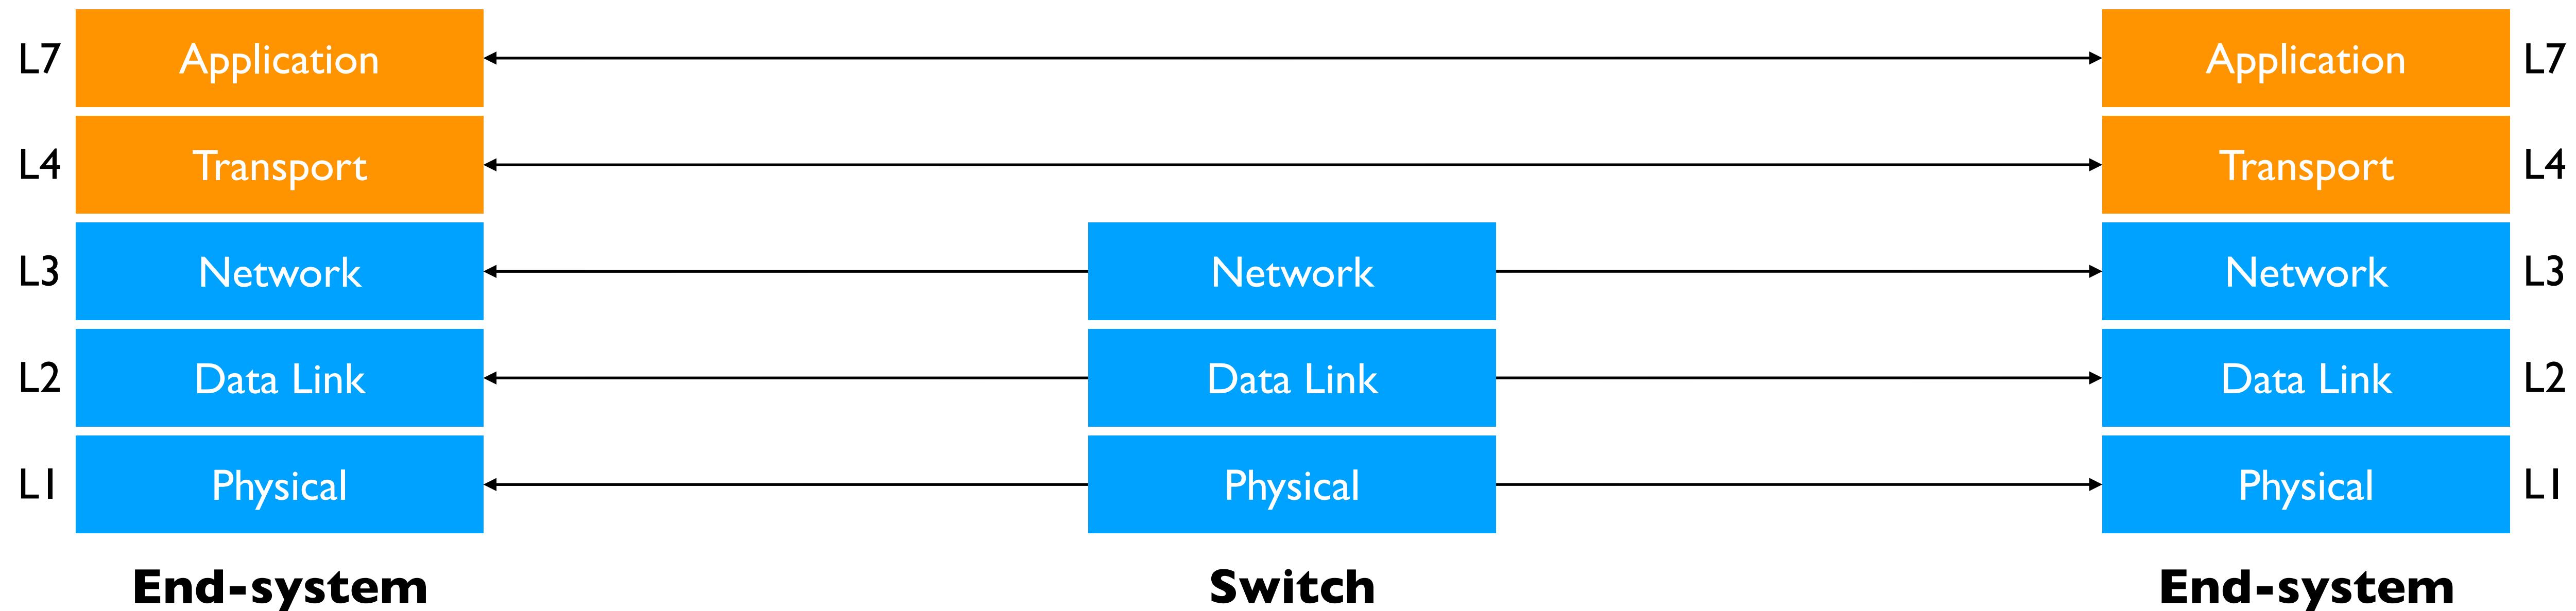
- Bits arrive on wire → physical layer (L1)
- Packets must be delivered across link and local networks → datalink layer (L2)
- Packets must be delivered b/w networks for global delivery → network layer (L3)
- The network does not support reliable delivery
  - *Transport layer (and above) not supported (**Why?**)*

# Path through Internet



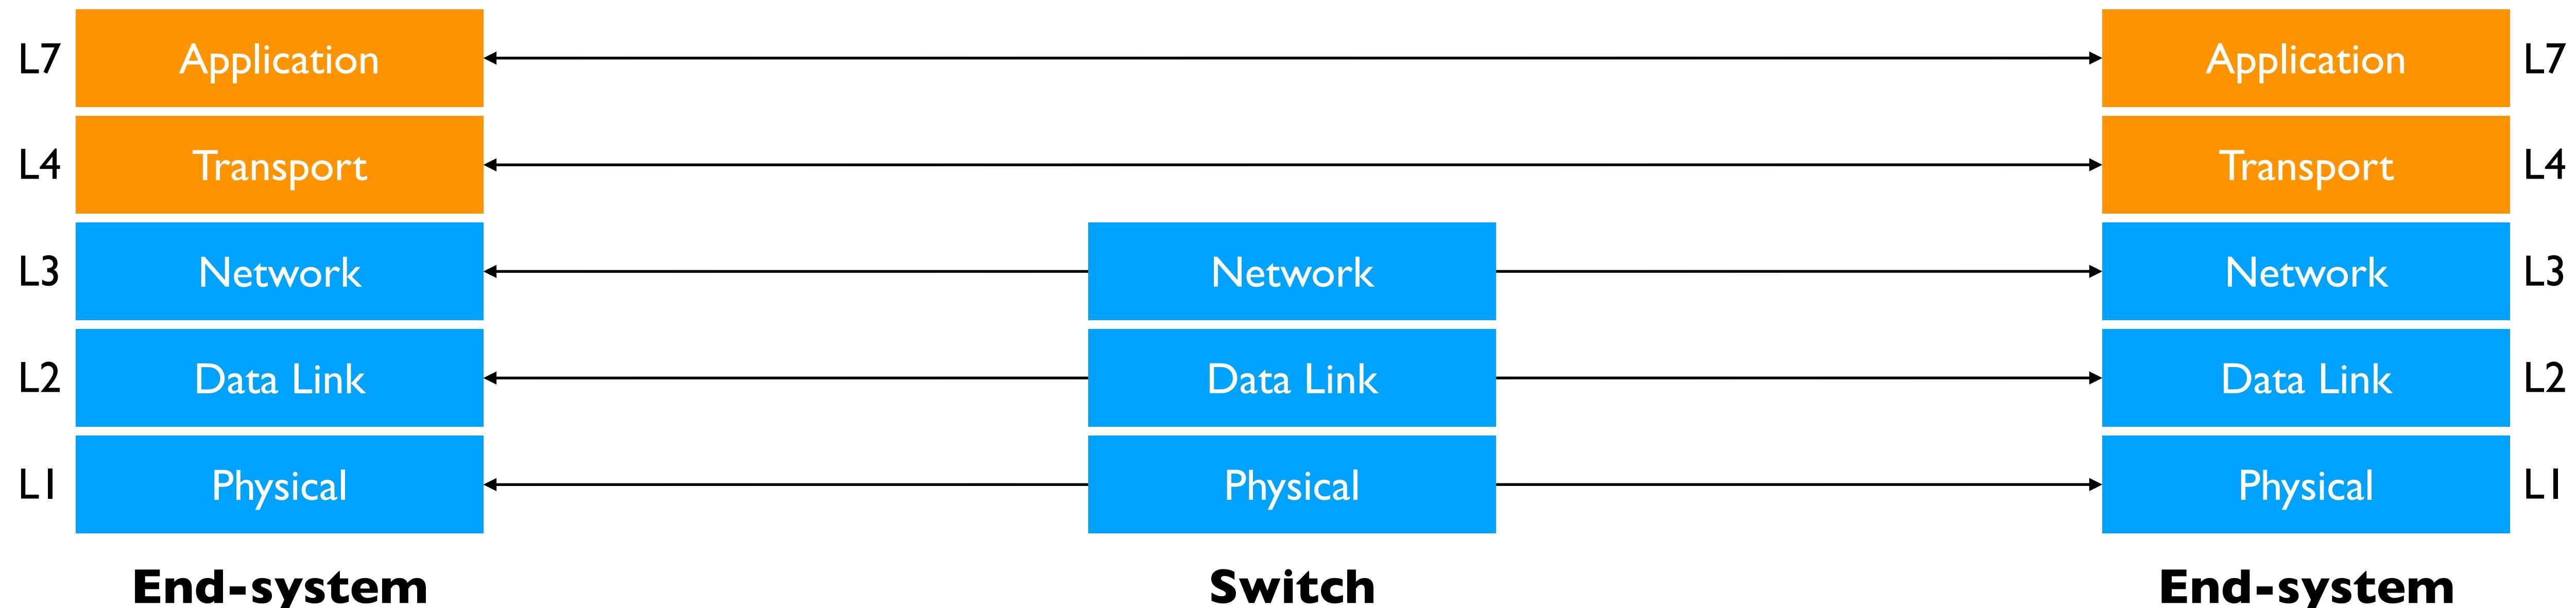
# Path through Internet

- Lower three layers implemented everywhere



# Path through Internet

- Lower three layers implemented everywhere
- Top two layers implemented only on hosts



# A closer look: end-system

# A closer look: end-system

- Application
  - *Web server, browser, mail, game*

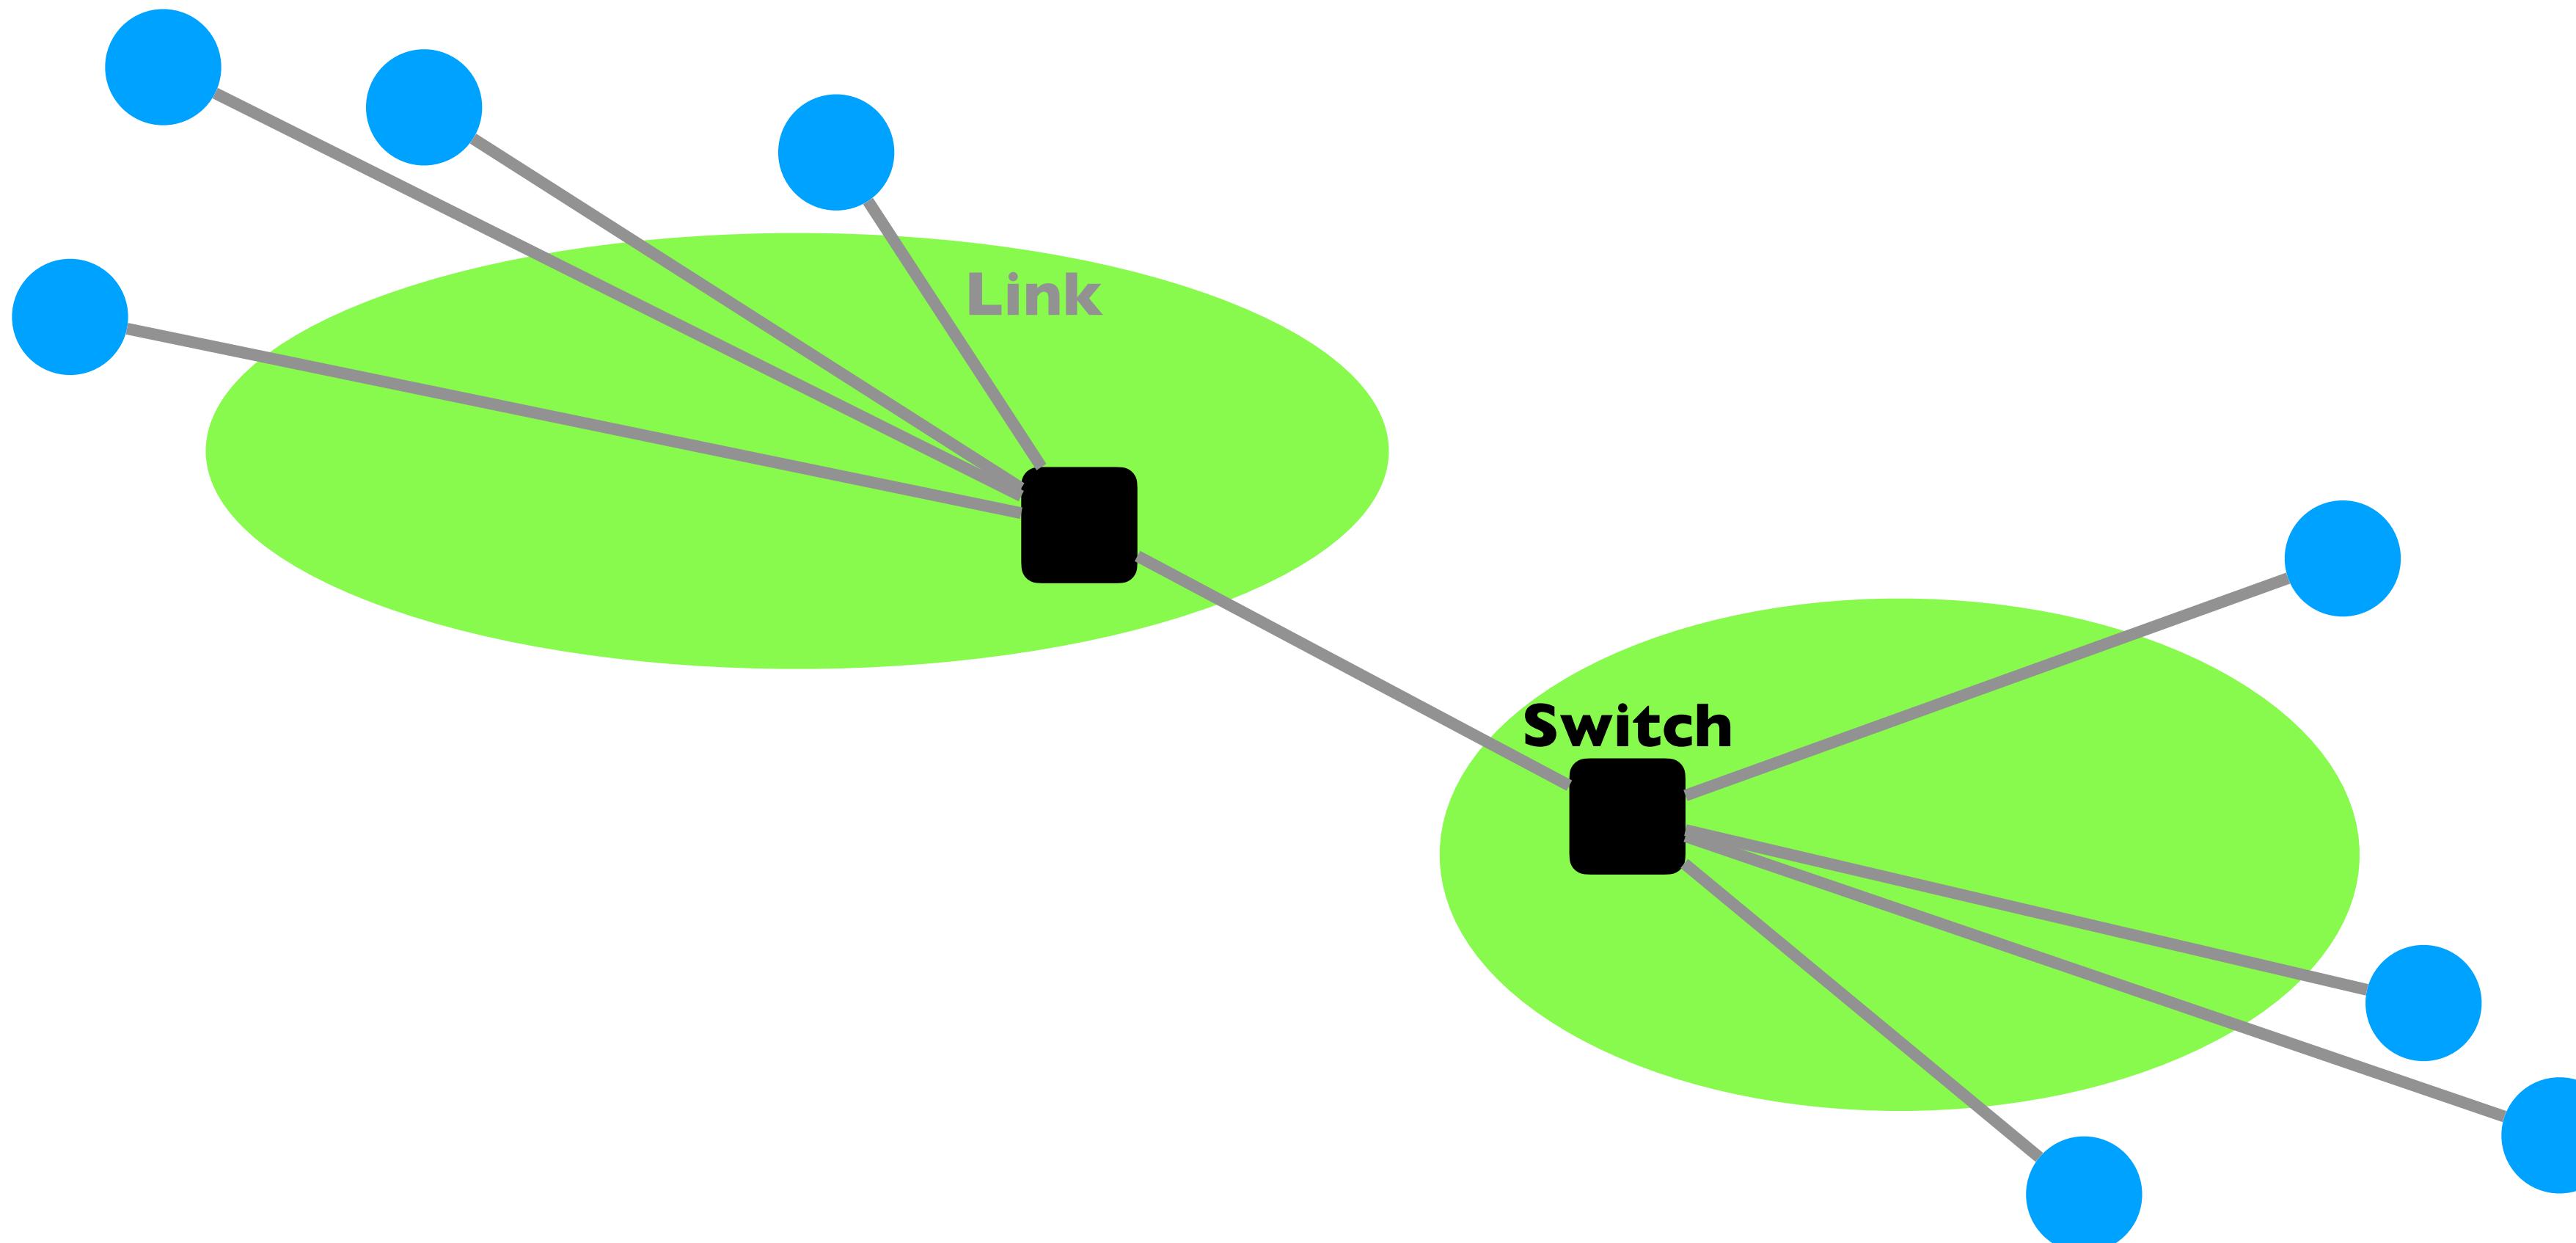
# A closer look: end-system

- Application
  - *Web server, browser, mail, game*
- Transport and network layer
  - *Typically part of the operating system*

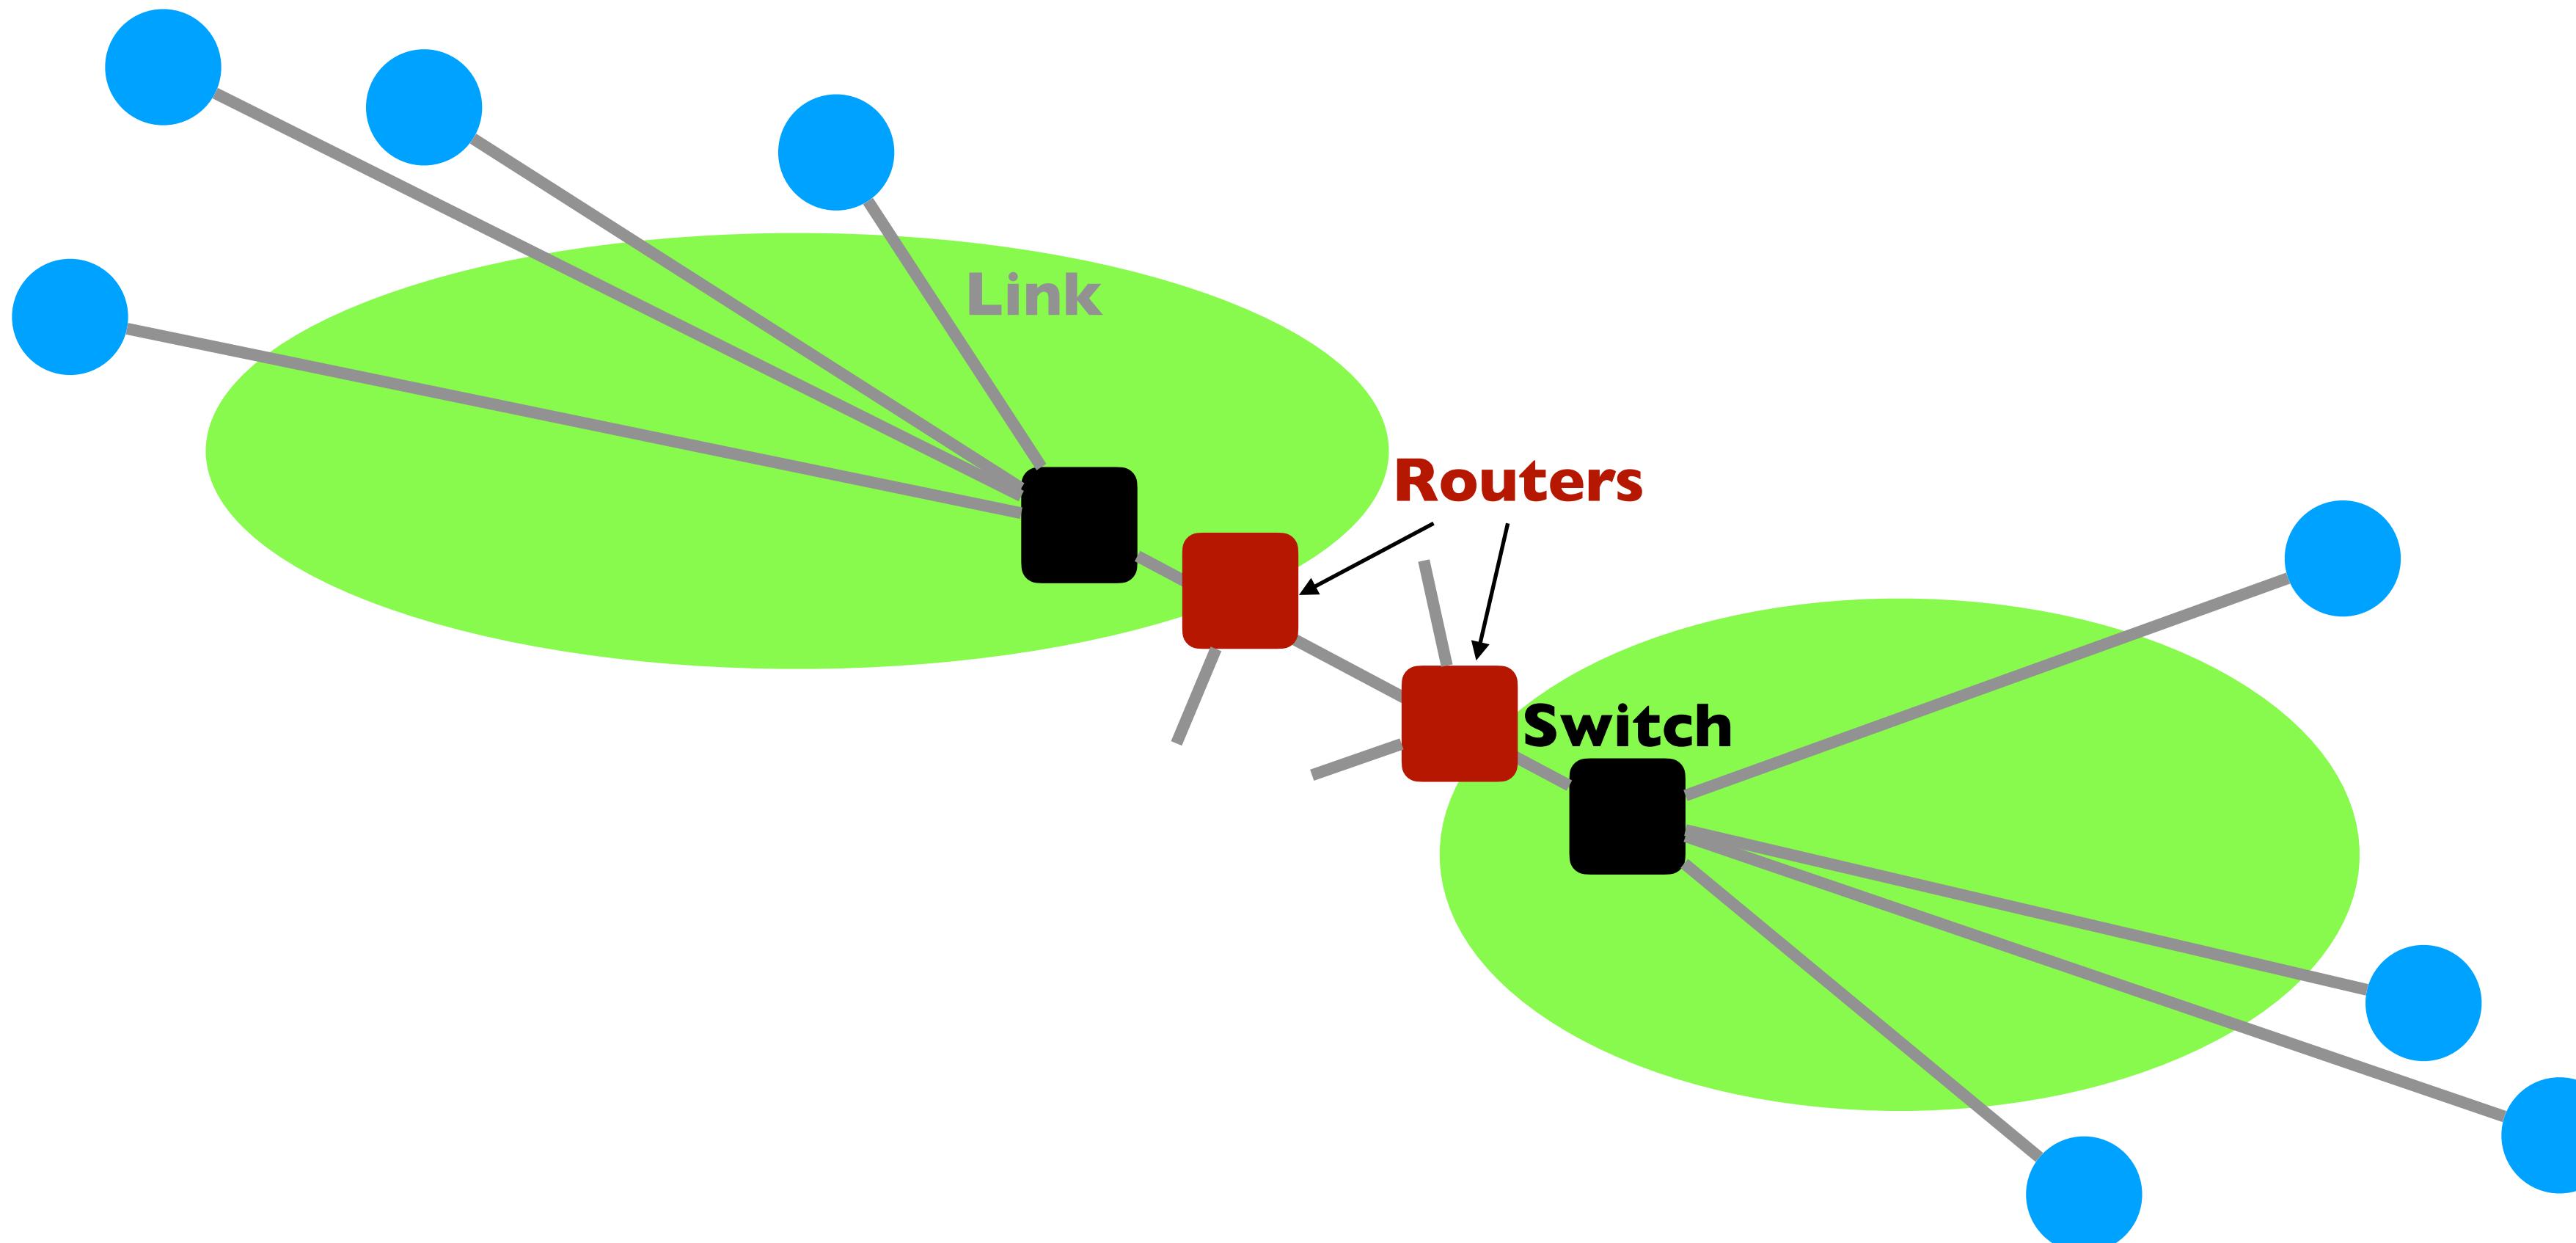
# A closer look: end-system

- Application
  - *Web server, browser, mail, game*
- Transport and network layer
  - *Typically part of the operating system*
- Datalink and physical layer
  - *Hardware/firmware/drivers*

# A closer look: Network



# A closer look: Network



# Switches vs. Routers

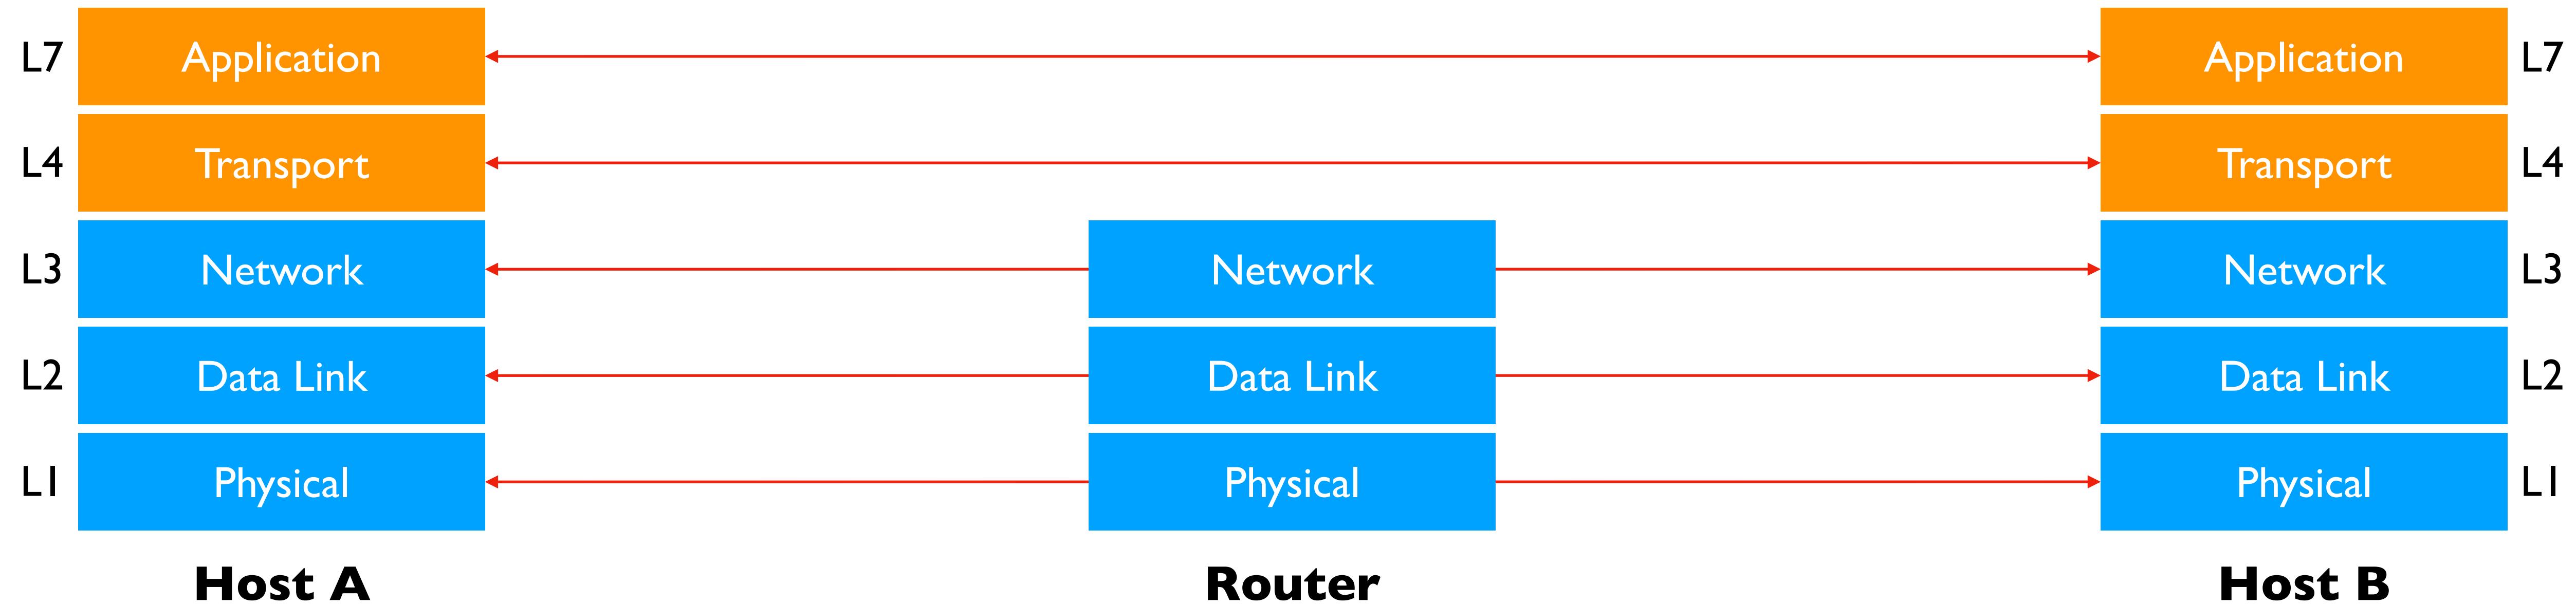
# Switches vs. Routers

- Switches do what routers do, but don't participate in global delivery, just local delivery
  - *Switches only need to support L1, L2*
  - *Routers support L1-L3*

# Switches vs. Routers

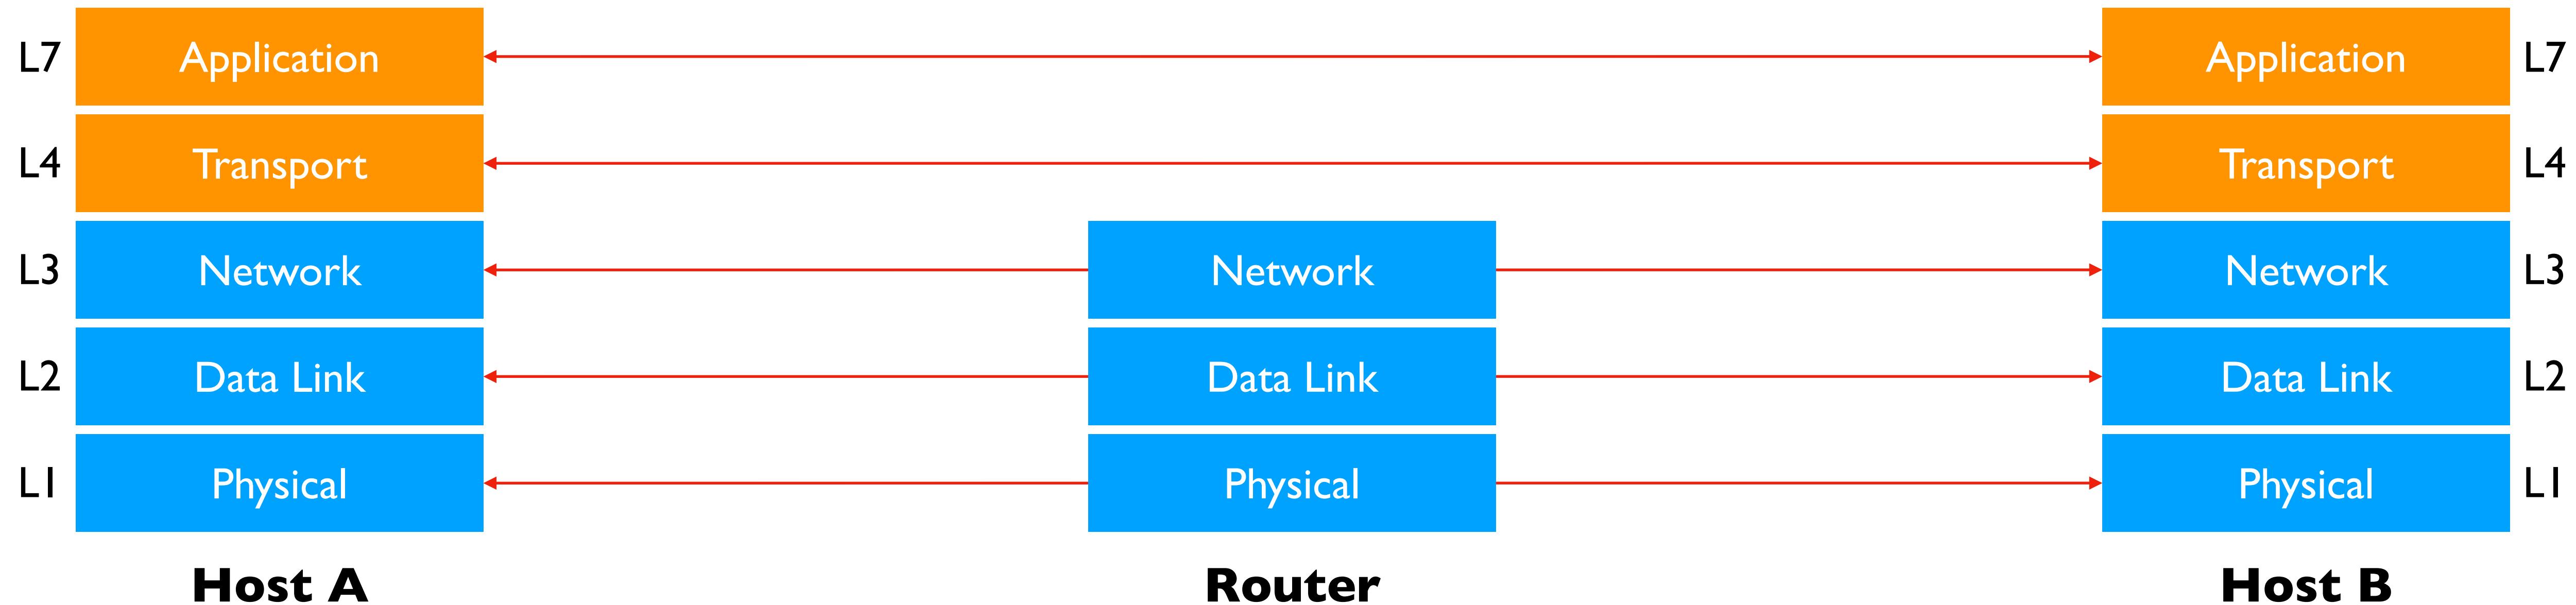
- Switches do what routers do, but don't participate in global delivery, just local delivery
  - *Switches only need to support L1, L2*
  - *Routers support L1-L3*
- Won't focus on the router/switch distinction
  - *When I say switch, I almost always mean router*
  - *Almost all boxes support network layer these days*

# Logical Communication

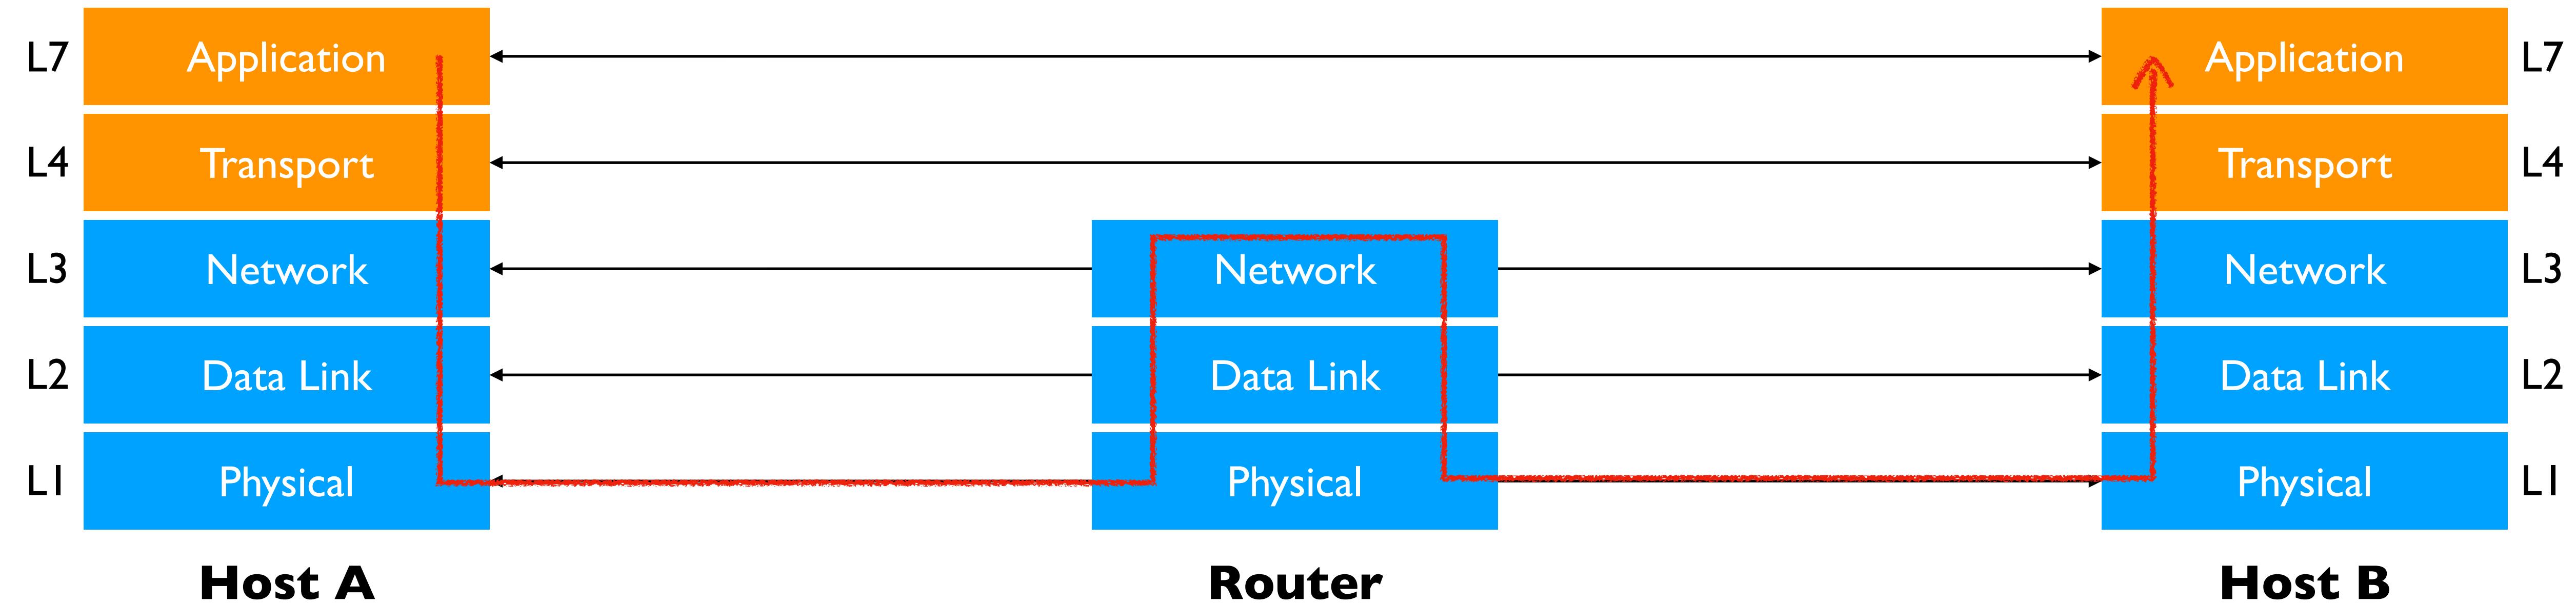


# Logical Communication

- Layers interact with peer's corresponding layer

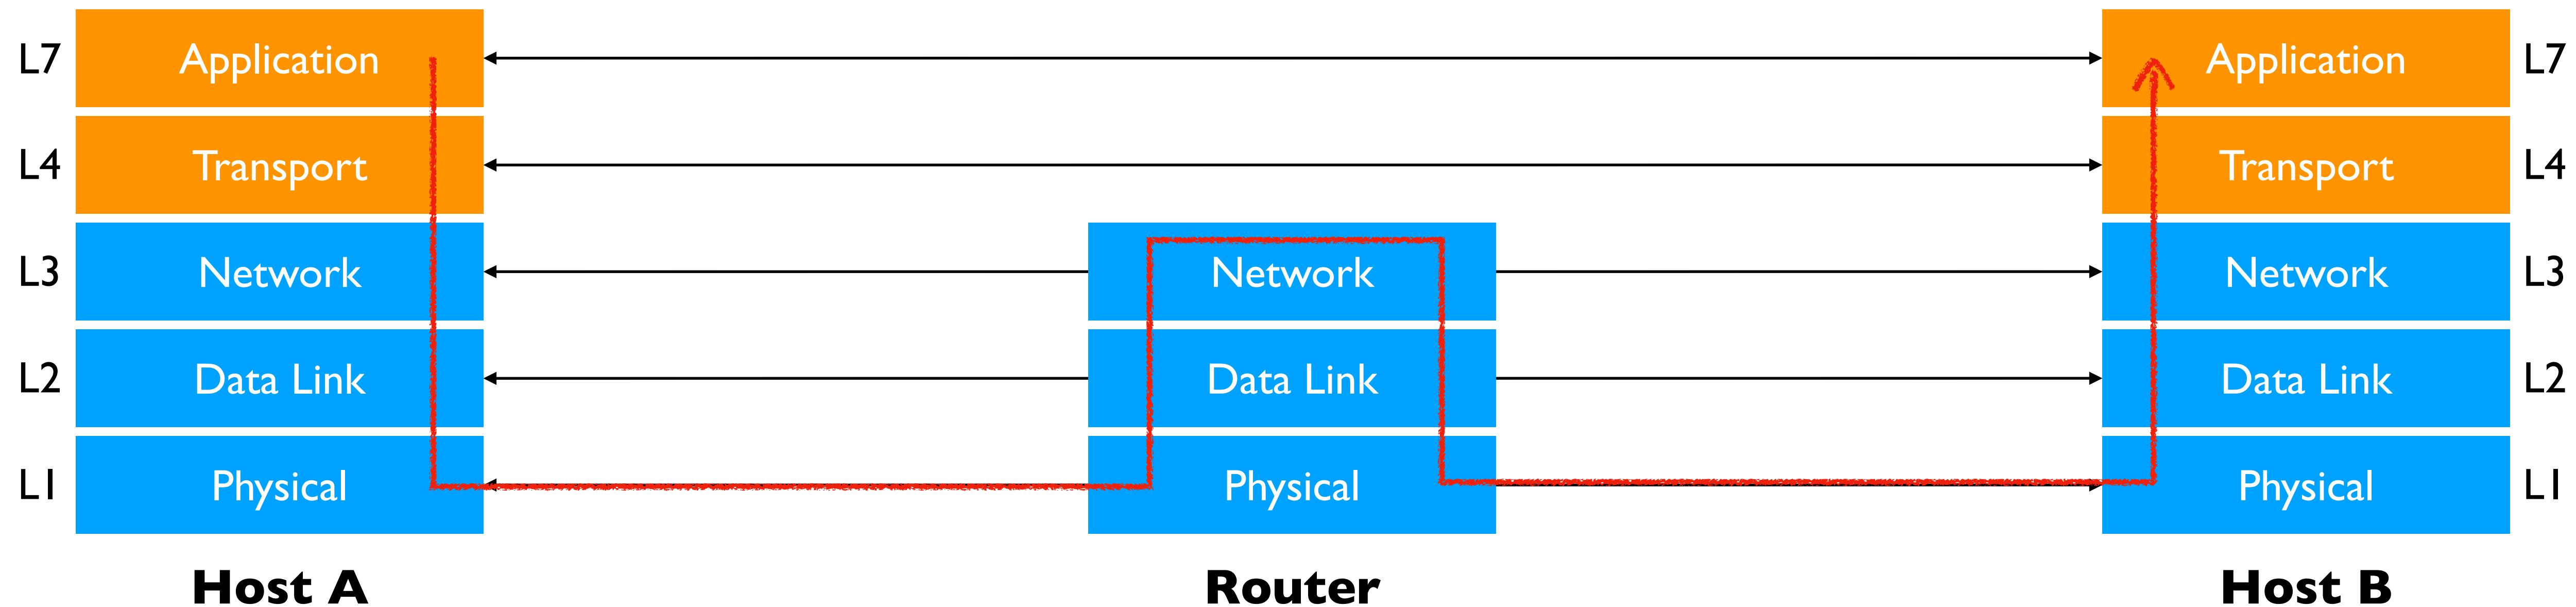


# Physical Communication



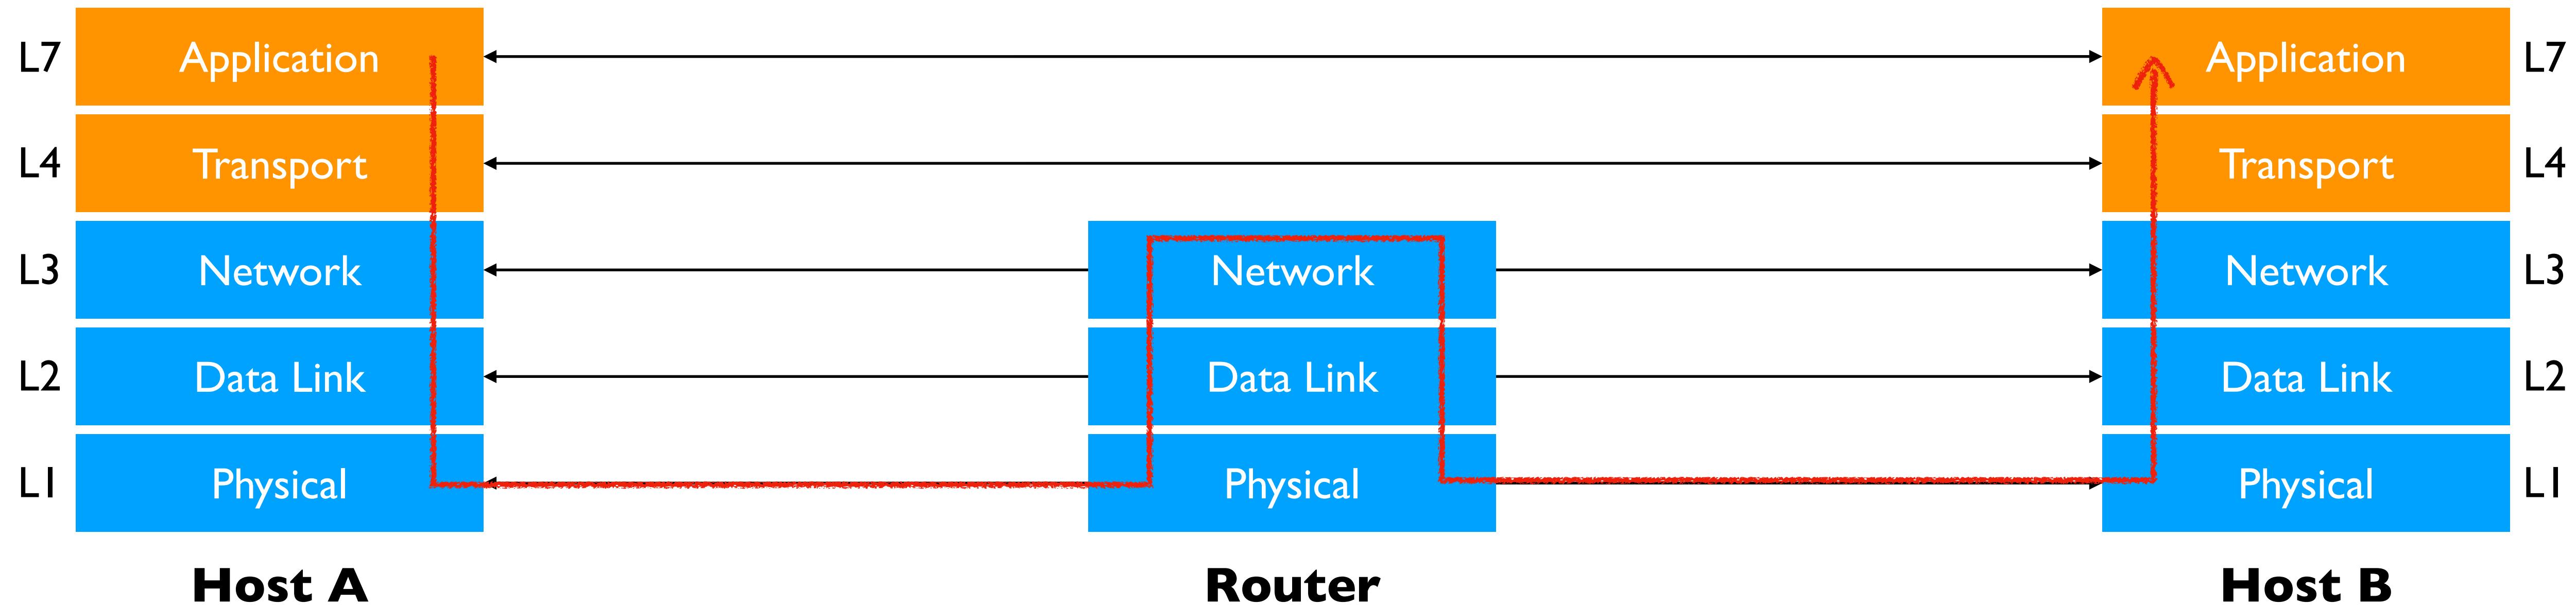
# Physical Communication

- Communication goes down to the physical network

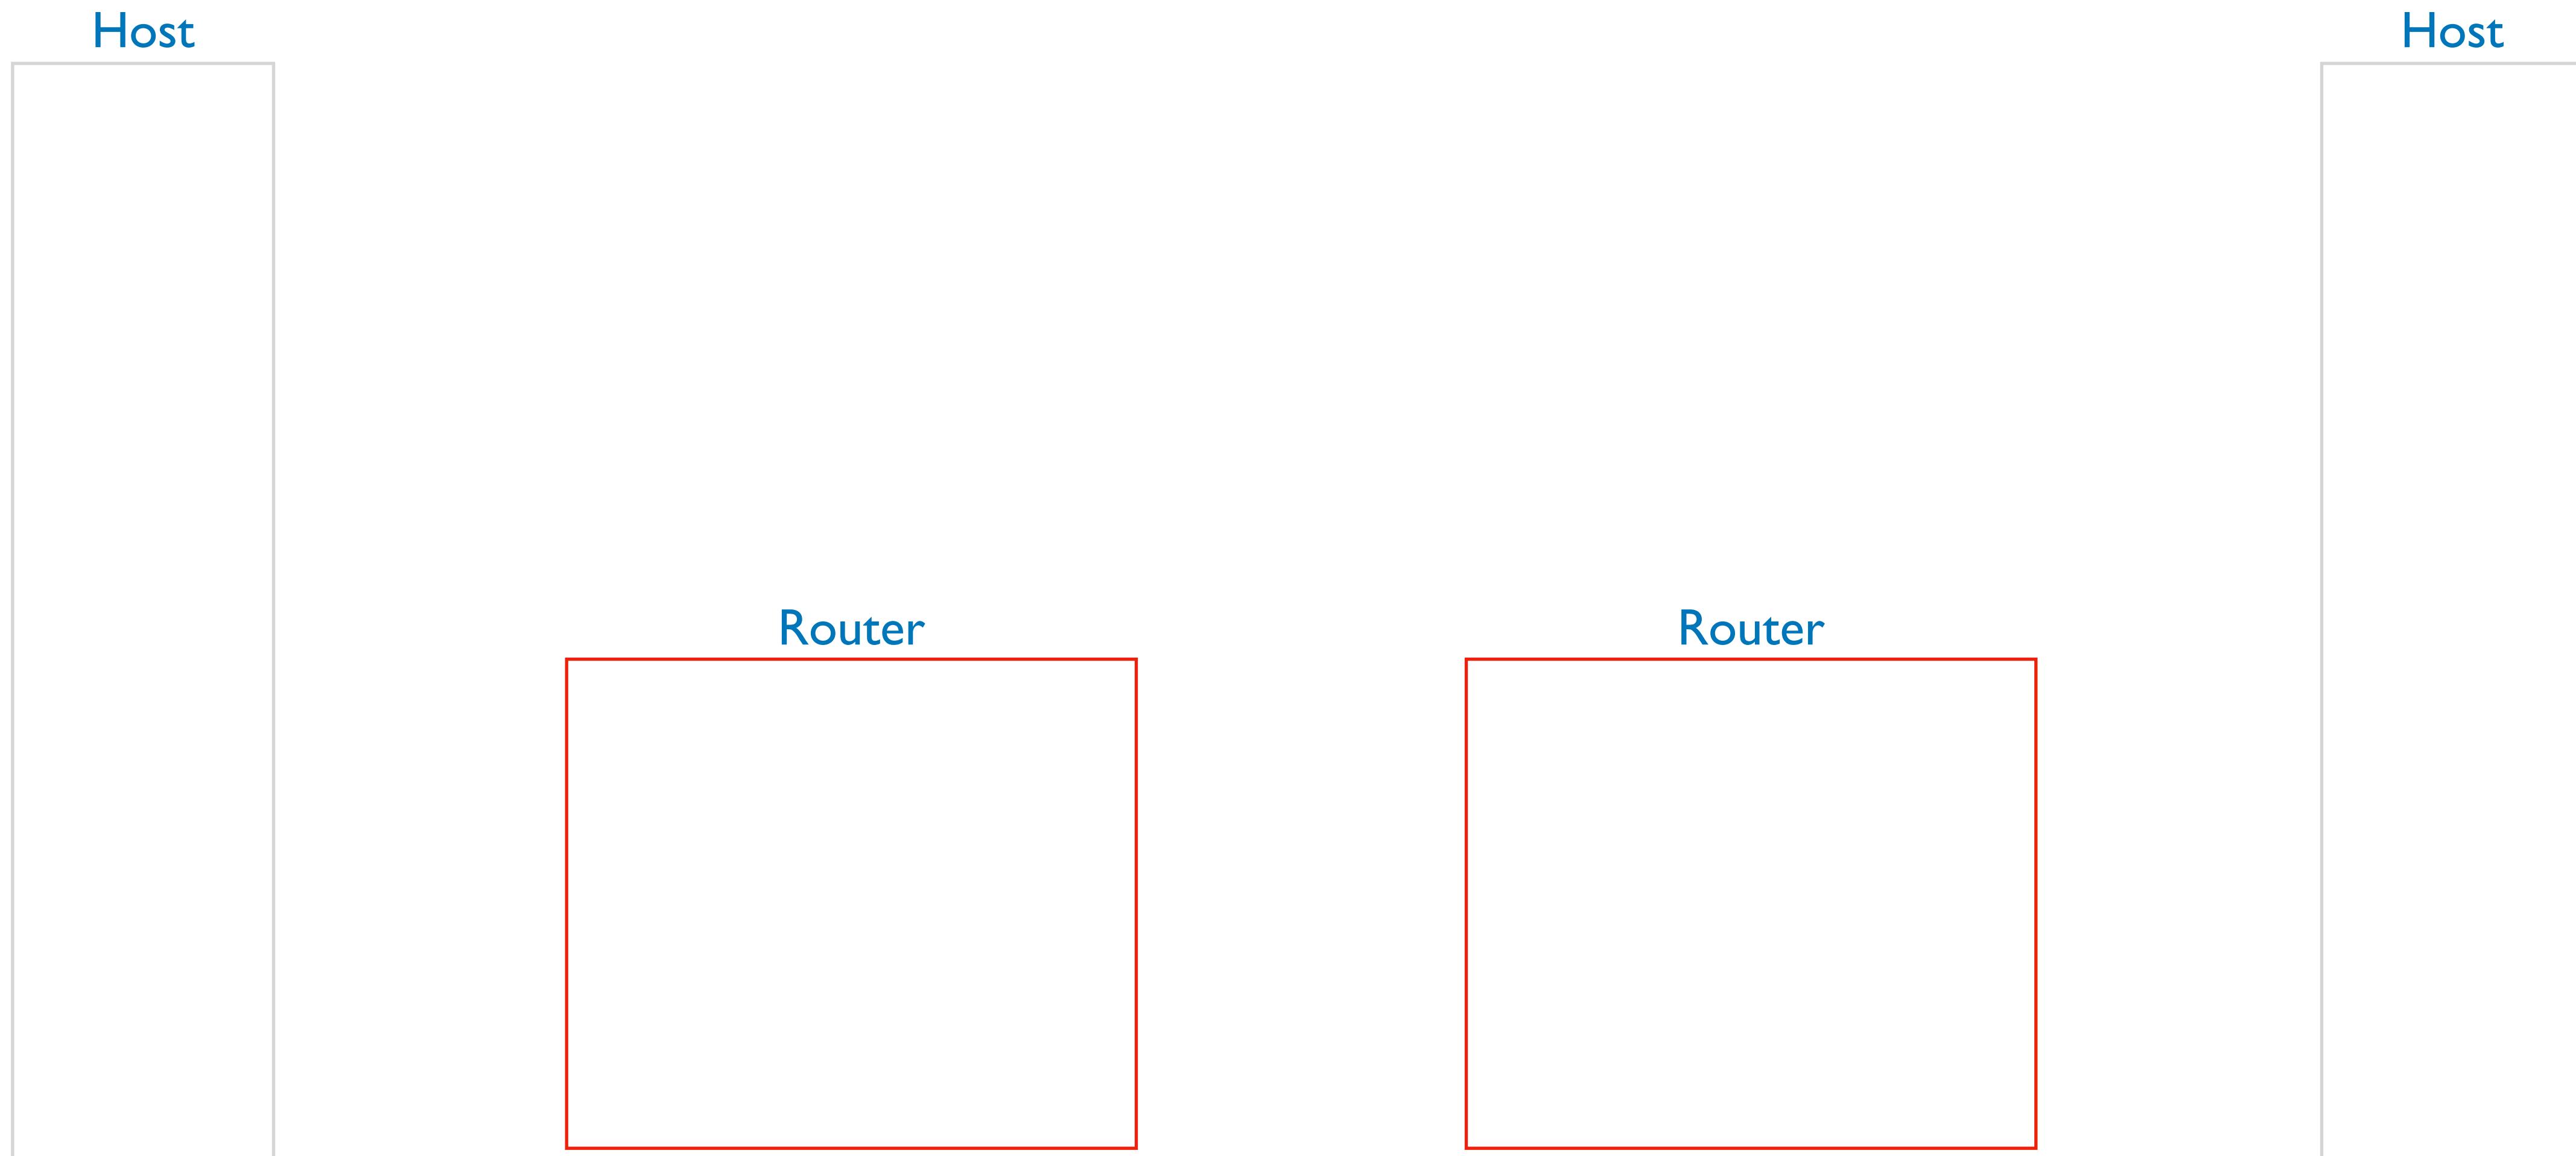


# Physical Communication

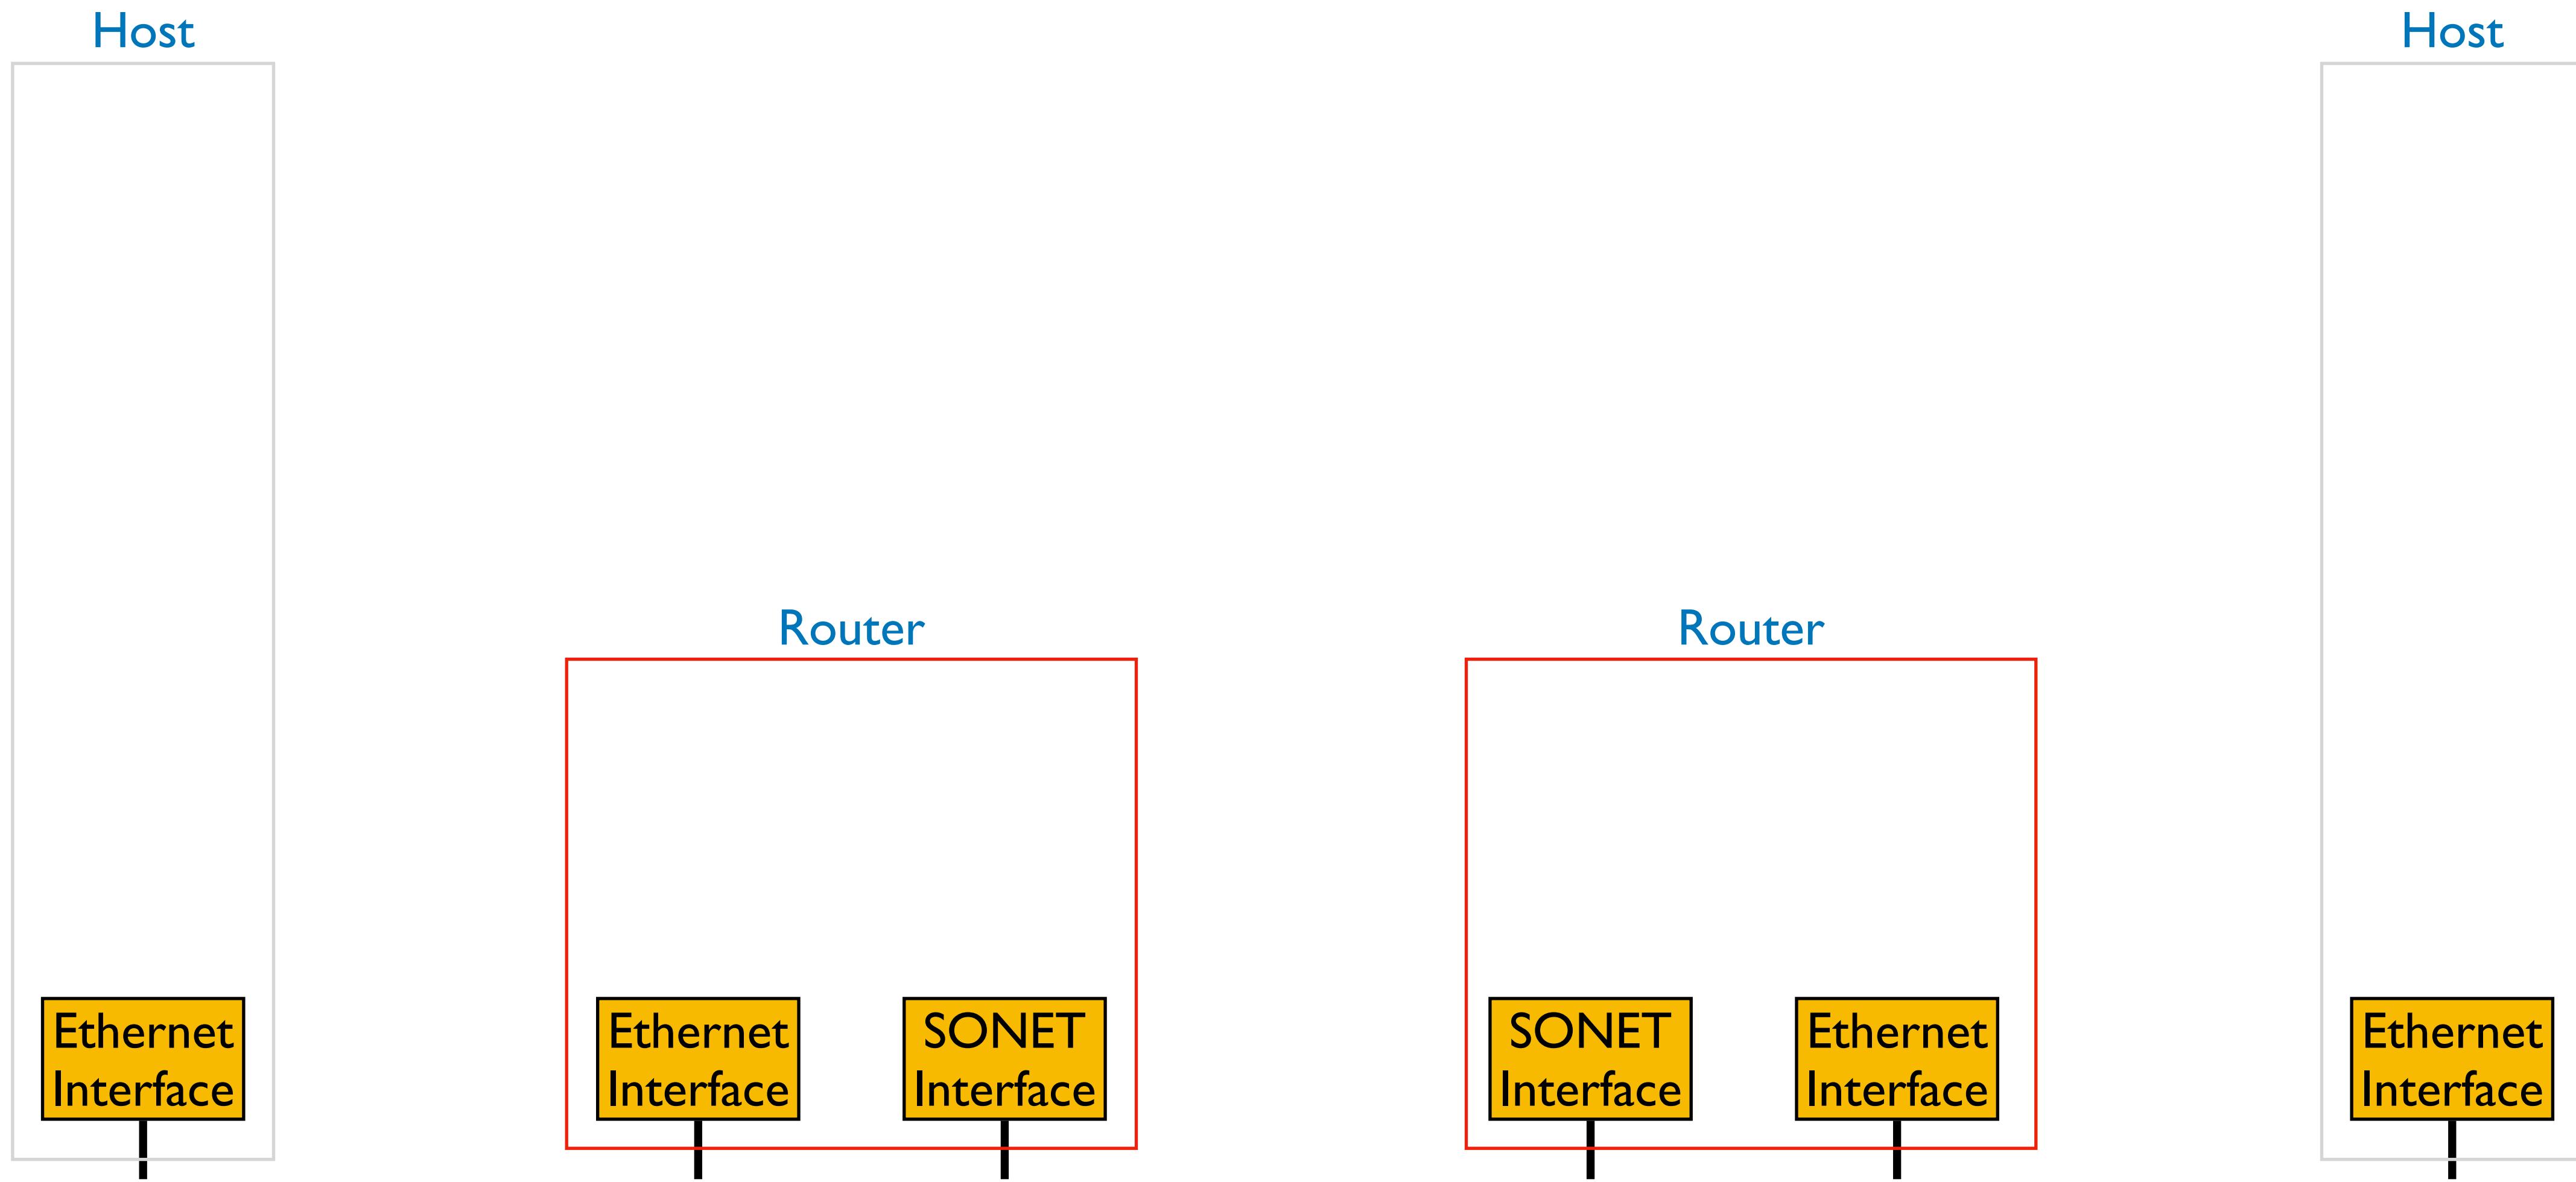
- Communication goes down to the physical network
- Then up to relevant layer



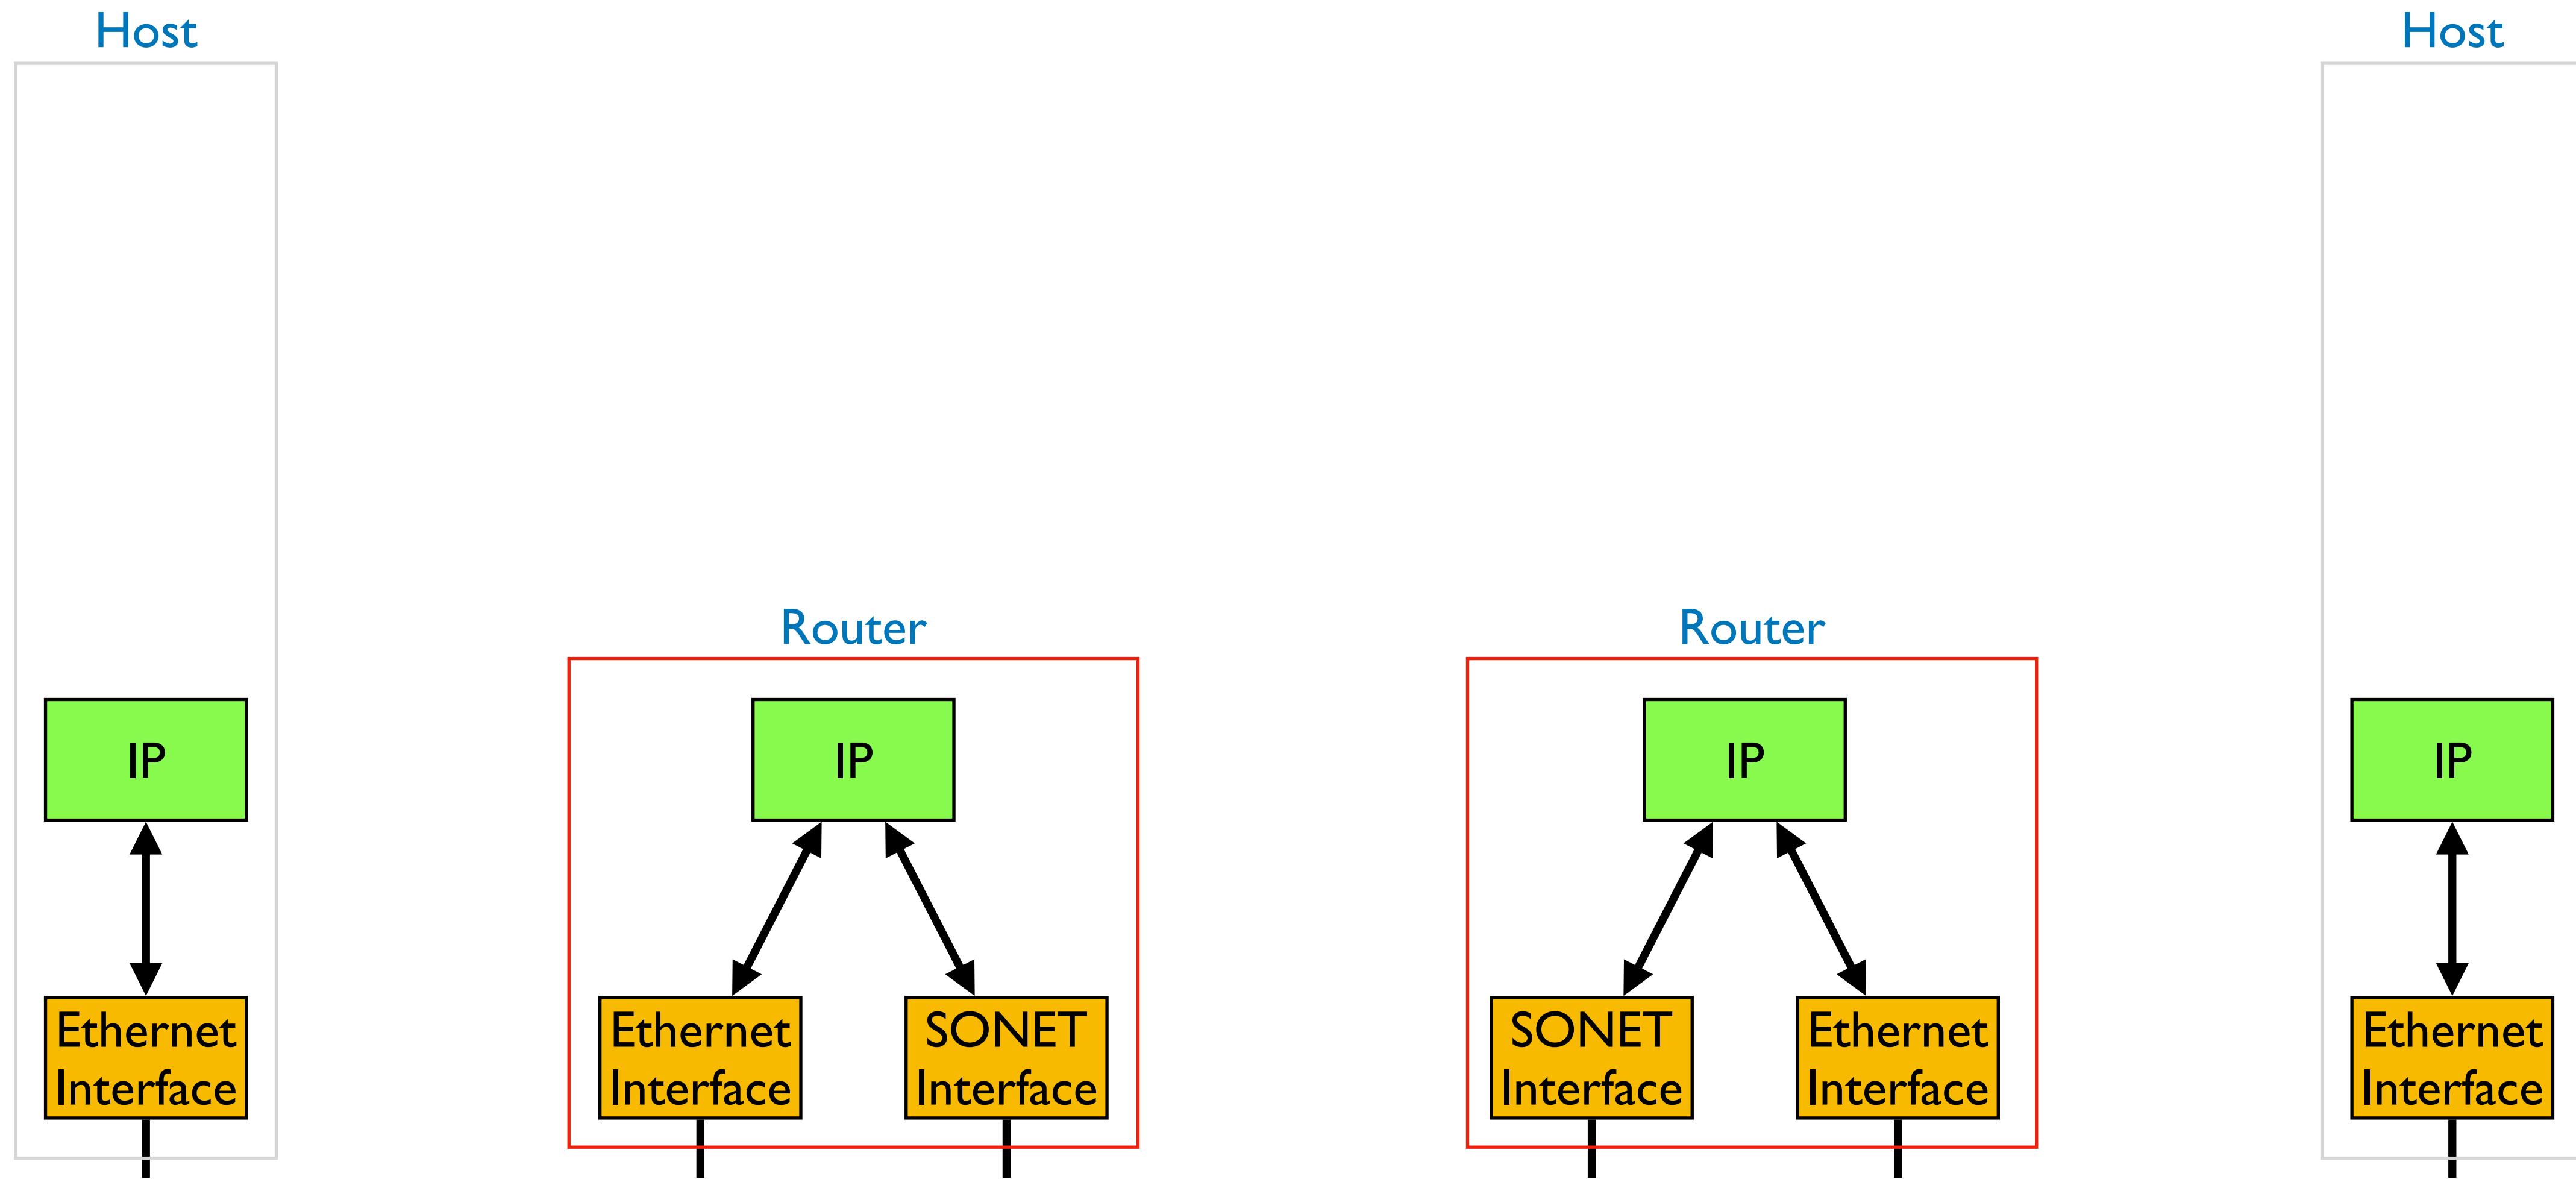
# A Protocol-centric Diagram



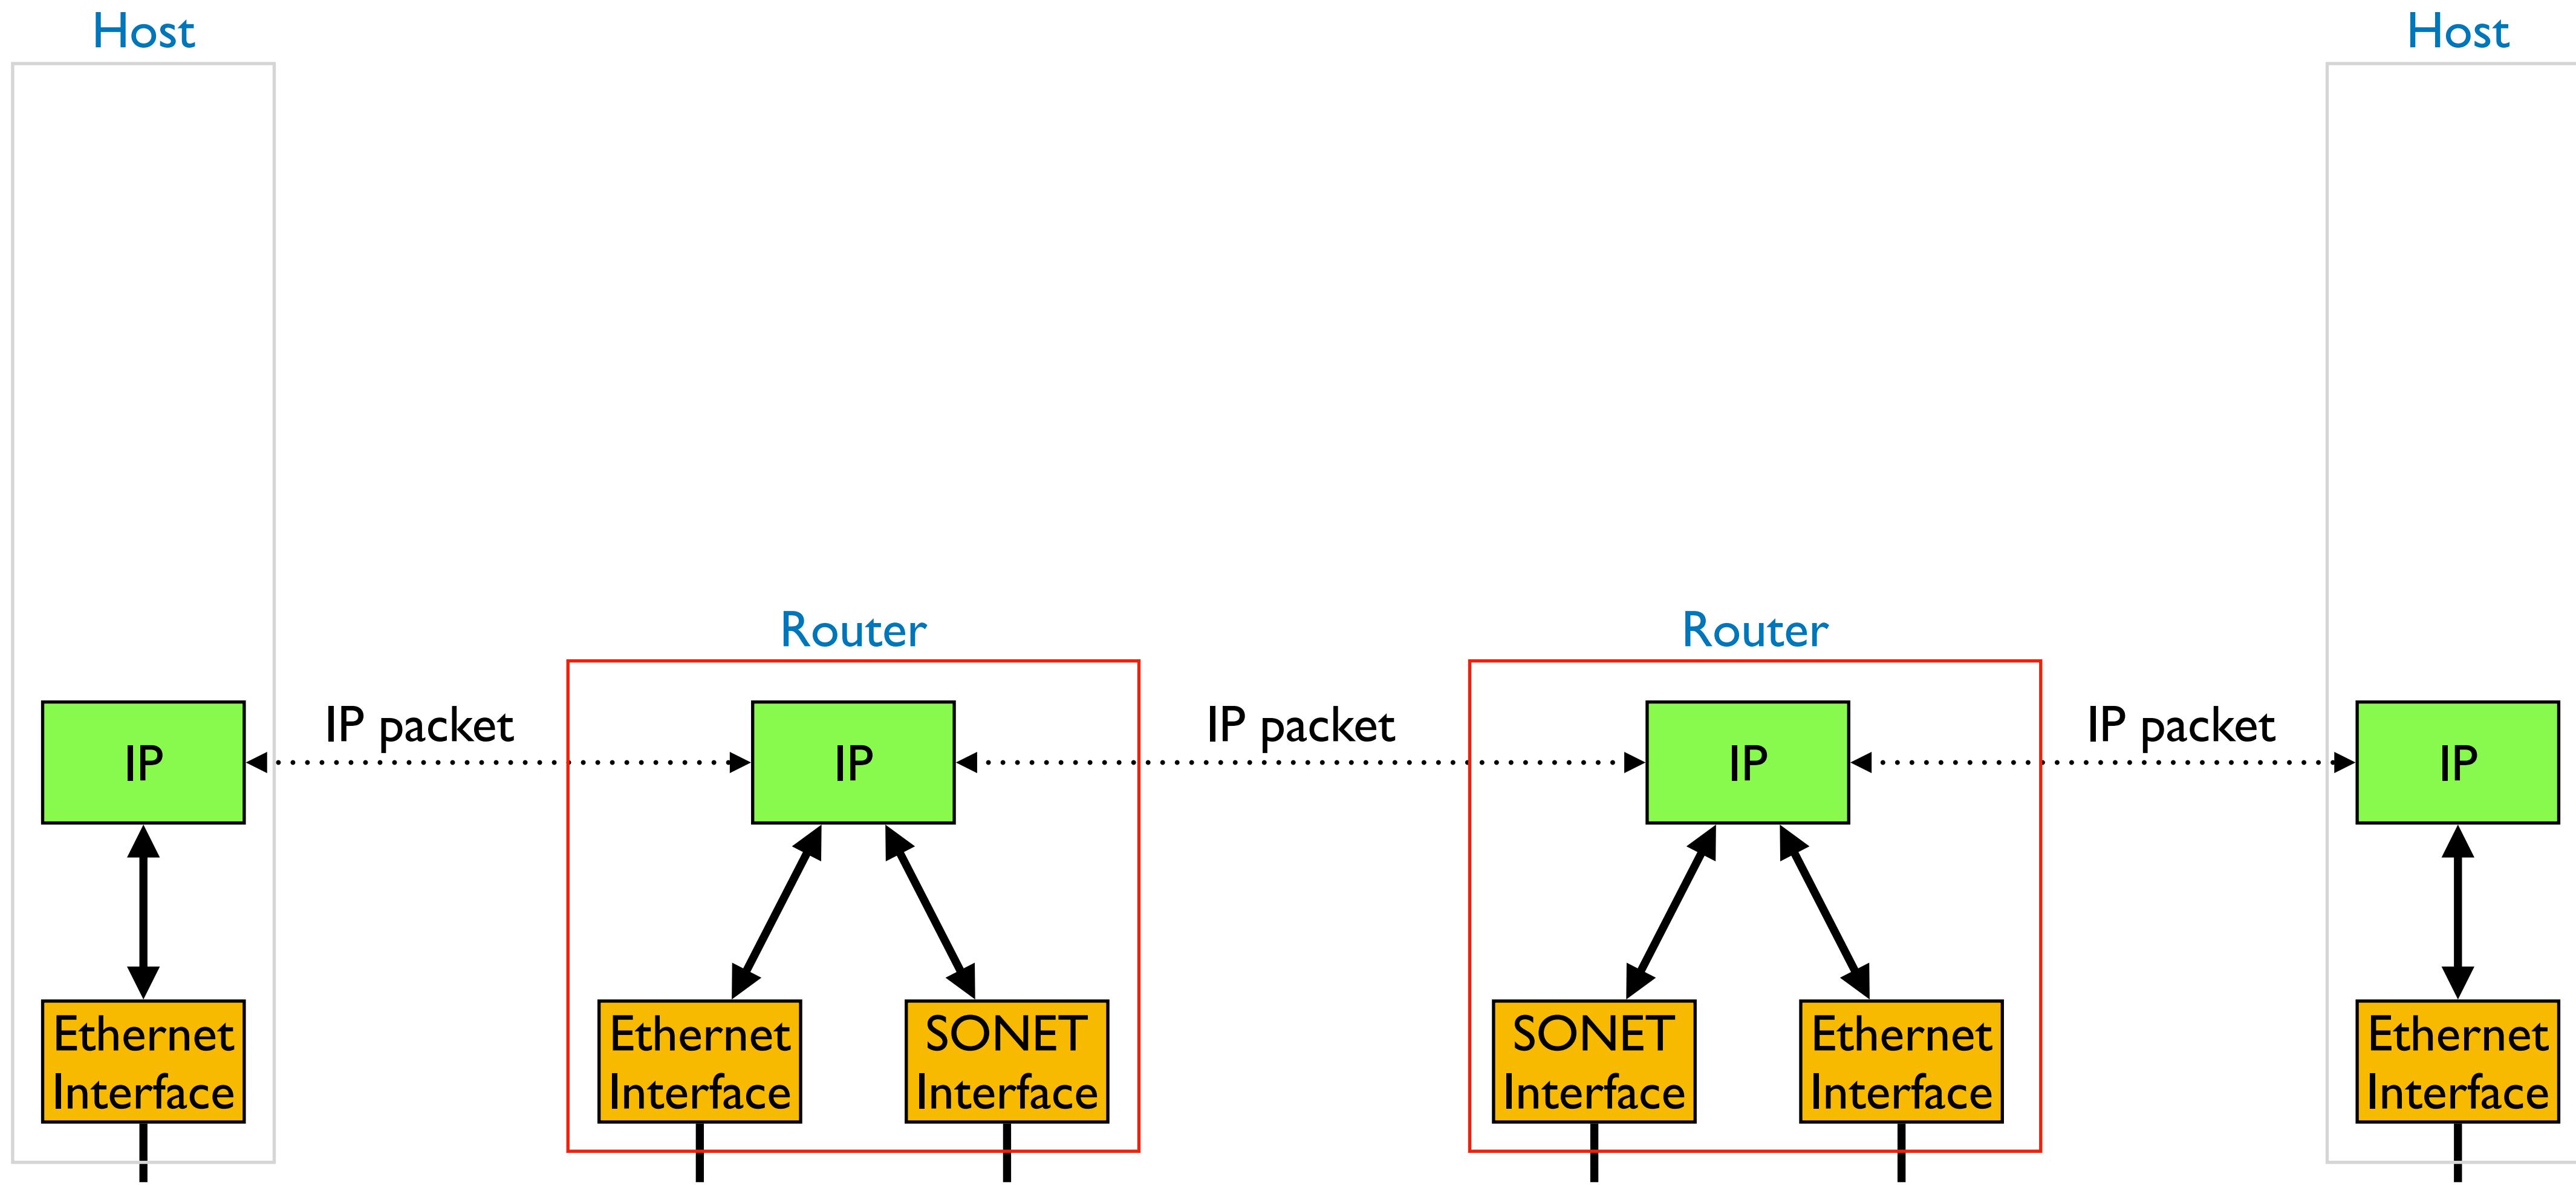
# A Protocol-centric Diagram



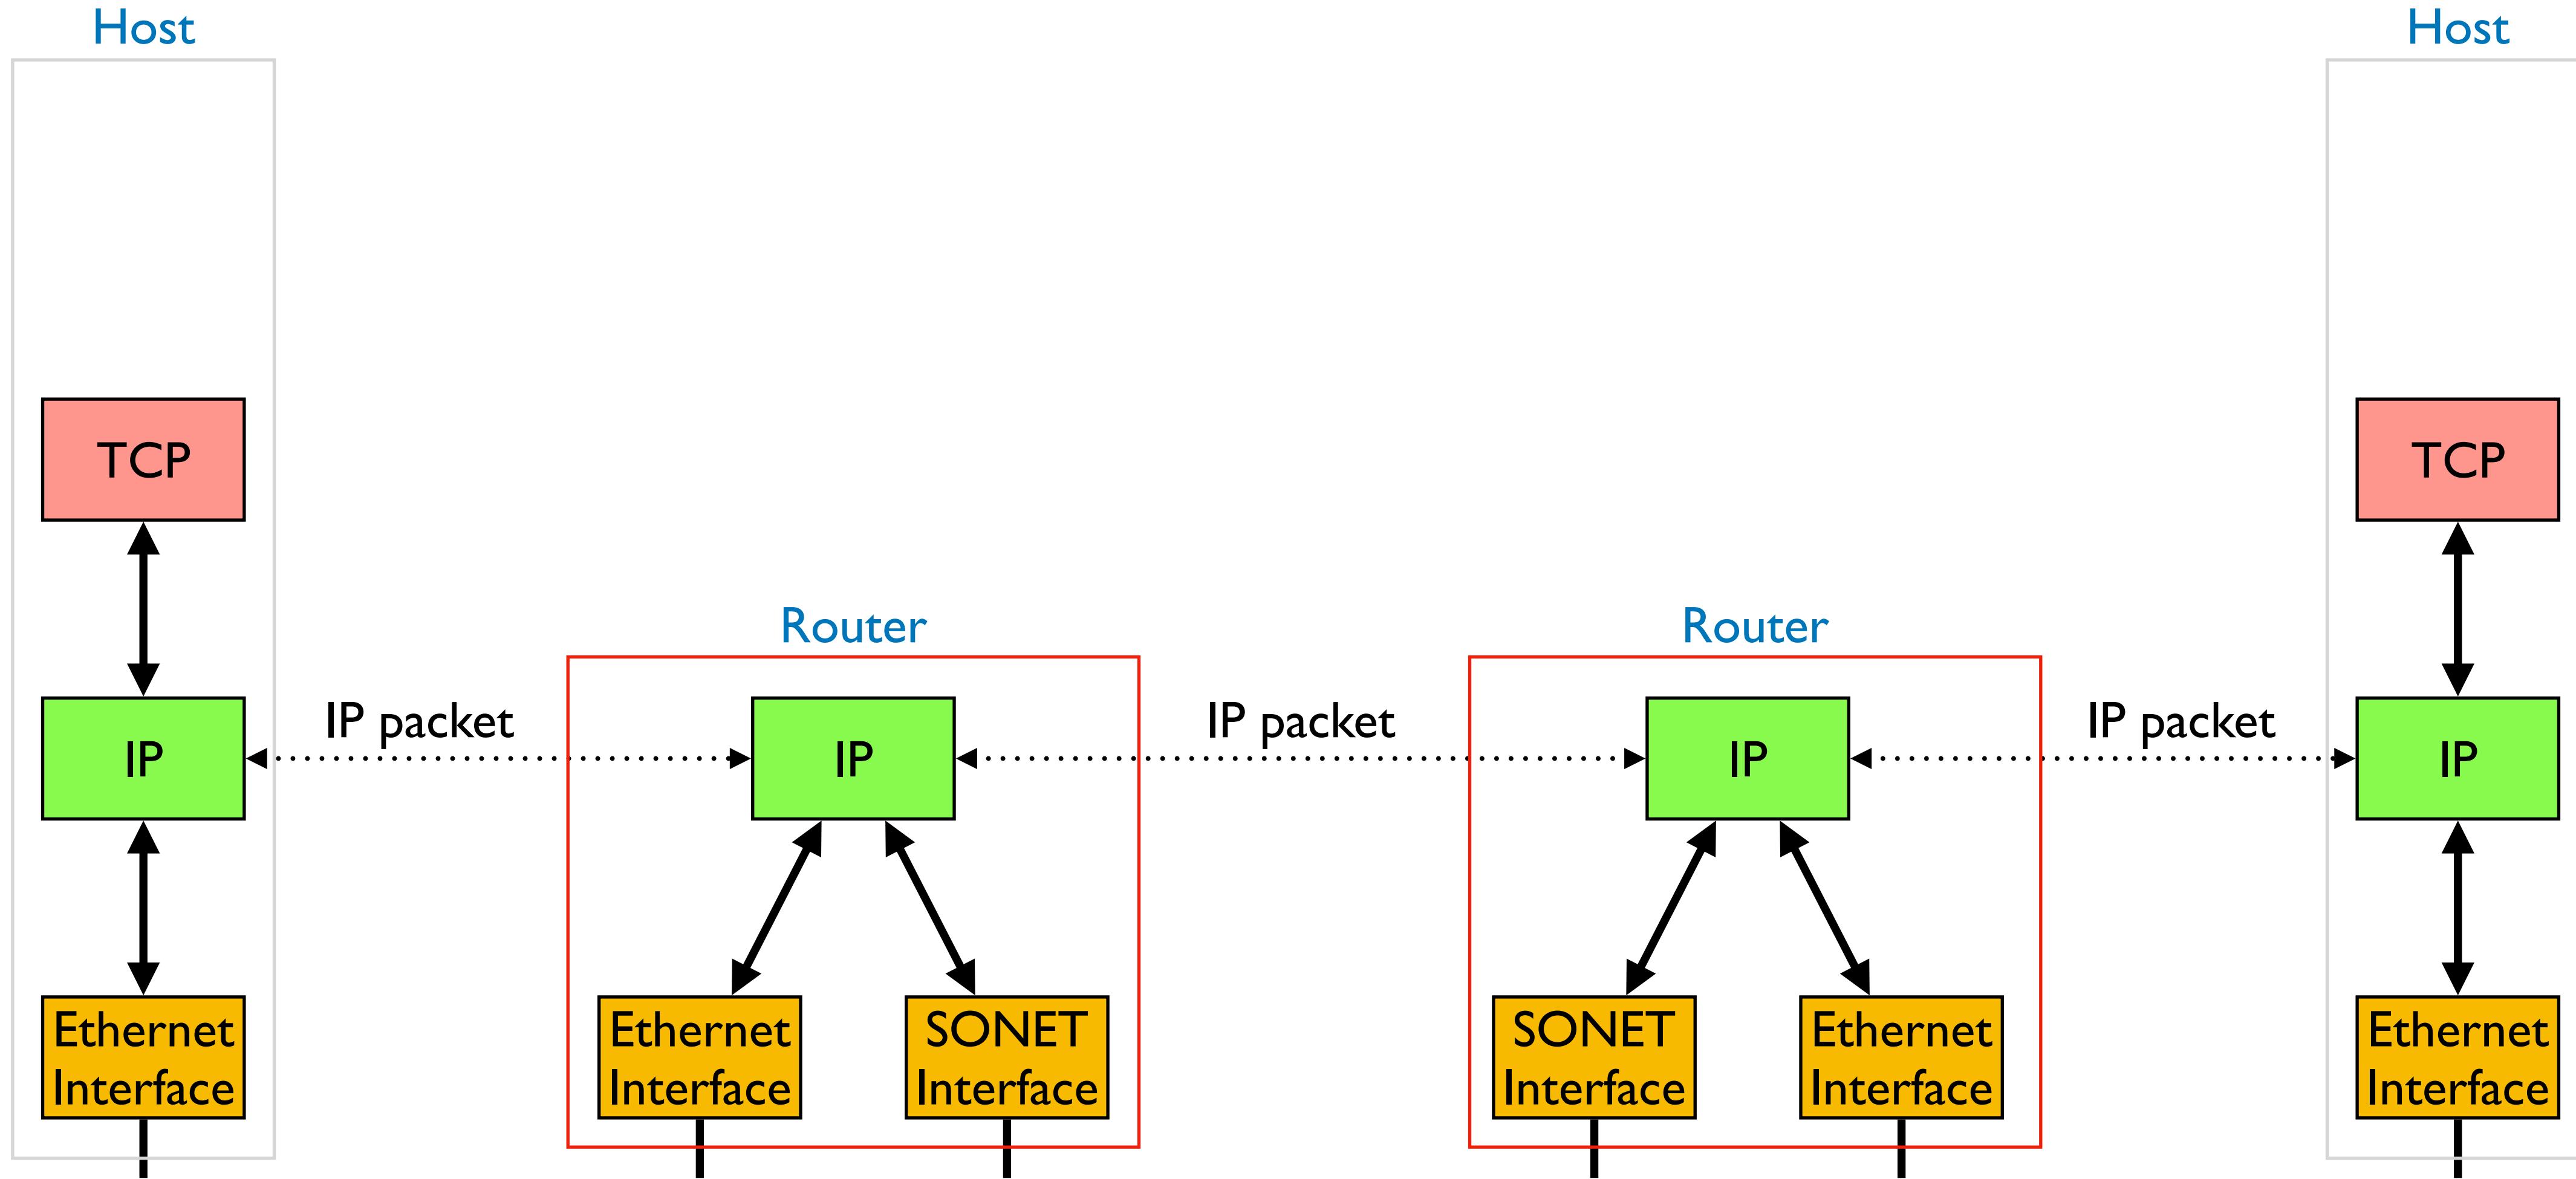
# A Protocol-centric Diagram



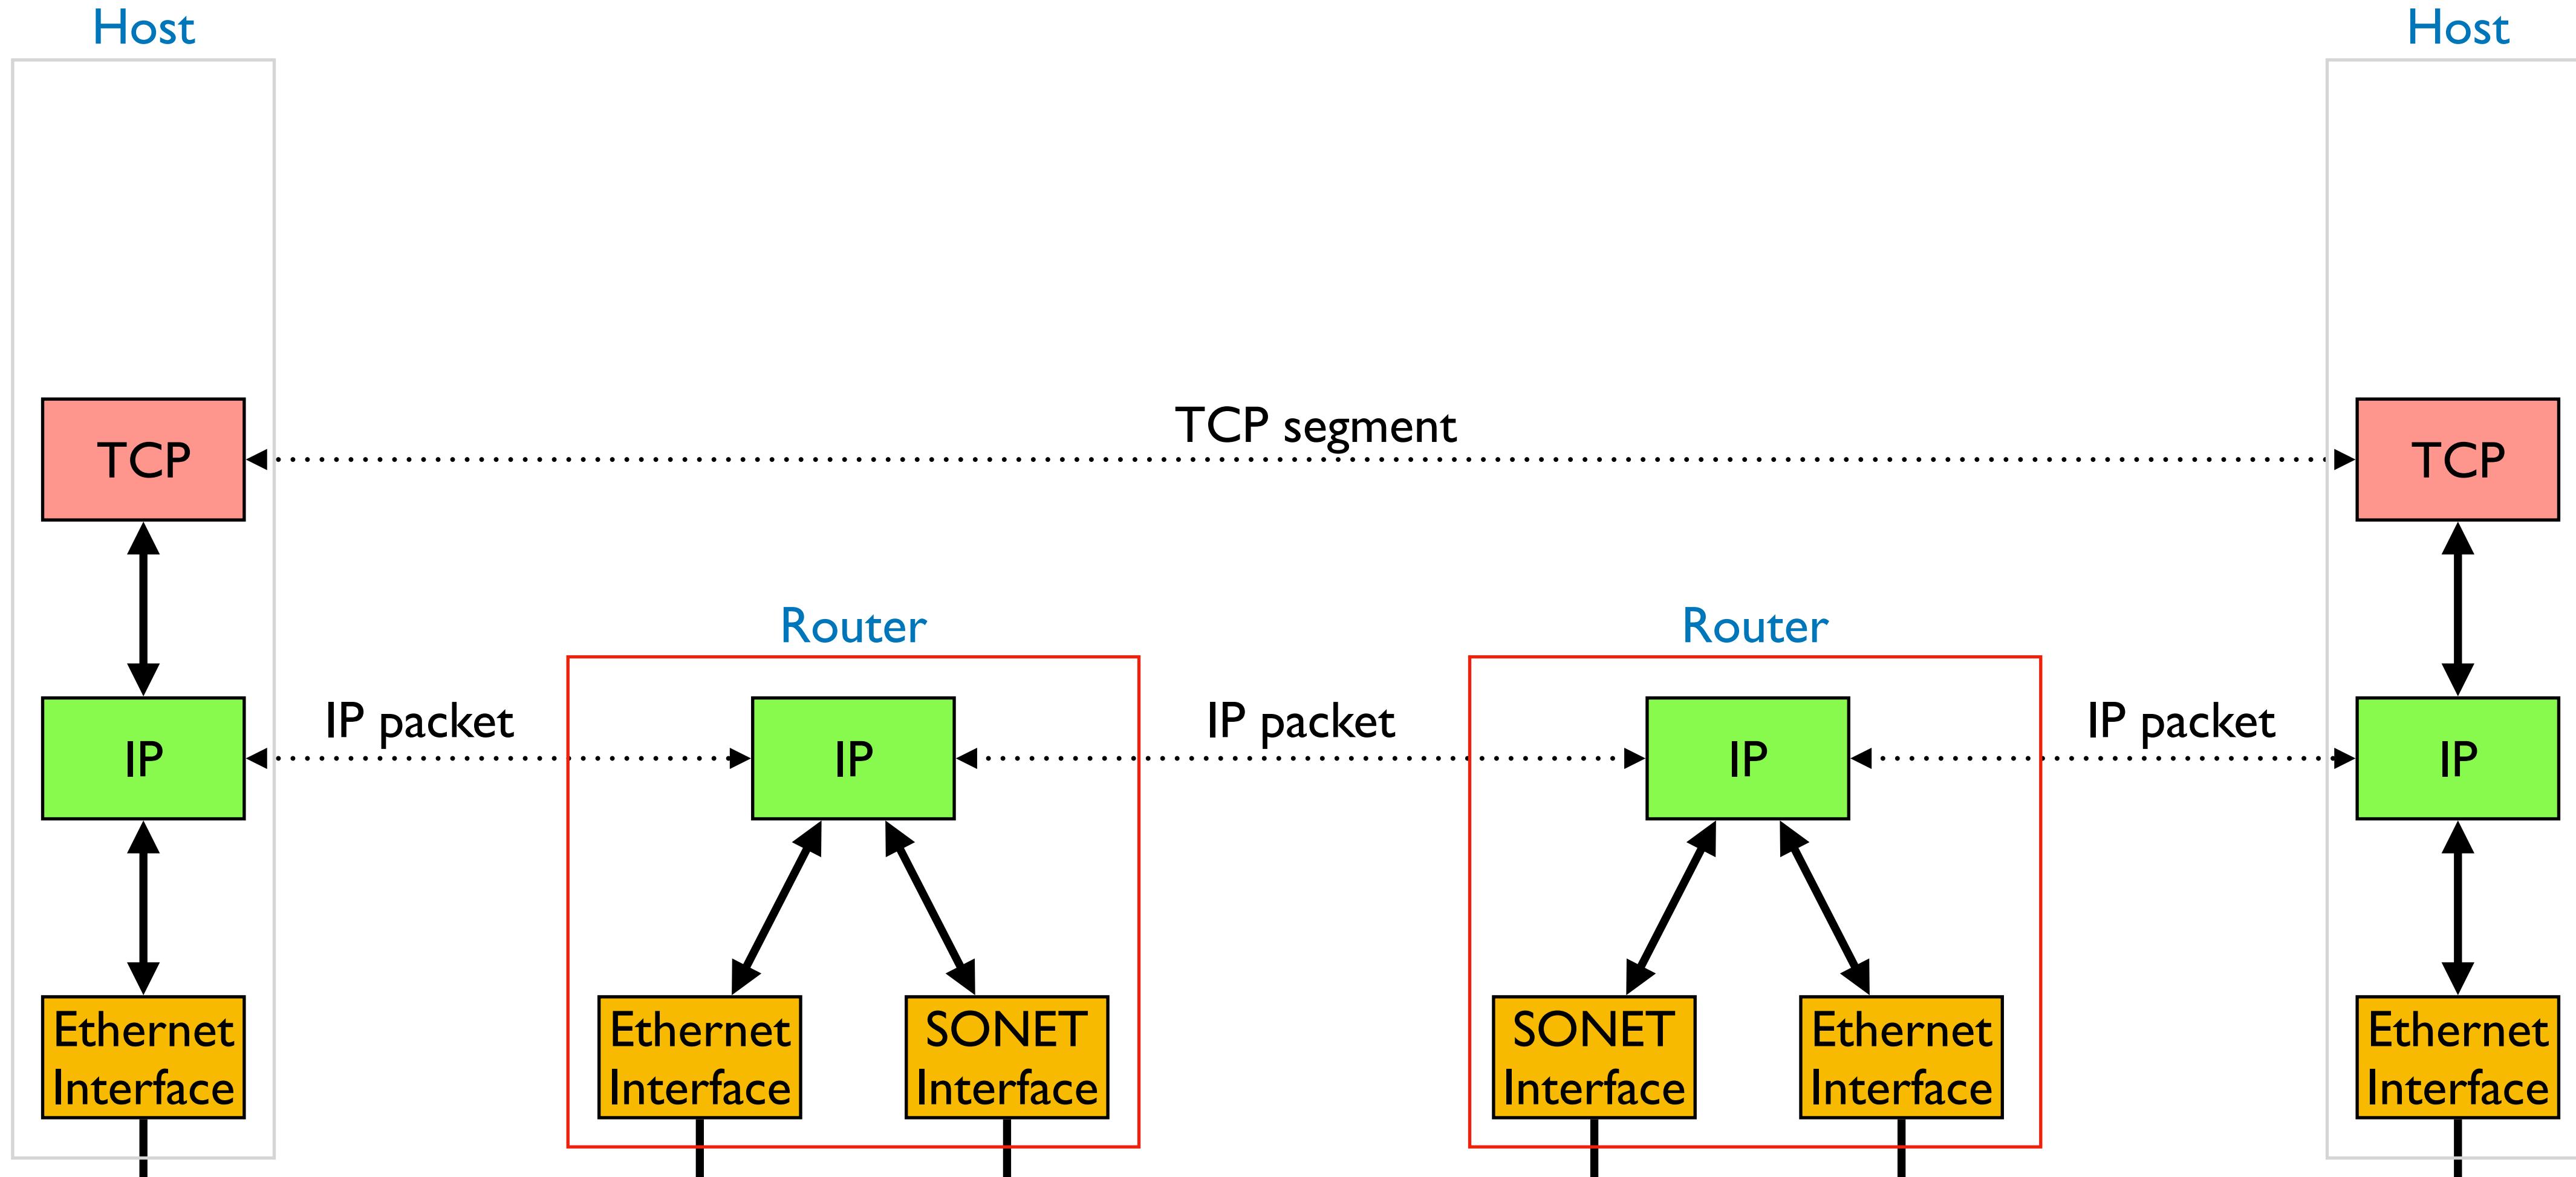
# A Protocol-centric Diagram



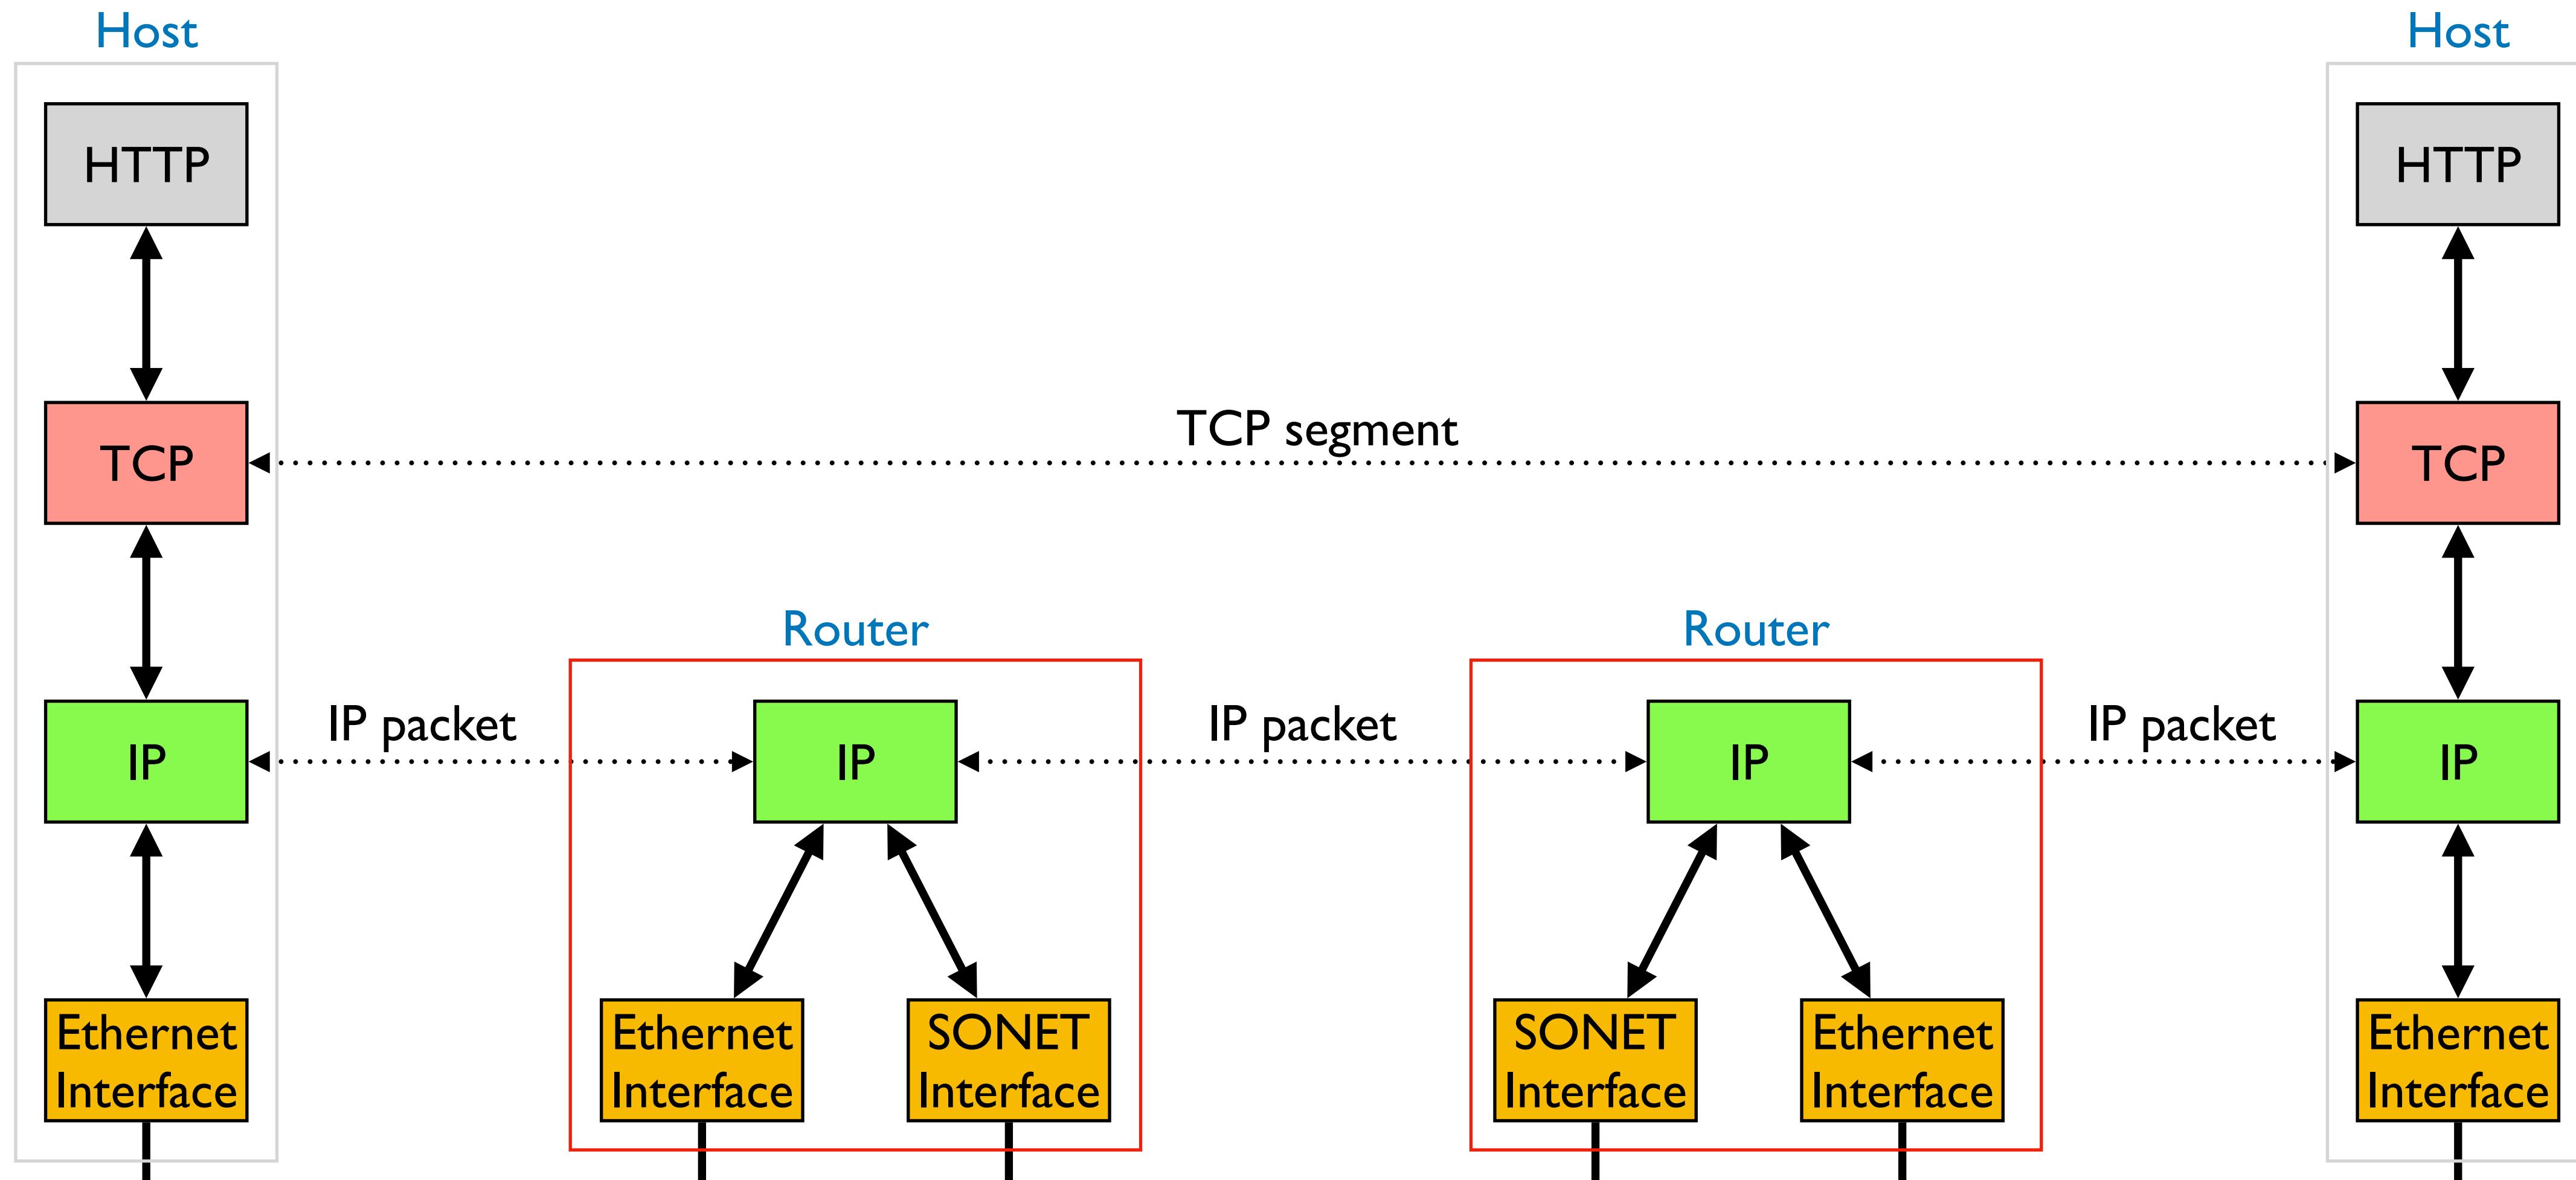
# A Protocol-centric Diagram



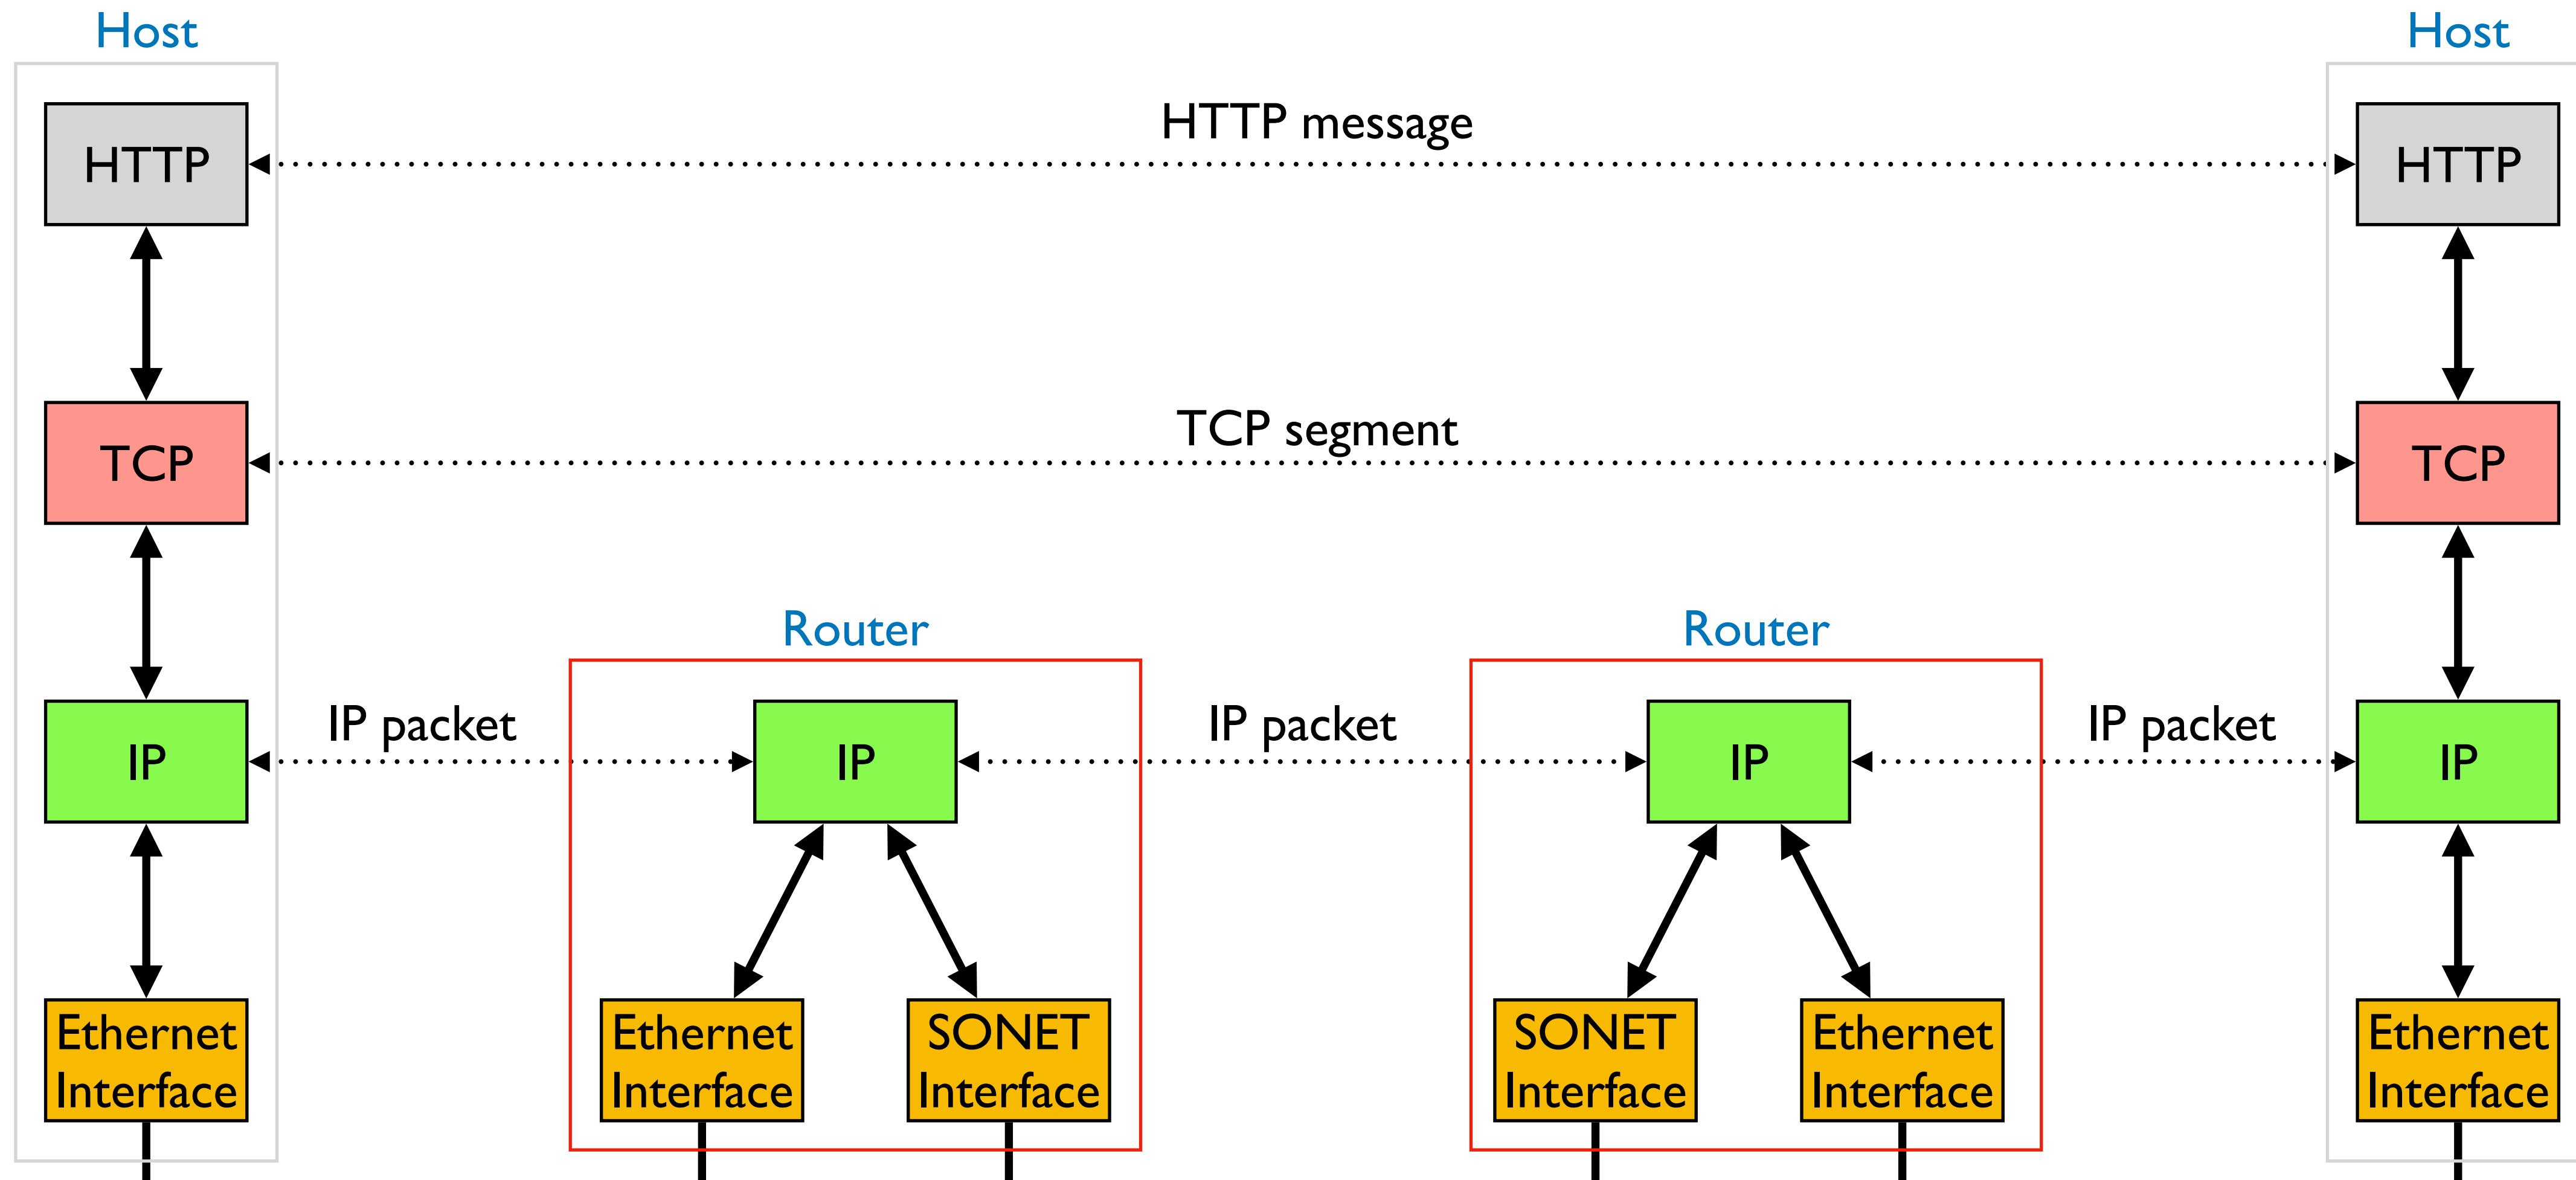
# A Protocol-centric Diagram



# A Protocol-centric Diagram



# A Protocol-centric Diagram



# Questions?

# Three steps

- **Decomposition**
- **Organization**
- **Assignment**

# Three steps

- **Decomposition**
- **Organization**

- **Assignment**

End-to-end principle

# Placing Network Functionality

# Placing Network Functionality

- Hugely influential paper: **“End-to-end Arguments in System Design” by Saltzer, Reed, and Clark (’84)**
  - Articulated the “End-to-end Principle” (E2E)

# Placing Network Functionality

- Hugely influential paper: **“End-to-end Arguments in System Design” by Saltzer, Reed, and Clark (’84)**
  - Articulated the “End-to-end Principle” (*E2E*)
  - Endless debate over what it means

# Placing Network Functionality

- Hugely influential paper: **“End-to-end Arguments in System Design” by Saltzer, Reed, and Clark (’84)**
  - Articulated the “End-to-end Principle” (E2E)
  - Endless debate over what it means
  - Many have used their own interpretation to support their position

# Placing Network Functionality

- Hugely influential paper: **“End-to-end Arguments in System Design” by Saltzer, Reed, and Clark (’84)**
  - Articulated the “End-to-end Principle” (E2E)
  - Endless debate over what it means
  - Many have used their own interpretation to support their position
  - But core argument holds value, esp. in understanding its impact on Internet design

# The Statement

# The Statement

- **If** a function **completely and correctly be implemented only with the knowledge and help of** the applications standing at the **end points** of the communication system,

# The Statement

- **If** a function **completely and correctly be implemented only with the knowledge and help of** the applications standing at the **end points** of the communication system,
- **Then,** providing that function as a **feature of the communication system** itself is not possible.

# The Statement

- **If** a function **completely and correctly be implemented only with the knowledge and help of** the applications standing at the **end points** of the communication system,
- **Then**, providing that function as a **feature of the communication system** itself is not possible.
- **Sometimes** an **incomplete version of the function** provided by the communication system **may be useful as a performance enhancement**

# Deconstructing the Statement

# Deconstructing the Statement

- Some application requirements can only be correctly implemented **at the end-system**

# Deconstructing the Statement

- Some application requirements can only be correctly implemented **at the end-system**
  - *Reliability, security, etc.*

# Deconstructing the Statement

- Some application requirements can only be correctly implemented **at the end-system**
  - *Reliability, security, etc.*
- Implementing these completely in the network may not only be hard, but impossible

# Deconstructing the Statement

- Some application requirements can only be correctly implemented **at the end-system**
  - *Reliability, security, etc.*
- Implementing these completely in the network may not only be hard, but impossible
- End-systems:

# Deconstructing the Statement

- Some application requirements can only be correctly implemented **at the end-system**
  - *Reliability, security, etc.*
  - Implementing these completely in the network may not only be hard, but impossible
  - End-systems:
    - **Can** satisfy the requirement without network's help

# Deconstructing the Statement

- Some application requirements can only be correctly implemented **at the end-system**
  - *Reliability, security, etc.*
  - Implementing these completely in the network may not only be hard, but impossible
  - End-systems:
    - **Can** satisfy the requirement without network's help
    - **Will/must** do so, since they cannot rely on the network

# Example: Reliable File Transfer



# Example: Reliable File Transfer



- **Solution I:** Make each step reliable, and string them together to make the end-to-end process reliable

# Example: Reliable File Transfer



- **Solution I:** Make each step reliable, and string them together to make the end-to-end process reliable

# Example: Reliable File Transfer



- **Solution I:** Make each step reliable, and string them together to make the end-to-end process reliable

# Example: Reliable File Transfer



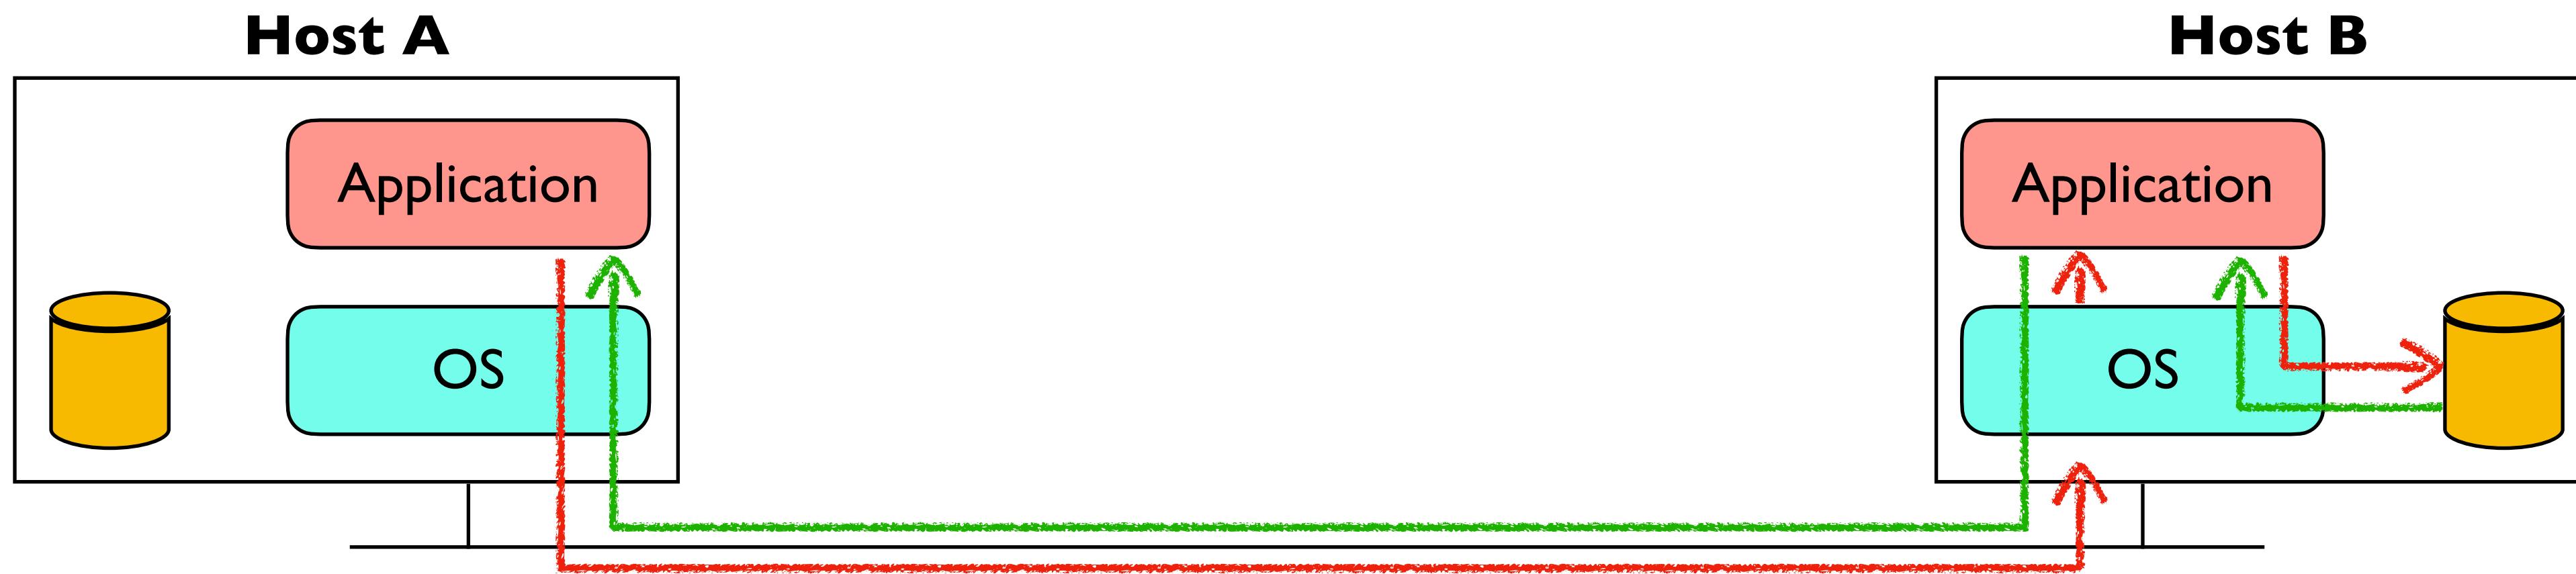
- **Solution I:** Make each step reliable, and string them together to make the end-to-end process reliable

# Example: Reliable File Transfer



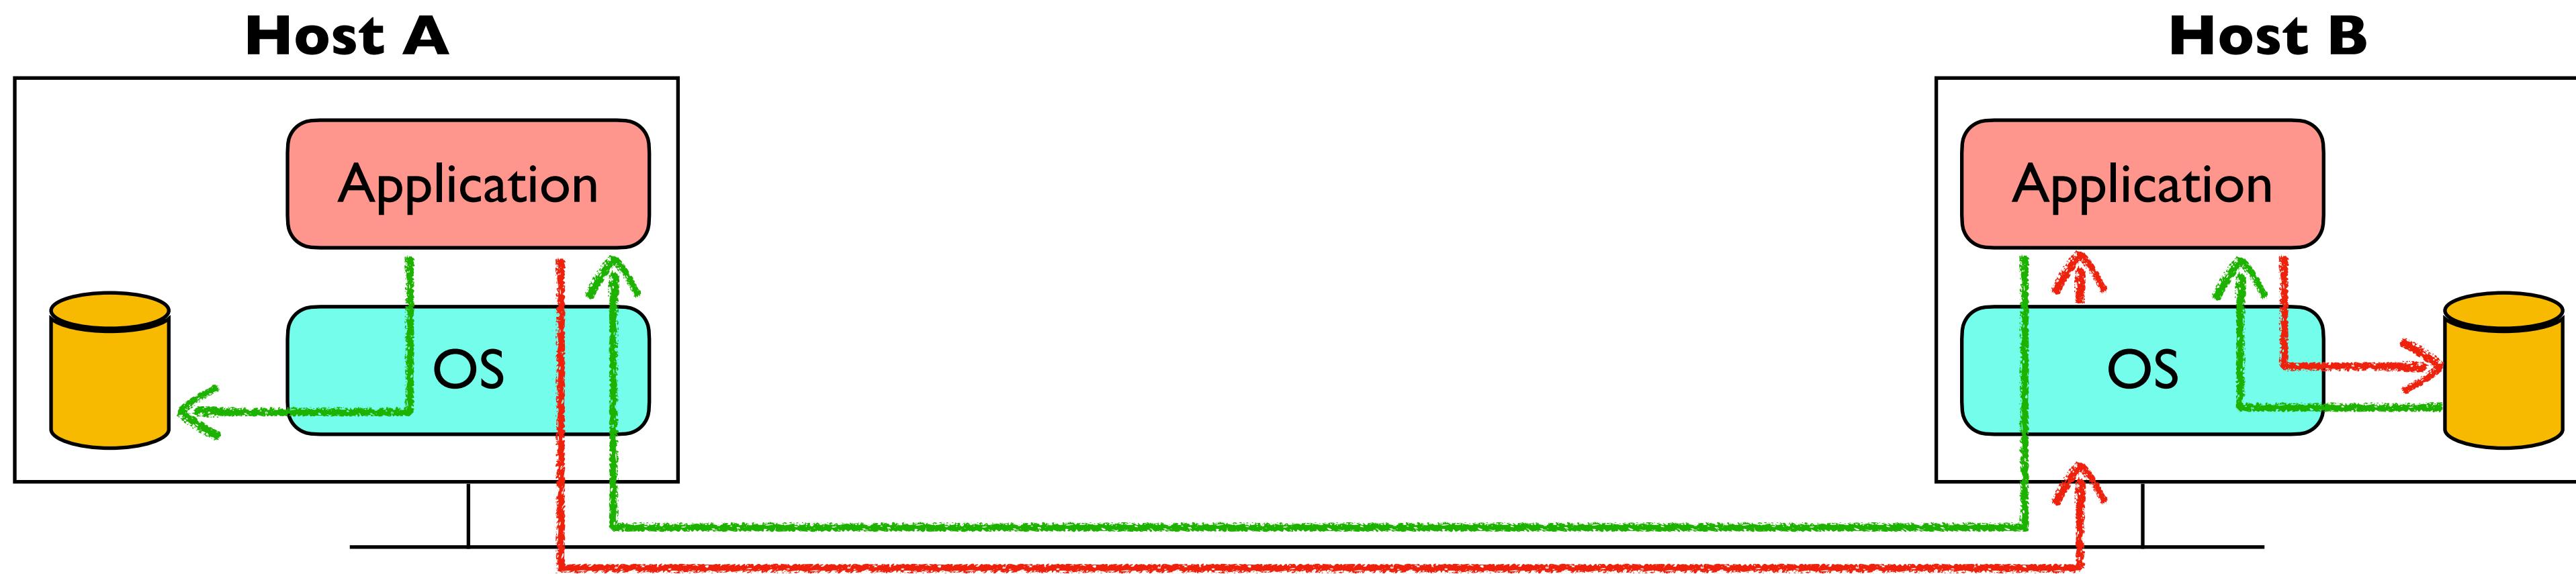
- **Solution I:** Make each step reliable, and string them together to make the end-to-end process reliable

# Example: Reliable File Transfer



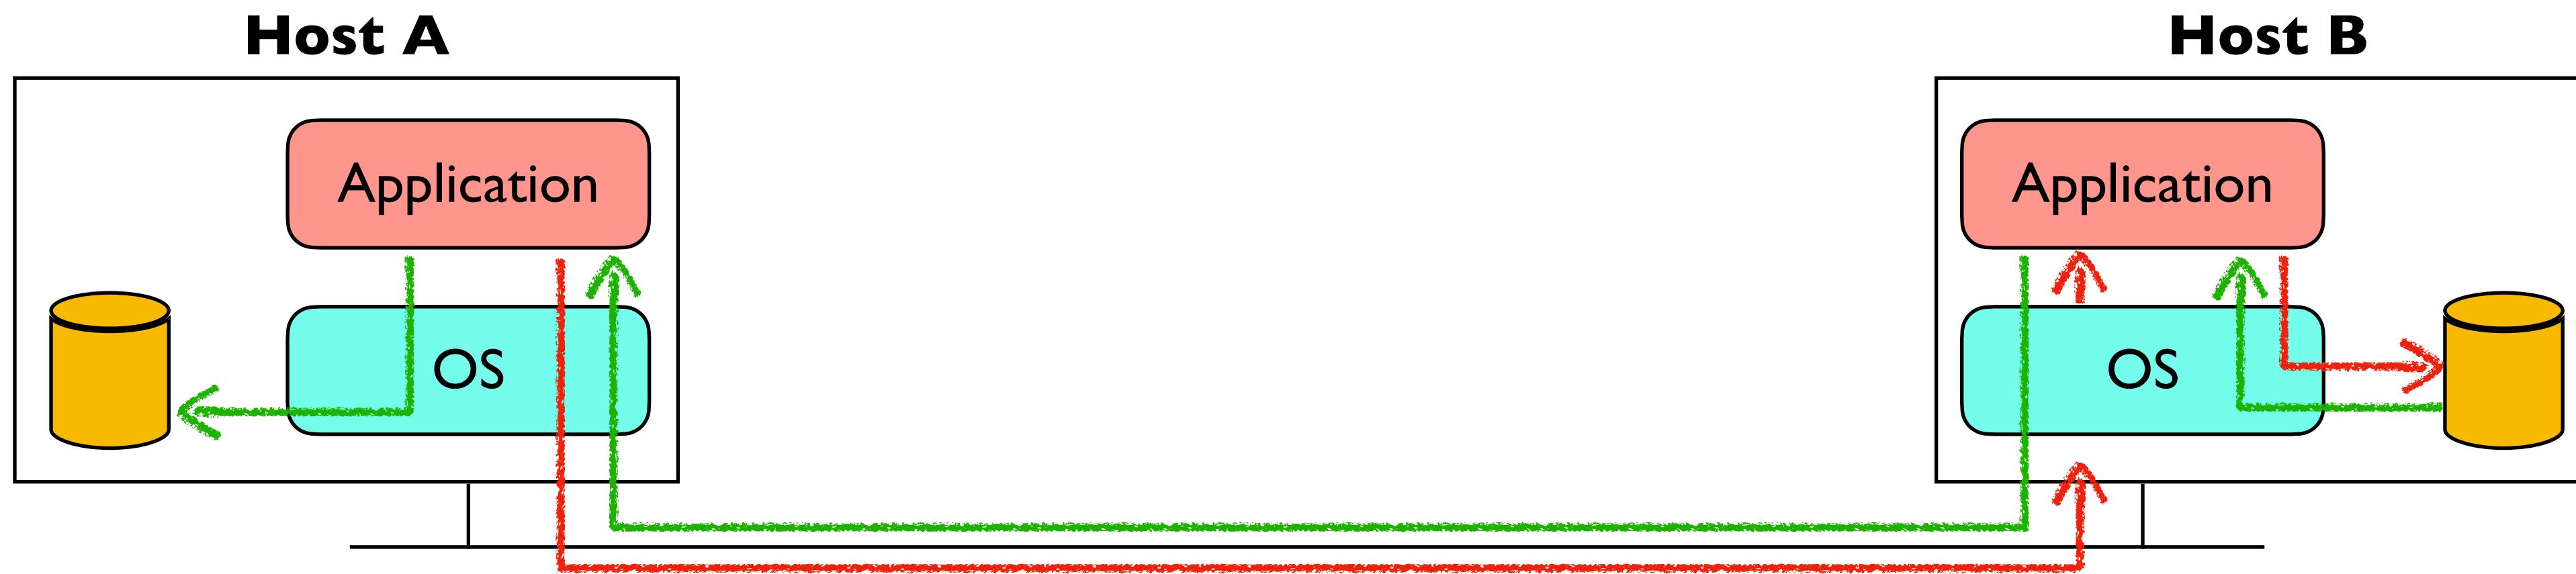
- **Solution I:** Make each step reliable, and string them together to make the end-to-end process reliable

# Example: Reliable File Transfer



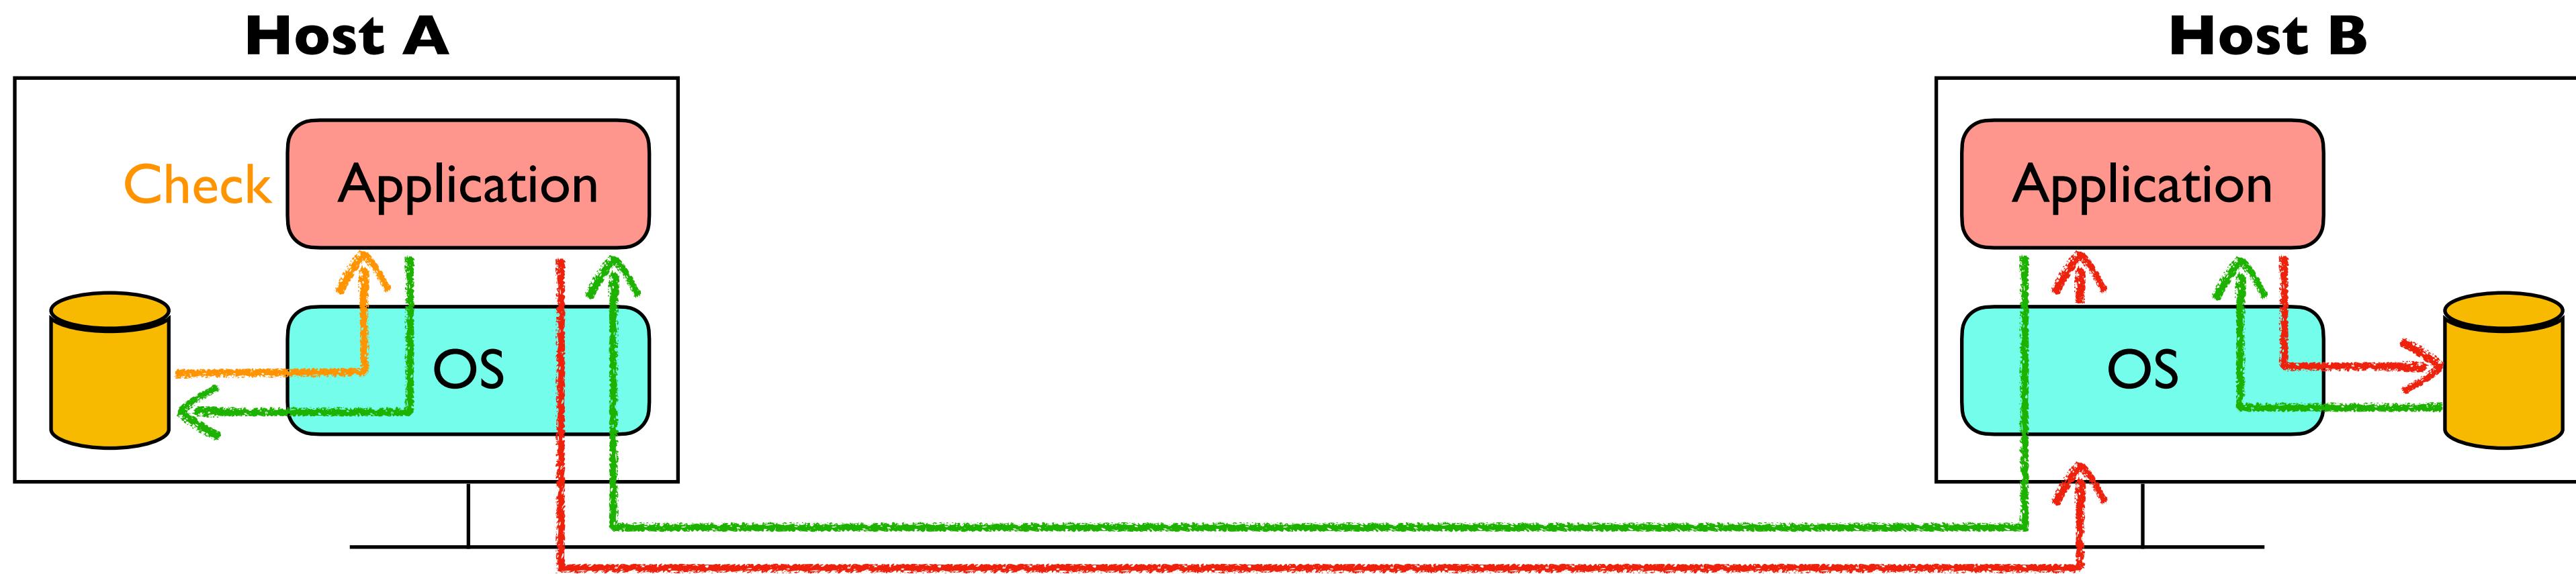
- **Solution I:** Make each step reliable, and string them together to make the end-to-end process reliable

# Example: Reliable File Transfer



- **Solution 1:** Make each step reliable, and string them together to make the end-to-end process reliable
- **Solution 2:** End-to-end check and retry

# Example: Reliable File Transfer



- **Solution 1:** Make each step reliable, and string them together to make the end-to-end process reliable
- **Solution 2:** End-to-end check and retry

# Discussion

# Discussion

- **Solution I** is incomplete
  - Receiver has to do the check anyway!

# Discussion

- **Solution 1** is incomplete
  - Receiver has to do the check anyway!
- **Solution 2** is complete
  - Full functionality can be entirely implemented at application layer with no need for reliability from lower layers

# Discussion

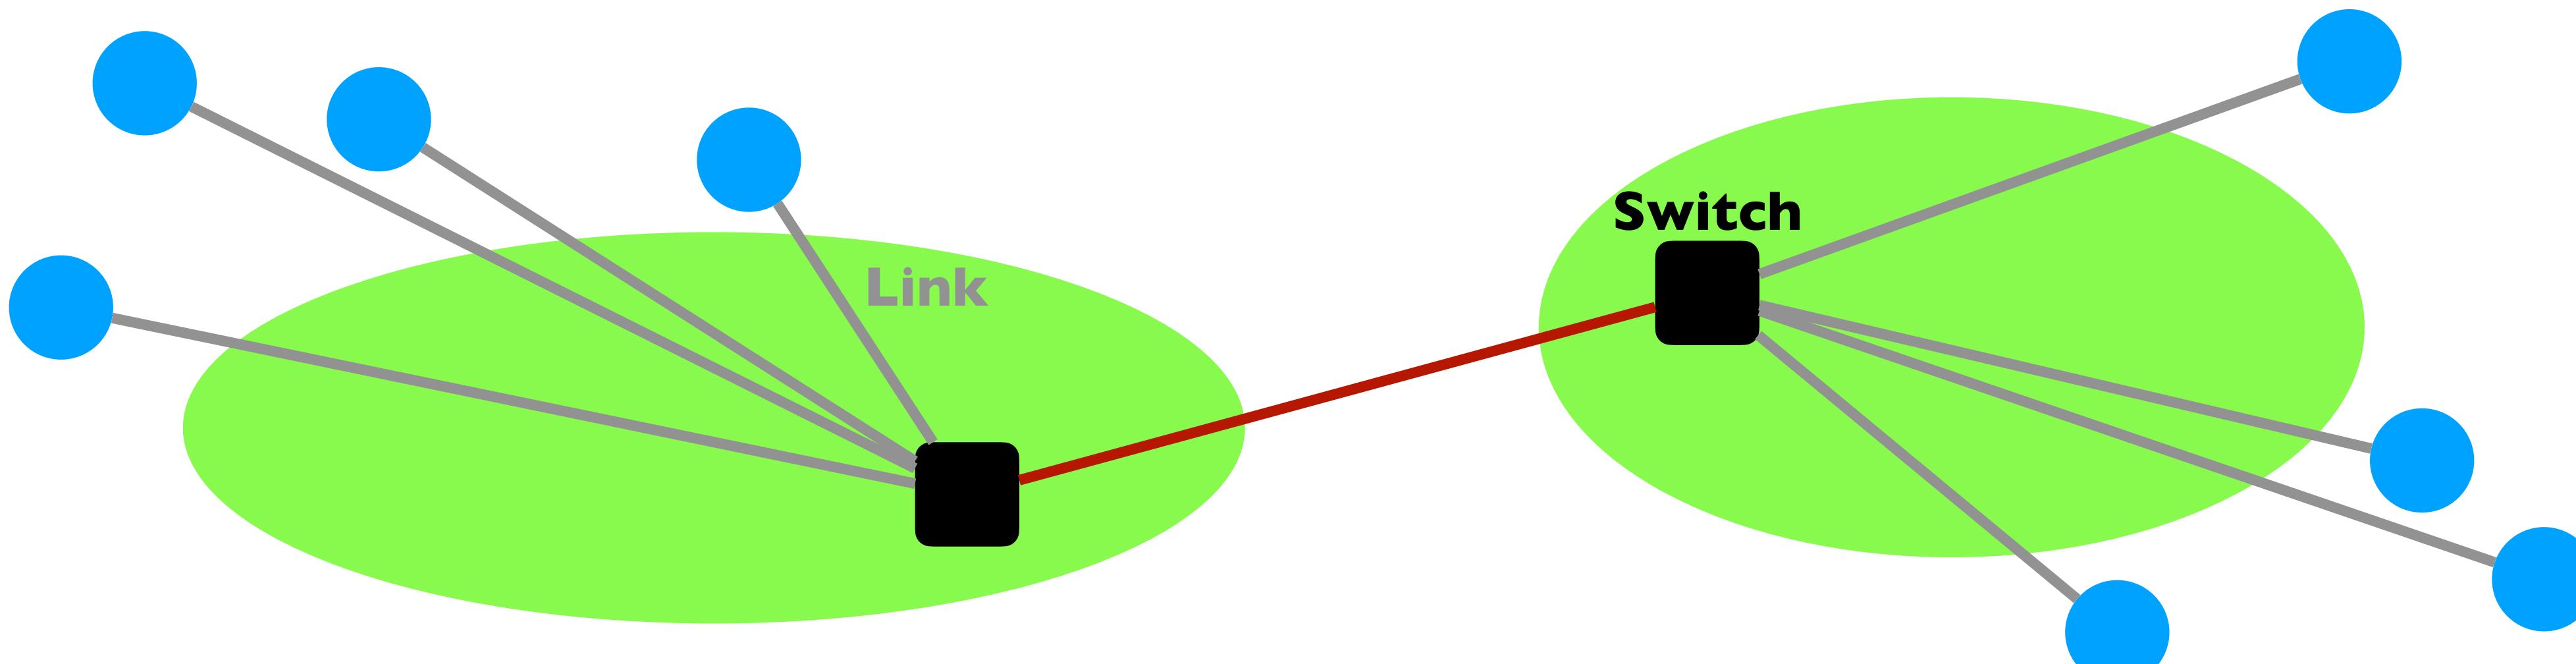
- **Solution 1** is incomplete
  - Receiver has to do the check anyway!
- **Solution 2** is complete
  - Full functionality can be entirely implemented at application layer with no need for reliability from lower layers
  - Is there any **need** to implement reliability at lower layers?

# Except...

**Sometimes** an **incomplete version of the function** provided by the communication system may be useful as a **performance enhancement**

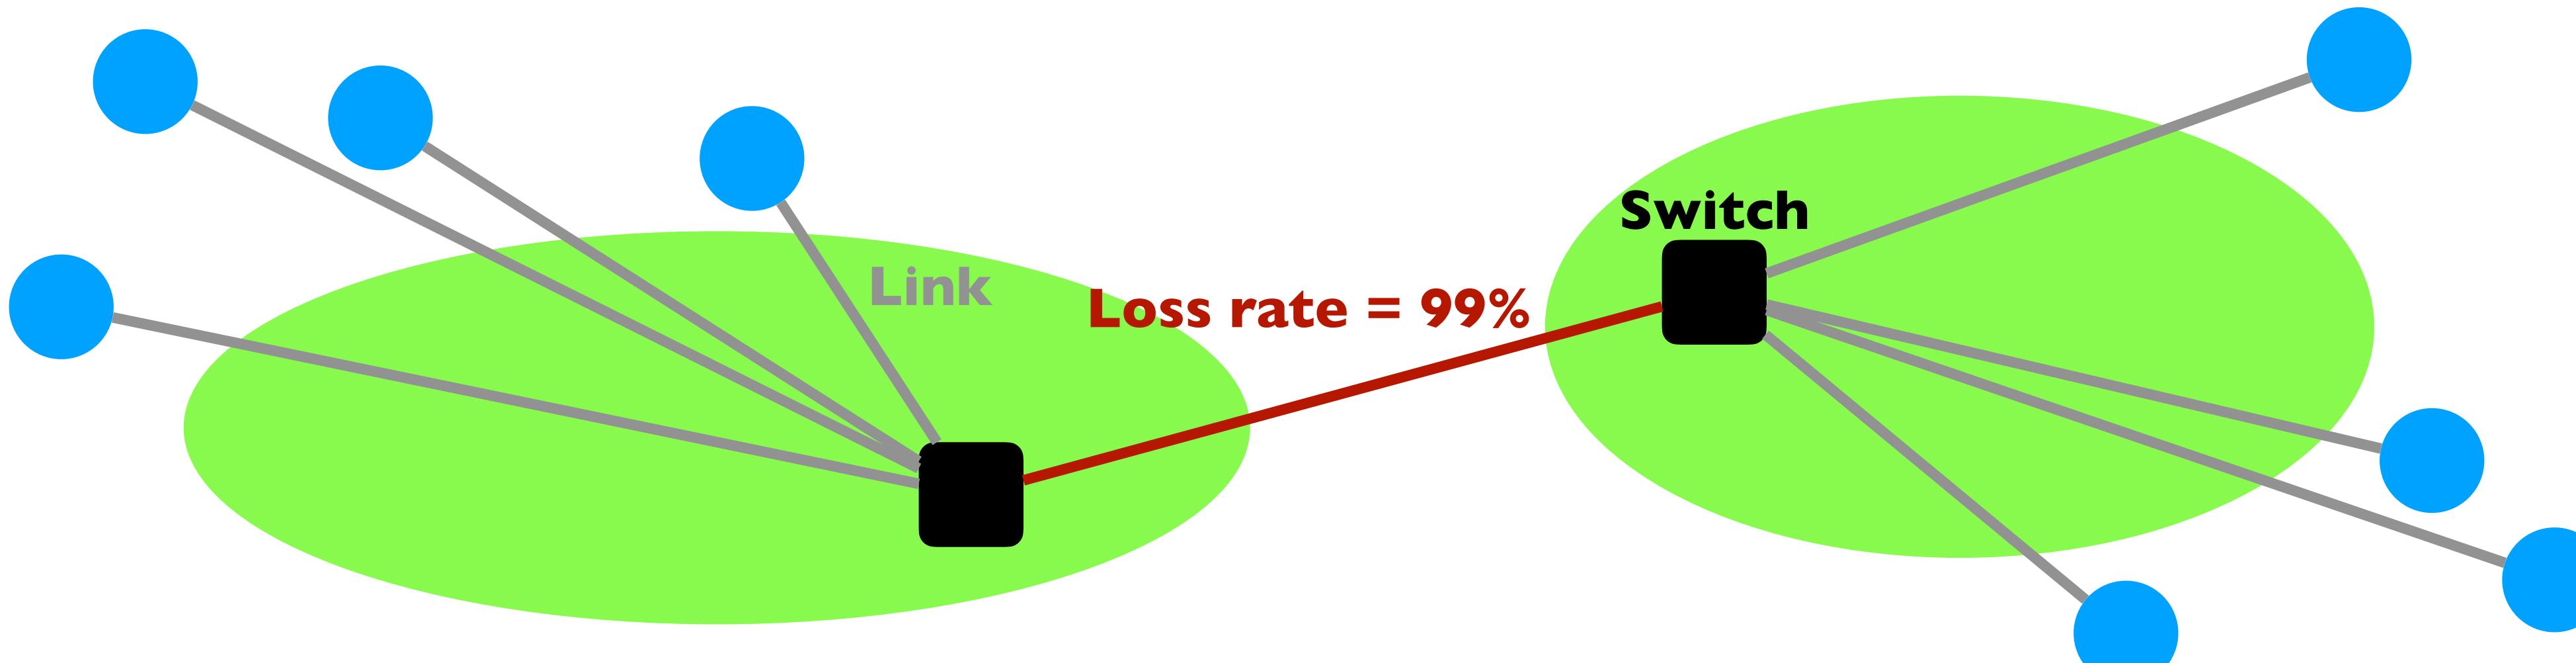
# Except...

**Sometimes** an **incomplete version of the function** provided by the communication system may be useful as a **performance enhancement**



# Except...

**Sometimes** an **incomplete version of the function** provided by the communication system may be useful as a **performance enhancement**



# Summary of the E2E principle

# Summary of the E2E principle

- Implementing functionality (e.g., reliability) in the network
  - Doesn't reduce host implementation complexity
  - Does increase network complexity
  - Probably increases delay and overhead on all applications even if they don't need the functionality (e.g., reliable data transfer vs. VoIP)

# Summary of the E2E principle

- Implementing functionality (e.g., reliability) in the network
  - Doesn't reduce host implementation complexity
  - Does increase network complexity
  - Probably increases delay and overhead on all applications even if they don't need the functionality (e.g., reliable data transfer vs. VoIP)
- However, implementing in the network can improve performance in some cases
  - E.g., consider a very lossy link

# Some consequences

# Some consequences

- In layered design, the E2E principle provides guidance on which layers are implemented where

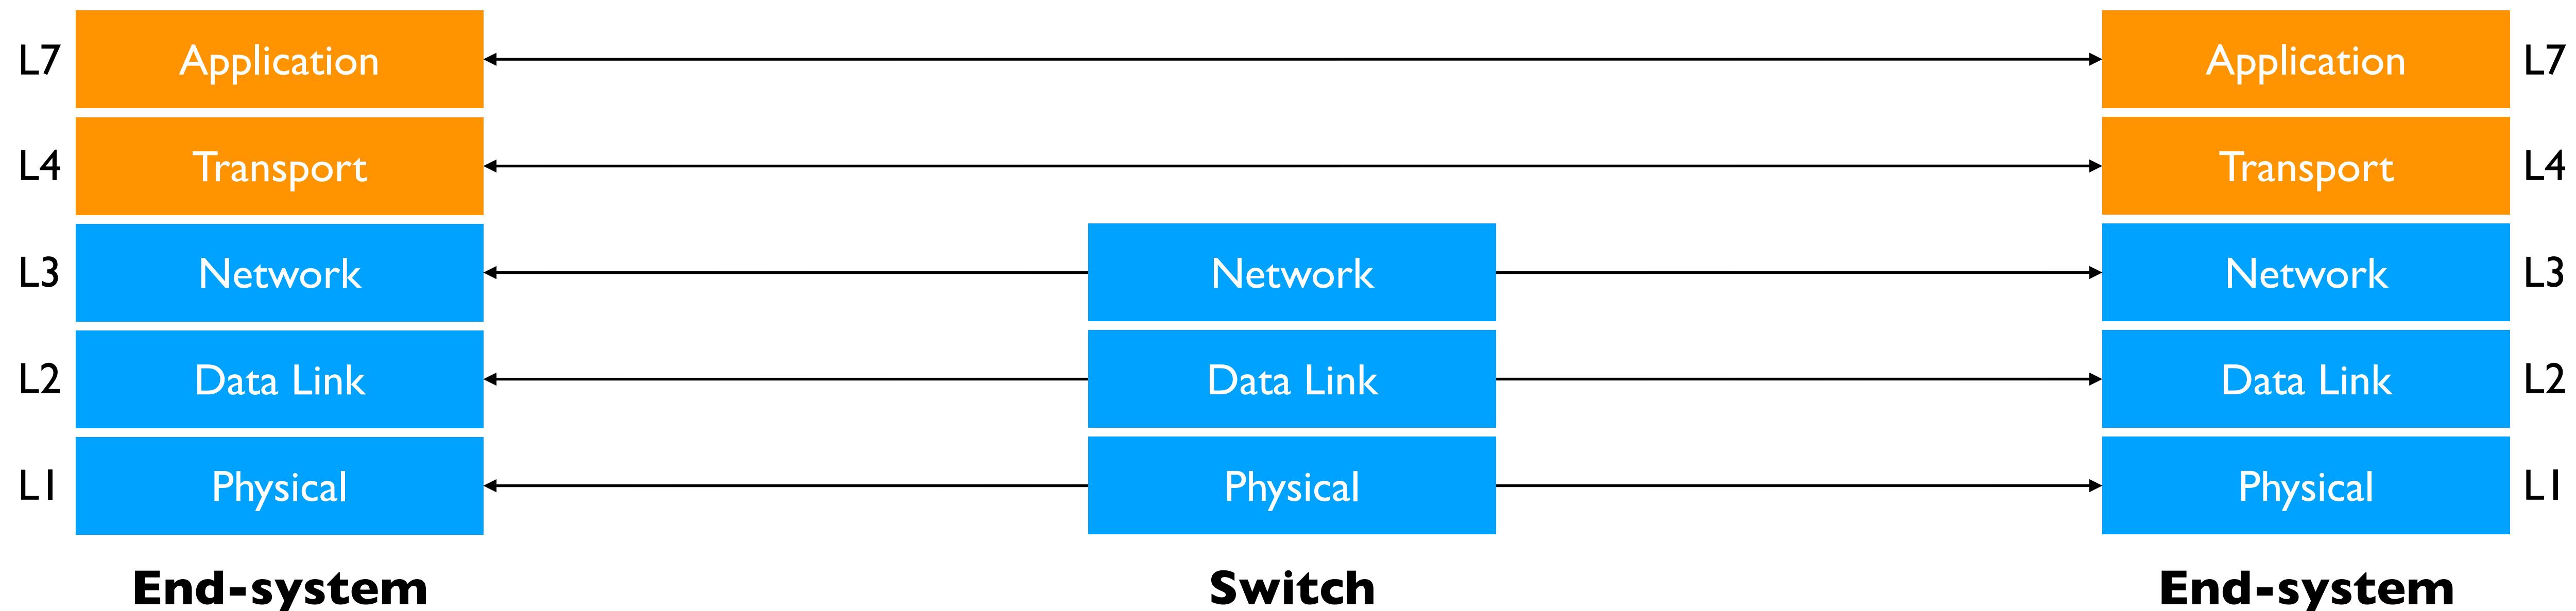
# Some consequences

- In layered design, the E2E principle provides guidance on which layers are implemented where
- “Dumb” network and “smart” end-systems
  - *Often credited as a key to the Internet’s success!*

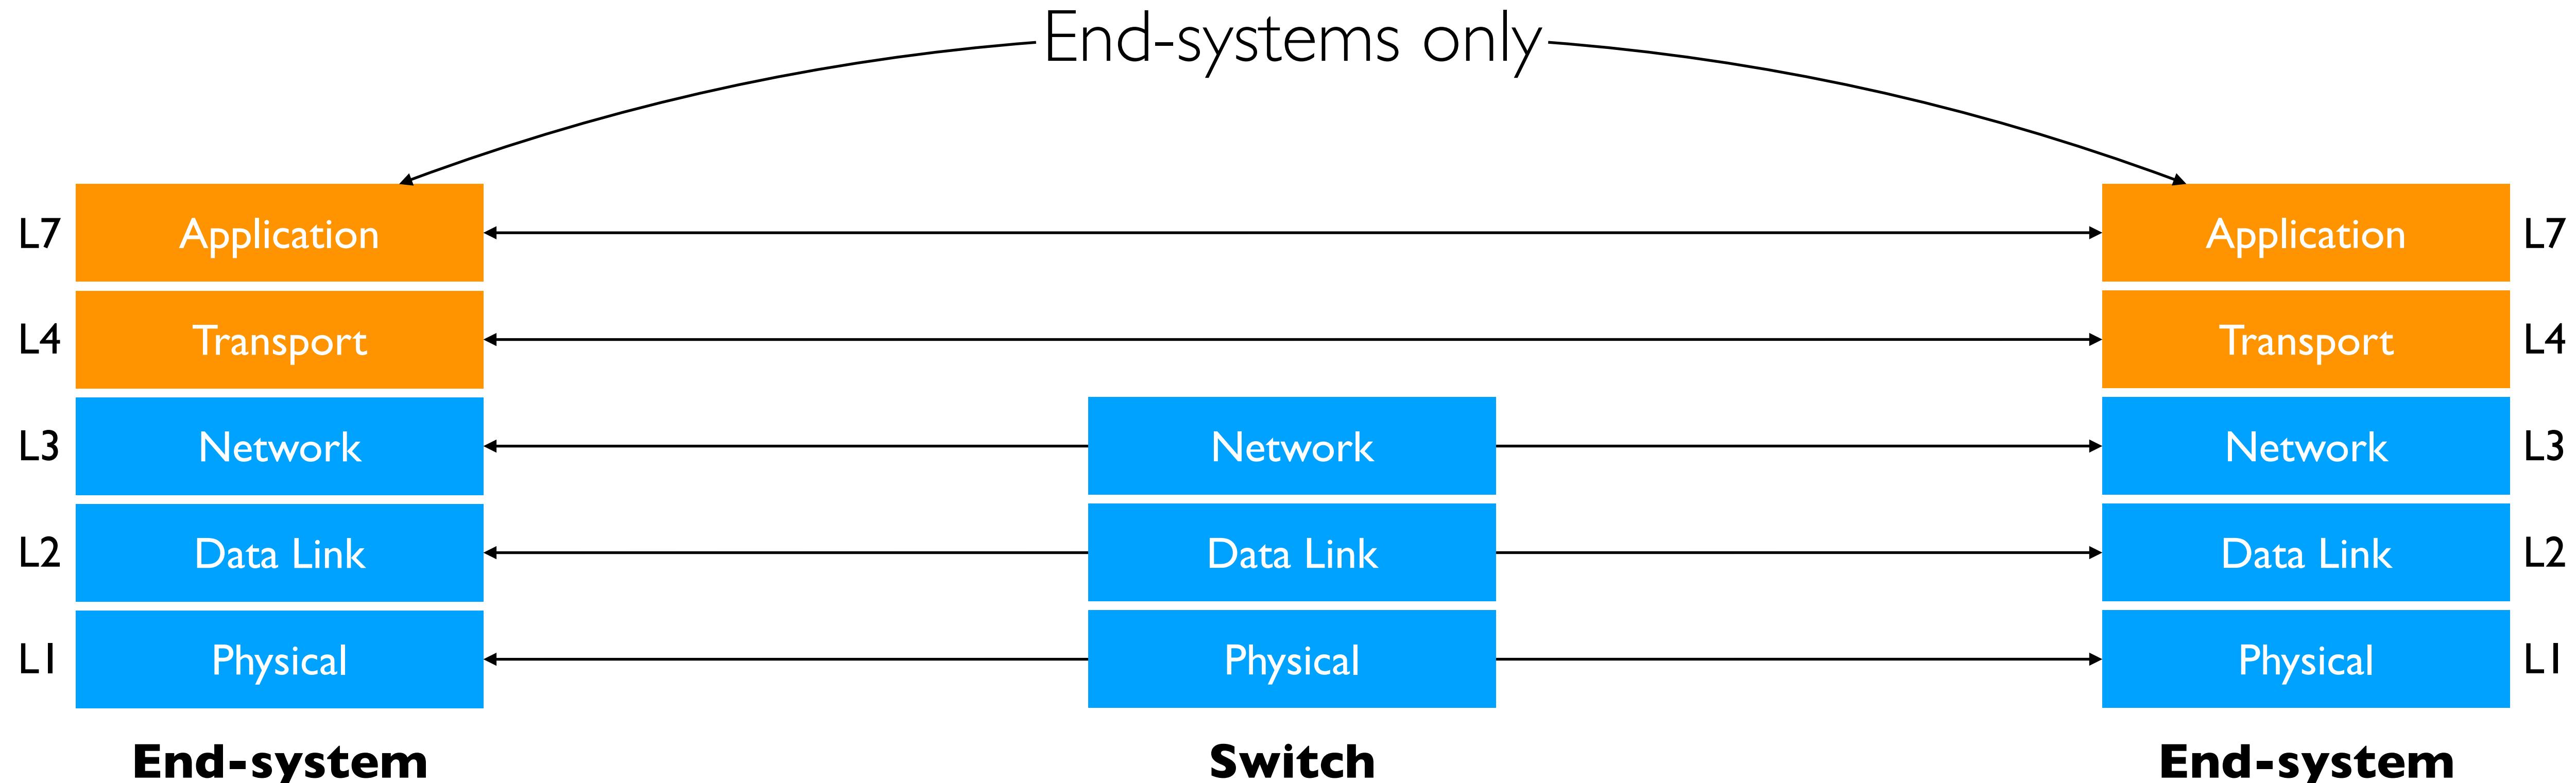
# Some consequences

- In layered design, the E2E principle provides guidance on which layers are implemented where
- “Dumb” network and “smart” end-systems
  - *Often credited as a key to the Internet’s success!*
- “Fate-sharing”
  - *Store state at the system entities that rely on the state*
  - *Often translated to keep end-host state out of routers*

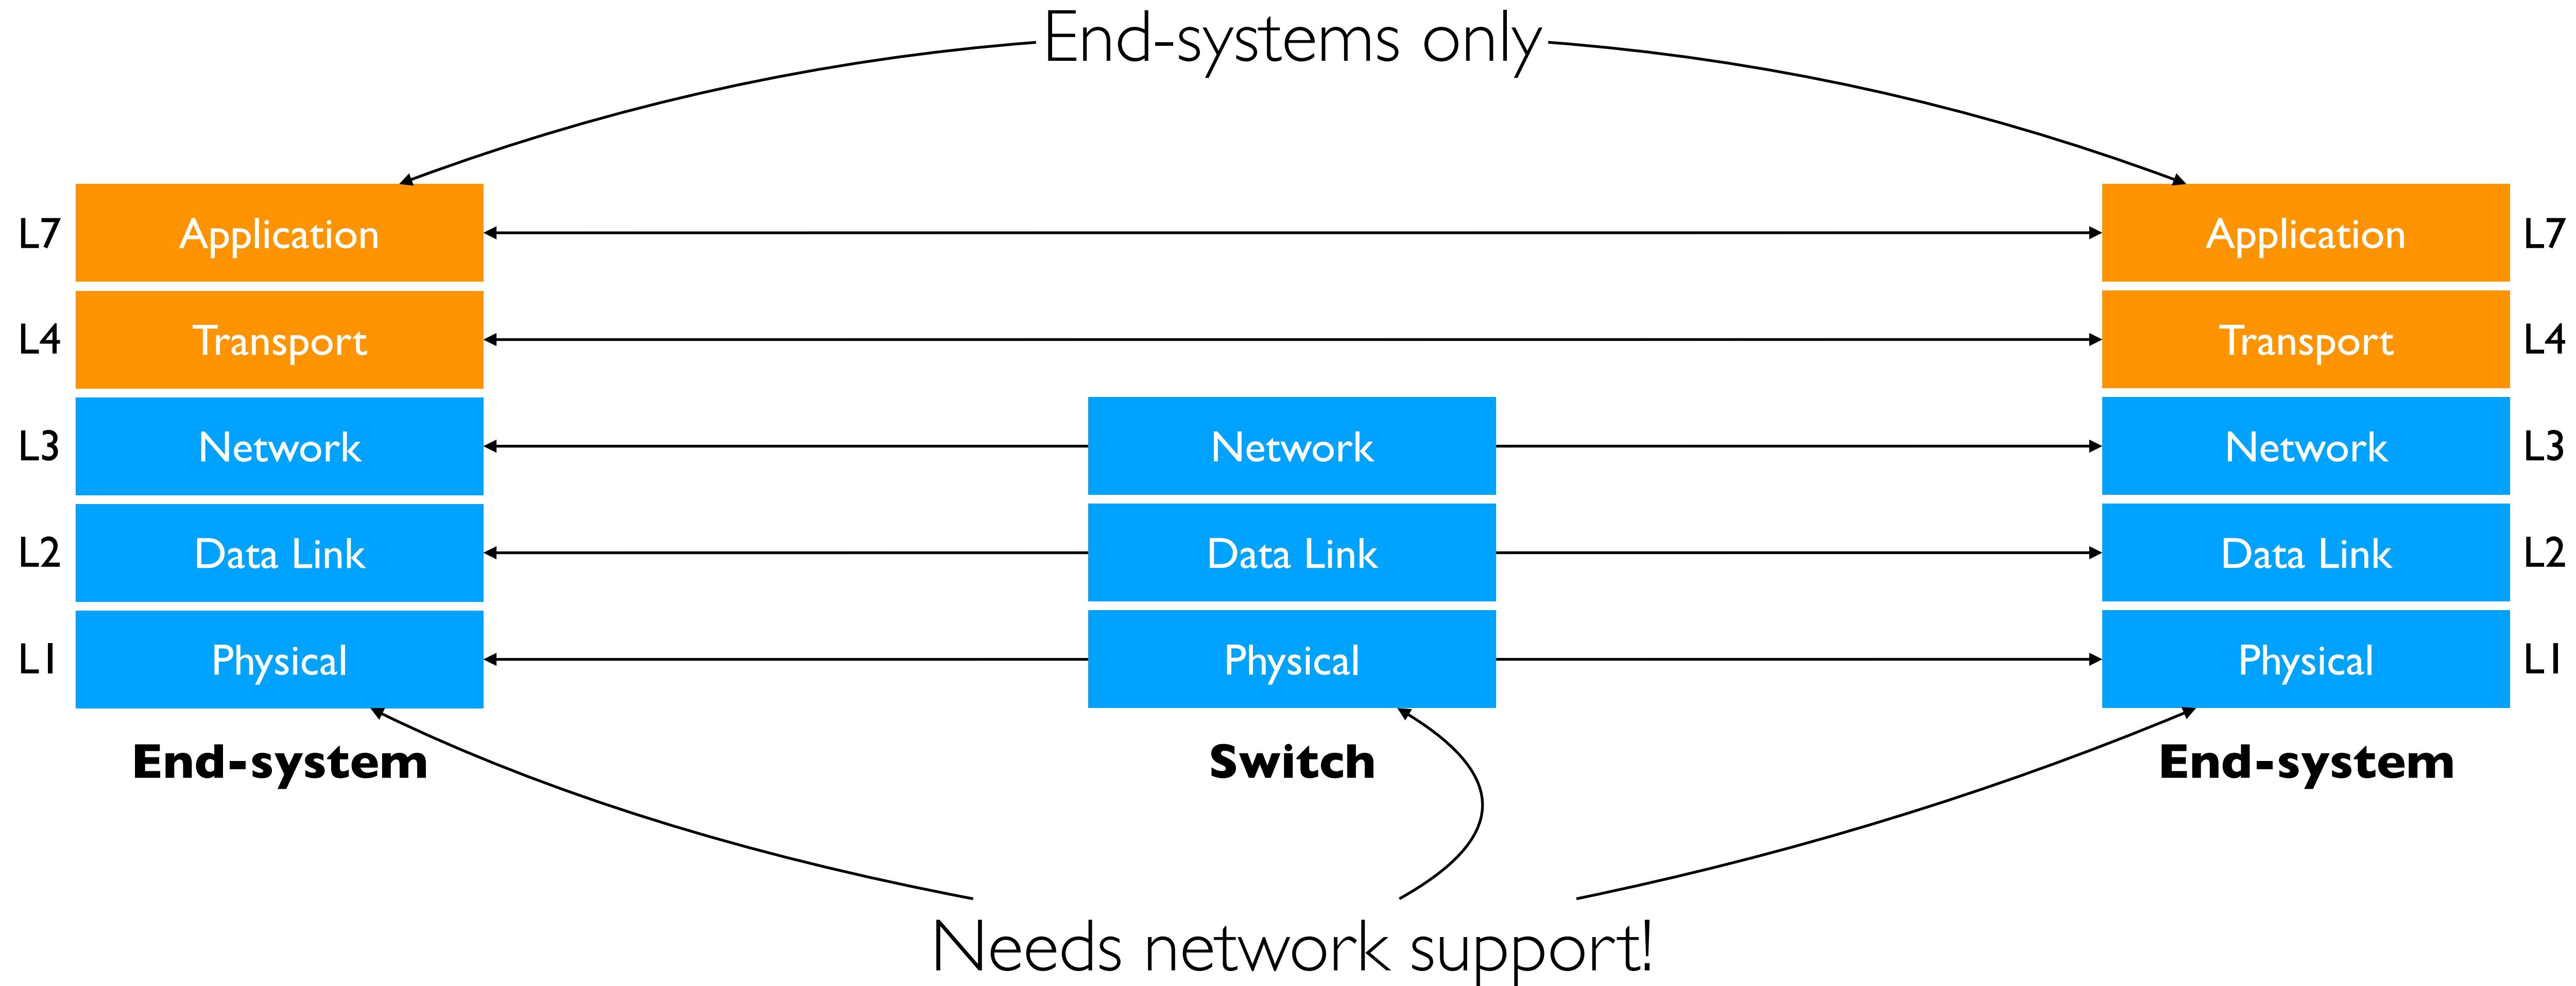
# E2E meets Layering



# E2E meets Layering



# E2E meets Layering



# Cracks in the E2E argument?

# Cracks in the E2E argument?

- Ignores incentives of different stakeholders

# Cracks in the E2E argument?

- Ignores incentives of different stakeholders
  - e.g., ISP looking for new revenue-generating services

# Cracks in the E2E argument?

- Ignores incentives of different stakeholders
  - e.g., ISP looking for new revenue-generating services
  - e.g., users looking for a performance boost

# Cracks in the E2E argument?

- Ignores incentives of different stakeholders
  - e.g., ISP looking for new revenue-generating services
  - e.g., users looking for a performance boost
- Sometimes we don't trust end-systems to do the job

# Cracks in the E2E argument?

- Ignores incentives of different stakeholders
  - e.g., ISP looking for new revenue-generating services
  - e.g., users looking for a performance boost
- Sometimes we don't trust end-systems to do the job
  - e.g., firewalls, intrusion detection systems

# Questions?

# Asking the right architectural questions...

# Asking the right architectural questions...

- How to break system into modules?

# Asking the right architectural questions...

- How to break system into modules?
  - *Classic decomposition into tasks*

# Asking the right architectural questions...

- How to break system into modules?
  - *Classic decomposition into tasks*
  - Where are modules implemented?

# Asking the right architectural questions...

- How to break system into modules?
  - *Classic decomposition into tasks*
- Where are modules implemented?
  - *Hosts? Routers? Both?*

# Asking the right architectural questions...

- How to break system into modules?
  - *Classic decomposition into tasks*
- Where are modules implemented?
  - Hosts? Routers? Both?
- Where is state stored?

# Asking the right architectural questions...

- How to break system into modules?
  - *Classic decomposition into tasks*
- Where are modules implemented?
  - *Hosts? Routers? Both?*
- Where is state stored?
  - *Hosts? Routers? Both?*

# ... leads to three design principles

- How to break system into modules?
  - *Layering*
- Where are modules implemented?
  - *End-to-end principle*
- Where is state stored?
  - *Fate-sharing*

# Taking Stock

# Taking Stock

- Layering is a good way to organize networks
  - *Unified IP layer decouples applications from networks*

# Taking Stock

- Layering is a good way to organize networks
  - *Unified IP layer decouples applications from networks*
- E2E argument encourages us to keep IP simple

# Taking Stock

- Layering is a good way to organize networks
  - *Unified IP layer decouples applications from networks*
- E2E argument encourages us to keep IP simple
- Fate-sharing principle keeps state out of routers

# Taking Stock: So Far...

# Taking Stock: So Far...

- What is a network made of?

# Taking Stock: So Far...

- What is a network made of?
- How is it shared?

# Taking Stock: So Far...

- What is a network made of?
- How is it shared?
- How do we evaluate a network?

# Taking Stock: So Far...

- What is a network made of?
- How is it shared?
- How do we evaluate a network?
- How is communication architected?

# Taking Stock

# Taking Stock

- **Basic concepts**

- Components: end-systems, links, switches, routers
- Performance: delay, loss, throughput

# Taking Stock

- **Basic concepts**

- Components: end-systems, links, switches, routers
- Performance: delay, loss, throughput

- **Basic design choices**

- Packets vs circuit switching
- Best-effort vs guaranteed service
- Layering
- “Dumb” network, “smart” ends
  - E2E principle, fate-sharing

# Coming up next...

- How, in detail, layer by layer (we will obsess a lot over “why” too...)



# Coming up next...

- How, in detail, layer by layer (we will obsess a lot over “why” too...)



# Coming up next...

- How, in detail, layer by layer (we will obsess a lot over “why” too...)



# Coming up next...

- How, in detail, layer by layer (we will obsess a lot over “why” too...)



# Coming up next...

- How, in detail, layer by layer (we will obsess a lot over “why” too...)



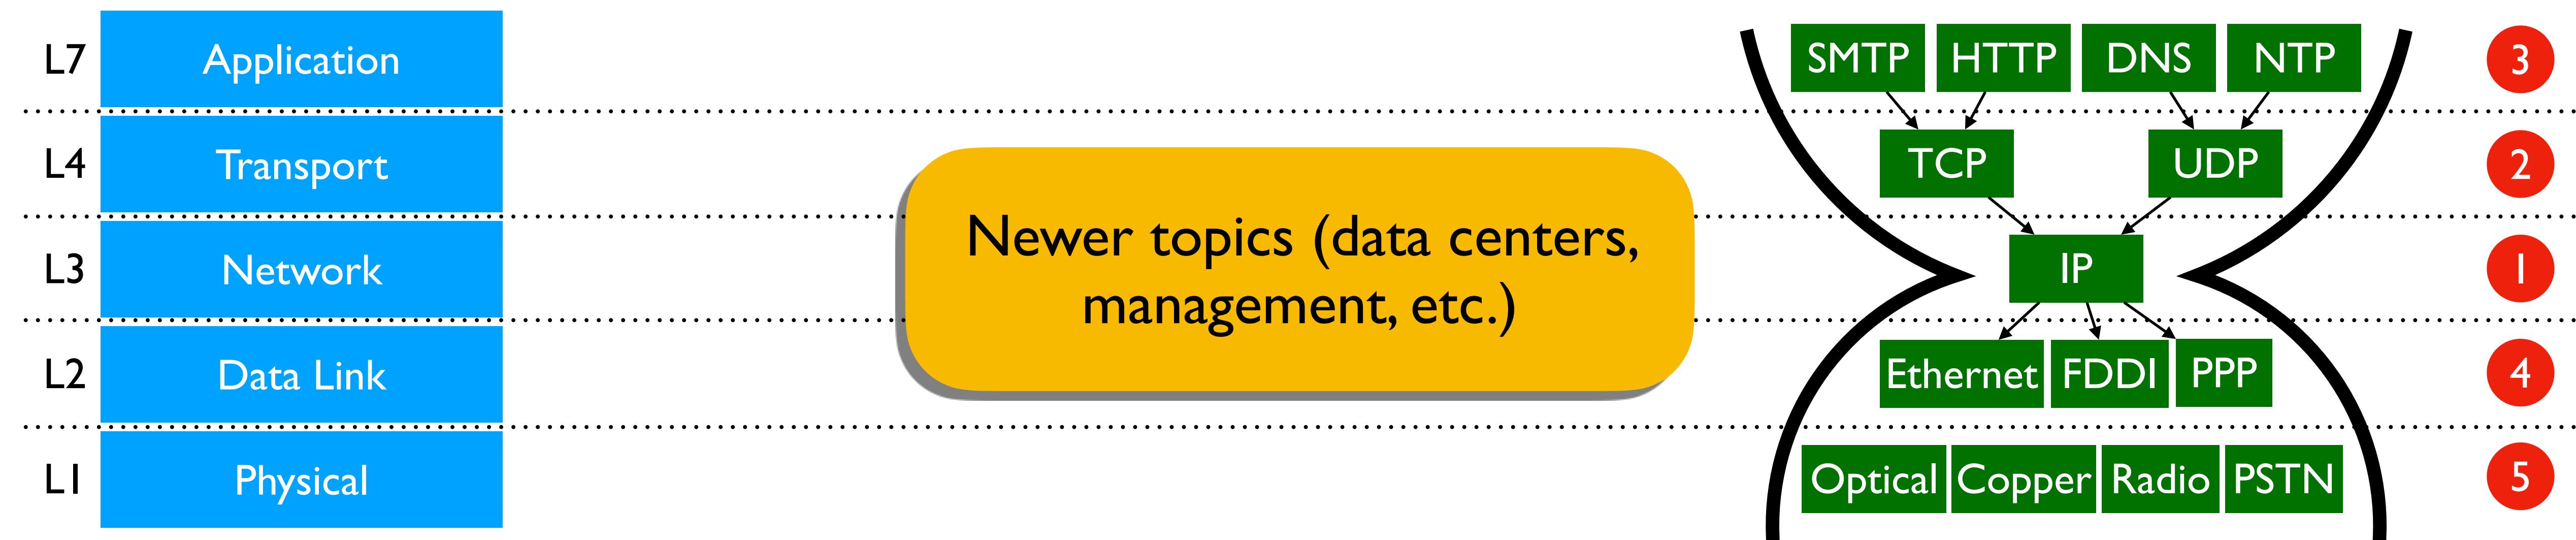
# Coming up next...

- How, in detail, layer by layer (we will obsess a lot over “why” too...)



# Coming up next...

- How, in detail, layer by layer (we will obsess a lot over “why” too...)



# The Network Layer

CPSC 433/533, Spring 2021

Anurag Khandelwal

# What does the Network Layer do?

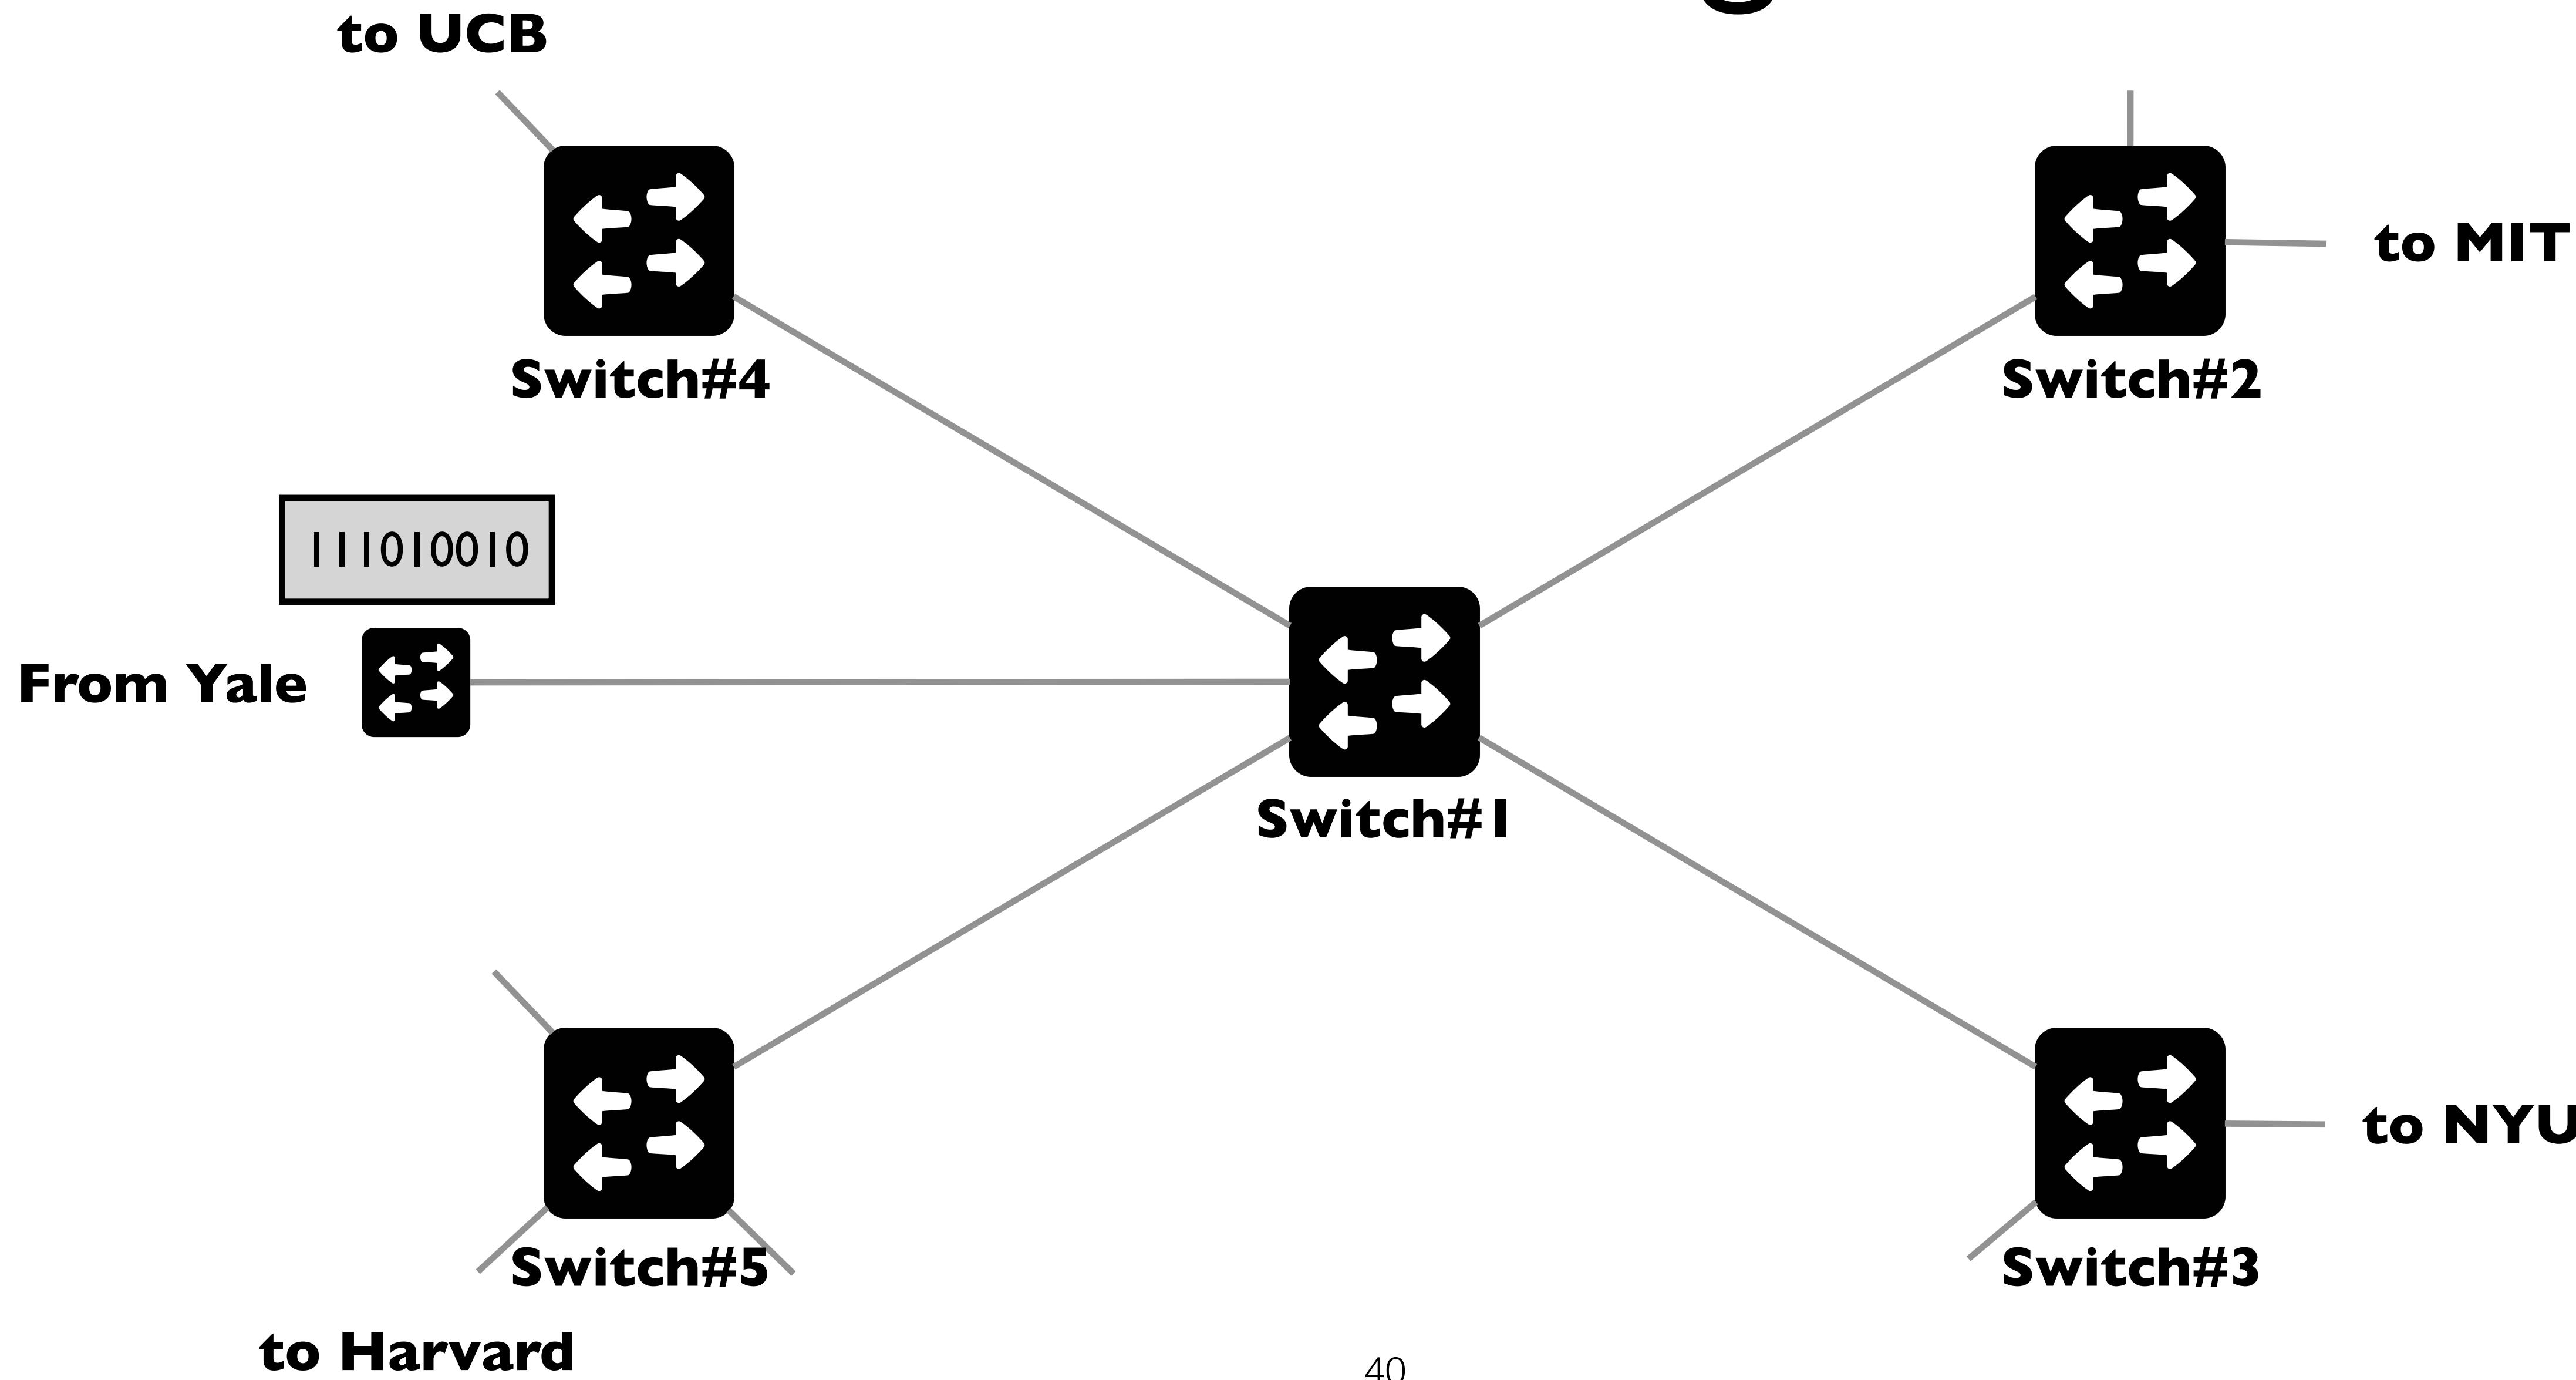
# What does the Network Layer do?

- ***Best-effort global delivery of packets***

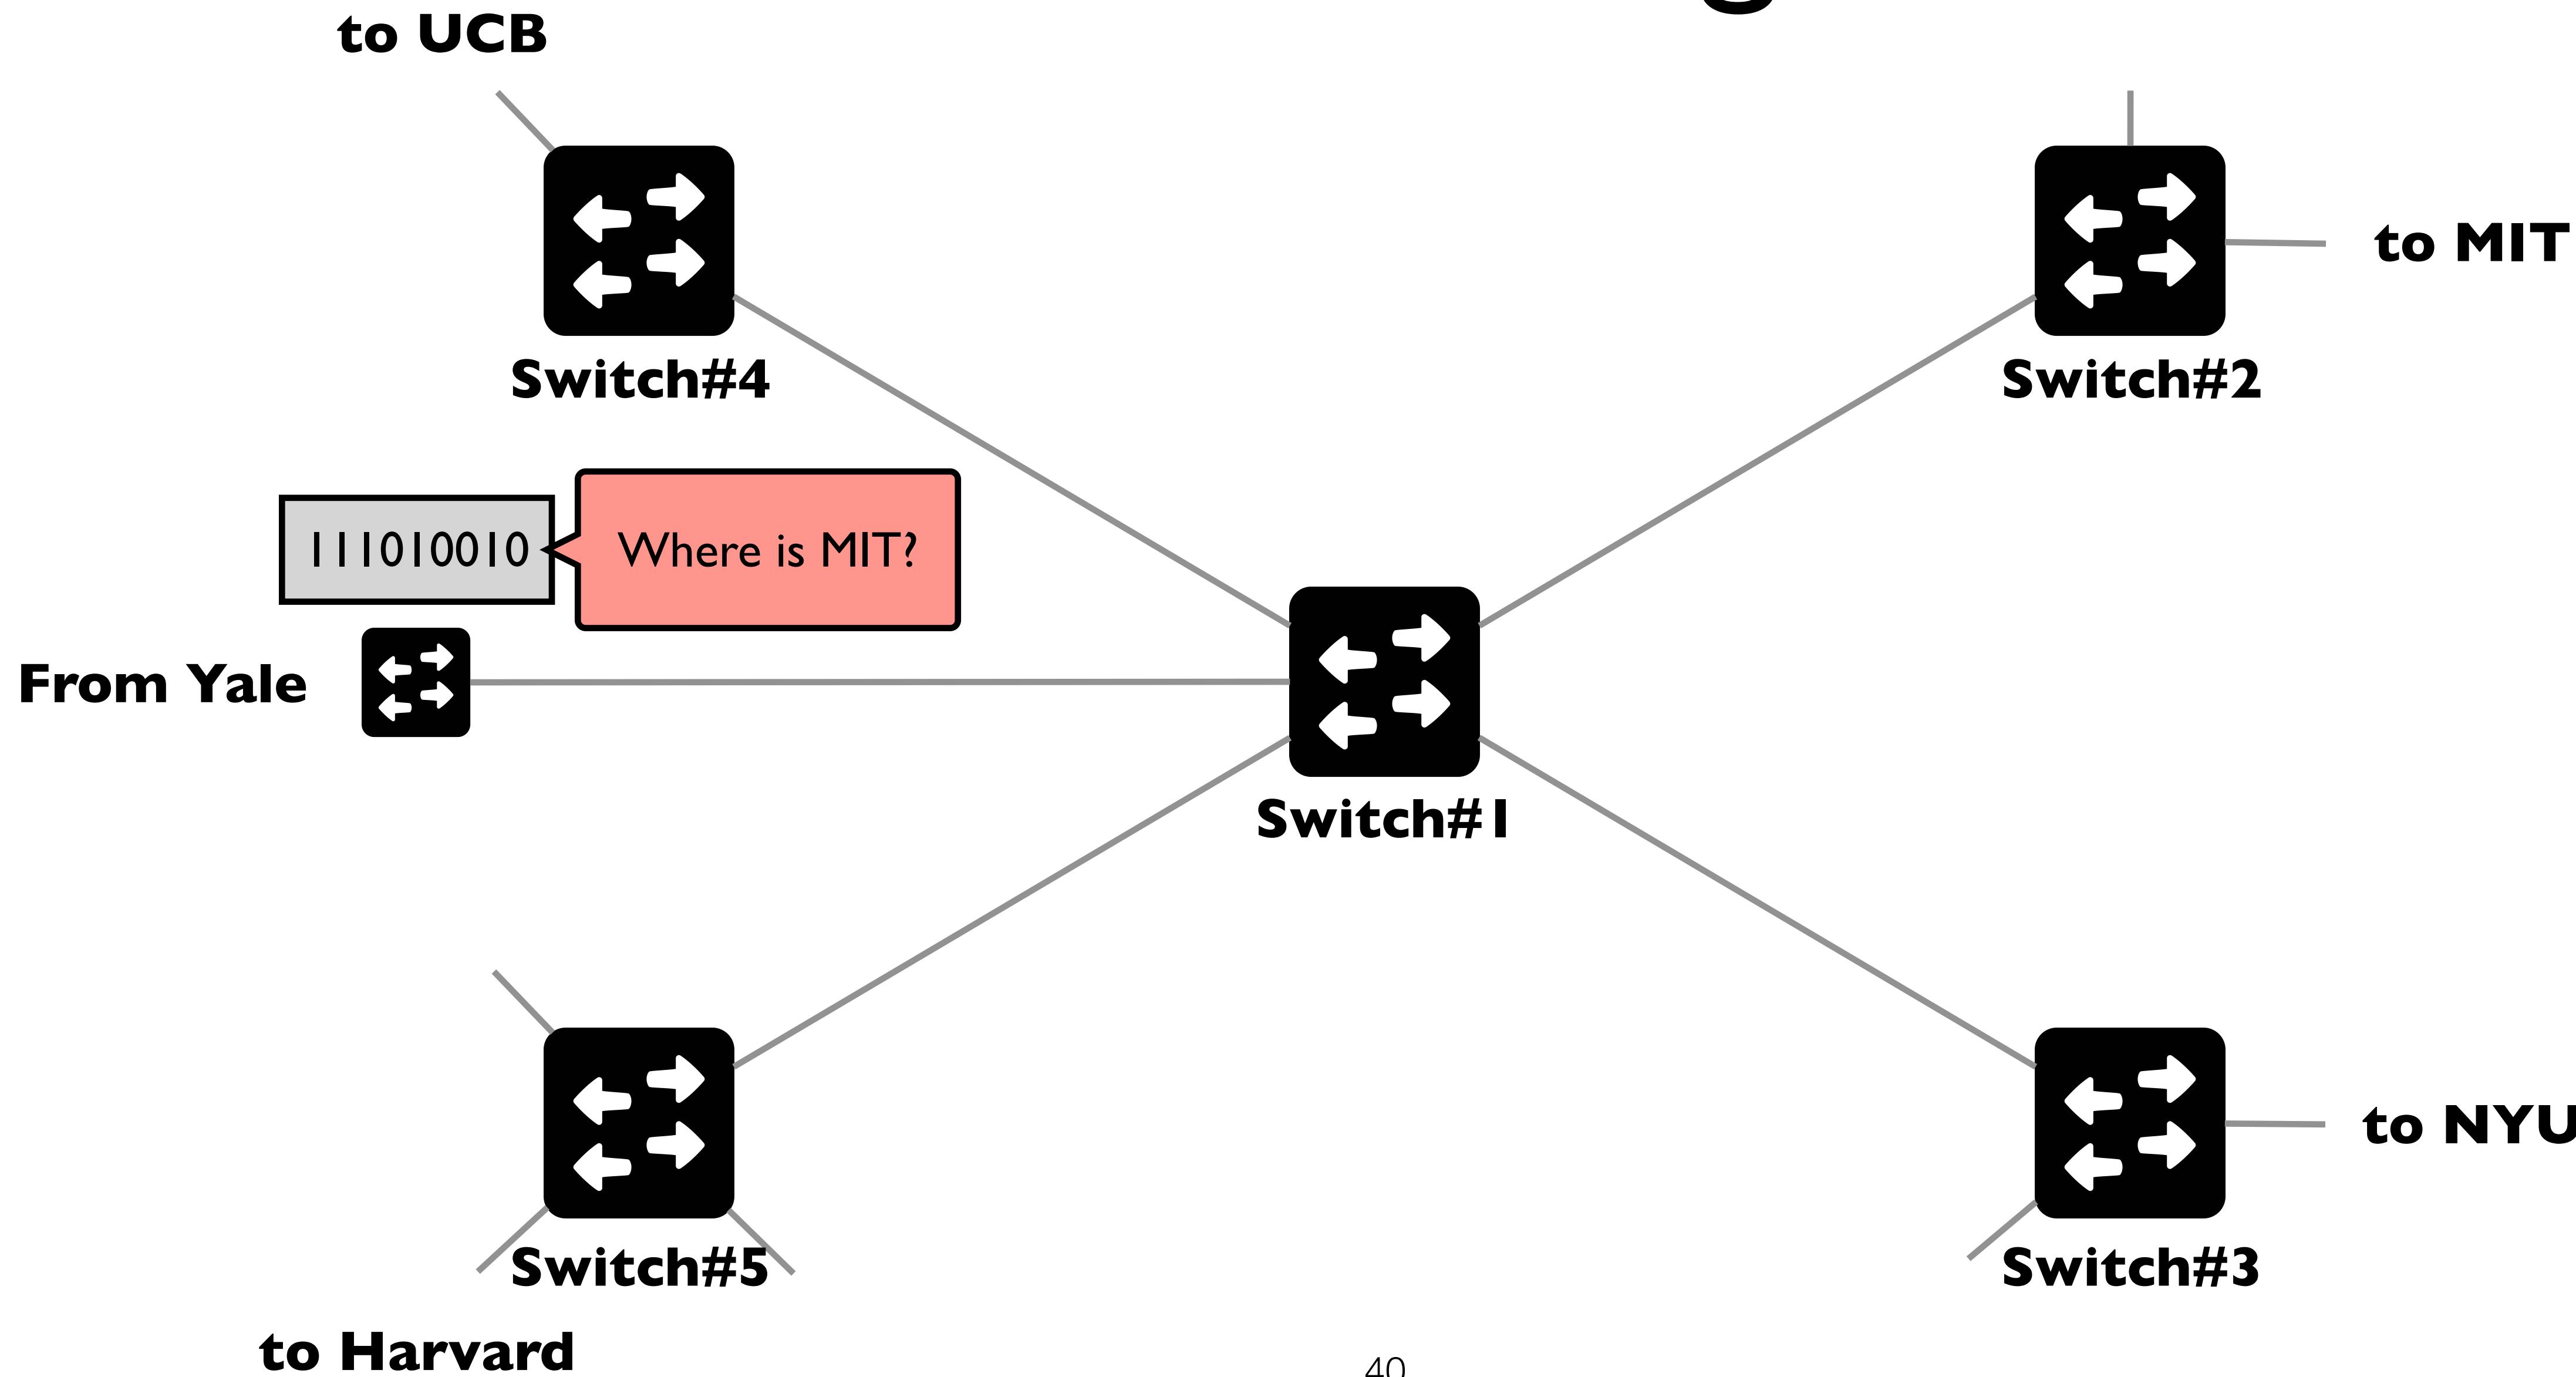
# What does the Network Layer do?

- ***Best-effort global delivery of packets***
- Ok, but what does that entail?

# Addressing



# Addressing



# Addressing (for now)

# Addressing (for now)

- Assume each end-host has a unique address

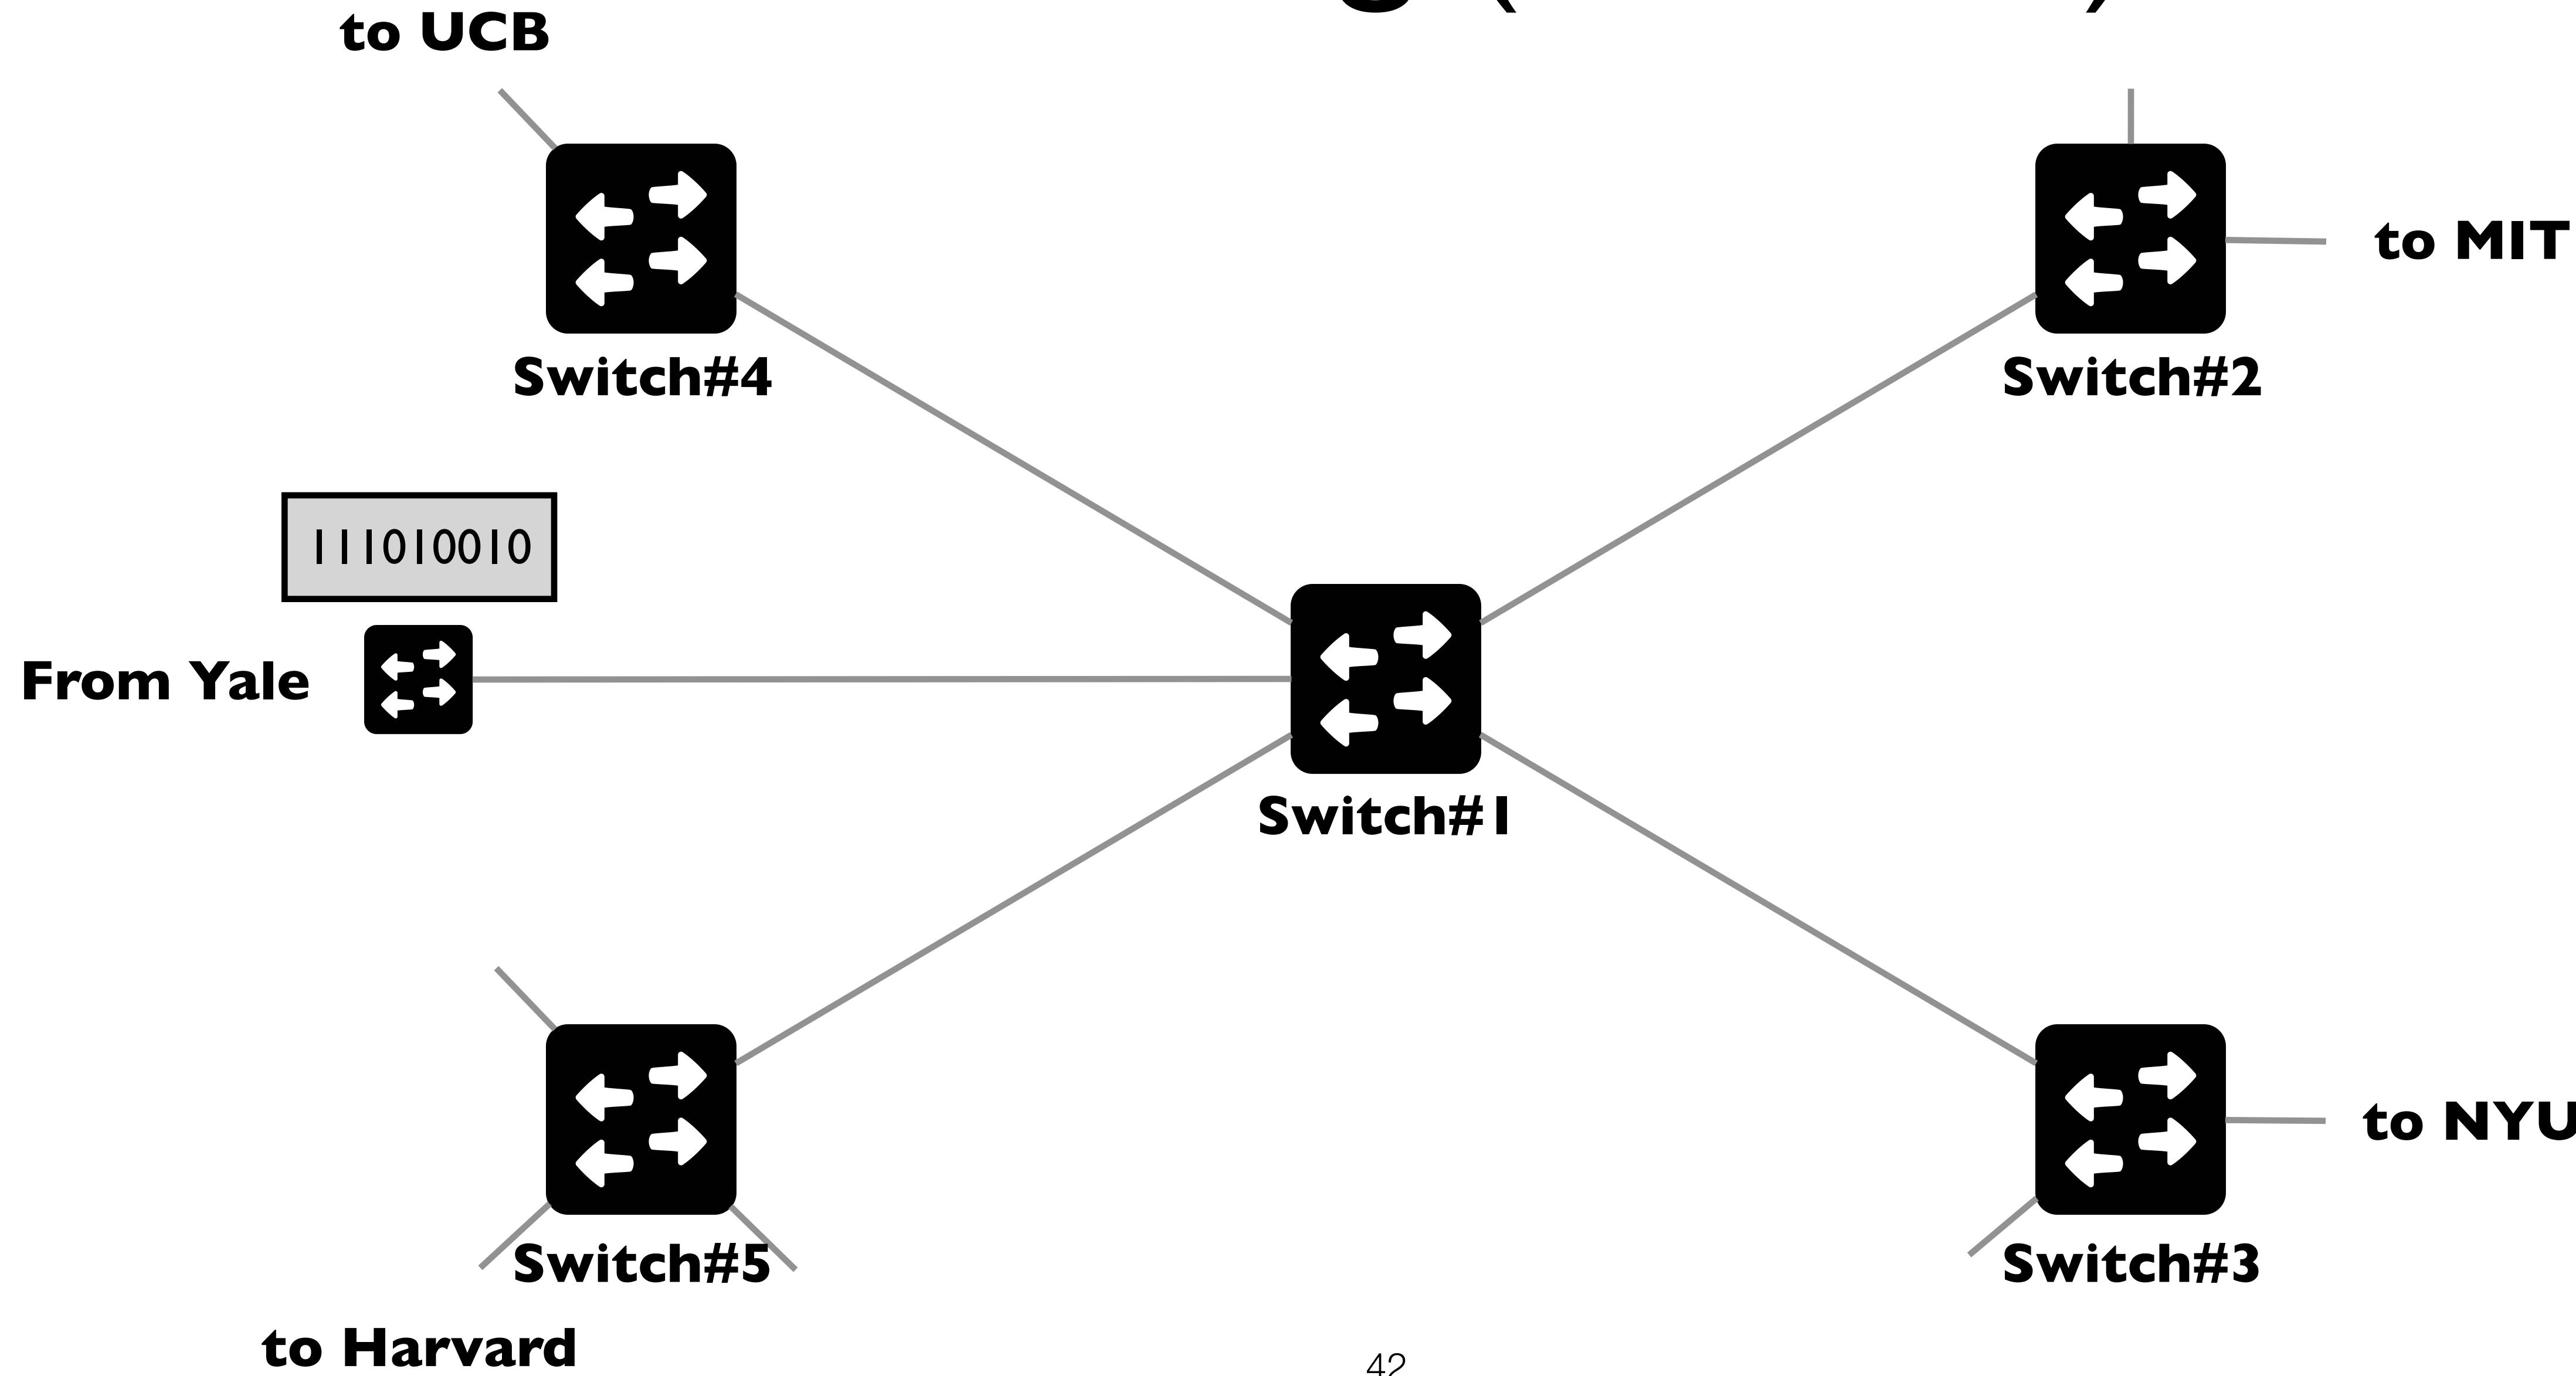
# Addressing (for now)

- Assume each end-host has a unique address
- No particular structure to these addresses

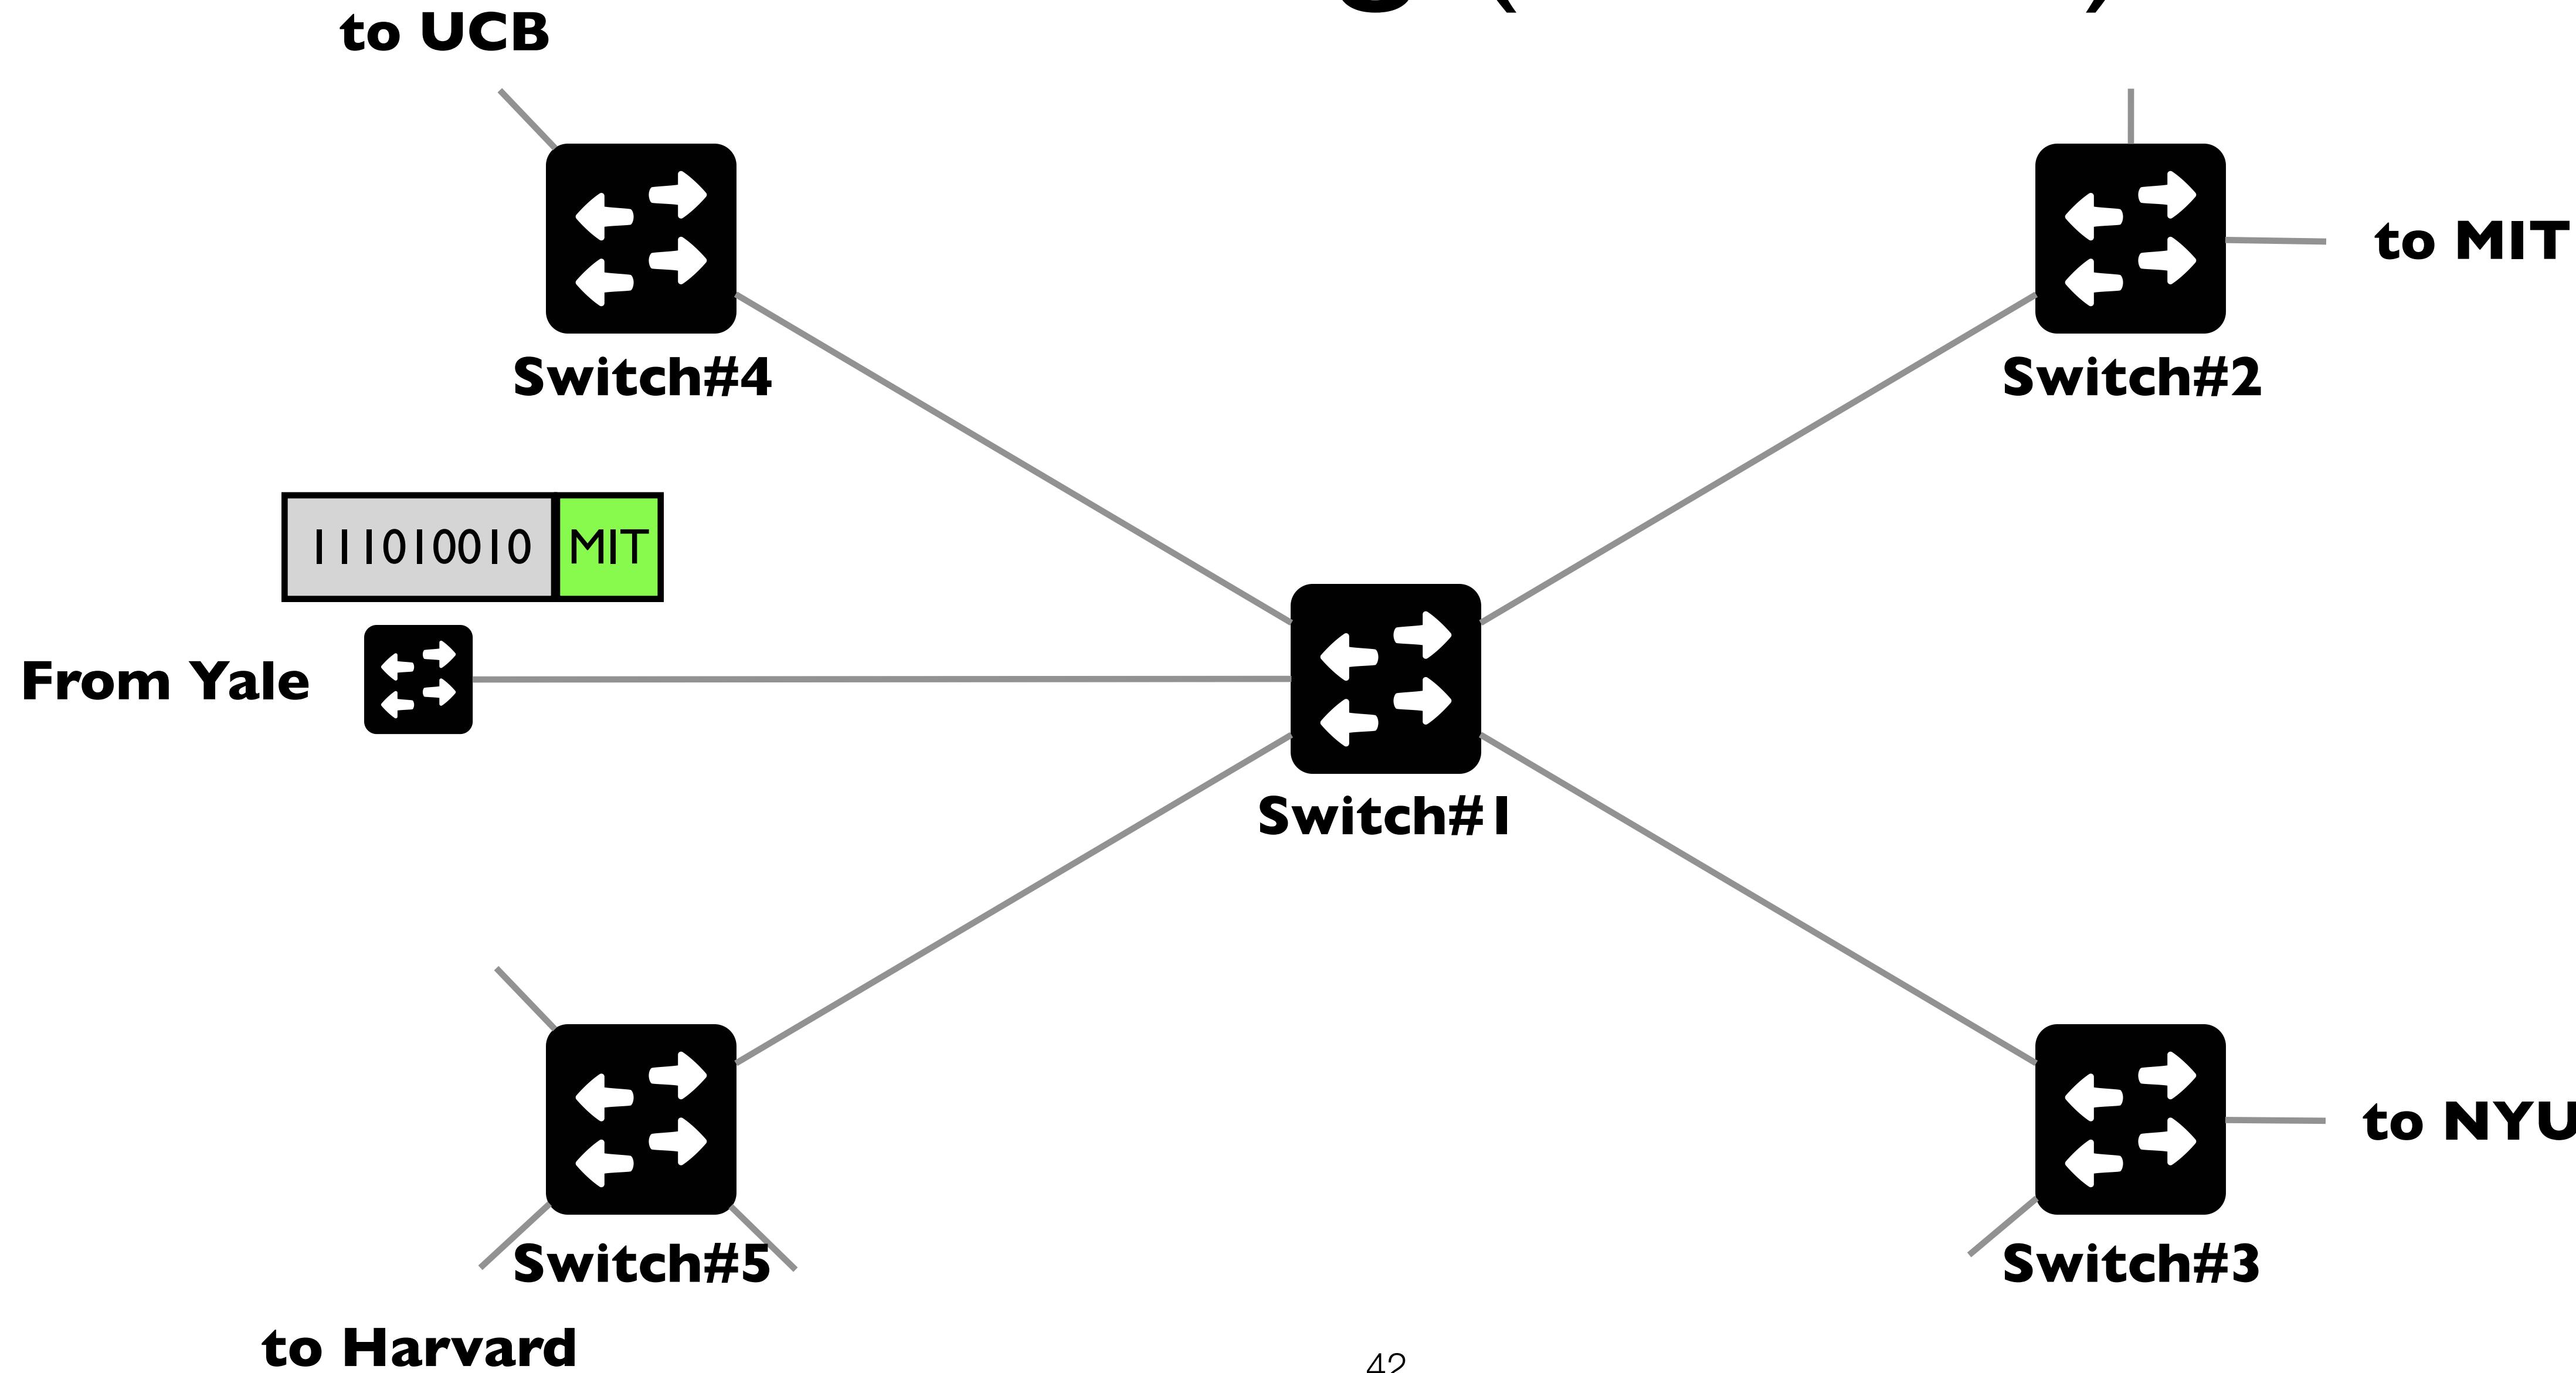
# Addressing (for now)

- Assume each end-host has a unique address
- No particular structure to these addresses
- Will cover IP addresses later in the course

# Addressing (for now)



# Addressing (for now)



# Forwarding

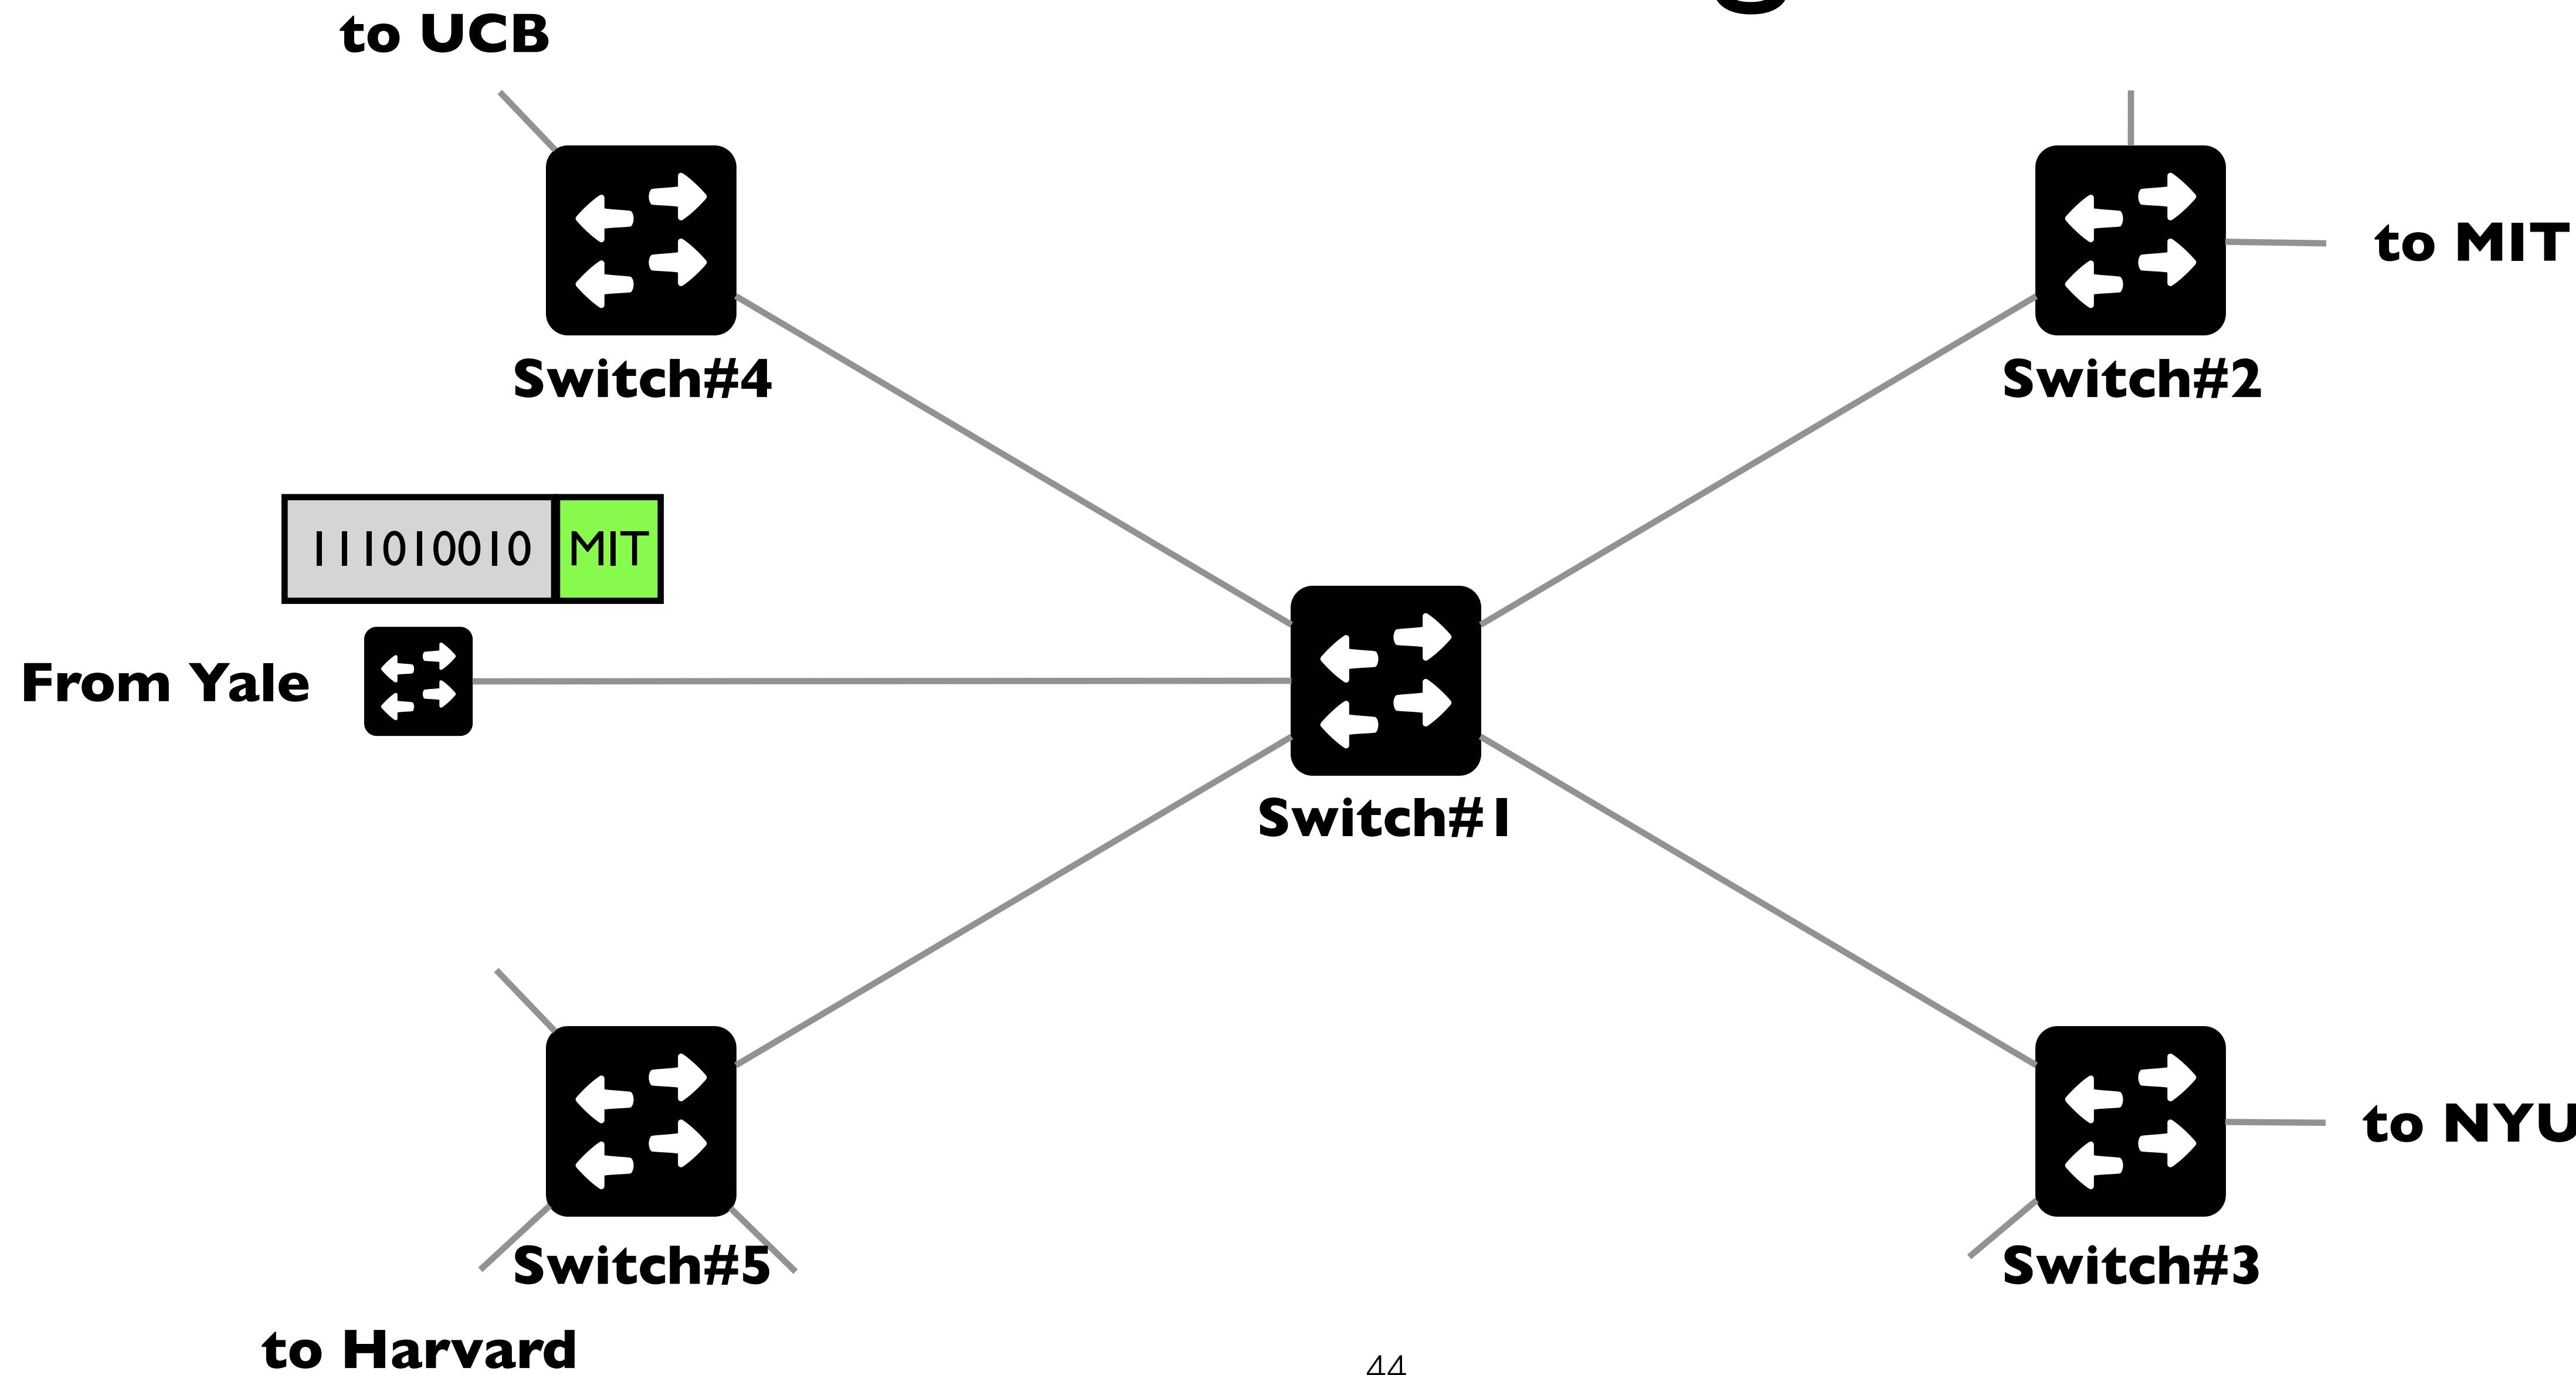
# Forwarding

- **Local** router process that determines the output link (aka “next hop”) for each packet

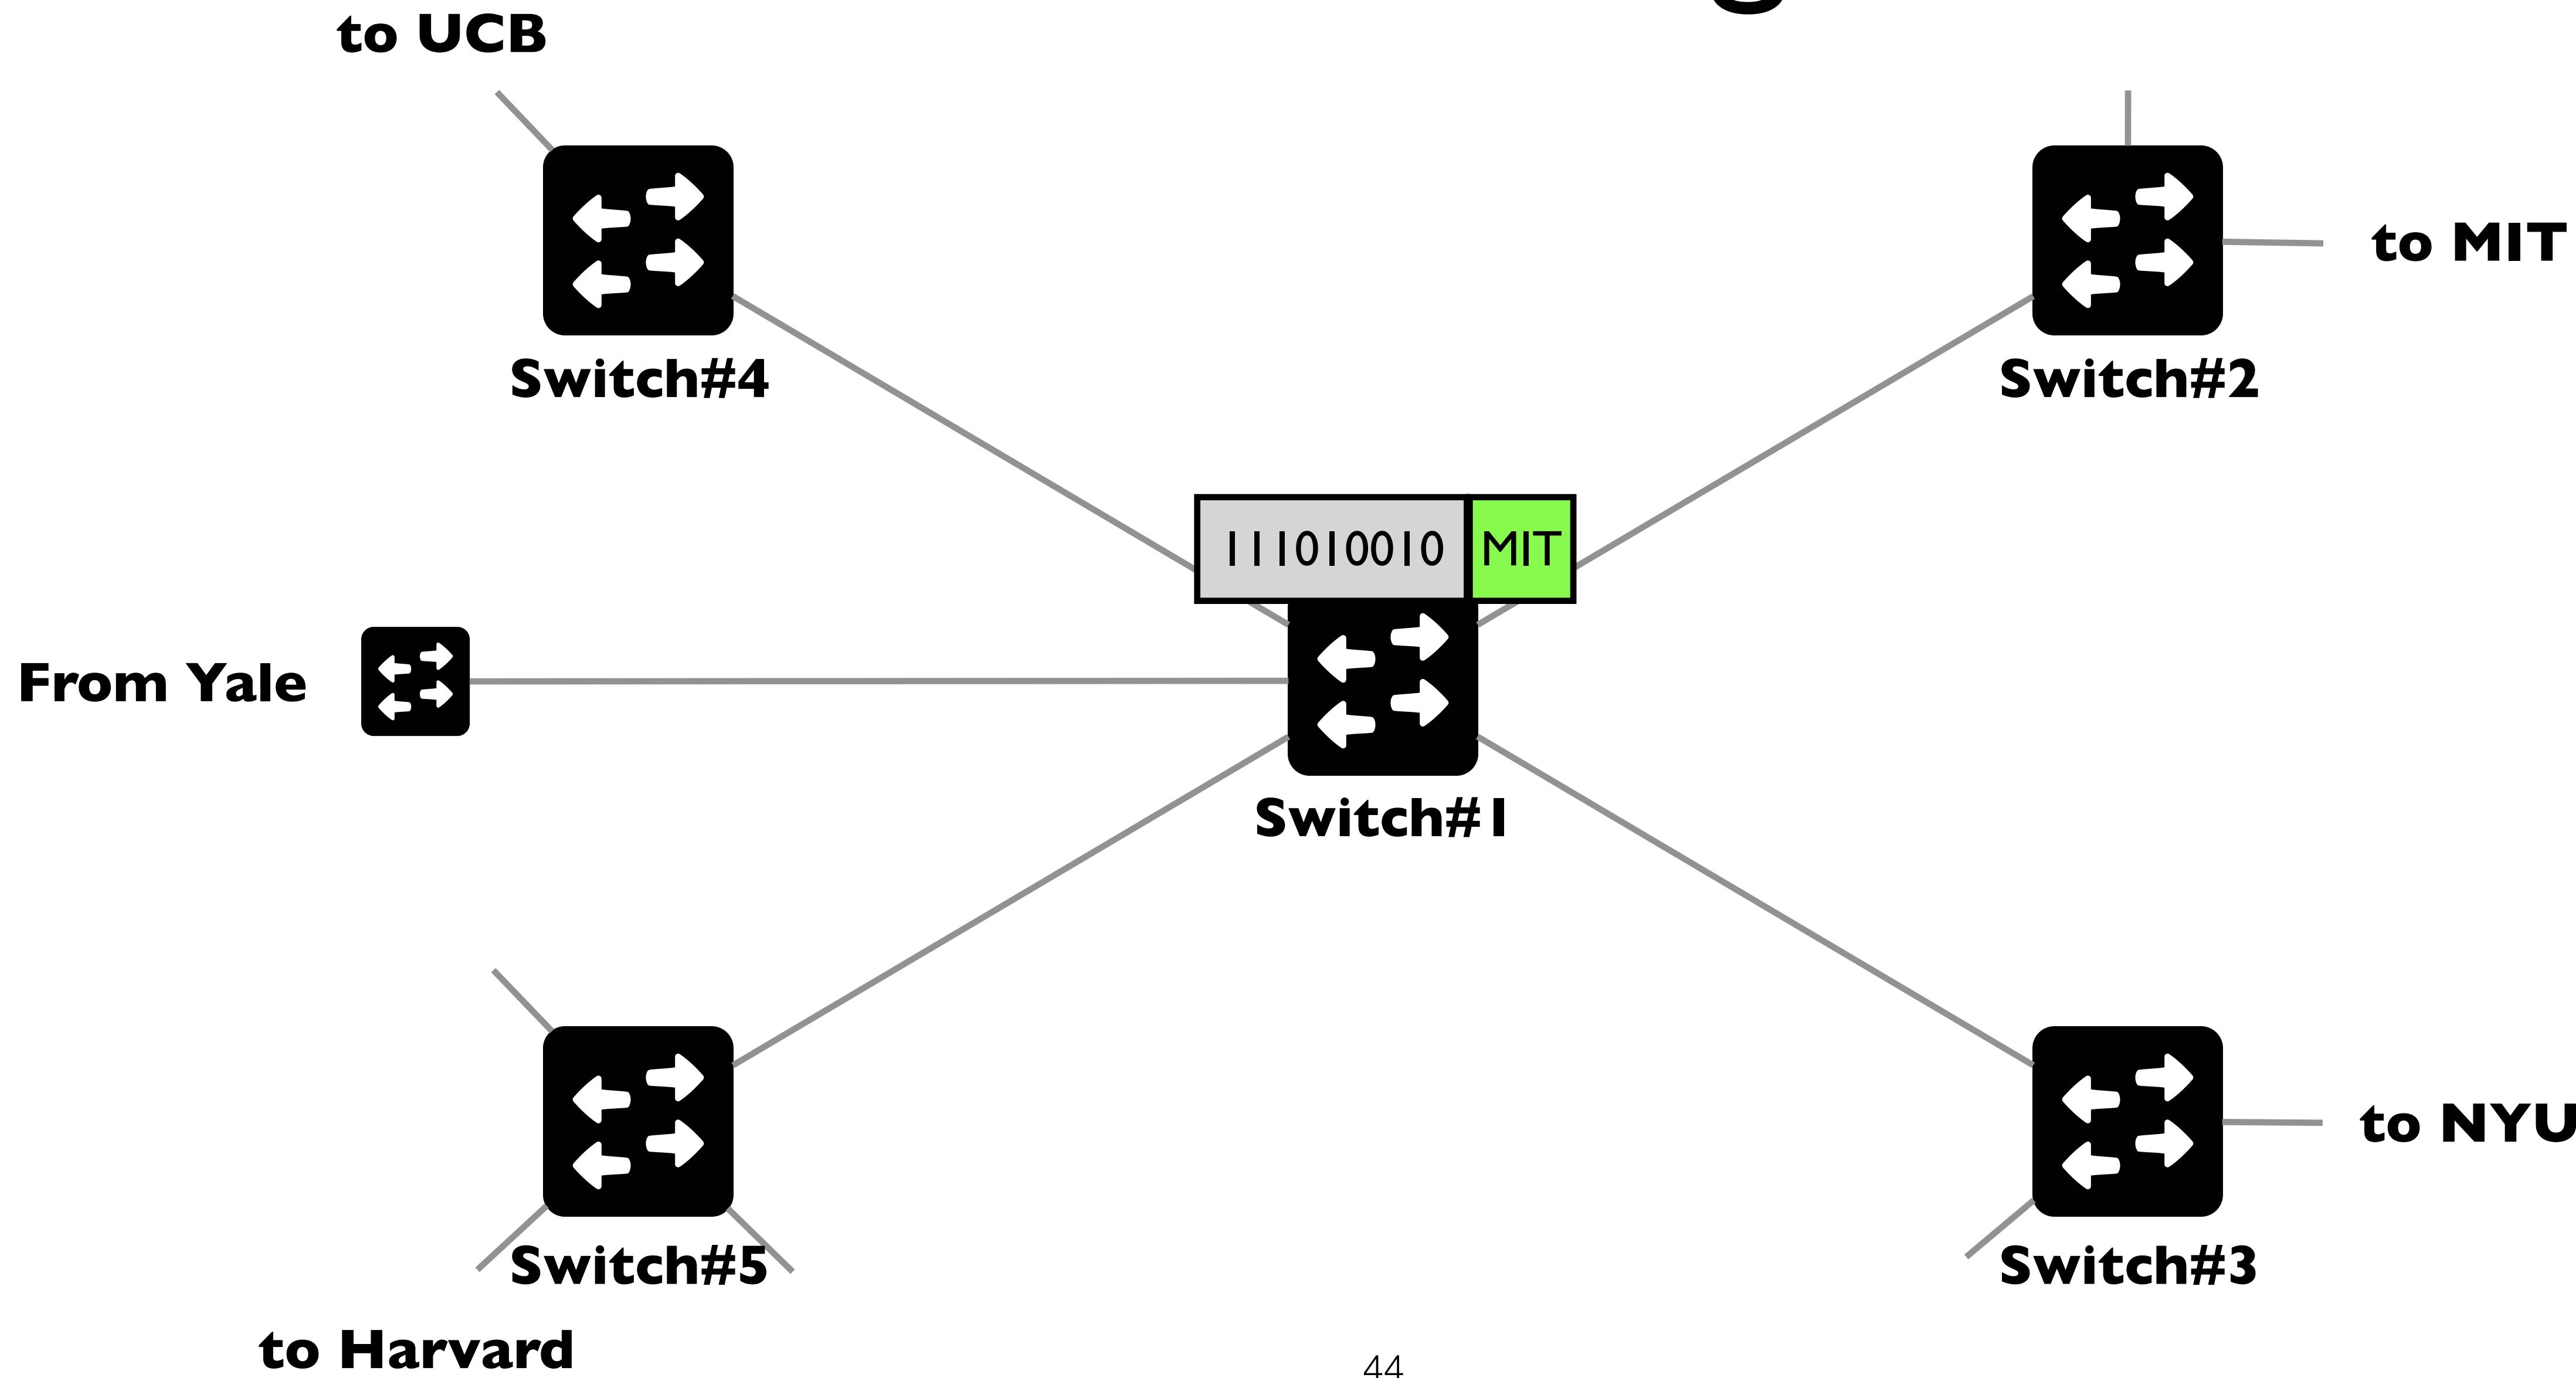
# Forwarding

- **Local** router process that determines the output link (aka “next hop”) for each packet
- **How?**
  - Read address from packet’s network layer header
  - Search forwarding table

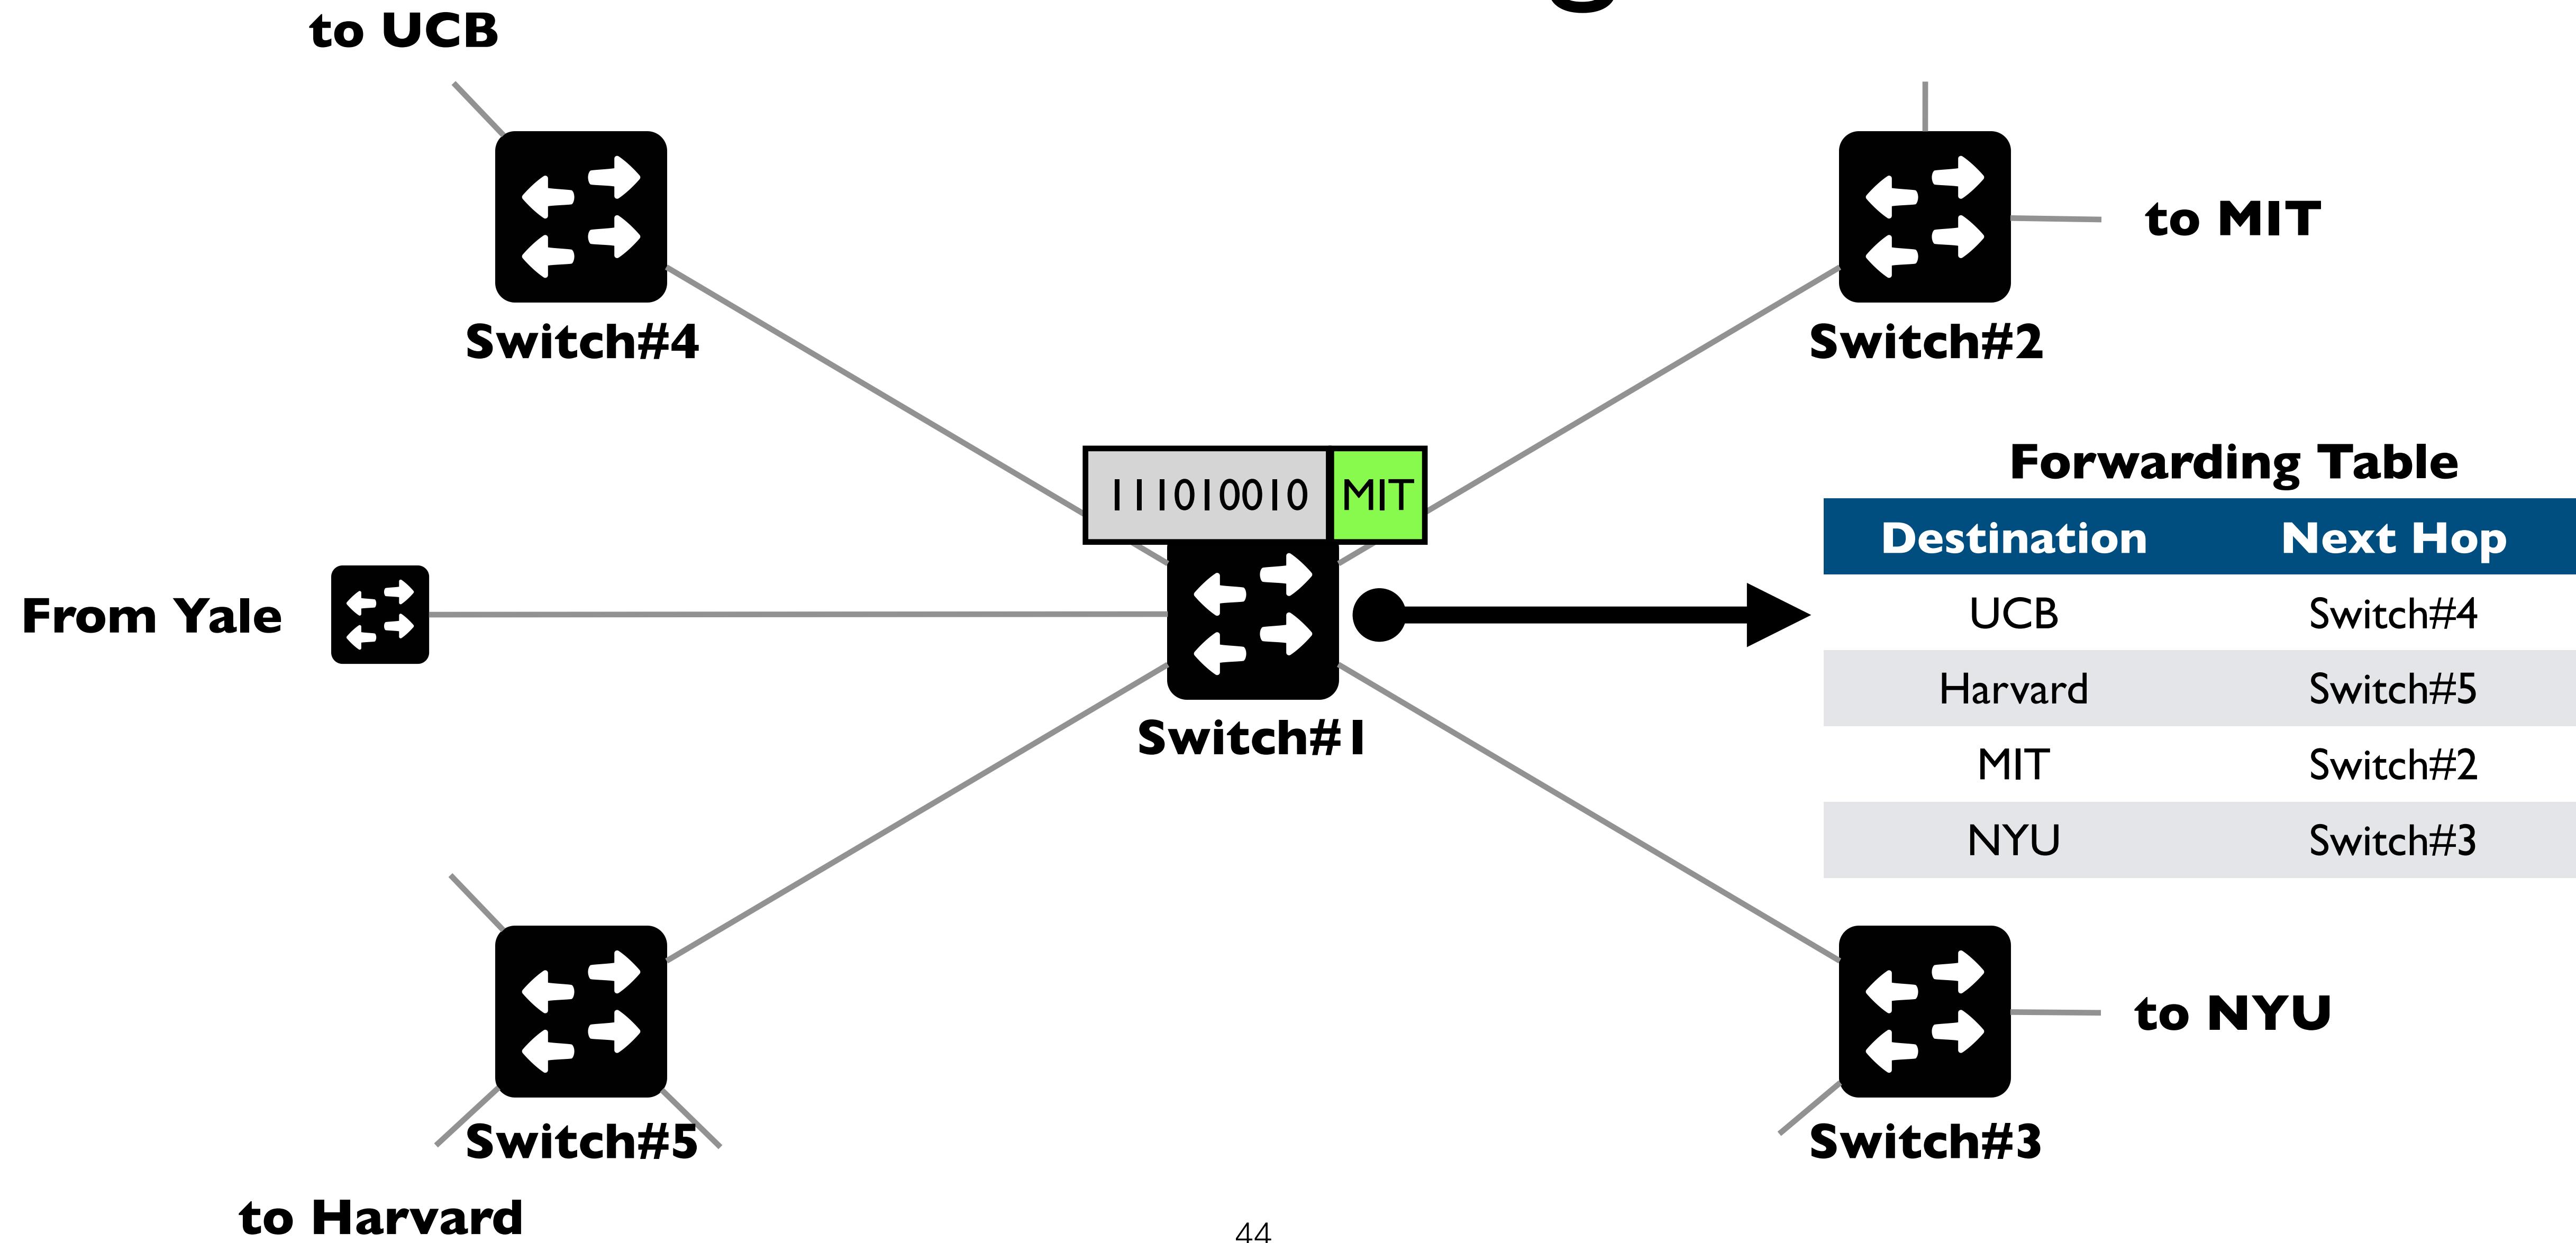
# Forwarding



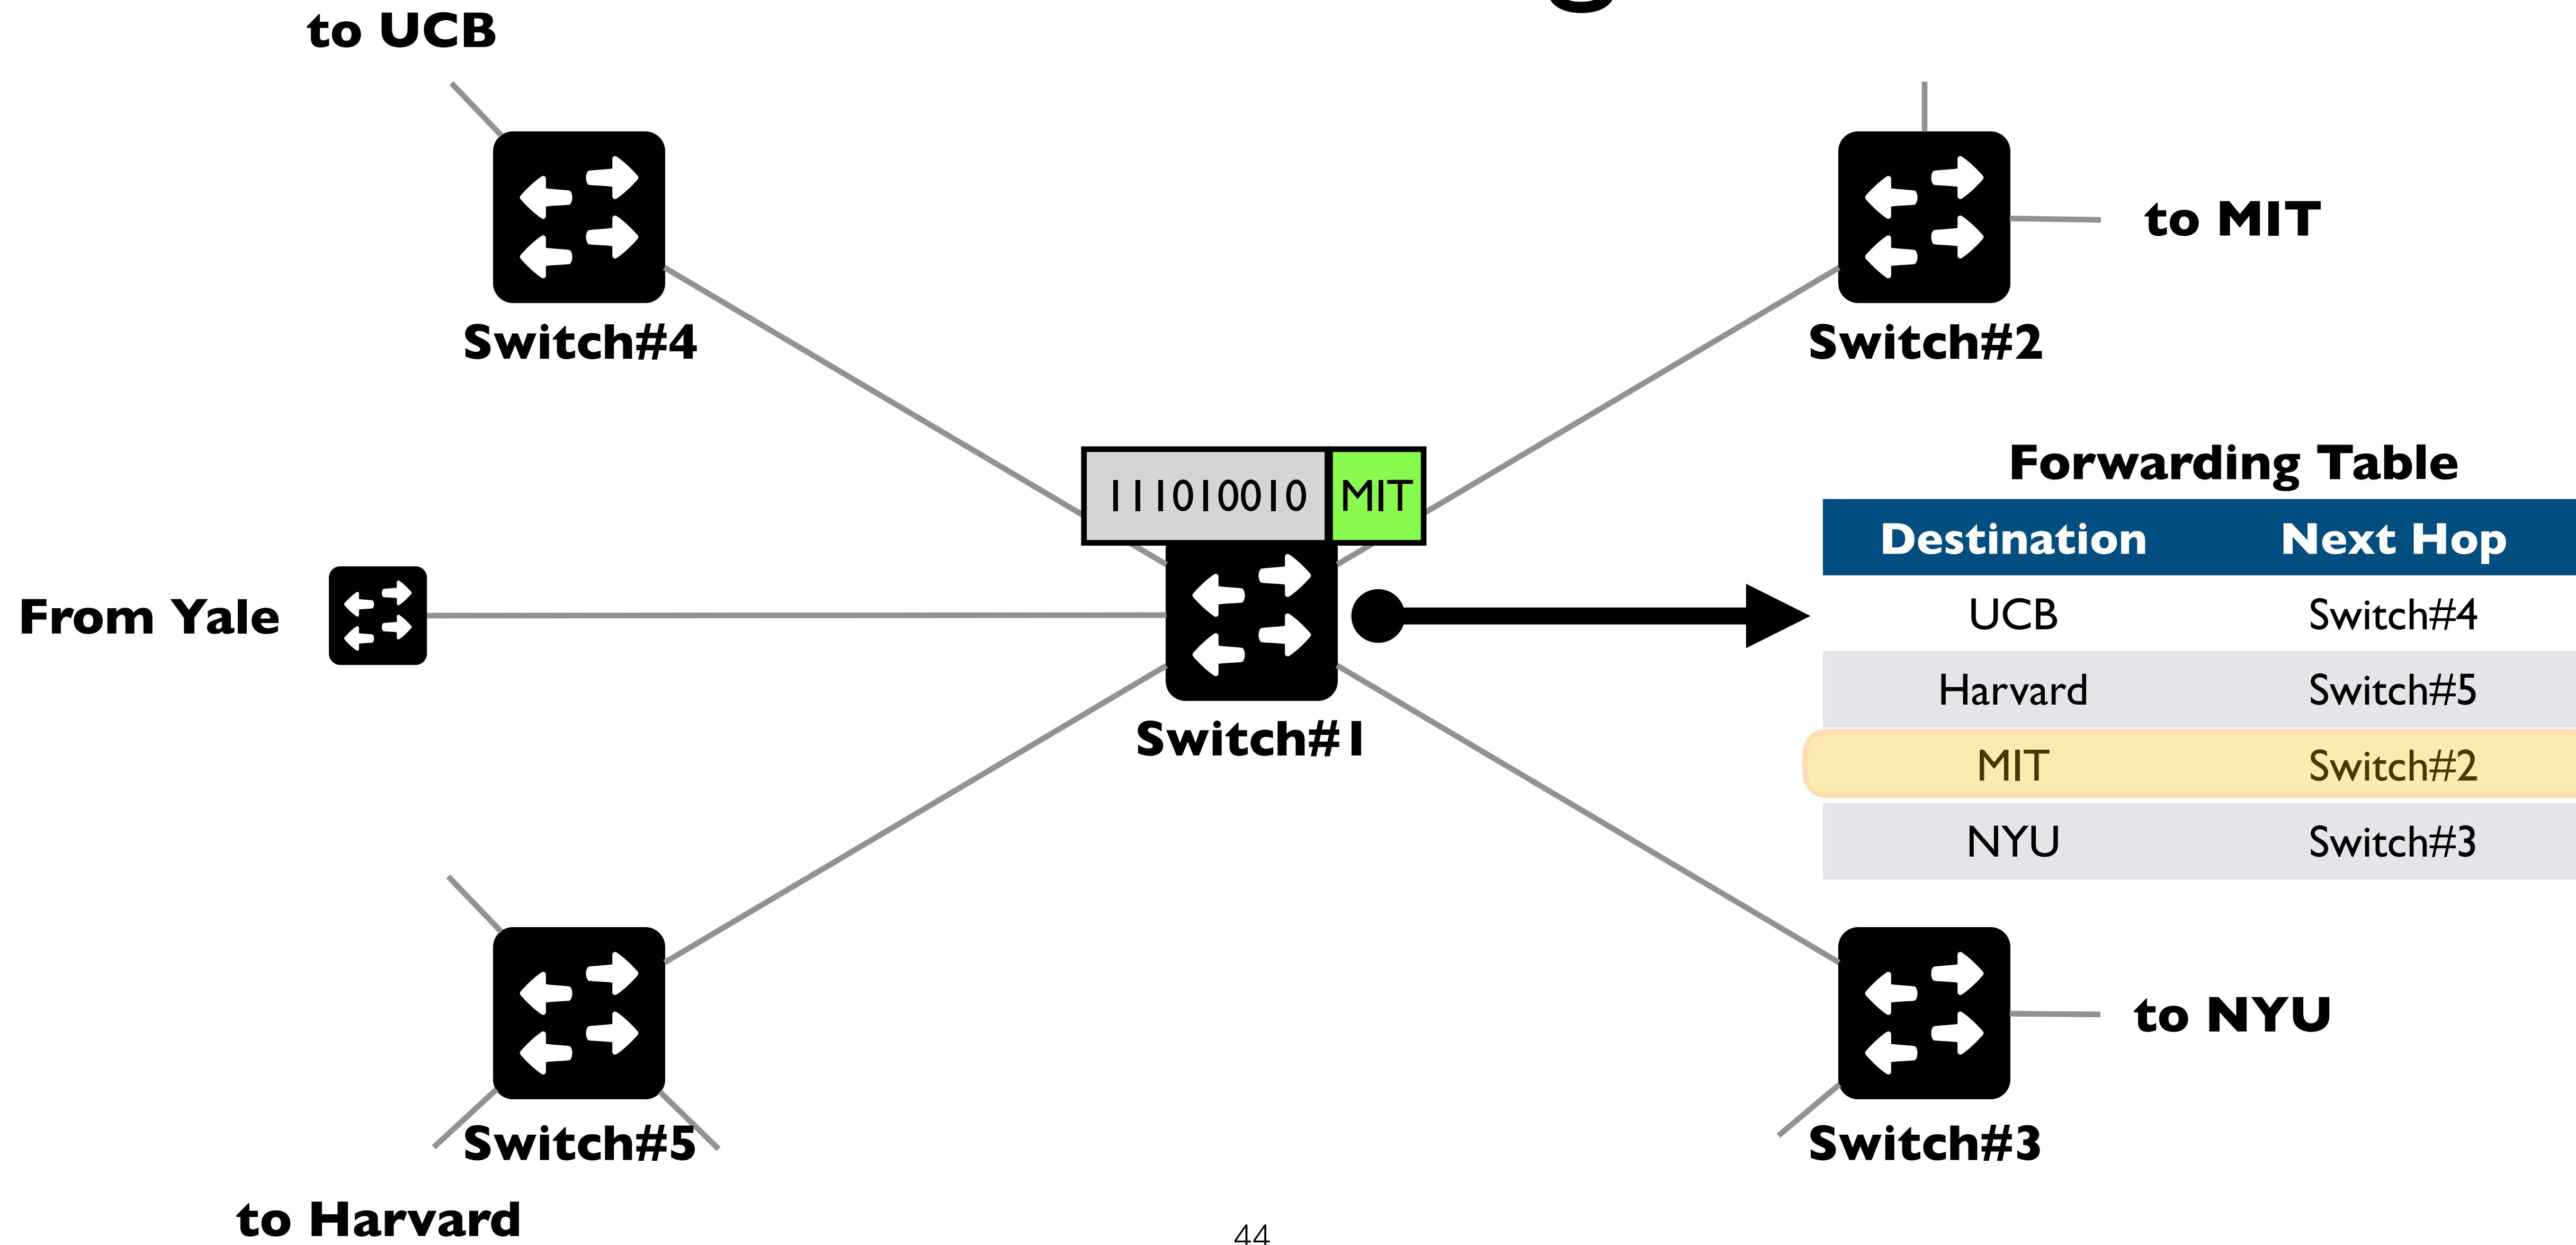
# Forwarding



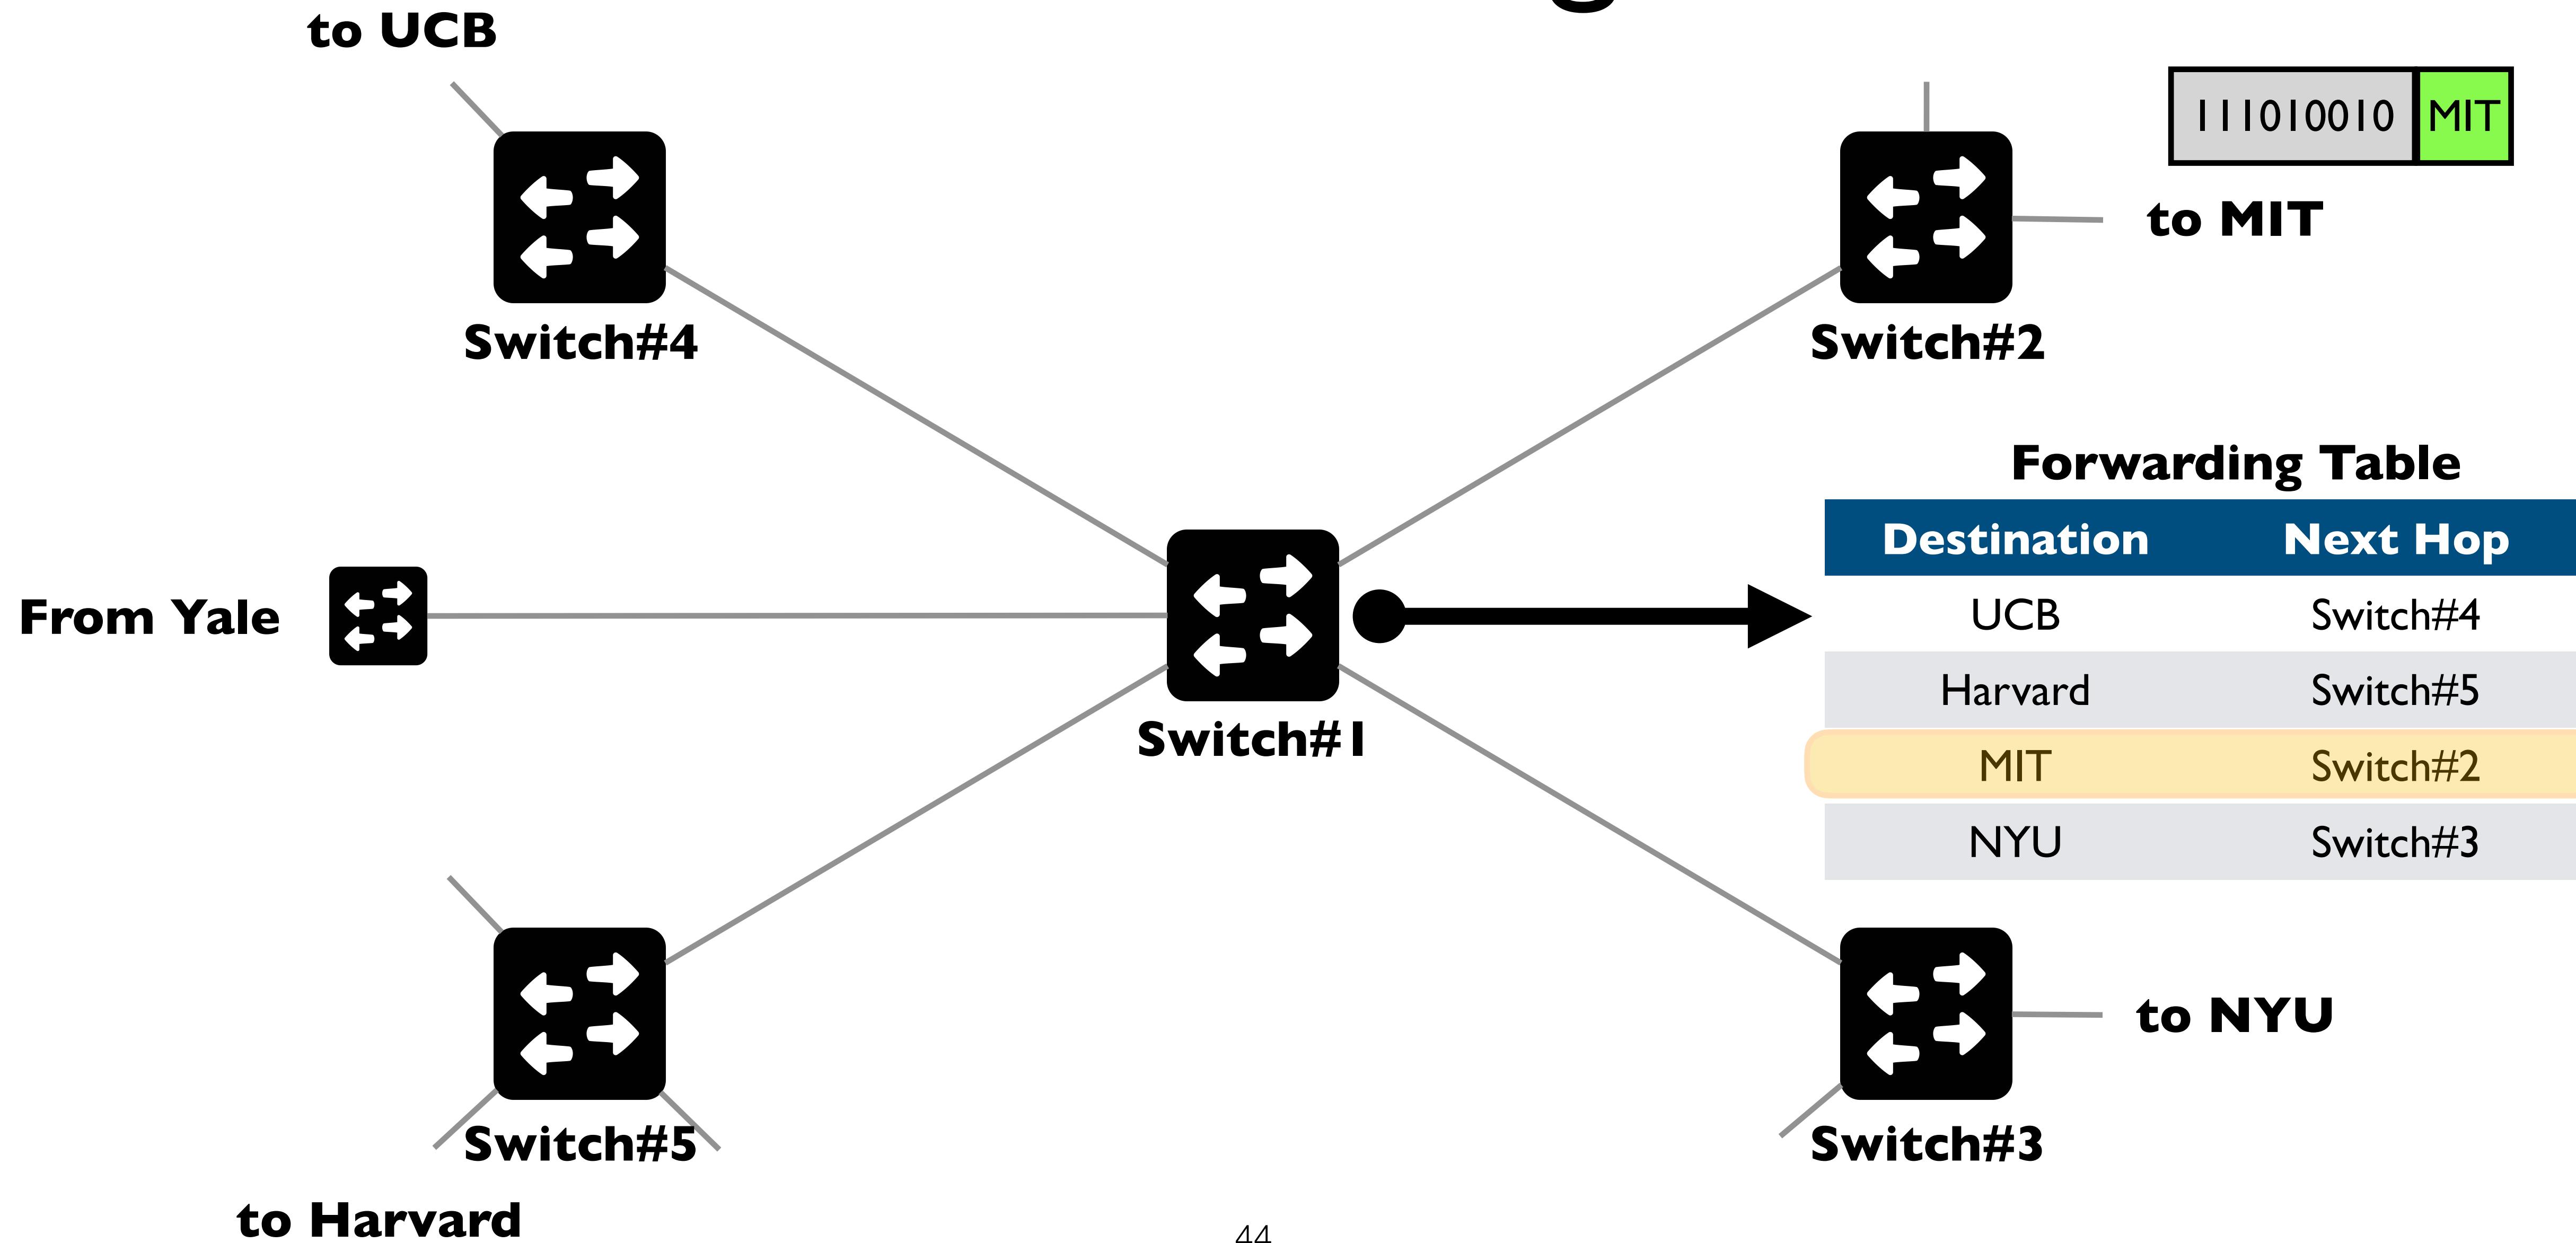
# Forwarding



# Forwarding



# Forwarding

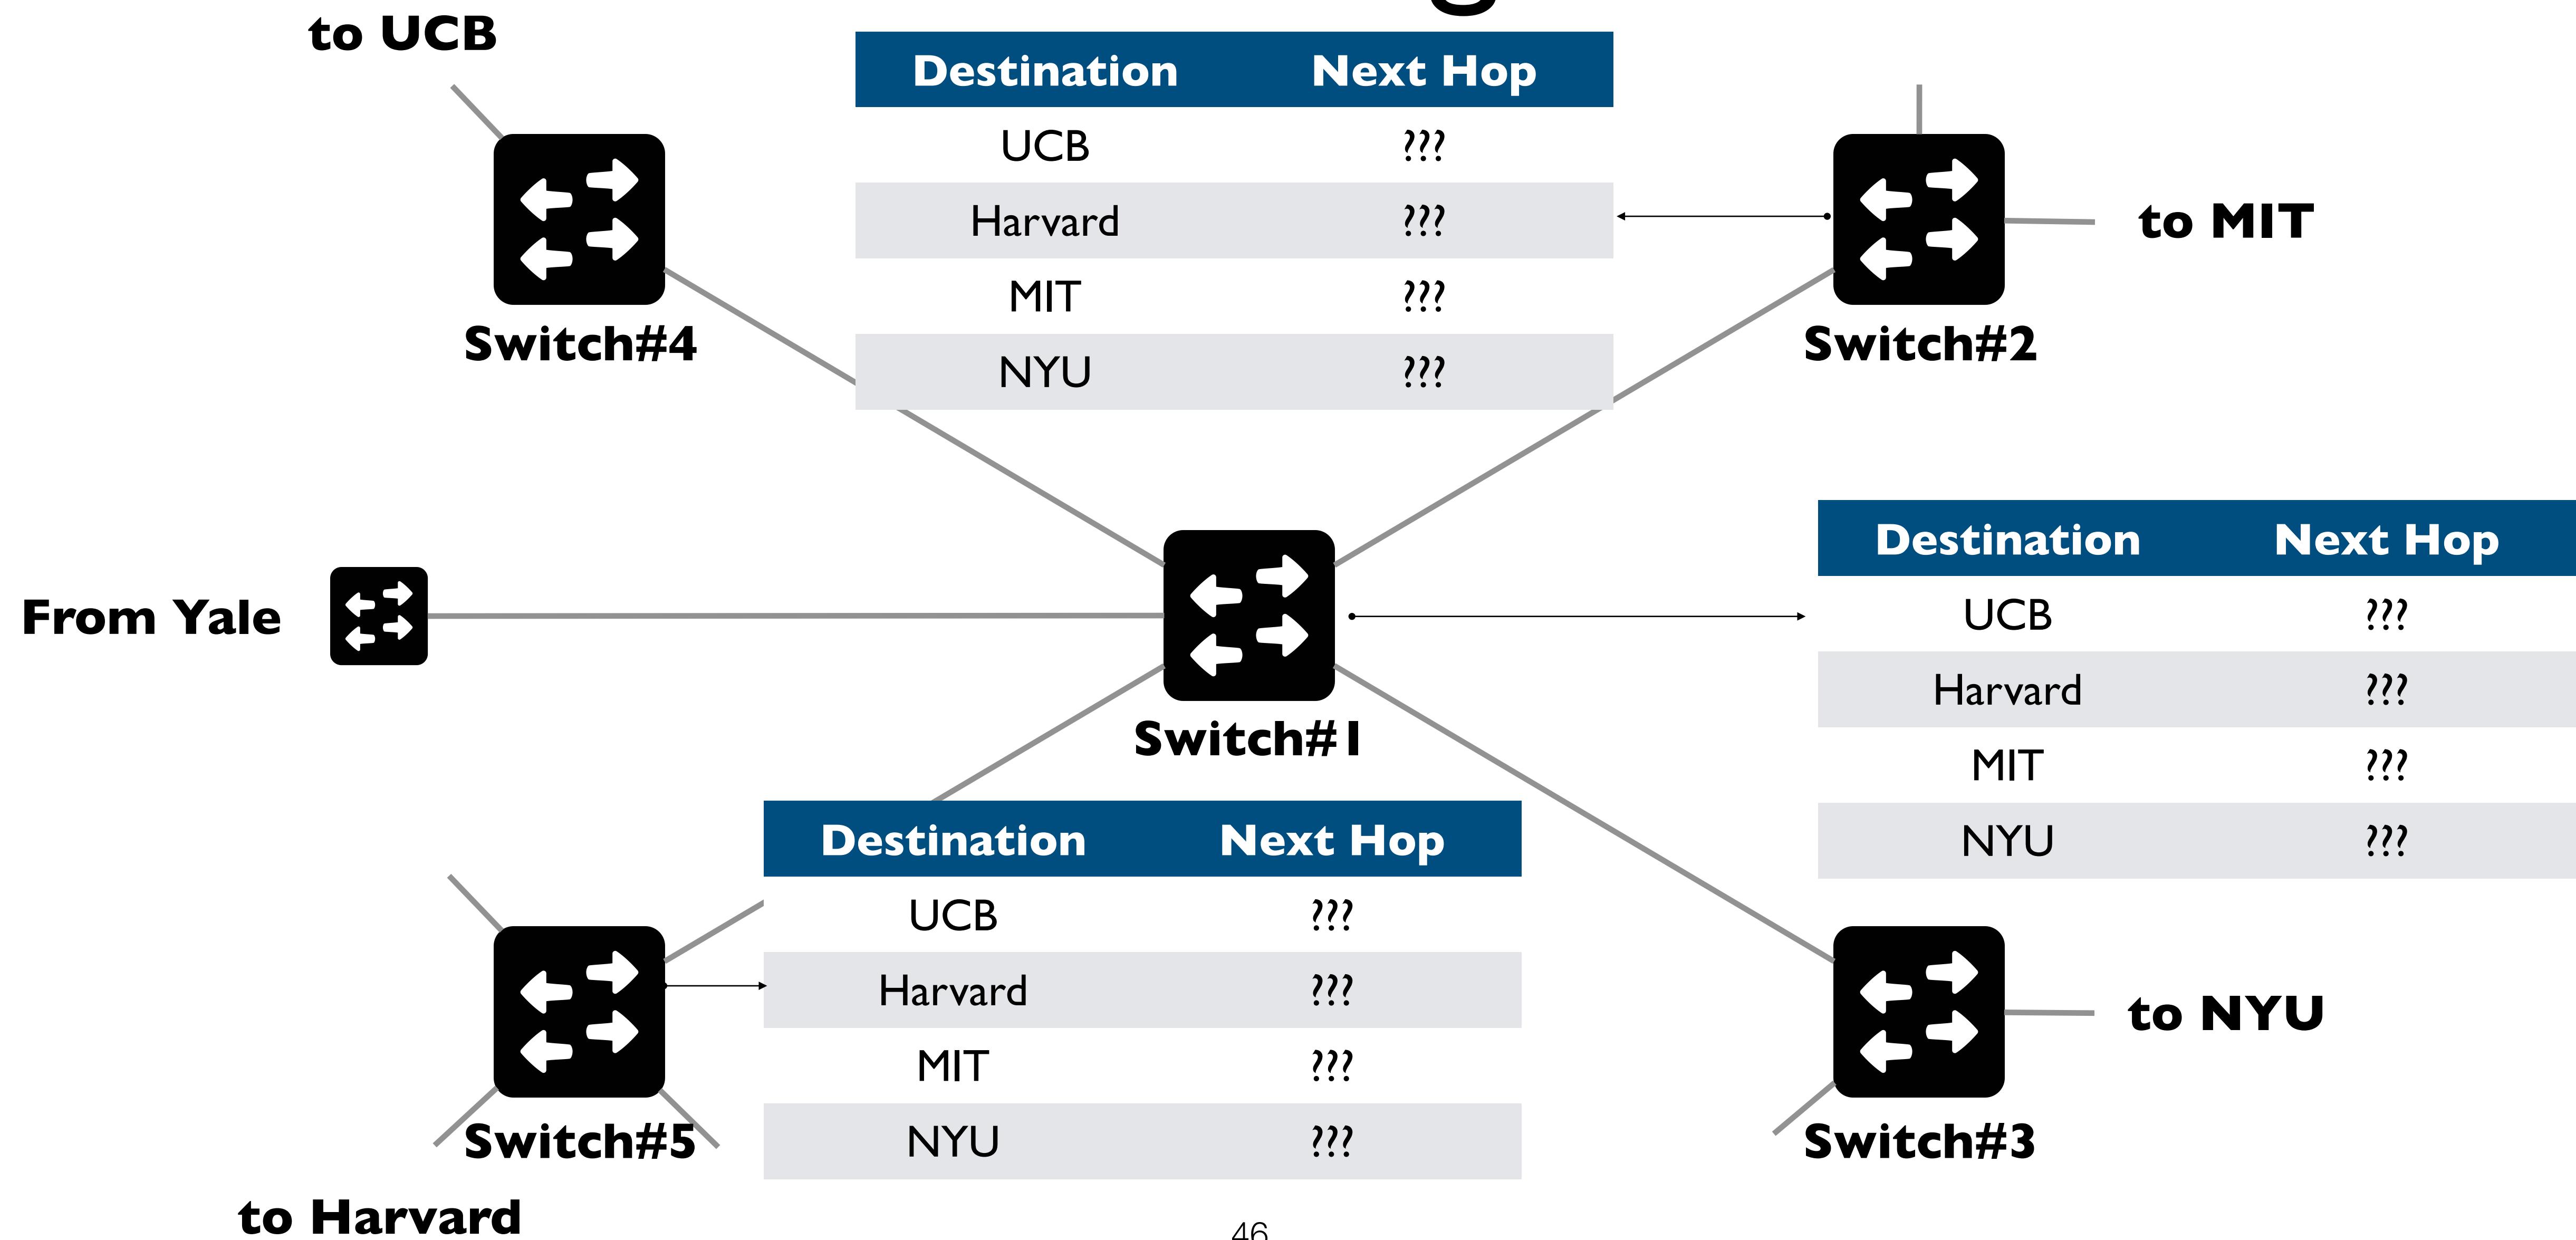


# Routing

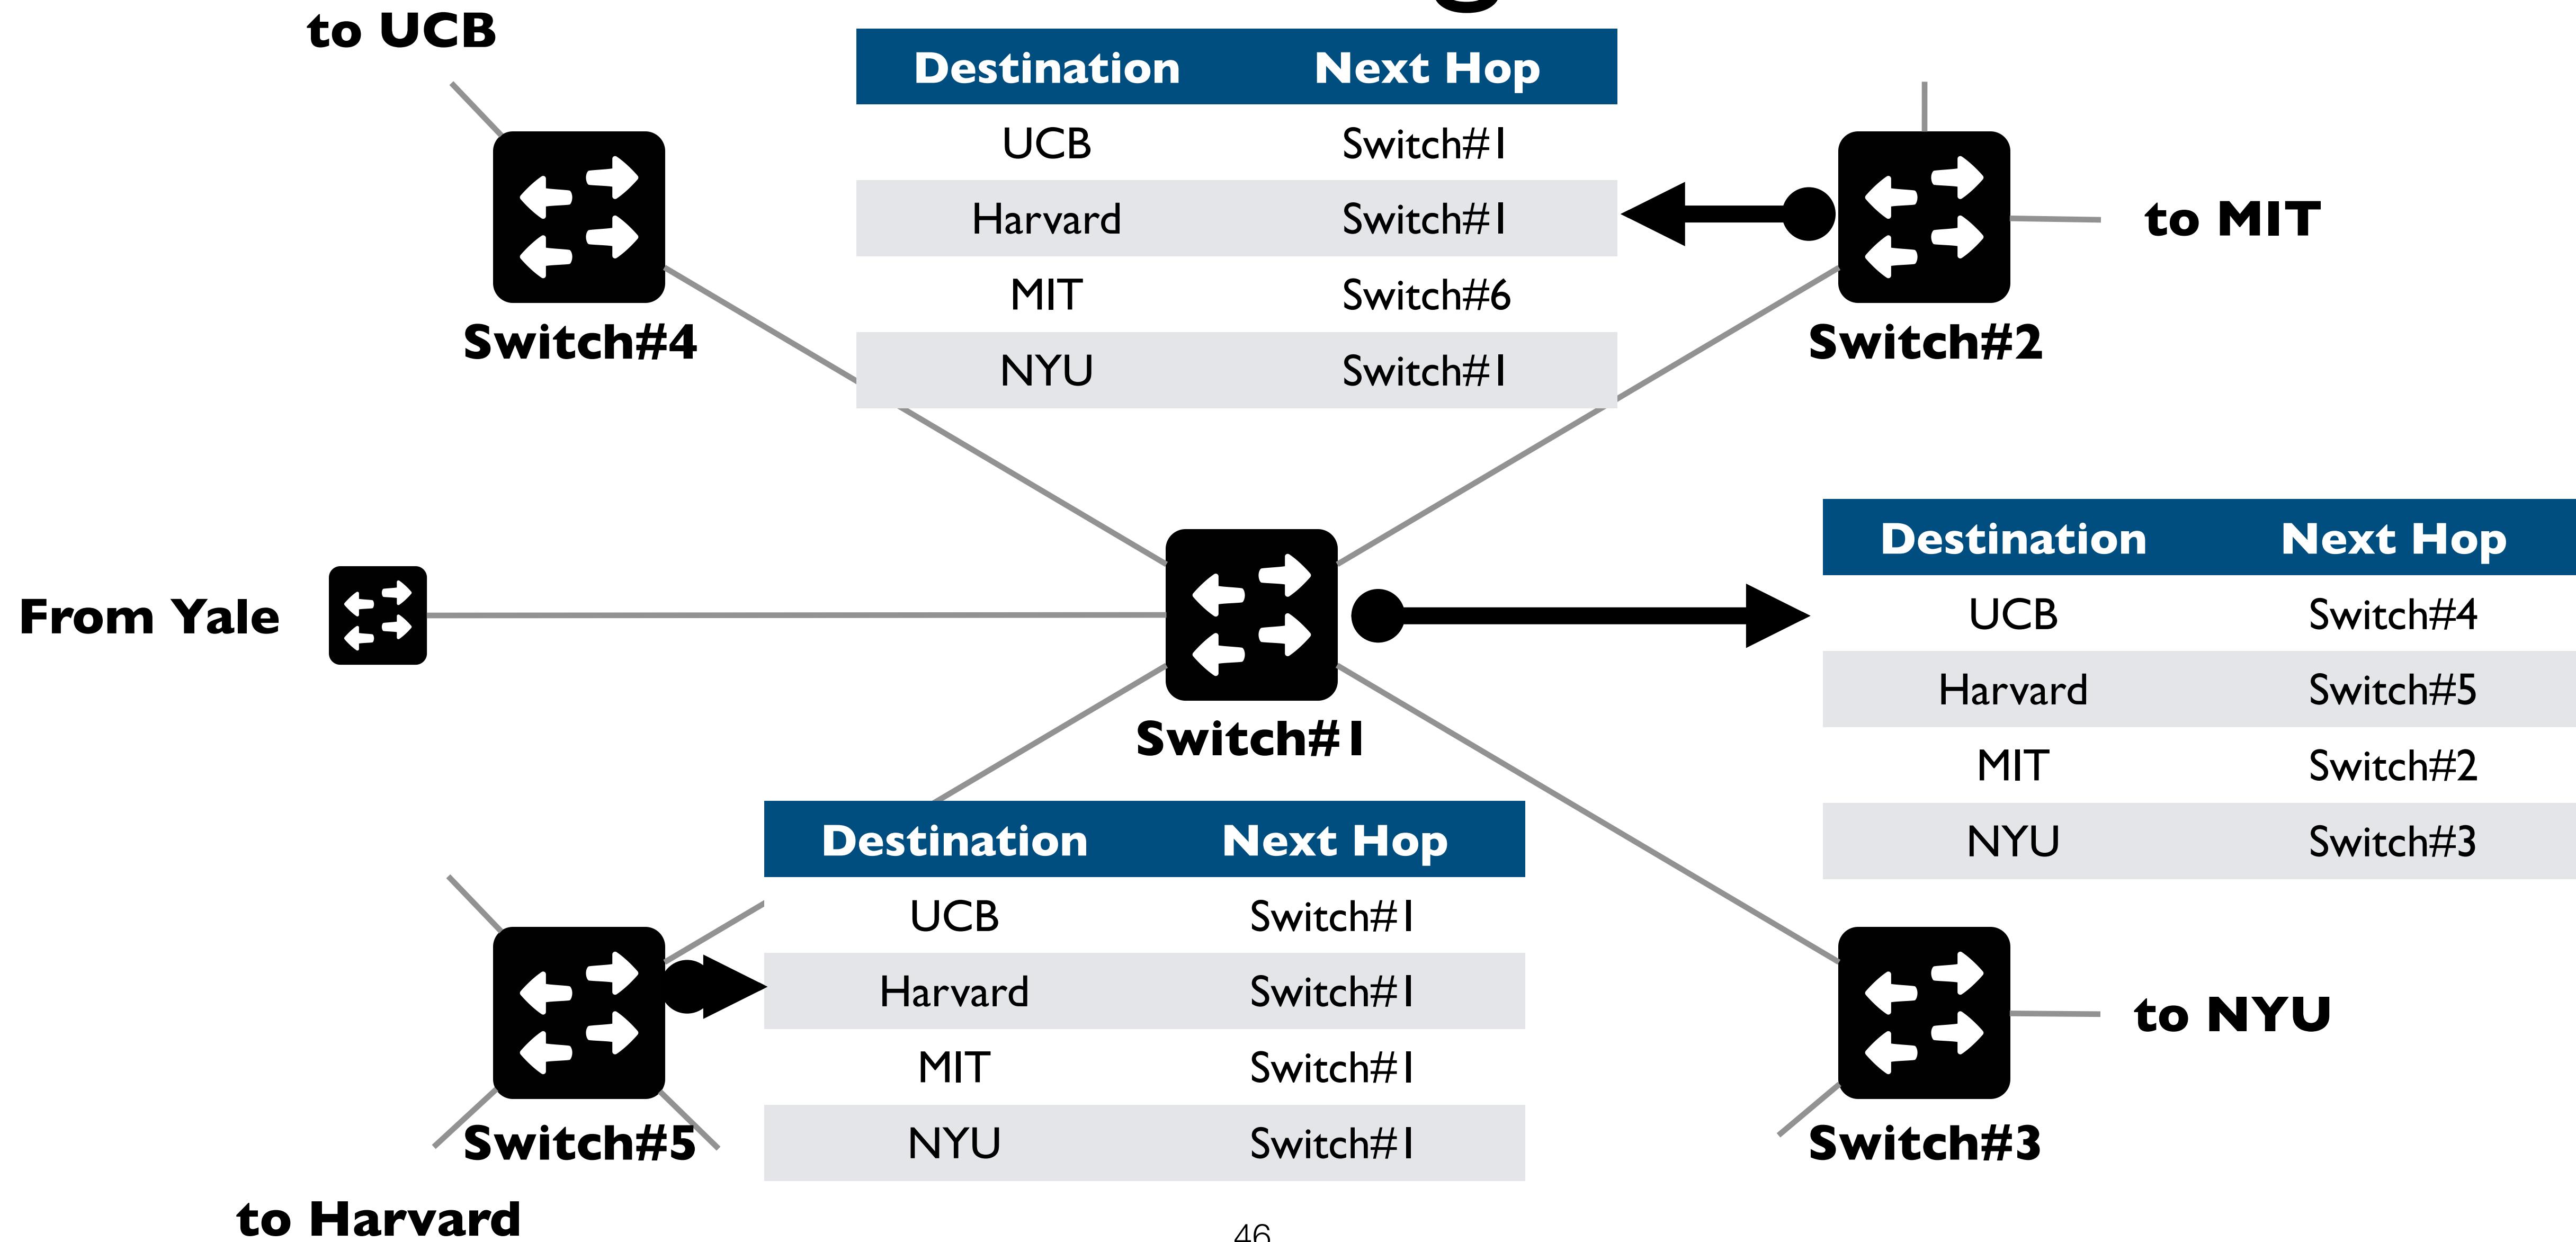
# Routing

- **Network-wide** process that determines the content of forwarding tables
  - Determines the end-to-end path for each destination

# Routing



# Routing



# Routing

# Routing

- **Network-wide** process that determines the content of forwarding tables
  - Determines the end-to-end path for each destination

# Routing

- **Network-wide** process that determines the content of forwarding tables
  - Determines the end-to-end path for each destination
- **How?**
  - Coming up soon...

# Forwarding vs. Routing

# Forwarding vs. Routing

- **Forwarding:** “data plane”
  - Directing one data packet
  - Each router using local routing state

# Forwarding vs. Routing

- **Forwarding:** “data plane”
  - Directing one data packet
  - Each router using local routing state
- **Routing:** “control plane”
  - Computing the forwarding tables that guide packets
  - Jointly computed by routers using a distributed algorithm

# Forwarding vs. Routing

- **Forwarding:** “data plane”
  - Directing one data packet
  - Each router using local routing state
- **Routing:** “control plane”
  - Computing the forwarding tables that guide packets
  - Jointly computed by routers using a distributed algorithm
- Very different timescales!

# Goals (for Routing)

# Goals (for Routing)

- **VI: Find a path to a given destination**

# Goals (for Routing)

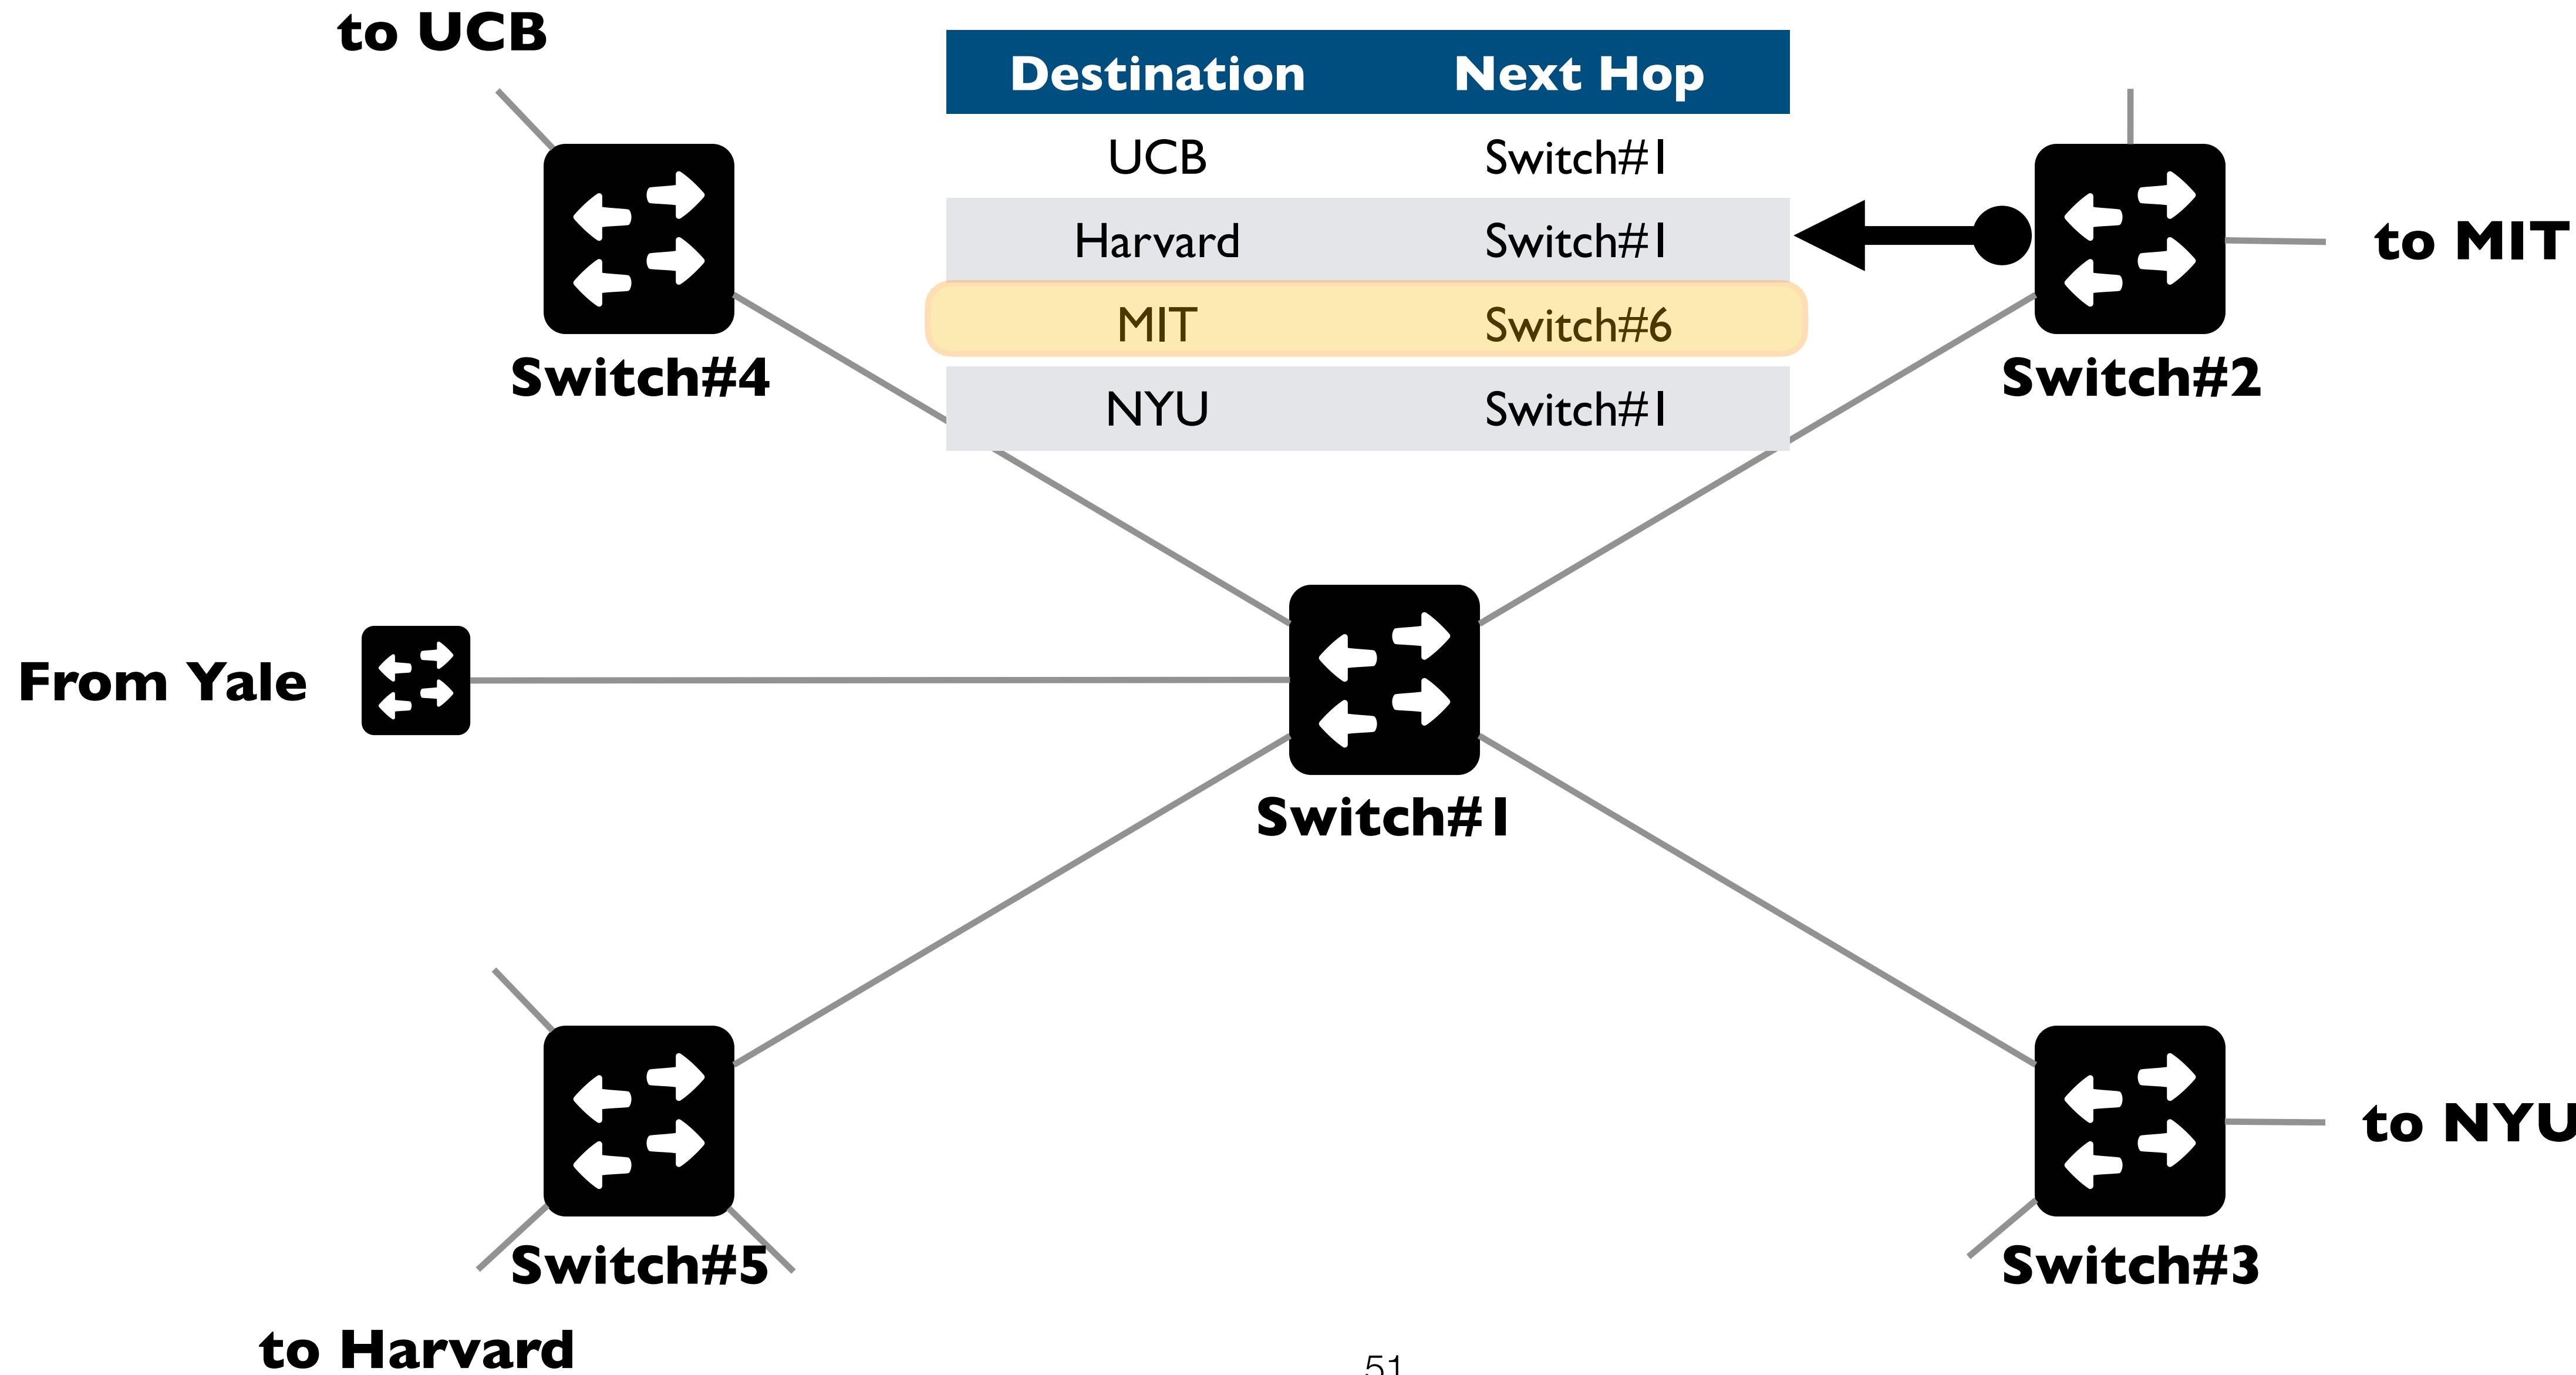
- **VI: Find a path to a given destination**
- How do we know that the state contained in forwarding tables meets our goal?
  - This is what “validity” of routing state tells us
  - [this is non-standard terminology]

# Local vs. Global view of state

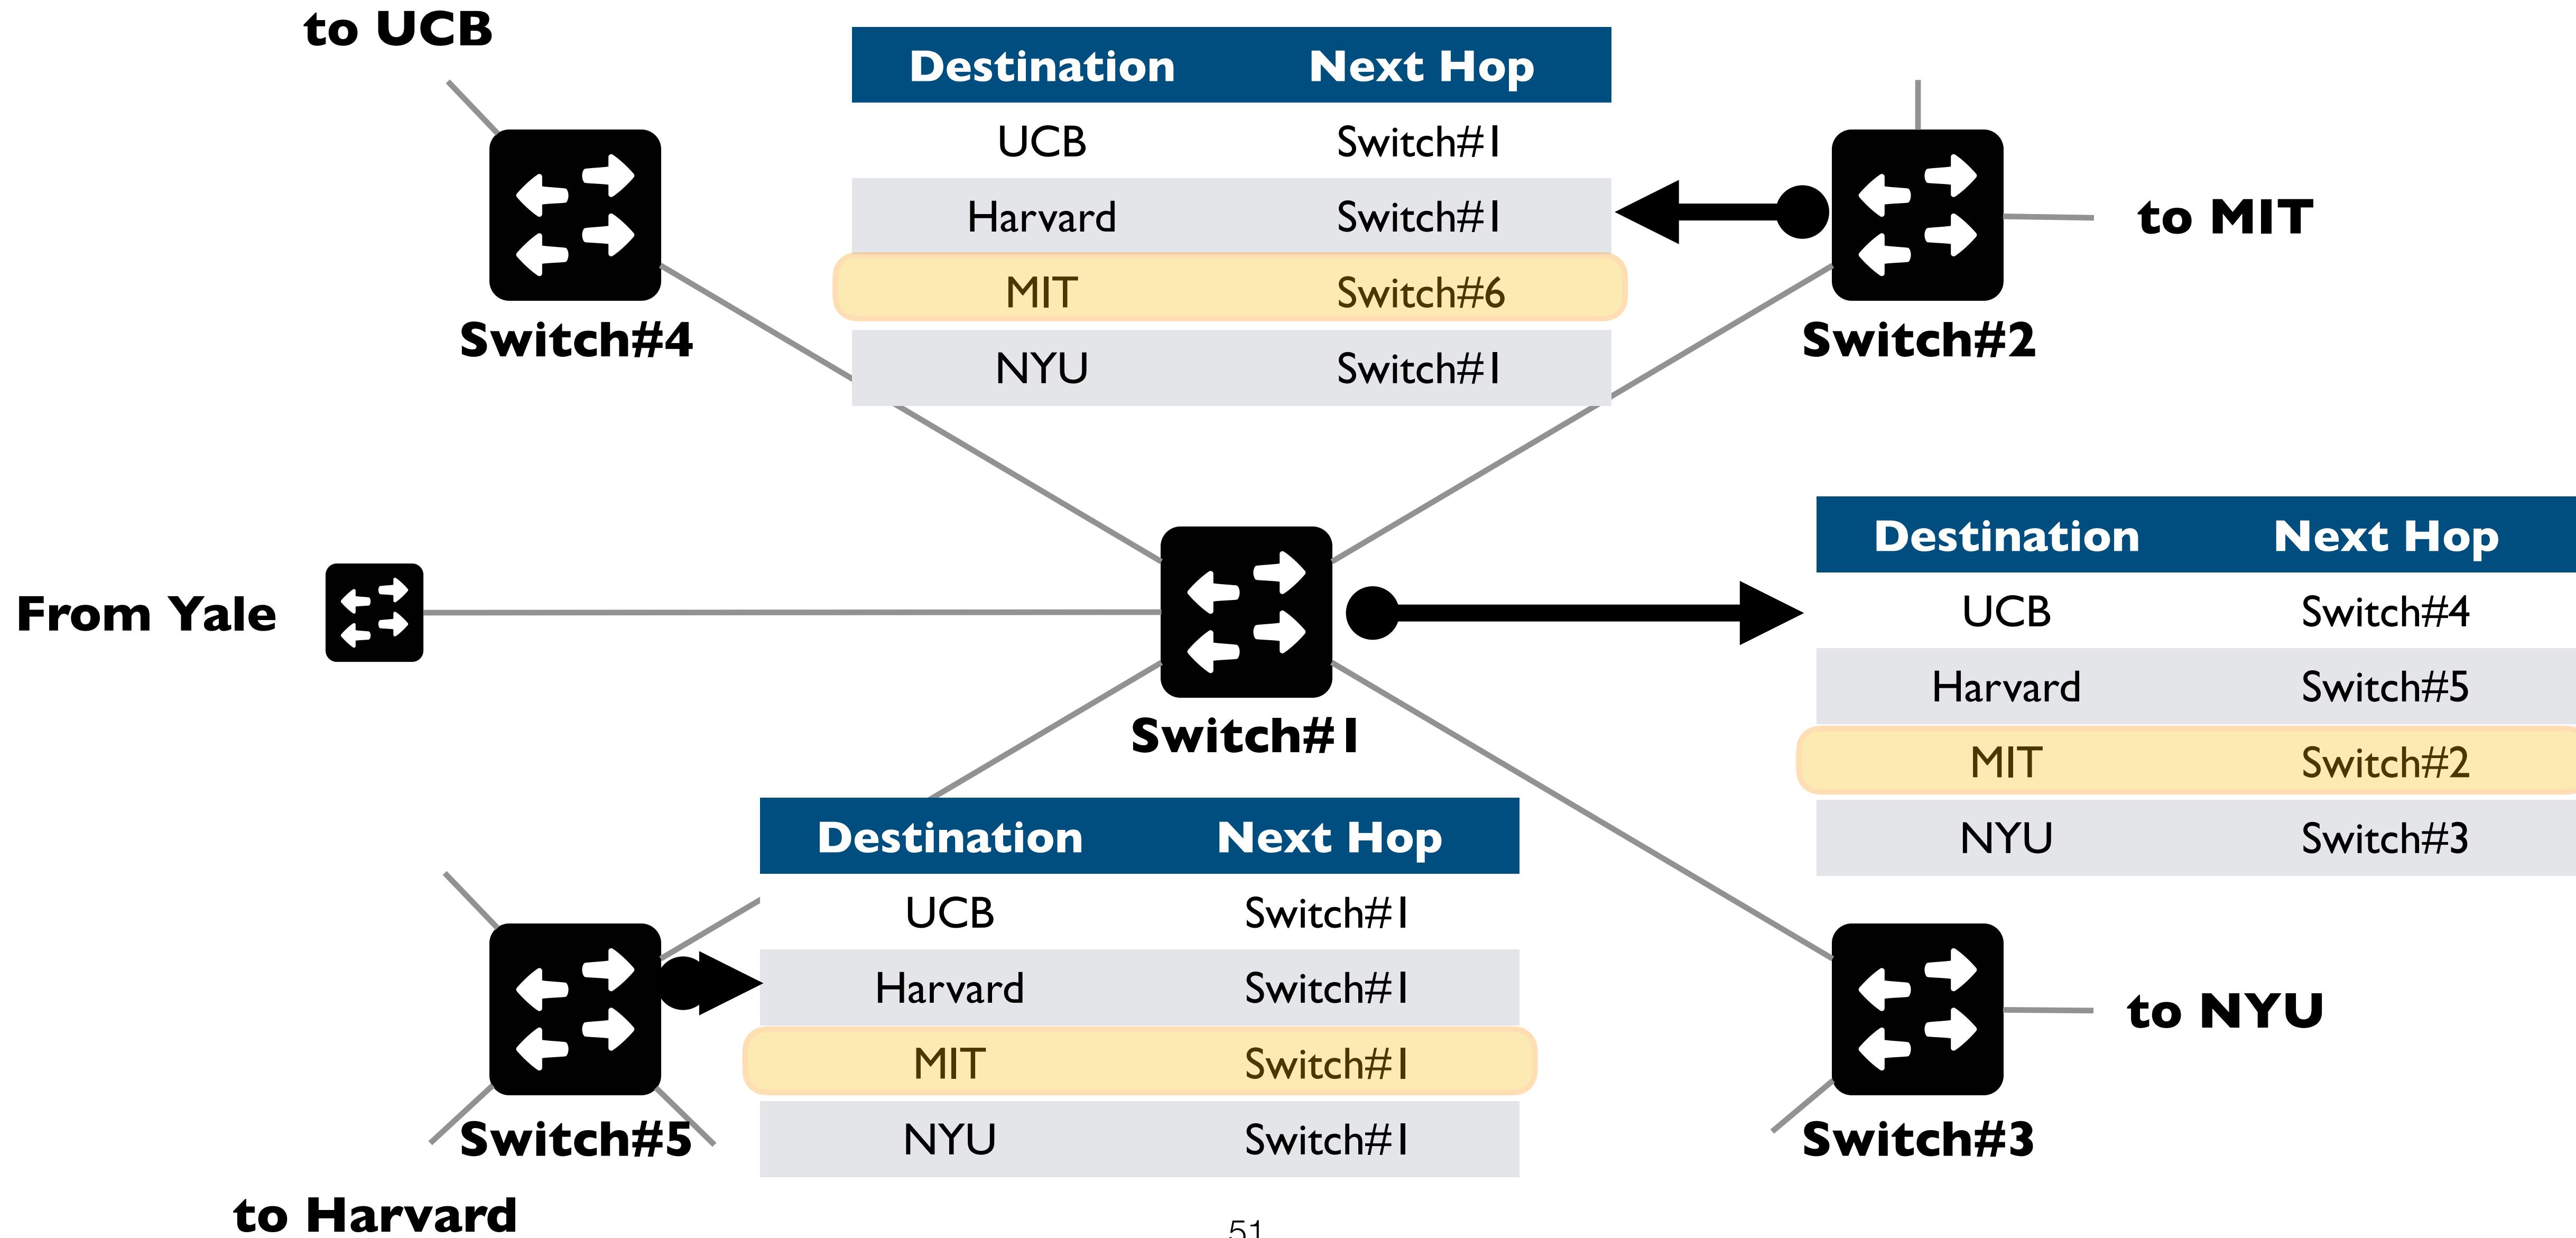
# Local vs. Global view of state

- **Local** routing state is the forwarding table in a single router
  - By itself, the state in a single router cannot be evaluated for validity
  - It must be evaluated in terms of the global context

# Local vs. Global view of state



# Local vs. Global view of state



# Local vs. Global view of state

# Local vs. Global view of state

- **Local** routing state is the forwarding table in a single router
  - By itself, the state in a single router cannot be evaluated for validity
  - It must be evaluated in terms of the global context

# Local vs. Global view of state

- **Local** routing state is the forwarding table in a single router
  - By itself, the state in a single router cannot be evaluated for validity
  - It must be evaluated in terms of the global context
- **Global** state refers to the collection of forwarding tables in each of the routers
  - Global state determines which paths packet take  
(Will discuss later where the state comes from)

# “Valid” Routing State

# “Valid” Routing State

- Global state is **“valid”** if it produces forwarding decisions that always deliver packets to their destinations

# “Valid” Routing State

- Global state is **“valid”** if it produces forwarding decisions that always deliver packets to their destinations
- **Goal of routing protocols:** compute valid state
  - But how can you tell if routing state is valid

# “Valid” Routing State

- Global state is **“valid”** if it produces forwarding decisions that always deliver packets to their destinations
- **Goal of routing protocols:** compute valid state
  - But how can you tell if routing state is valid
- Need a **succinct** correctness condition for routing
  - Suggestions?

# Necessary and Sufficient Condition

# Necessary and Sufficient Condition

- Global routing state is “valid” ***if and only if:***
  - There are no dead-ends (other than the destination)
  - There are no loops

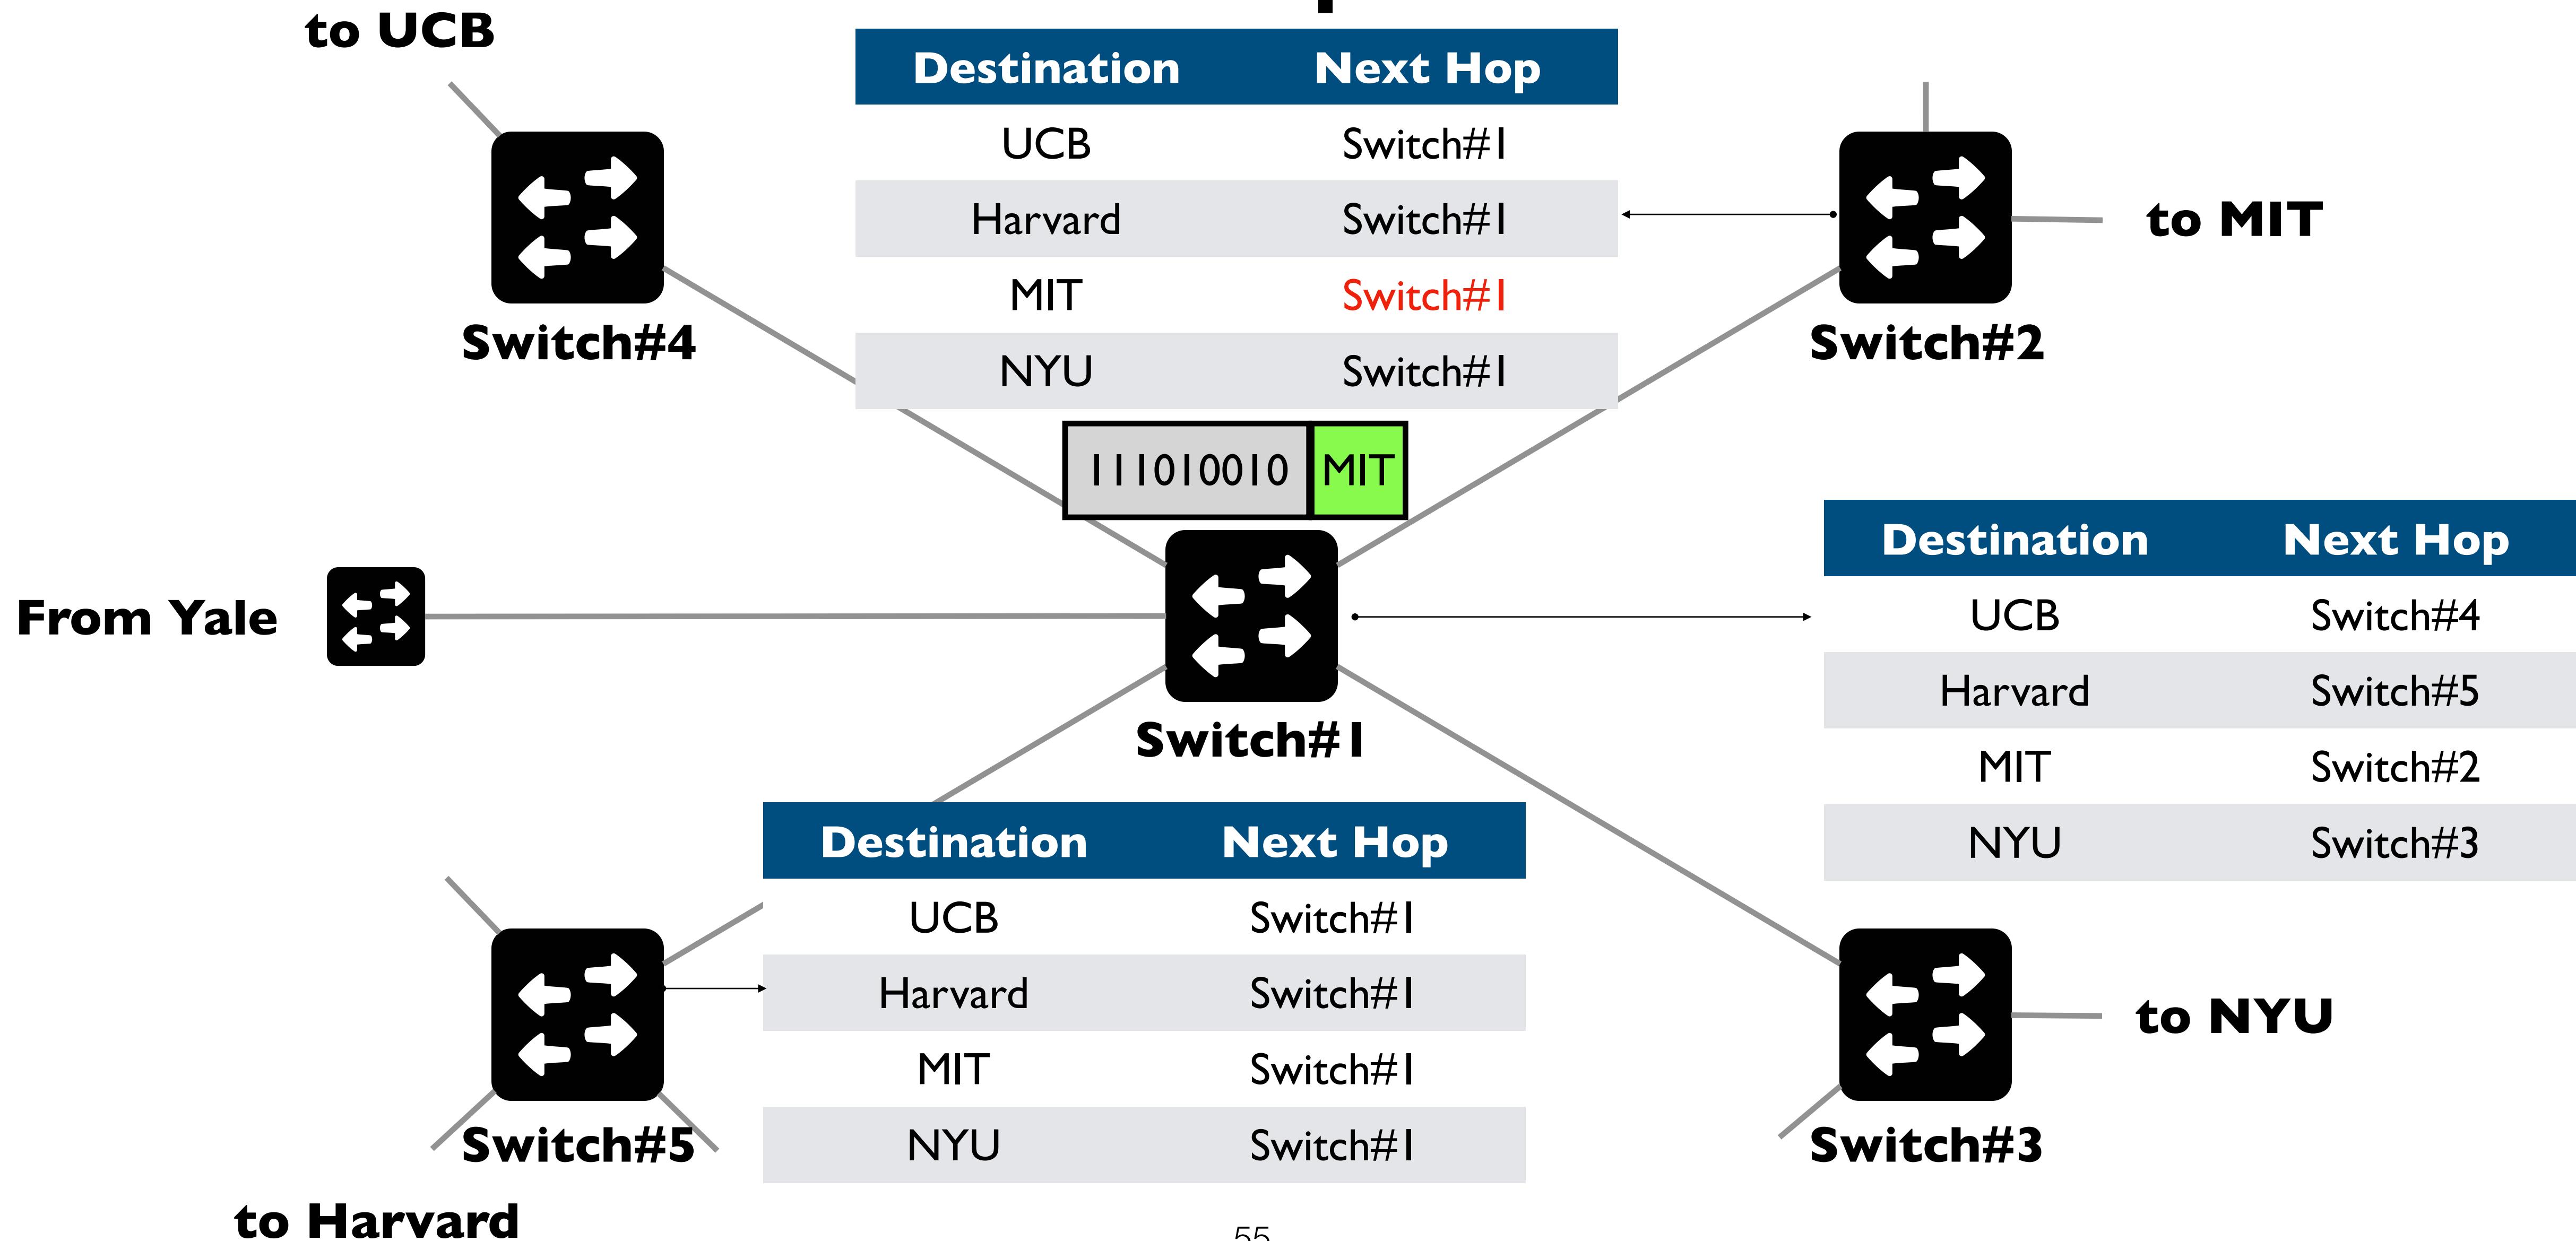
# Necessary and Sufficient Condition

- Global routing state is “valid” ***if and only if:***
  - There are no dead-ends (other than the destination)
  - There are no loops
- Dead-end: when there is no outgoing link (next-hop)
  - A packet arrives, but the forwarding decision does not yield any outgoing link

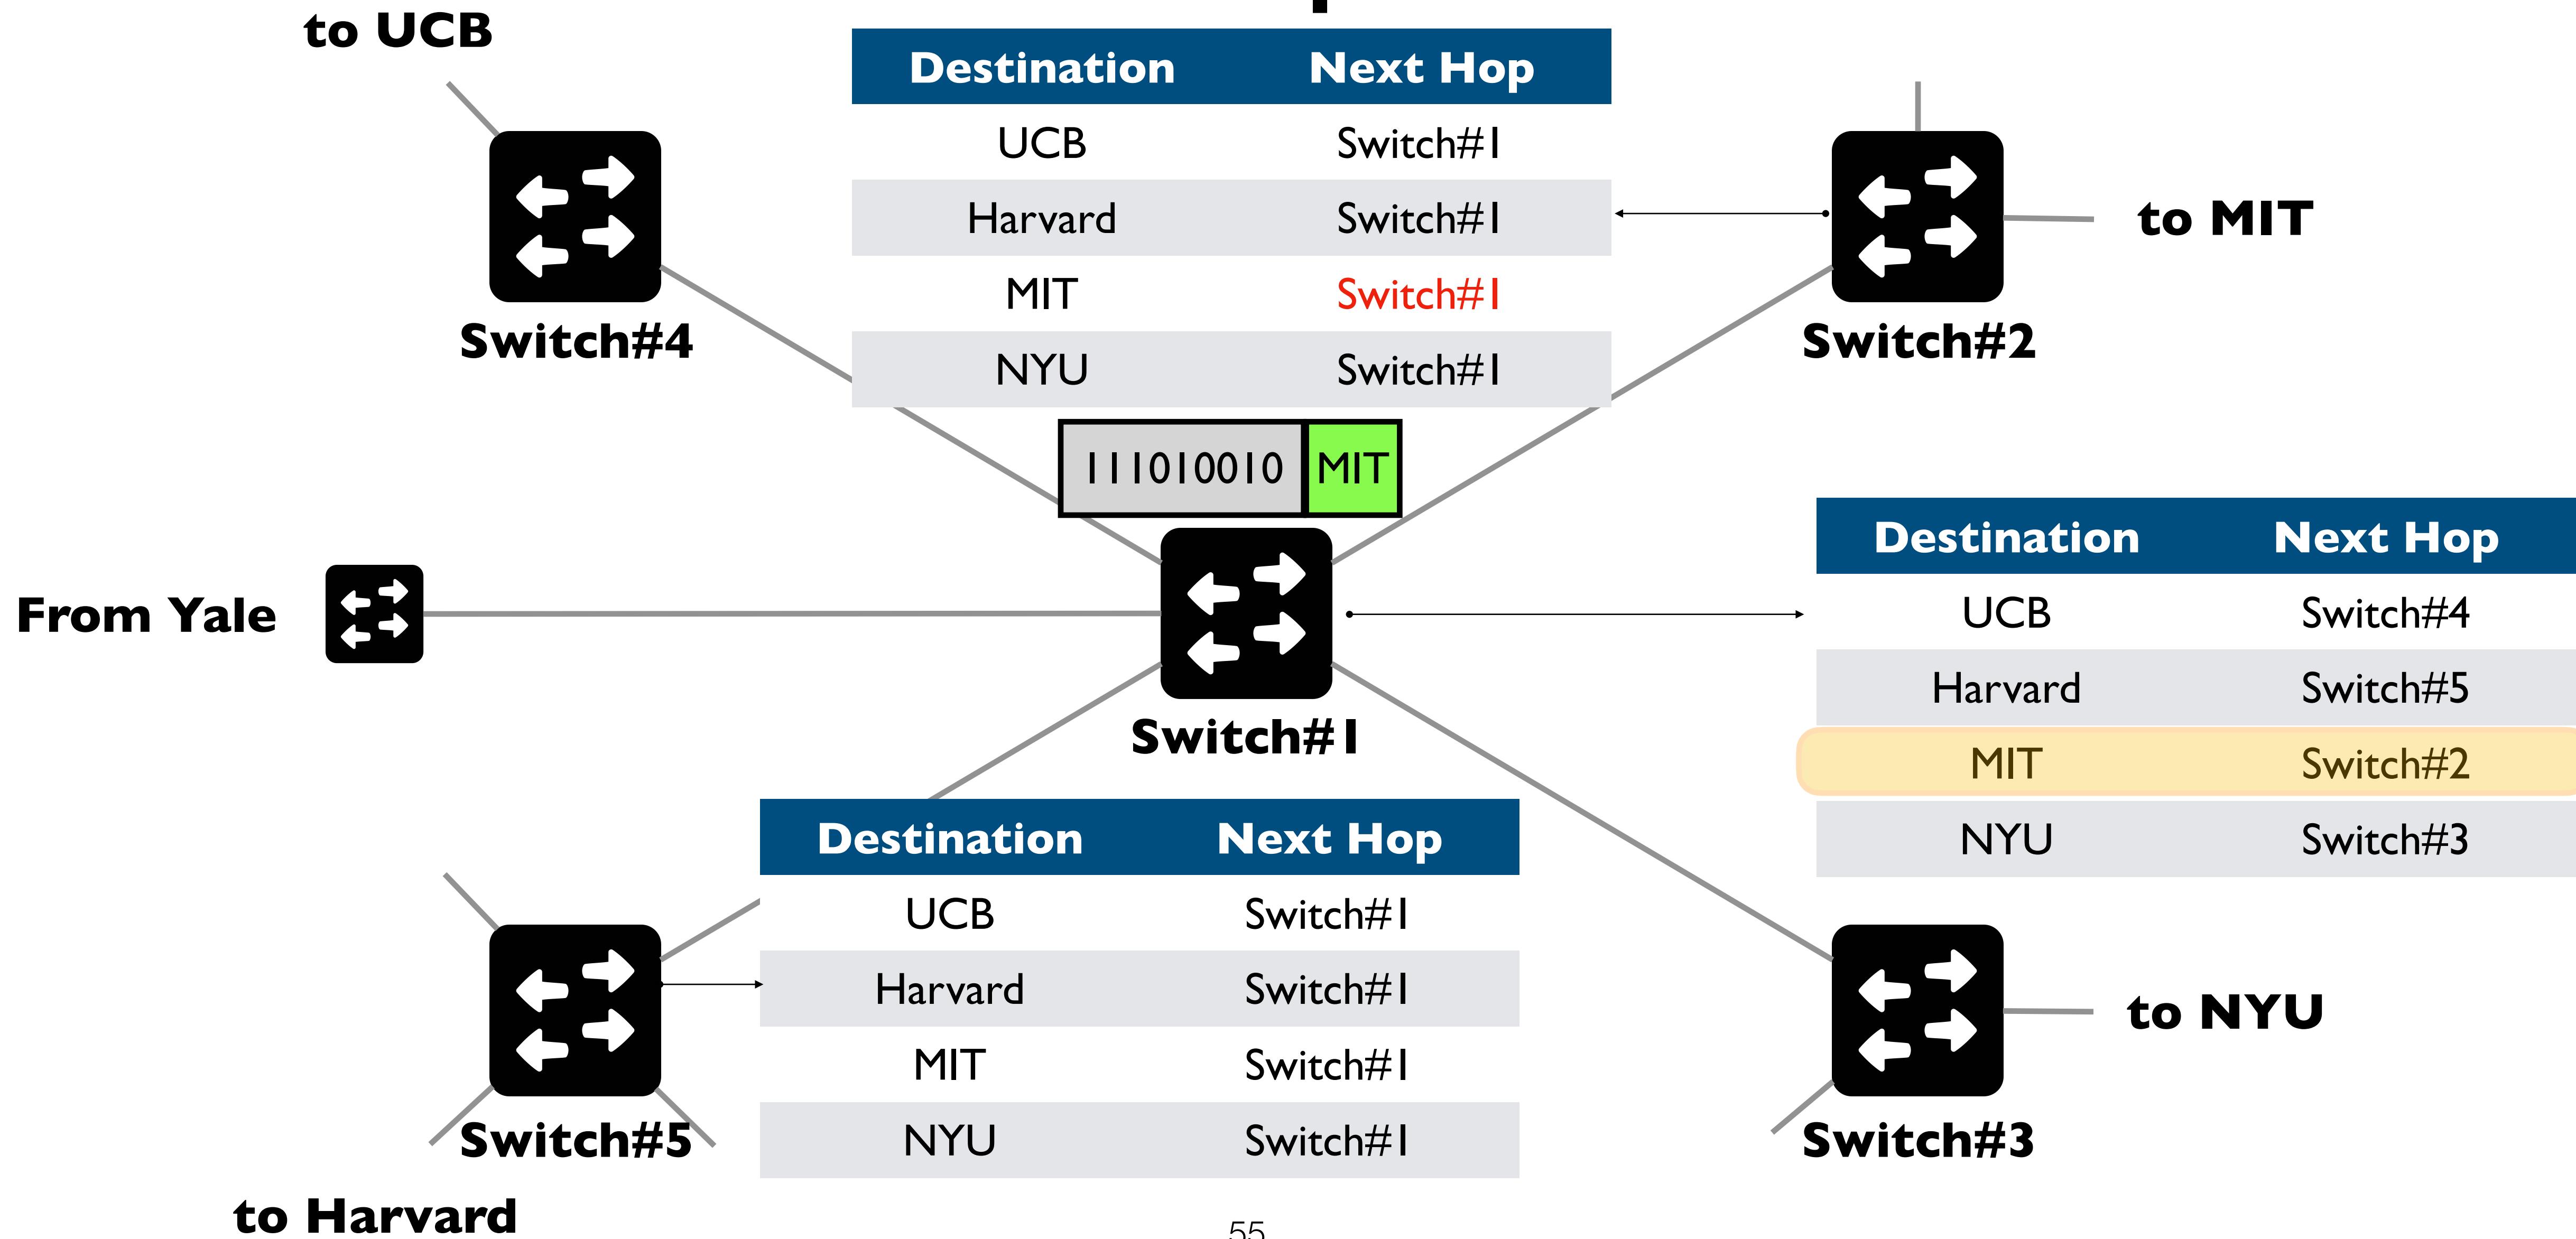
# Necessary and Sufficient Condition

- Global routing state is “valid” ***if and only if:***
  - There are no dead-ends (other than the destination)
  - There are no loops
- Dead-end: when there is no outgoing link (next-hop)
  - A packet arrives, but the forwarding decision does not yield any outgoing link
- Loop: when a packet cycles around the same set of nodes forever

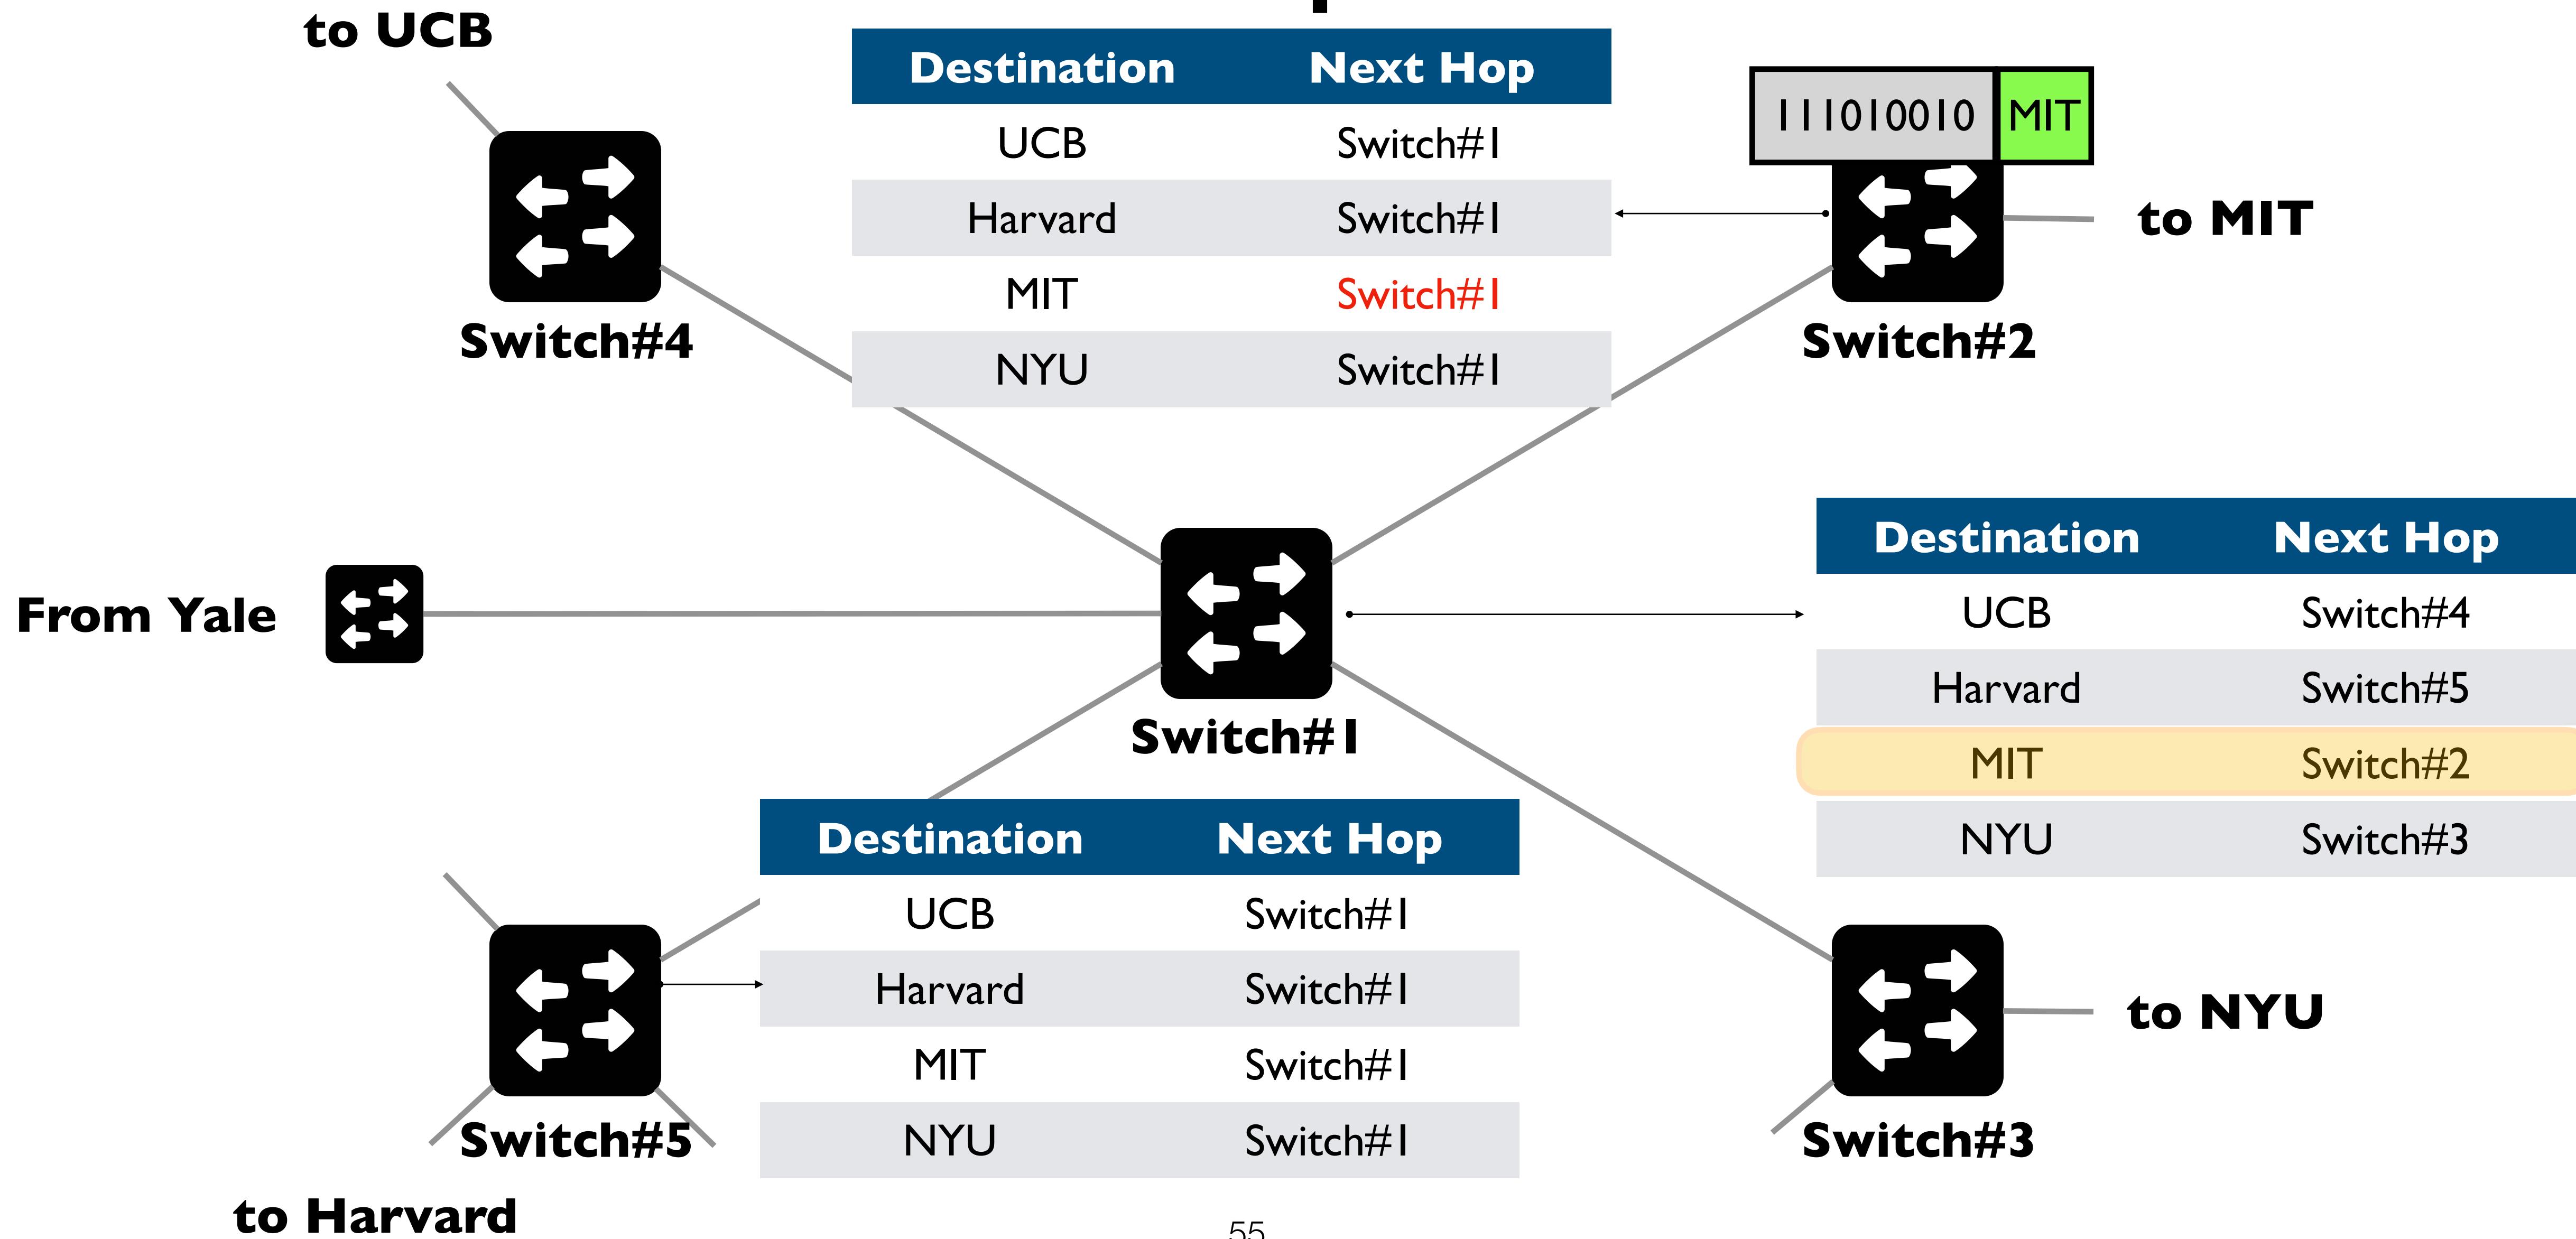
# Loops



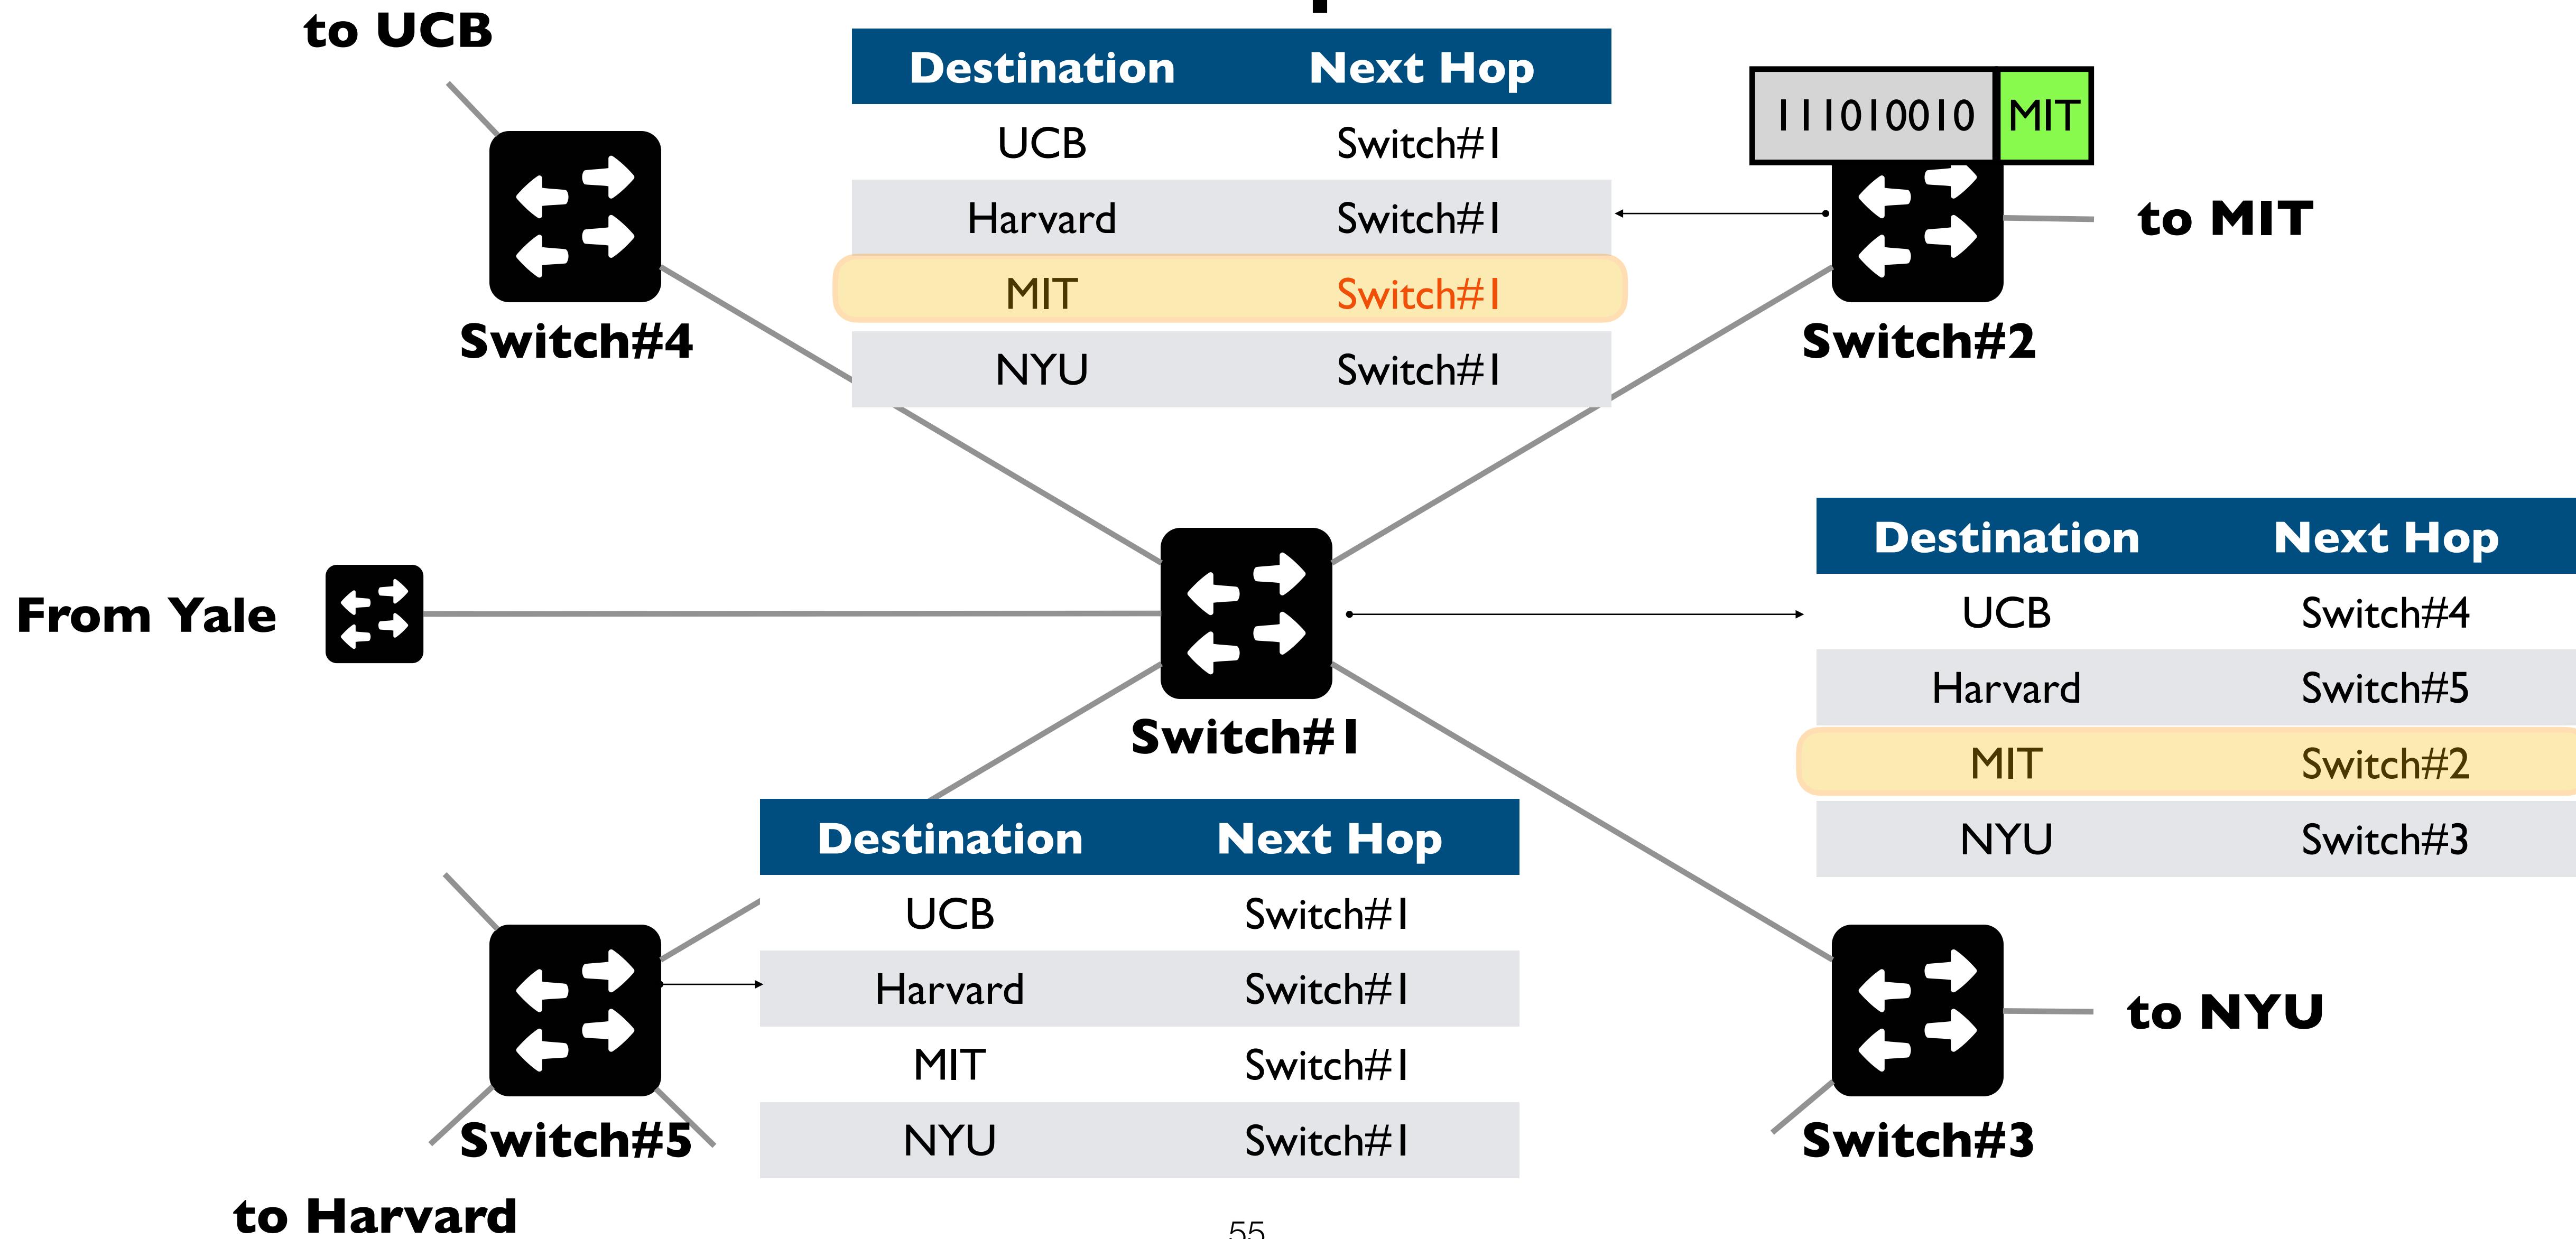
# Loops



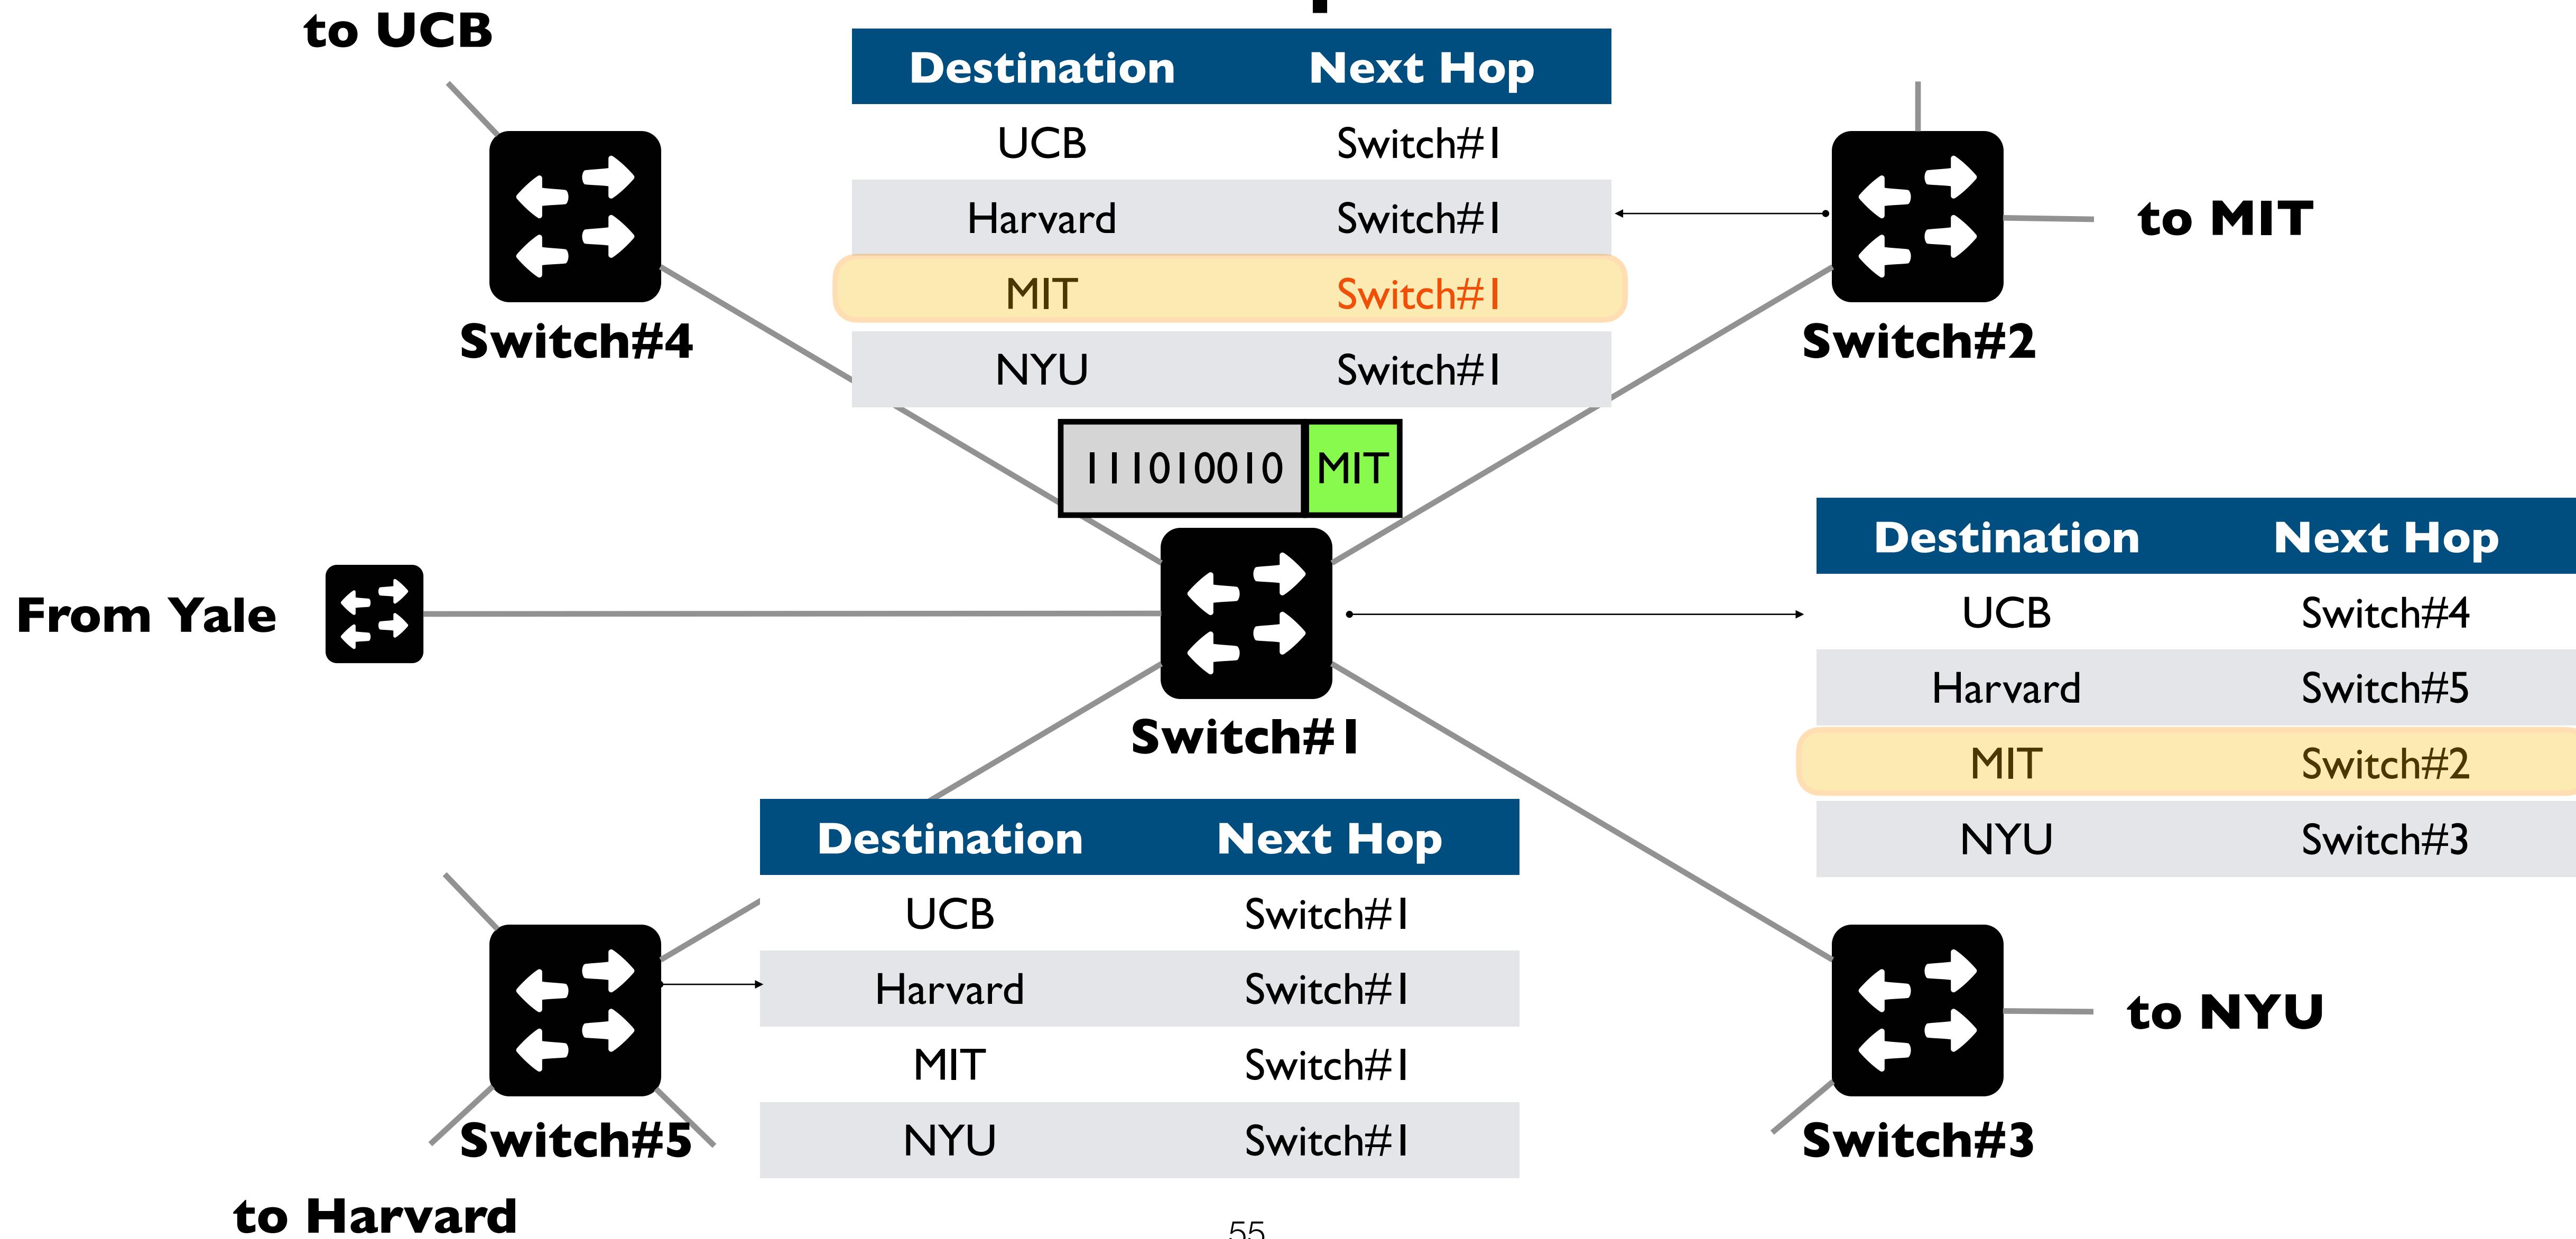
# Loops



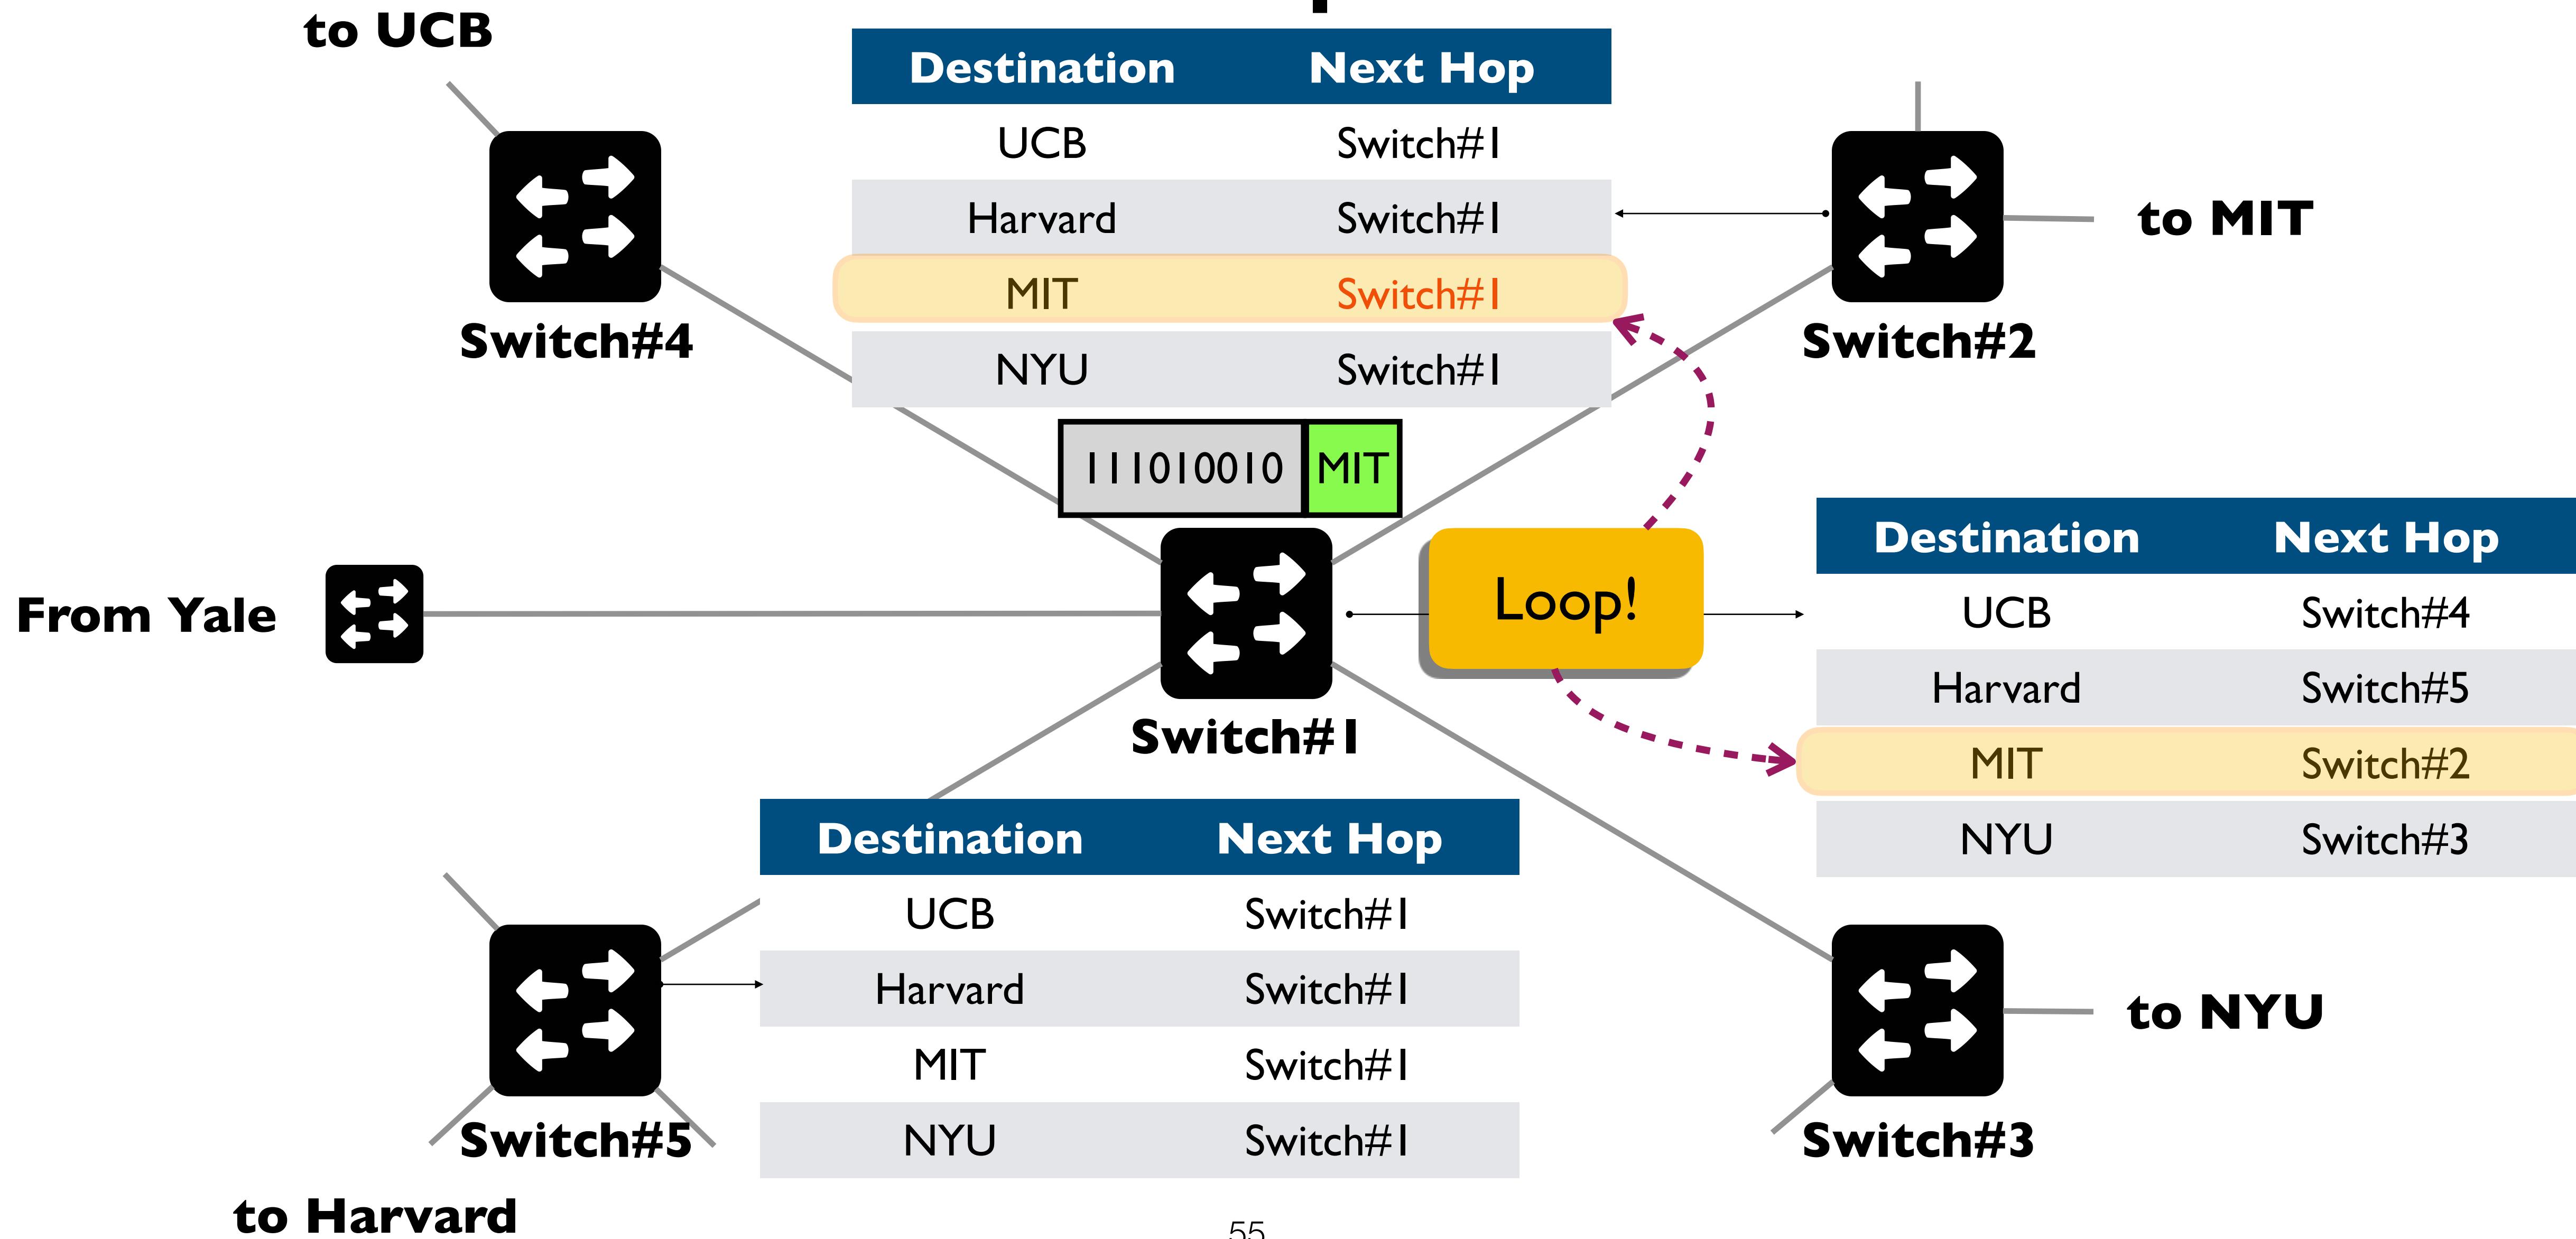
# Loops



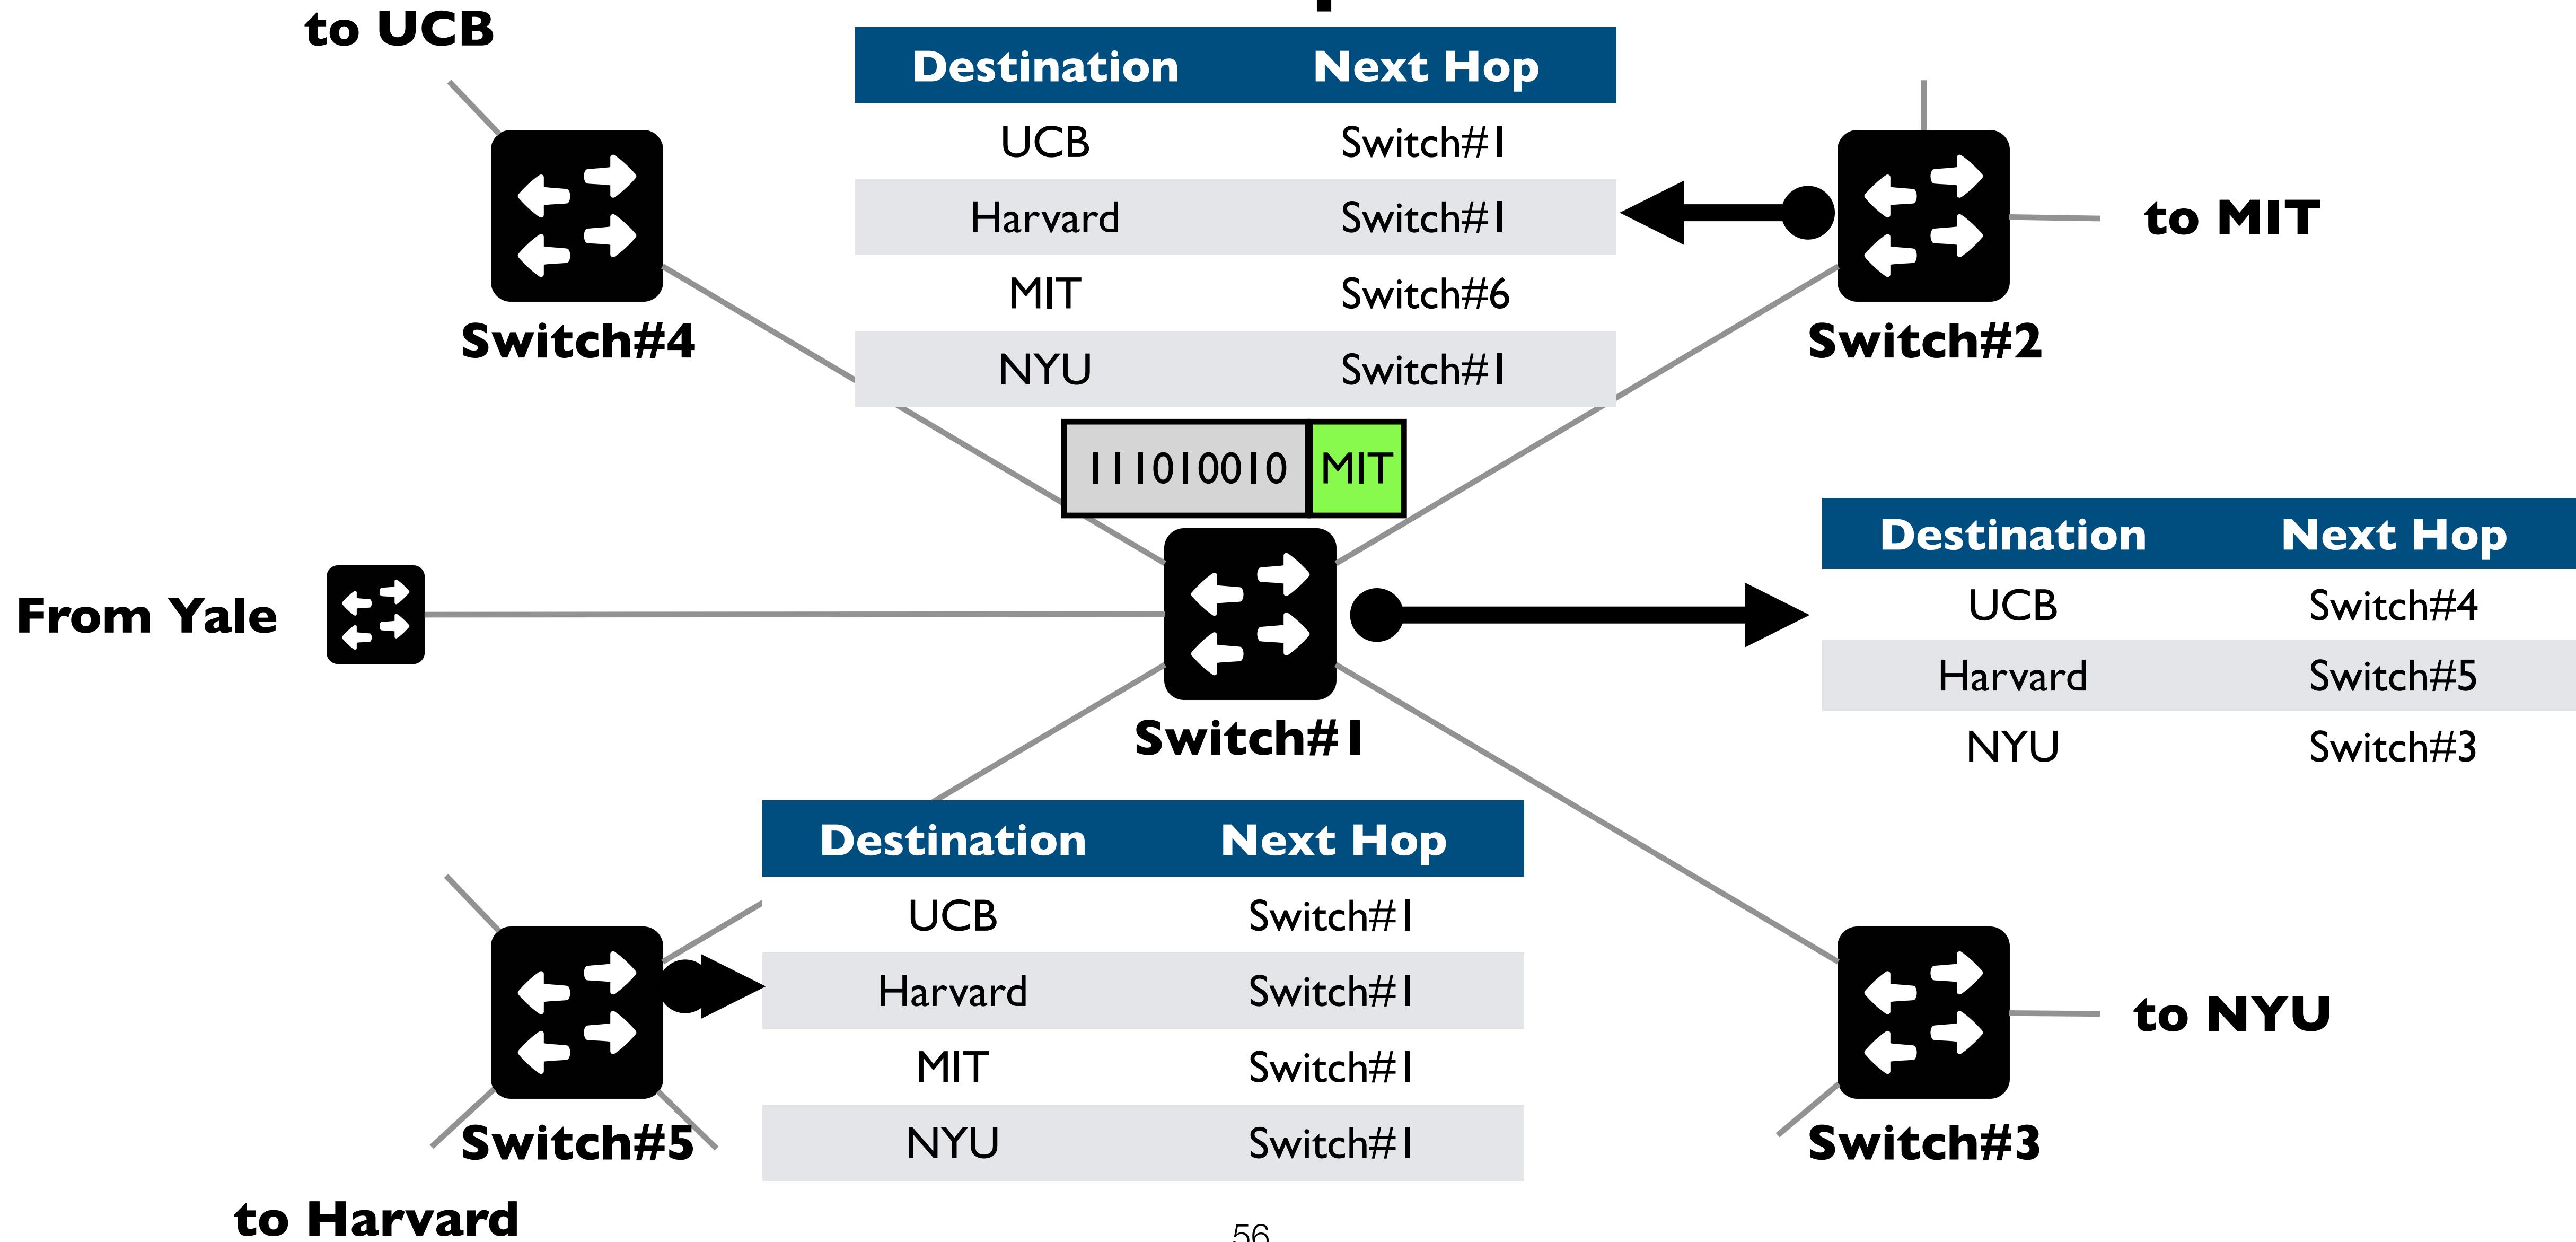
# Loops



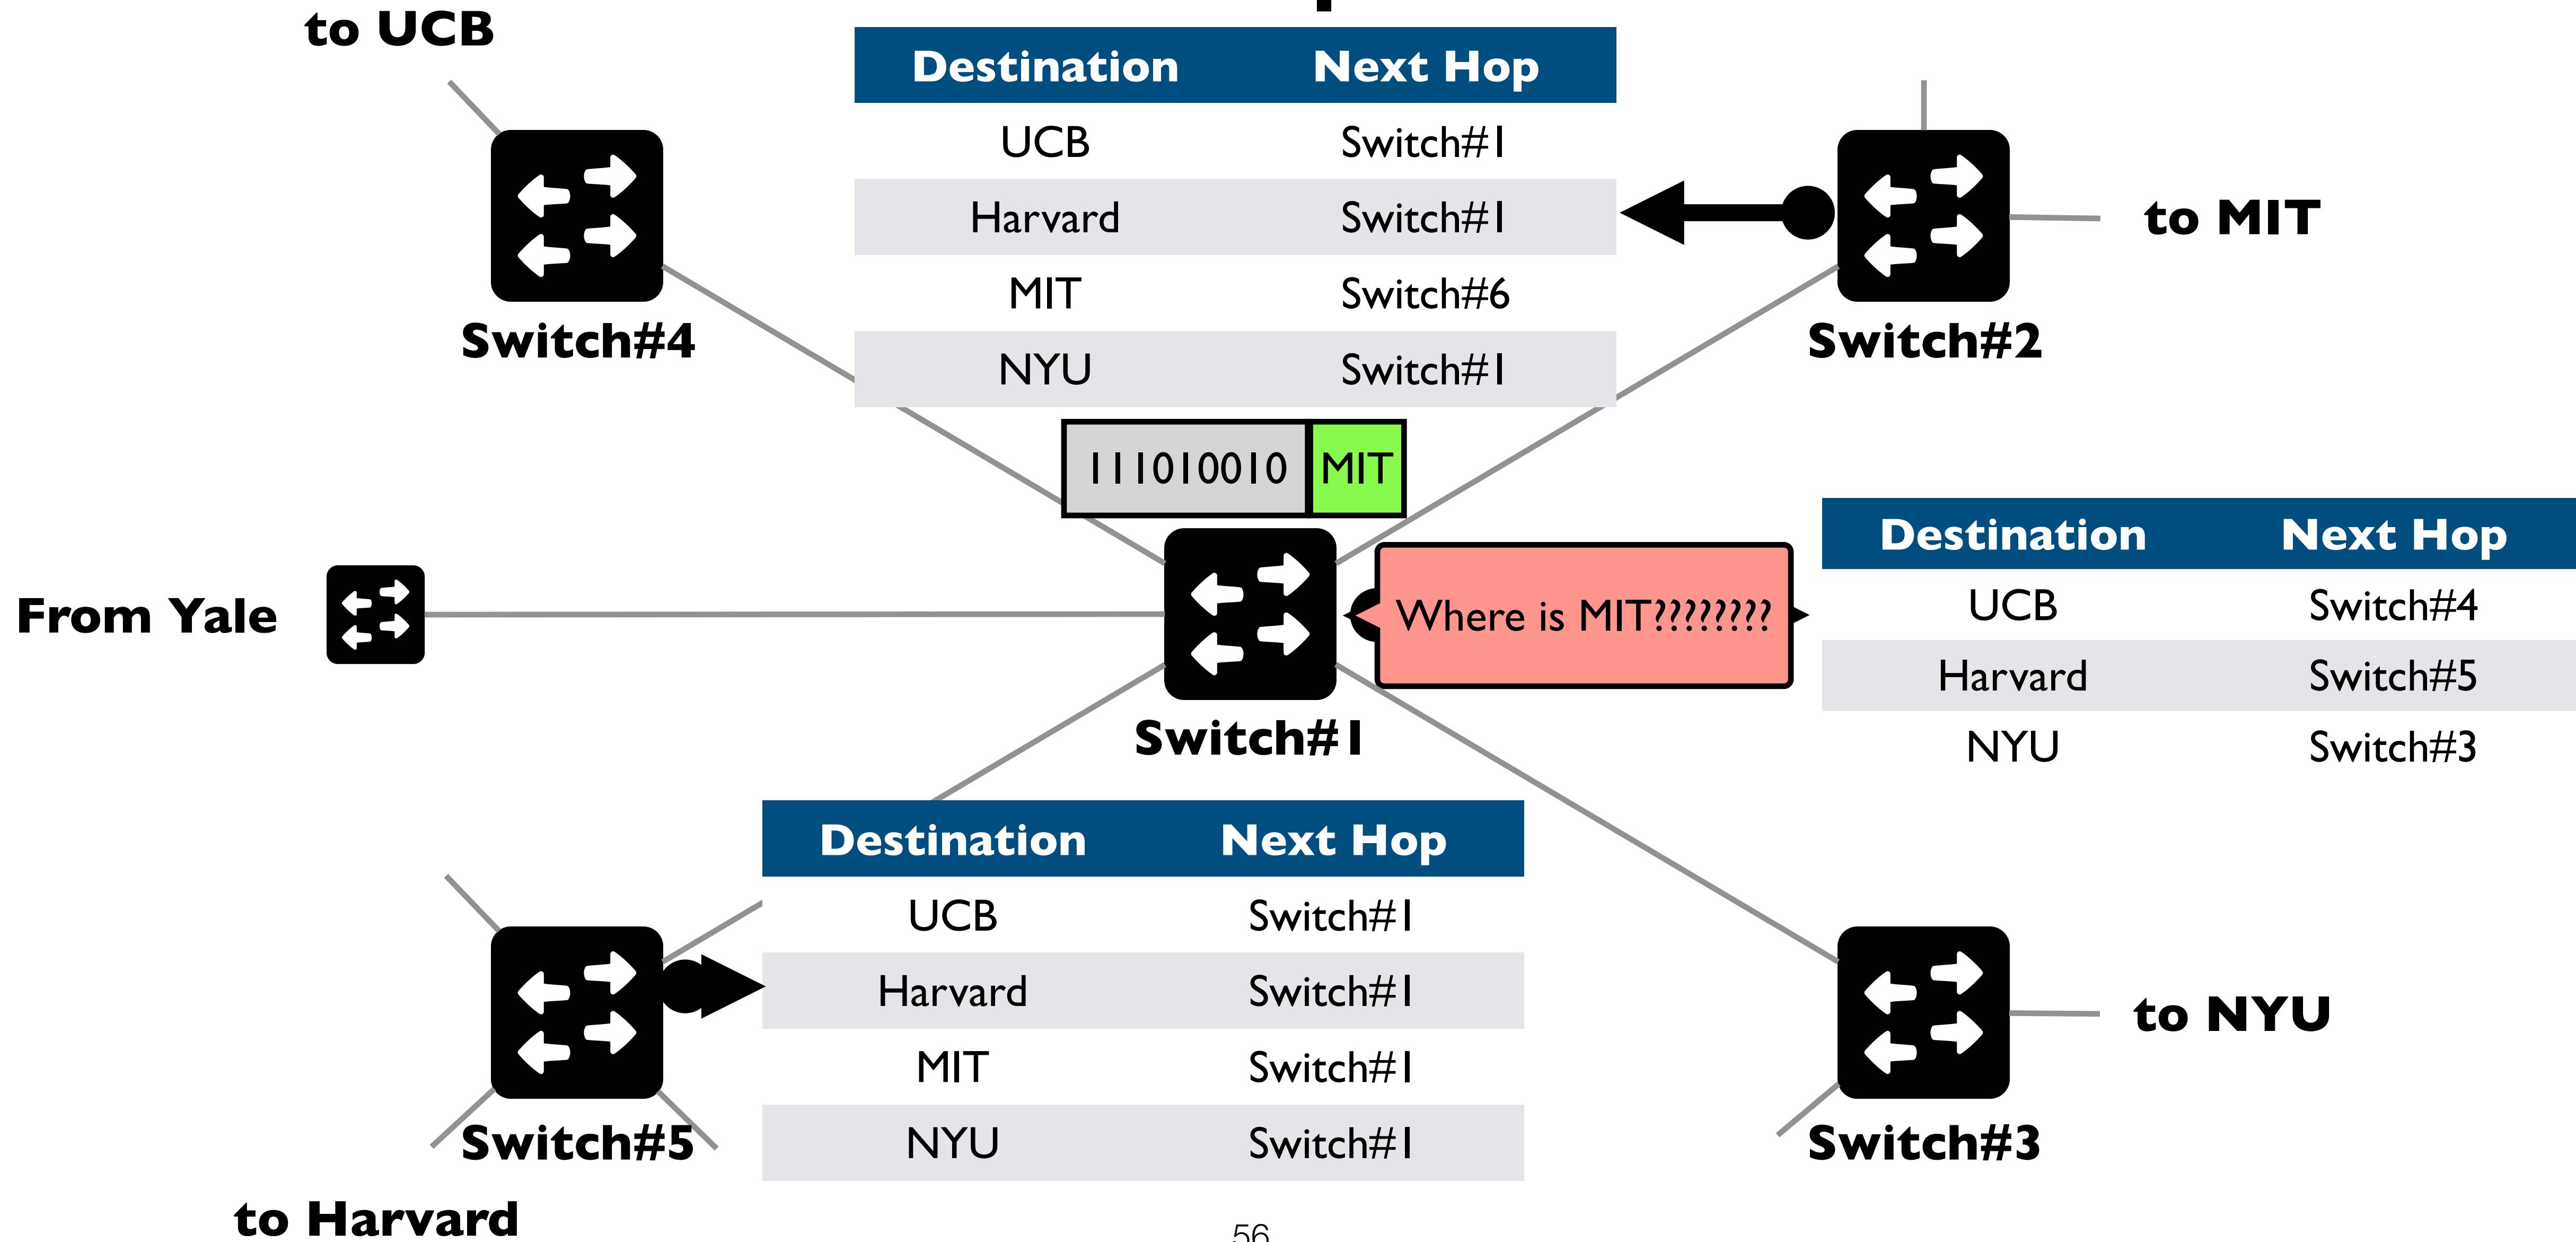
# Loops



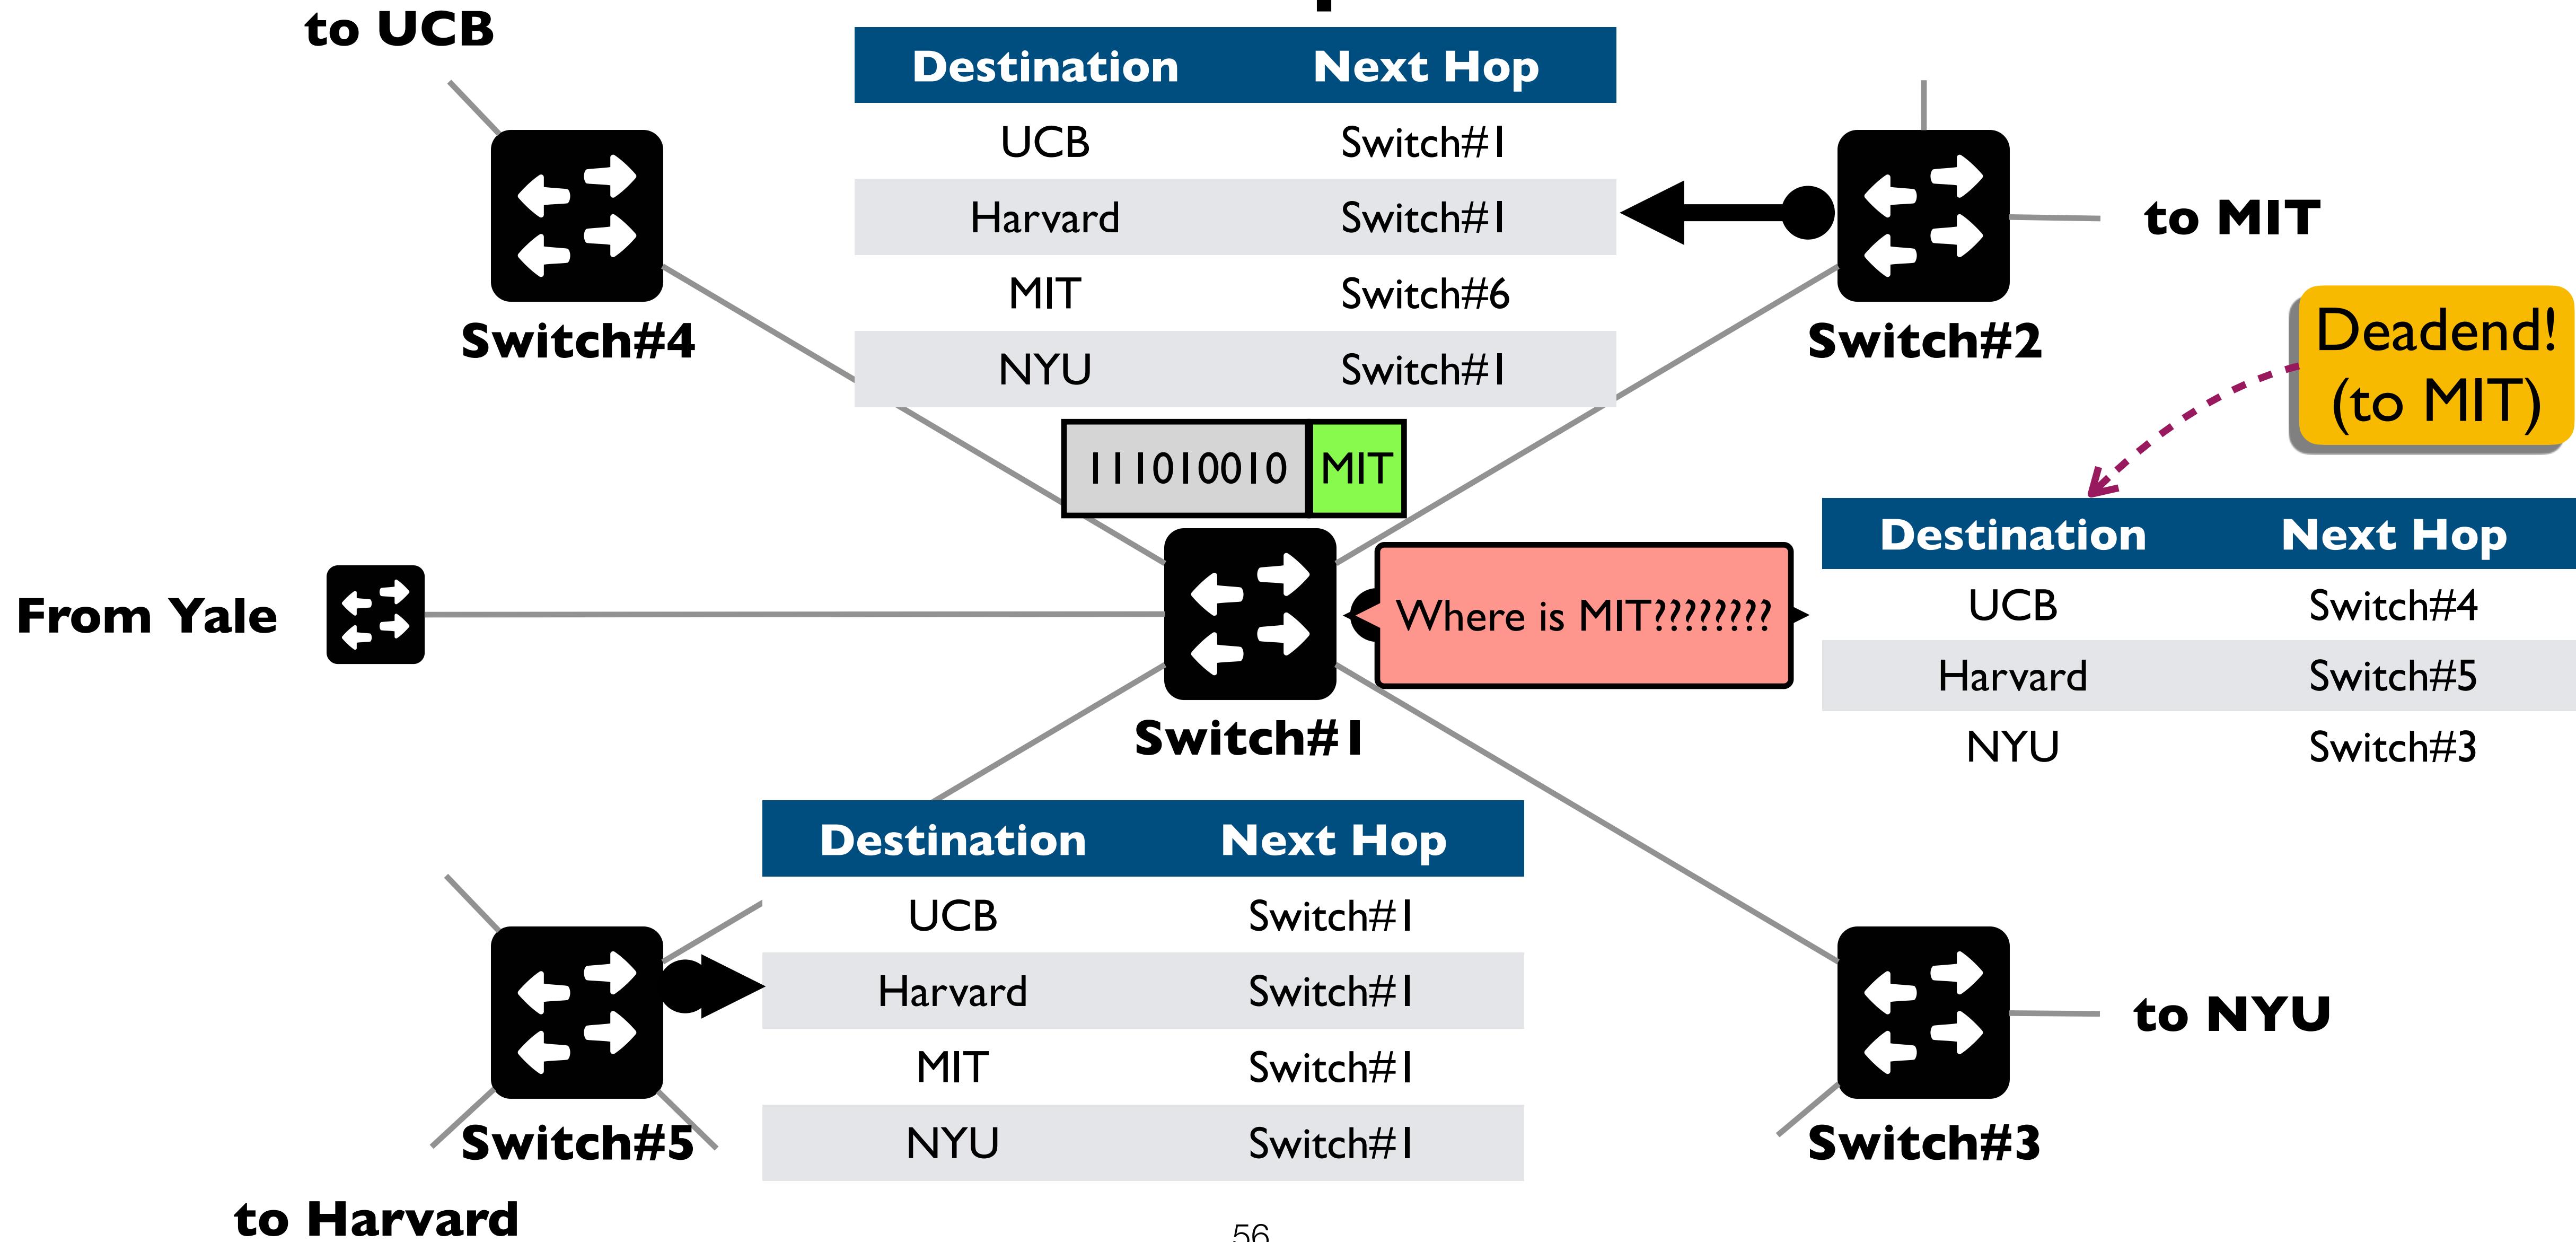
# Loops



# Loops



# Loops



# Necessary and Sufficient Condition

# Necessary and Sufficient Condition

- Global routing state is “valid” ***if and only if:***
  - There are no dead-ends (other than the destination)
  - There are no loops

# Necessary (“only if”): Easy

# Necessary (“only if”): Easy

- If you run into a dead-end before reaching destination, you'll **never** reach the destination

# Necessary (“only if”): Easy

- If you run into a dead-end before reaching destination, you'll **never** reach the destination
- If you run into a loop before reaching destination, you'll **never** reach the destination

# Sufficient (“if”)

# Sufficient (“if”)

- Assume no dead-ends, no loops

# Sufficient (“if”)

- Assume no dead-ends, no loops
- Packets may keep wandering, but without repeating link
  - If ever a packet enters the same switch from the same link, there must be a loop

# Sufficient (“if”)

- Assume no dead-ends, no loops
- Packets may keep wandering, but without repeating link
  - If ever a packet enters the same switch from the same link, there must be a loop
- Only a finite number of possible links for it to visit
  - It cannot keep wandering forever without looping
  - Must eventually hit destination!

# Checking Validity of Routing State

# Checking Validity of Routing State

- Focus only on a single destination
  - Ignore all other routing state

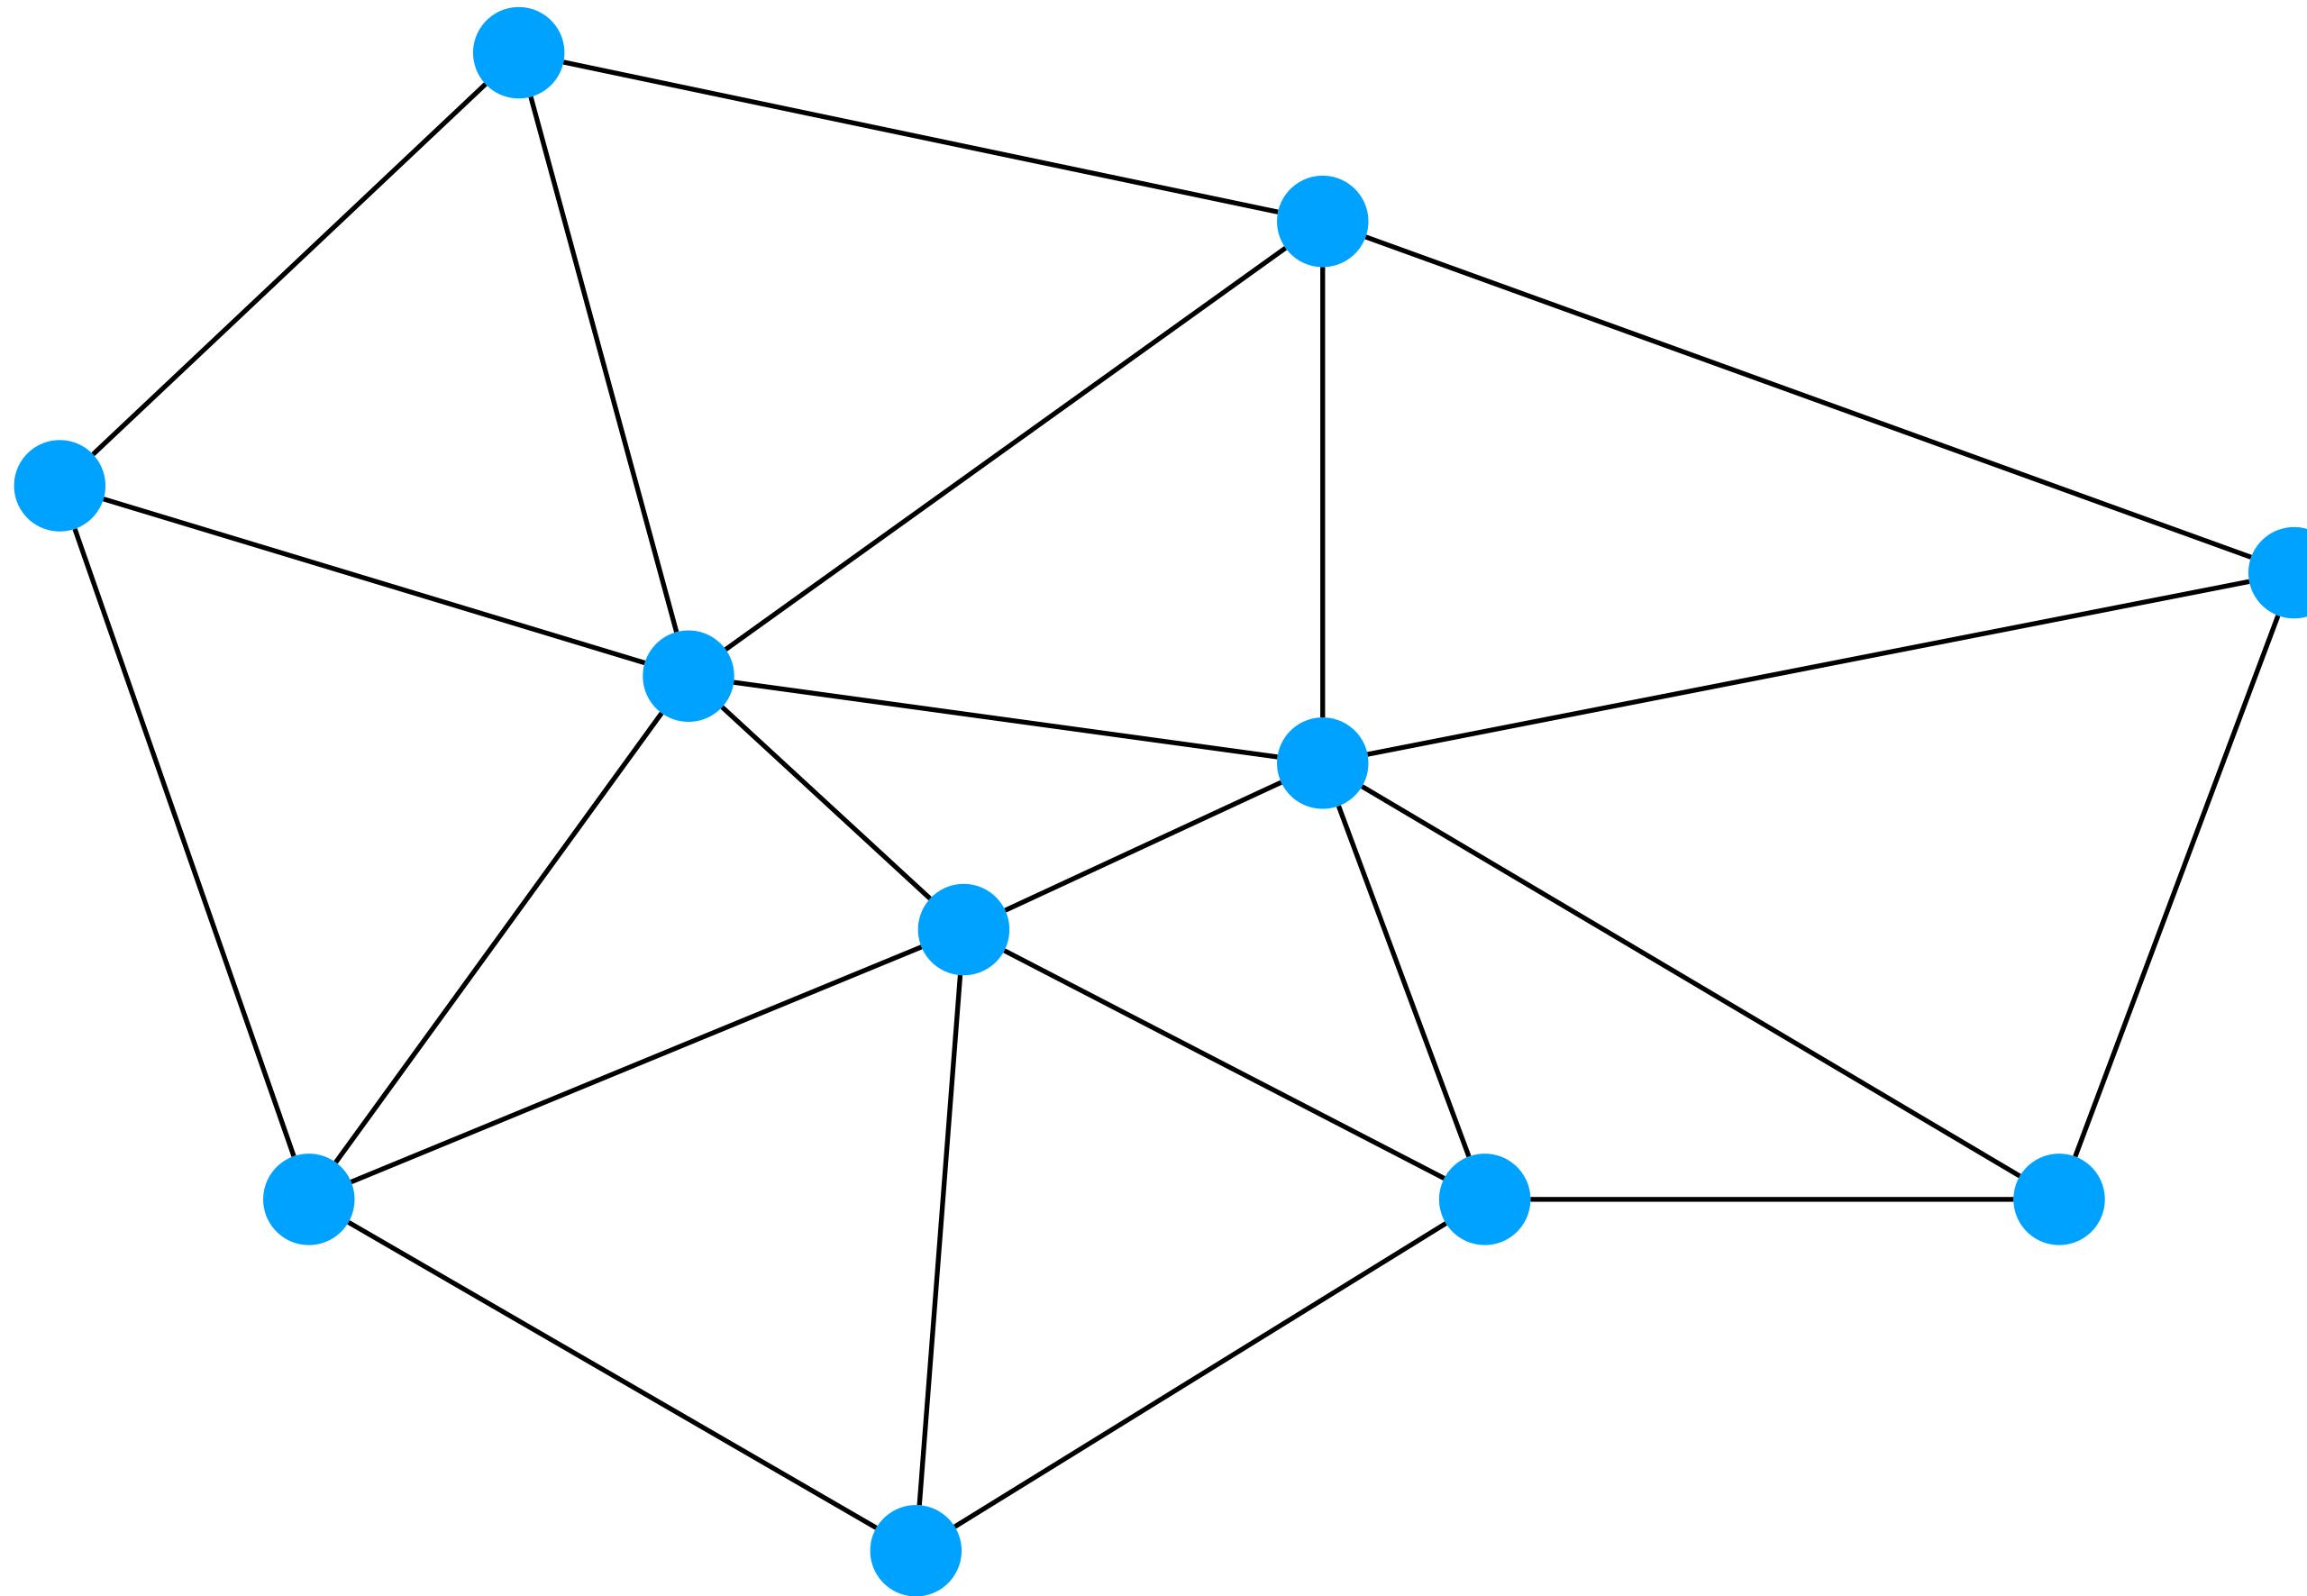
# Checking Validity of Routing State

- Focus only on a single destination
  - Ignore all other routing state
- Mark outgoing link to that destination (“next hop”) with arrow
  - There is only one at each node

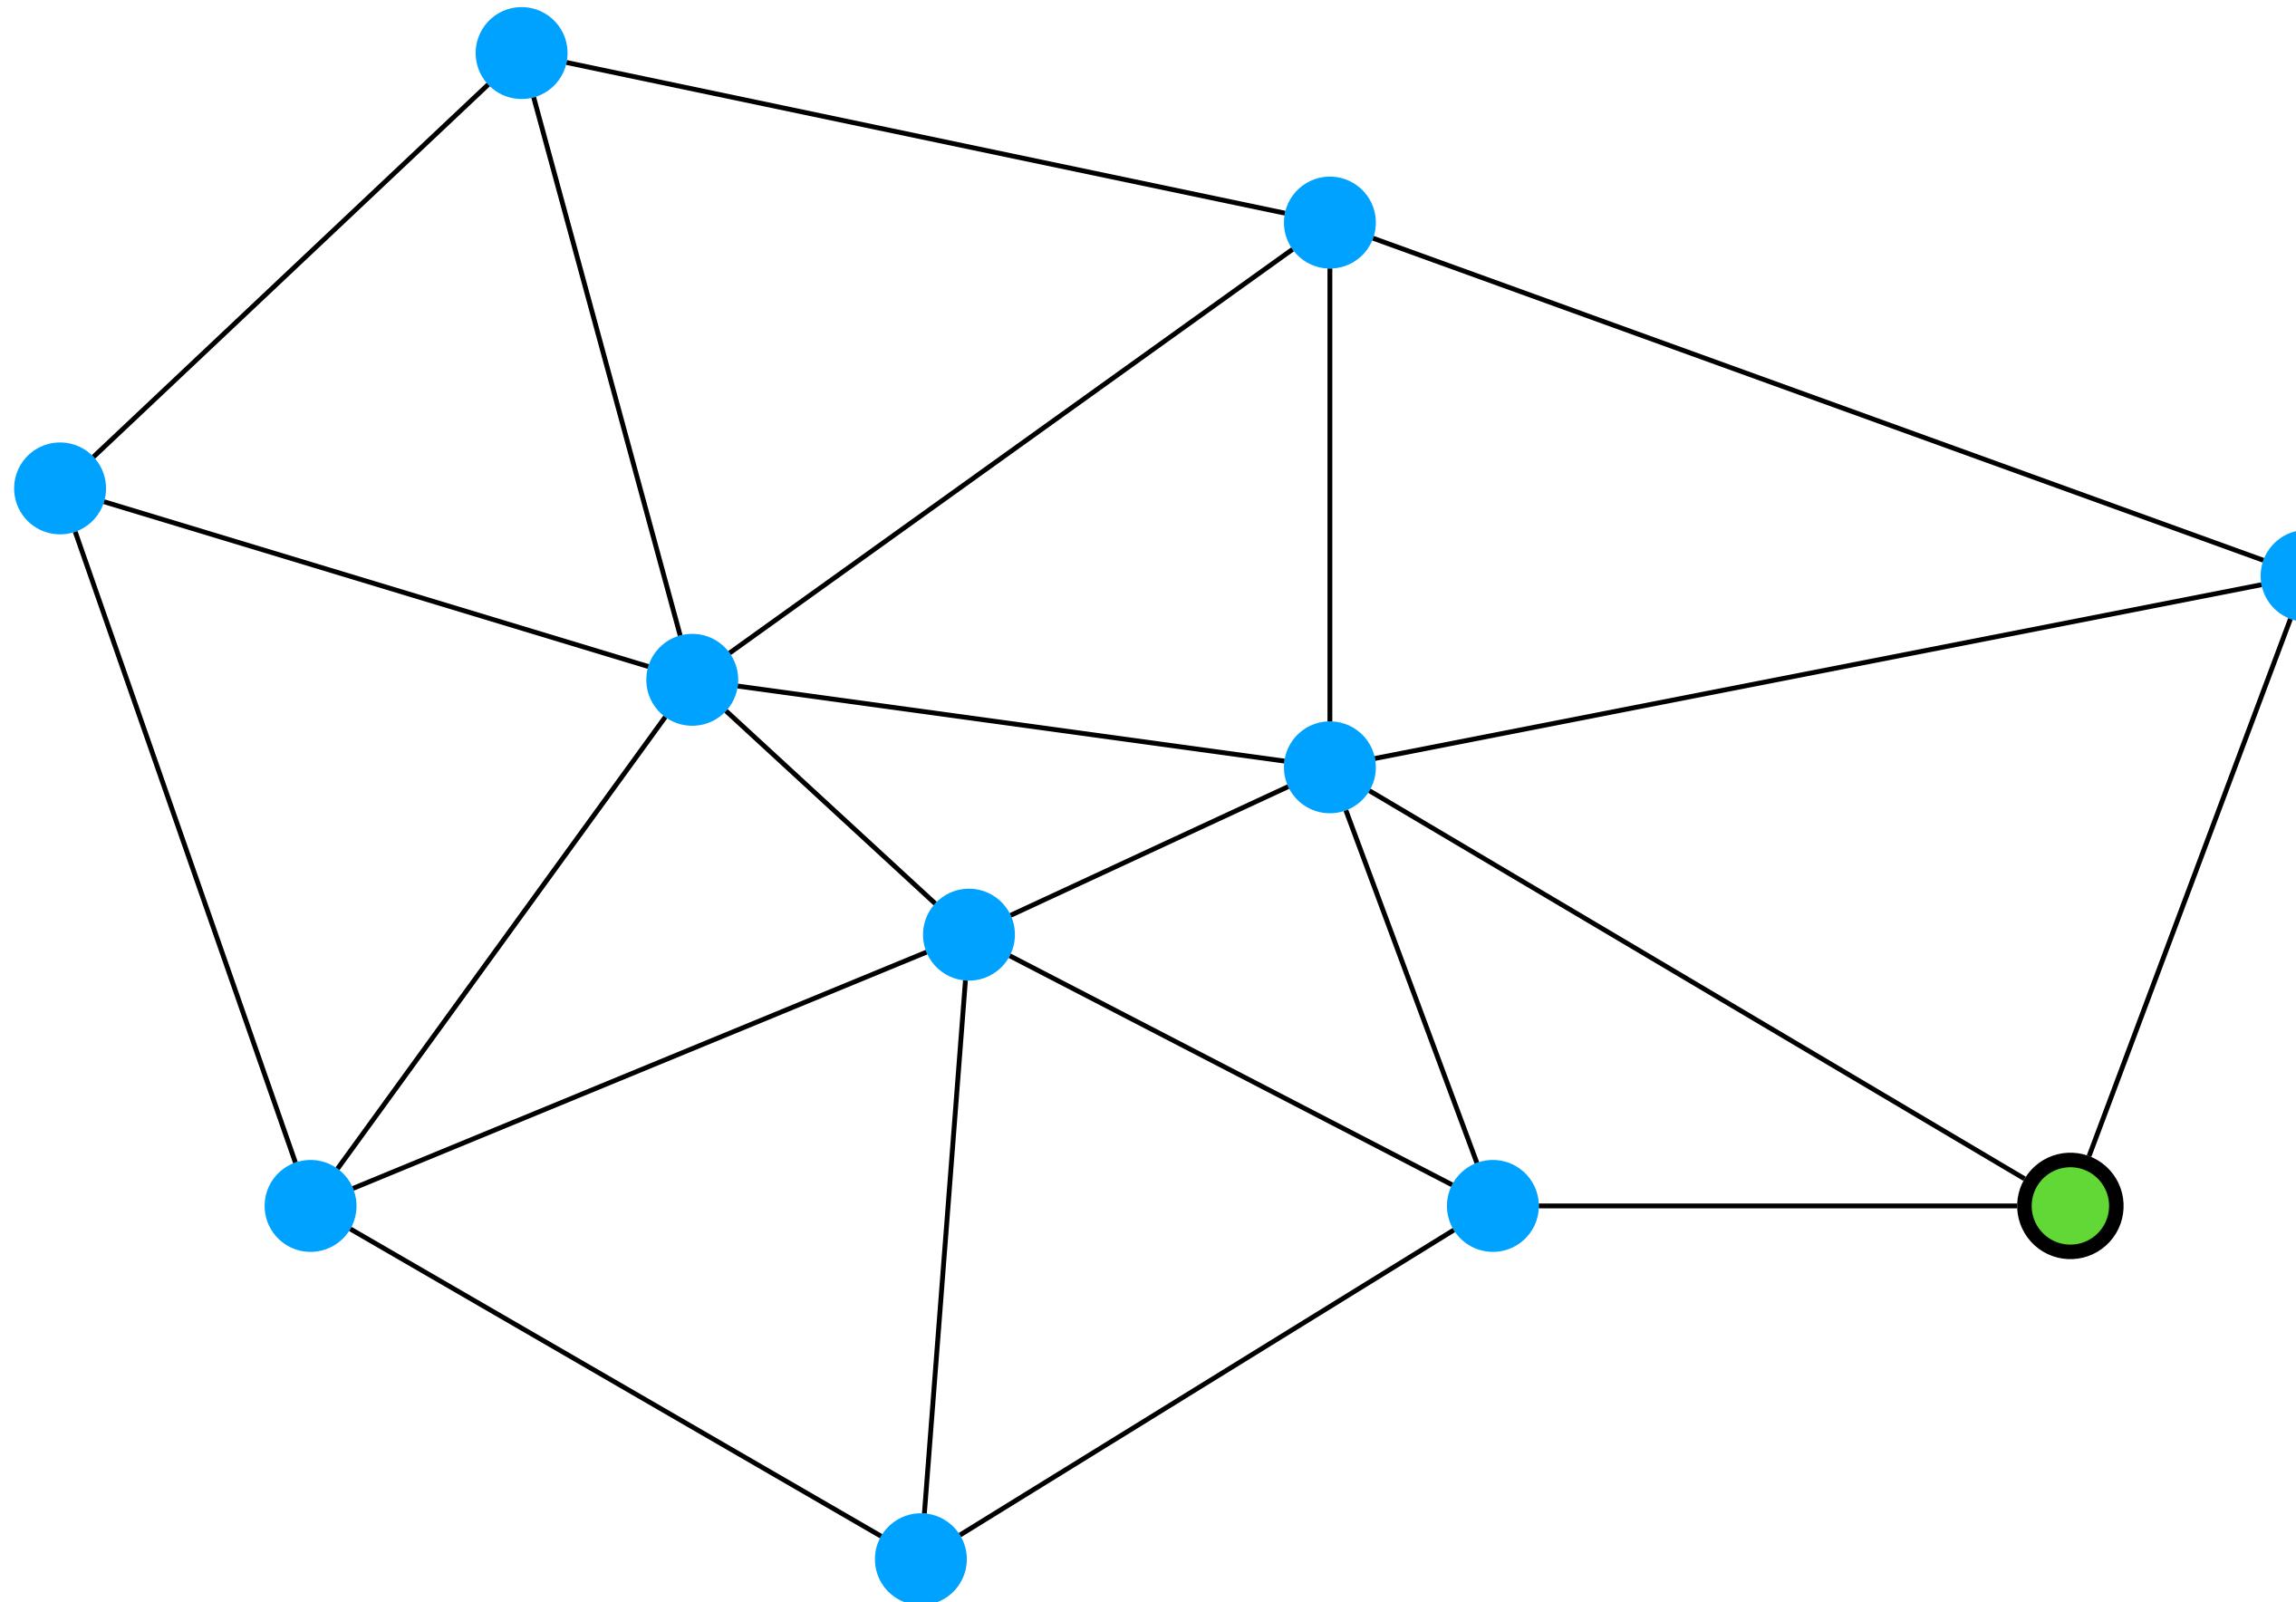
# Checking Validity of Routing State

- Focus only on a single destination
  - Ignore all other routing state
- Mark outgoing link to that destination (“next hop”) with arrow
  - There is only one at each node
- Eliminate all links with no arrows
  - Look at what's left...

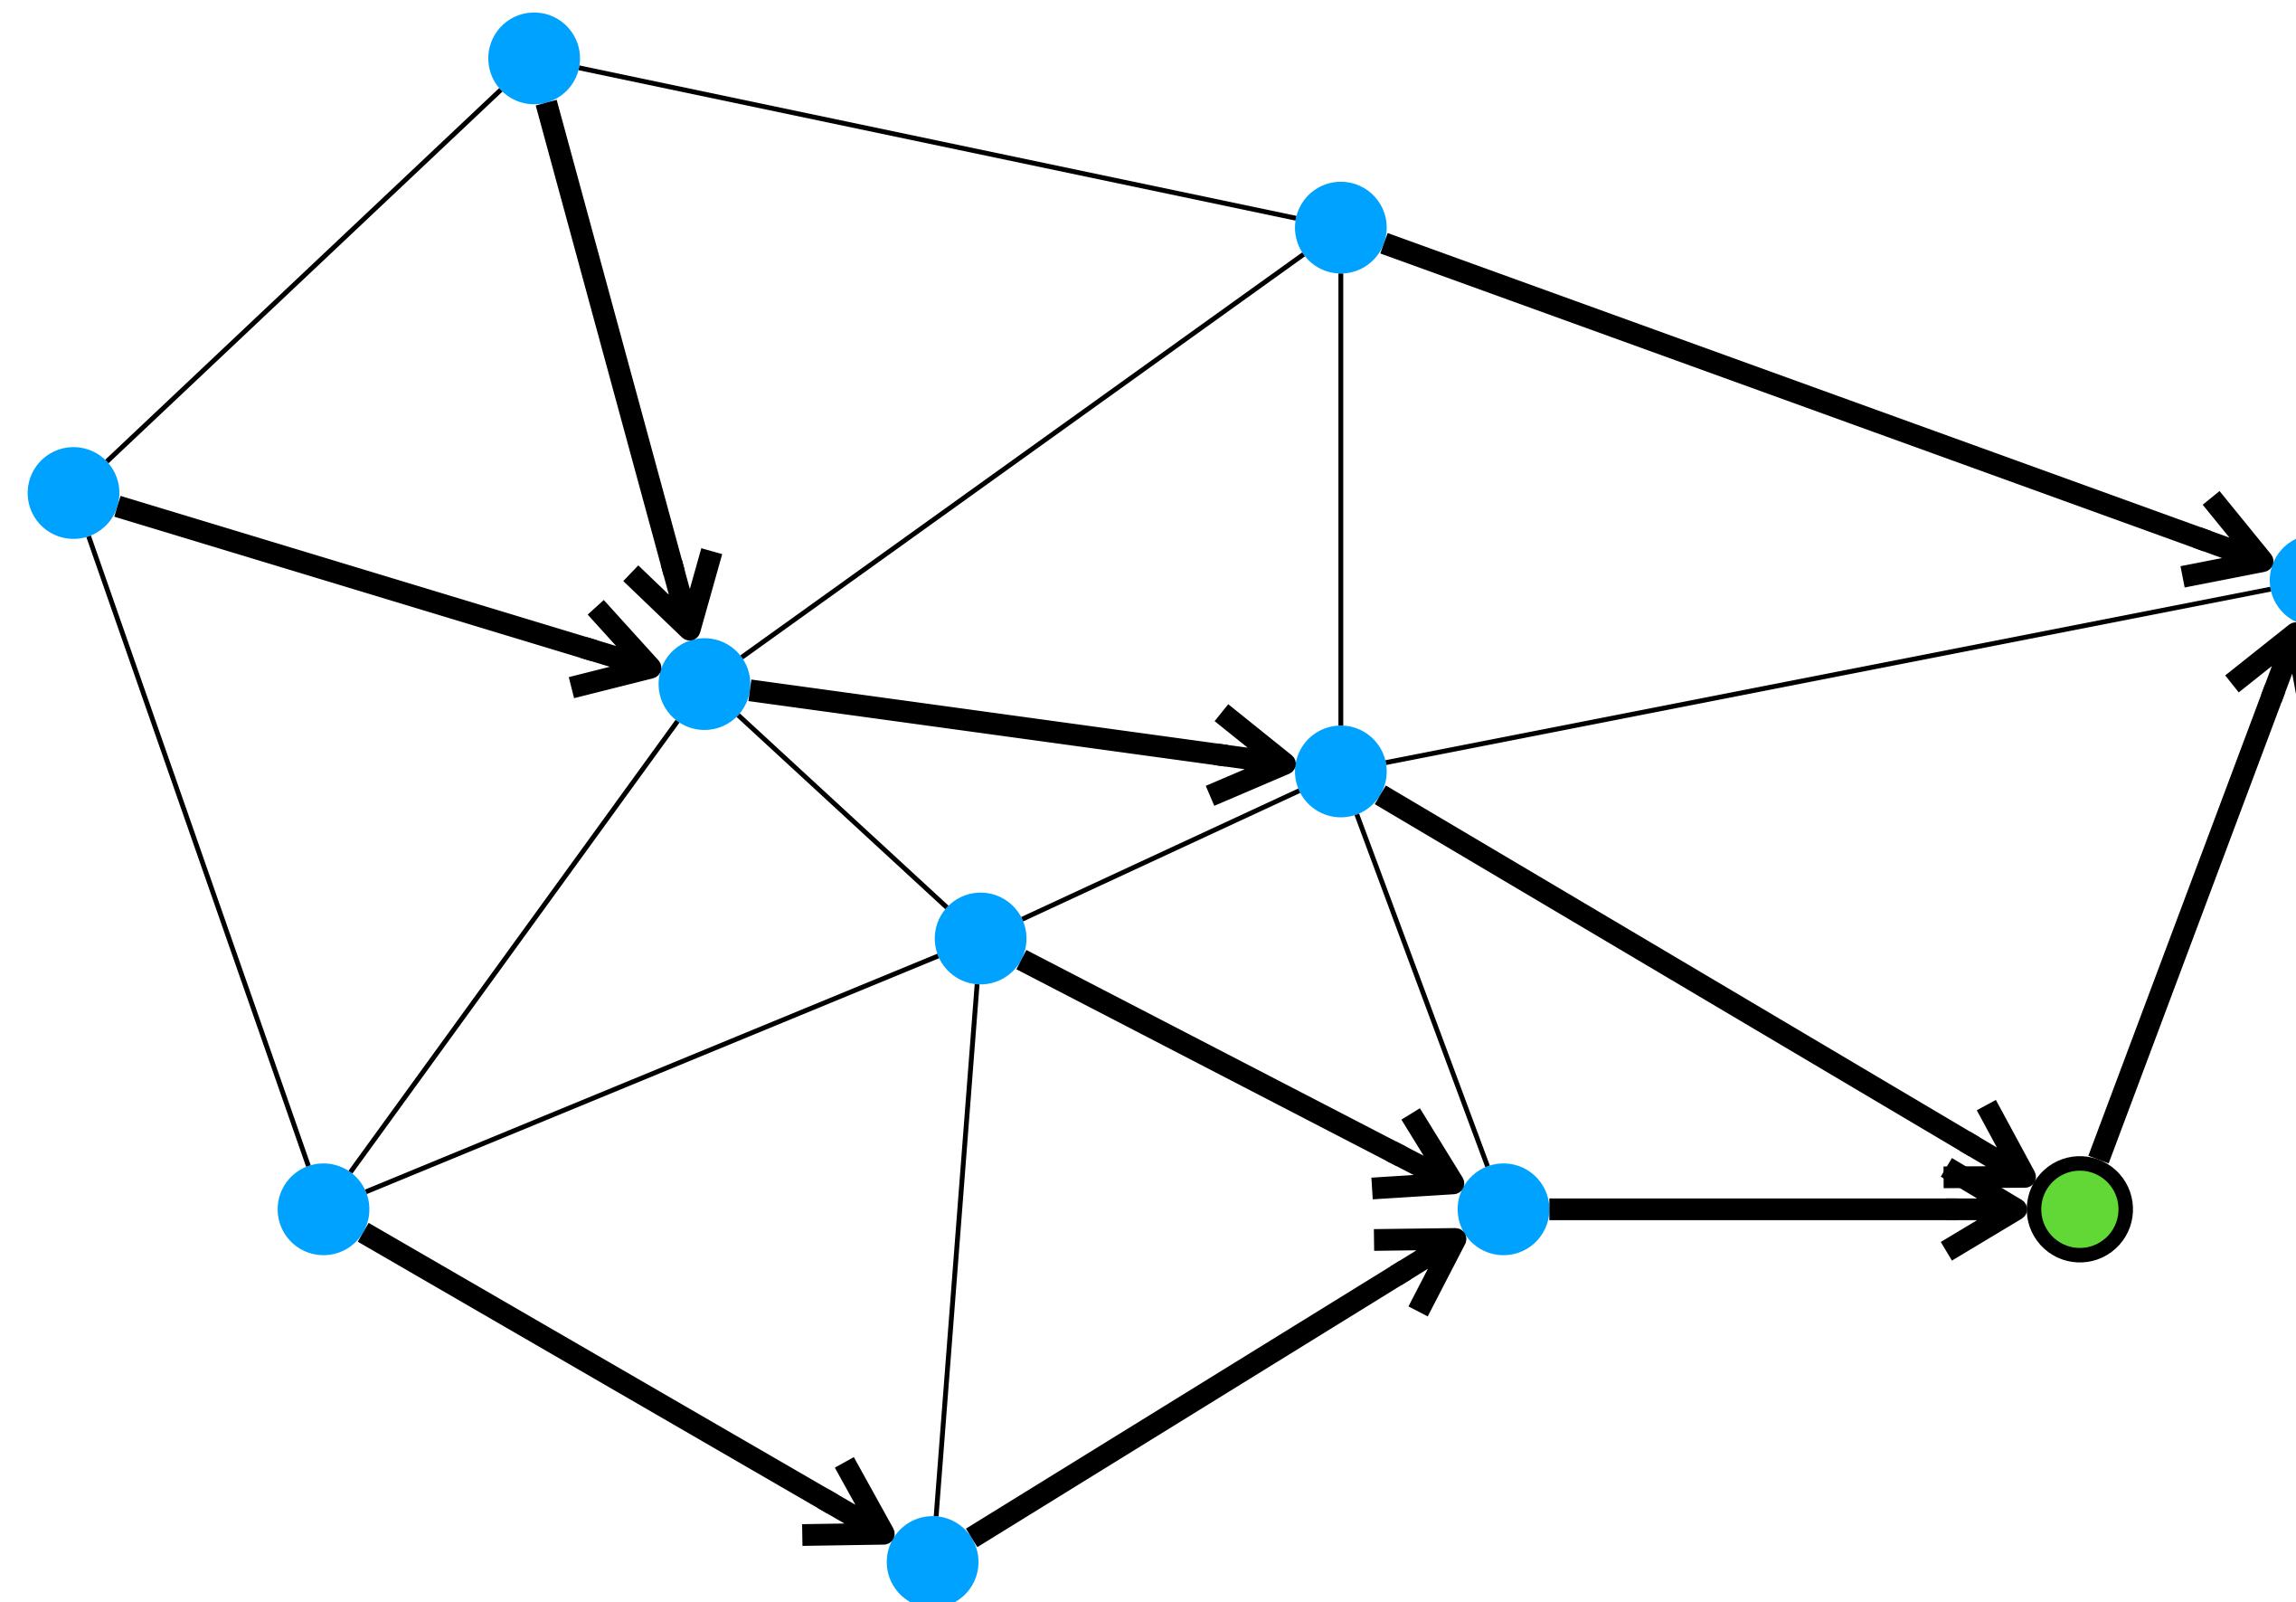
# Example I



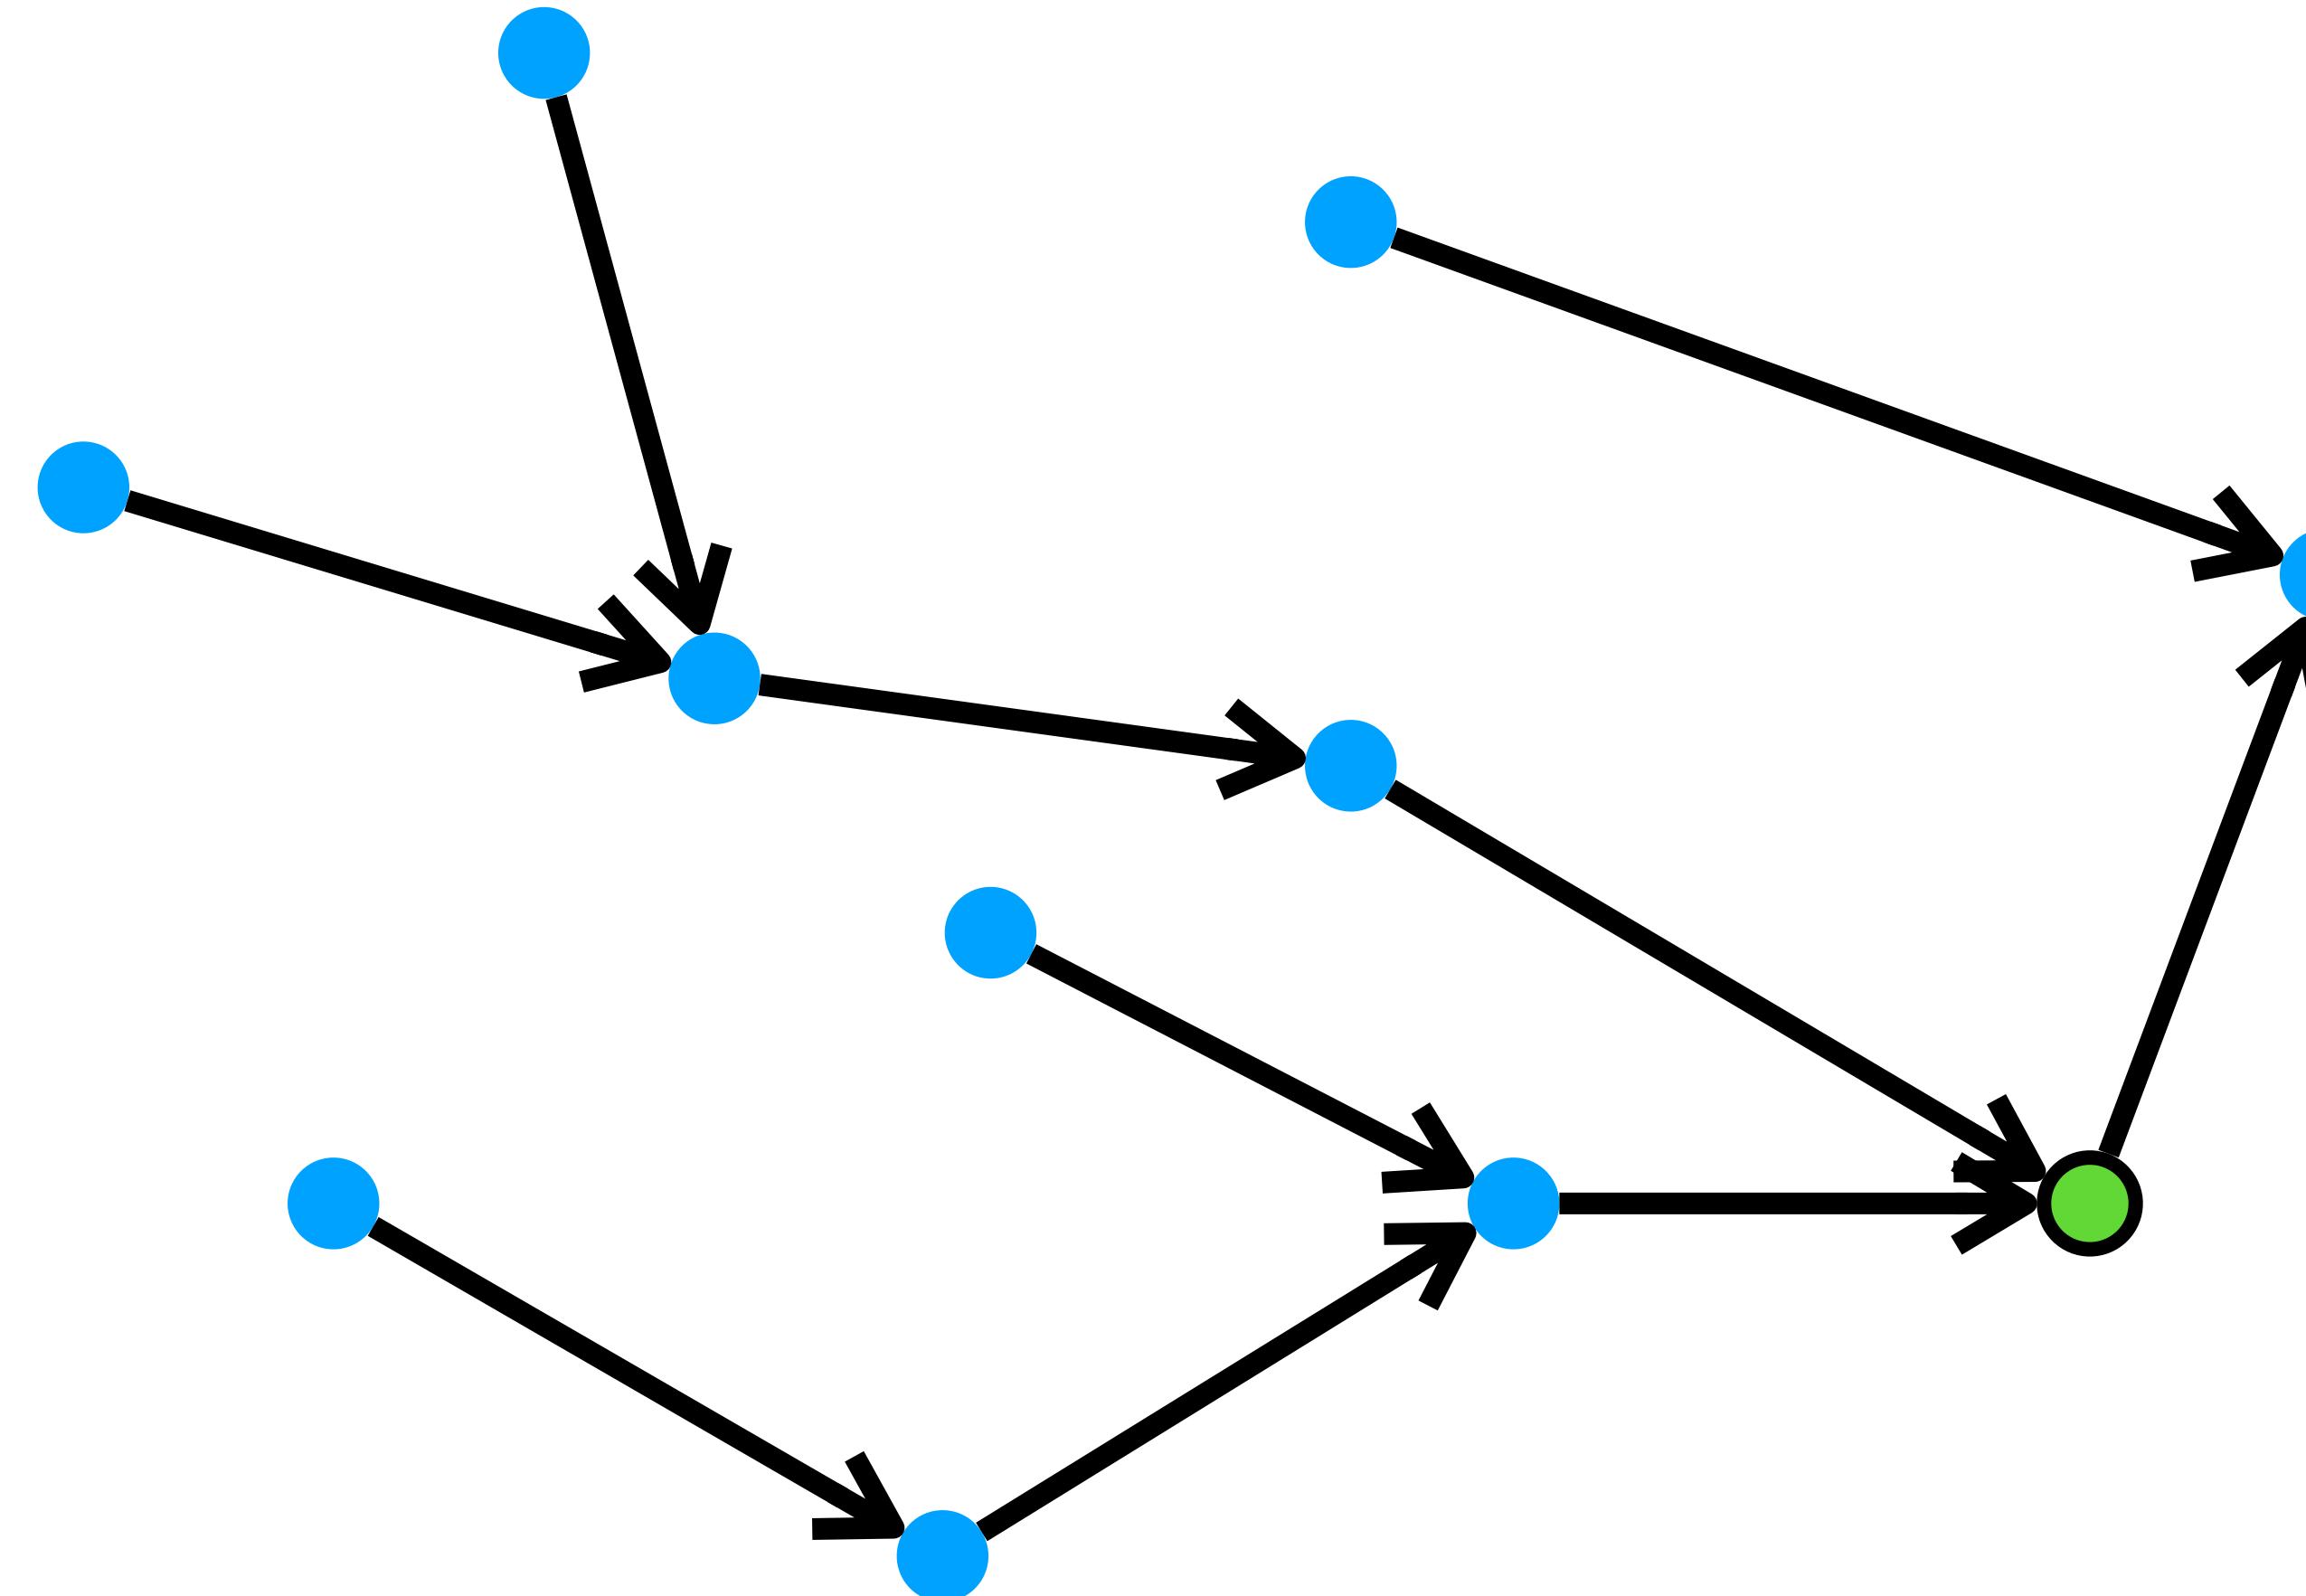
# Pick Destination



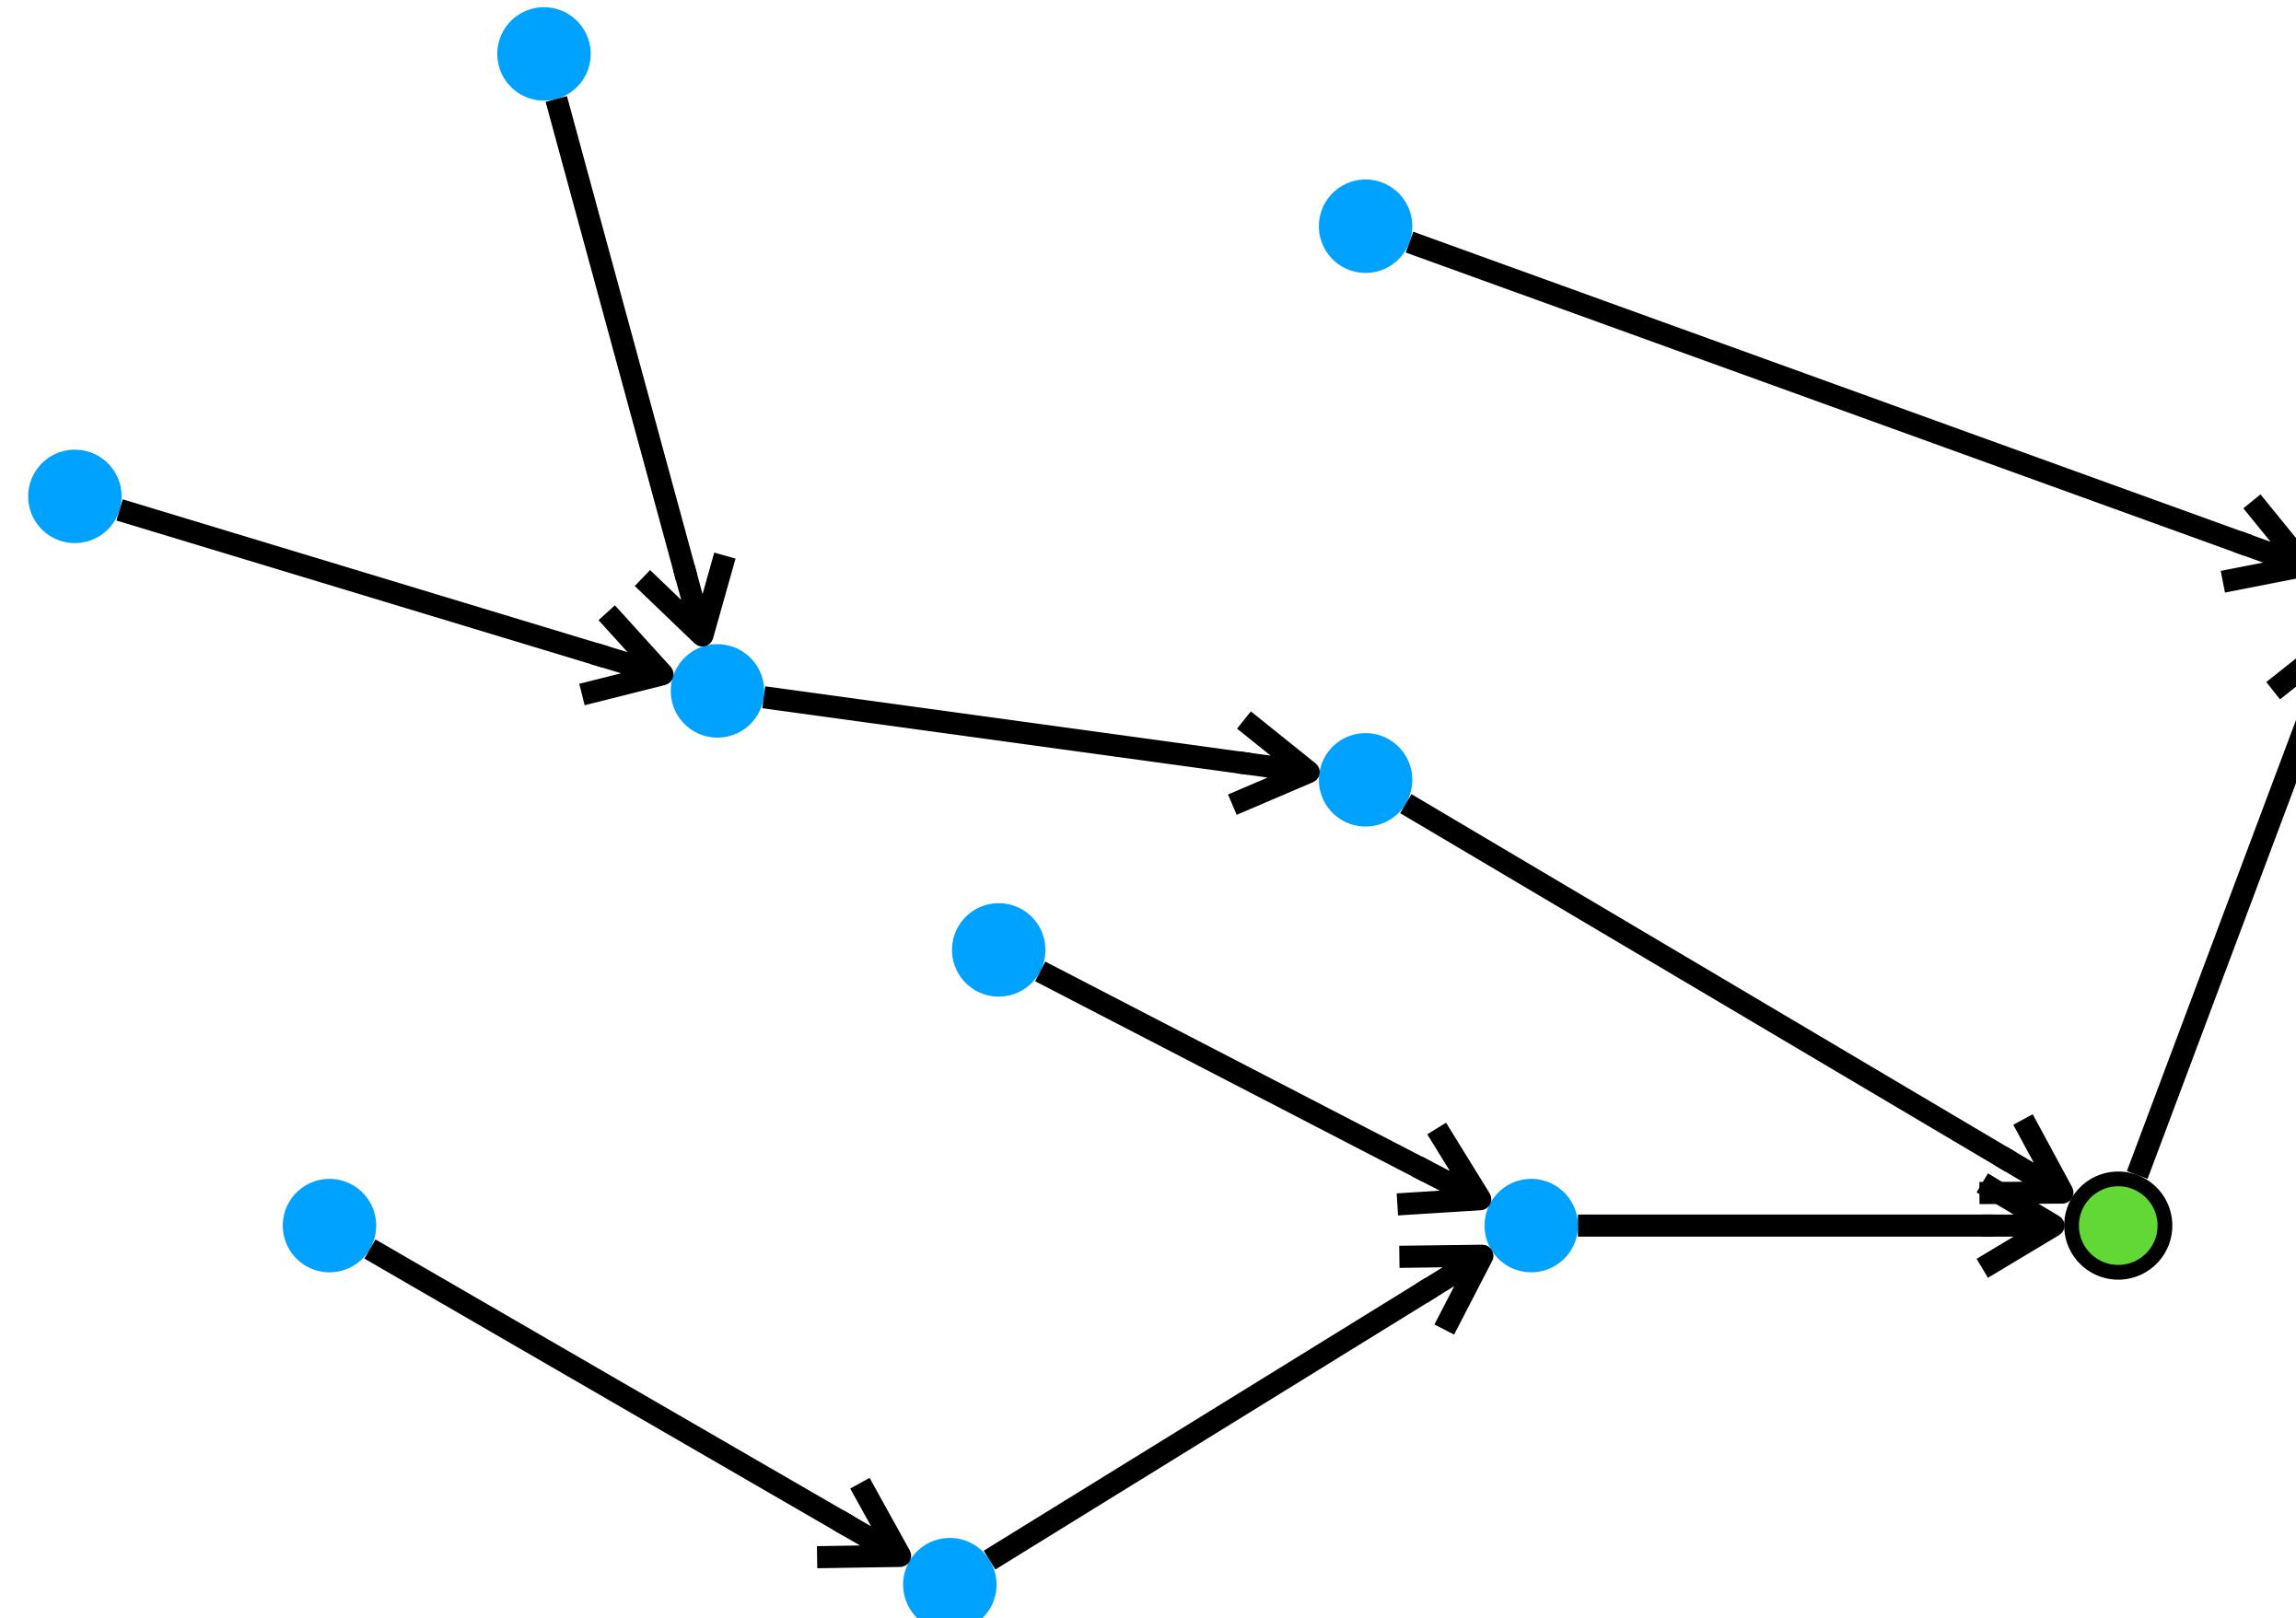
Put arrows on outgoing links (to green dot)



# Remove Unused Links

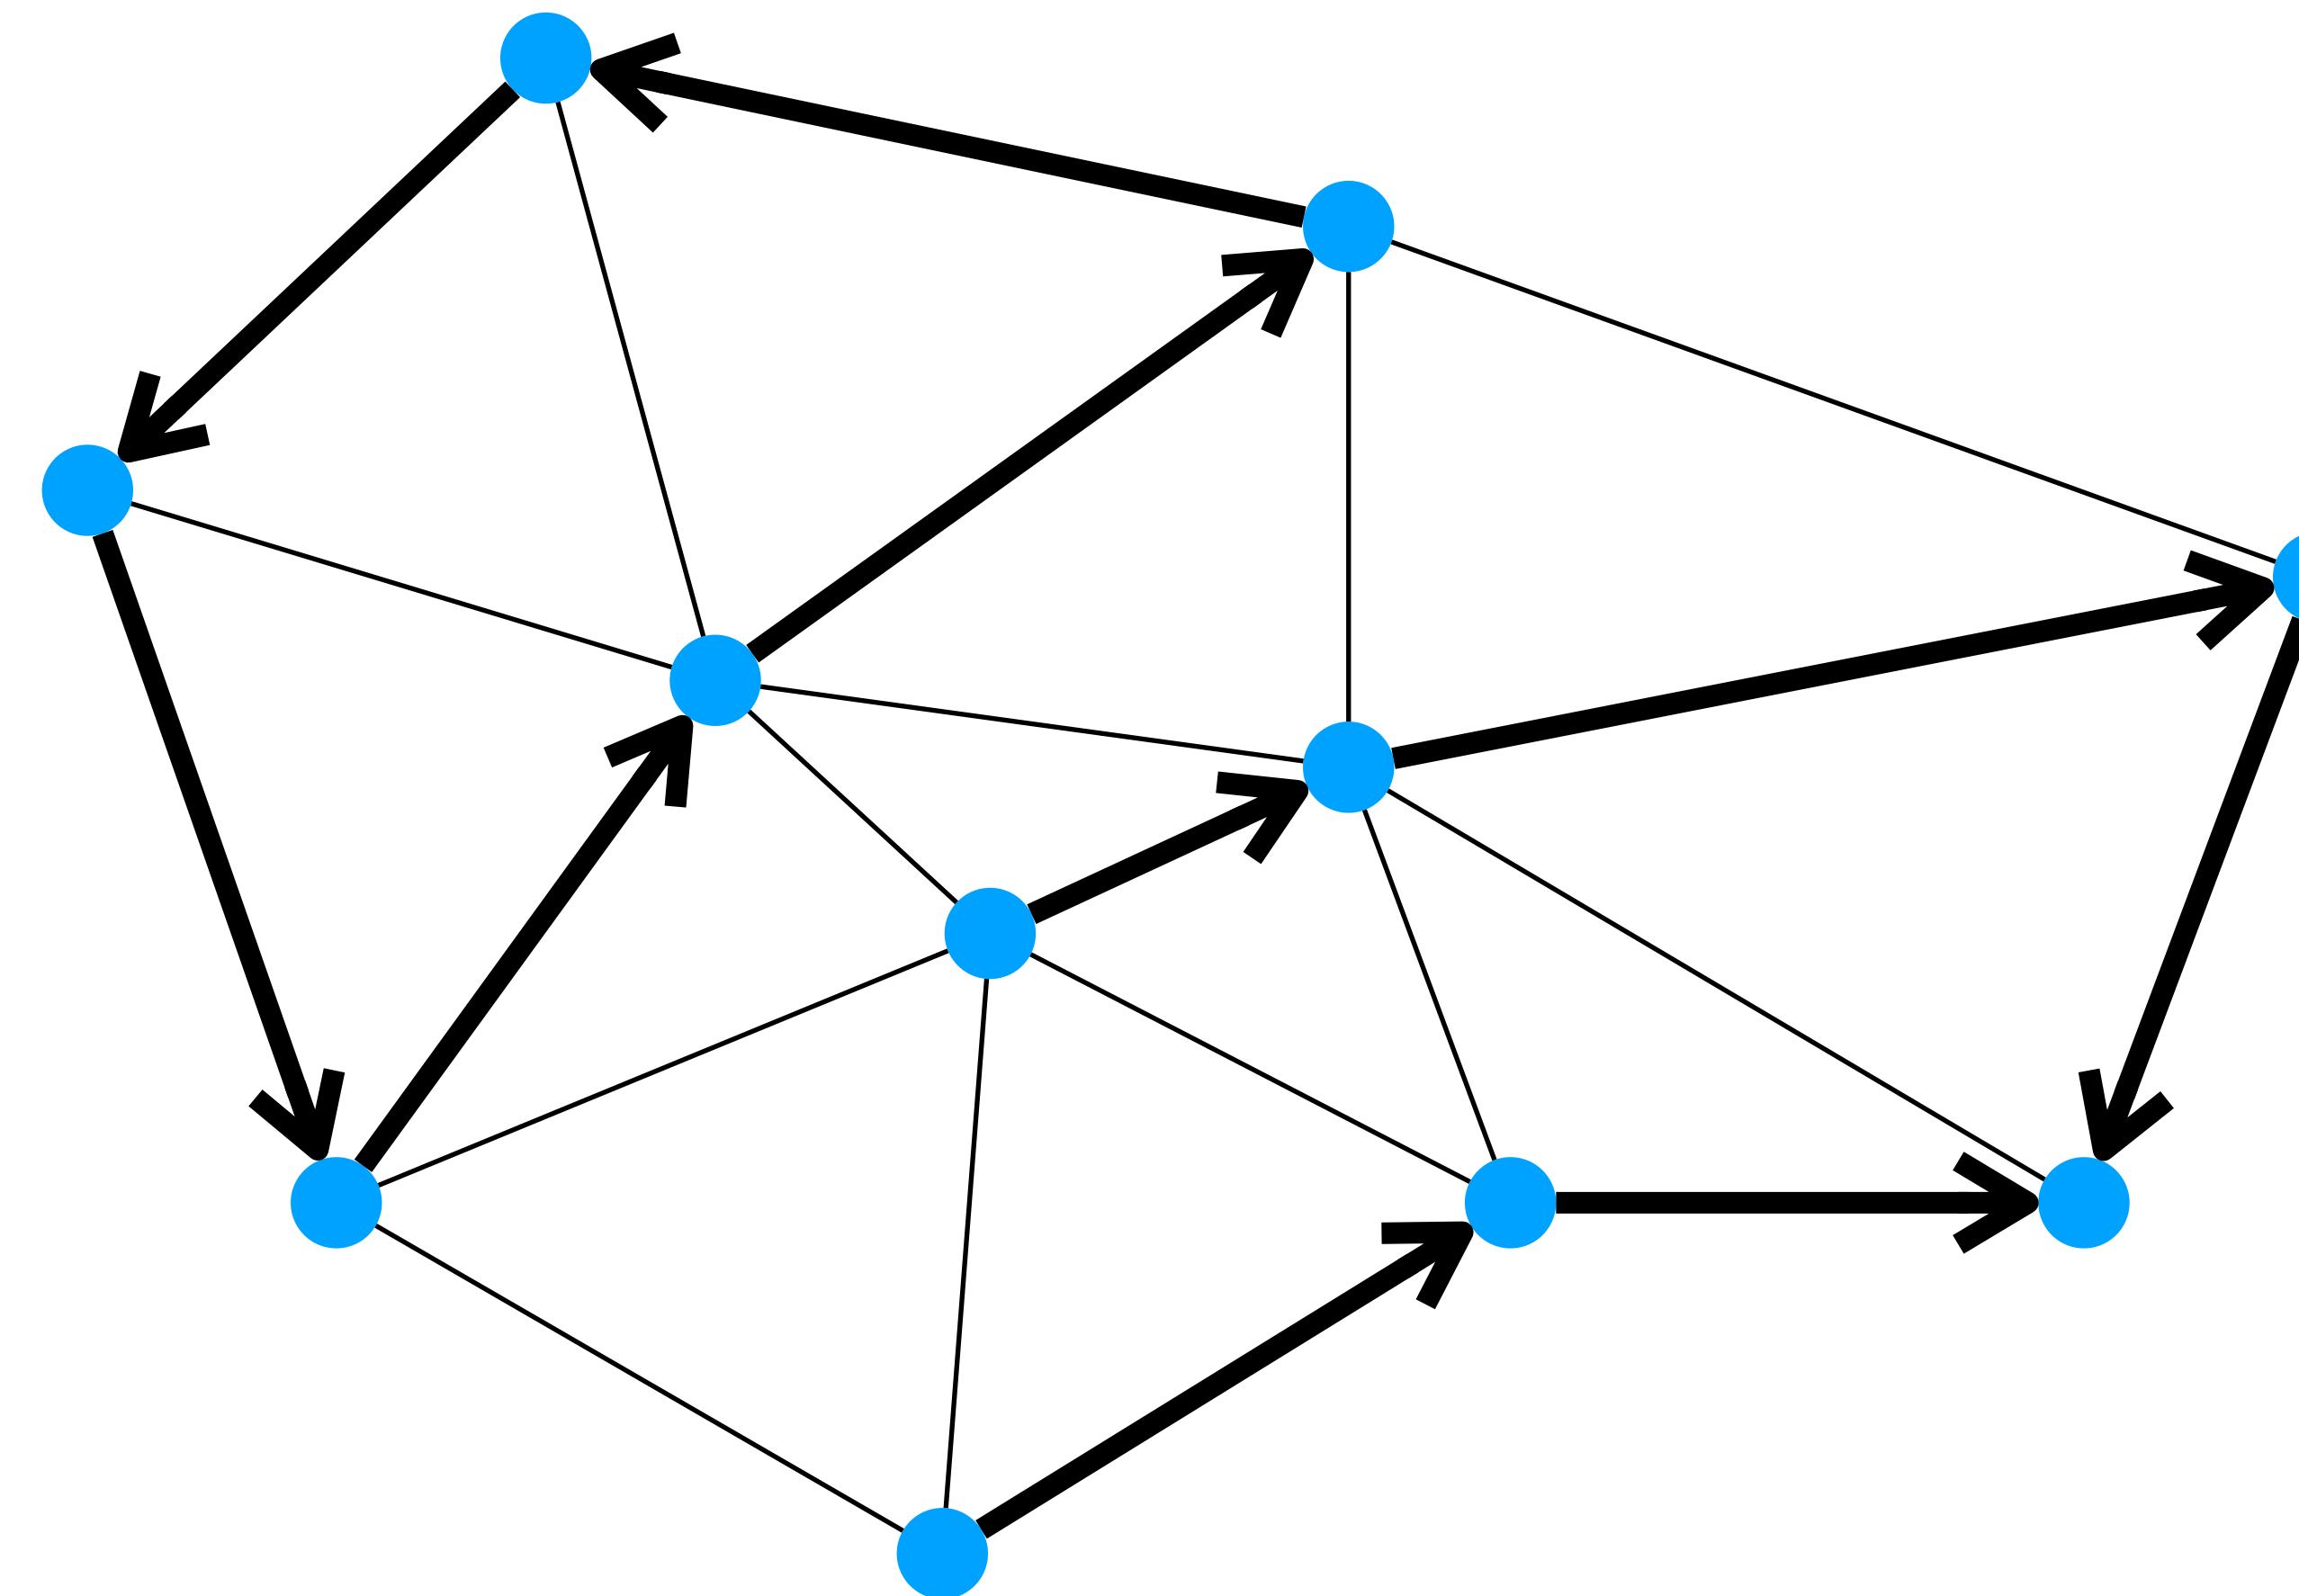


# Remove Unused Links

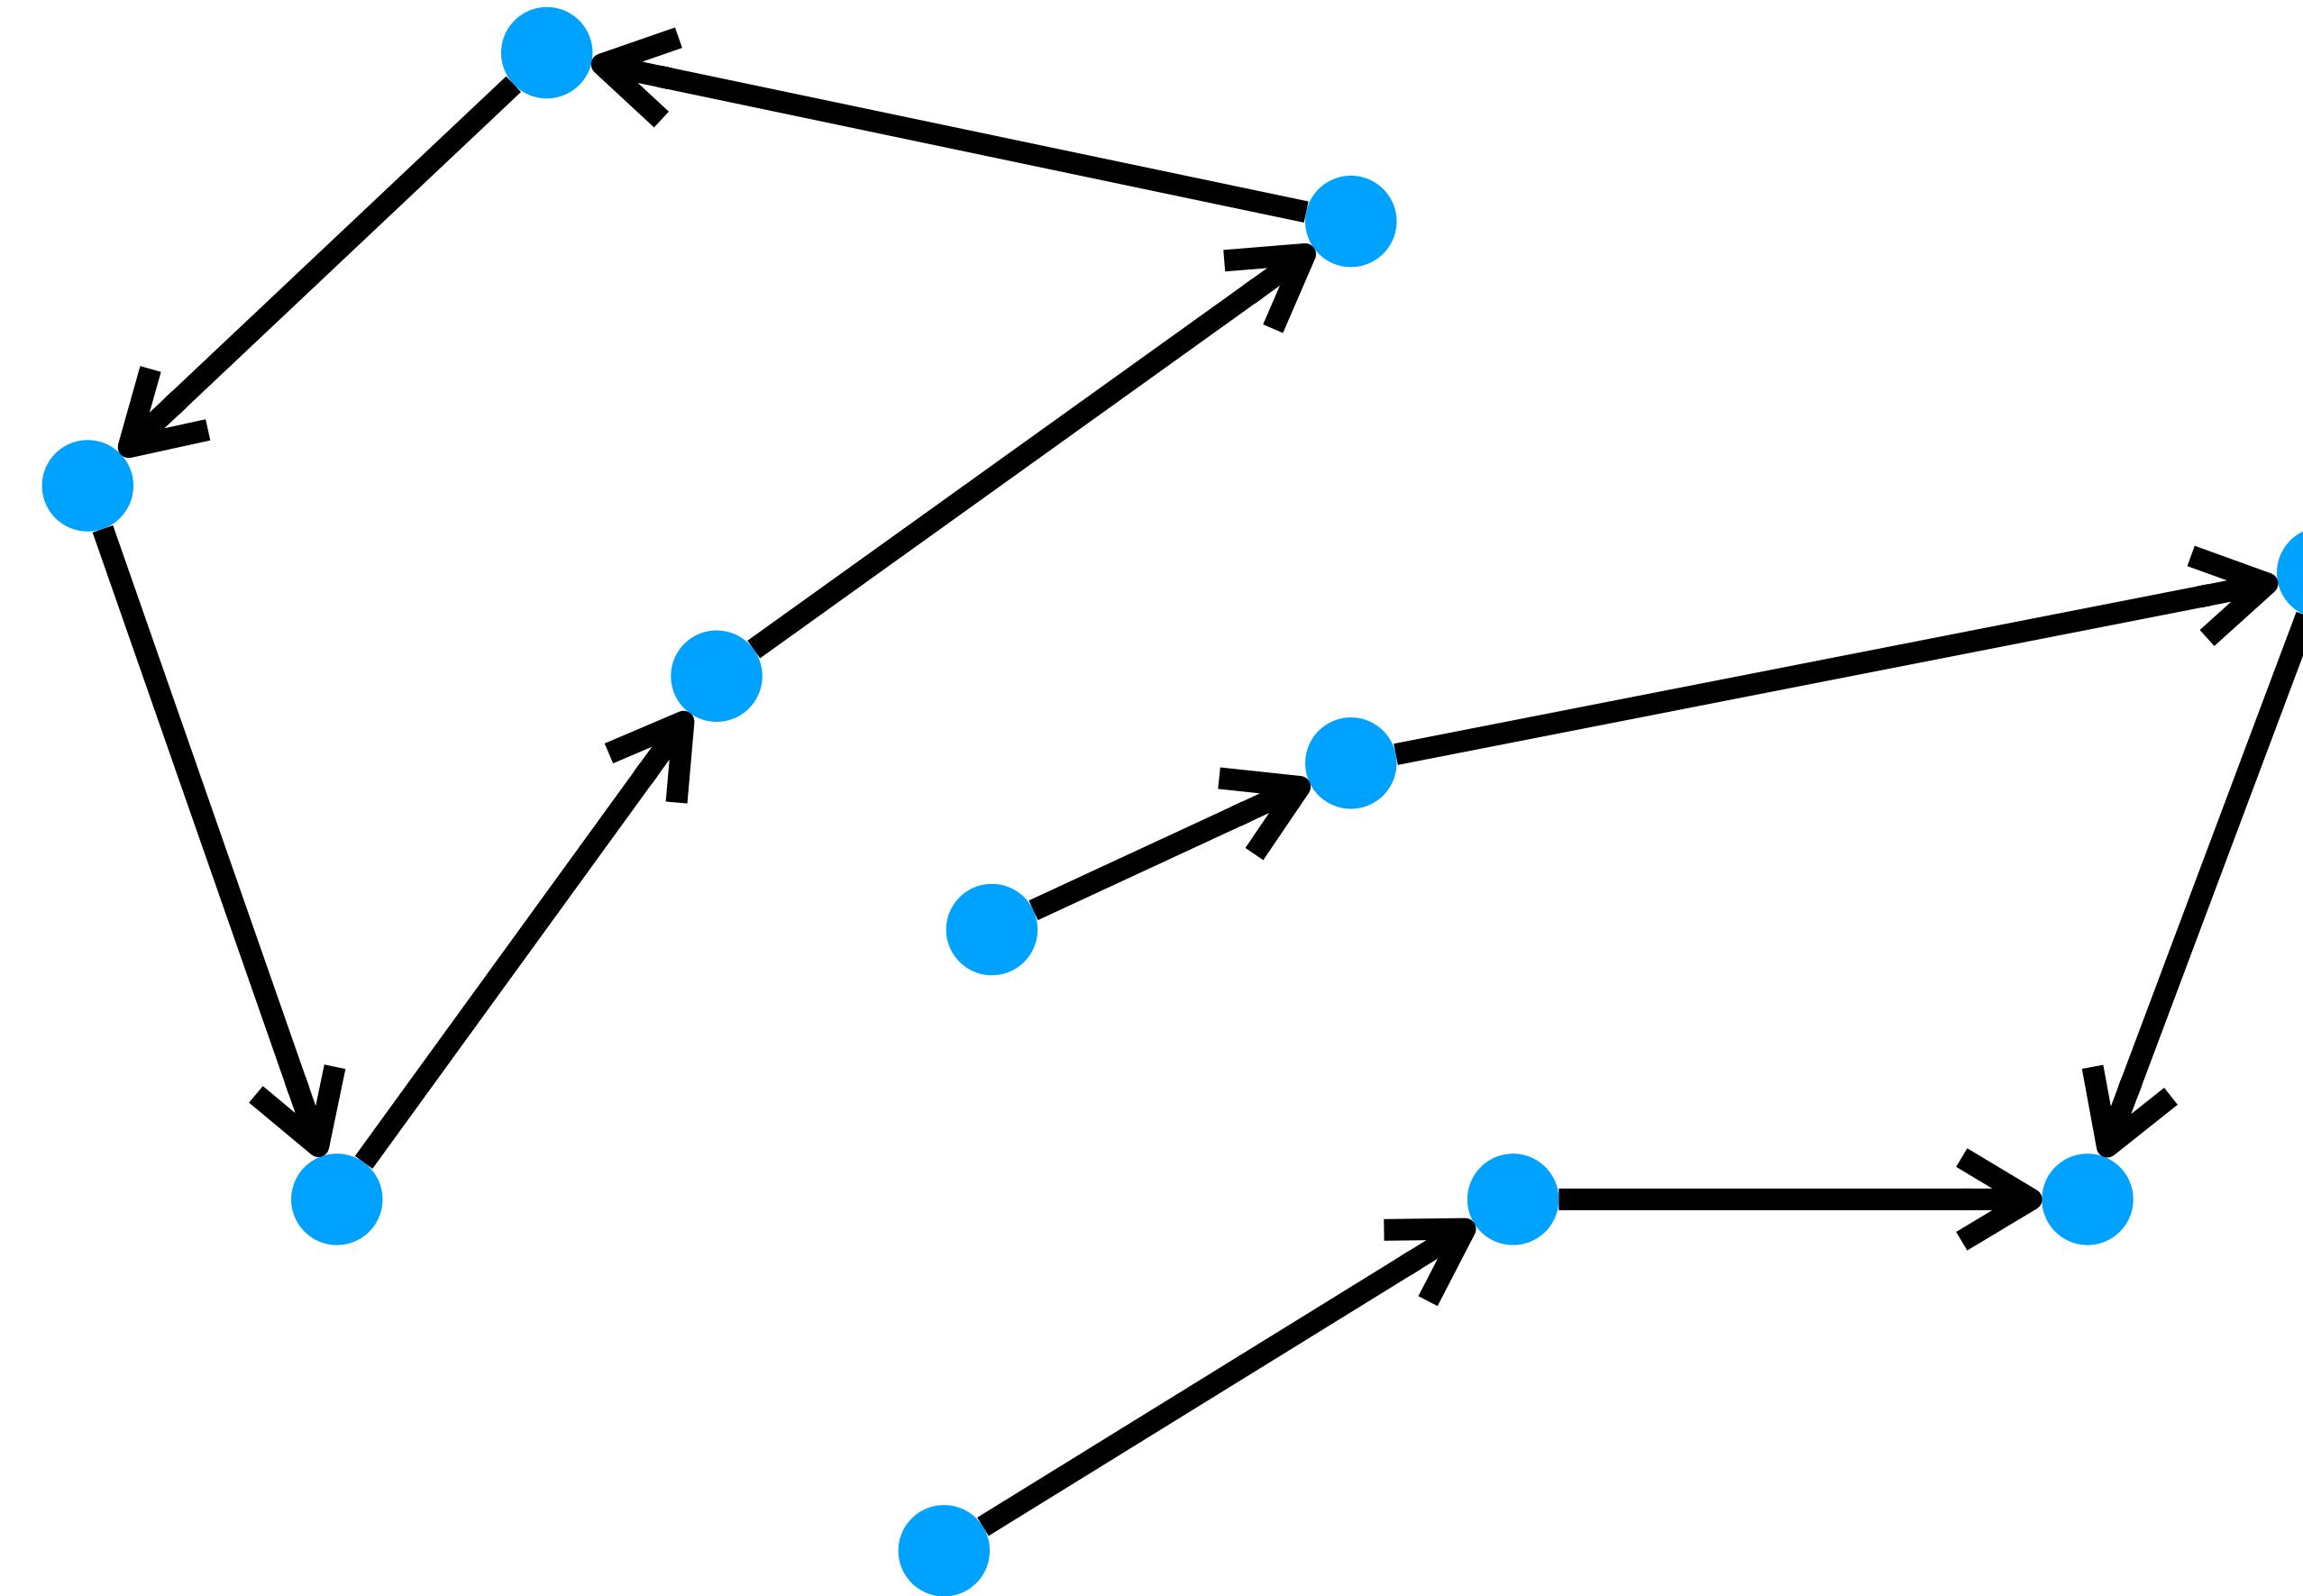


Leaves Spanning Tree: **Valid**

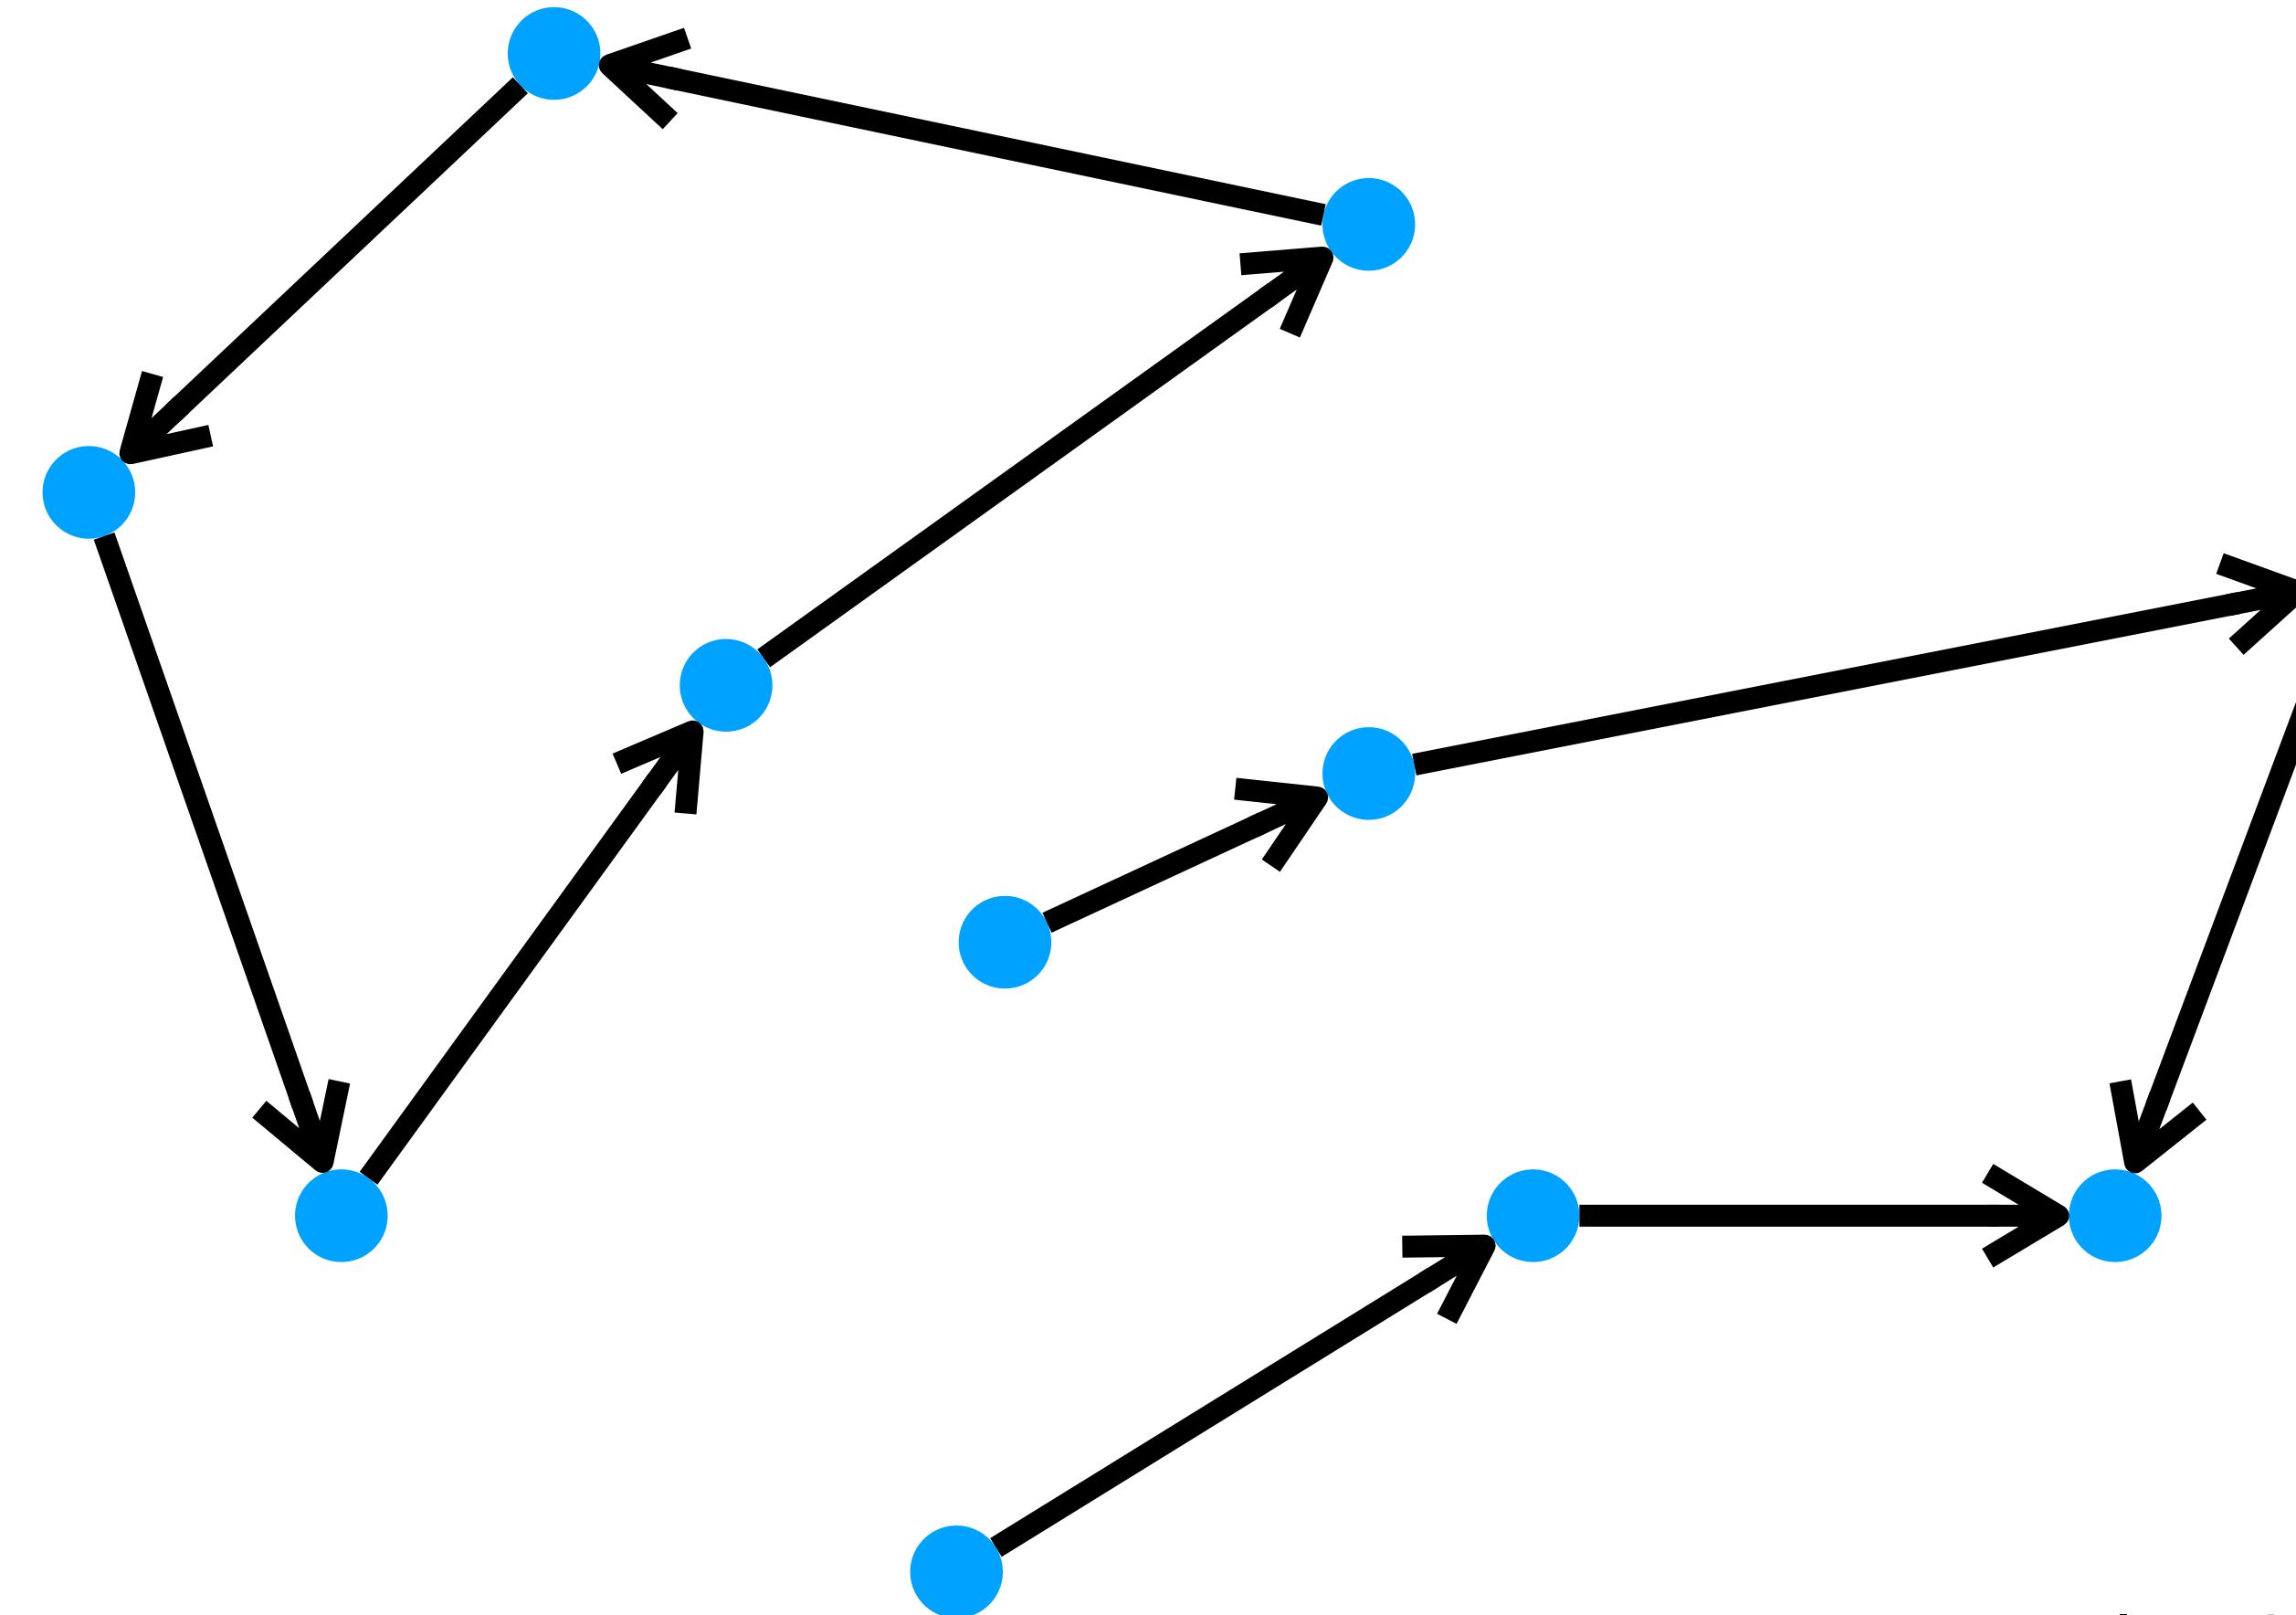
# Example 2



# Example 2



# Example 2



Is this **Valid**?

# Lesson?

- Very easy to check validity of routing state for a particular destination
- Dead-ends are obvious
  - Node without outgoing arrow
- Loops are obvious
  - Disconnected from the rest of the graph

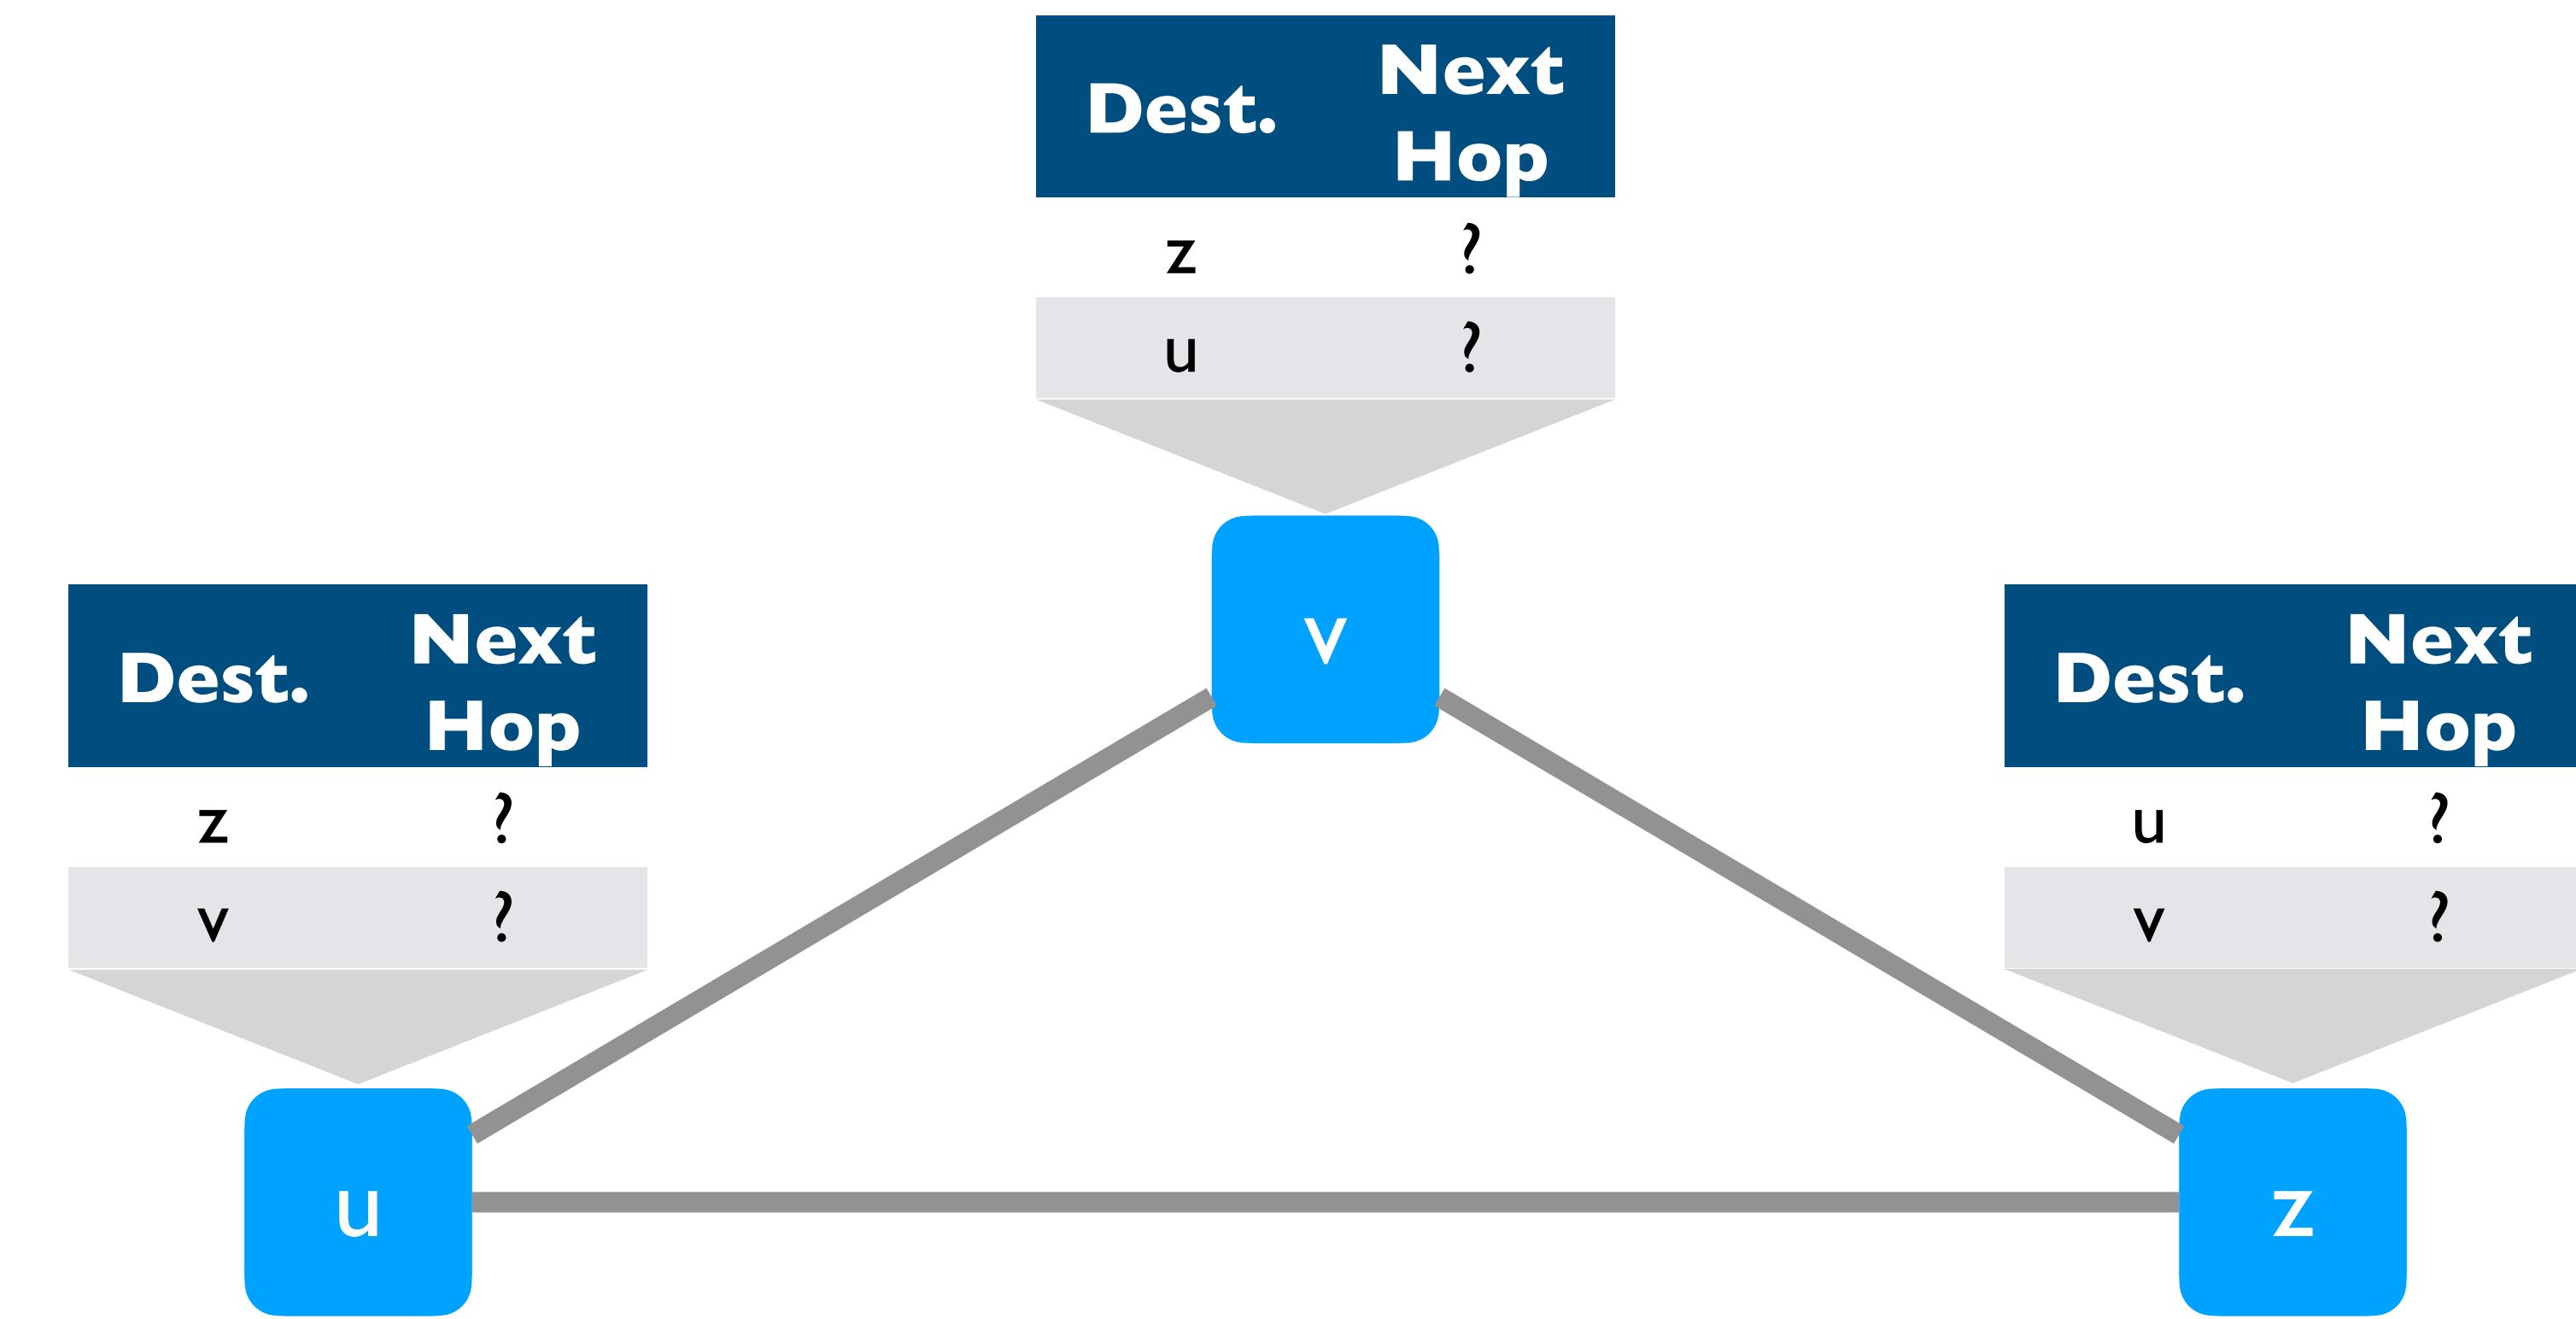
# Goal (for routing)

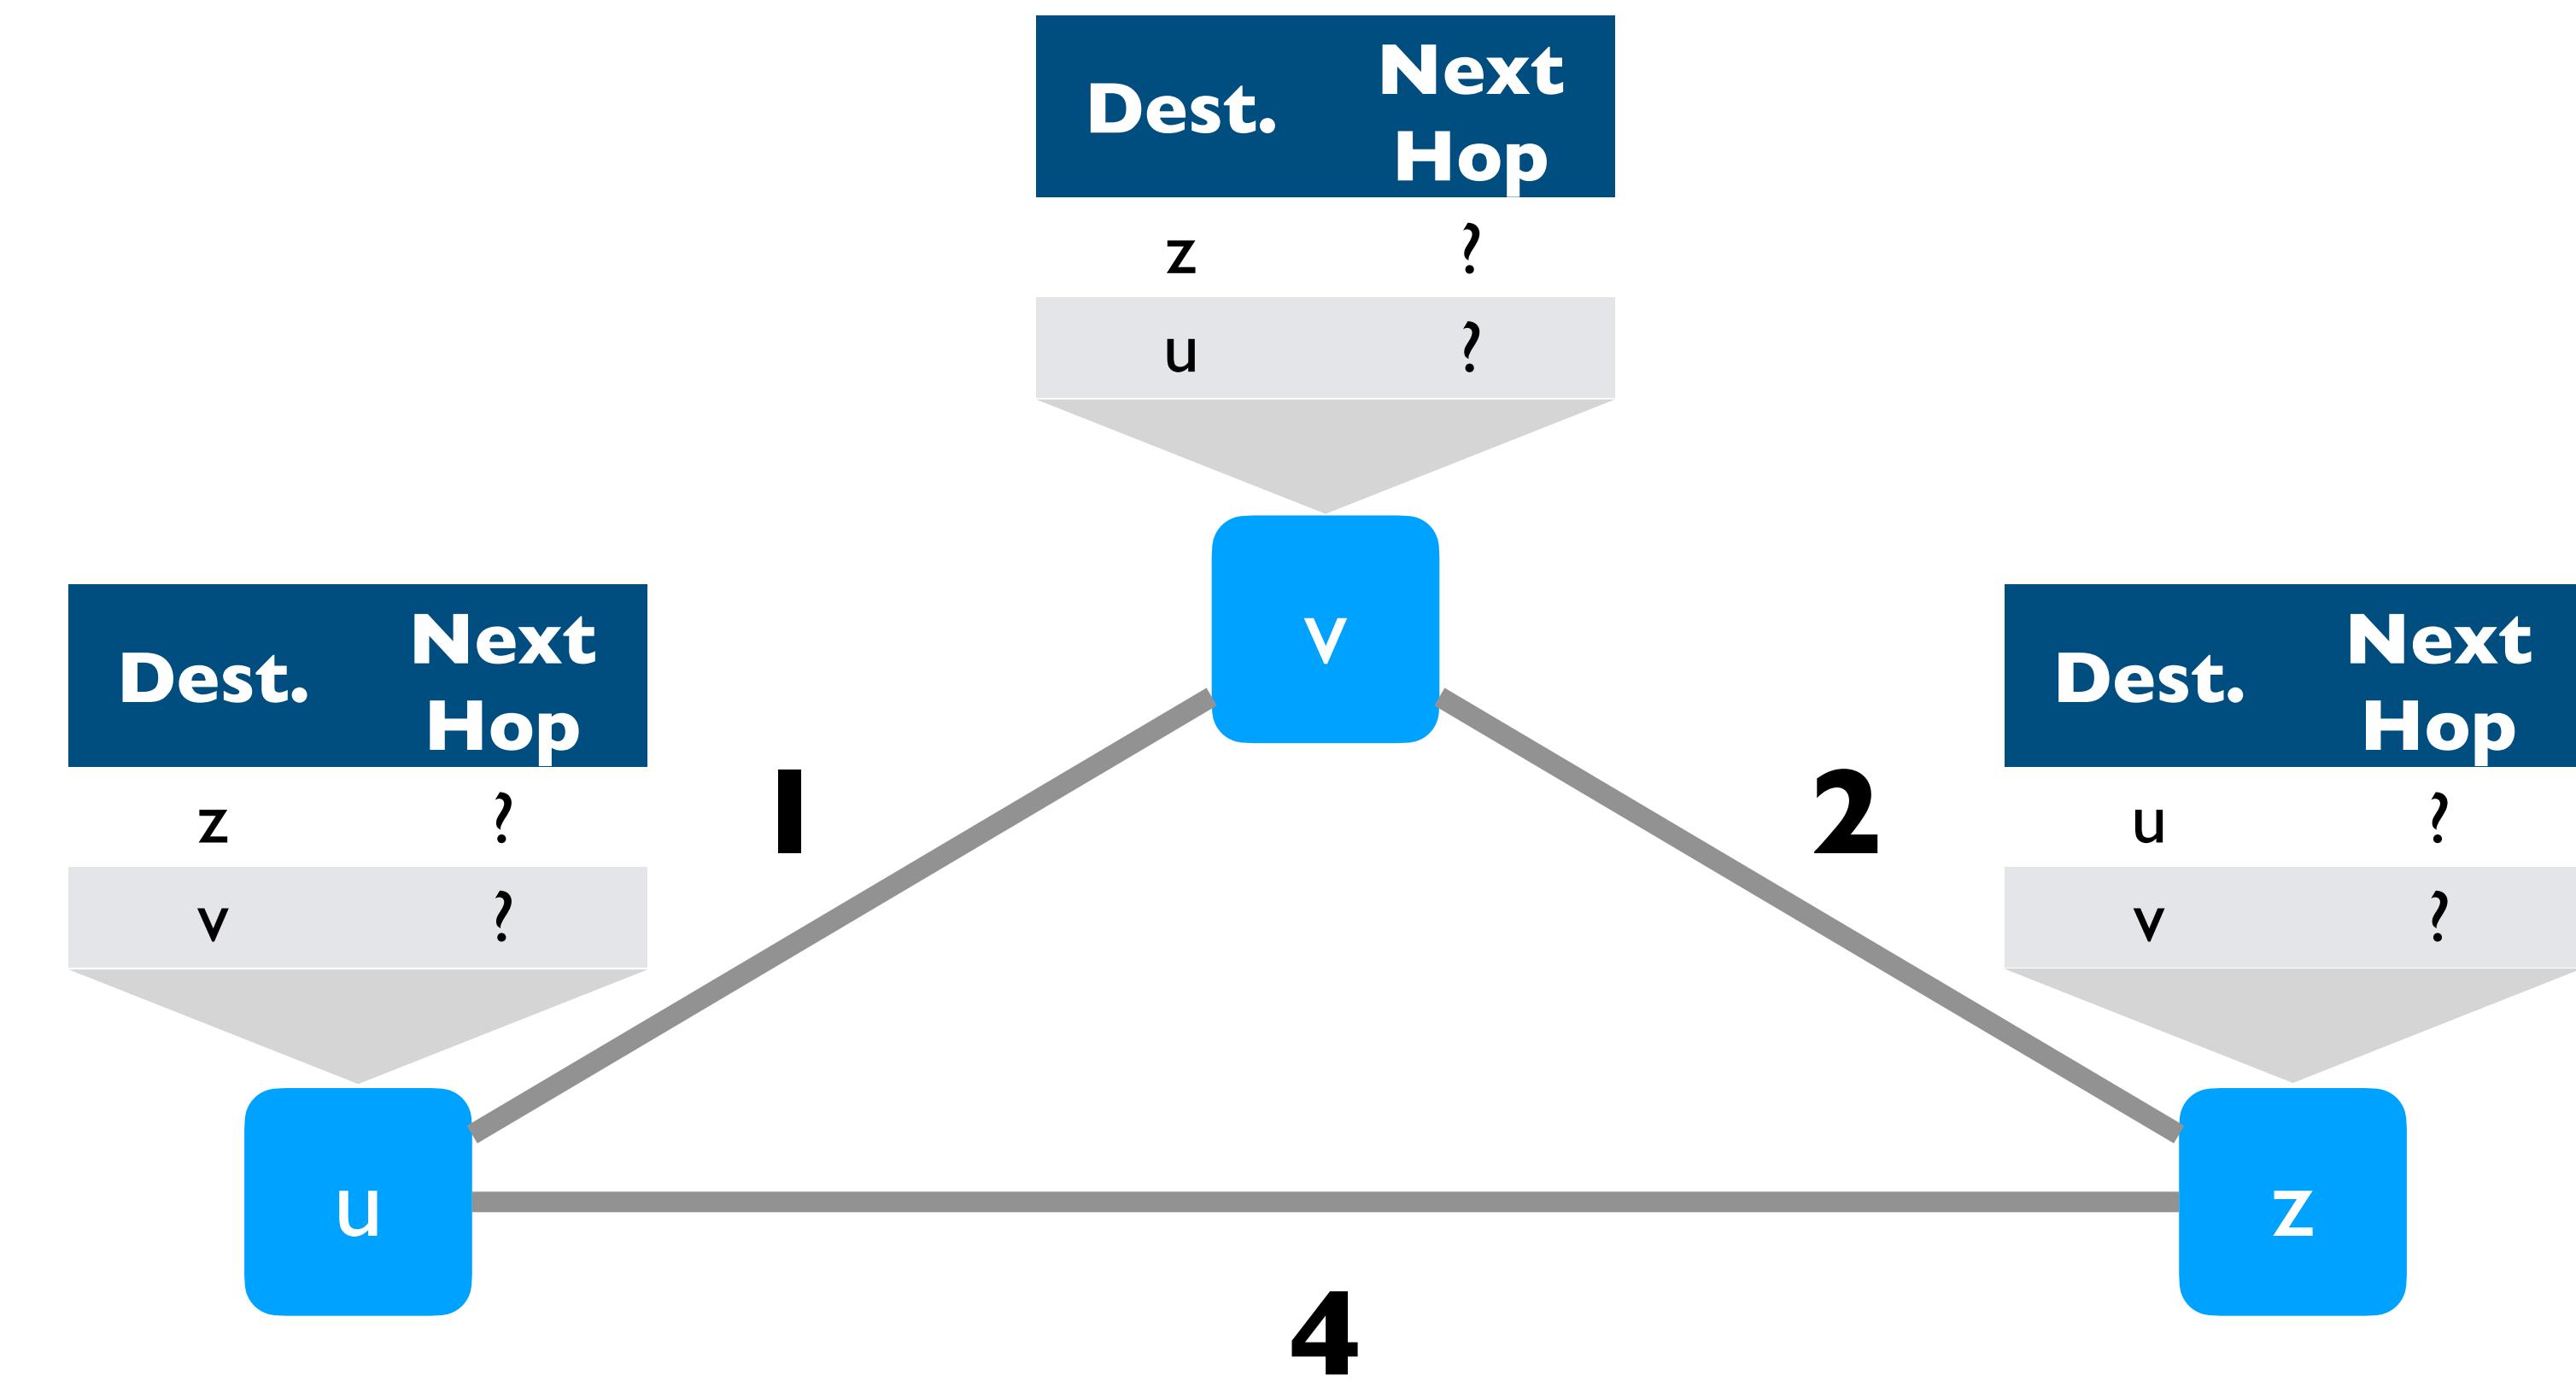
# Goal (for routing)

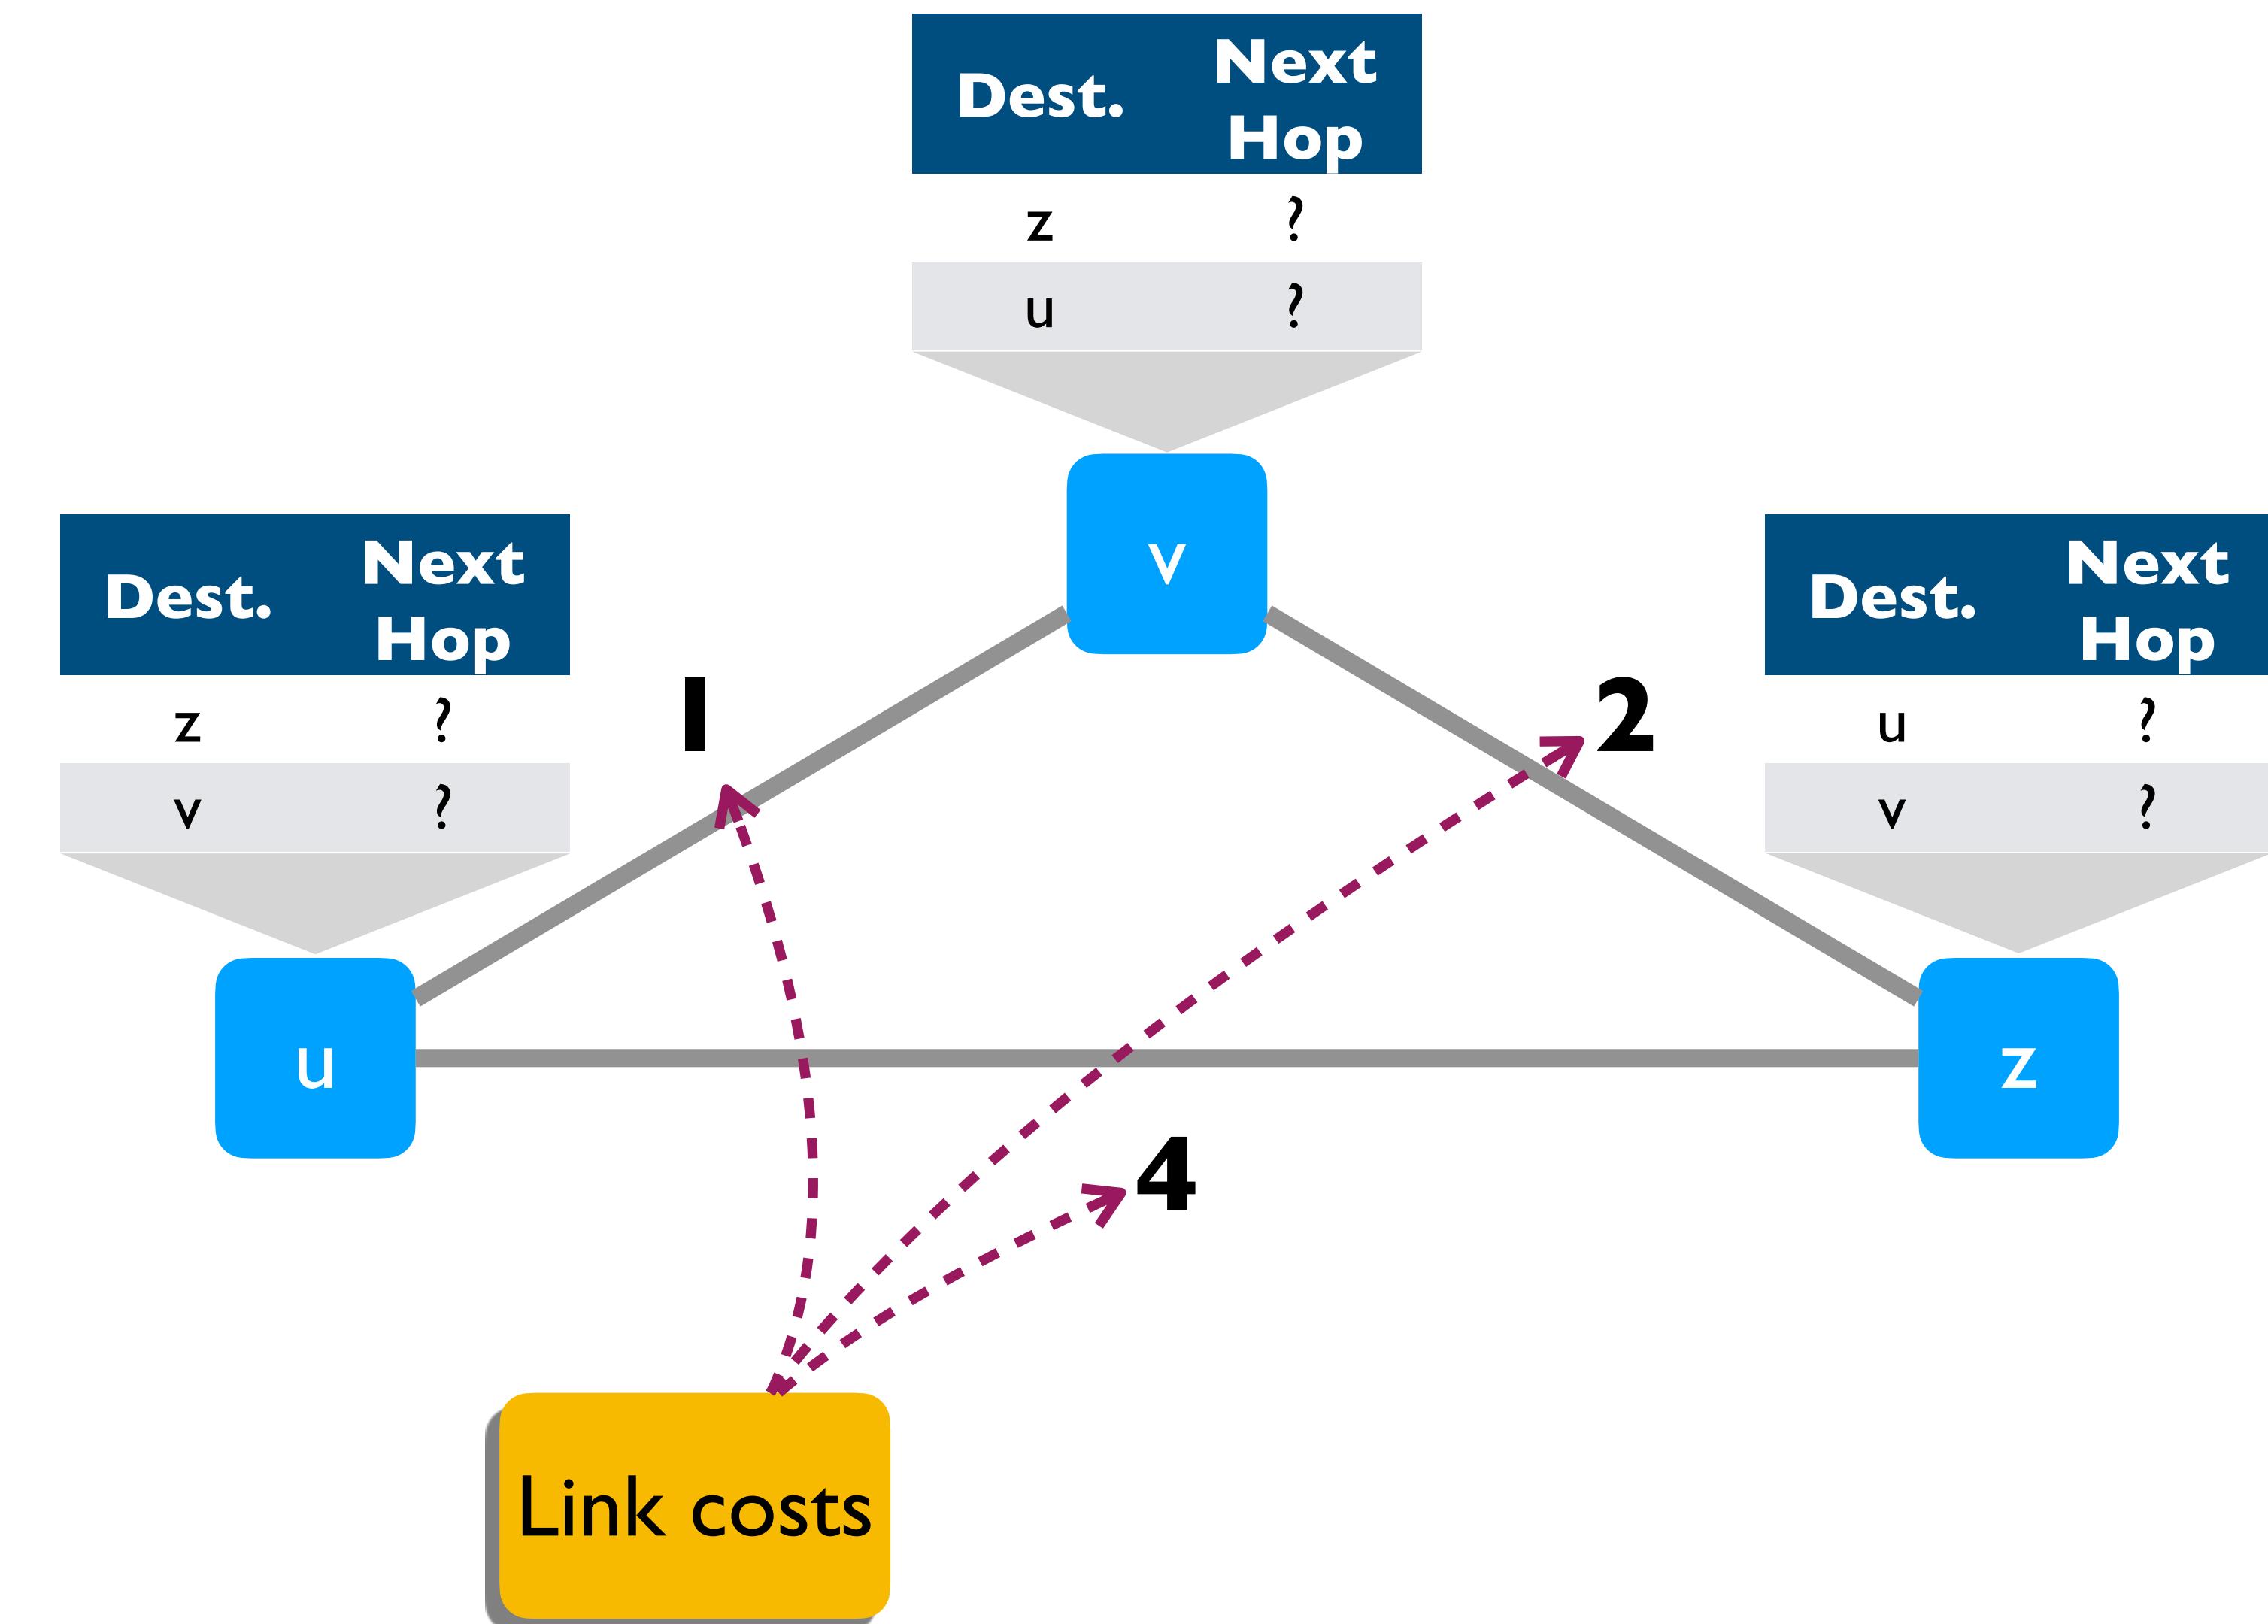
- **VI: Find a path to a given destination**

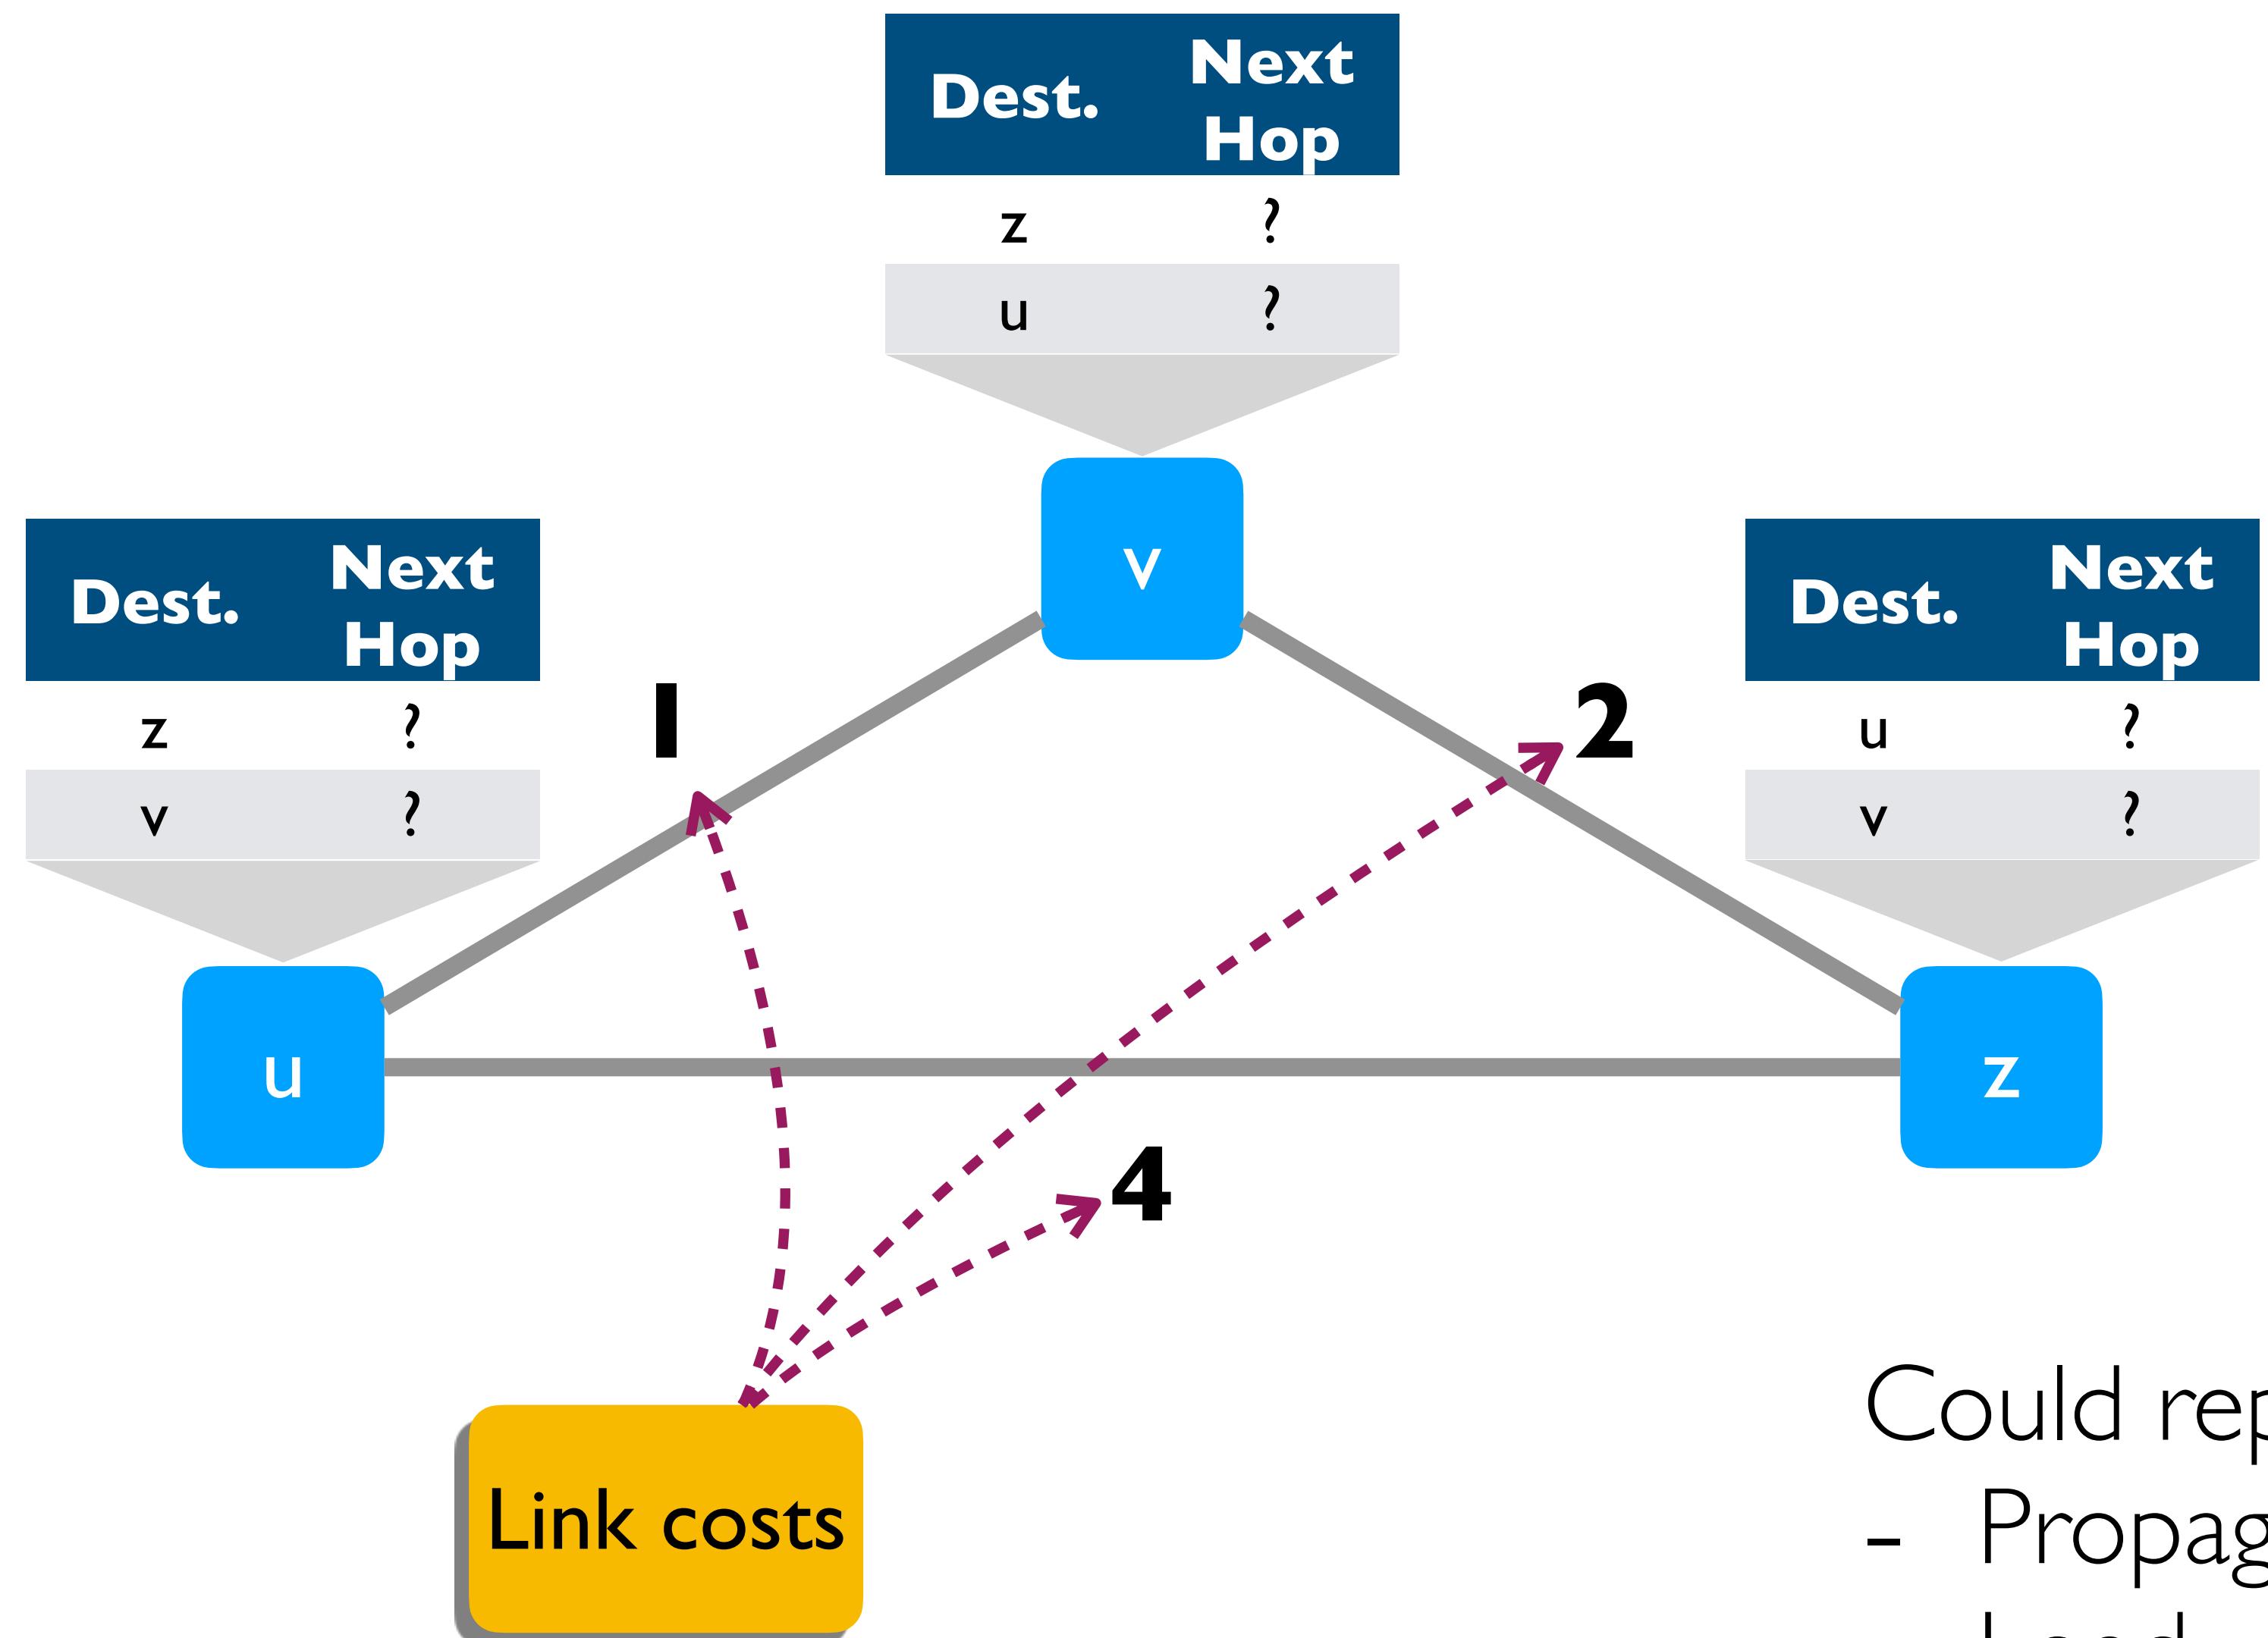
# Goal (for routing)

- **V1: Find a path to a given destination**
- **V2: Find a *least cost path* to a given destination**



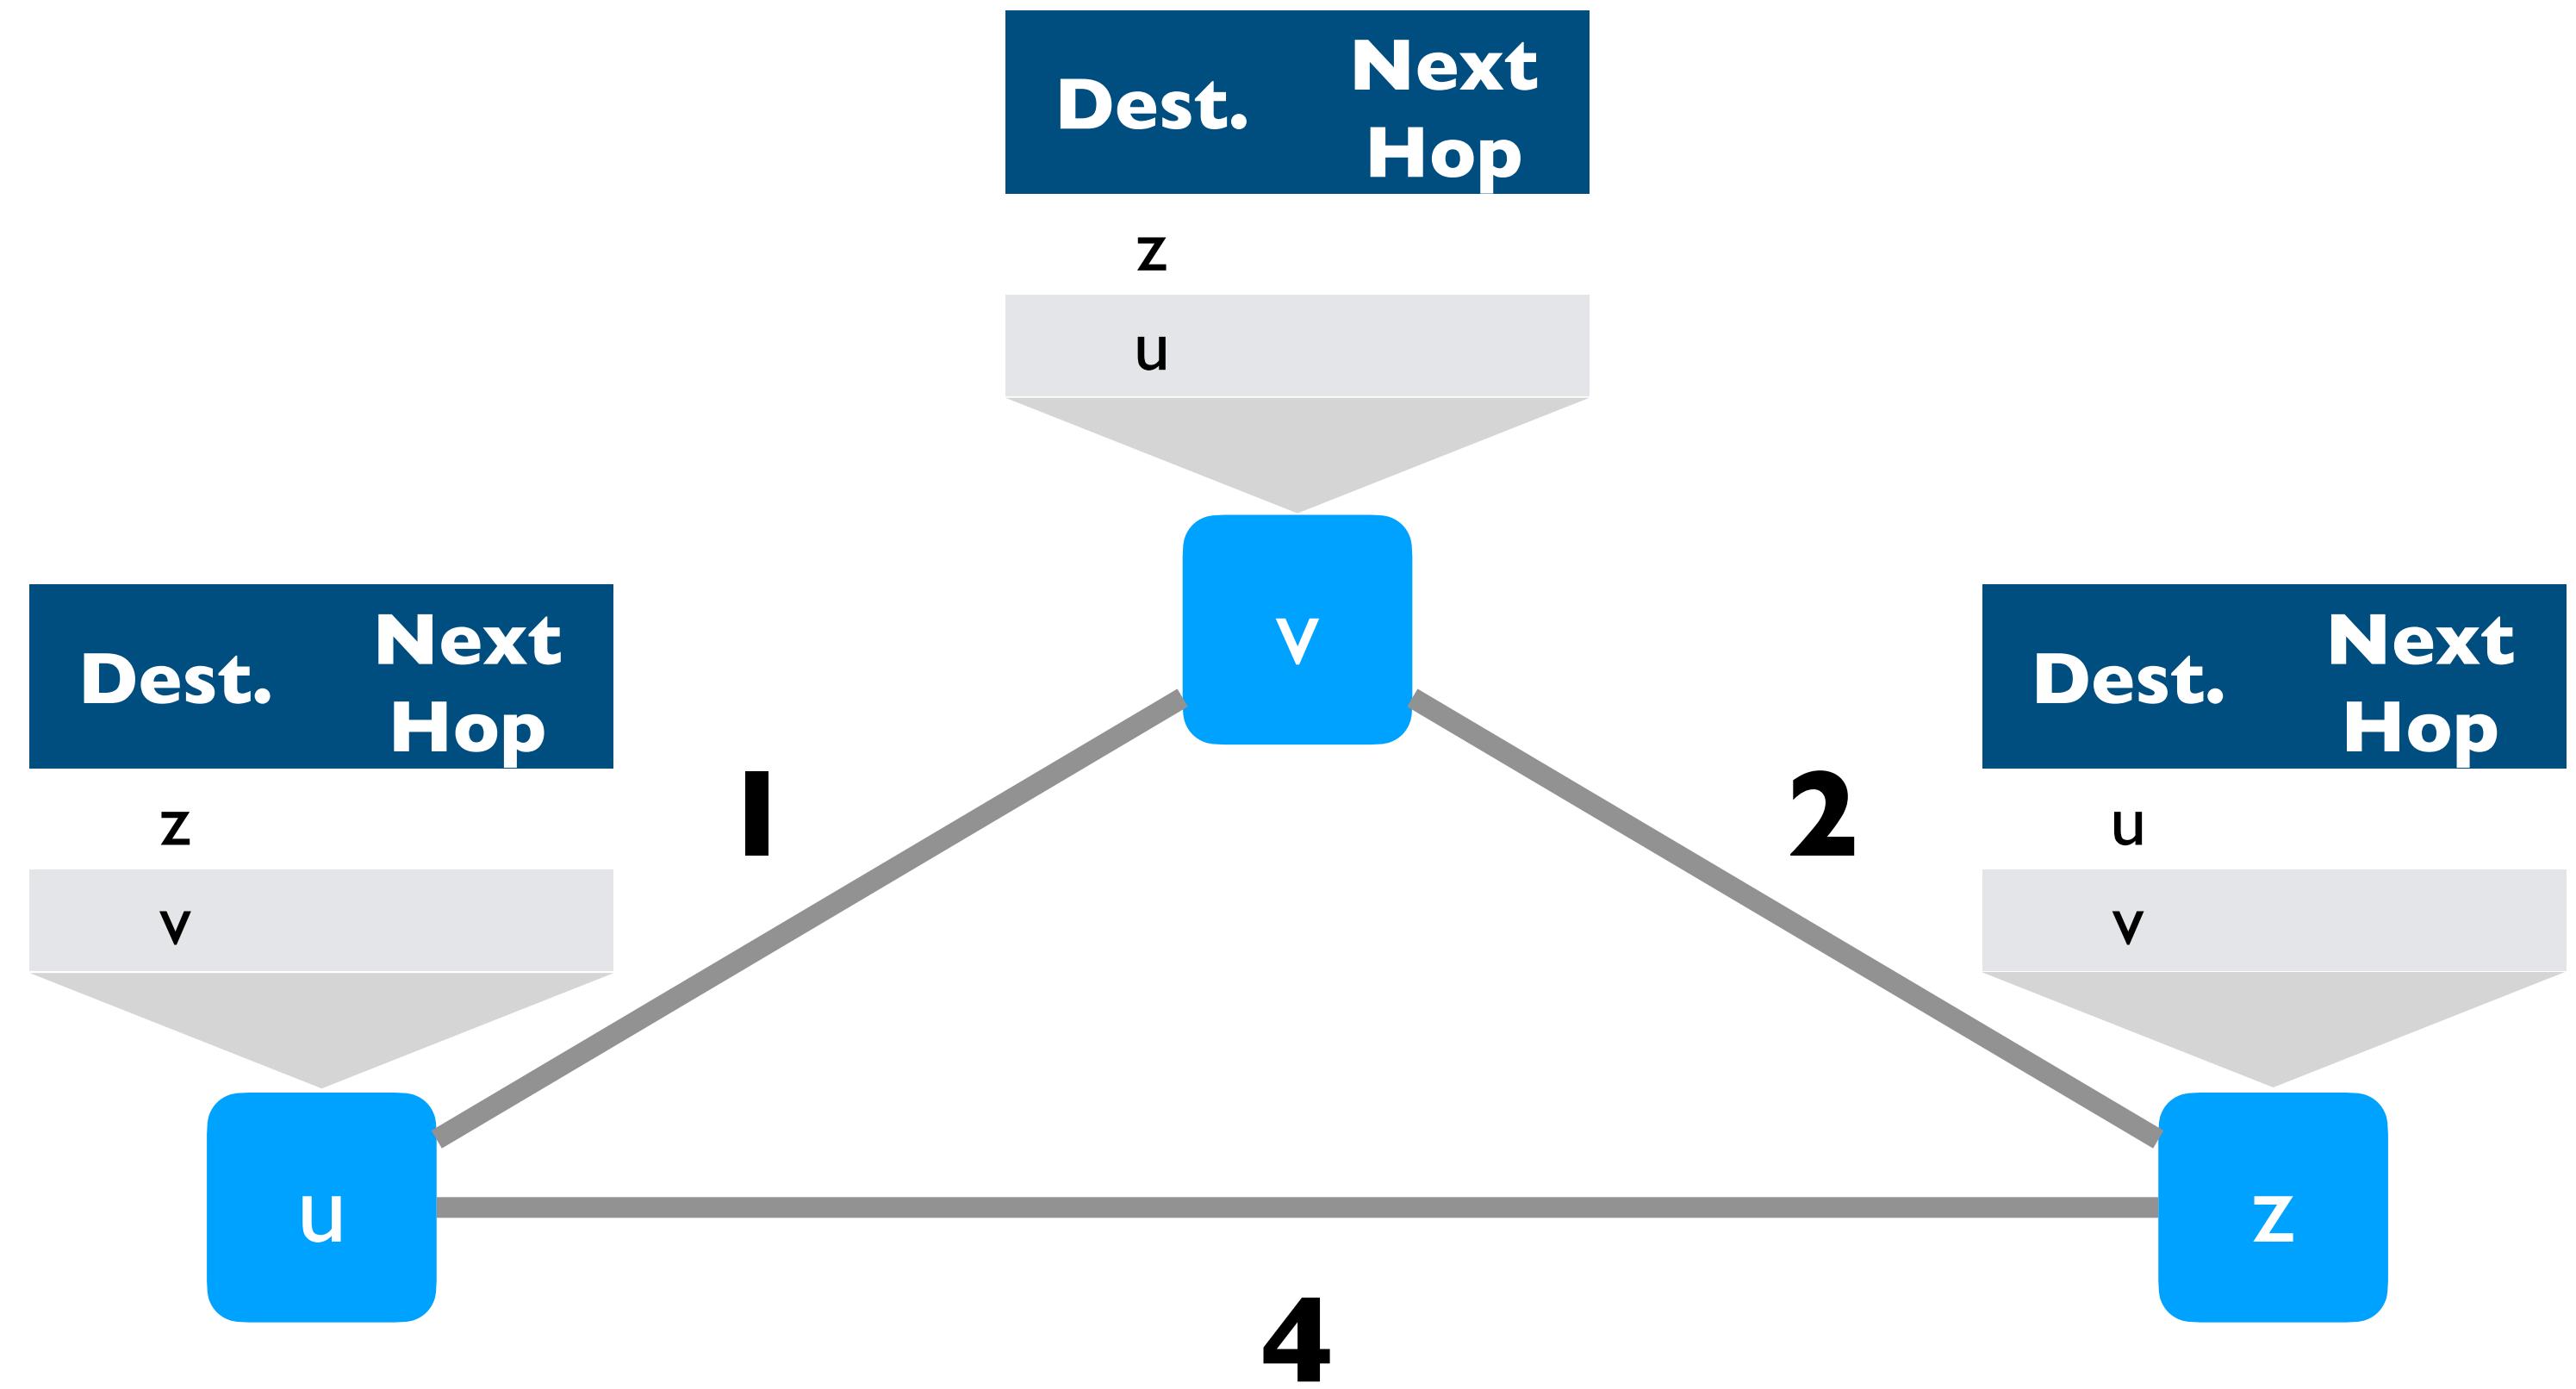


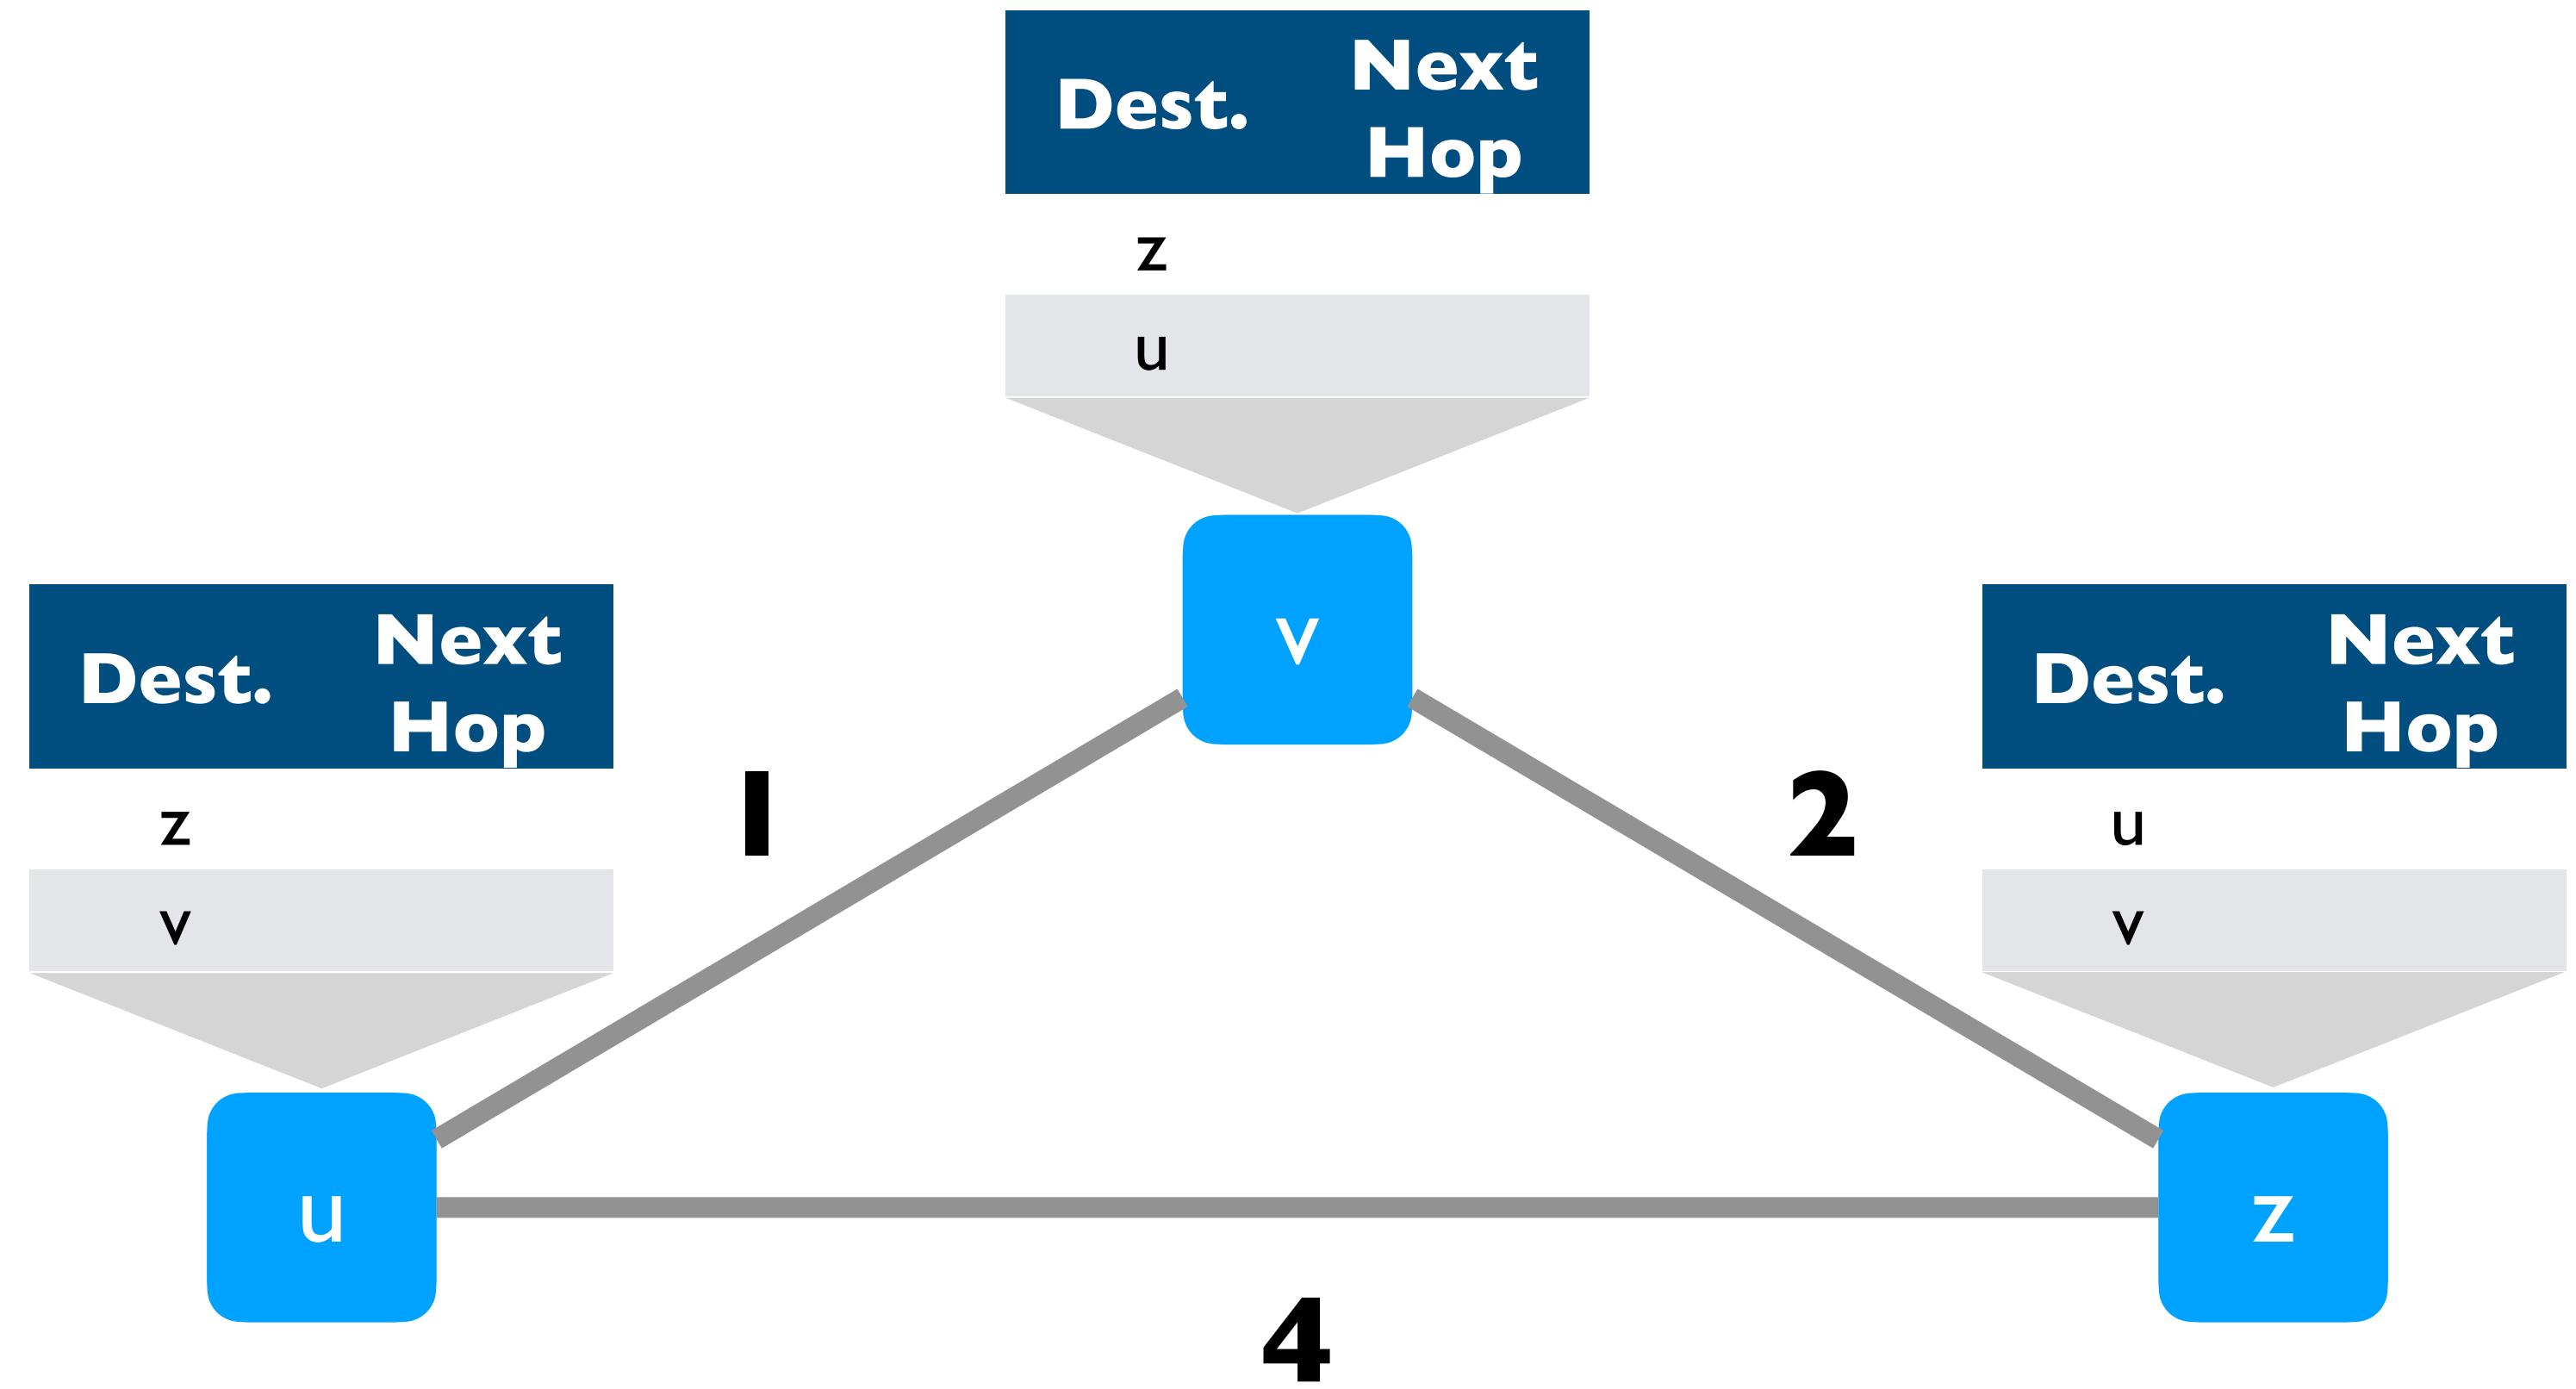




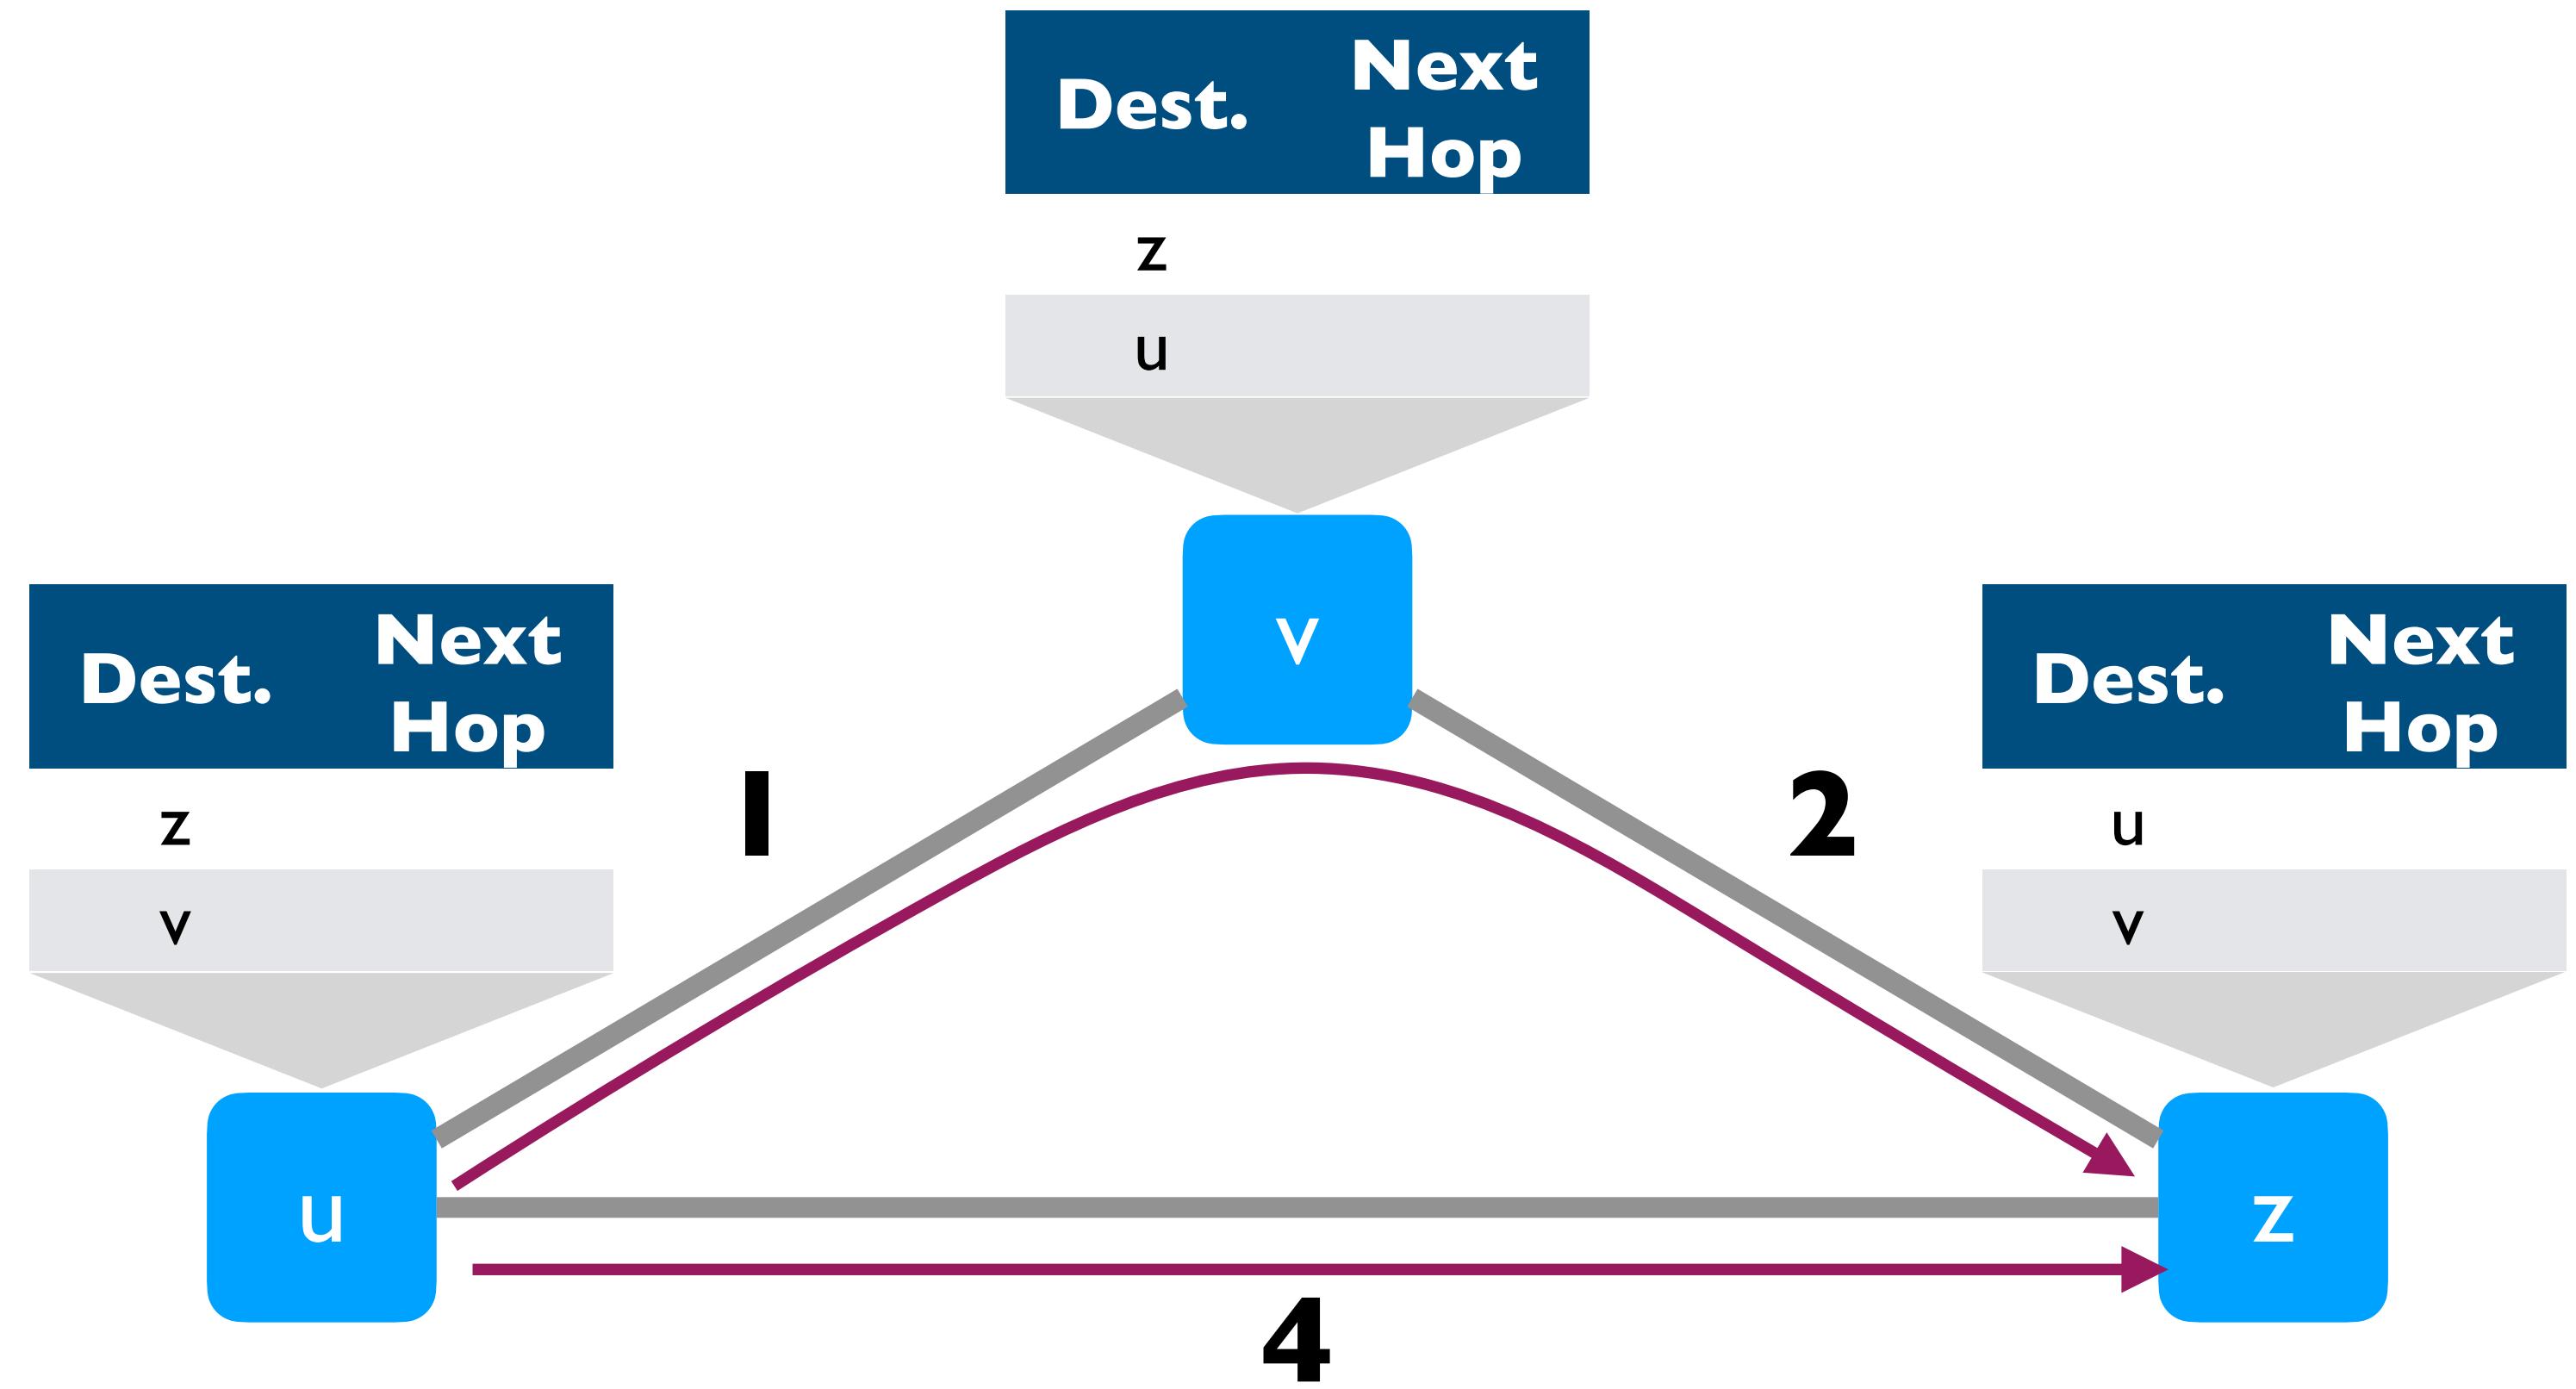
Could represent:

- Propagation delay
- Load
- Cost

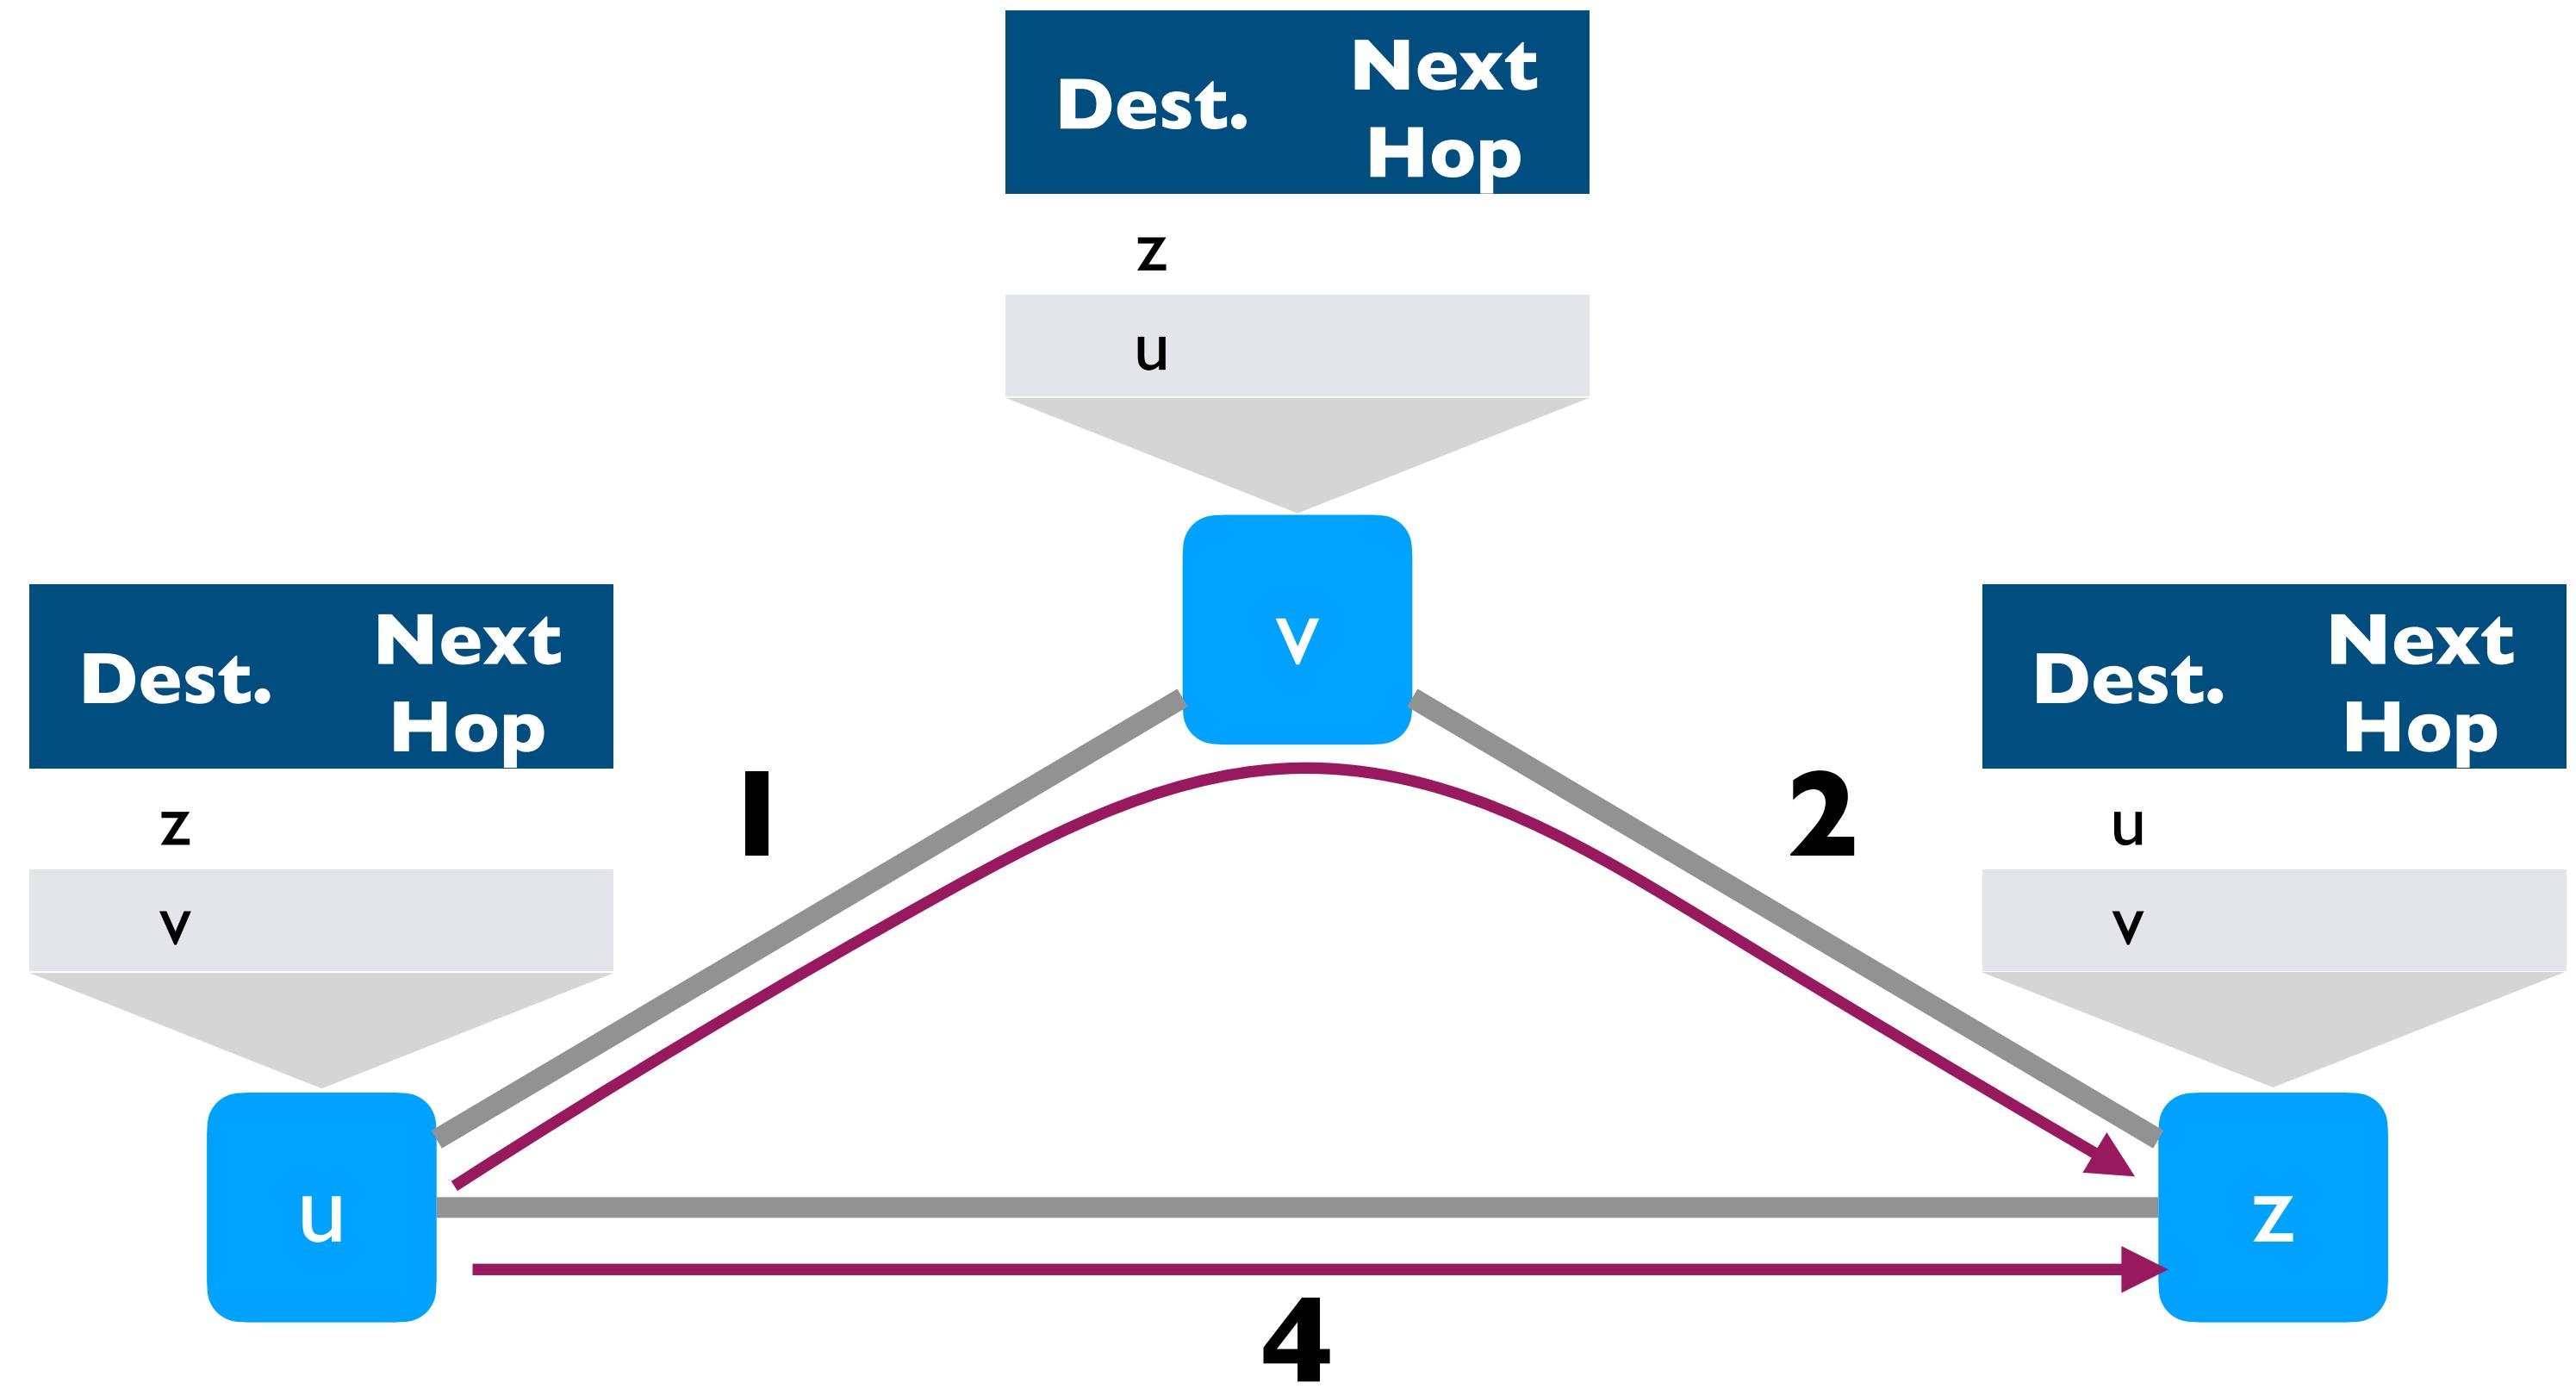




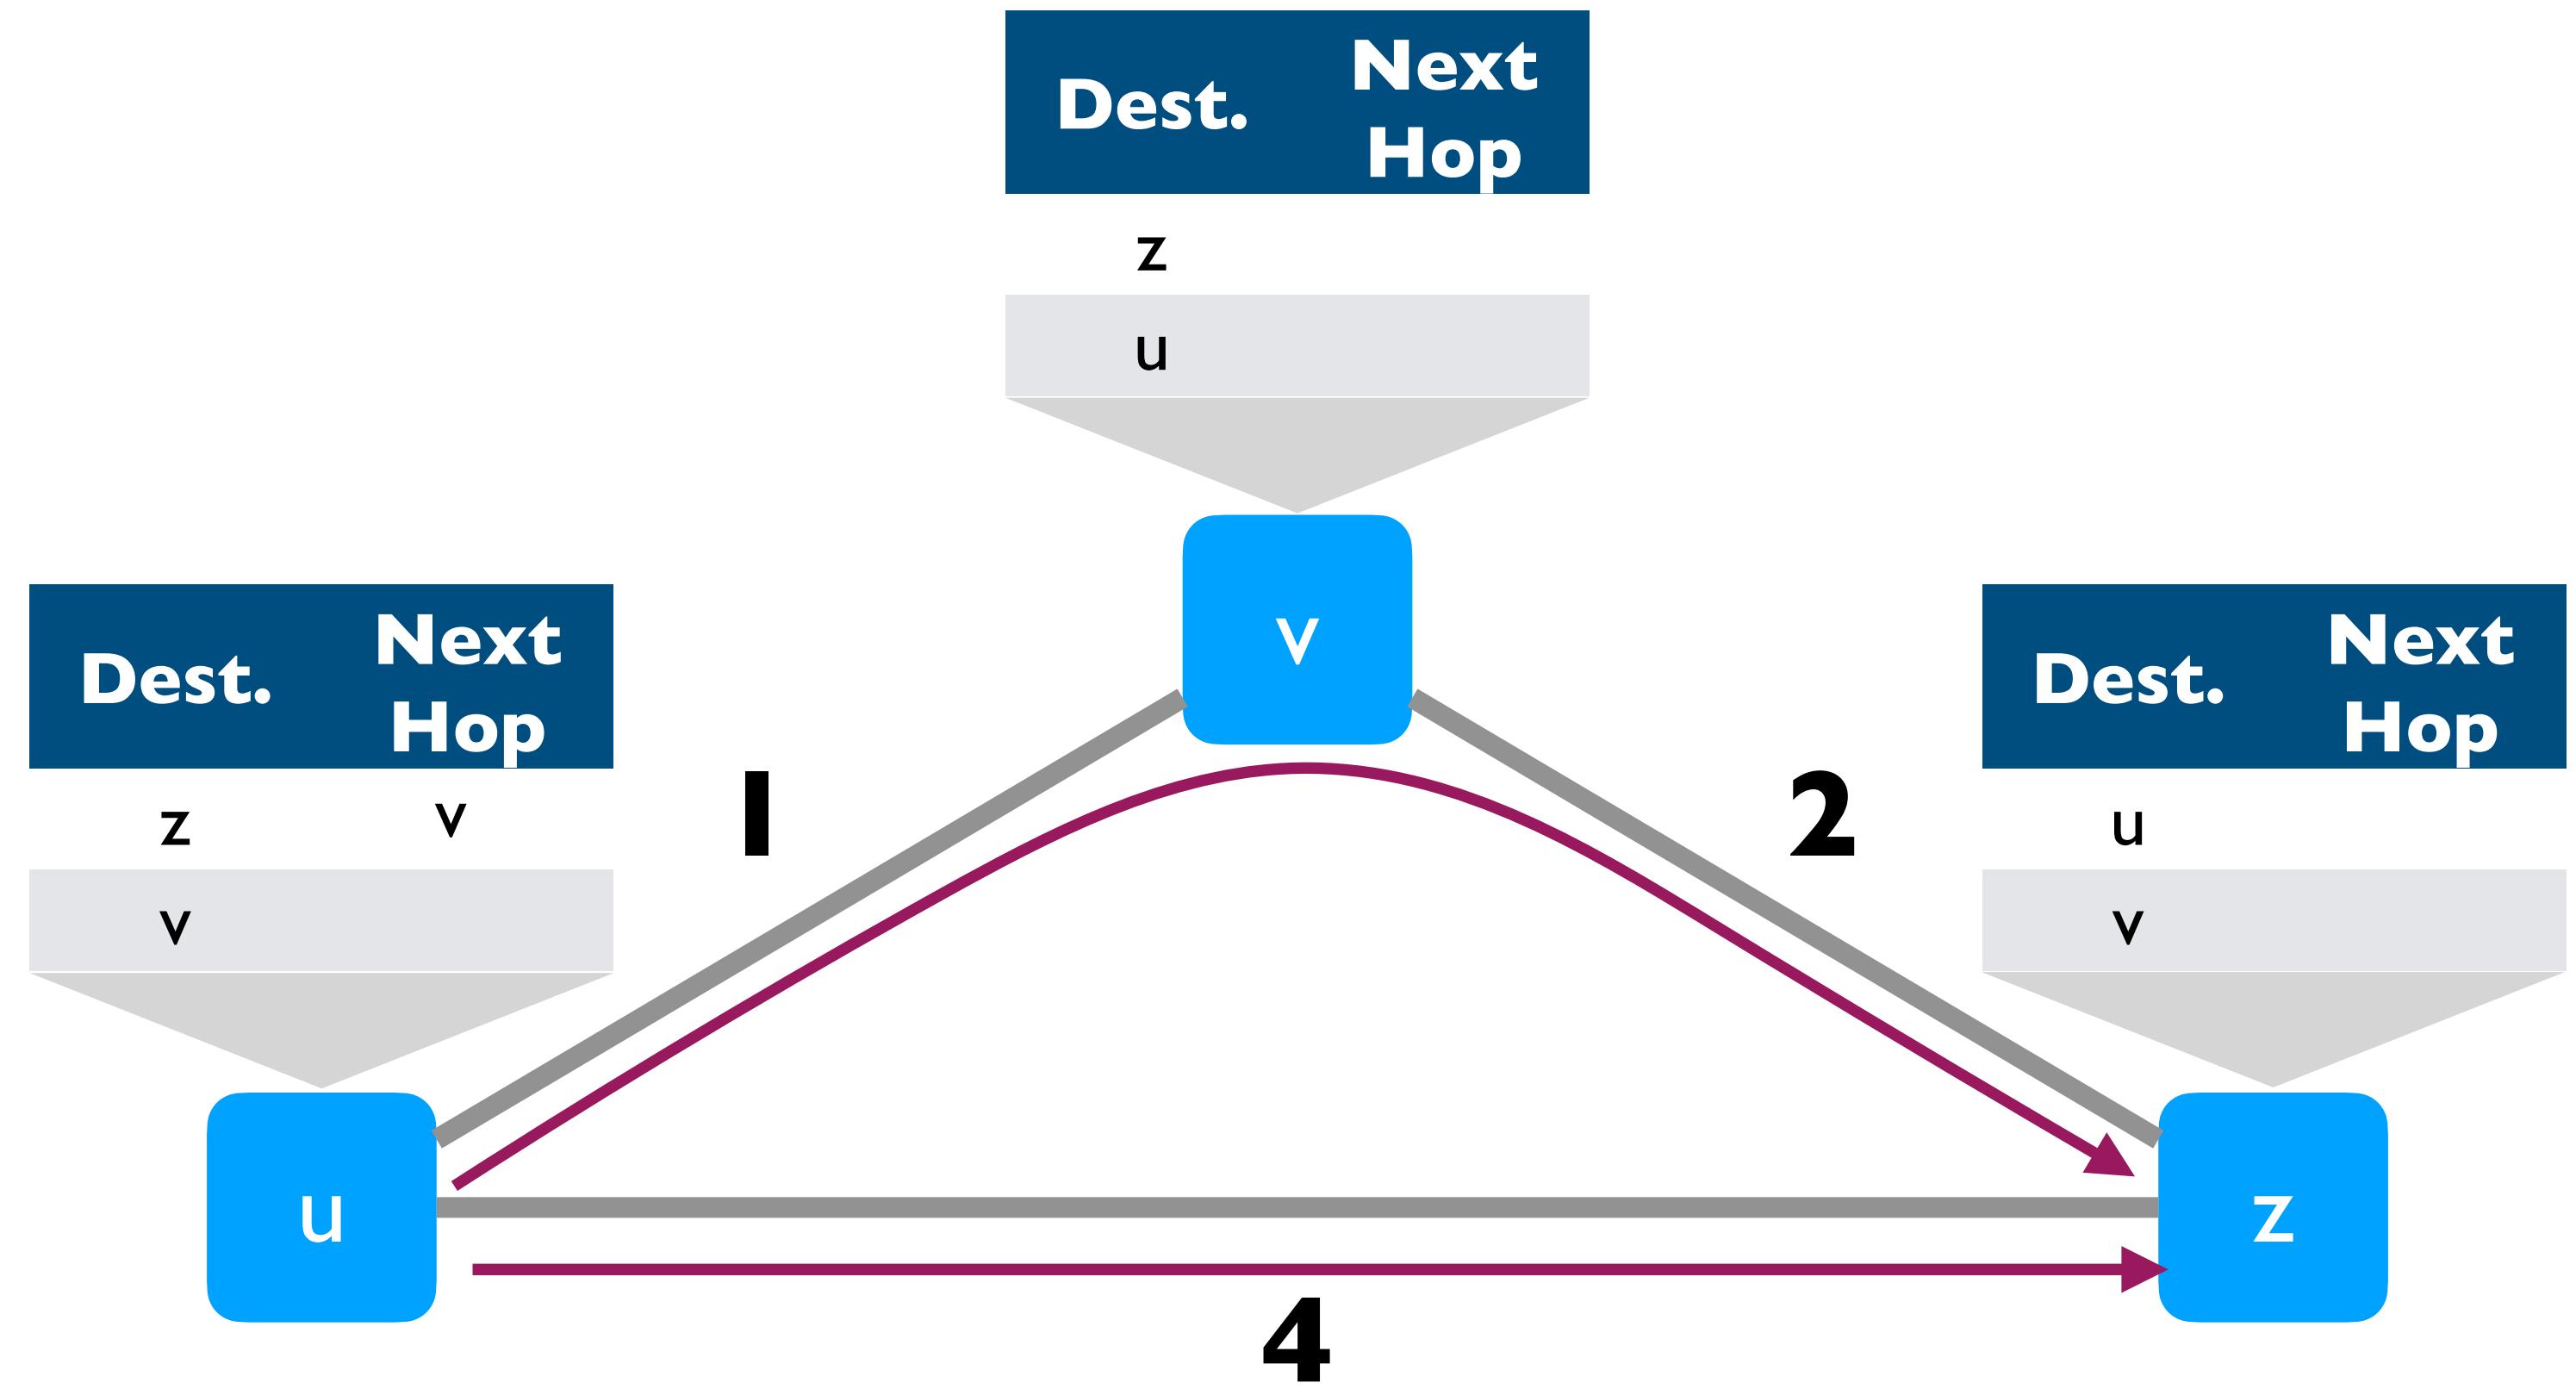
Least-cost path from **u** to **z**:



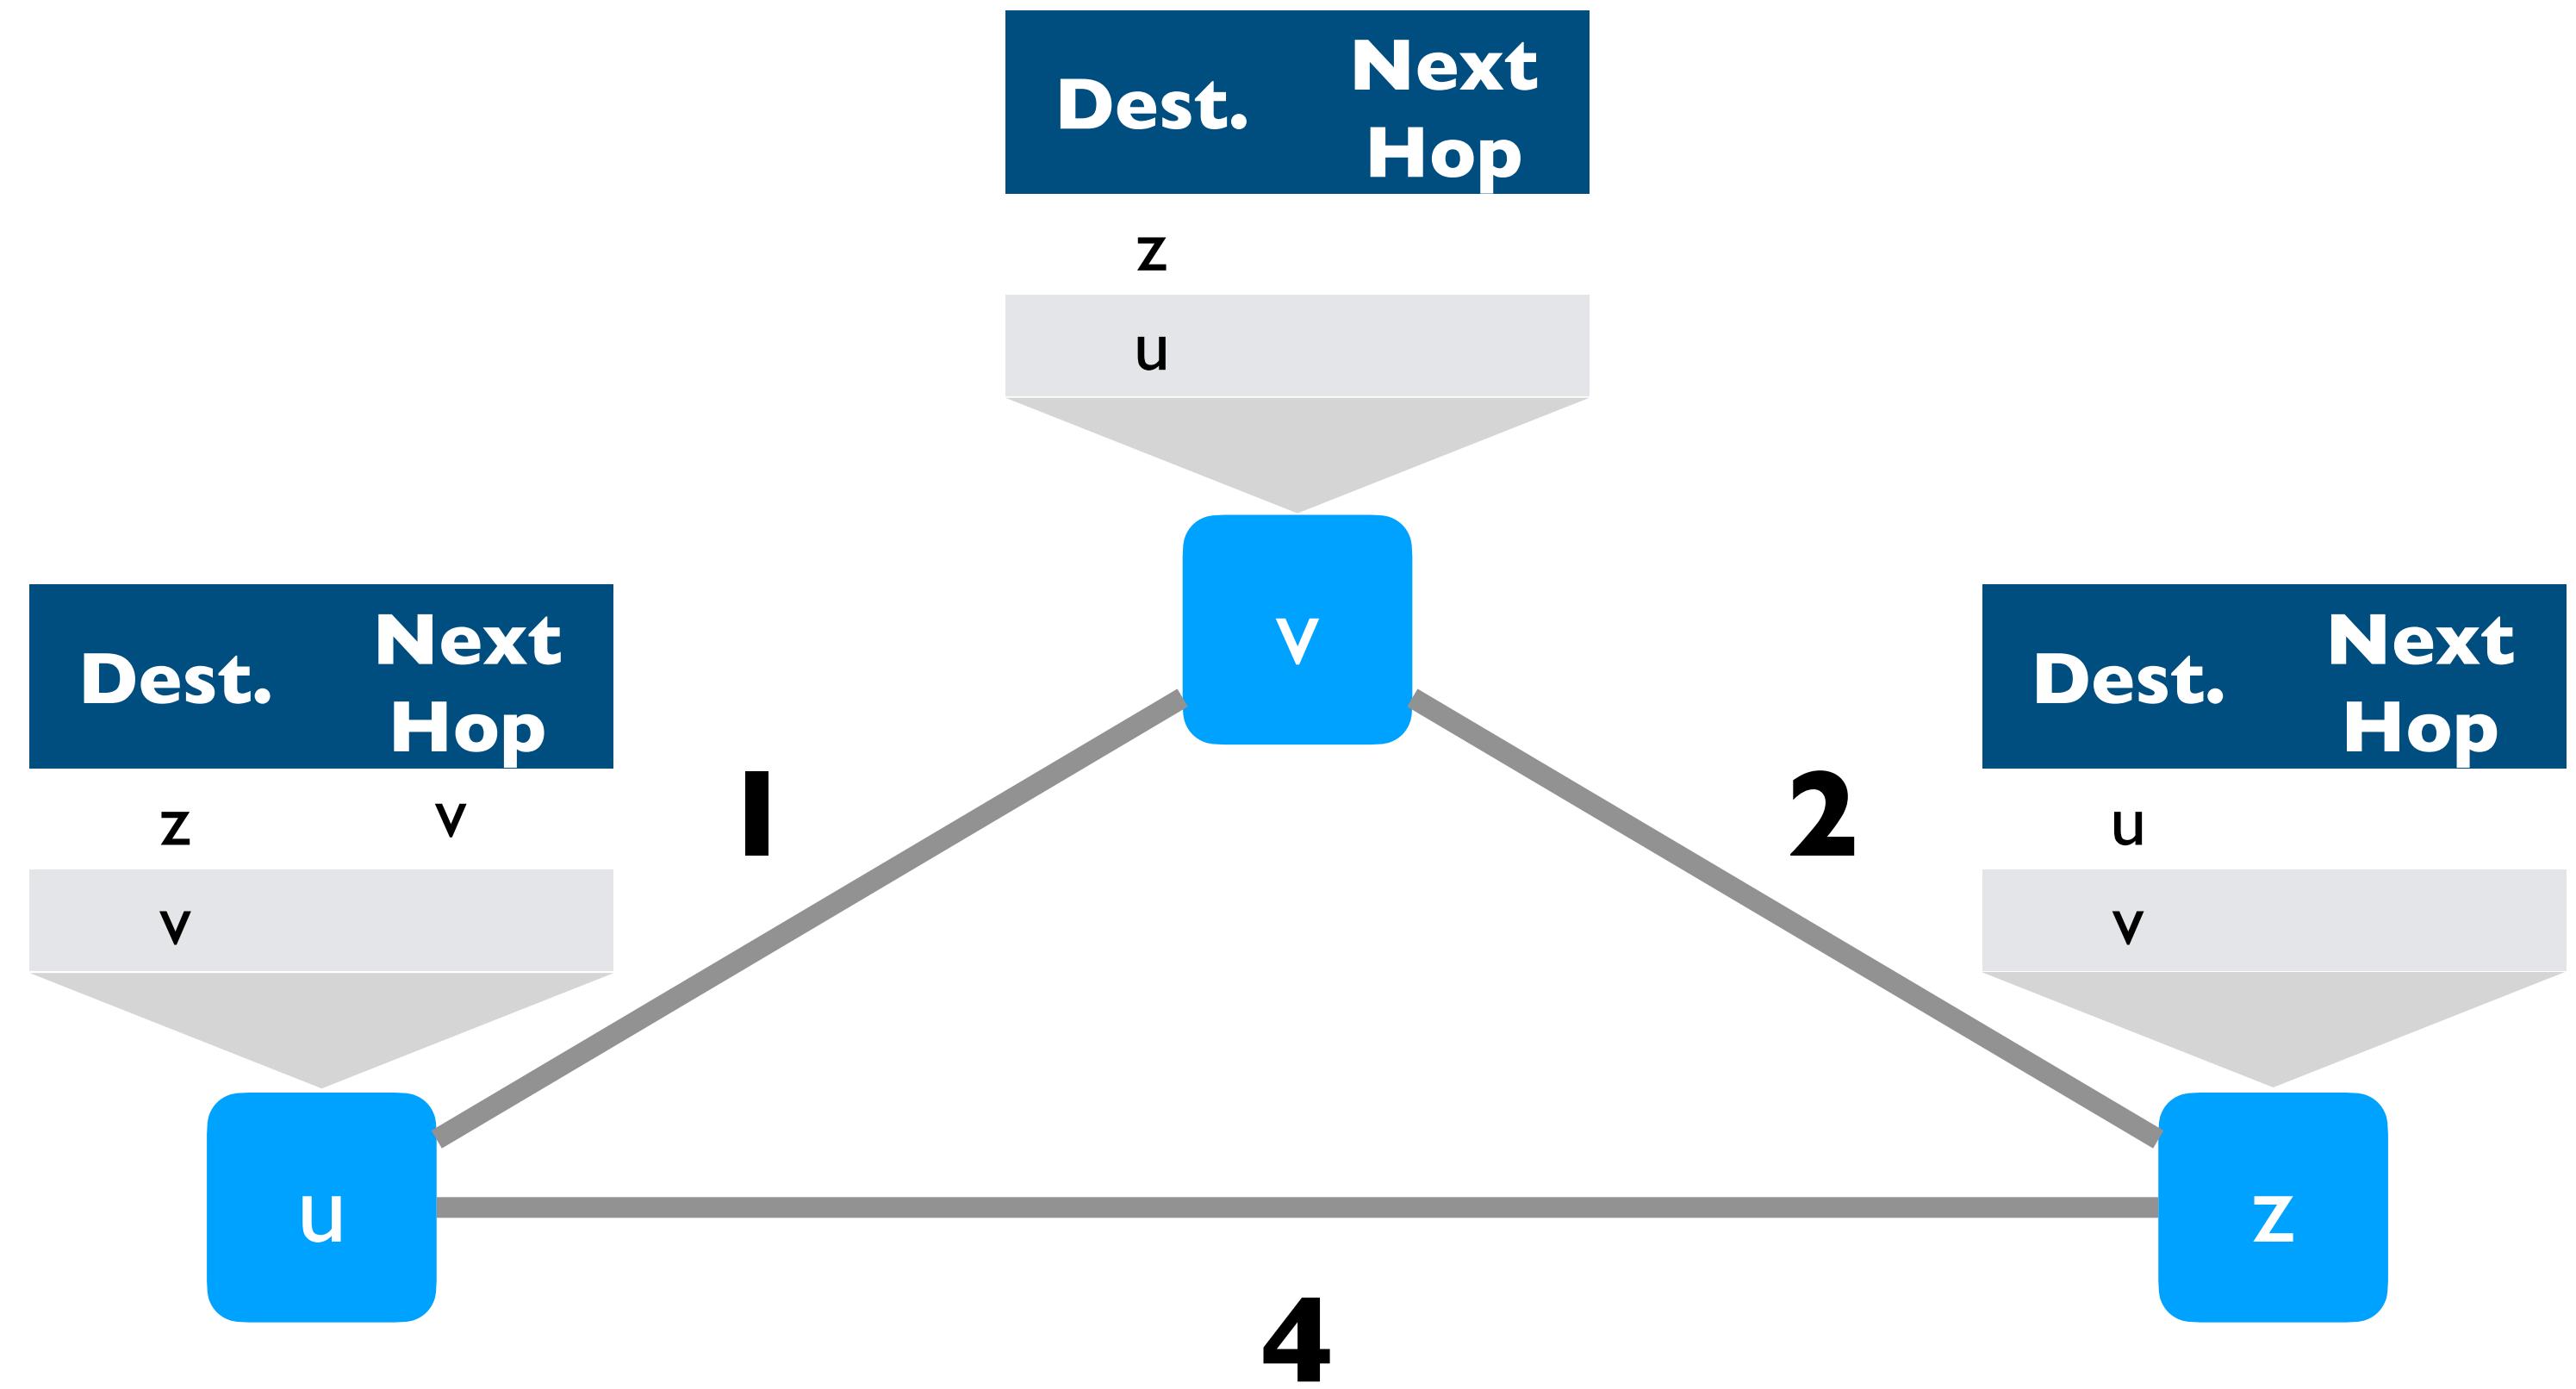
Least-cost path from **u** to **z**:



Least-cost path from **u** to **z**: **u v z**

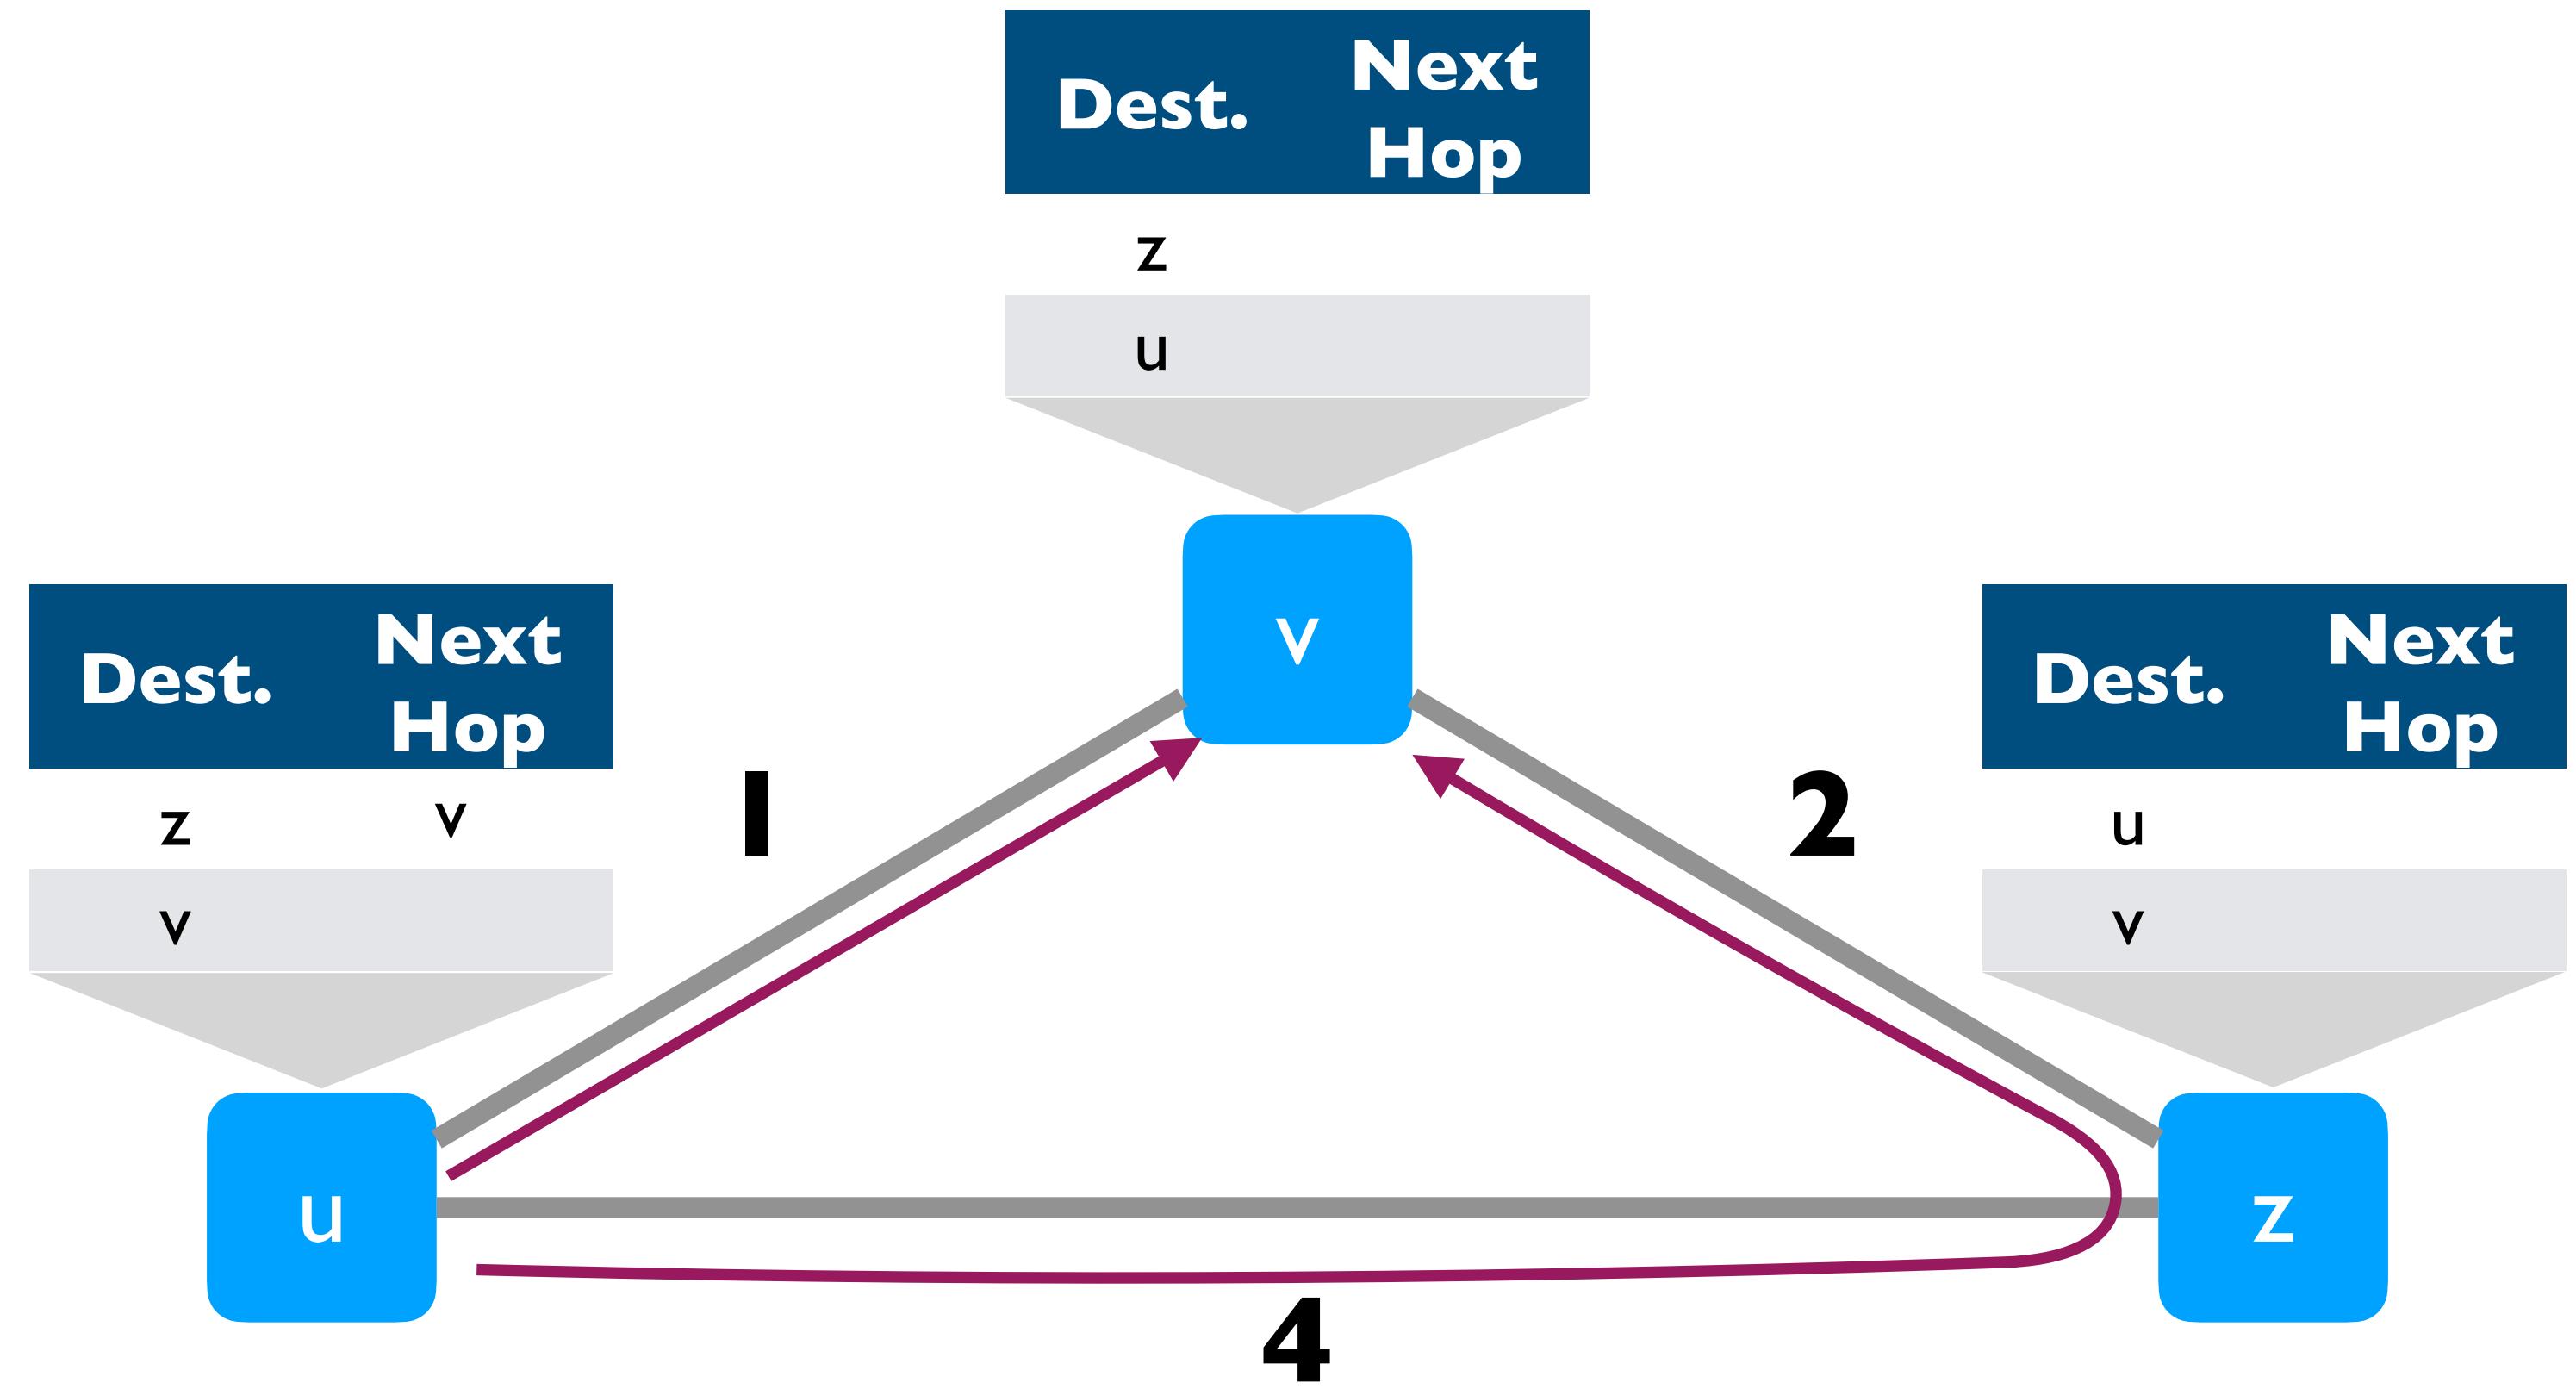


Least-cost path from **u** to **z**: **u v z**



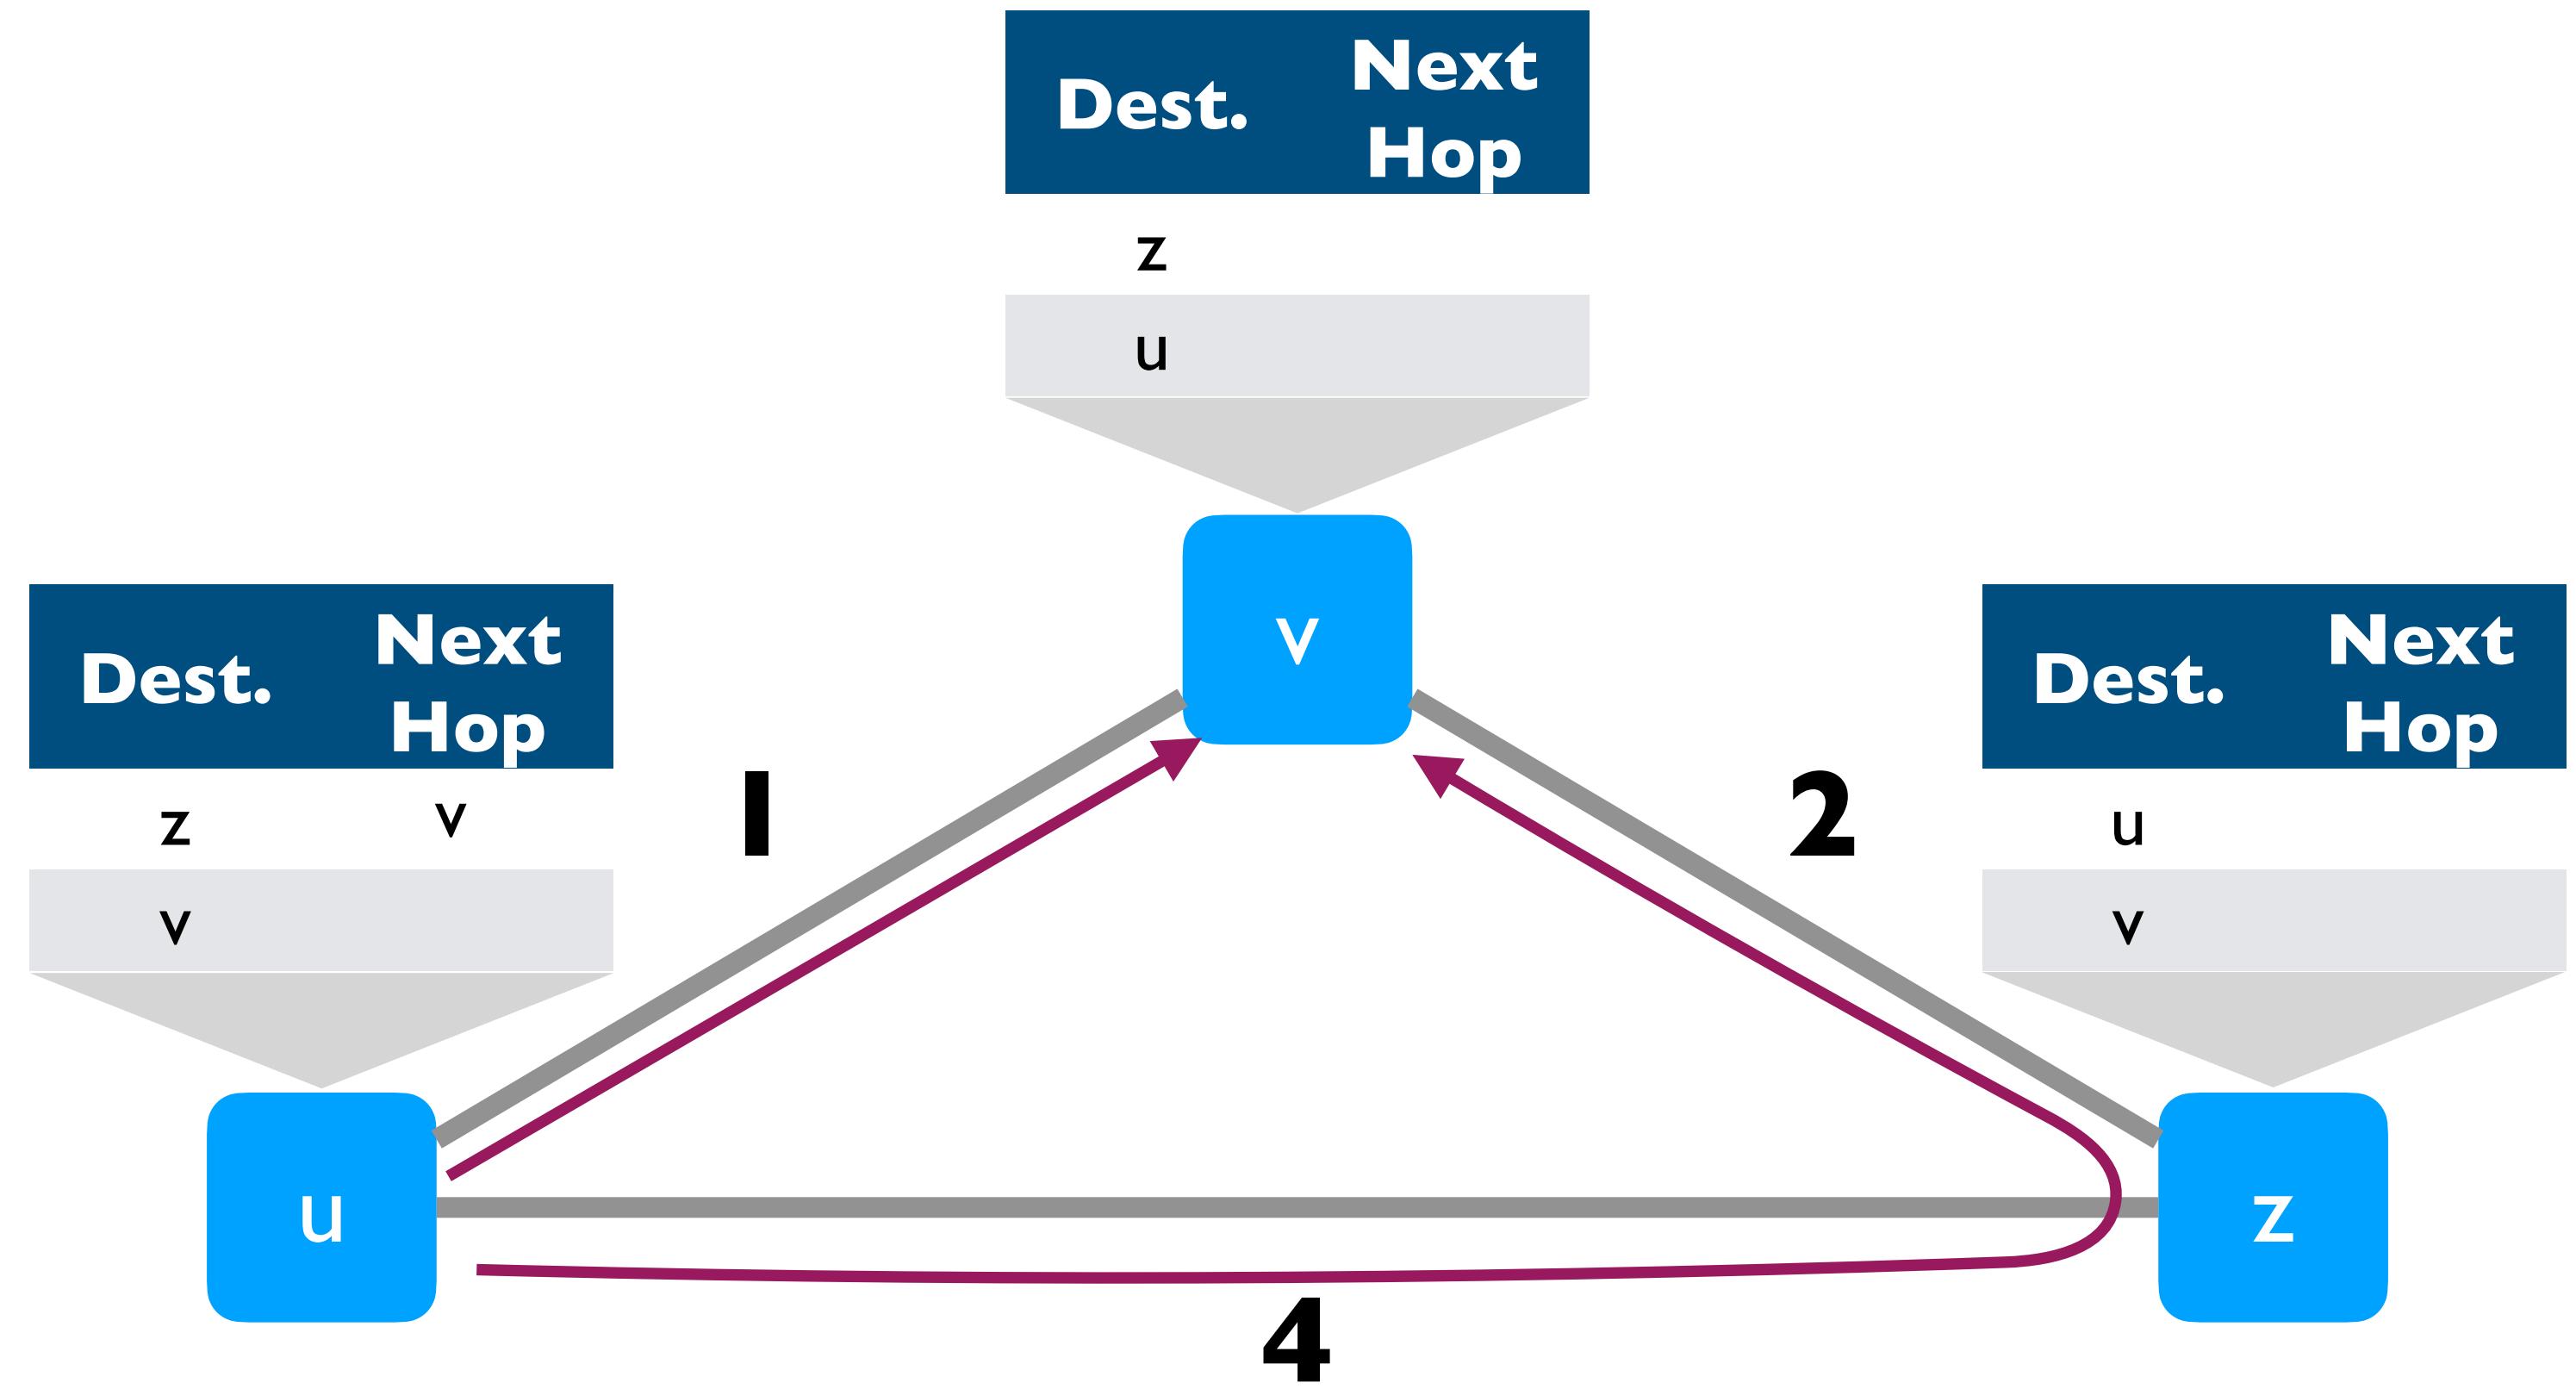
Least-cost path from **u** to **z**: **u v z**

Least-cost path from **u** to **v**:



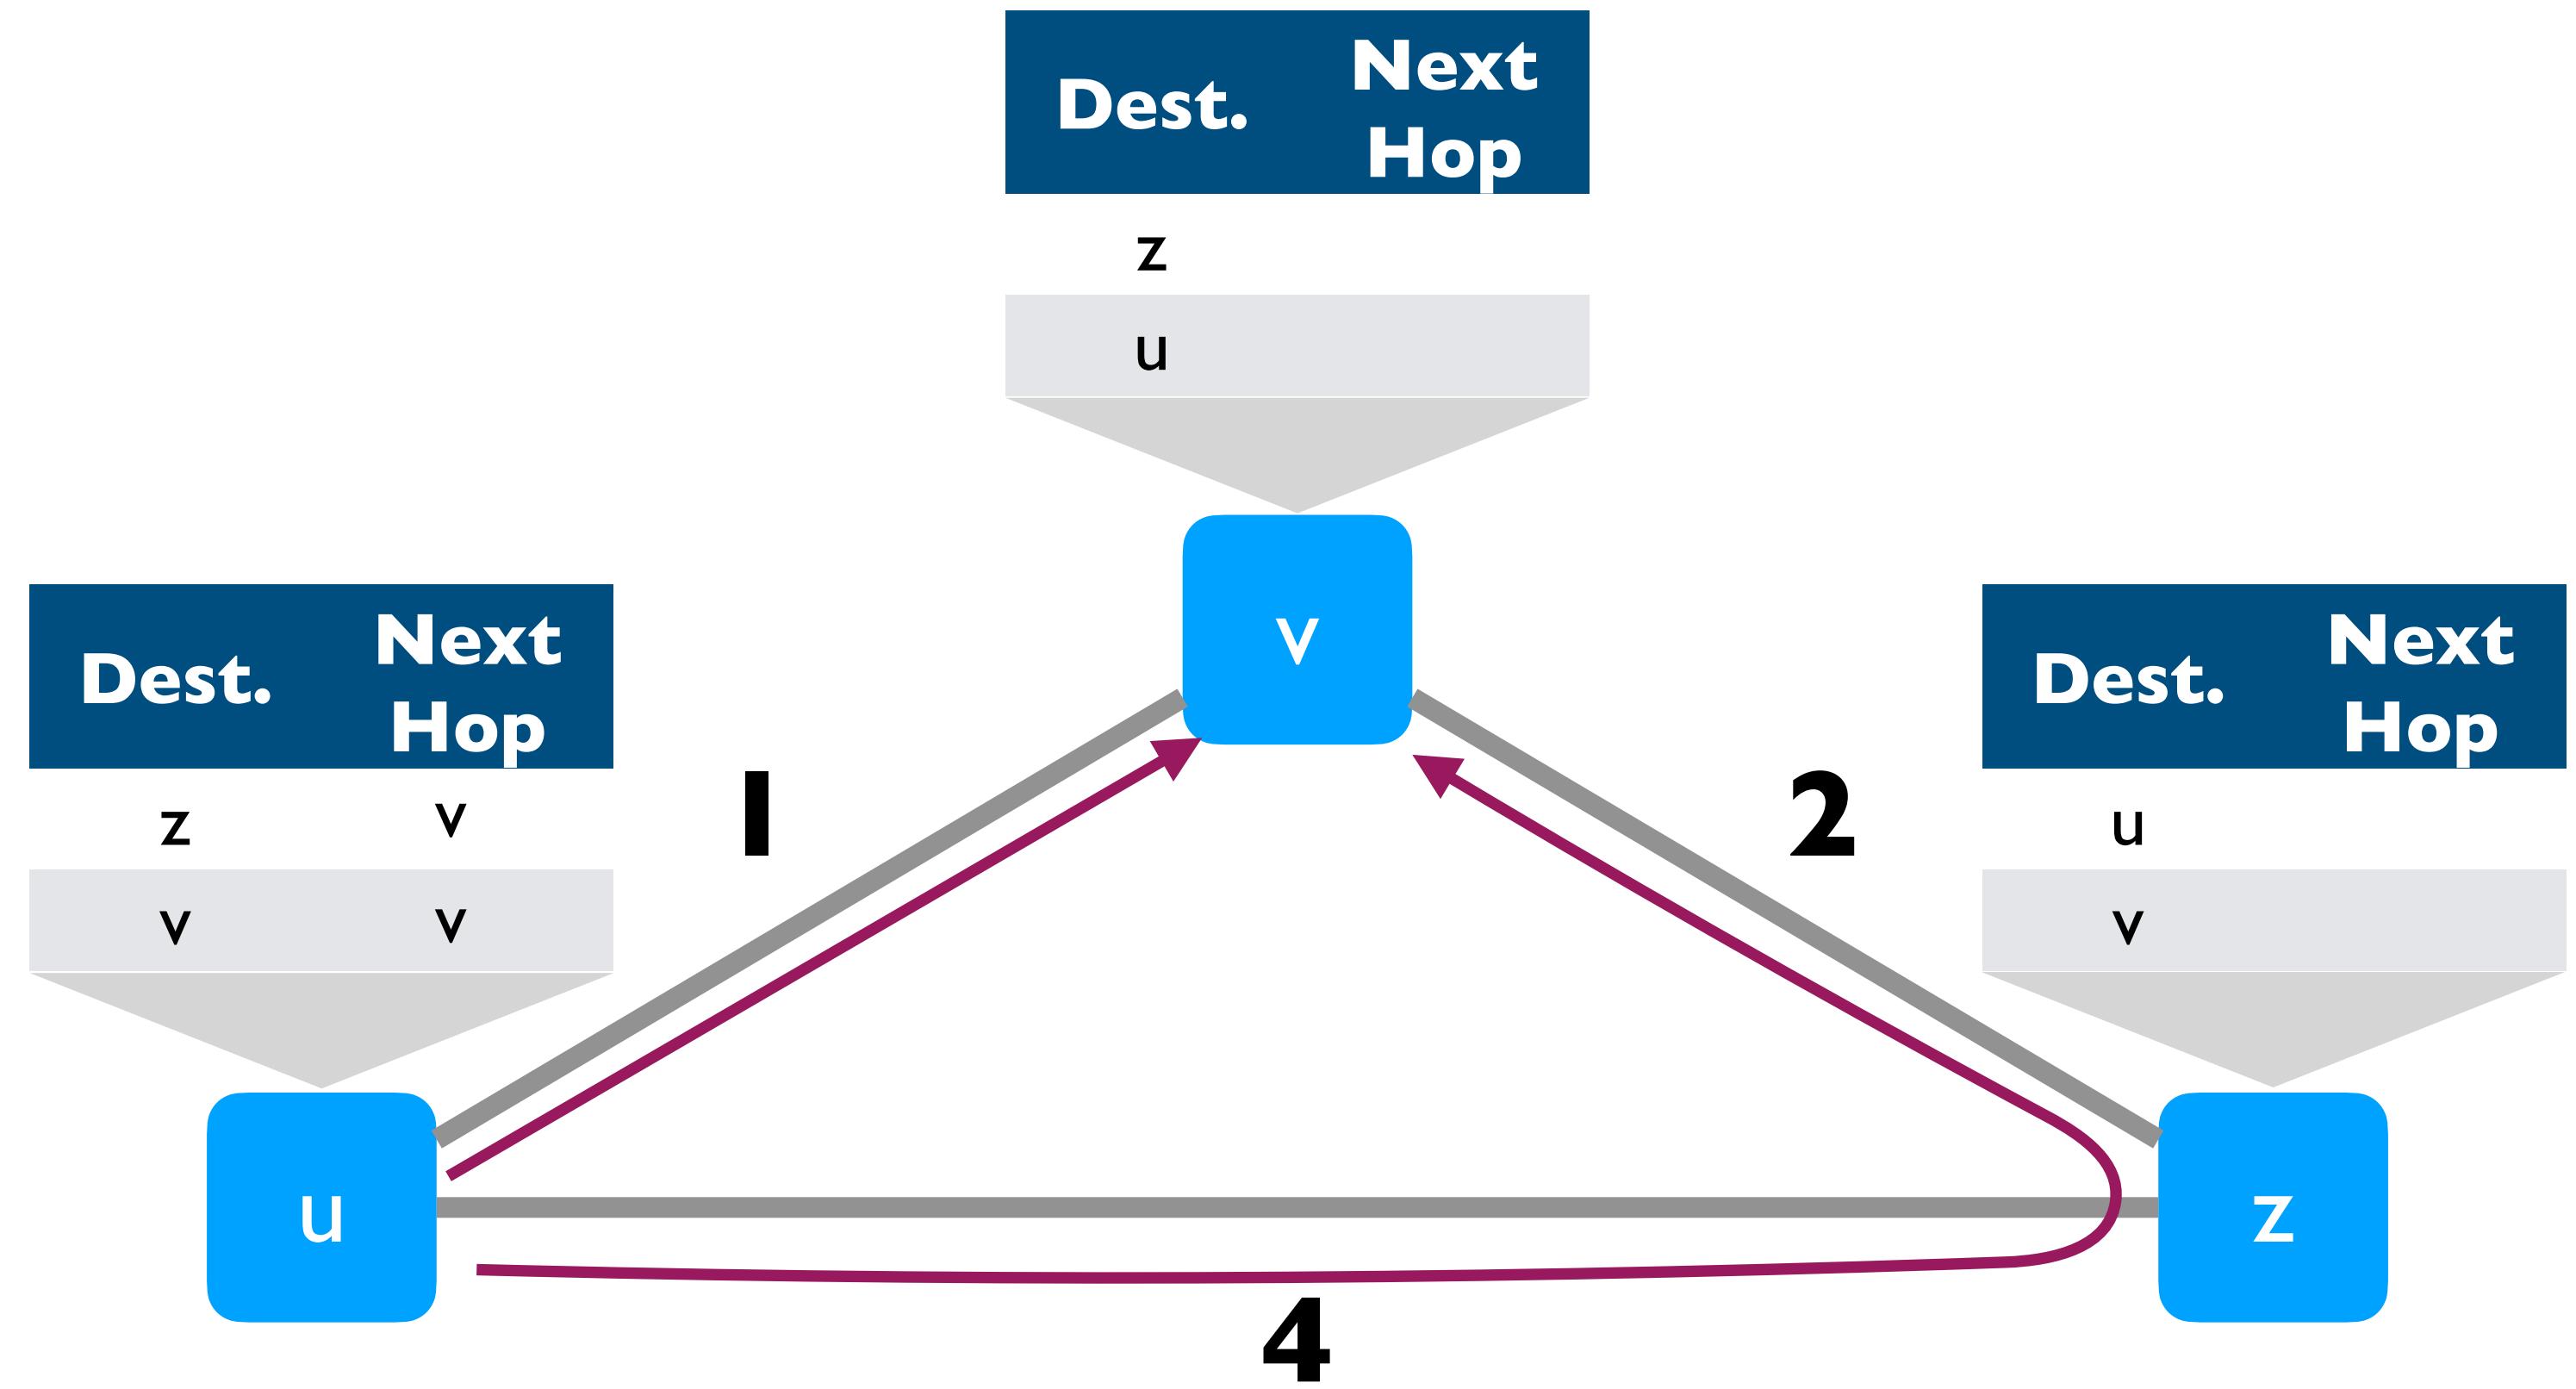
Least-cost path from **u** to **z**: **u v z**

Least-cost path from **u** to **v**:



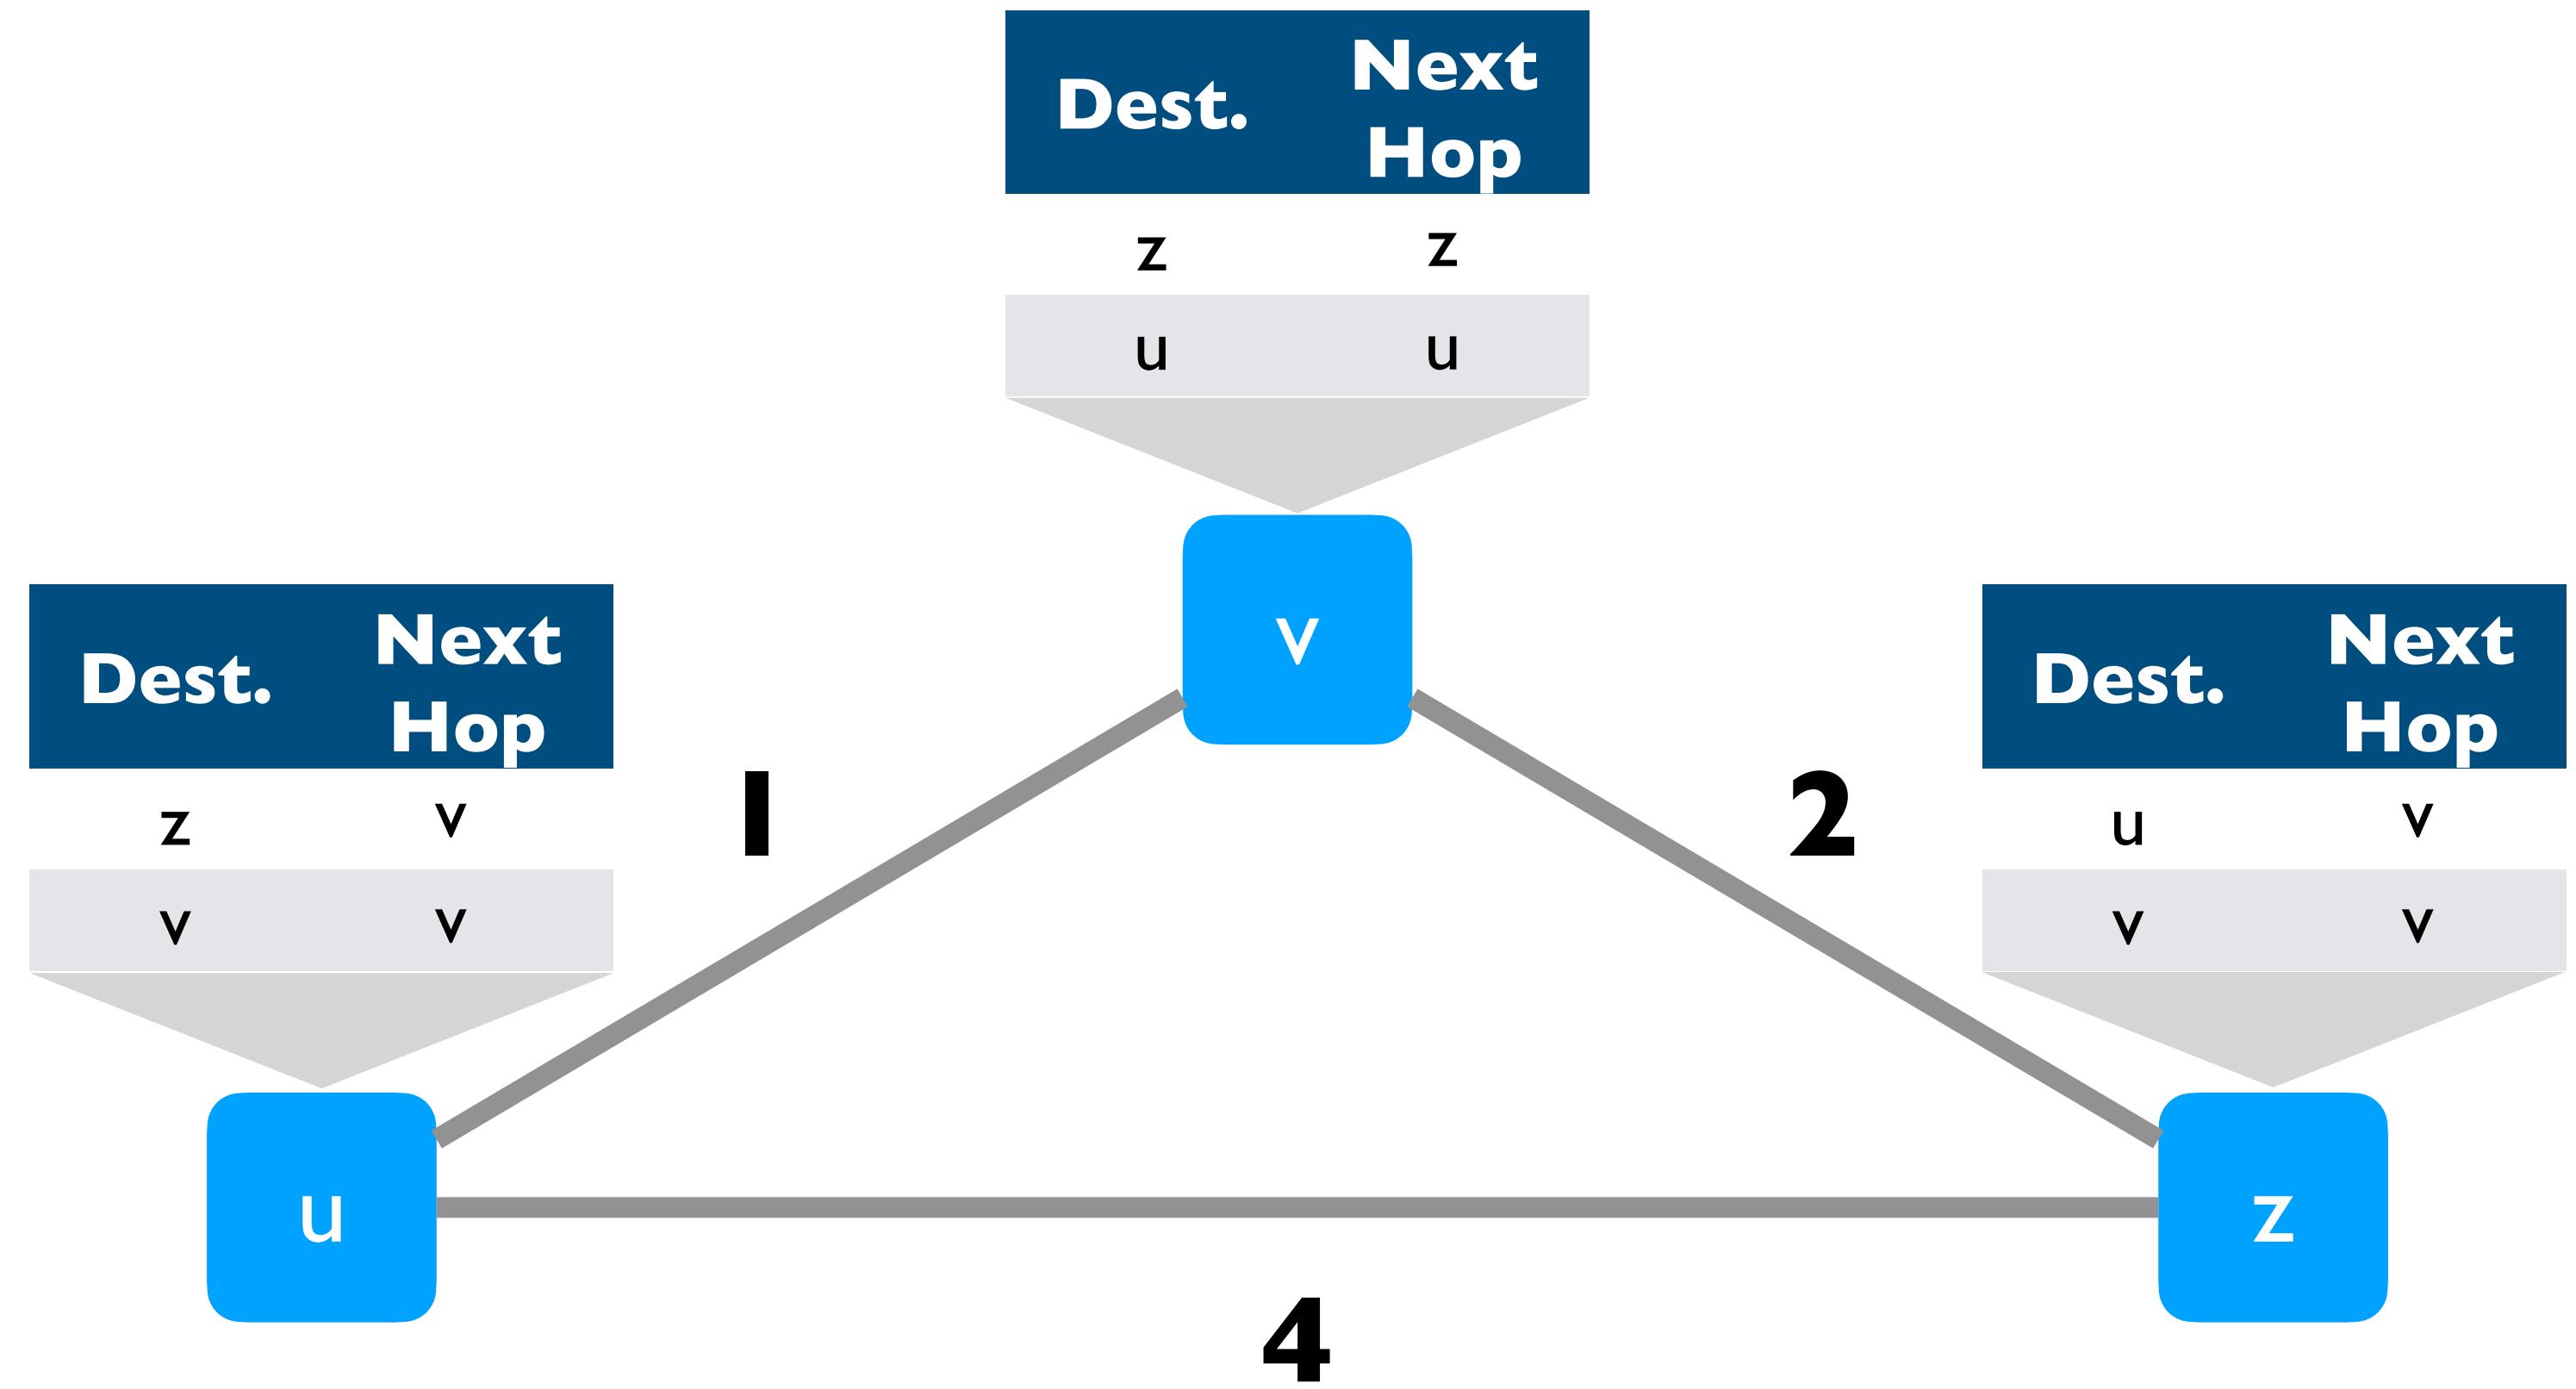
Least-cost path from **u** to **z**: **u v z**

Least-cost path from **u** to **v**: **u v**



Least-cost path from **u** to **z**: **u v z**

Least-cost path from **u** to **v**: **u v**



Least-cost path from **u** to **z**: **u v z**

Least-cost path from **u** to **v**: **u v**

# Least-cost path routing

# Least-cost path routing

- **Given:** Router graph & link costs

# Least-cost path routing

- **Given:** Router graph & link costs
- **Goal:** find least-cost path  
from each source router  
to each destination router

# Least-cost path routing

- **Given:** Router graph & link costs
- **Goal:** find least-cost path  
from each source router  
to each destination router
- **Next time on CPSC 433/533:** How we compute these!