

# TCP Congestion Control Advanced Techniques (Contd.)

CPSC 433/533, Spring 2021

Anurag Khandelwal

# Router-assisted Congestion Control

- Three tasks for CC:
  - Isolation/fairness
  - Adjustment
  - Detecting congestion

# Max-Min Fairness

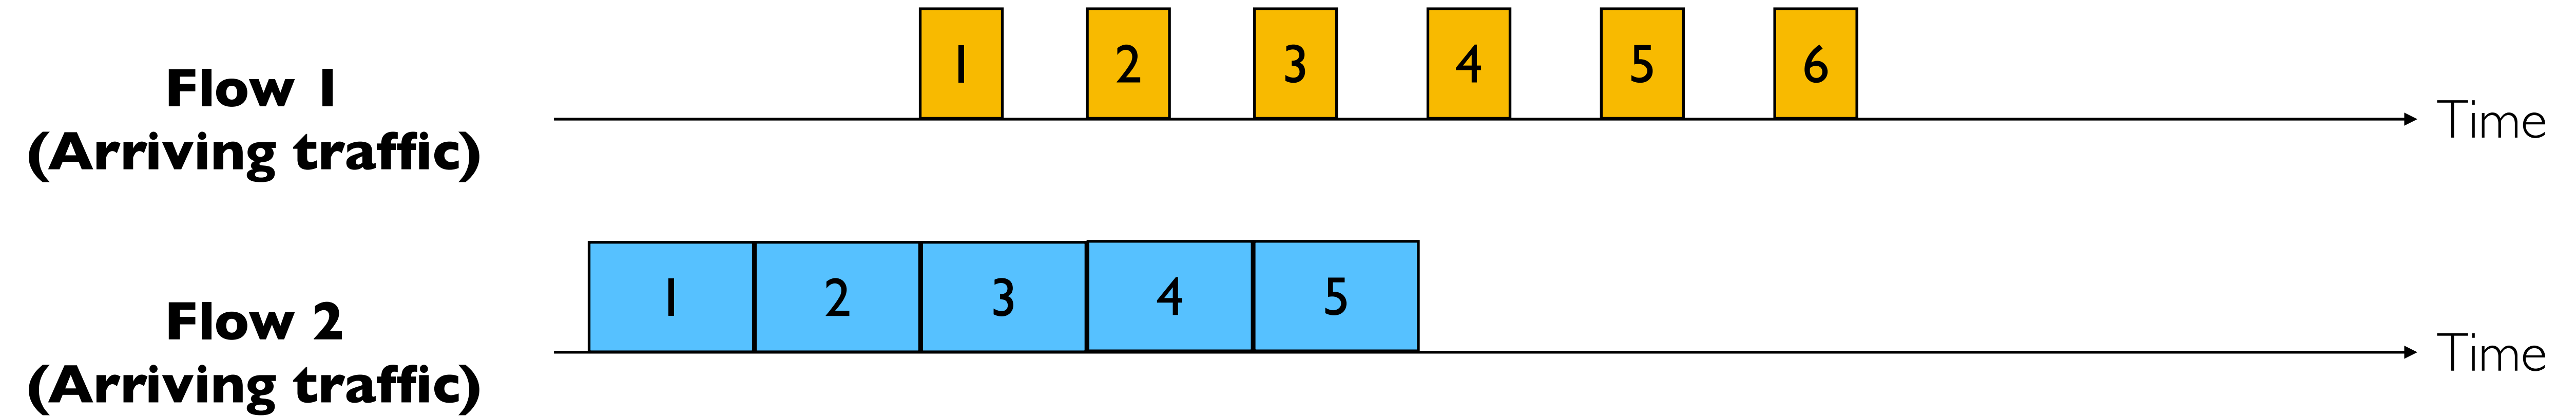
- Given a set of bandwidth demands  $r_i$  and total bandwidth  $C$ , max-min bandwidth allocations are:

$$a_i = \min(f, r_i)$$

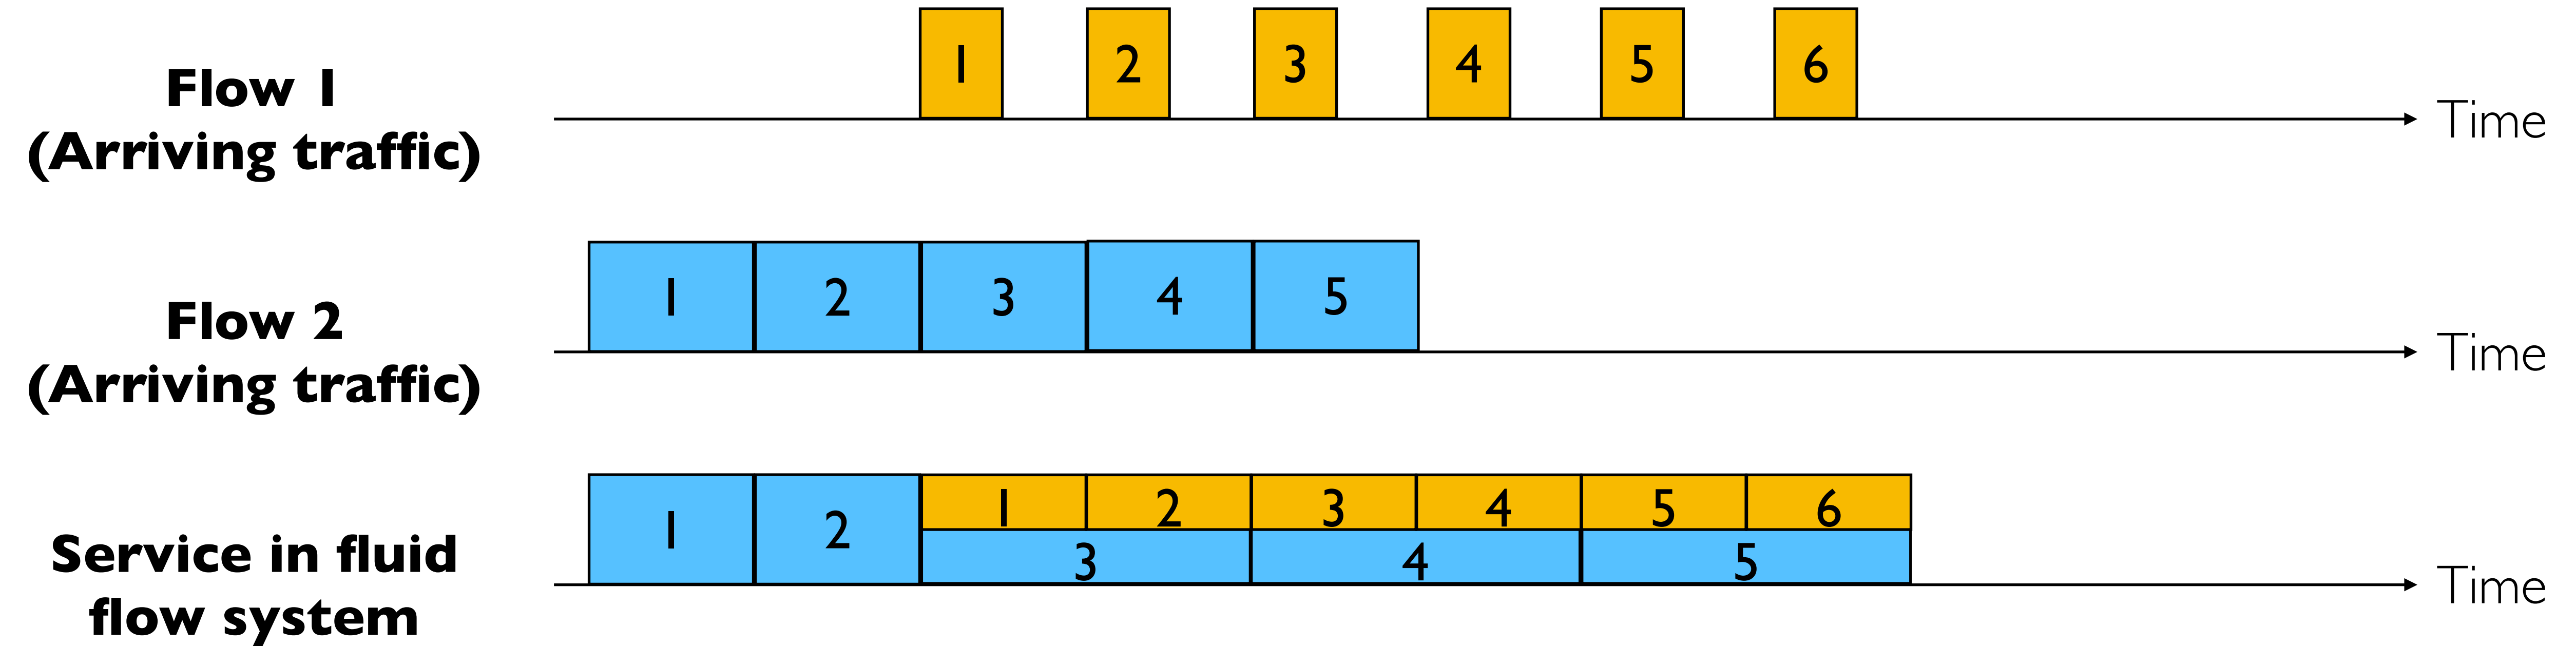
where  $f$  is the unique value such that  $\sum_i a_i = C$

- Property:
  - If you don't get full demand, no one gets more than you
- This is what round robin service gives if all packets are the same size

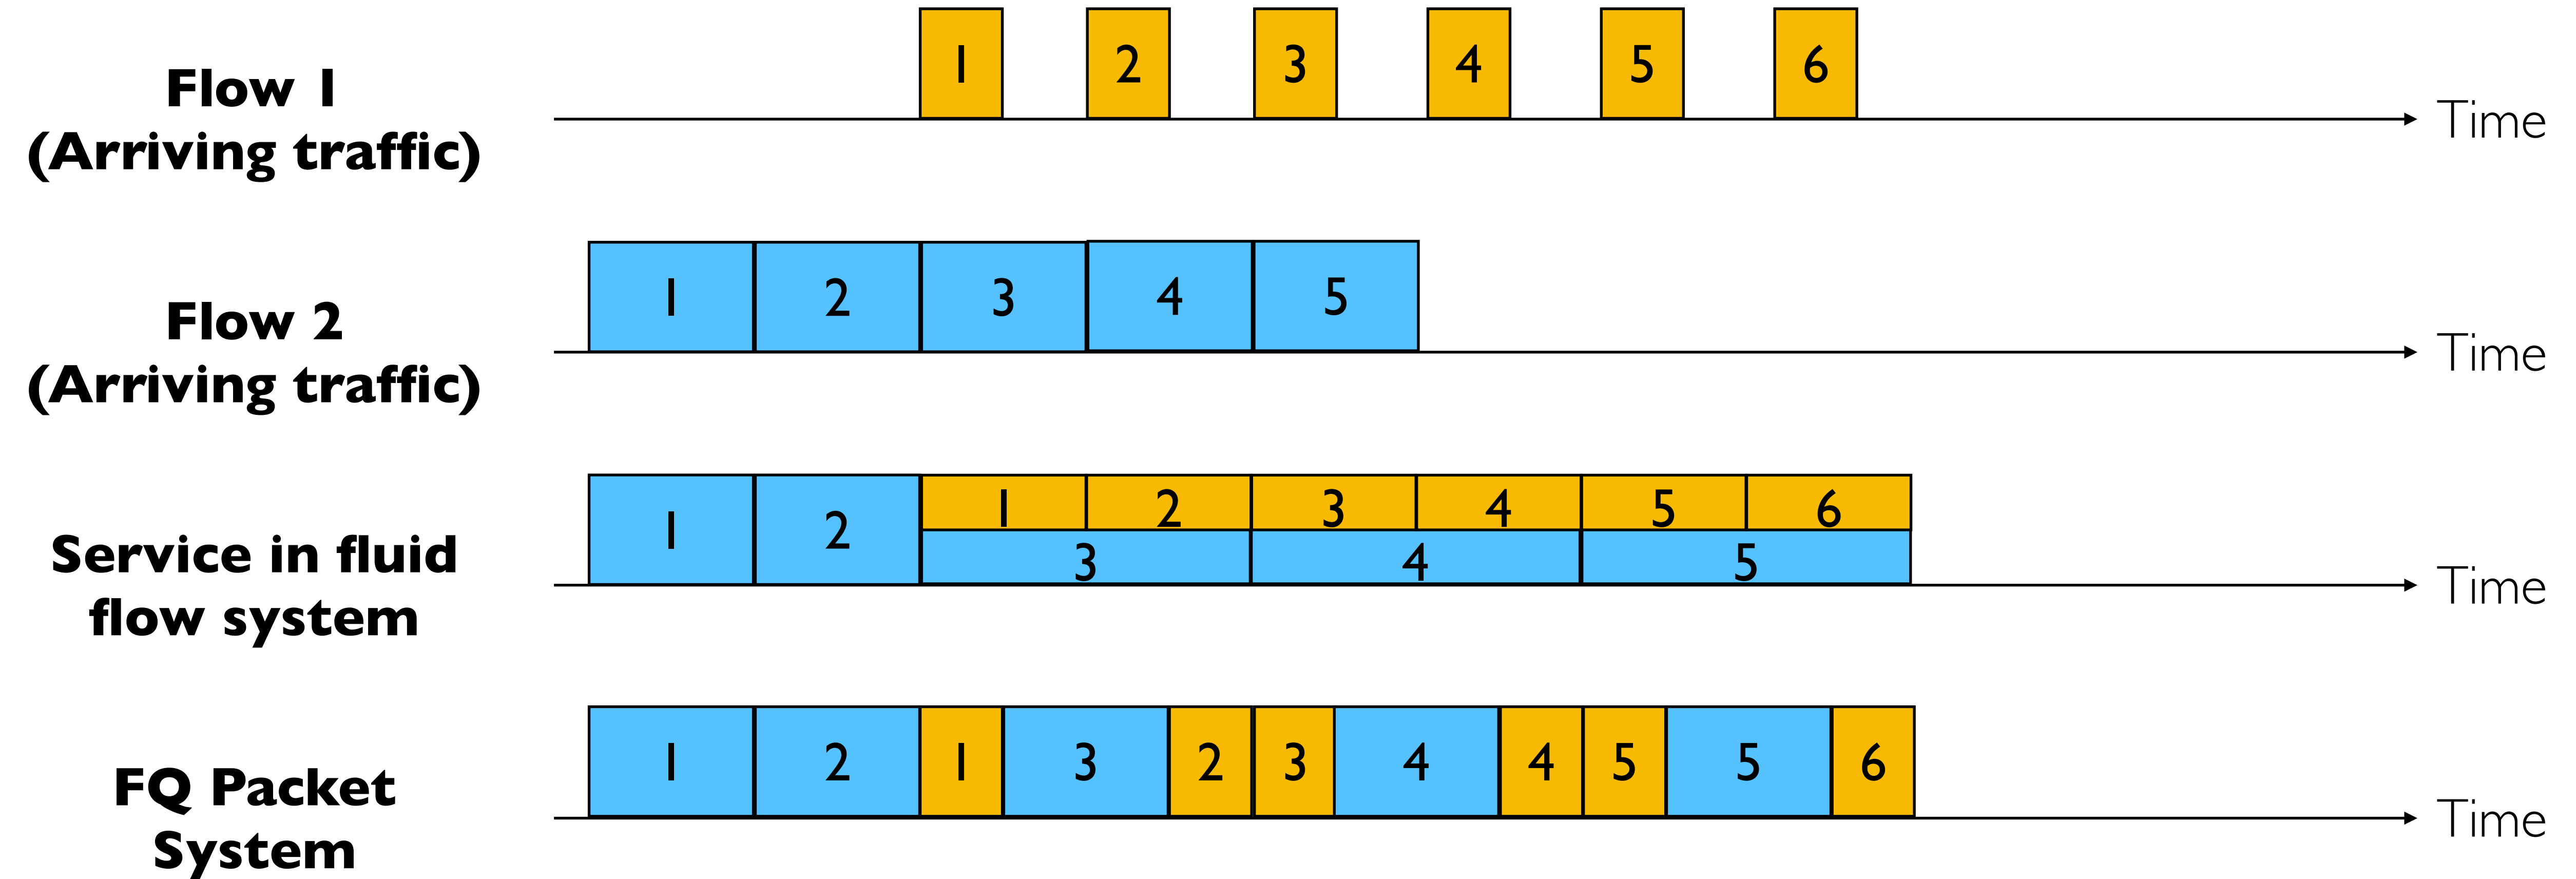
# Practical Fair Queueing (FQ)



# Practical Fair Queueing (FQ)



# Practical Fair Queueing (FQ)



# FQ vs. FIFO

# FQ vs. FIFO

- **FQ Advantages:**

- Isolation: cheating flows don't benefit
- Bandwidth share does not depend on RTT
- Flows can pick any rate adjustment scheme they want



# FQ vs. FIFO

- **FQ Advantages:**

- Isolation: cheating flows don't benefit
- Bandwidth share does not depend on RTT
- Flows can pick any rate adjustment scheme they want

- **Disadvantages**

- More complex than FIFO: per flow queue/state, additional per-packet book-keeping

# Router-assisted Congestion Control

- Three tasks for CC:
  - Isolation/fairness
  - Adjustment
  - Detecting congestion

Why not just let routers tell end-systems what rate they should use?

# Why not just let routers tell end-systems what rate they should use?

- Packets carry “rate field”

# Why not just let routers tell end-systems what rate they should use?

- Packets carry “rate field”
- Routers insert “fair share”  $f$  in the packet header

# Why not just let routers tell end-systems what rate they should use?

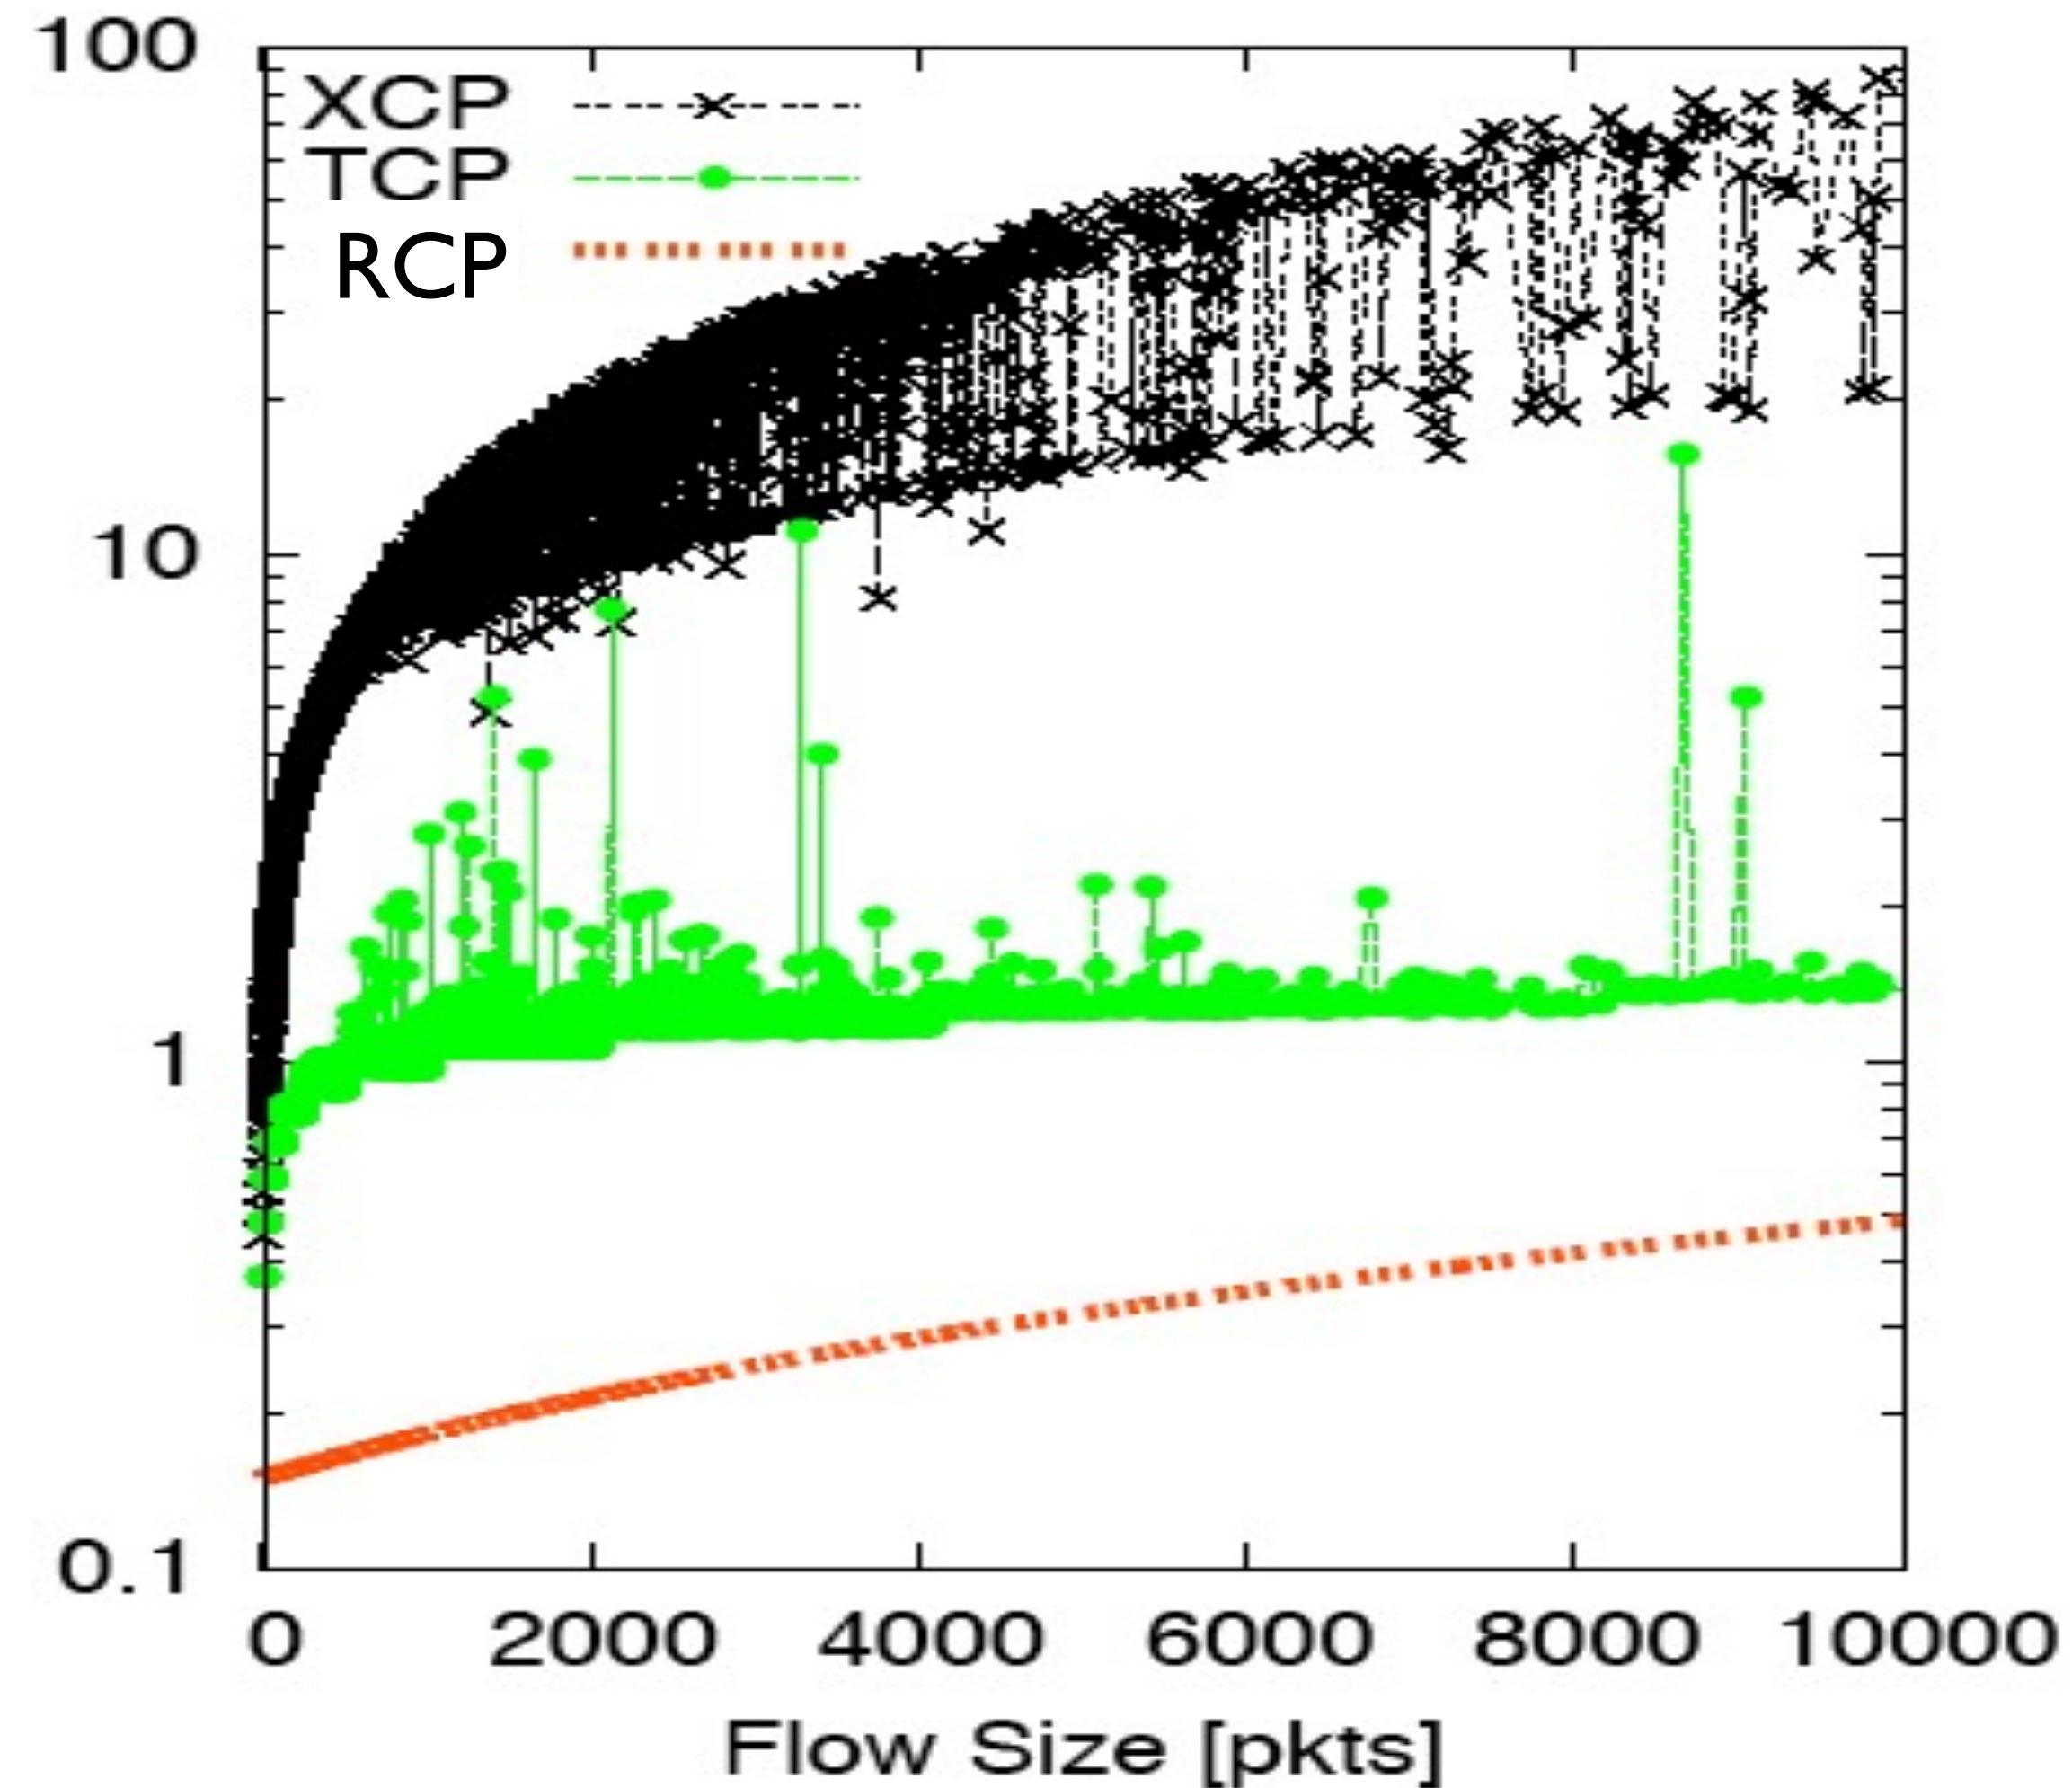
- Packets carry “rate field”
- Routers insert “fair share”  $f$  in the packet header
- End-systems set sending rate (or window size) to  $f$ 
  - Hopefully (still need some policing of end-systems!)

# Why not just let routers tell end-systems what rate they should use?

- Packets carry “rate field”
- Routers insert “fair share”  $f$  in the packet header
- End-systems set sending rate (or window size) to  $f$ 
  - Hopefully (still need some policing of end-systems!)
- *Basic idea behind the “Rate Control Protocol” (RCP) from Dukkiperi et. al. ‘07*



# Flow Completion Time: TCP vs. RCP (Ignore XCP)

## Flow Duration (seconds) vs. Flow Size





# Why the Improvement?



# Router-assisted Congestion Control

- Three tasks for CC:
  - Isolation/fairness
  - Adjustment
  - Detecting congestion

# Explicit Congestion Notification (ECN)

# Explicit Congestion Notification (ECN)

- Single bit in packet header; set by congested routers
  - If data packet has bit set, then ACK has ECN bit set

# Explicit Congestion Notification (ECN)

- Single bit in packet header; set by congested routers
  - If data packet has bit set, then ACK has ECN bit set
- Congestion semantics can be exactly like that of drop
  - i.e., end-system reacts as though it saw a drop

# Explicit Congestion Notification (ECN)

- Single bit in packet header; set by congested routers
  - If data packet has bit set, then ACK has ECN bit set
- Congestion semantics can be exactly like that of drop
  - i.e., end-system reacts as though it saw a drop
- Advantages:
  - Don't confuse corruption with congestion; recovery with rate adjustment
  - Can serve as an early indicator of congestion to avoid delays
  - Easy (easier) to incrementally deploy
    - Today: defined in RFC 3168 using **ToS/DSCP** bits in the IP header

**One final proposal: Charge people for congestion!**

# One final proposal: Charge people for congestion!

- Use ECN as congestion markers



# One final proposal: Charge people for congestion!

- Use ECN as congestion markers
- Whenever I get an ECN bit set, I have to pay \$\$

# One final proposal: Charge people for congestion!

- Use ECN as congestion markers
- Whenever I get an ECN bit set, I have to pay \$\$
- Now there's no debate over what a flow is, or what fair is...

# One final proposal: Charge people for congestion!

- Use ECN as congestion markers
- Whenever I get an ECN bit set, I have to pay \$\$
- Now there's no debate over what a flow is, or what fair is...
- Idea started by Frank Kelly at Cambridge
  - “Optimal” solution, backed by much math
  - Great idea: simple, elegant, effective
  - Unclear that it will impact practice

# Recap: TCP

# Recap: TCP

- TCP
  - Somewhat hacky
  - But practical/deployable
  - Good enough for Internet traffic
  - Needs of data centers might change status quo (future lecture)

# Taking Stock: Transport Layer

# Taking Stock: Transport Layer

- **Communication between application processes**
  - Multiplex and demultiplex from/to application processes using ***ports***

# Taking Stock: Transport Layer

- Communication between application processes
- **Provide common end-to-end services for application layer [optional]**
  - Reliable, in-order data delivery
  - Well-placed data delivery
    - Too fast may overwhelm the receiver/network
    - Too slow is not efficient



# Taking Stock: Transport Layer

- Communication between application processes
- Provide common end-to-end services for application layer [optional]
- **TCP and UDP are the common transport protocols**
  - Also SCTP, MTCP, SST, RDP, DCCP, ...

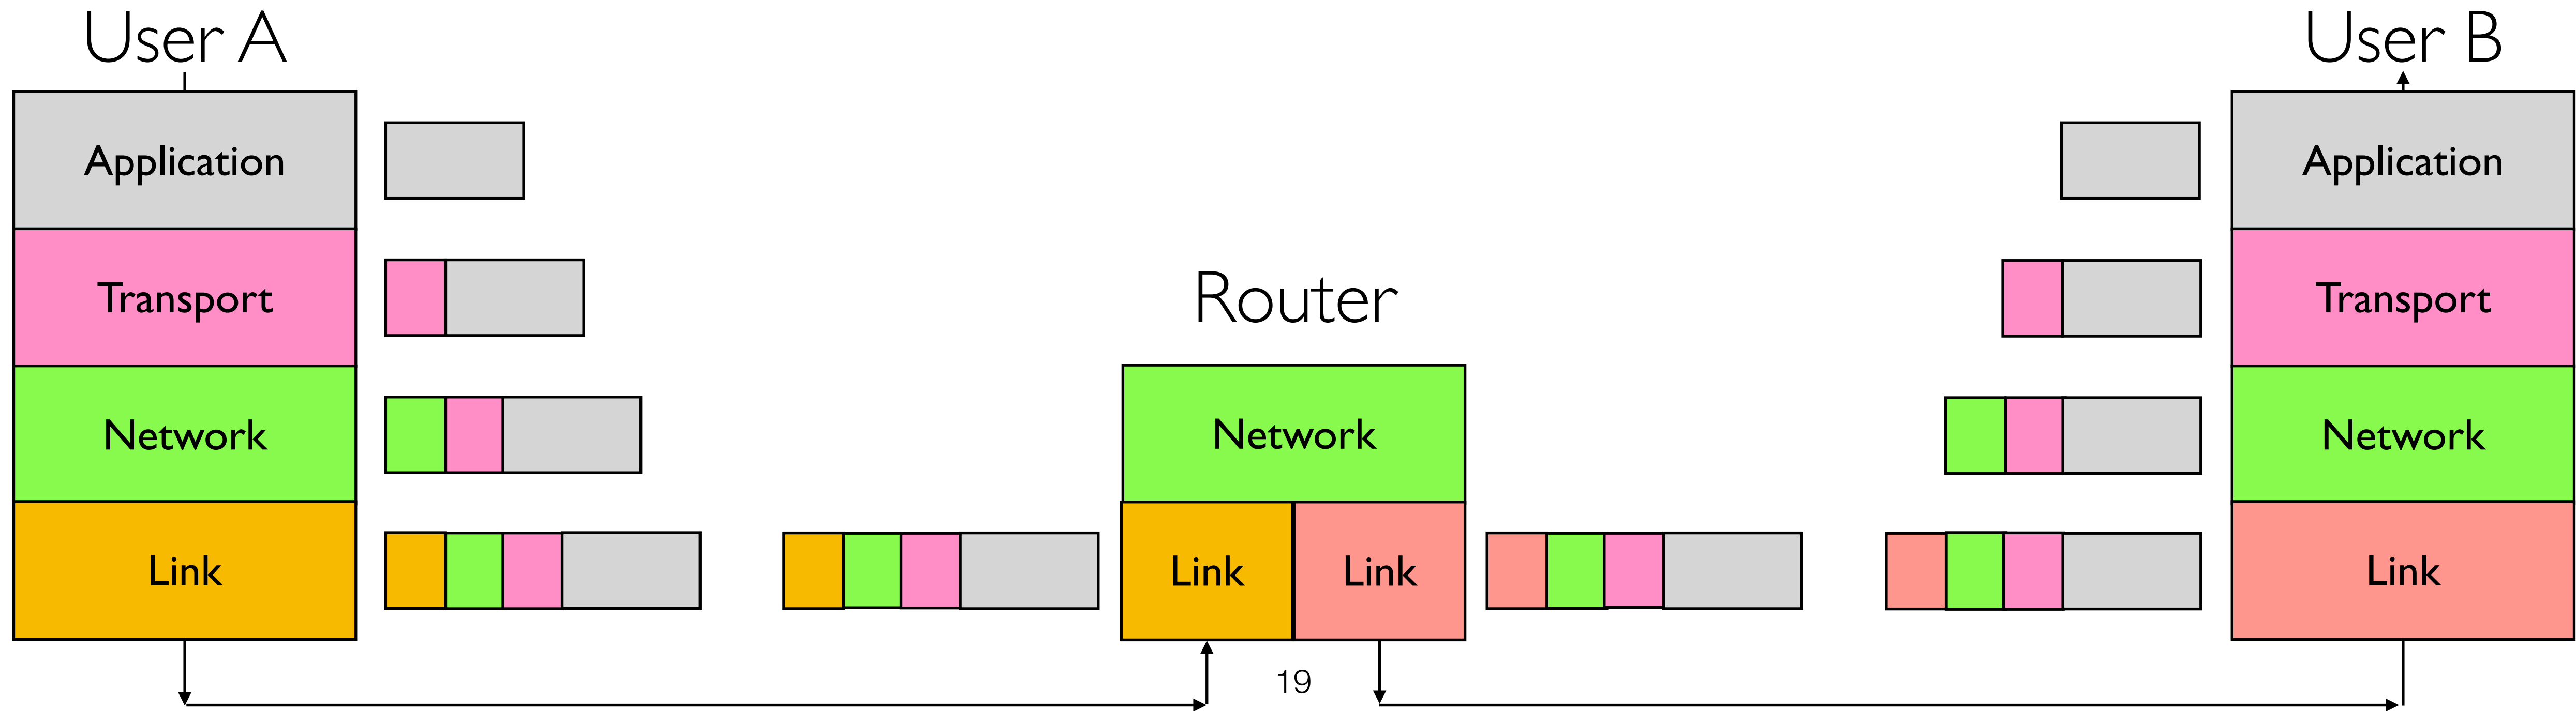
# Taking Stock: Transport Layer

- Communication between application processes
- Provide common end-to-end services for application layer [optional]
- TCP and UDP are the common transport protocols
- **UDP is a minimalist, no-frills transport protocol**
  - Only provides multiplex/demultiplex capabilities

# Taking Stock: Transport Layer


- Communication between application processes
- Provide common end-to-end services for application layer [optional]
- TCP and UDP are the common transport protocols
- UDP is a minimalist, no-frills transport protocol
- **TCP is the whole-hog protocol**
  - Offers applications a reliable, in-order, byte stream abstraction
  - With congestion control
  - But no performance guarantees (delay, bw, etc.)

# Taking Stock



# Taking Stock

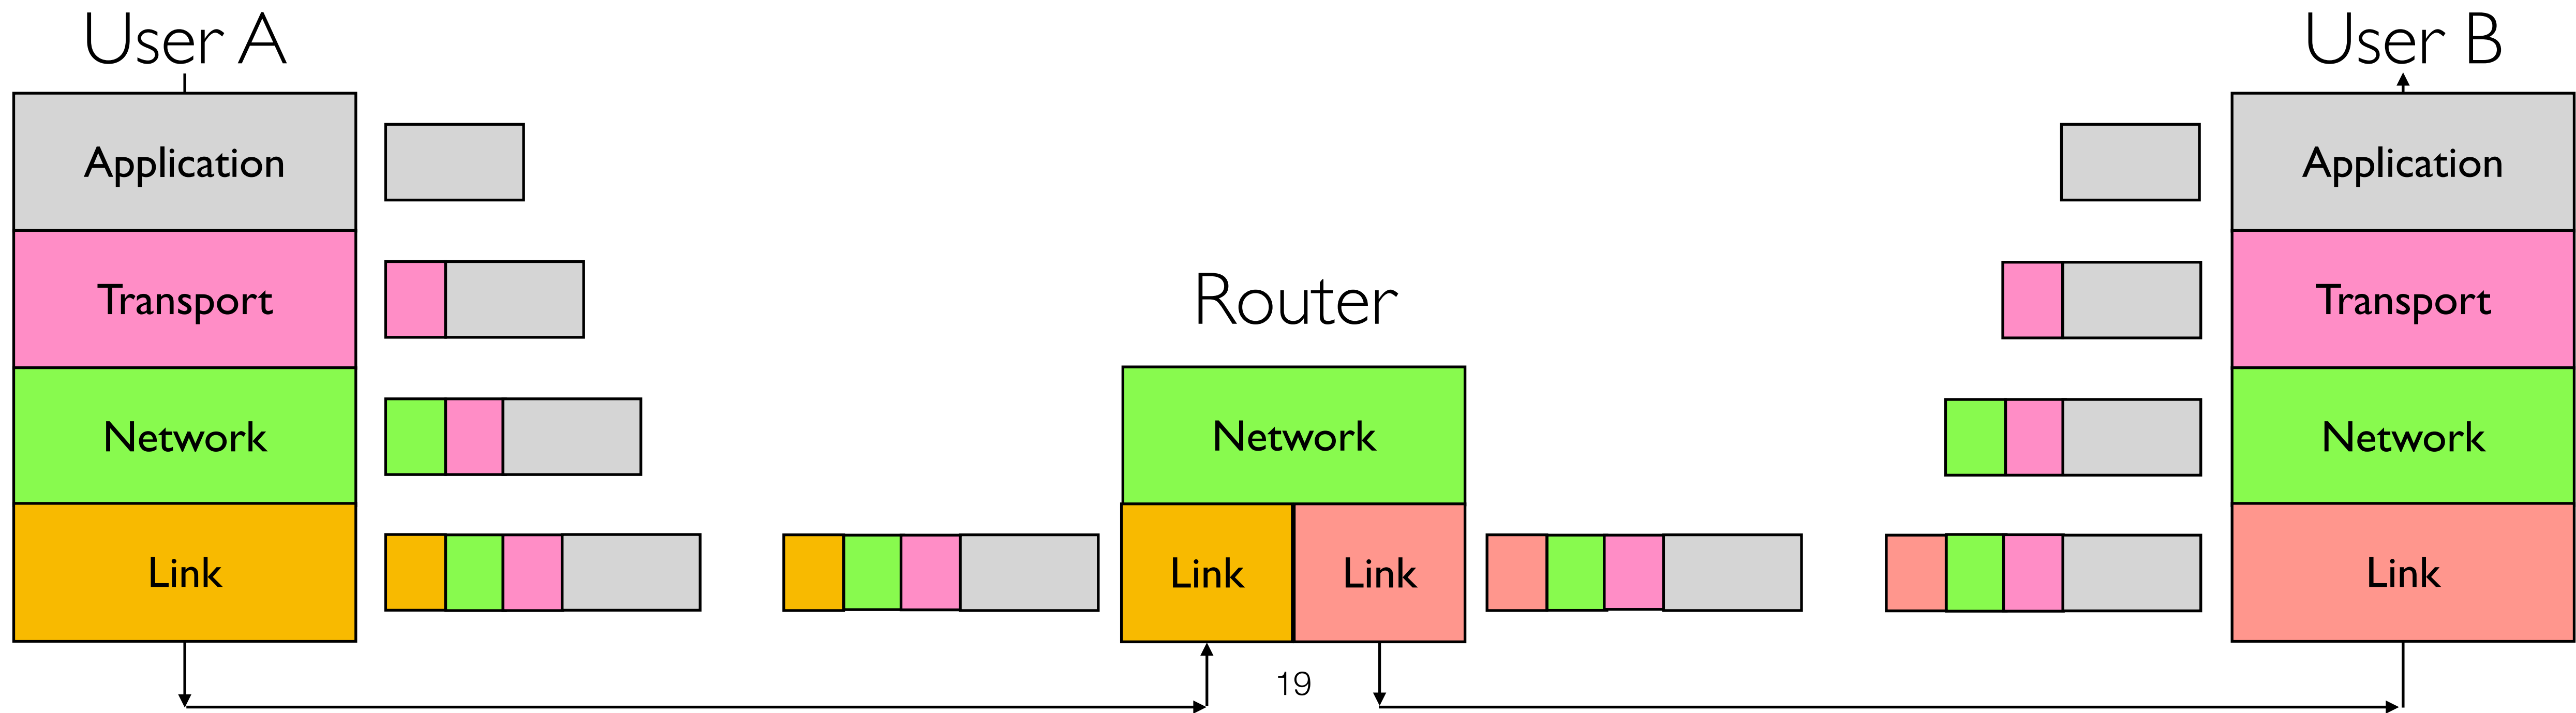
**Course so far:**



# Taking Stock

## Course so far:


- **Concepts,** *Links, delays, switches*



# Taking Stock

## Course so far:

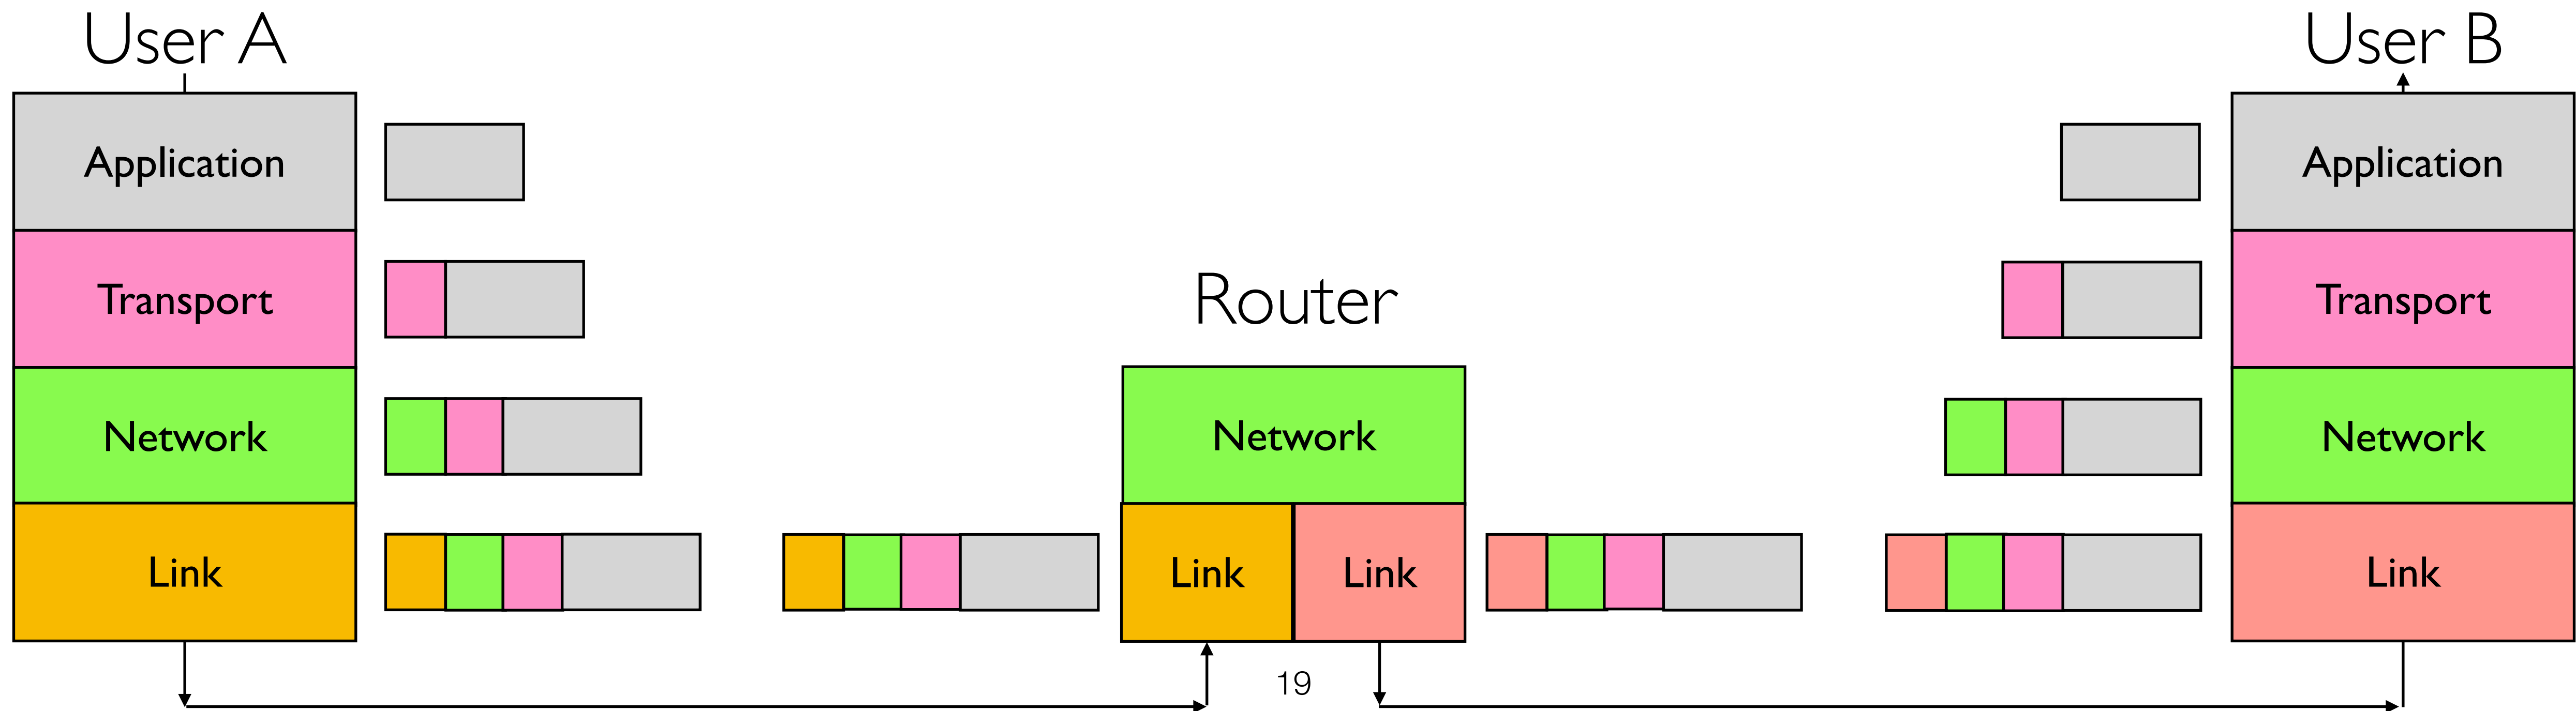
- **Concepts,** *Links, delays, switches*
- **Overall Architecture,** *Layers, protocols principles*



# Taking Stock

## Course so far:

- **Concepts,** *Links, delays, switches*
- **Overall Architecture,** *Layers, protocols principles*
- **Network Layer,** *Best-effort global delivery of packets*

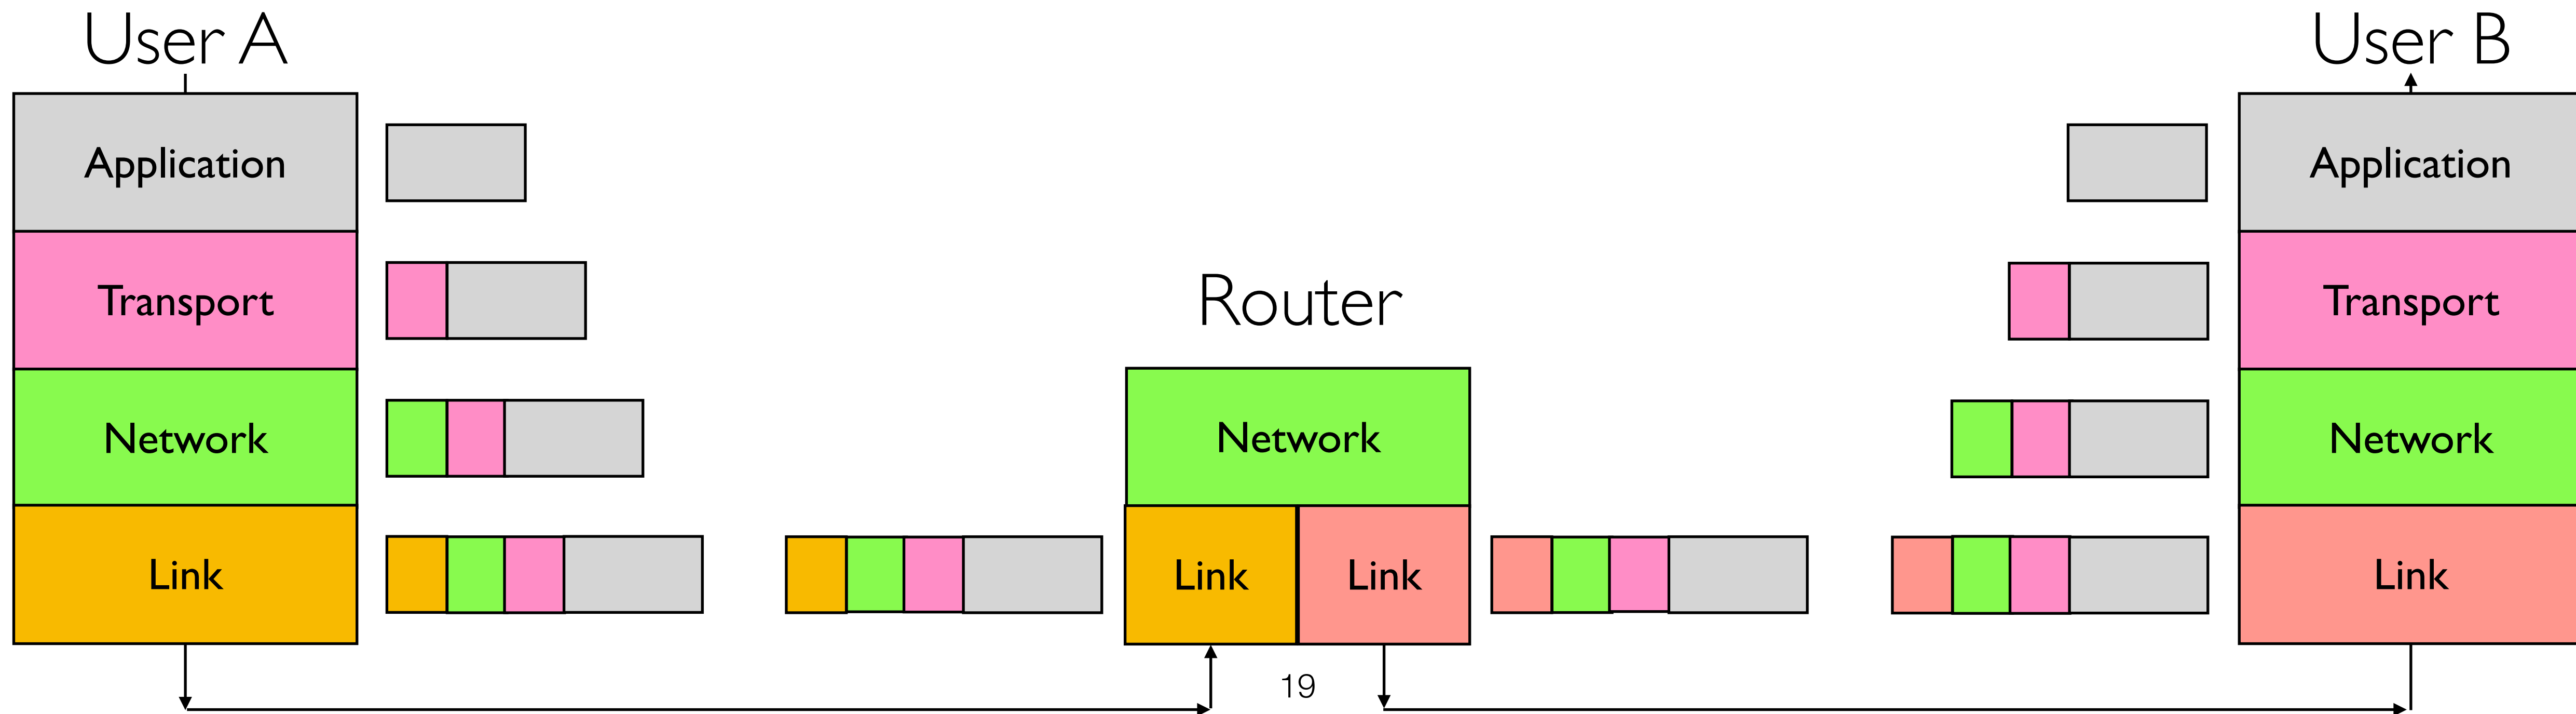




# Taking Stock

## Course so far:

- **Concepts**, *Links, delays, switches*
- **Overall Architecture**, *Layers, protocols principles*
- **Network Layer**, *Best-effort global delivery of packets*
- **Transport Layer**, *Reliable (or unreliable) delivery of data*

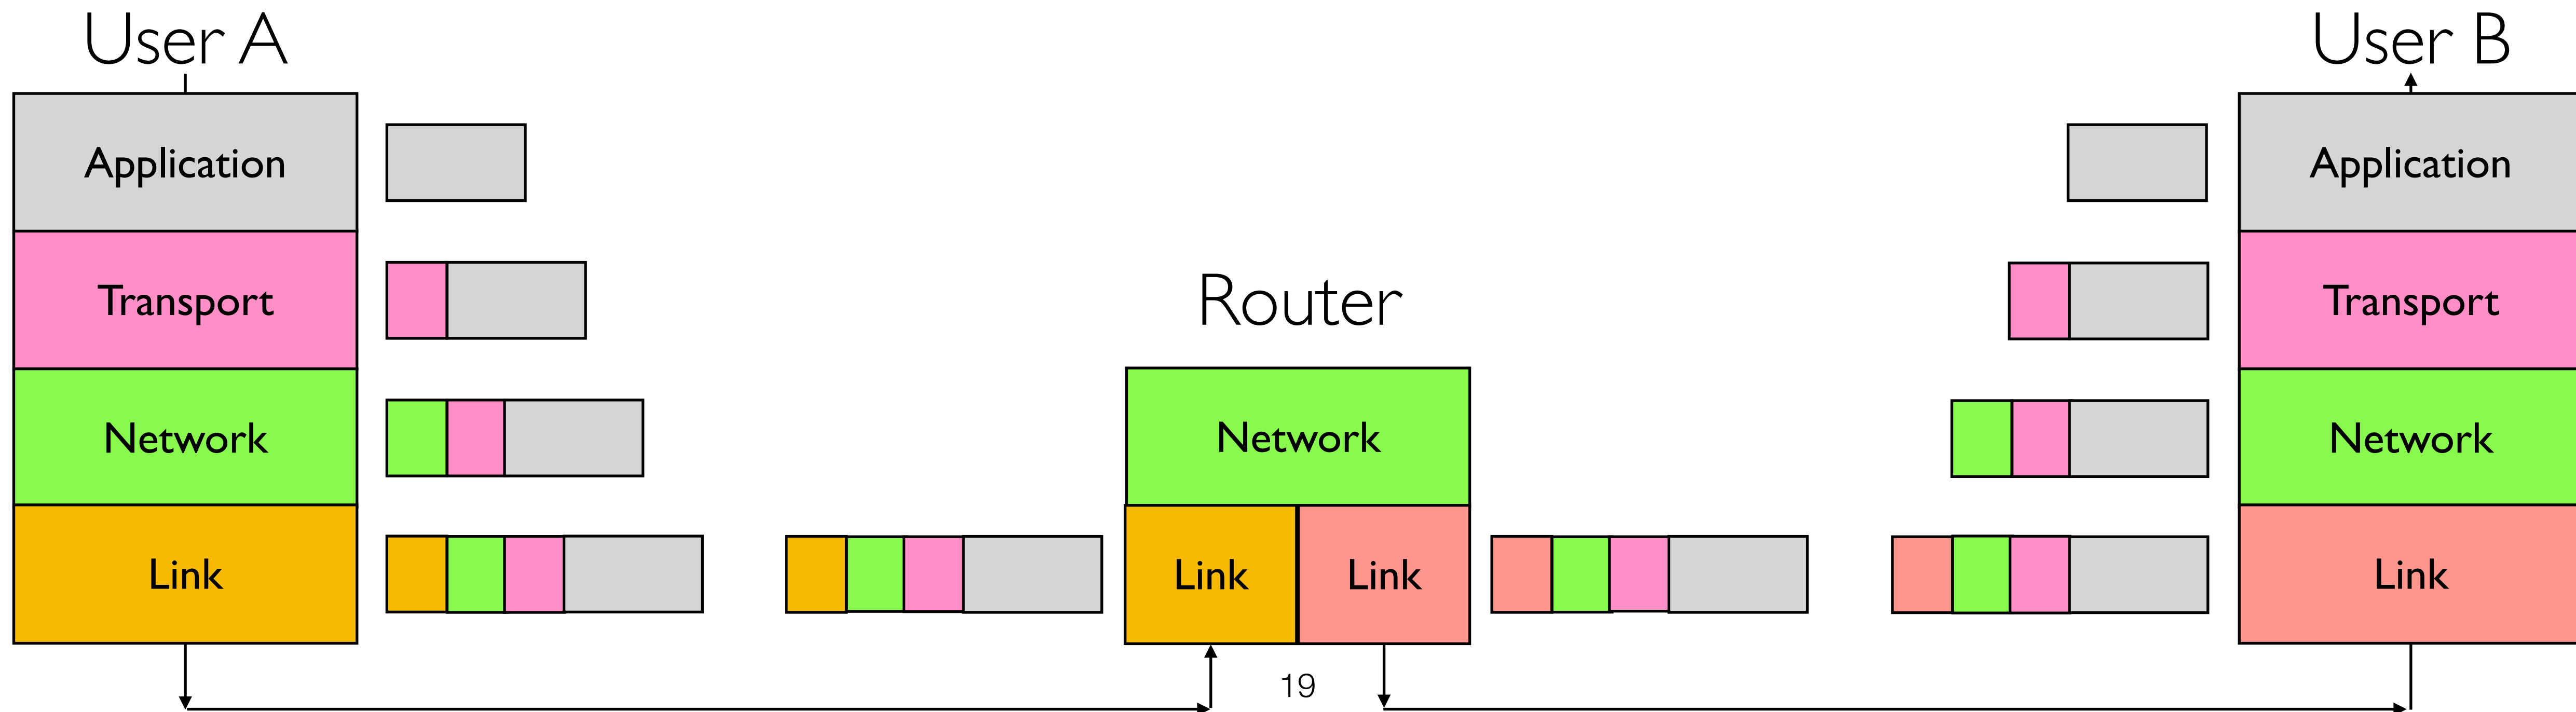


# Taking Stock

## Course so far:

- **Concepts,** *Links, delays, switches*
- **Overall Architecture,** *Layers, protocols principles*
- **Network Layer,** *Best-effort global delivery of packets*
- **Transport Layer,** *Reliable (or unreliable) delivery of data*

## What's left?



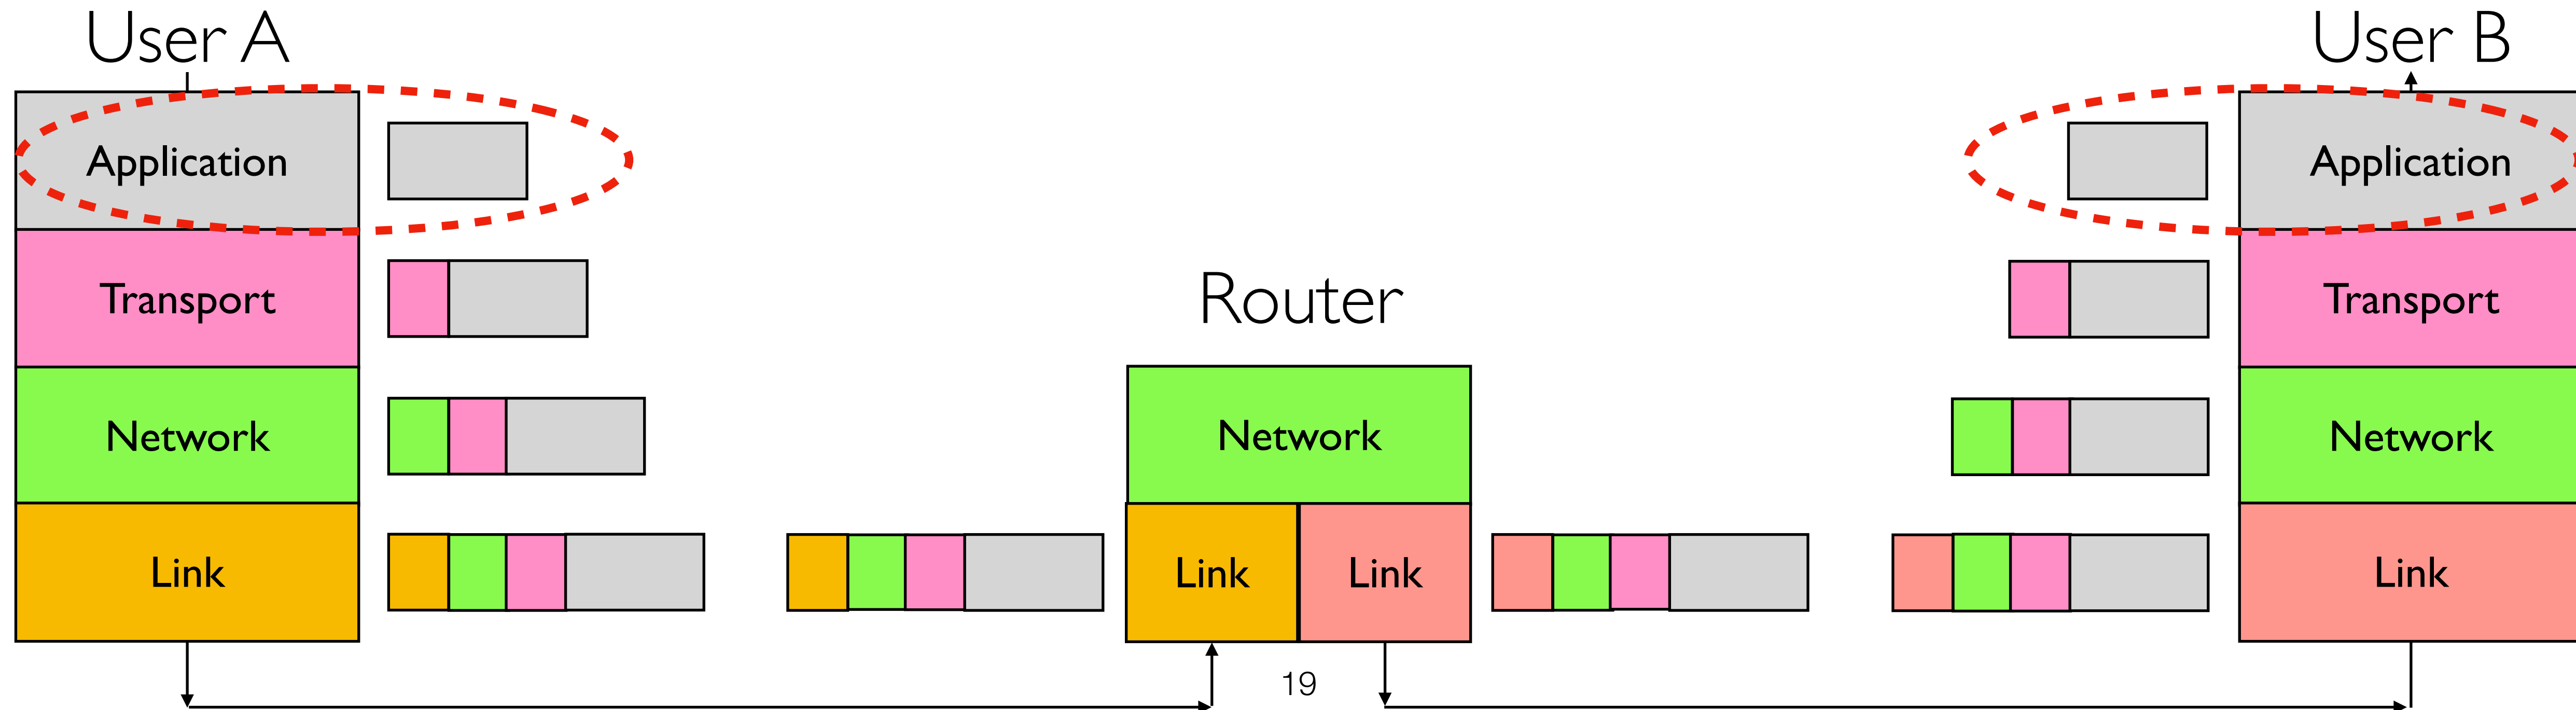
# Taking Stock

## Course so far:

- **Concepts**, *Links, delays, switches*
- **Overall Architecture**, *Layers, protocols principles*
- **Network Layer**, *Best-effort global delivery of packets*
- **Transport Layer**, *Reliable (or unreliable) delivery of data*

## What's left?

- **Application Layer**, *DNS, HTTP (today)*



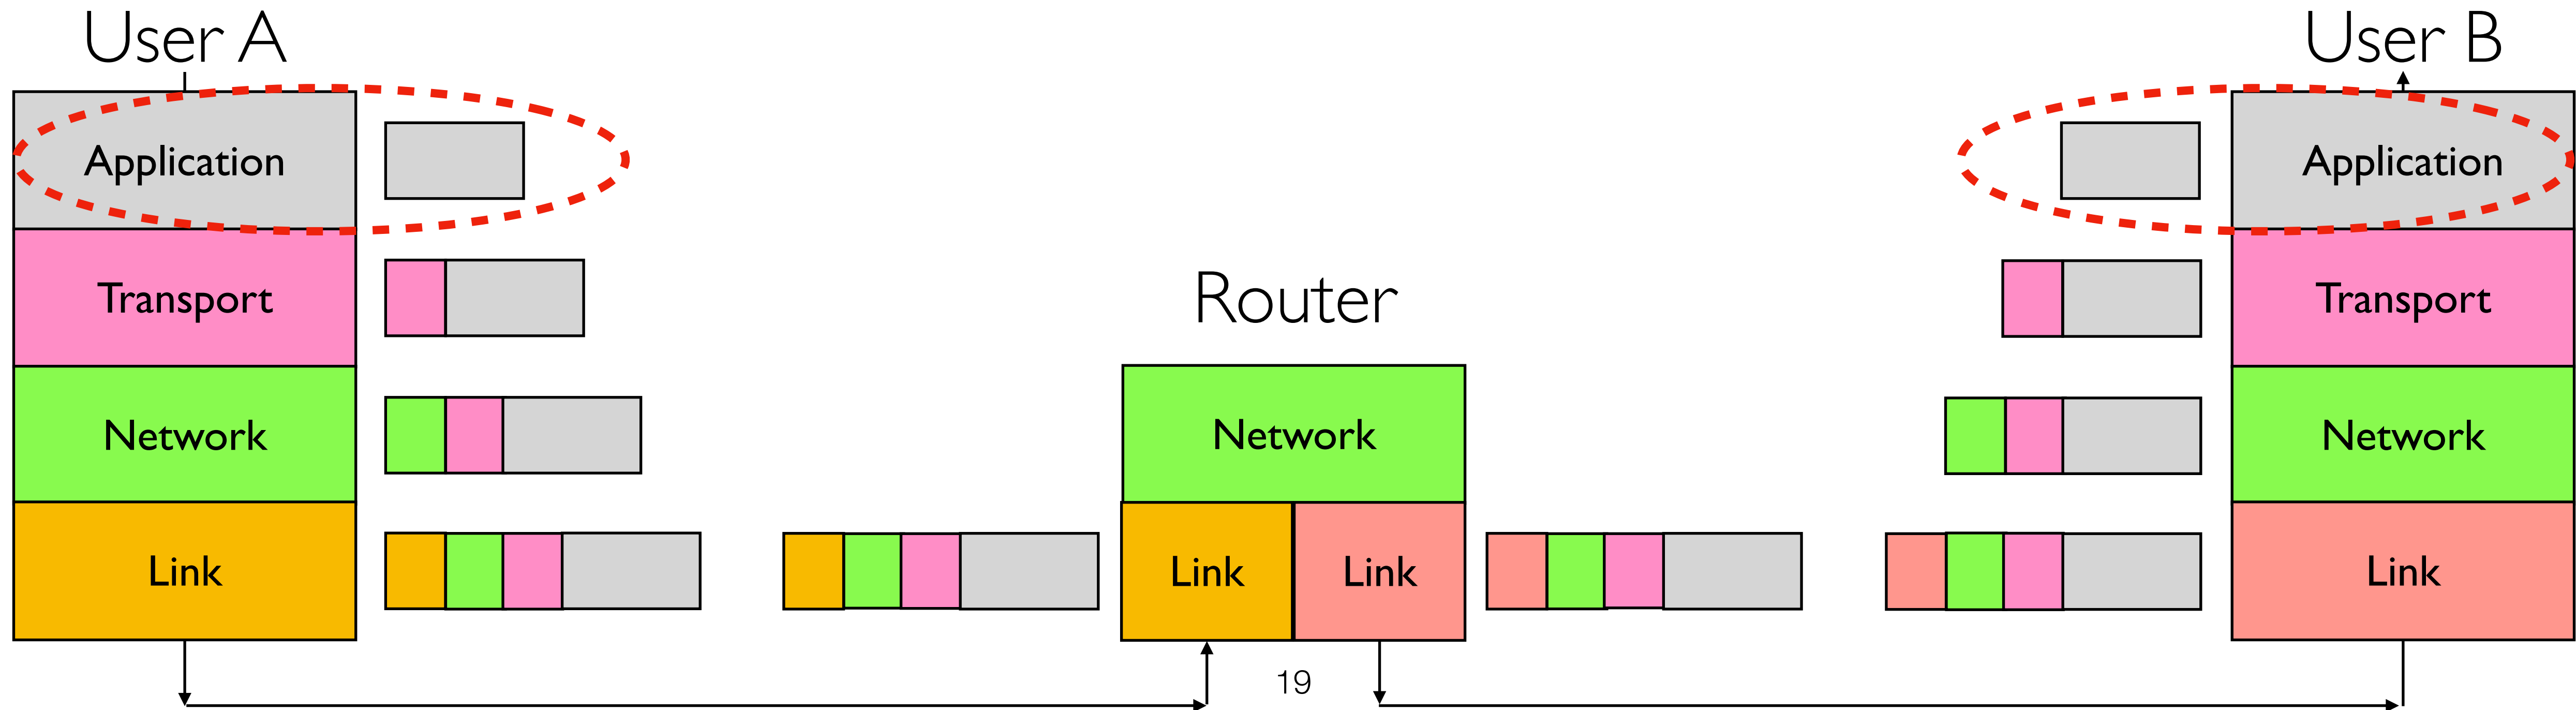
# Taking Stock

## Course so far:

- **Concepts**, *Links, delays, switches*
- **Overall Architecture**, *Layers, protocols principles*
- **Network Layer**, *Best-effort global delivery of packets*
- **Transport Layer**, *Reliable (or unreliable) delivery of data*

## What's left?

- **Application Layer**, *DNS, HTTP (today)*
- **Lower Layers**, *Ethernet, Wireless*



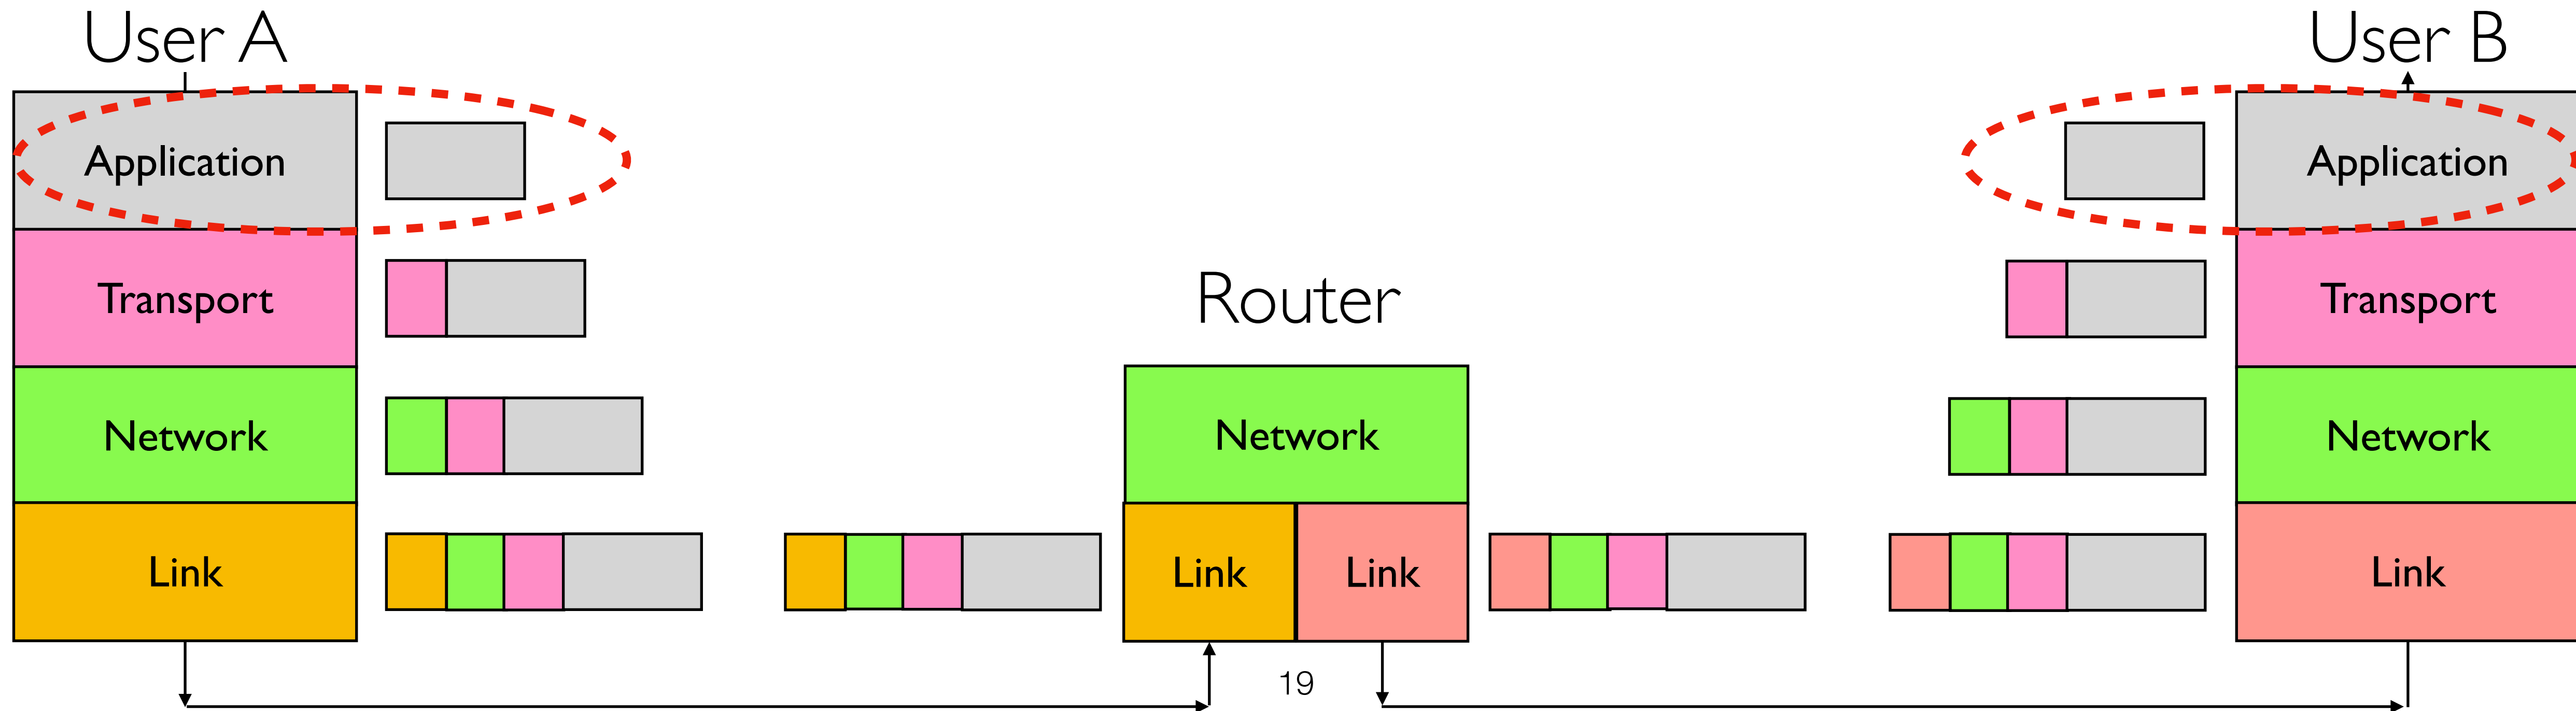
# Taking Stock

## Course so far:

- **Concepts**, *Links, delays, switches*
- **Overall Architecture**, *Layers, protocols principles*
- **Network Layer**, *Best-effort global delivery of packets*
- **Transport Layer**, *Reliable (or unreliable) delivery of data*

## What's left?

- **Application Layer**, *DNS, HTTP (today)*
- **Lower Layers**, *Ethernet, Wireless*
- **Advanced Topics**, *Datacenters, SDN*



# On to the Application Layer!

# Application Layer

# Application Layer

- Domain Name System (DNS)
  - What's behind (e.g.) zoo.cs.yale.edu?



# Application Layer

- **Domain Name System (DNS)**
  - What's behind (e.g.) zoo.cs.yale.edu?
- **HTTP and the Web**
  - What happens when you click on a link?

# Application Layer

- **Domain Name System (DNS)**
  - What's behind (e.g.) zoo.cs.yale.edu?
- HTTP and the Web
  - What happens when you click on a link?

# Host Names & Addresses

# Host Names & Addresses

- Host address: e.g., *169.229.131.109*
  - A number used by protocols
  - Conforms to network structure (*the “where”*)

# Host Names & Addresses

- **Host address: e.g., 169.229.131.109**
  - A number used by protocols
  - Conforms to network structure (*the “where”*)
- **Host names: e.g., zoo.cs.yale.edu**
  - Mnemonic name usable by humans
  - Conforms to organizational structure (*the “who”*)

# Host Names & Addresses

- **Host address: e.g., 169.229.131.109**
  - A number used by protocols
  - Conforms to network structure (*the “where”*)
- **Host names: e.g., zoo.cs.yale.edu**
  - Mnemonic name usable by humans
  - Conforms to organizational structure (*the “who”*)
- **The Domain Name System (DNS) is how we map from one to the other**
  - A *directory* service for hosts on the Internet

# Why Bother?

# Why Bother?

- Convenience
  - Easier to remember www.google.com than 172.217.8.174



# Why Bother?

- **Convenience**
  - Easier to remember www.google.com than 172.217.8.174
- **Provides a level of indirection!**
  - Decoupled names from addresses
  - Many uses beyond just naming a specific host

# DNS: Early Days

# DNS: Early Days

- **Mappings stored in a `hosts.txt` file (in `/etc/hosts`)**
  - Maintained by the Stanford Research Institute (SRI)
  - New versions periodically copied from SRI (via FTP)

# DNS: Early Days

- **Mappings stored in a hosts.txt file (in /etc/hosts)**
  - Maintained by the Stanford Research Institute (SRI)
  - New versions periodically copied from SRI (via FTP)
- **As the Internet grew, this system broke down (obviously)**
  - SRI couldn't handle the load
  - Conflicts in selecting names
  - Hosts had inaccurate copies of hosts.txt

# DNS: Early Days

- **Mappings stored in a hosts.txt file (in /etc/hosts)**
  - Maintained by the Stanford Research Institute (SRI)
  - New versions periodically copied from SRI (via FTP)
- **As the Internet grew, this system broke down (obviously)**
  - SRI couldn't handle the load
  - Conflicts in selecting names
  - Hosts had inaccurate copies of hosts.txt
- **The Domain Name Service (DNS) was invented to fix this**

# Goals?

# Goals?

- Scalable
  - Many *names*
  - Many *updates*
  - Many *users* creating names
  - Many *users* looking up names

# Goals?

- **Scalable**
  - Many *names*
  - Many *updates*
  - Many *users* creating names
  - Many *users* looking up names
- **Highly available**



# Goals?

- **Scalable**
  - Many *names*
  - Many *updates*
  - Many *users* creating names
  - Many *users* looking up names
- **Highly available**
- **Correct**
  - No naming conflicts (uniqueness)
  - Consistency → observe the latest update

# Goals?

- **Scalable**
  - Many *names*
  - Many *updates*
  - Many *users* creating names
  - Many *users* looking up names
- **Highly available**
- **Correct**
  - No naming conflicts (uniqueness)
  - Consistency → observe the latest update
- **Lookups are fast**

# How Do We Scale?

# How Do We Scale?

- Partition the namespace

# How Do We Scale?

- Partition the namespace
- Distribute the administration of each partition
  - Autonomy to update my partition's machines' names
  - Don't have to track other partition's updates

# How Do We Scale?

- Partition the namespace
- **Distribute the administration of each partition**
  - Autonomy to update my partition's machines' names
  - Don't have to track other partition's updates
- **Distribute name resolution for each partition**
  - If I need to figure out the address for a name, I ask the partition that it belongs to

# How Do We Scale?

- Partition the namespace
- Distribute the administration of each partition
  - Autonomy to update my partition's machines' names
  - Don't have to track other partition's updates
- Distribute name resolution for each partition
  - If I need to figure out the address for a name, I ask the partition that it belongs to
- ***How should we partition things?***

# Key Idea: Hierarchical Distribution



# Key Idea: Hierarchical Distribution

Three intertwined hierarchies

# Key Idea: Hierarchical Distribution

Three intertwined hierarchies

- Hierarchical naming
  - As opposed to flat namespace

# Key Idea: Hierarchical Distribution

Three intertwined hierarchies

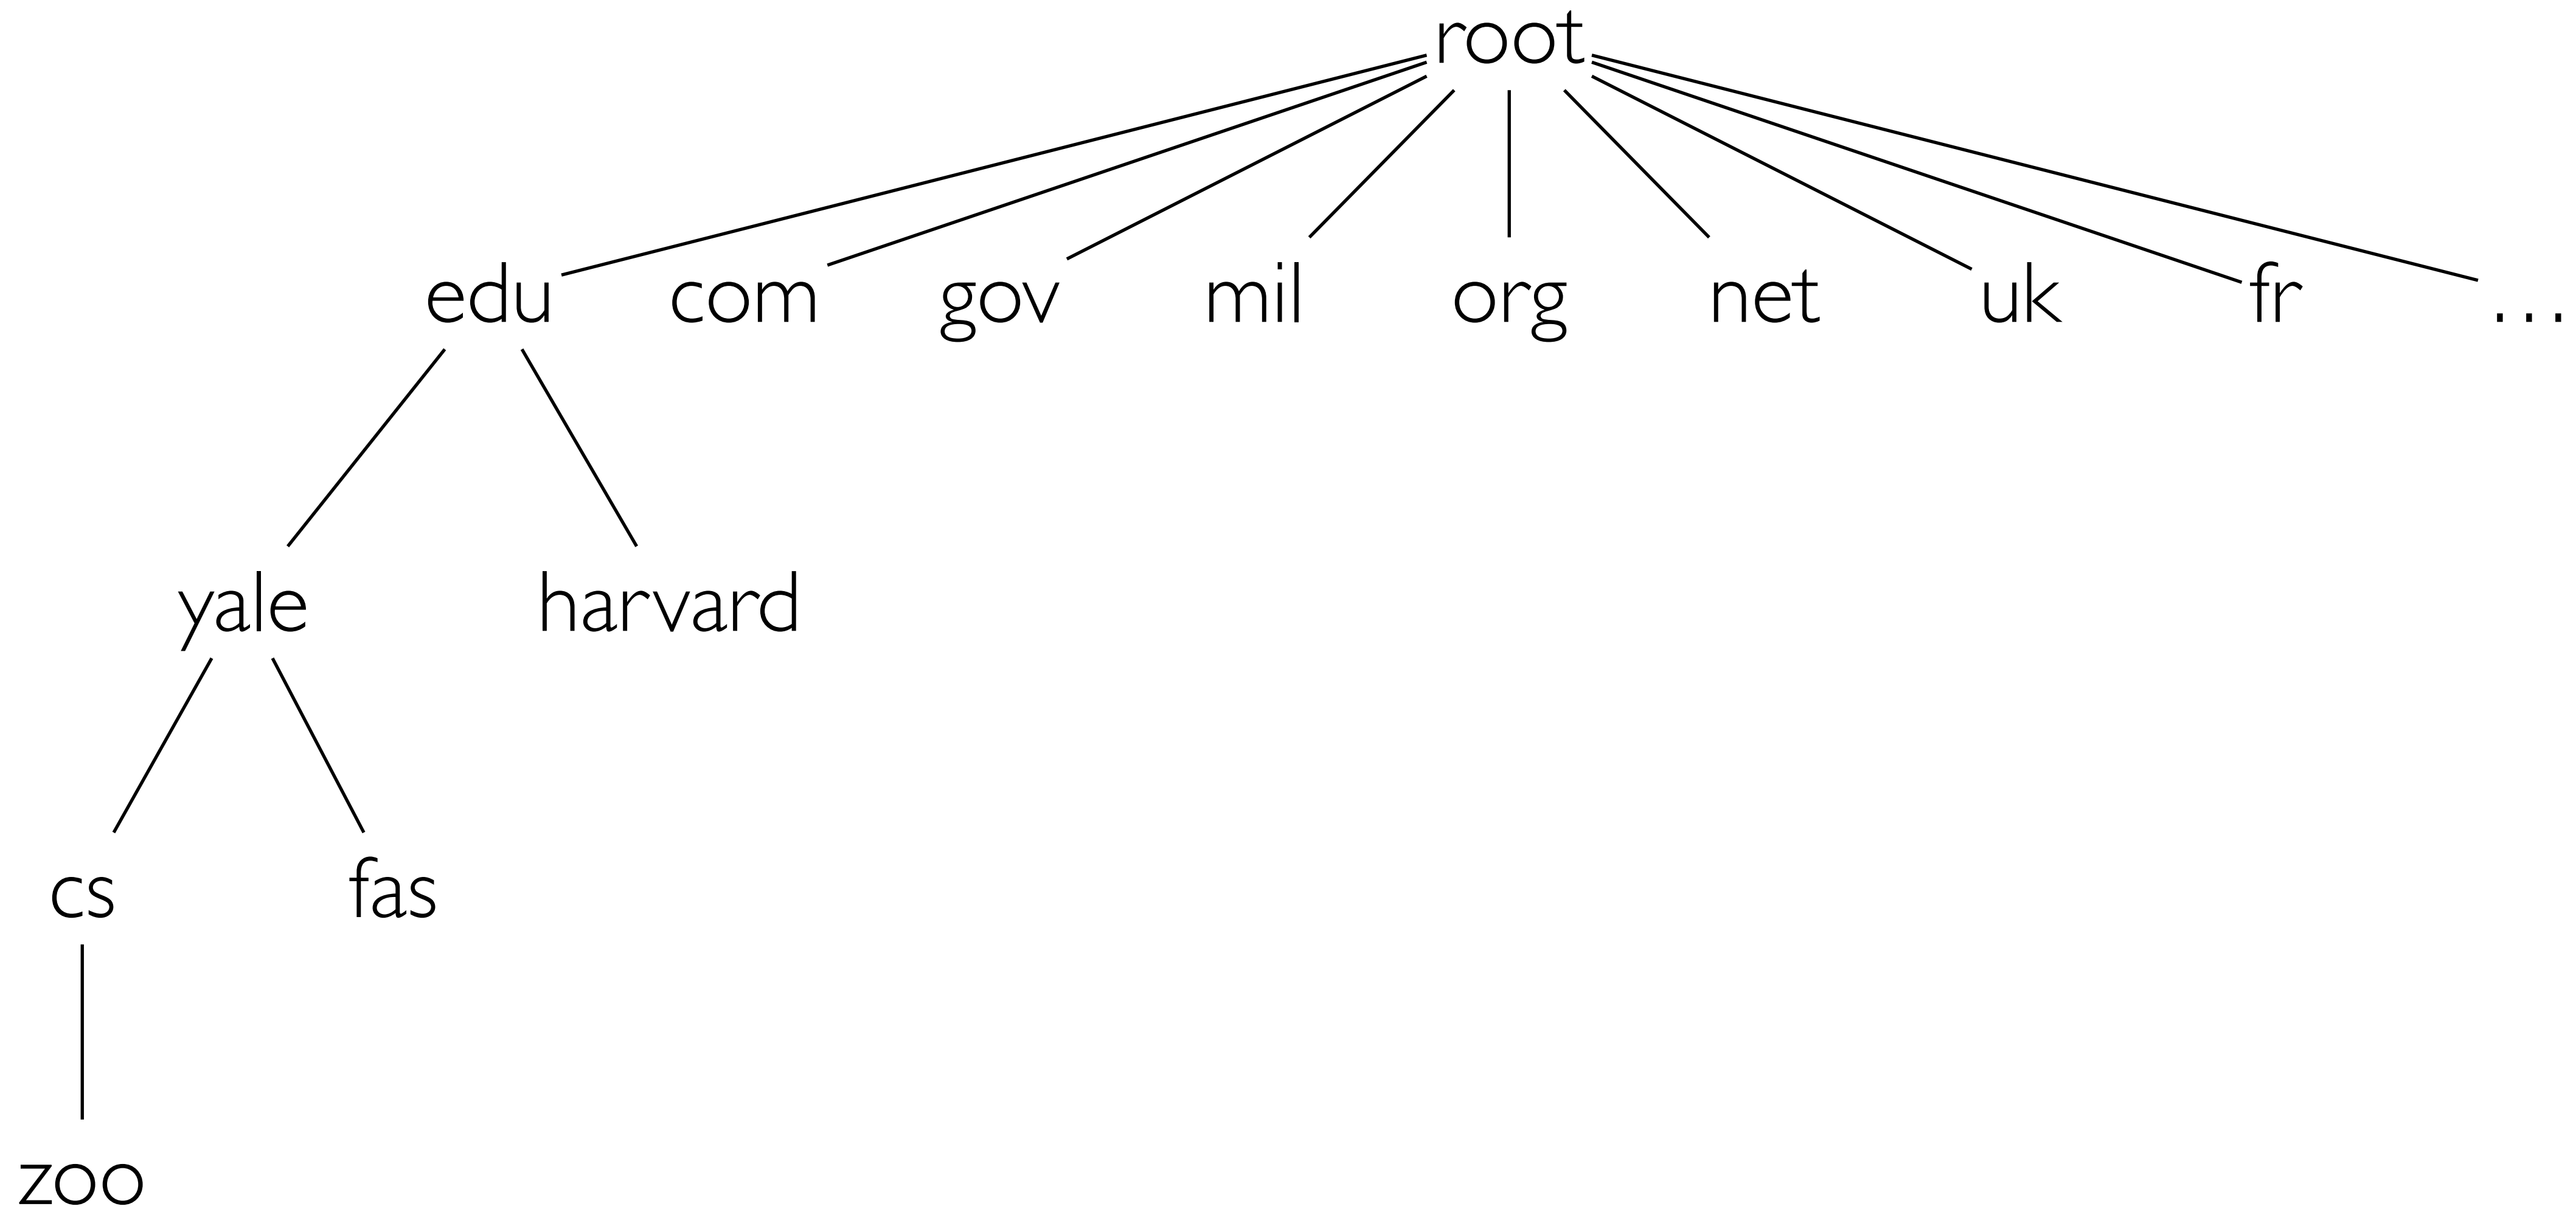
- Hierarchical naming
  - As opposed to flat namespace
- Hierarchical administration
  - As opposed to centralized administration

# Key Idea: Hierarchical Distribution

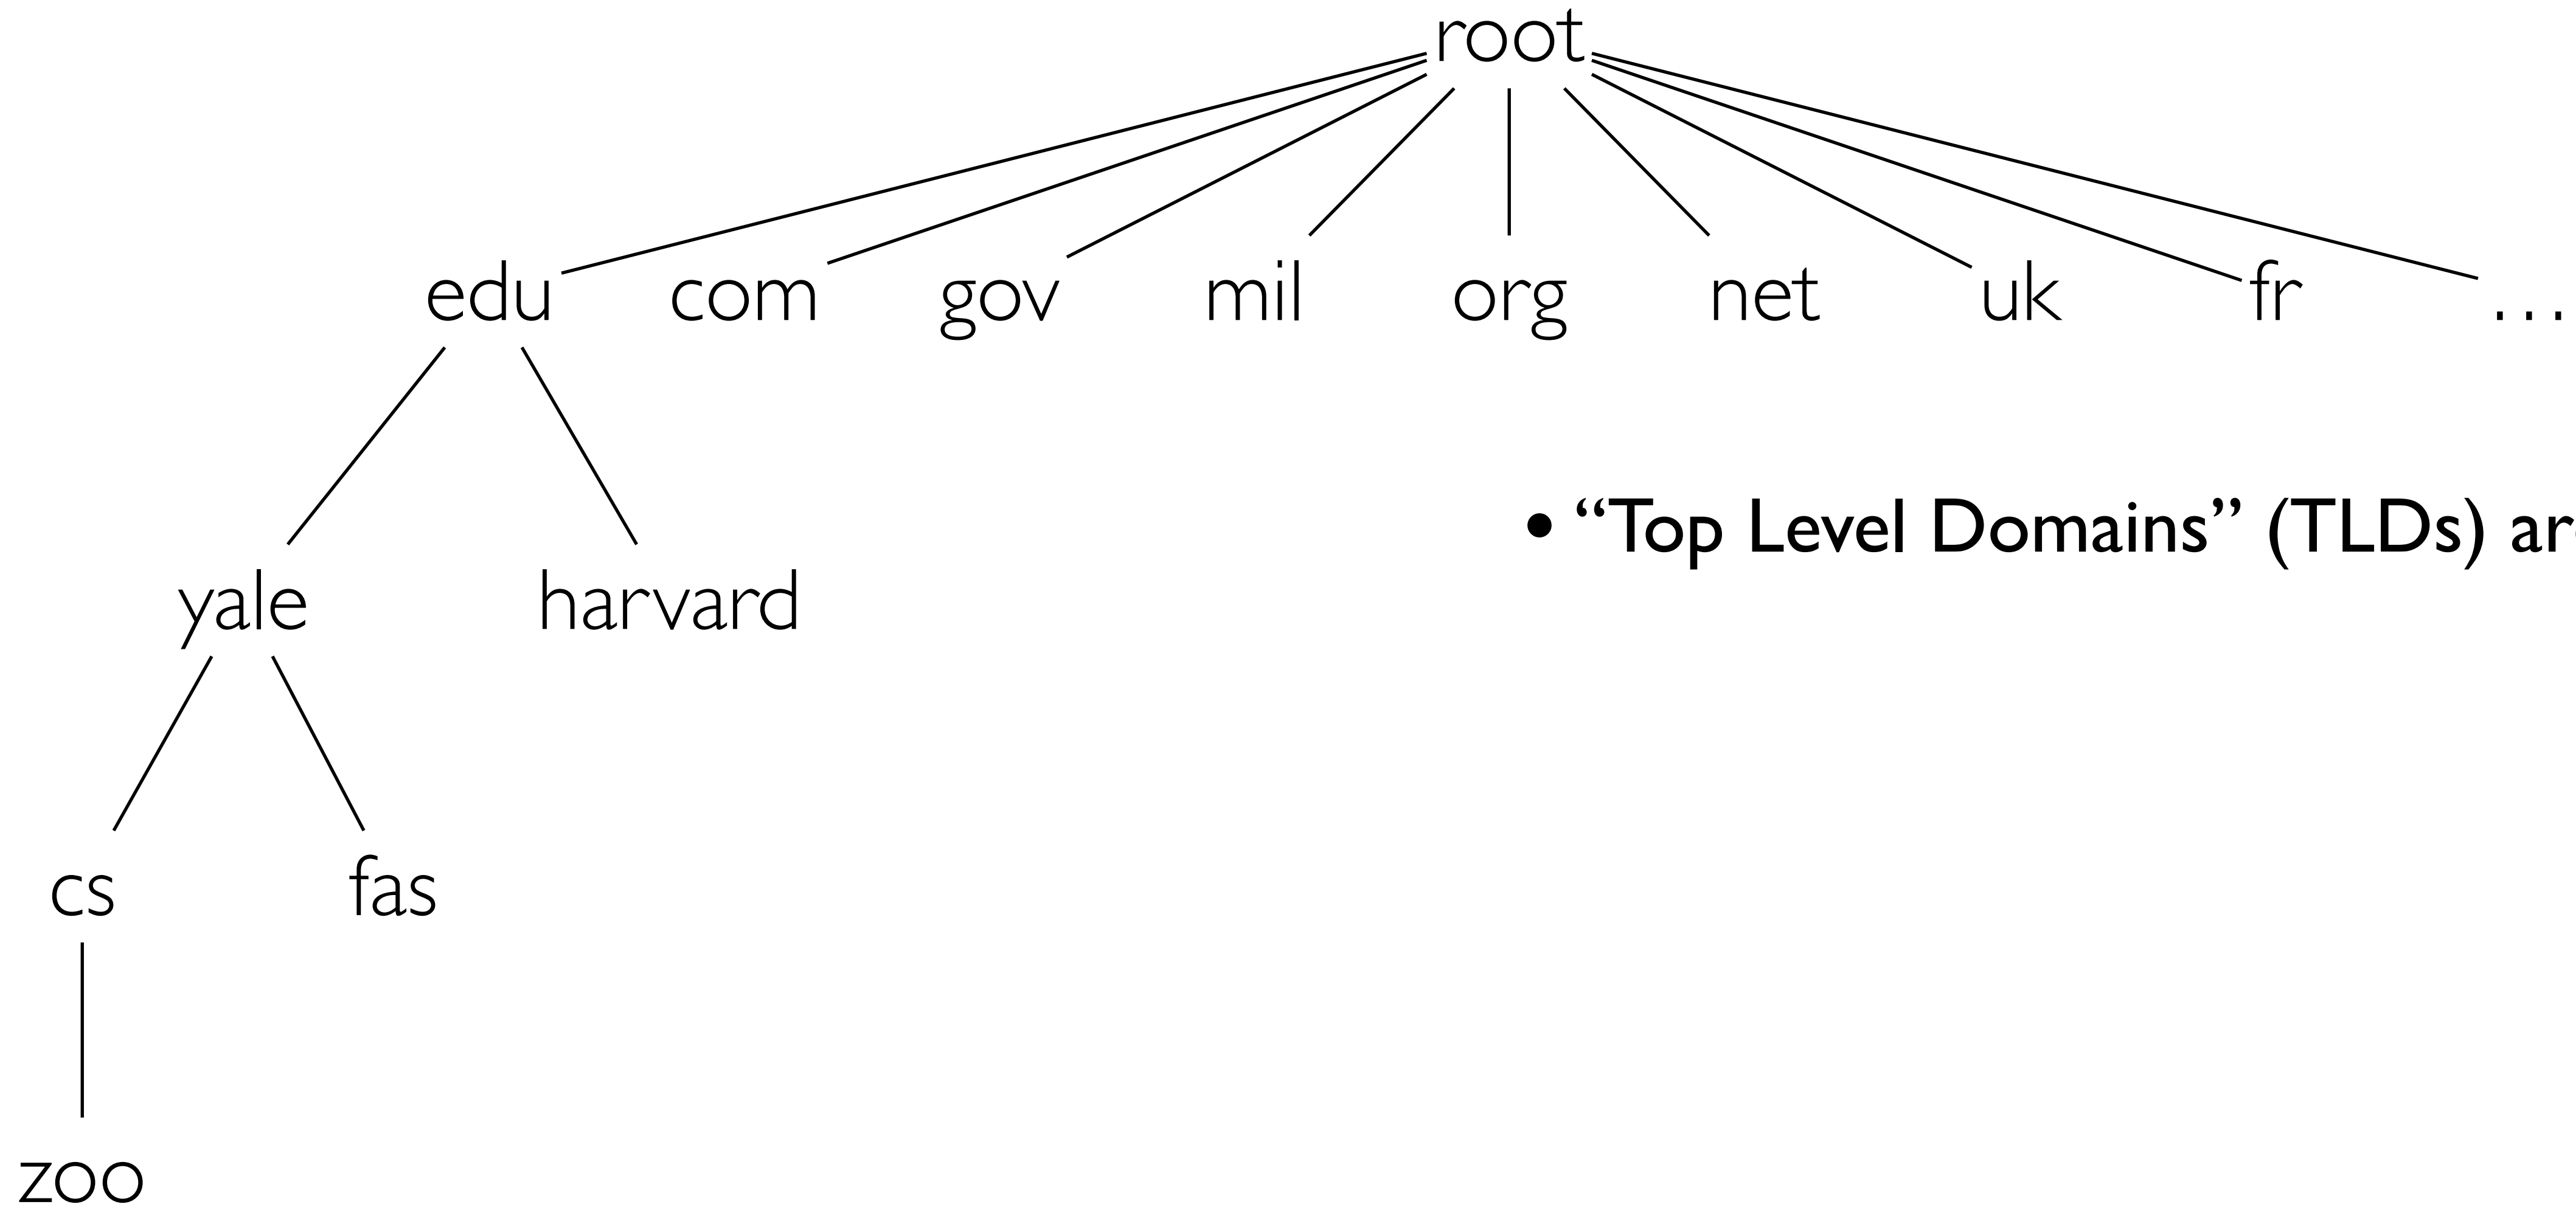
Three intertwined hierarchies

- **Hierarchical naming**
  - As opposed to flat namespace
- **Hierarchical administration**
  - As opposed to centralized administration
- **Hierarchical storage**
  - As opposed to centralized storage

# Hierarchical Naming

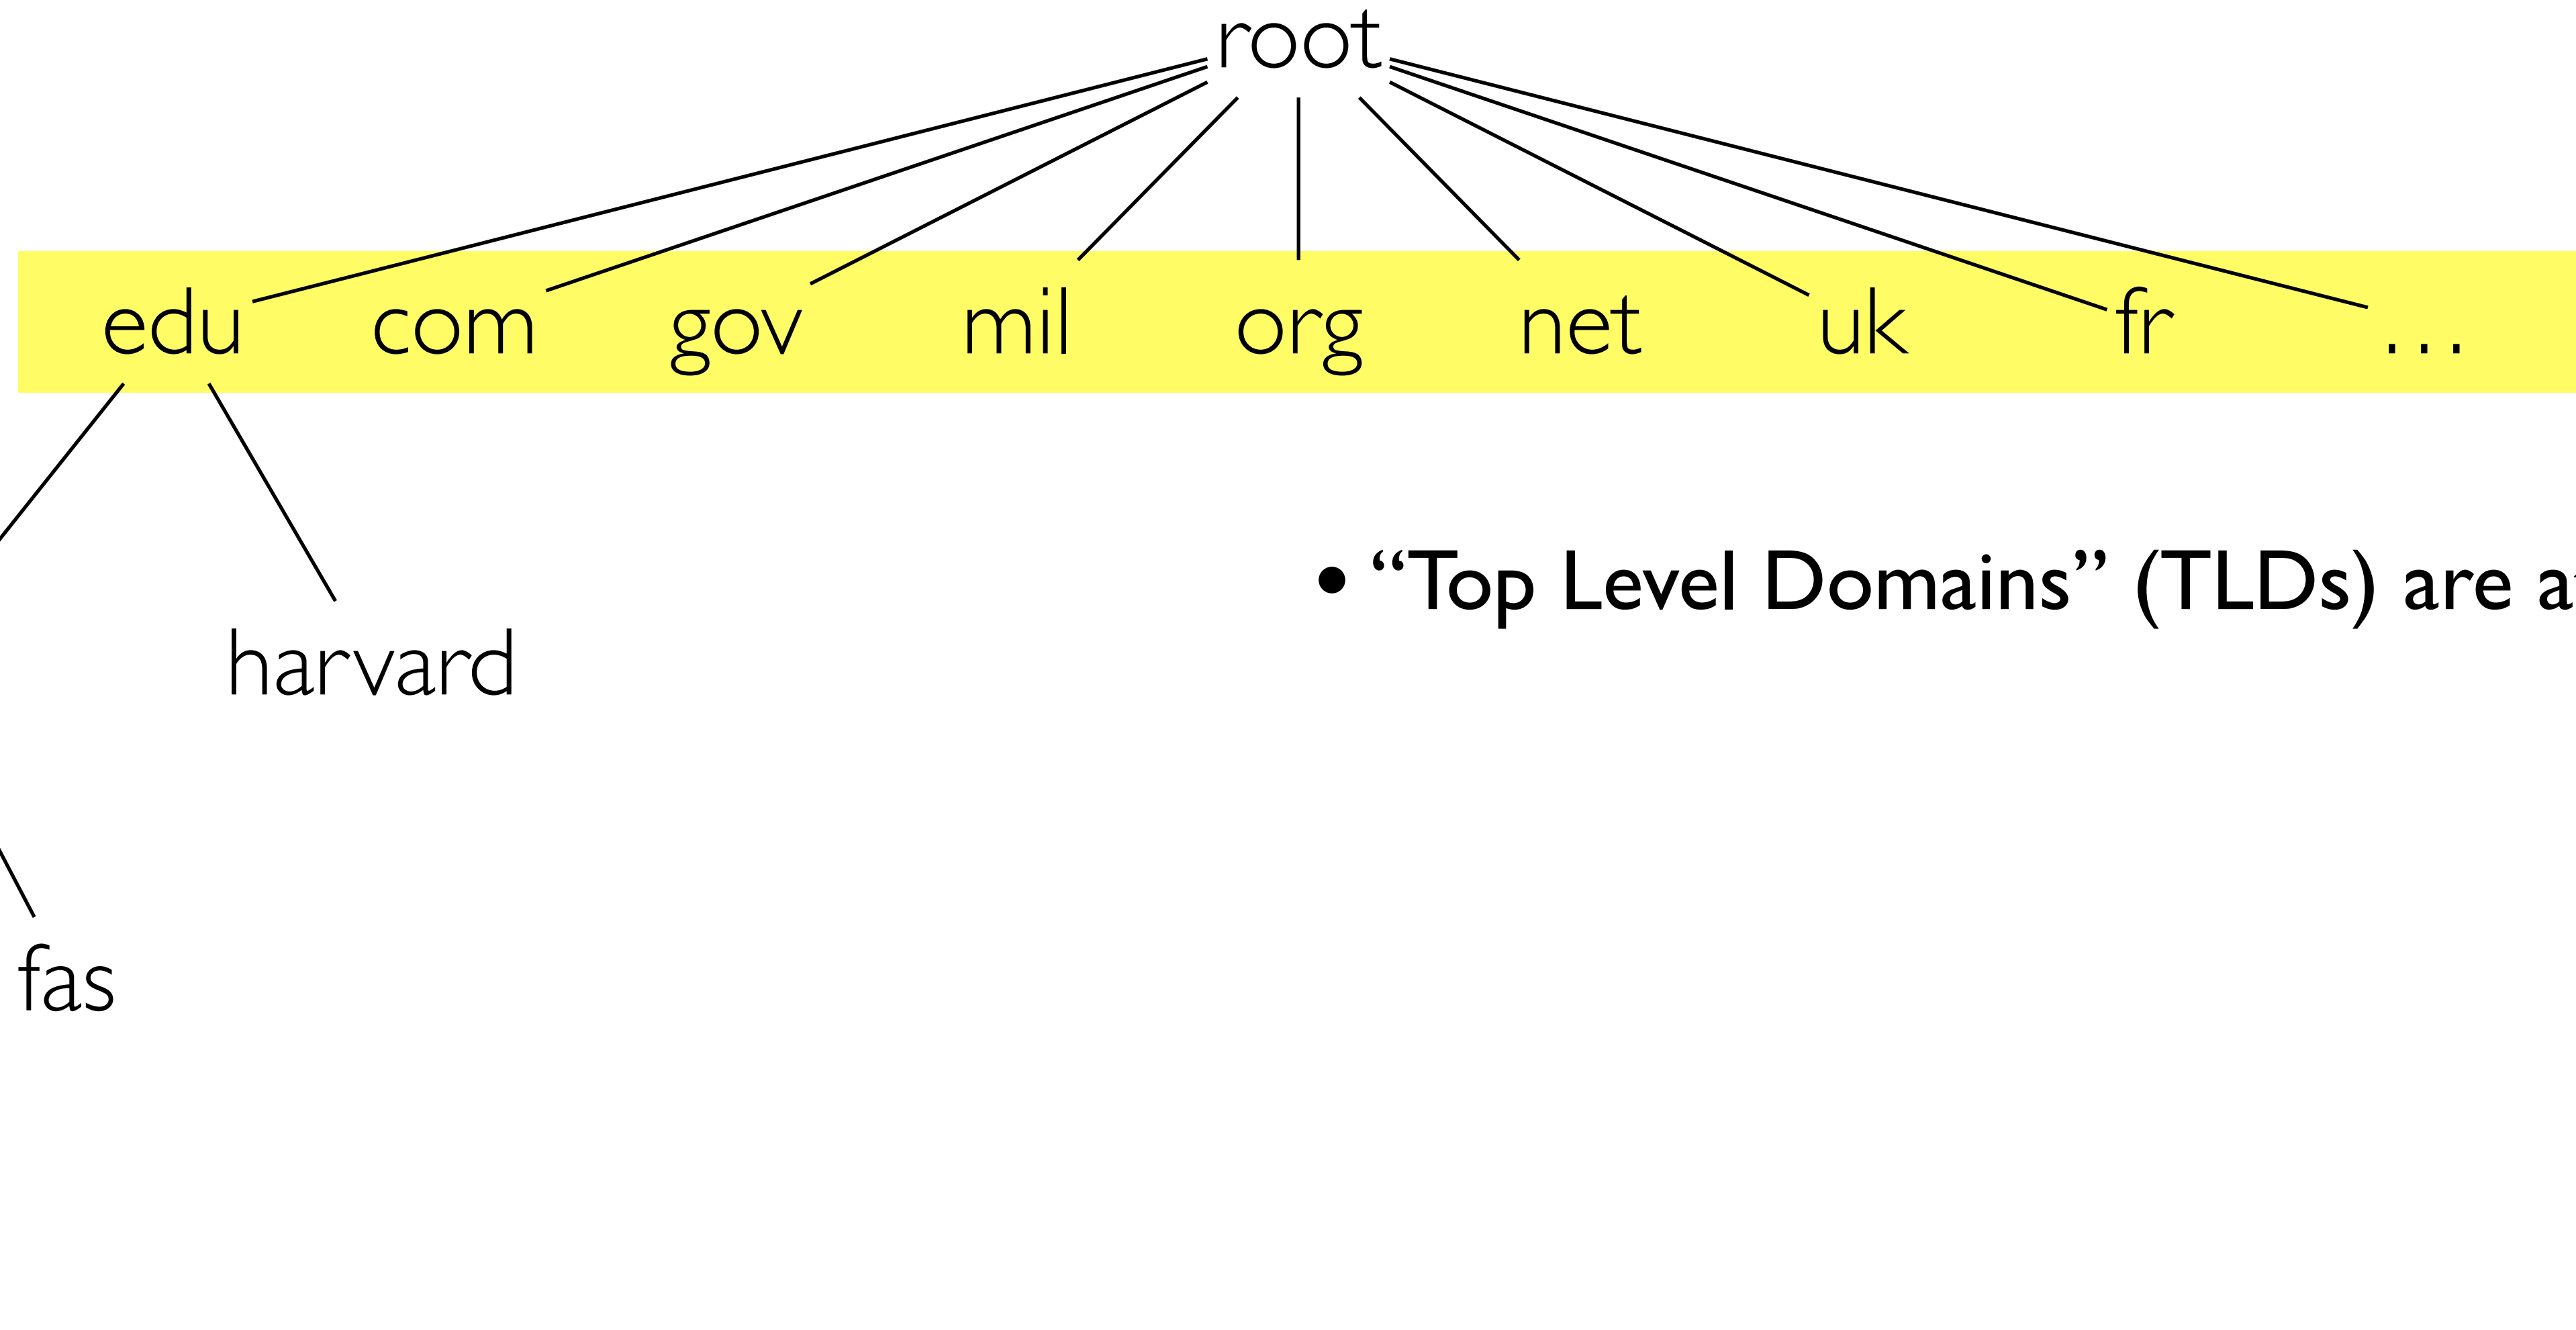


# Hierarchical Naming



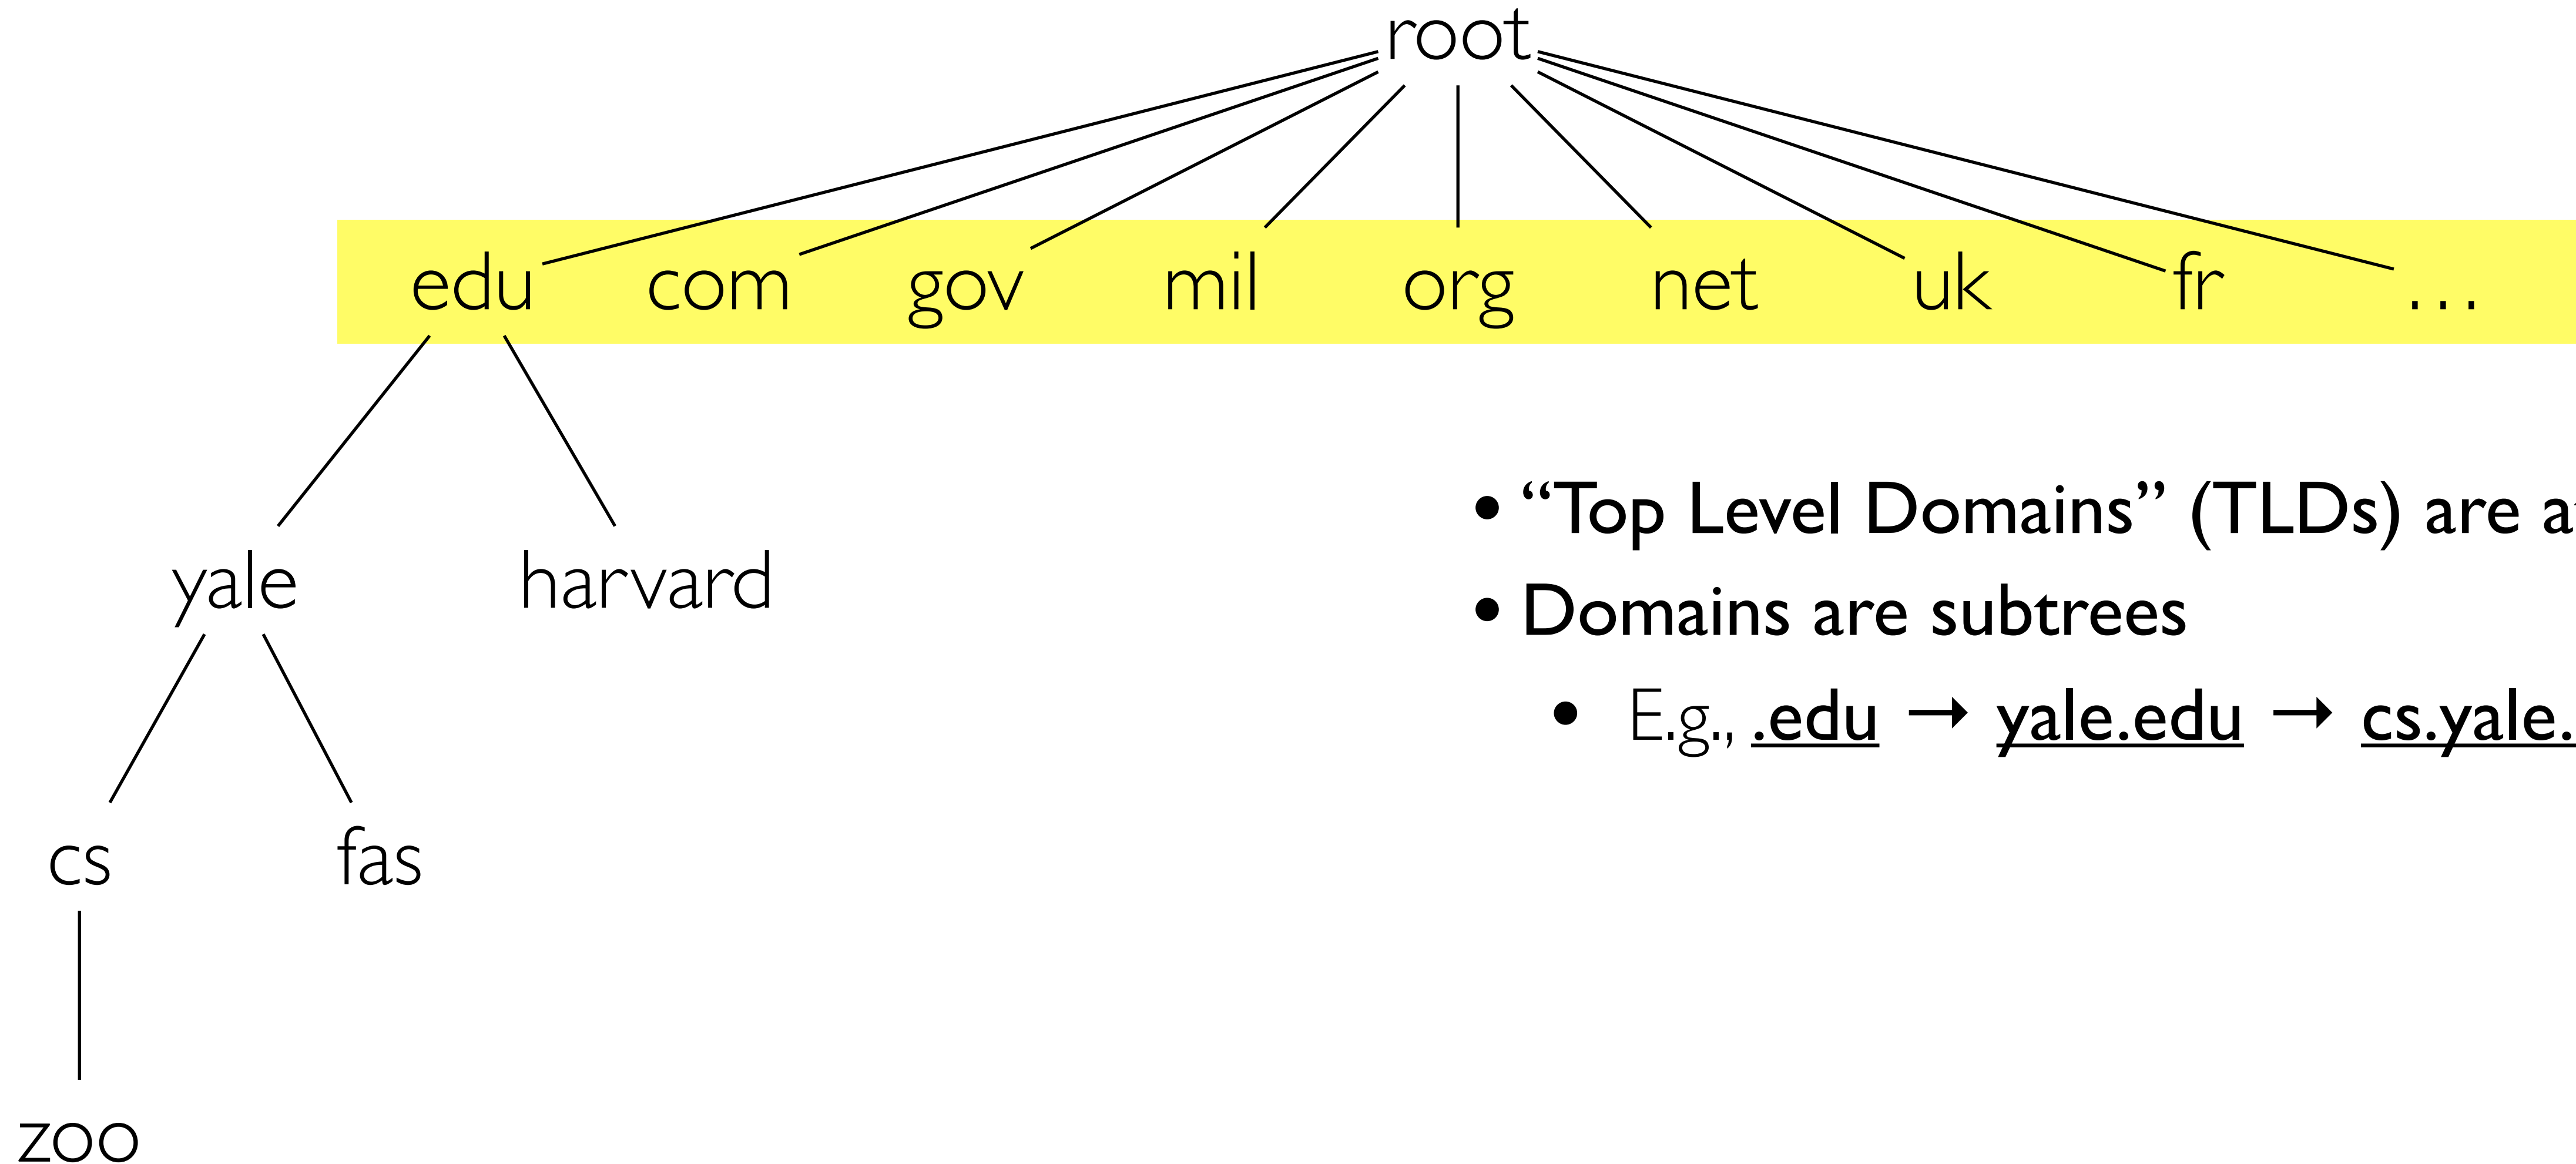
- “Top Level Domains” (TLDs) are at the top

# Hierarchical Naming



- “Top Level Domains” (TLDs) are at the top

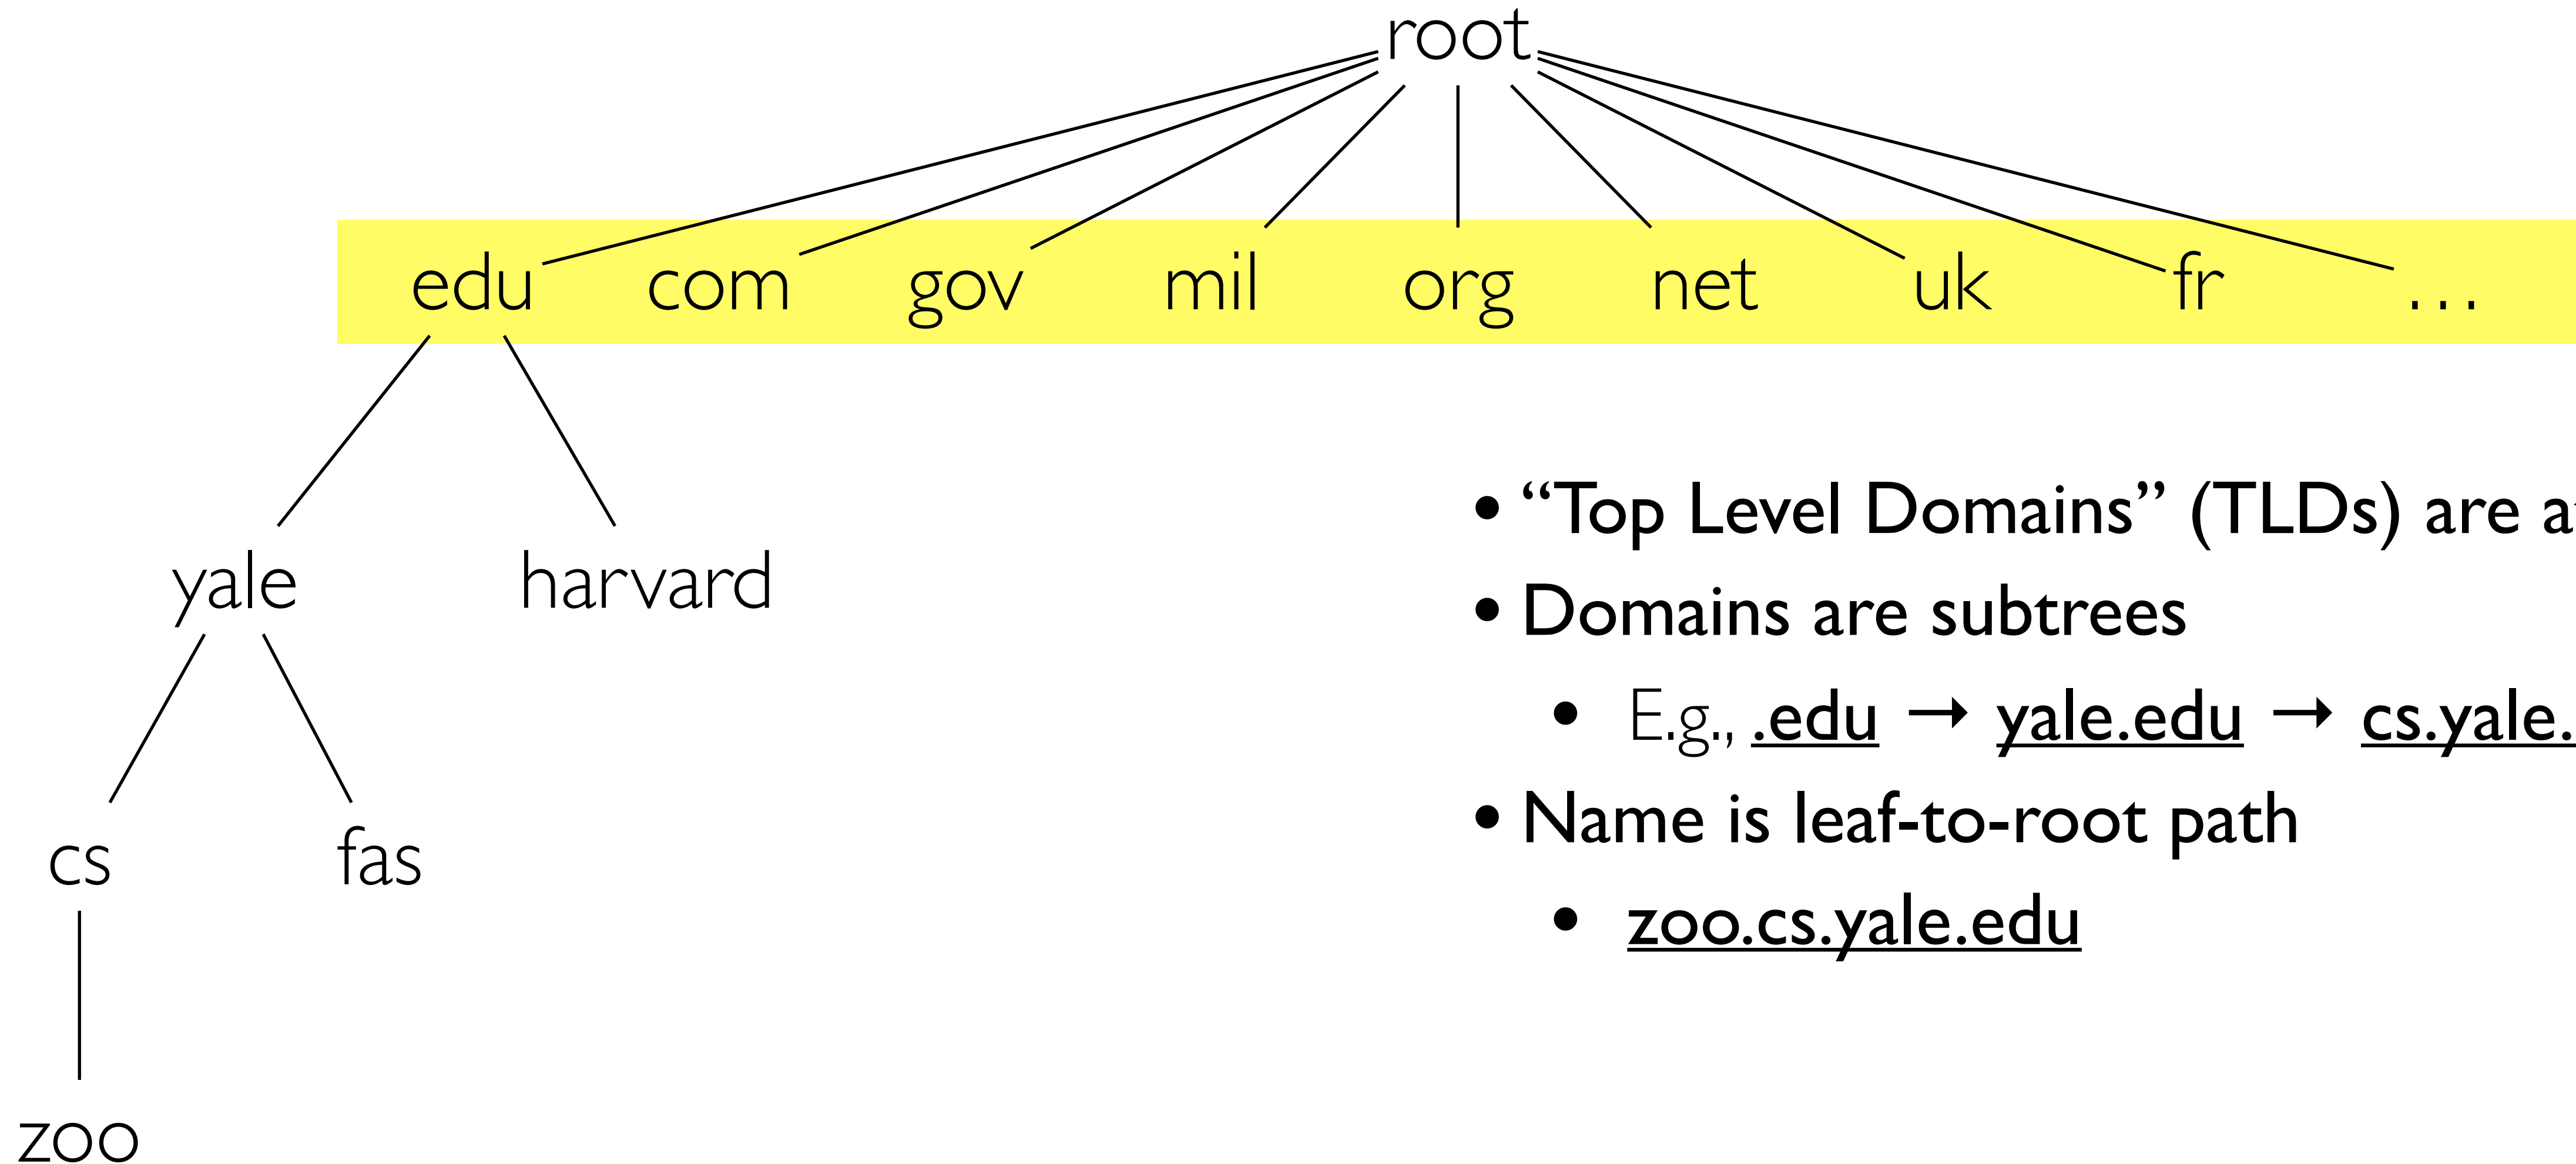
# Hierarchical Naming



- “Top Level Domains” (TLDs) are at the top
- Domains are subtrees
  - E.g., .edu → yale.edu → cs.yale.edu → ...

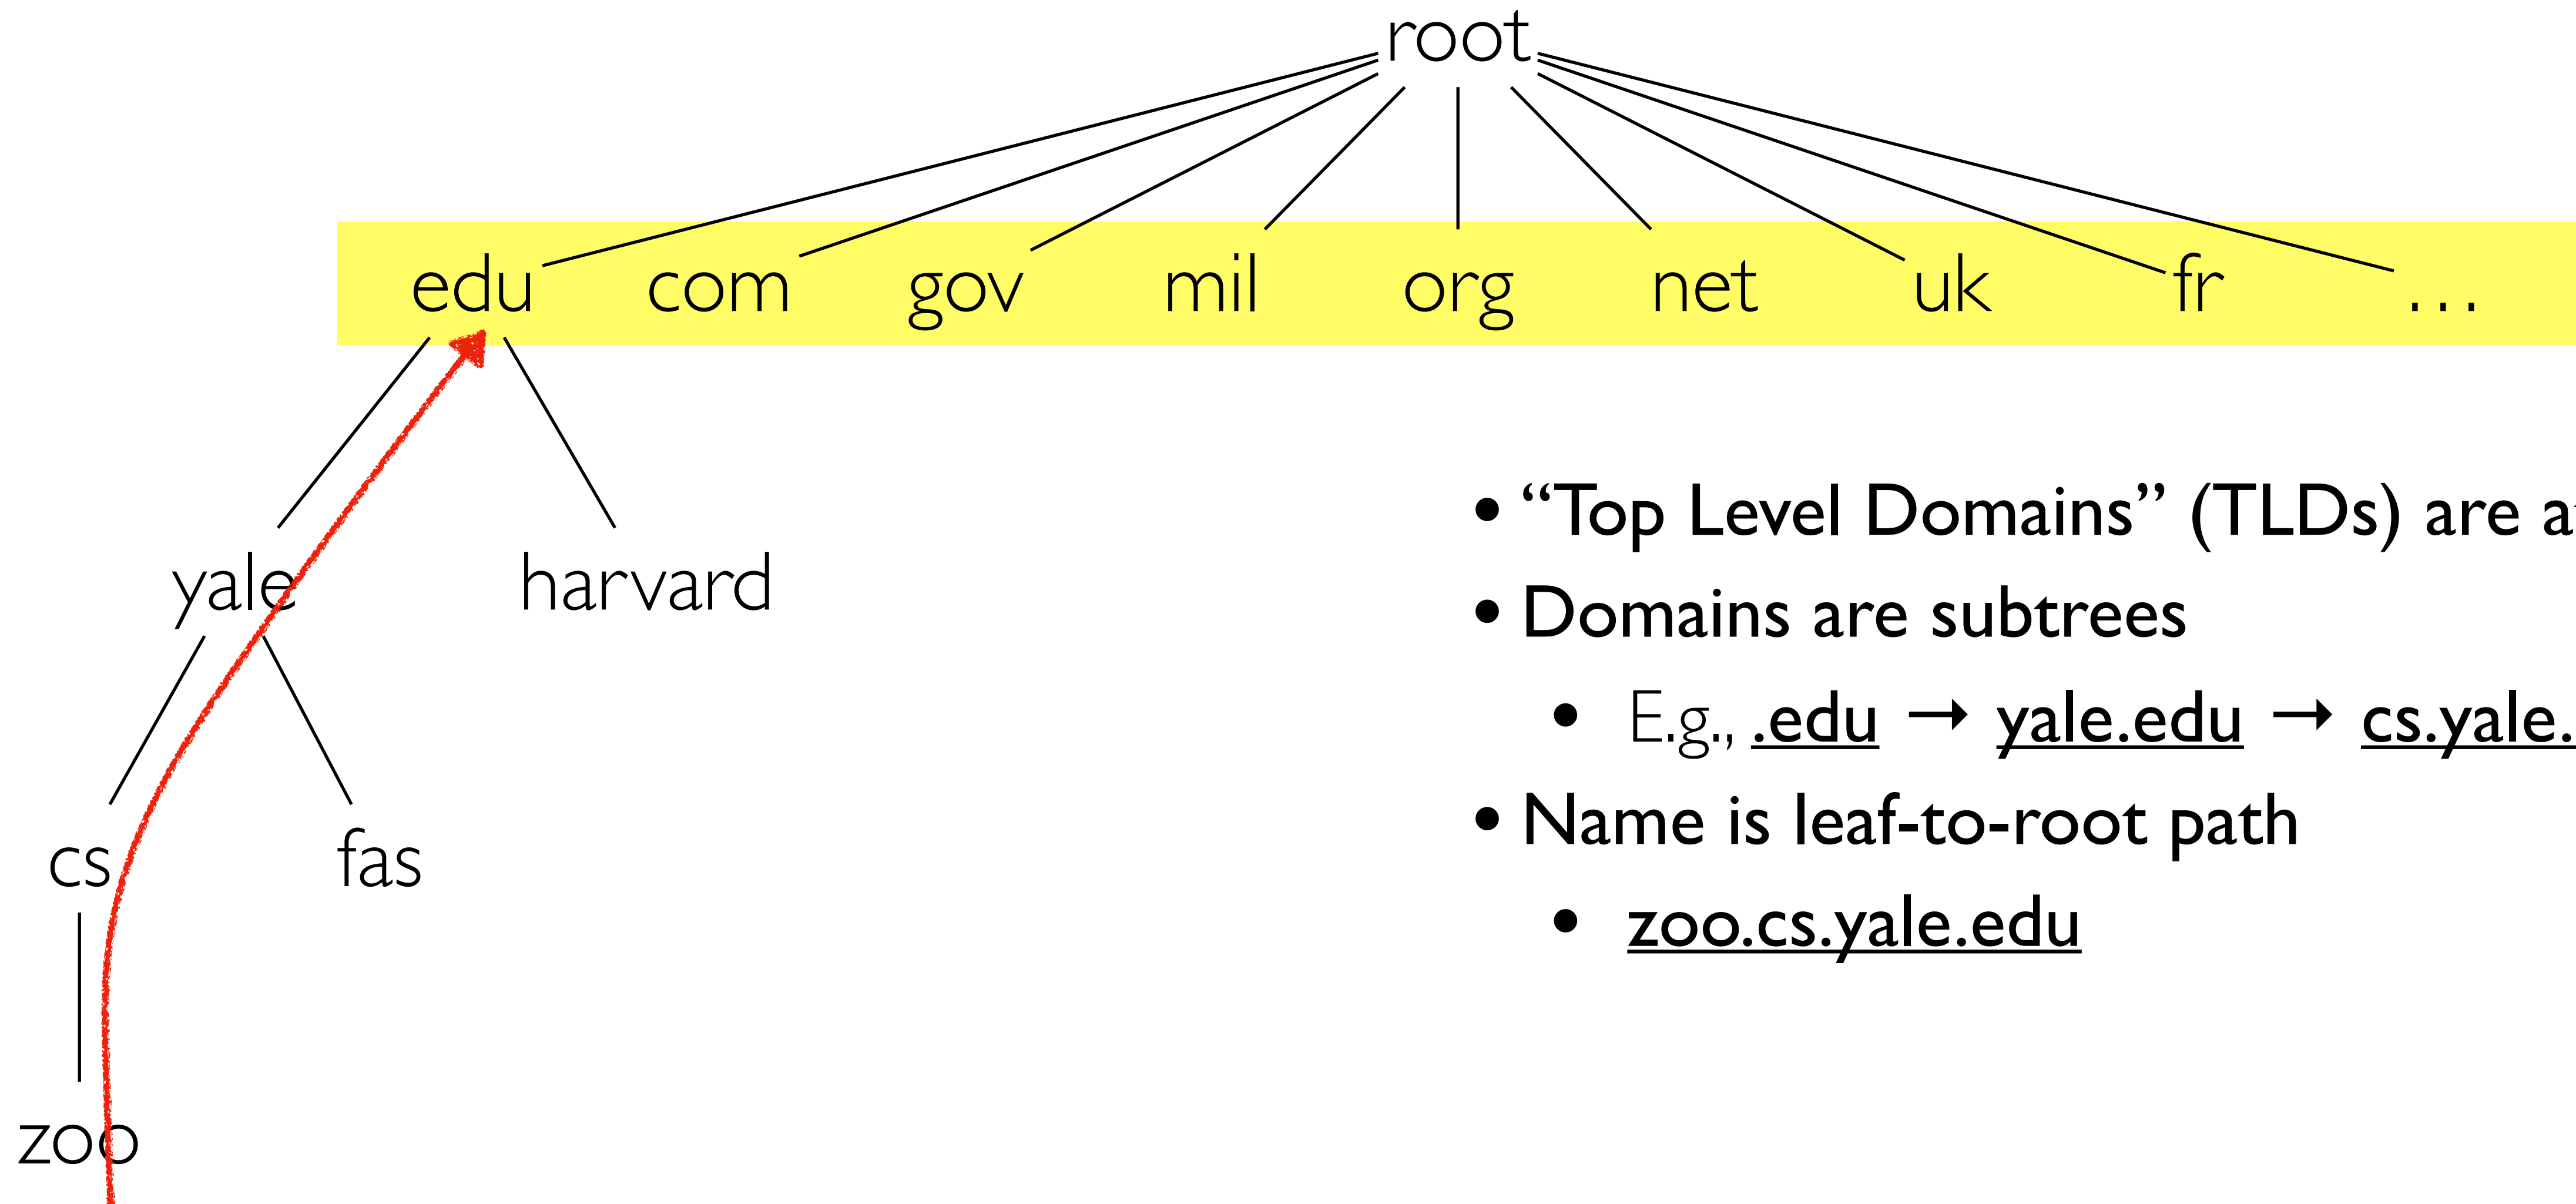


# Hierarchical Naming



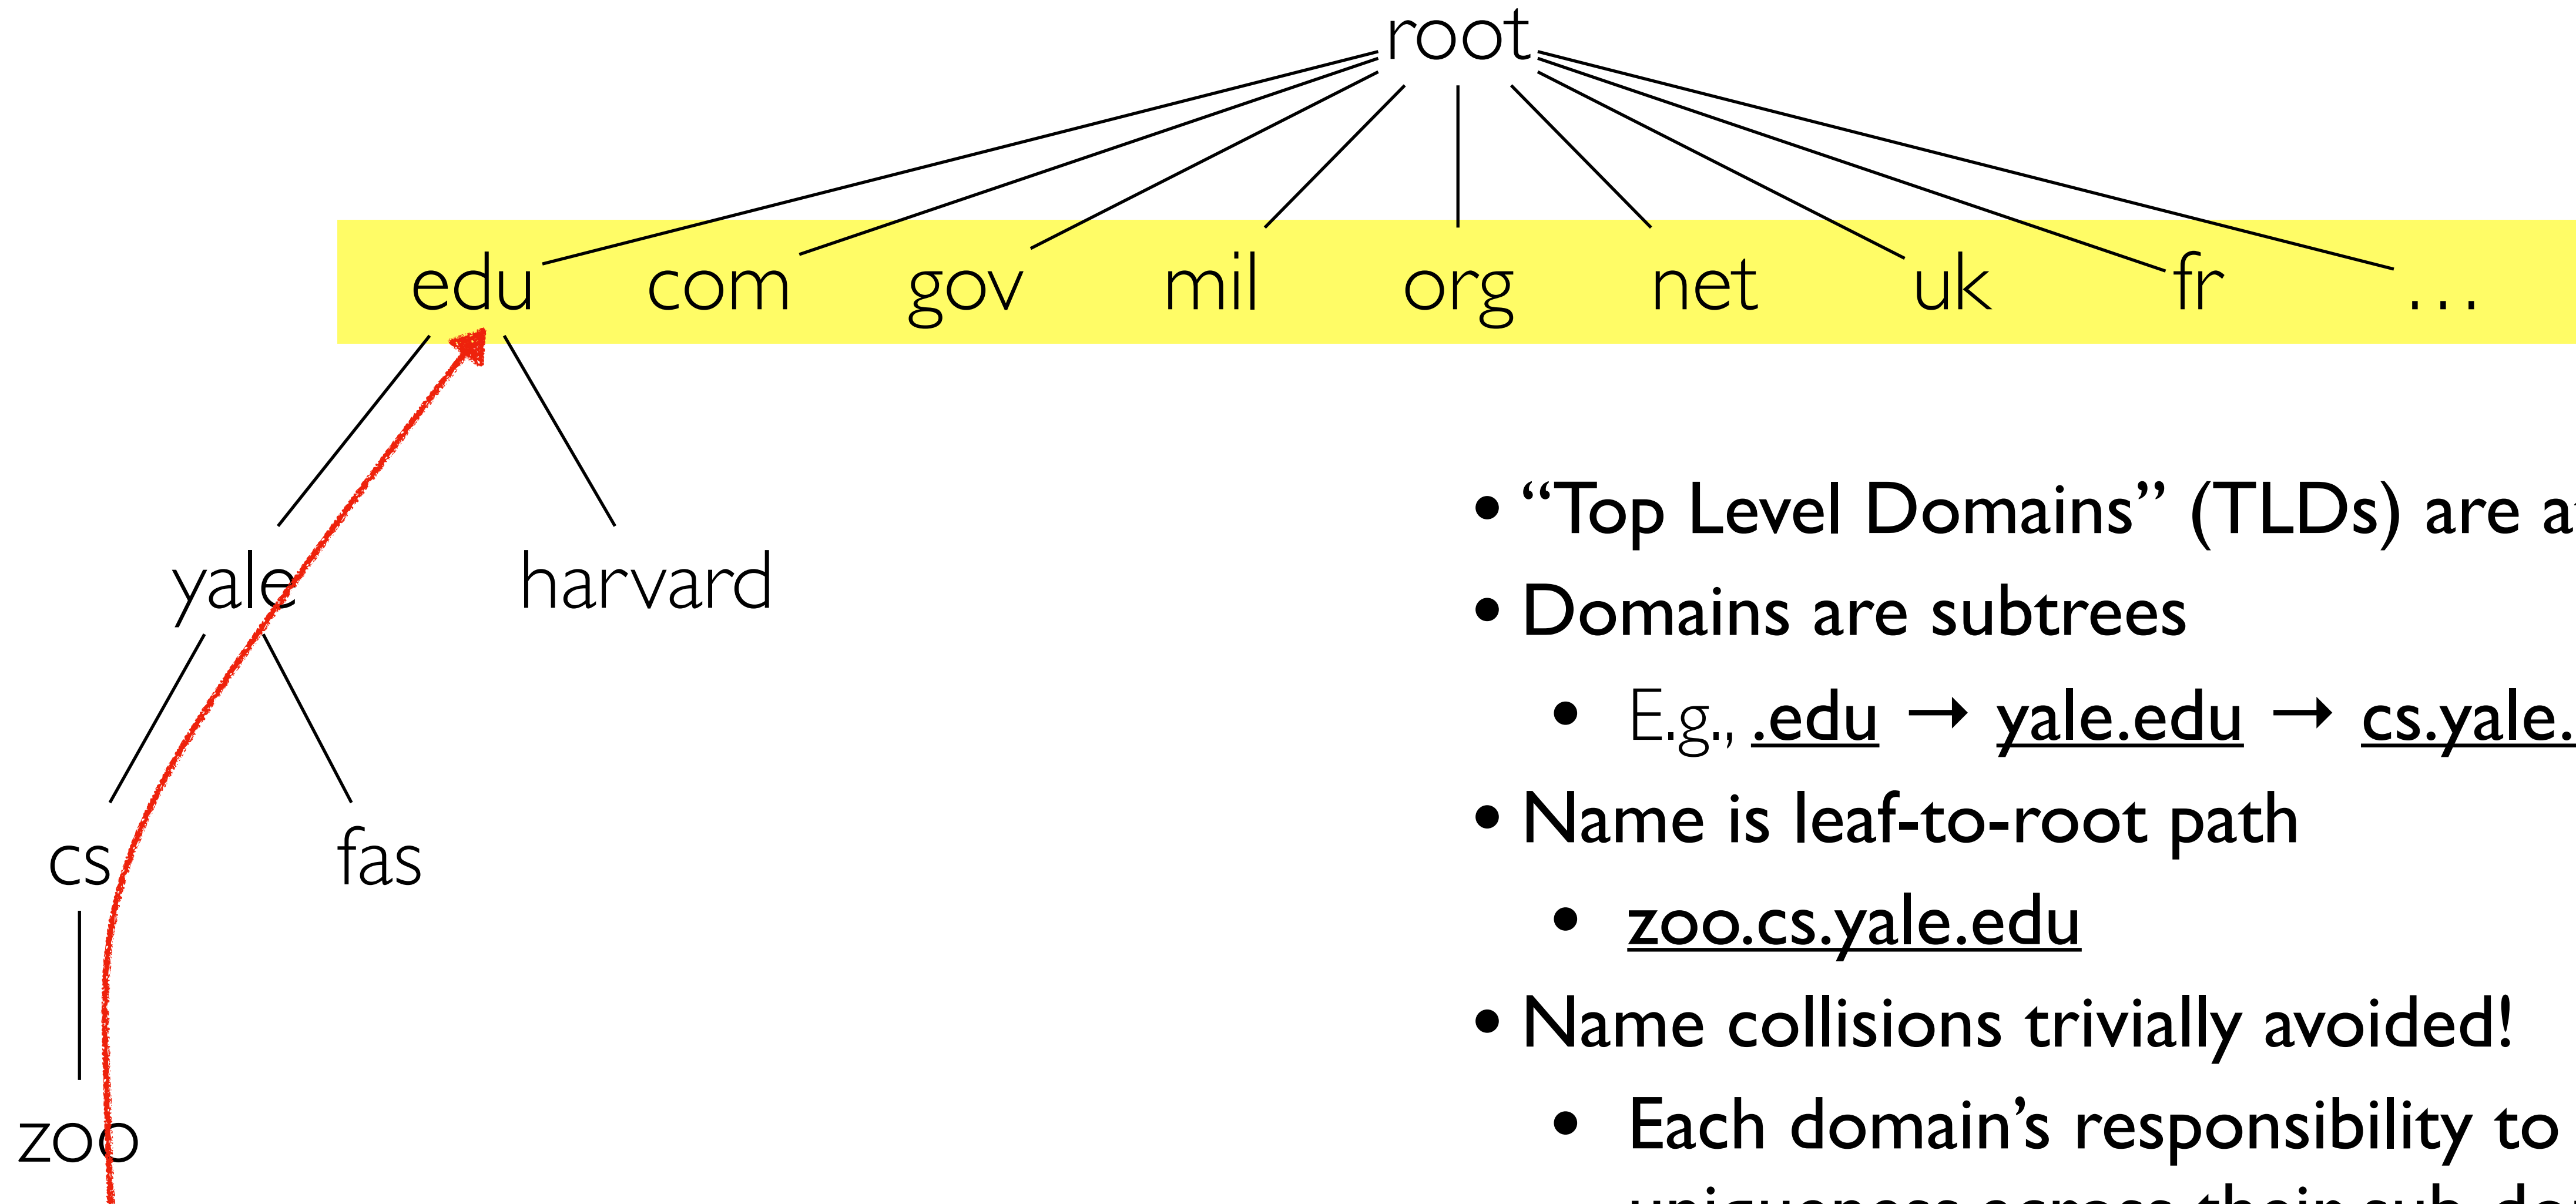
- “Top Level Domains” (TLDs) are at the top
- Domains are subtrees
  - E.g., .edu → yale.edu → cs.yale.edu → ...
- Name is leaf-to-root path
  - zoo.cs.yale.edu

# Hierarchical Naming



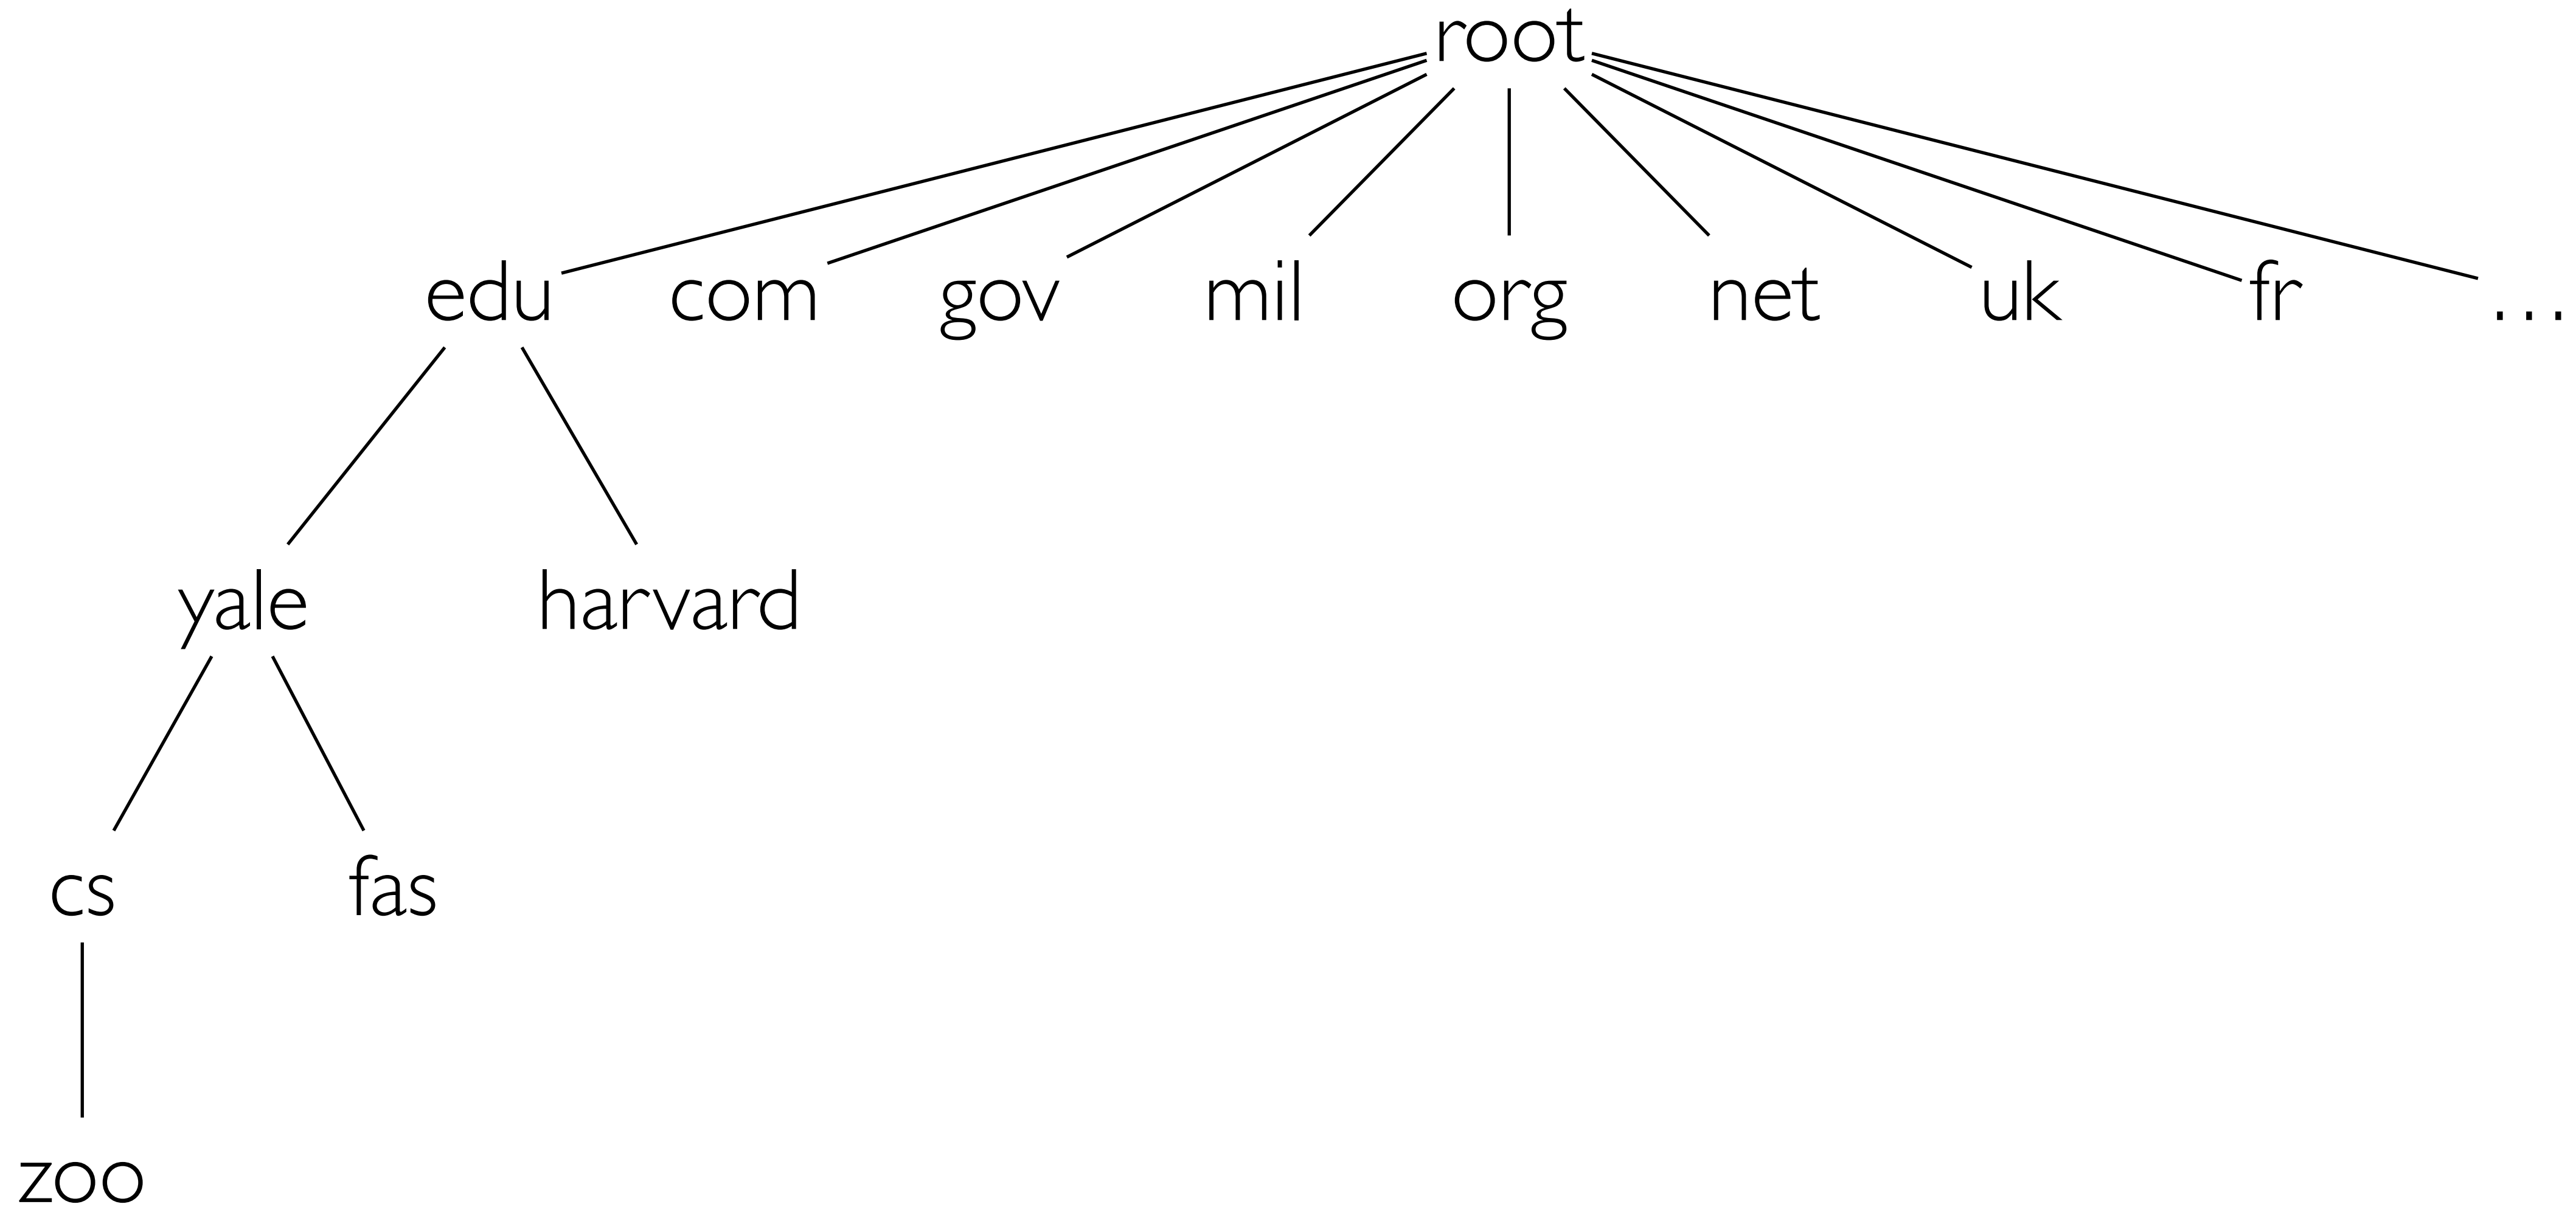
- “Top Level Domains” (TLDs) are at the top
- Domains are subtrees
  - E.g., .edu → yale.edu → cs.yale.edu → ...
- Name is leaf-to-root path
  - zoo.cs.yale.edu

# Hierarchical Naming

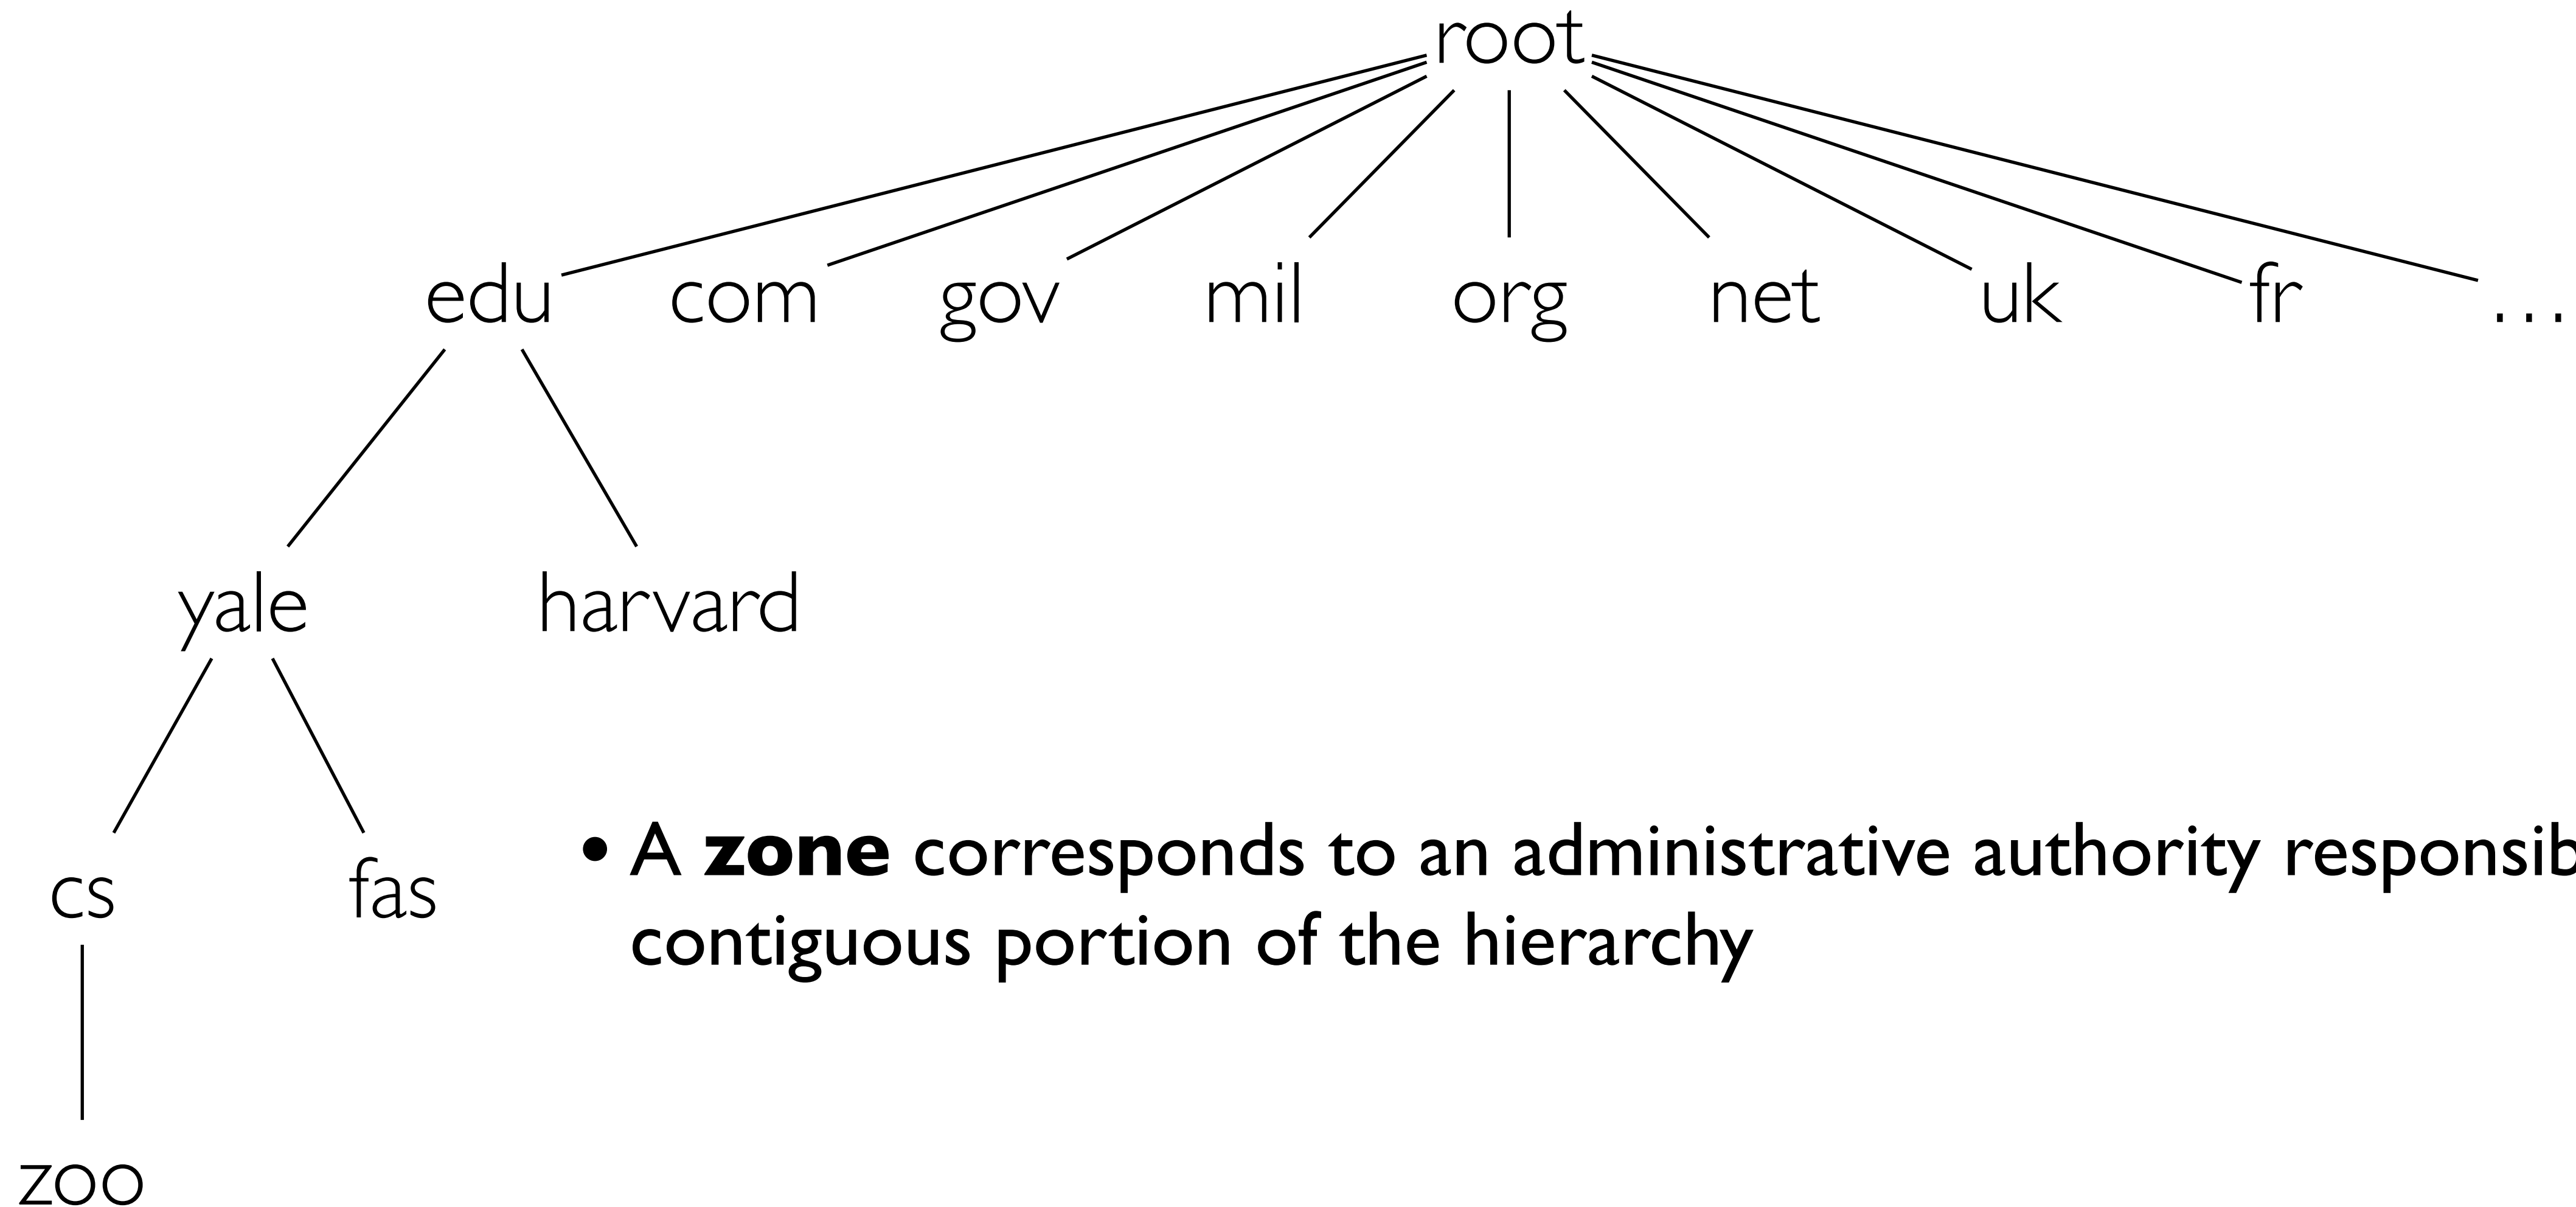


- “Top Level Domains” (TLDs) are at the top
- Domains are subtrees
  - E.g., .edu → yale.edu → cs.yale.edu → ...
- Name is leaf-to-root path
  - zoo.cs.yale.edu
- Name collisions trivially avoided!
  - Each domain’s responsibility to maintain uniqueness across their sub-domains



# Hierarchical Administration



# Hierarchical Administration

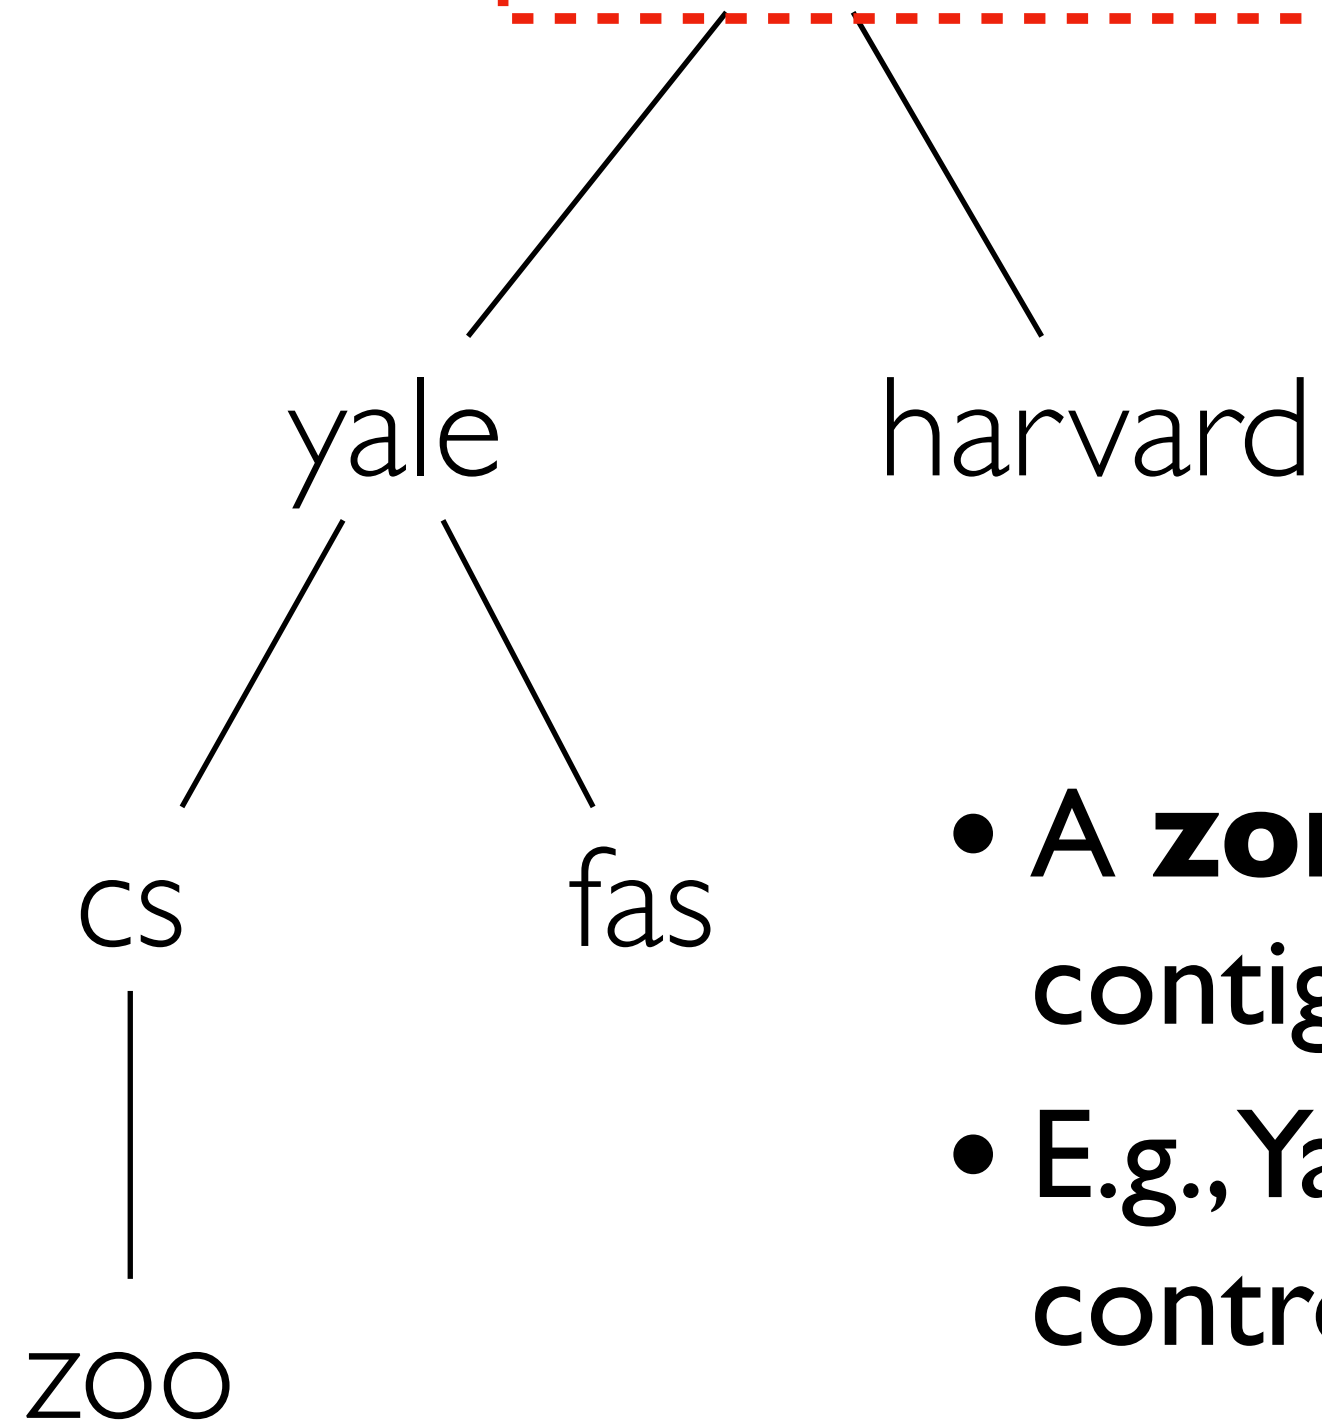



# Hierarchical Administration



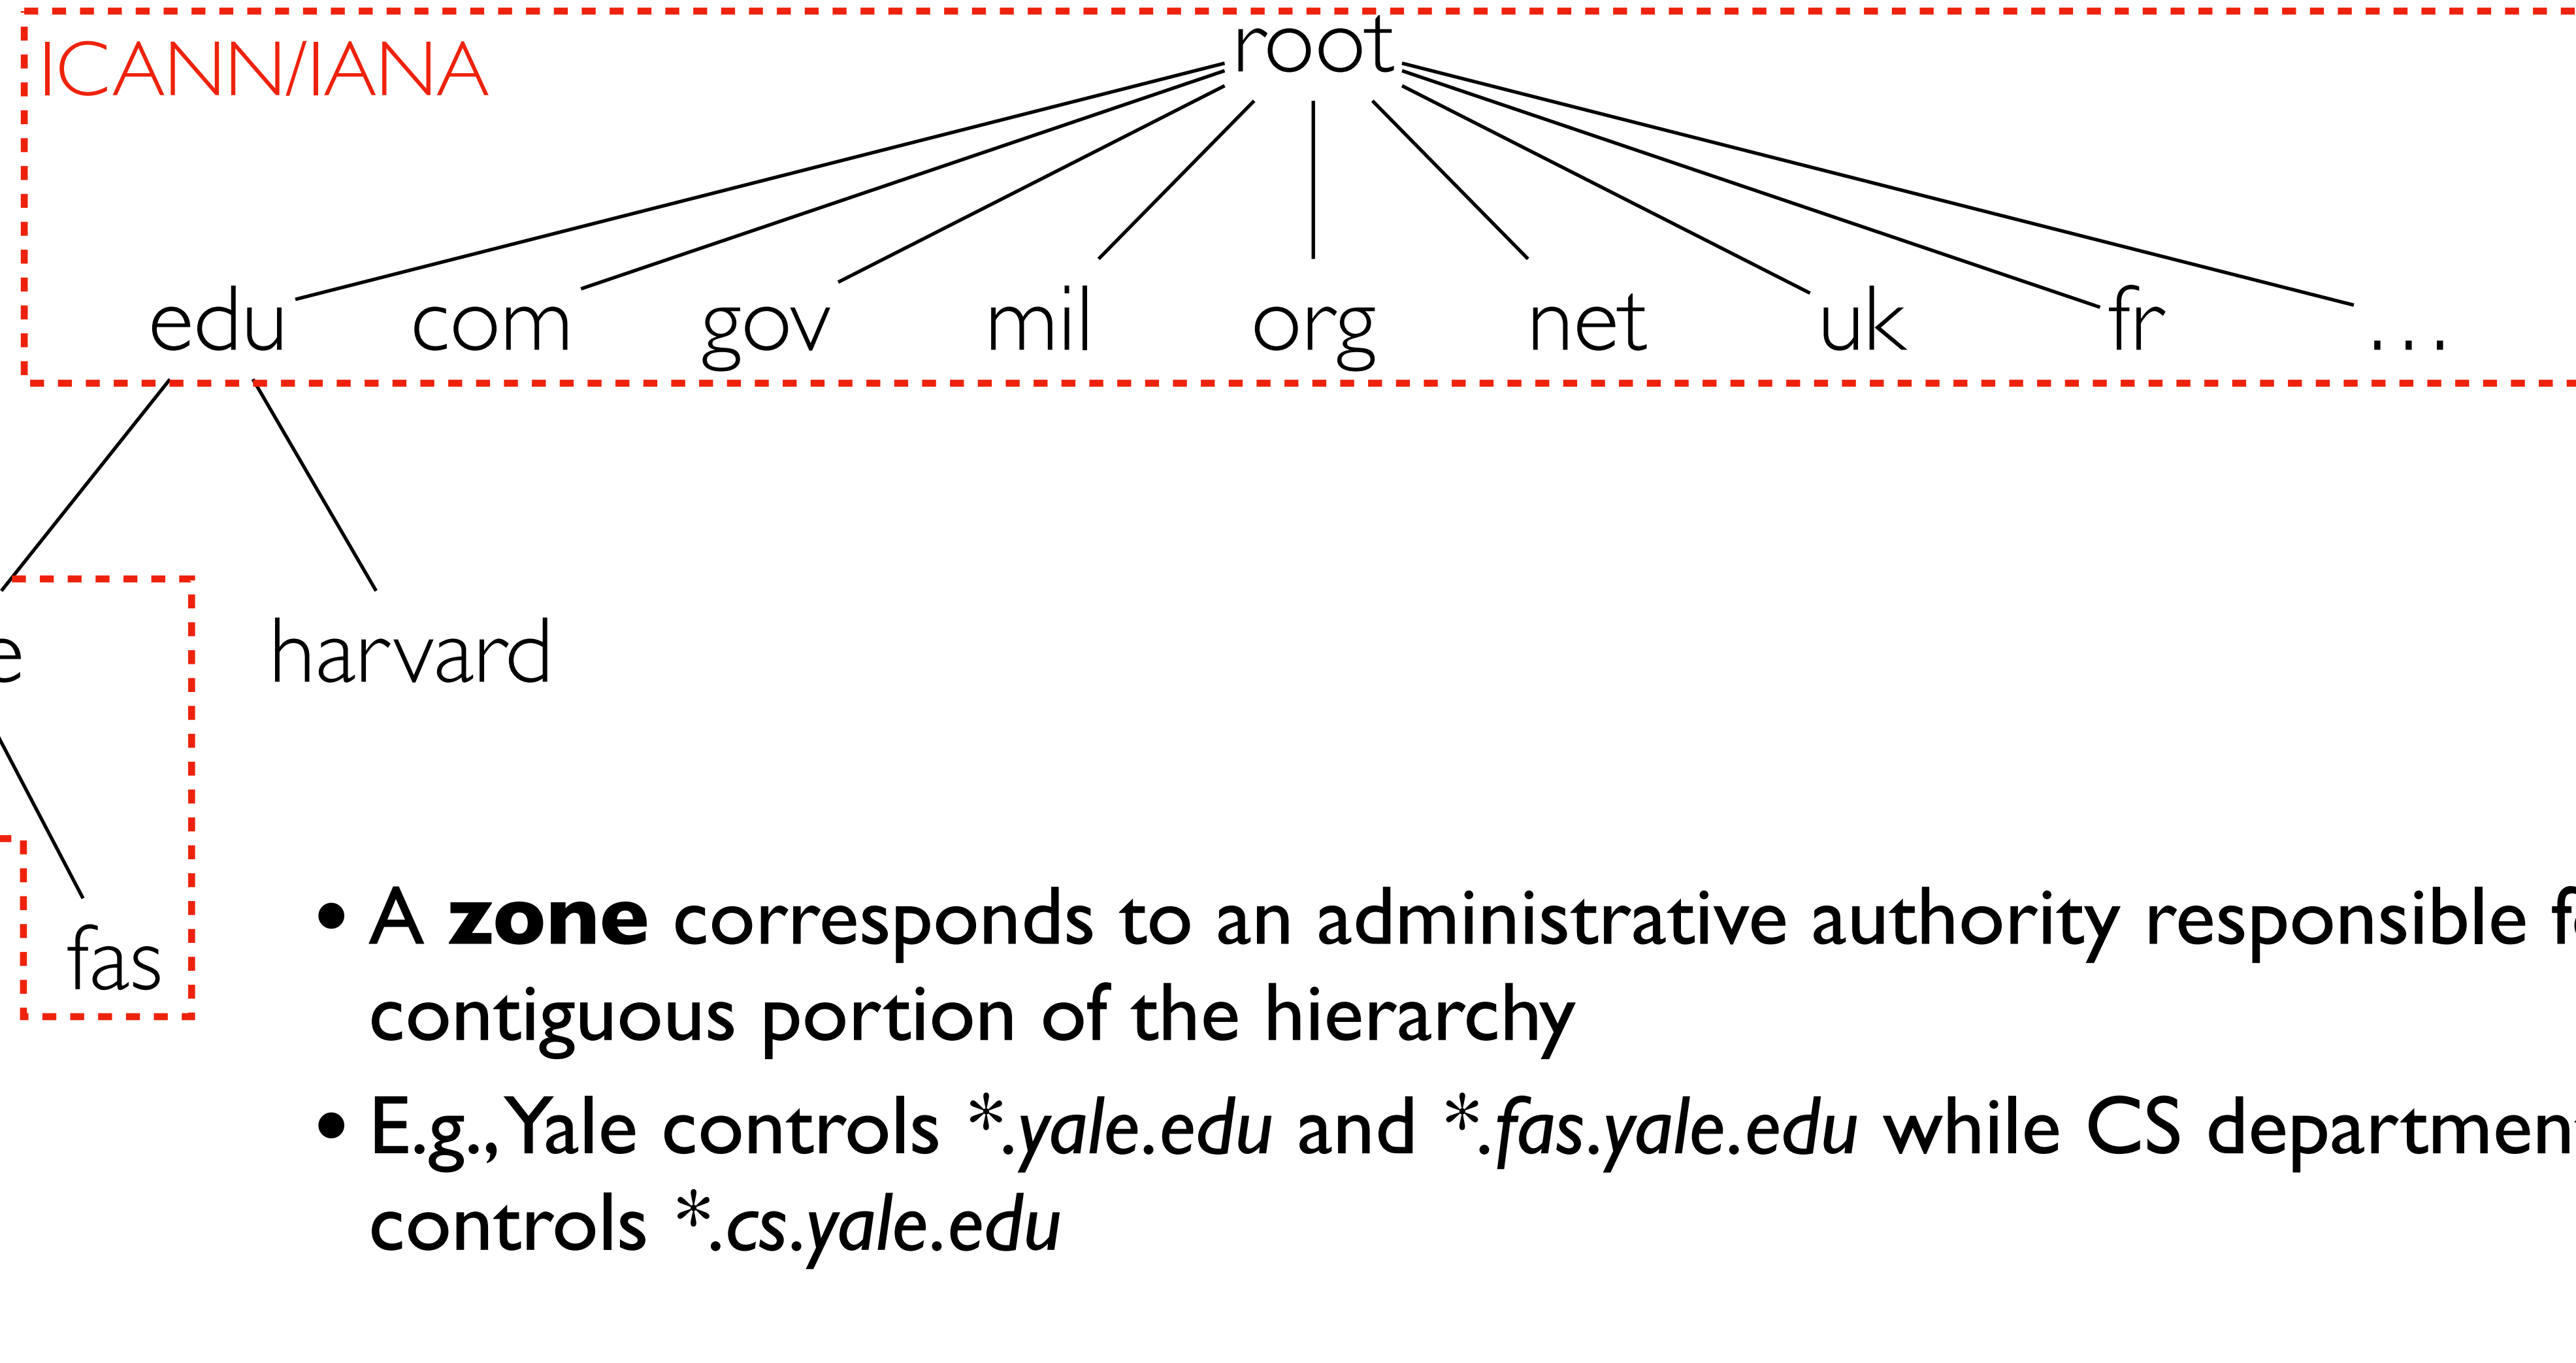
- **A zone** corresponds to an administrative authority responsible for a contiguous portion of the hierarchy

# Hierarchical Administration



- A **zone** corresponds to an administrative authority responsible for a contiguous portion of the hierarchy
- E.g., Yale controls *\*.yale.edu* and *\*.fas.yale.edu* while CS department controls *\*.cs.yale.edu*

# Hierarchical Administration





# Storage Hierarchy

# Storage Hierarchy

- Top of hierarchy: Root servers
  - Location hardwired into relevant servers
  - Store addresses **all** next level DNS servers...

# Storage Hierarchy

- **Top of hierarchy: Root servers**
  - Location hardwired into relevant servers
  - Store addresses **all** next level DNS servers...
- **Next Level: Top-Level Domain (TLD) servers**
  - .com, .edu, etc.; managed professionally
  - Store addresses of next level DNS servers **under them...**

# Storage Hierarchy

- **Top of hierarchy: Root servers**
  - Location hardwired into relevant servers
  - Store addresses **all** next level DNS servers...
- **Next Level: Top-Level Domain (TLD) servers**
  - .com, .edu, etc.; managed professionally
  - Store addresses of next level DNS servers **under them...**
- **Bottom Level: Authoritative DNS servers**
  - Store the **actual name-to-address mappings**
  - Maintained by the corresponding administrative authority

# Storage Hierarchy

# Storage Hierarchy

- Every server knows the address of the root name server

# Storage Hierarchy

- Every server knows the address of the root name server
- Root servers know the address of all TLD servers

# Storage Hierarchy

- Every server knows the address of the root name server
- Root servers know the address of all TLD servers
- ...



# Storage Hierarchy

- Every server knows the address of the root name server
- Root servers know the address of all TLD servers
- ...
- An authoritative DNS server stores name-to-address mappings (“resource records”) for all DNS names in the domain that it has authority for

# Storage Hierarchy

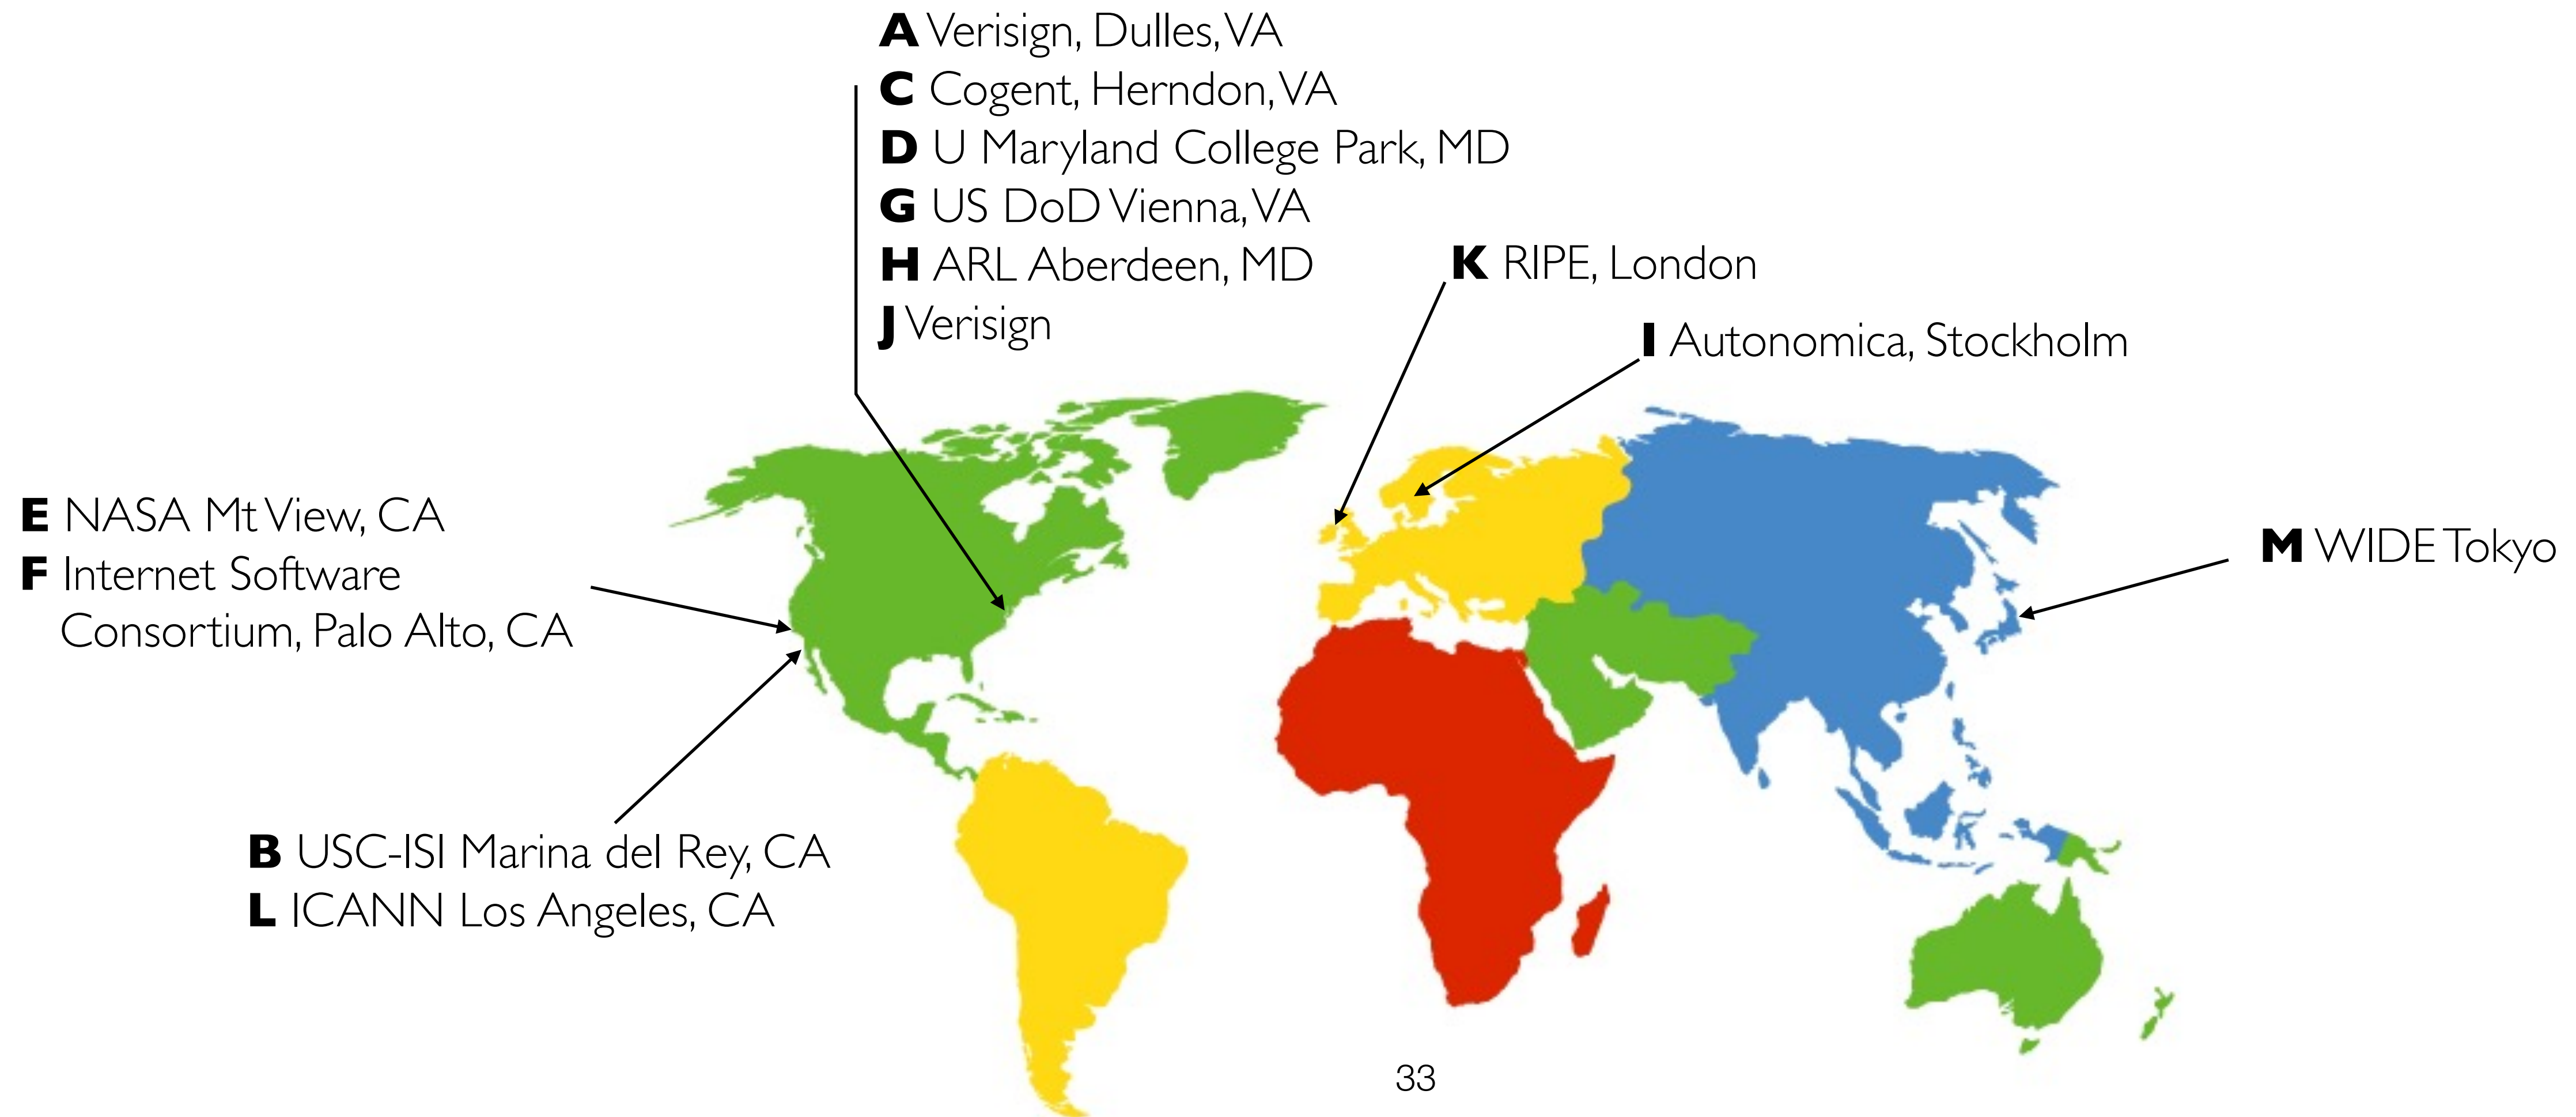
- Every server knows the address of the root name server
- Root servers know the address of all TLD servers
- ...
- An authoritative DNS server stores name-to-address mappings (“resource records”) for all DNS names in the domain that it has authority for
- Each server stores a subset of the total DNS database

# Storage Hierarchy

- Every server knows the address of the root name server
- Root servers know the address of all TLD servers
- ...
- An authoritative DNS server stores name-to-address mappings (“resource records”) for all DNS names in the domain that it has authority for
- Each server stores a subset of the total DNS database
- Each server can discover the server(s) responsible for any portion of the hierarchy

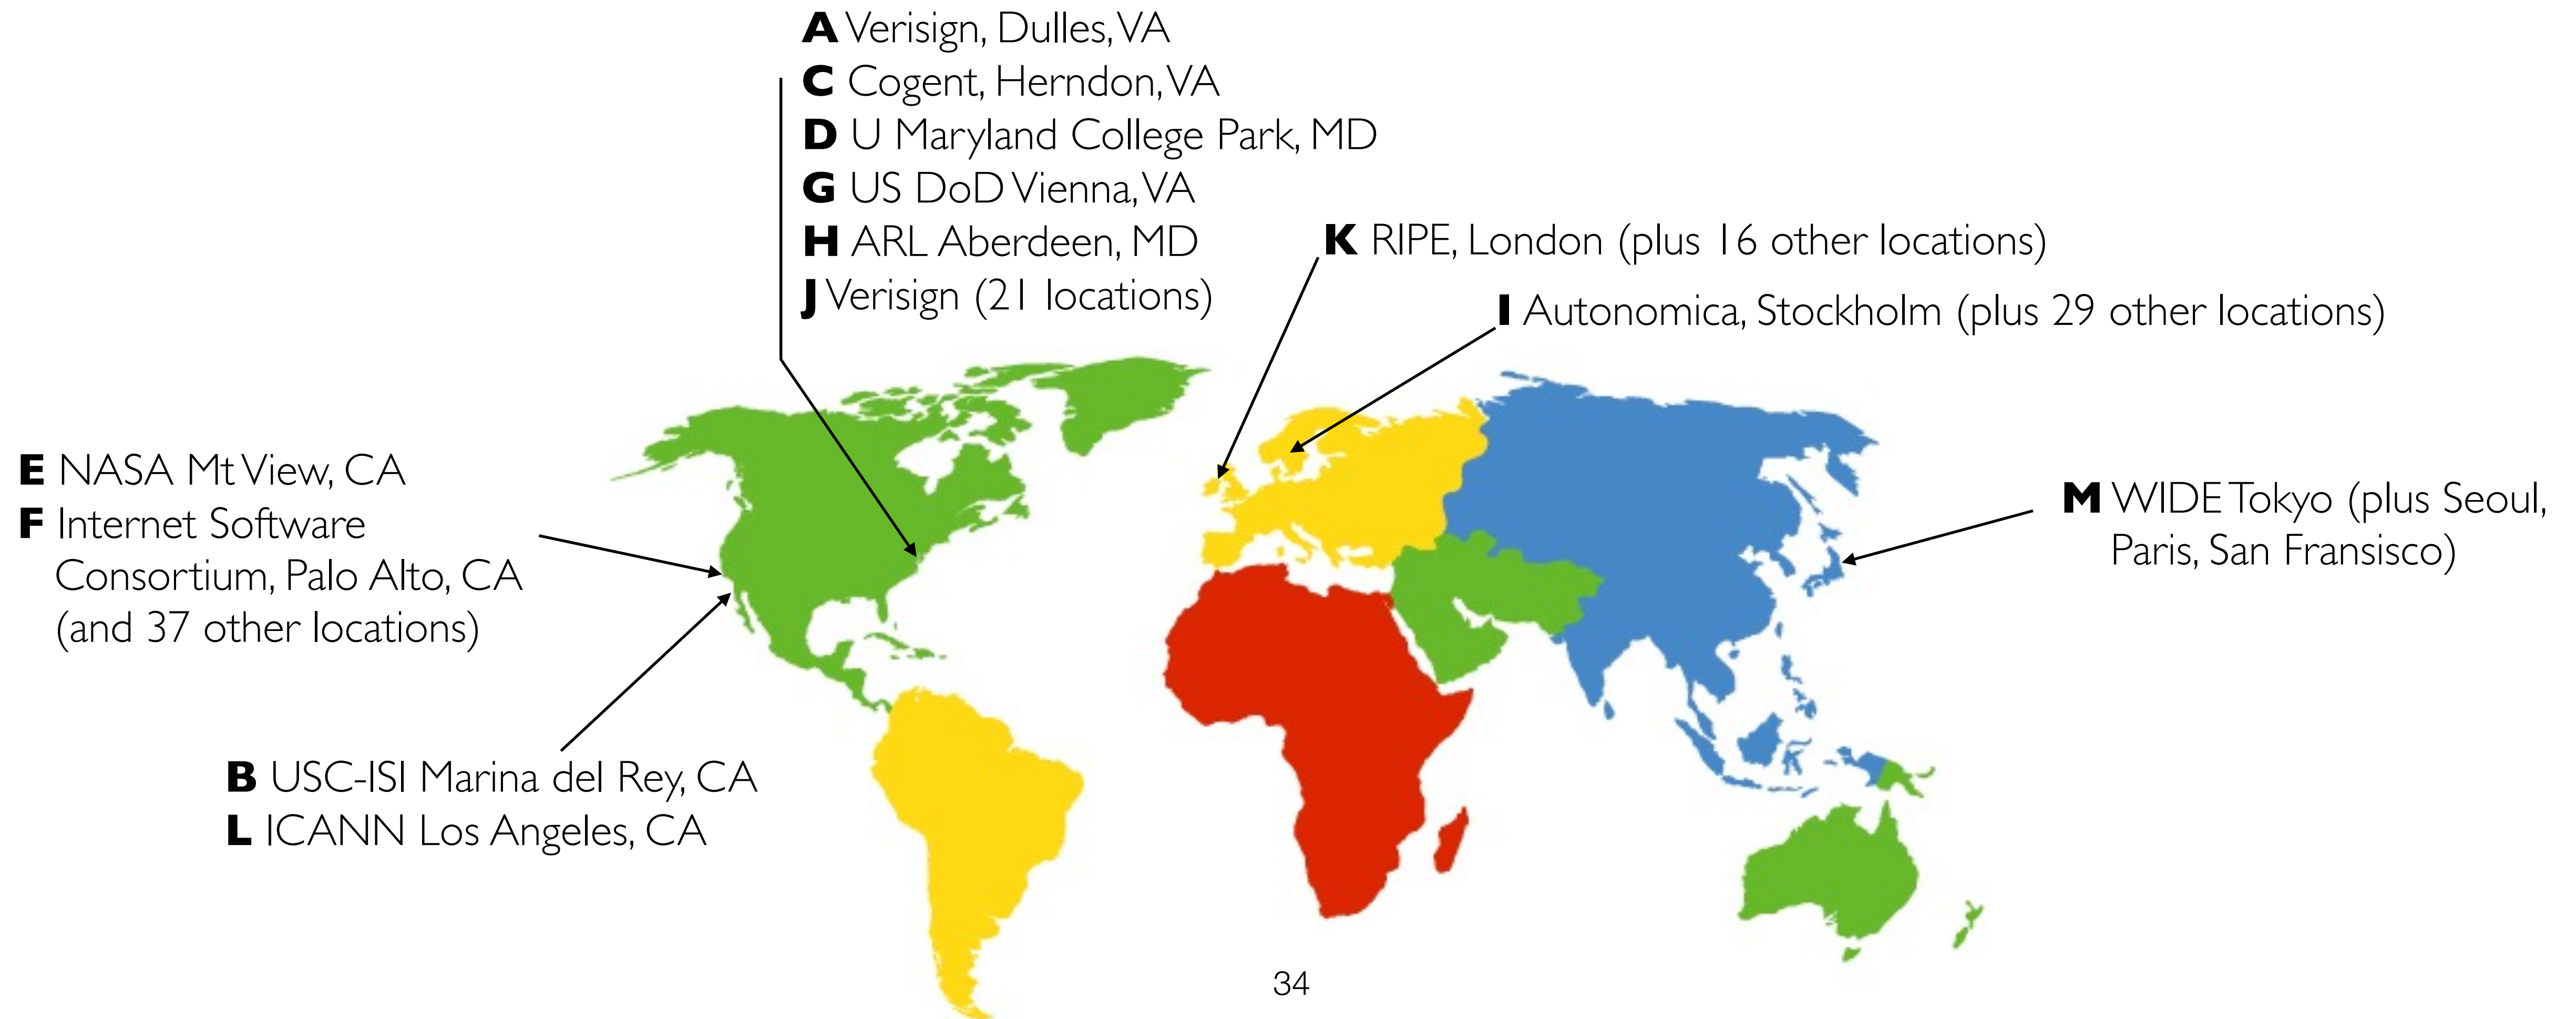
# DNS Root Servers

- 13 root servers (labeled A-M; see <http://www.root-servers>)



# DNS Root Servers

- 13 root servers (labeled A-M; see <http://www.root-servers>)
- Replicated via **any-casting**



# Anycast in a nutshell

# Anycast in a nutshell

- Routing finds shortest paths to destination

# Anycast in a nutshell

- Routing finds shortest paths to destination
- What happens if multiple machines advertise the same address?



# Anycast in a nutshell

- Routing finds shortest paths to destination
- What happens if multiple machines advertise the same address?
- The network will deliver the packet to the closest machine with that address

# Anycast in a nutshell

- Routing finds shortest paths to destination
- What happens if multiple machines advertise the same address?
- The network will deliver the packet to the closest machine with that address
- This is called “anycast”
  - Very robust
  - Requires no modification to routing algorithms

# DNS Records

# DNS Records

- DNS Servers store **resource records (RRs)**
  - RR is (name, value, type, TTL)

# DNS Records

- DNS Servers store **resource records (RRs)**
  - RR is (name, value, type, TTL)
- Type = A: ( $\rightarrow$  Address)
  - Name = hostname
  - Value = IP address

# DNS Records

- DNS Servers store **resource records (RRs)**
  - RR is (name, value, type, TTL)
- Type = A: ( $\rightarrow$  Address)
  - Name = hostname
  - Value = IP address
- Type = NS: ( $\rightarrow$  Name Server)
  - Name = domain
  - Value = name of DNS server for domain

# DNS Records

- DNS Servers store **resource records (RRs)**
  - RR is (name, value, type, TTL)
- Type = A: ( $\rightarrow$  Address)
  - Name = hostname
  - Value = IP address
- Type = NS: ( $\rightarrow$  Name Server)
  - Name = domain
  - Value = name of DNS server for domain
- Type = MX: ( $\rightarrow$  Mail eXchanger)
  - Name = domain in email address
  - Value = name(s) of mail server(s)

# Inserting Resource Records into DNS



# Inserting Resource Records into DNS

- Example: you just created company “FooBar”

# Inserting Resource Records into DNS

- Example: you just created company “FooBar”
- You get a block of IP addresses from your ISP
  - Say 212.44.9.128/25

# Inserting Resource Records into DNS

- Example: you just created company “FooBar”
- You get a block of IP addresses from your ISP
  - Say 212.44.9.128/25
- Register foobar.com at registrar (e.g., GoDaddy)
  - Provide registrar with names and IP addresses of your authoritative name server(s)
  - Registrar inserts RR pairs into the .com TLD server
    - (foobar.com, dns1.foobar.com, NS)
    - (dns1.foobar.com, 212.44.9.129, A)

# Inserting Resource Records into DNS

- Example: you just created company “FooBar”
- You get a block of IP addresses from your ISP
  - Say 212.44.9.128/25
- Register foobar.com at registrar (e.g., GoDaddy)
  - Provide registrar with names and IP addresses of your authoritative name server(s)
  - Registrar inserts RR pairs into the .com TLD server
    - (foobar.com, dns1.foobar.com, NS)
    - (dns1.foobar.com, 212.44.9.129, A)
- Store resource records in your server dns1.foobar.com
  - e.g., type A records: (foobar.com, 212.44.9.130, A), (social.foobar.com, 212.44.9.131, A), etc.
  - e.g., type MX records for foobar.com

# Using DNS (Client/Application view)

# Using DNS (Client/Application view)

- Two components
  - Local DNS server
  - Resolver software on hosts

# Using DNS (Client/Application view)

- **Two components**
  - Local DNS server
  - Resolver software on hosts
- **Local DNS server (“default name server”)**
  - Clients configured with the default server’s address or learn it via a host configuration protocol (e.g., DHCP — future lecture)

# Using DNS (Client/Application view)



- **Two components**
  - Local DNS server
  - Resolver software on hosts
- **Local DNS server (“default name server”)**
  - Clients configured with the default server’s address or learn it via a host configuration protocol (e.g., DHCP — future lecture)
- **Client application:**
  - Obtain DNS name (e.g., from URL)
  - Triggers DNS request to its local DNS server




**Root  
Servers**

**.edu  
Servers**

**nyu.edu  
Servers**






**Root  
Servers**

**.edu  
Servers**

**nyu.edu  
Servers**




**Root  
Servers**

**.edu  
Servers**

**nyu.edu  
Servers**

**Local DNS server**  
**(mydns.yale.edu)**




**Root  
Servers**

**.edu  
Servers**

**nyu.edu  
Servers**


**Local DNS server**  
**(mydns.yale.edu)**

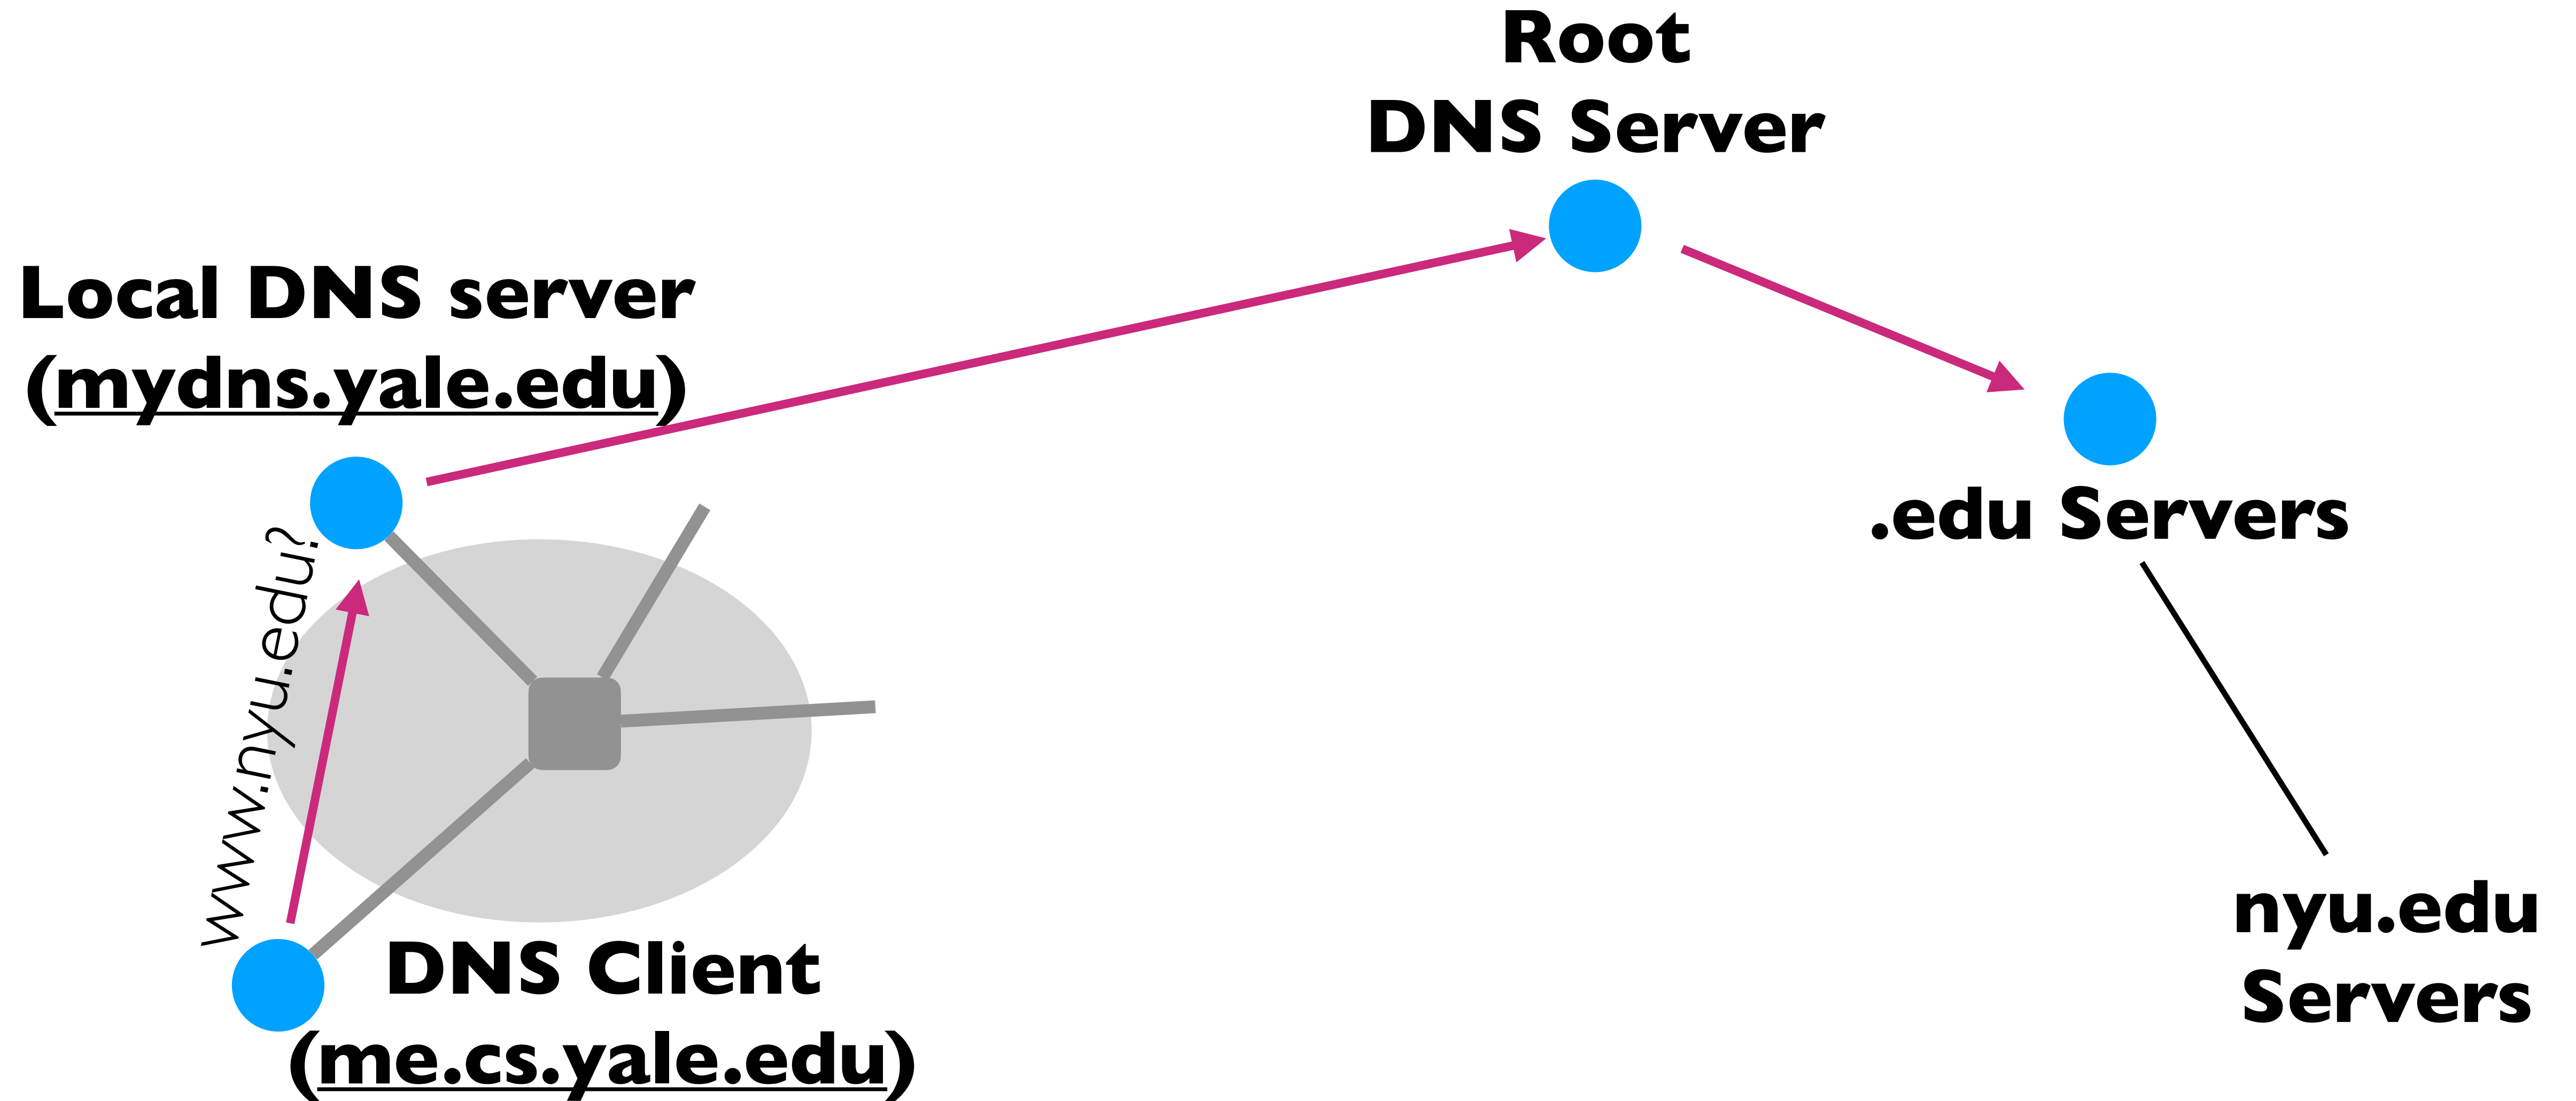



**Root  
Servers**

**.edu  
Servers**

**nyu.edu  
Servers**












## Recursive DNS Query




**Root  
DNS Server**




**.edu Servers**





**nyu.edu Servers**




**Local DNS server  
(mydns.yale.edu)**

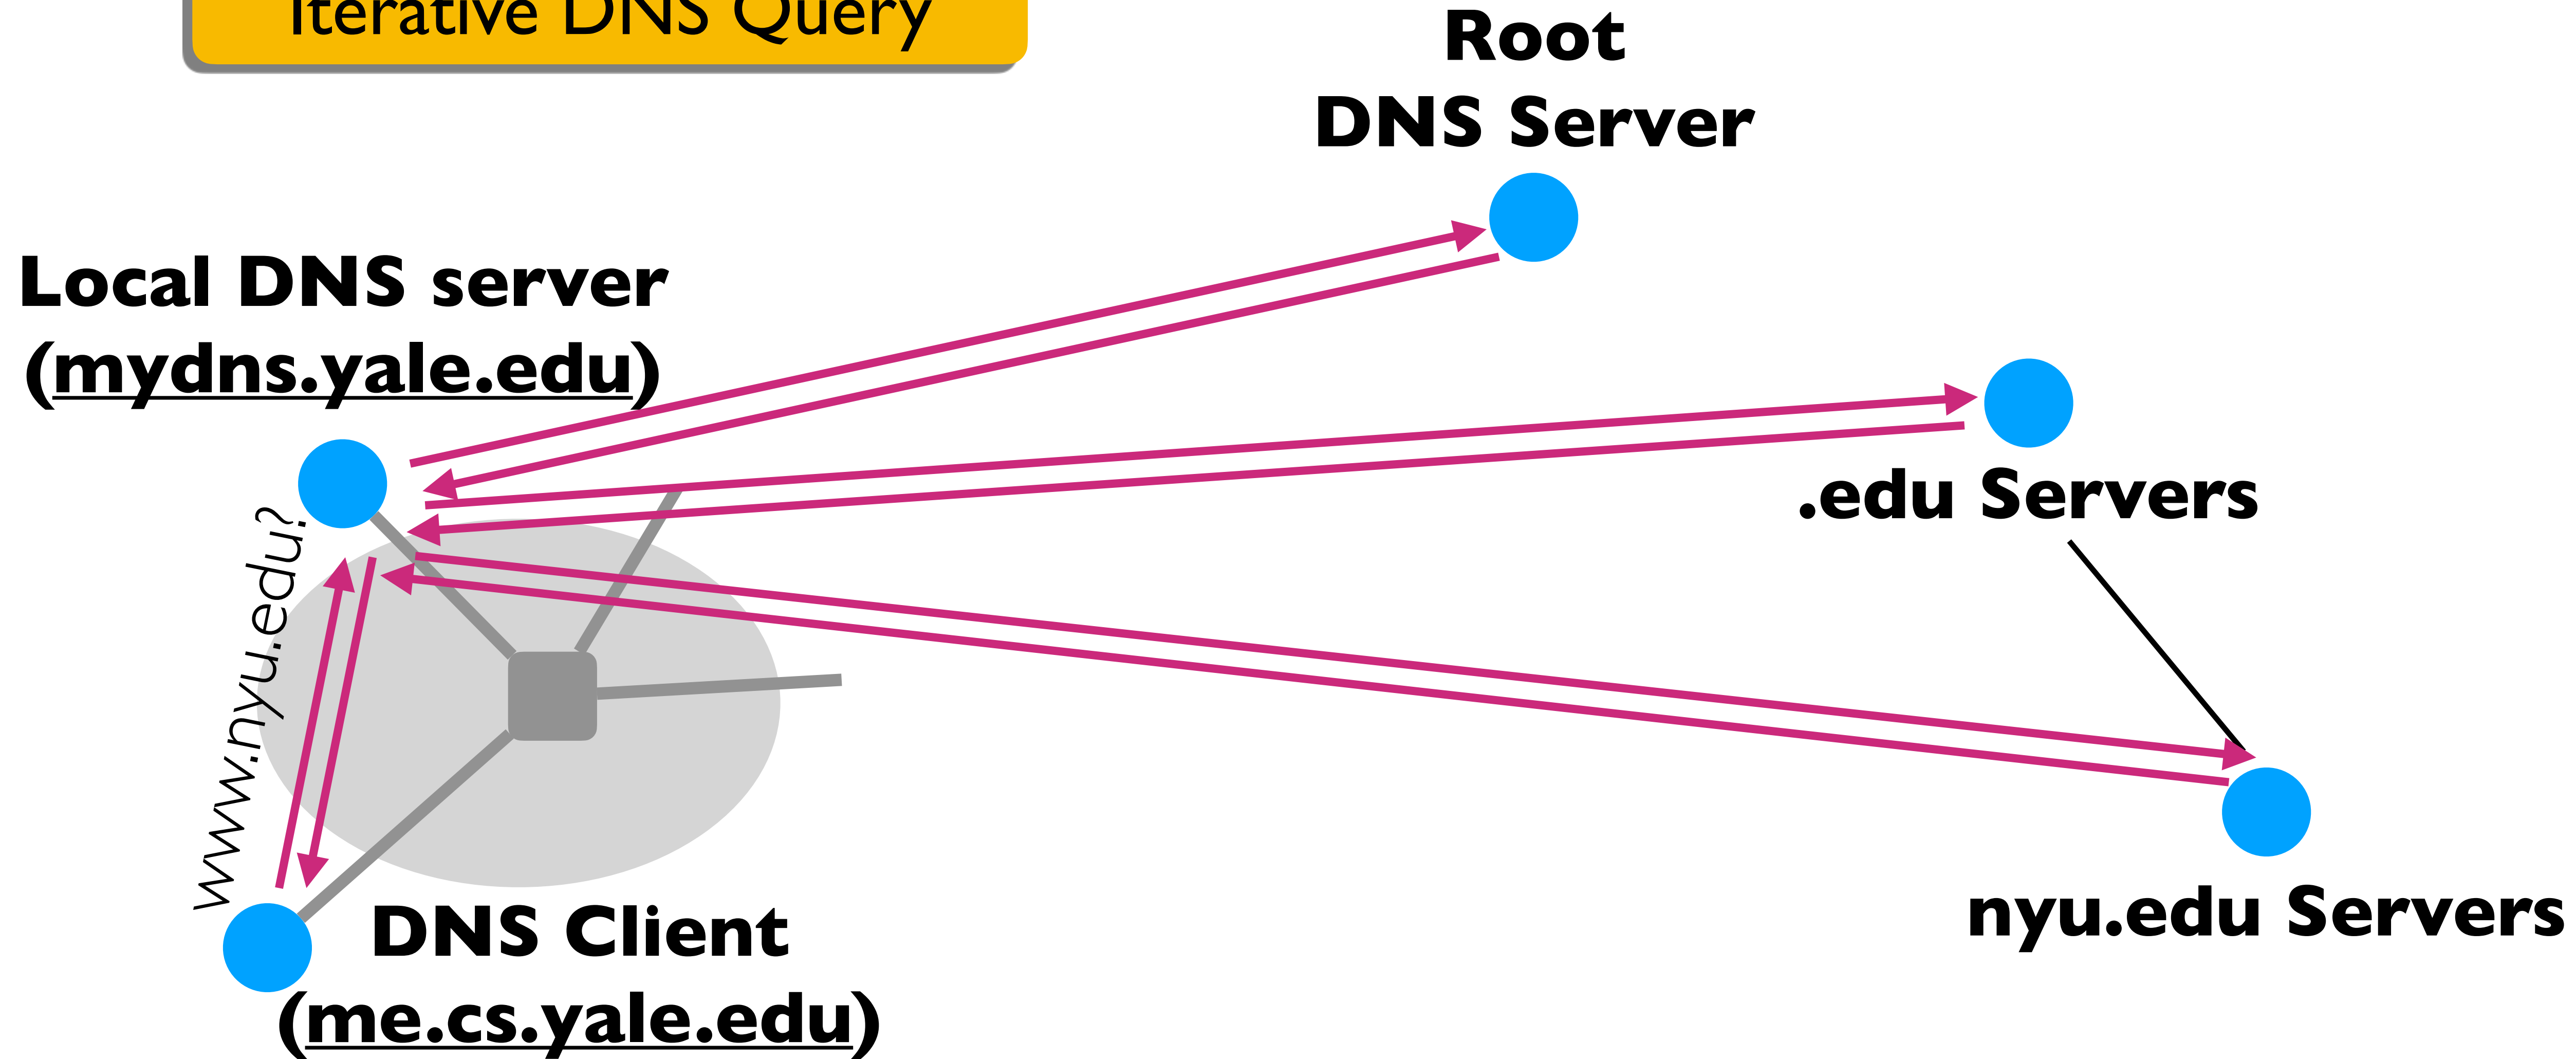








## Iterative DNS Query



# DNS Protocol

# DNS Protocol

- Query and Reply messages; both with the same message format
  - *See text for details*



# DNS Protocol

- Query and Reply messages; both with the same message format
  - *See text for details*
- Client-Server Interaction on UDP Port 53
  - Spec. supports TCP too, but not always implemented

# Goals: How are we doing?

- Scalable
  - Many names
  - Many updates
  - Many users creating names
  - Many users looking up names
- Highly available
- Correct
  - No naming conflicts (uniqueness)
  - Consistency
- Lookups are fast

# Per-domain Availability

# Per-domain Availability


- DNS servers are **replicated**
  - Primary and secondary name servers required
  - Name service available if at least one replica is up
  - Queries can be load-balanced between replicas

# Per-domain Availability

- DNS servers are **replicated**
  - Primary and secondary name servers required
  - Name service available if at least one replica is up
  - Queries can be load-balanced between replicas
- Try alternate servers on timeout
  - *Exponential backoff* when retrying same server

# Goals: How are we doing?

- Scalable
  - Many names
  - Many updates
  - Many users creating names
  - Many users looking up names
- Highly available
- Correct
  - No naming conflicts (uniqueness)
  - Consistency
- Lookups are fast



# Caching



# Caching

- Caching of DNS responses at all levels

# Caching

- Caching of DNS responses at all levels
- Reduces load at all levels

# Caching

- Caching of DNS responses at all levels
- Reduces load at all levels
- Reduces delay experienced by DNS client

# DNS Caching

# DNS Caching

- **How DNS caching works**
  - DNS servers cache responses to queries
  - Responses include a “time-to-live” (TTL) field
  - Server deletes cached entry after TTL expires

# DNS Caching

- **How DNS caching works**
  - DNS servers cache responses to queries
  - Responses include a “time-to-live” (TTL) field
  - Server deletes cached entry after TTL expires
- **Why caching is effective**
  - The top-level servers very rarely change
  - Popular sites visited often → local DNS server often has the information cached

# Negative Caching

# Negative Caching

- Remember things that don't work
  - Misspellings like [www.cnn.comm](#) and [www.cnnn.com](#)
  - These can take a long time to fail the first time
  - Good to remember that they don't work
  - ... so the failure takes less time the next time around




# Negative Caching

- Remember things that don't work
  - Misspellings like [www.cnn.comm](#) and [www.cnnn.com](#)
  - These can take a long time to fail the first time
  - Good to remember that they don't work
  - ... so the failure takes less time the next time around
- Negative caching is optional


Time to put your malicious hats on...

How can one attach DNS?


**Local DNS server**  
**(mydns.yale.edu)**




**Root**  
**DNS Server**




**.edu Servers**




**nyu.edu Servers**




**Local DNS server**  
**(mydns.yale.edu)**




**Root**  
**DNS Server**



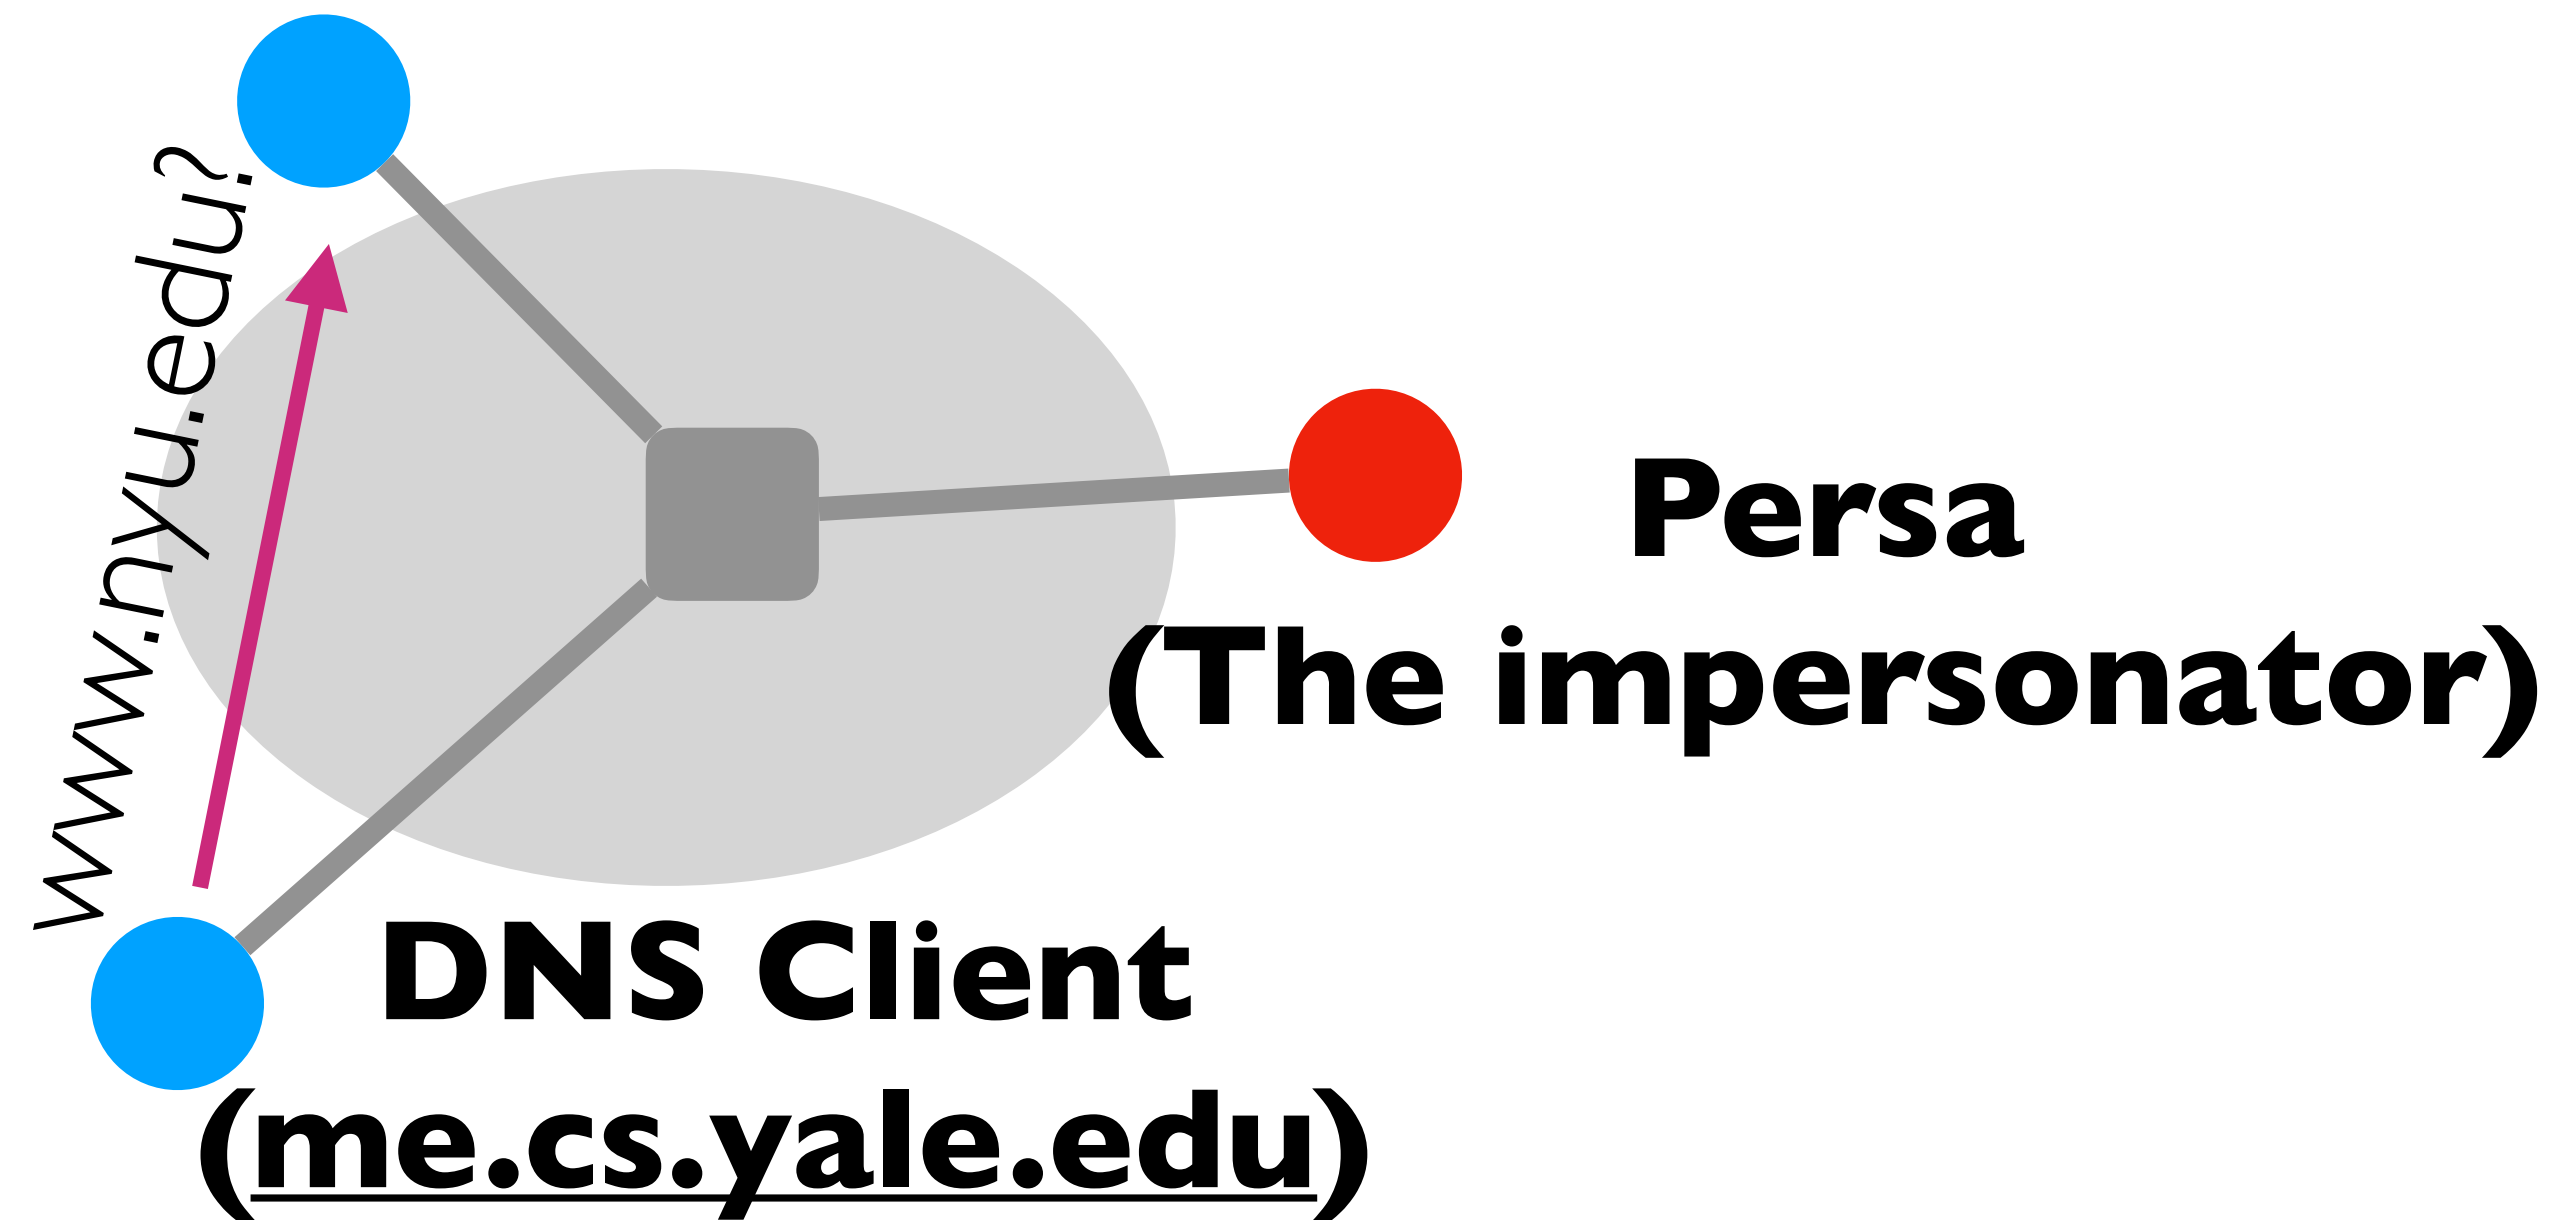
**.edu Servers**




**nyu.edu Servers**



**Local DNS server**  
**(mydns.yale.edu)**



**Local DNS server**  
**(mydns.yale.edu)**



**Root**  
**DNS Server**

**.edu Servers**

**nyu.edu Servers**


# How Can One Attack DNS?

# How Can One Attack DNS?


- Impersonate the local DNS server
  - Give the wrong IP address to the DNS client




**Local DNS server**  
**(mydns.yale.edu)**




**Root**  
**DNS Server**

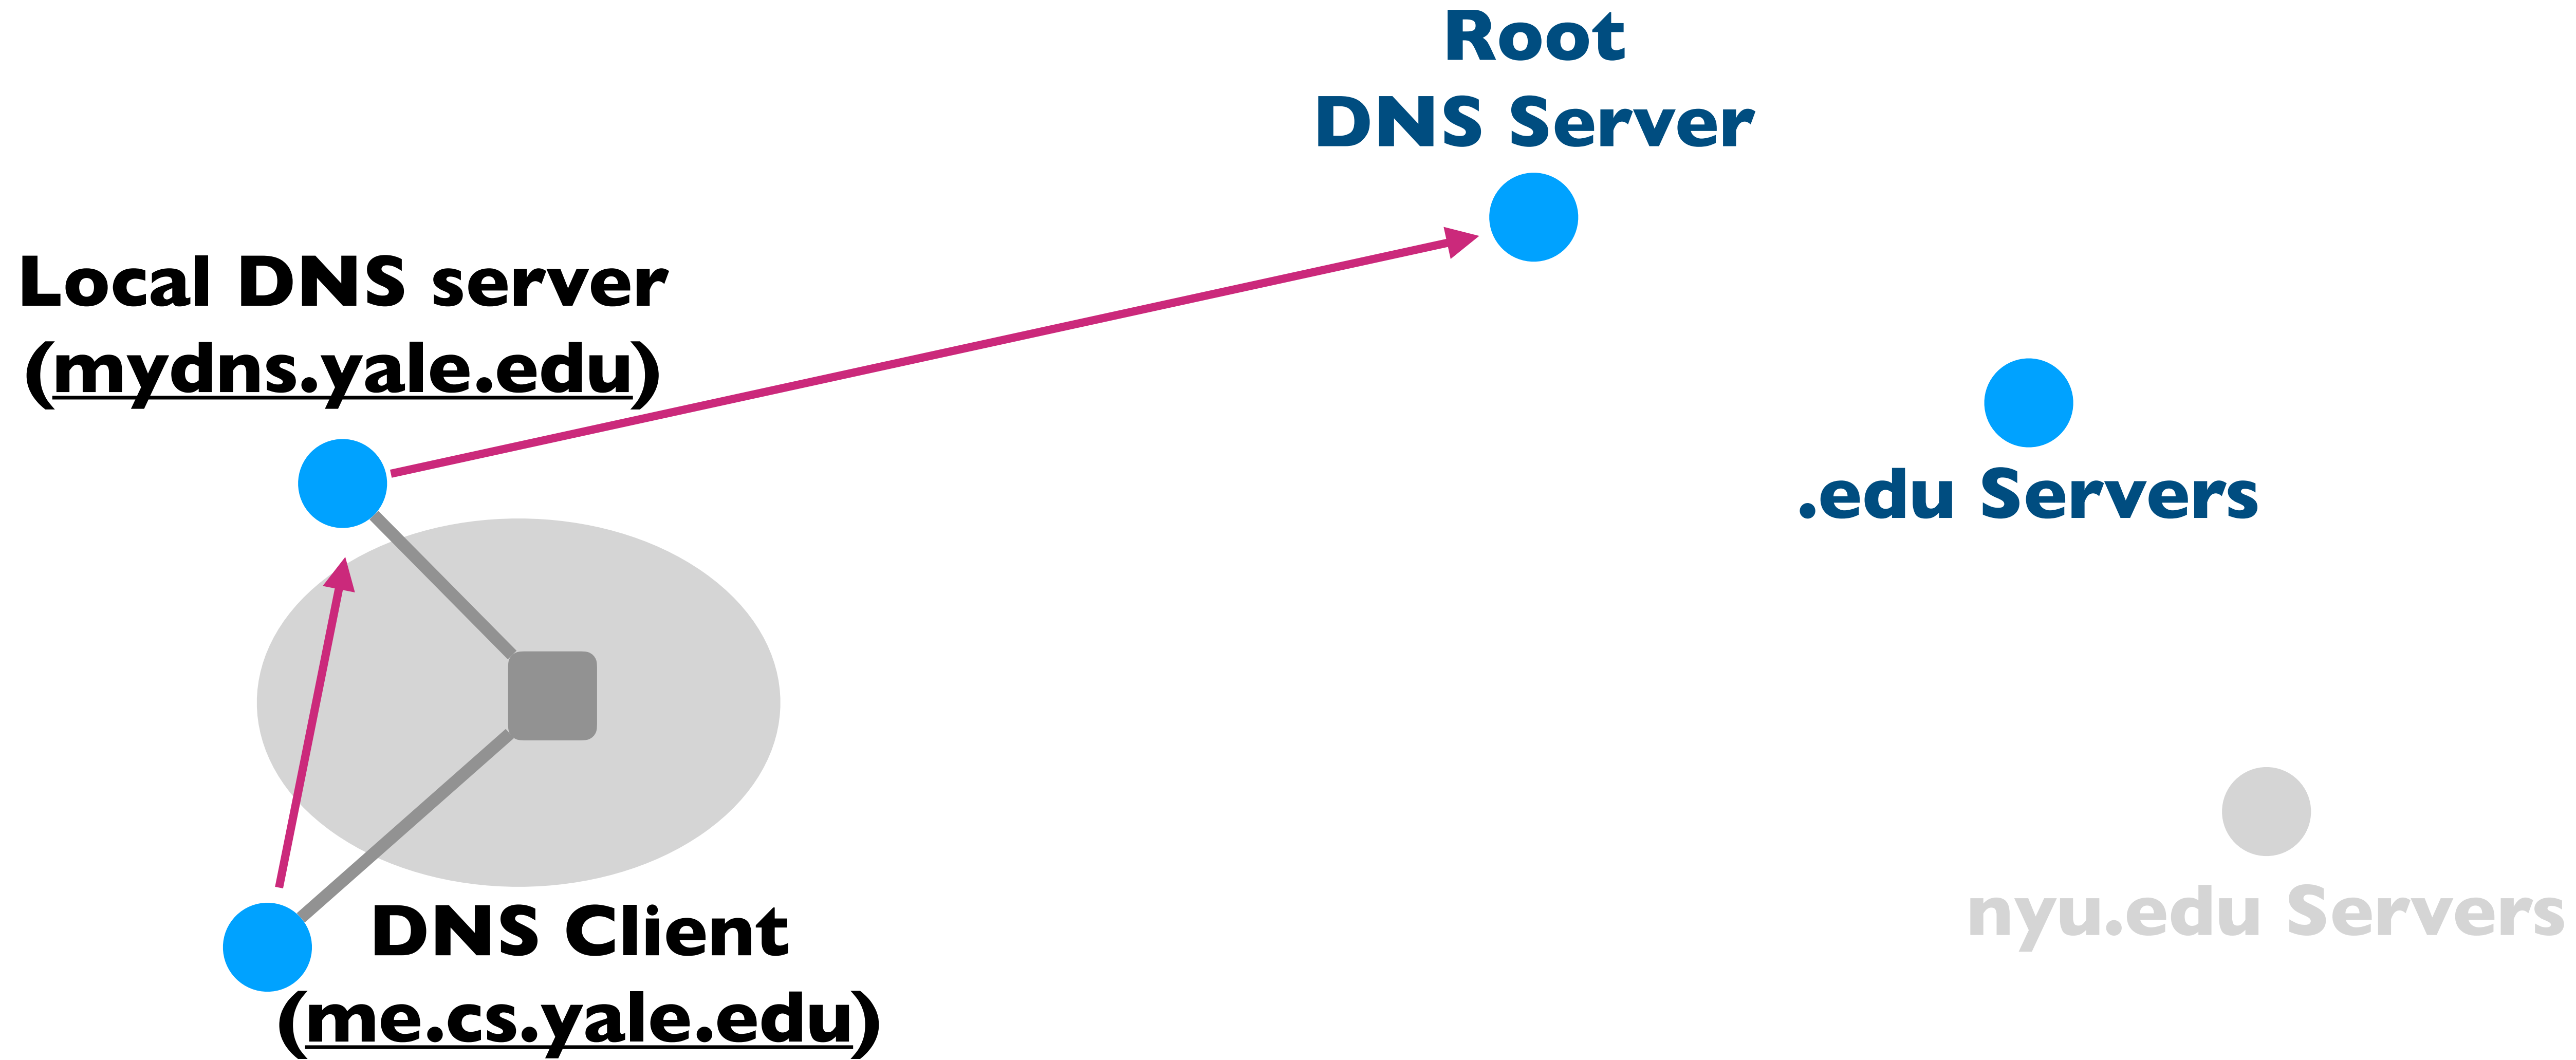


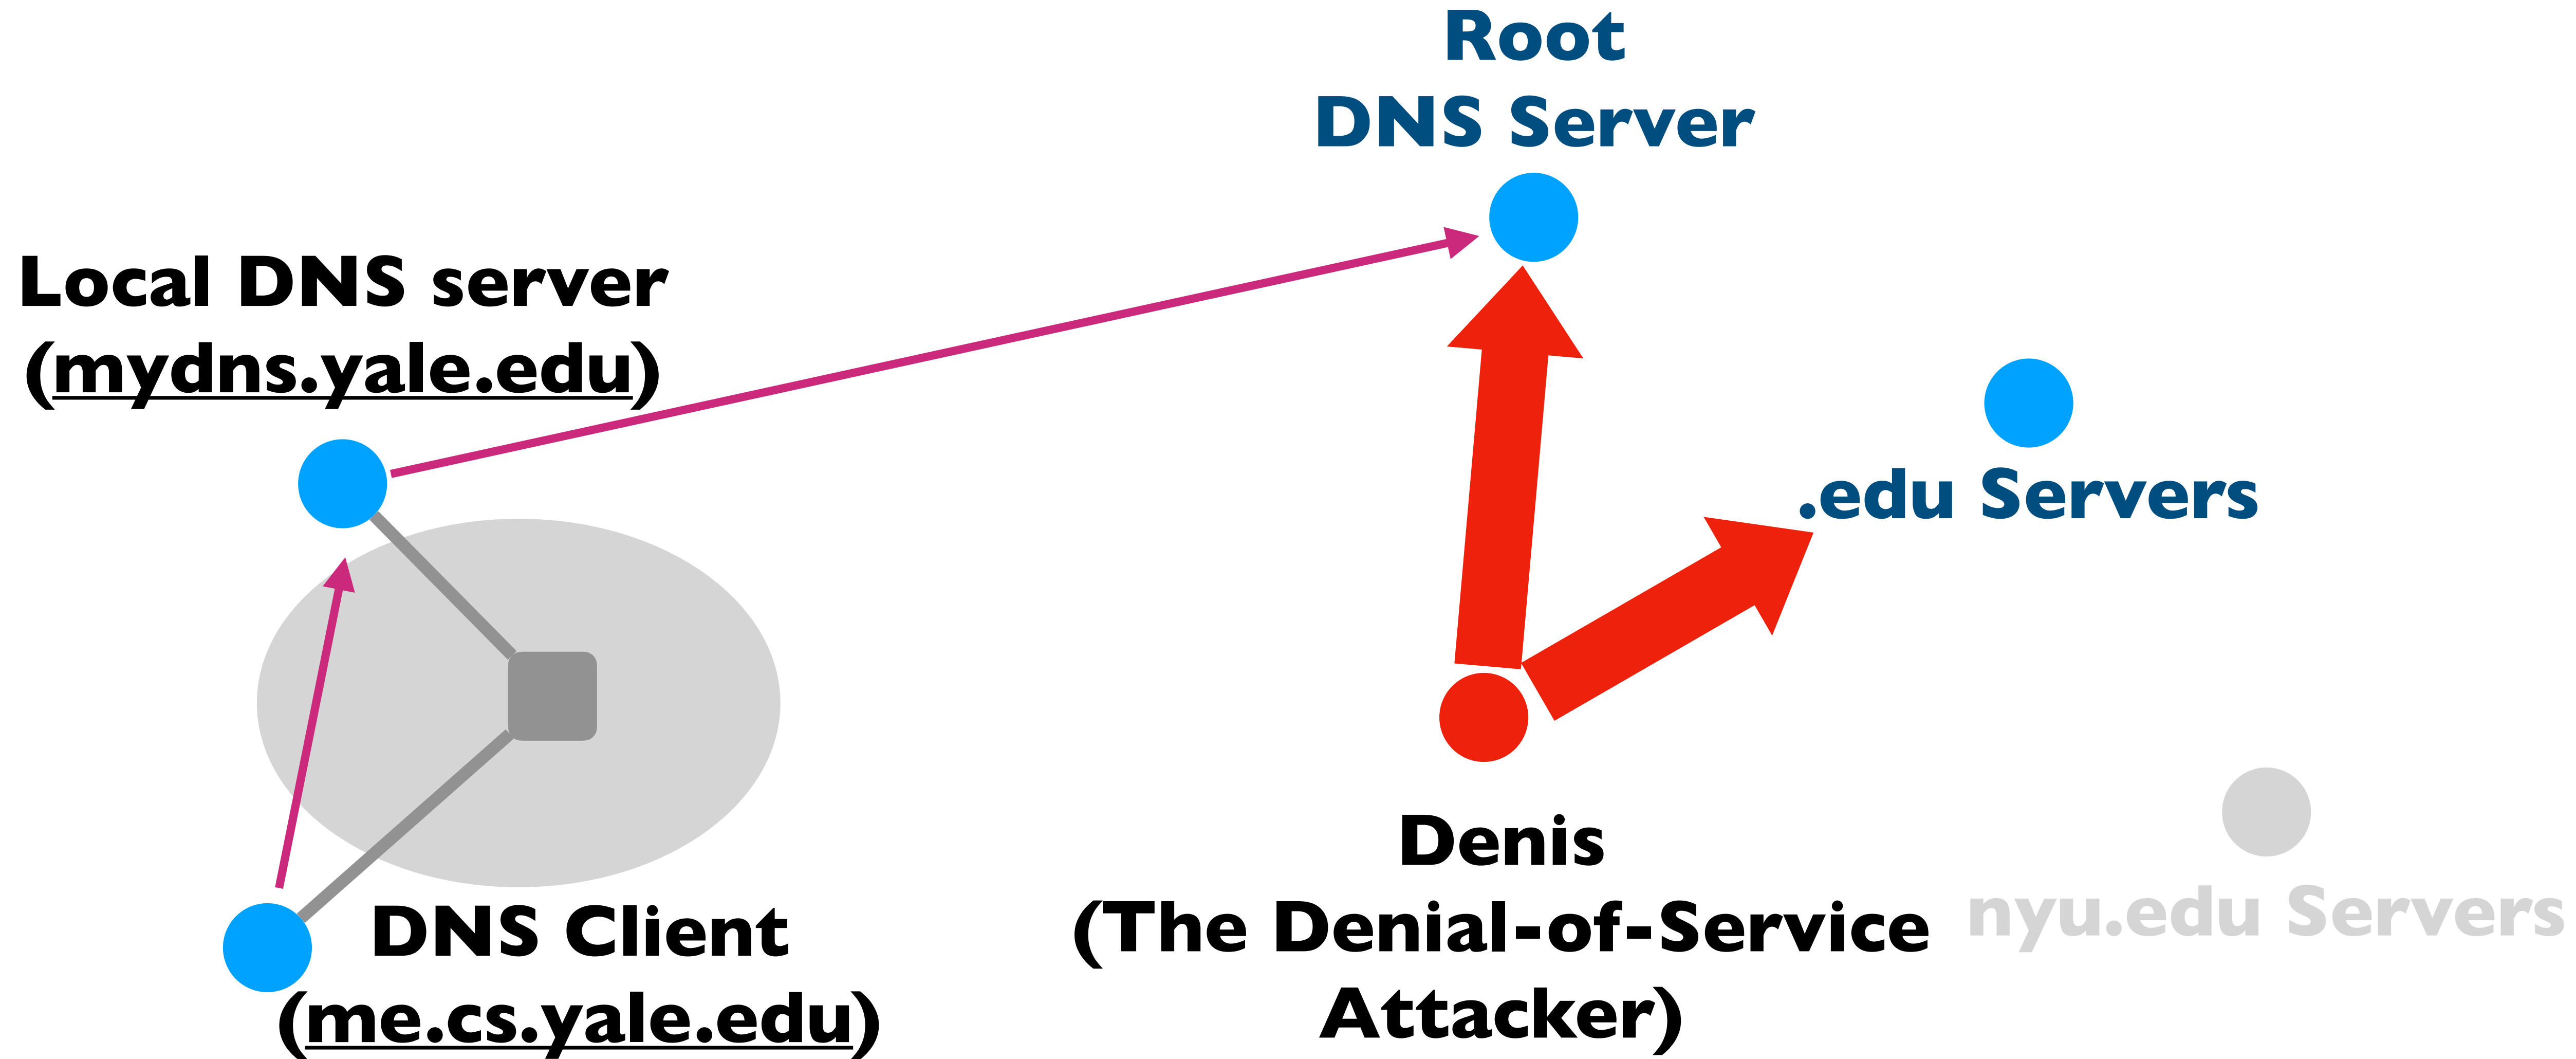
**.edu Servers**



**nyu.edu Servers**









# How Can One Attack DNS?

- Impersonate the local DNS server
  - Give the wrong IP address to the DNS client
- Denial-of-service the root or TLD servers
  - Make them unavailable to the rest of the world


**Local DNS server**  
**(mydns.yale.edu)**




**Root**  
**DNS Server**




**.edu Servers**




**nyu.edu Servers**



**Local DNS server**  
**(mydns.yale.edu)**




**Root**  
**DNS Server**




**.edu Servers**



**nyu.edu Servers**





# How Can One Attack DNS?

- **Impersonate the local DNS server**
  - Give the wrong IP address to the DNS client
- **Denial-of-service the root or TLD servers**
  - Make them unavailable to the rest of the world
- **Poison the cache of a DNS server**
  - Increase the delay experienced by DNS clients



# Important Properties of DNS

# Important Properties of DNS

Administrative delegation and hierarchy results in:

# Important Properties of DNS

Administrative delegation and hierarchy results in:

- Easy unique naming

# Important Properties of DNS

Administrative delegation and hierarchy results in:

- Easy unique naming
- “Fate sharing” for network failures

# Important Properties of DNS

Administrative delegation and hierarchy results in:

- Easy unique naming
- “Fate sharing” for network failures
- Reasonable trust model

# Important Properties of DNS

Administrative delegation and hierarchy results in:

- Easy unique naming
- “Fate sharing” for network failures
- Reasonable trust model
- Caching lends scalability, performance

# DNS Provides Indirection

# DNS Provides Indirection

- Addresses can **change** underneath
  - Move [www.cnn.com](http://www.cnn.com) to a new IP address
  - Humans/applications are unaffected



# DNS Provides Indirection

- Addresses can **change** underneath
  - Move [www.cnn.com](http://www.cnn.com) to a new IP address
  - Humans/applications are unaffected
- Name could map to **multiple** IP addresses
  - Enables load-balancing

# DNS Provides Indirection

- Addresses can **change** underneath
  - Move [www.cnn.com](http://www.cnn.com) to a new IP address
  - Humans/applications are unaffected
- Name could map to **multiple** IP addresses
  - Enables load-balancing
- Multiple names for the same addresses
  - E.g., many services (mail, www, ftp) on same machine

# DNS Provides Indirection

- Addresses can **change** underneath
  - Move [www.cnn.com](http://www.cnn.com) to a new IP address
  - Humans/applications are unaffected
- Name could map to **multiple** IP addresses
  - Enables load-balancing
- Multiple names for the same addresses
  - E.g., many services (mail, www, ftp) on same machine
- Allowing “host” names to evolve into “service” names