

Application Layer: HTTP (Contd.)

CPSC 433/533, Spring 2021

Anurag Khandelwal

Meeting HTTP Goals

- **Users**

Improve HTTP to compensate for TCP weak-spots

1

- Fast downloads (not the same as low-latency communication!)
- High availability

- **Content provider**

Caching & Replication

2

- Happy users (hence, above)
- Cost-effective delivery infrastructure

- **Network (secondary)**

Exploit economies of scale
(Webhosting, CDNs, datacenters)

3

- Avoid overload

Improving HTTP Performance:

Content Distribution Networks

- Caching and replication as a service
- Large-scale distributed storage infrastructure (usually) administrated by one entity
 - e.g., Akamai has servers in 20,000+ locations
- Combination of (pull) caching and (push) replication
 - **Pull:** Direct result of client requests
 - **Push:** Expectation of high access rate
- Also do some processing
 - Handle *dynamic* webpages
 - *Transcoding*

Improving HTTP Performance:

CDN Example - Akamai

- Akamai creates new domain names for each of their clients
 - e.g., *a128.g.akamai.net* for *cnn.com*
- The CDN's DNS servers are authoritative for the new domains
- The client content provider modifies its content so that embedded URLs reference the new domains
 - “Akamaize” content
 - e.g., *http://www.cnn.com/image-of-the-day.gif* becomes *http://a128.g.akamai.net/image-of-the-day.gif*
- Requests now sent to CDN's infrastructure...

Cost-effective Content Delivery

- **General theme:** multiple sites hosted on shared physical infrastructure
 - Efficiency of statistical multiplexing
 - Economies of scale (volume pricing, etc.)
 - Amortization of human operator costs
- **Examples:**
 - Web hosting companies
 - CDNs
 - Cloud infrastructure

Questions?

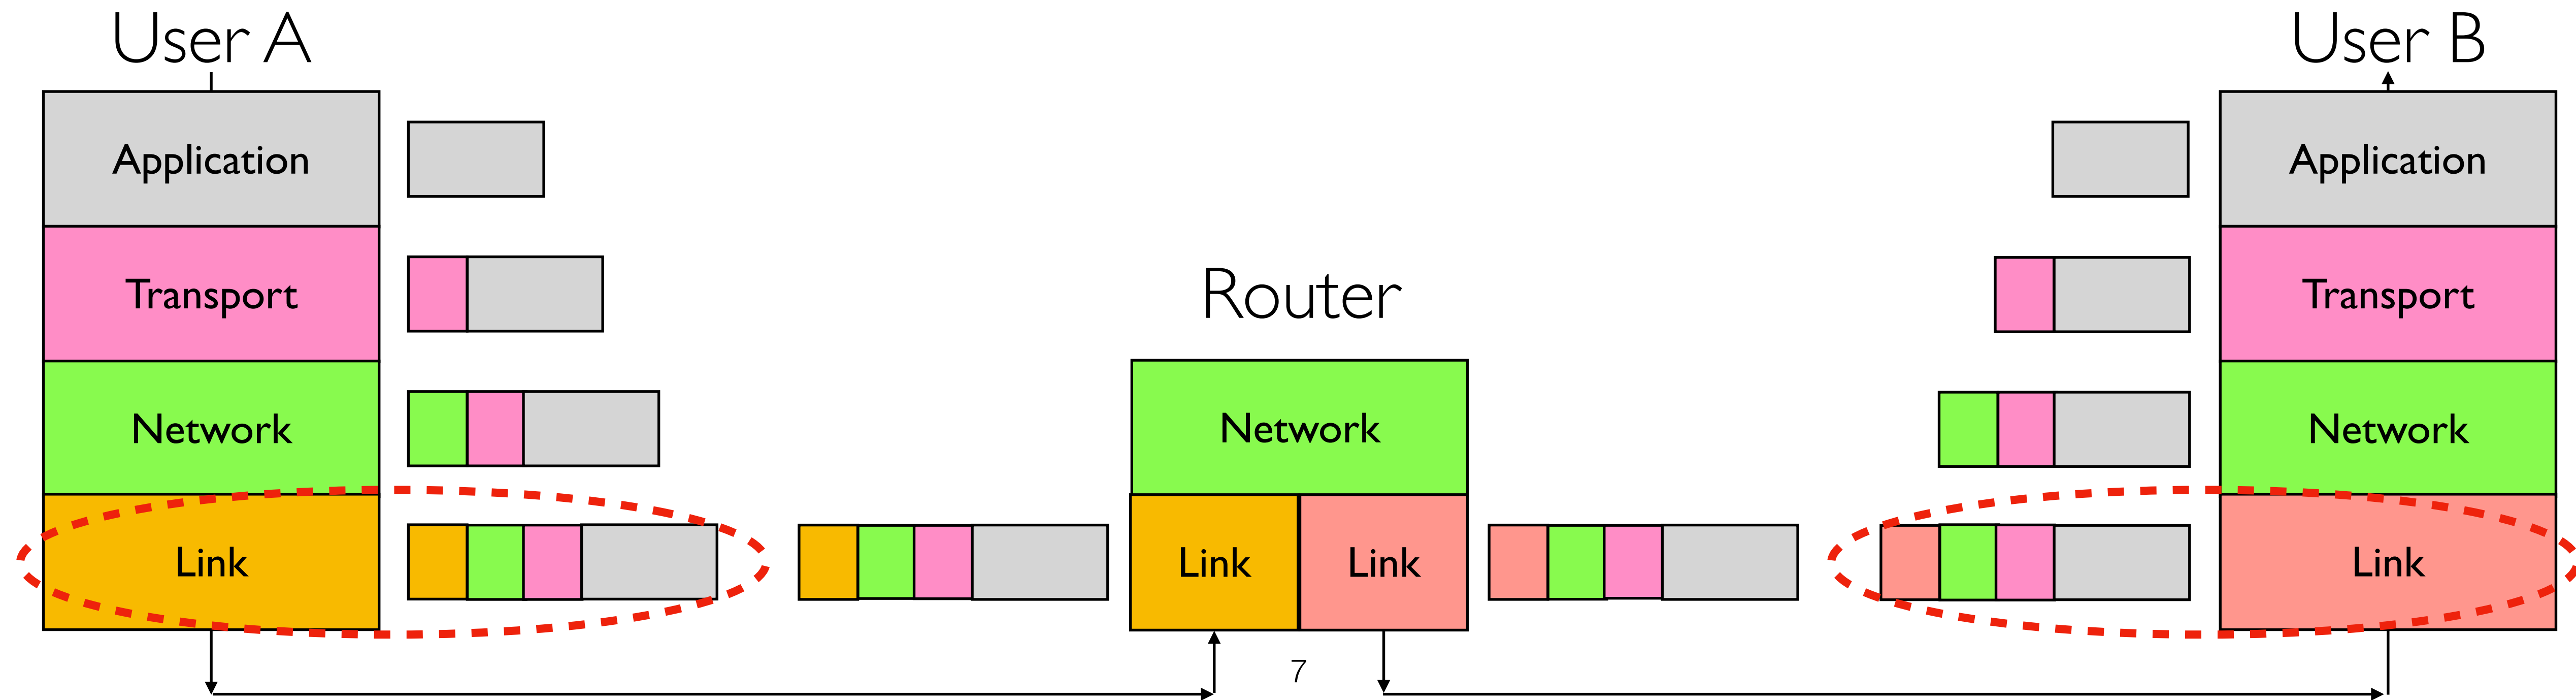
Taking Stock

Course so far:

- **Concepts, Overall Architecture**
- **Network Layer**, *Best-effort global delivery of packets*
- **Transport Layer**, *Reliable (or unreliable) delivery of data*
- **Application Layer**, *DNS, HTTP*

Next:

- **Link Layer**, *Ethernet*





The Link Layer: Ethernet

CPSC 433/533, Spring 2021

Anurag Khandelwal

Point-to-Point vs. Broadcast Media

- **Point-to-point:** *dedicated* pairwise communication
 - E.g., long-distance fiber link
 - E.g., point-to-point link between Ethernet switch and host
- **Broadcast:** *shared* wire or medium
 - Traditional Ethernet (pre ~2000)
 - 802.11 Wireless LAN



Context: Shared Broadcast Channel

- Must avoid having multiple nodes speaking at once
- Otherwise, collisions can lead to garbled data
- Need distributed algorithm for sharing the channel
- Algorithm determines which node can transmit


Multiple Access Algorithm

- ***Channel partitioning:*** divide channel by frequency (Frequency Division Multiplexing)
 - Can be wasteful! Only so much EM spectrum to go around, and many frequencies likely to be idle often (traffic is bursty)
- ***Taking turns:*** scheme for trading off who gets to transmit
 - Same drawbacks as FDM...
- ***Random access:*** allow collisions, and then recover
 - More in the Internet style!

Random Access MAC Protocols

- **When node has packet to send**
 - Transmit at full channel data rate
 - No a priori coordination among nodes
- **Two or more transmitted nodes → *collision***
 - Data lost
- **Random access MAC protocol specifies:**
 - How to detect collisions
 - How to recover from collisions
- **Examples**
 - ALOHA and Slotted ALOHA
 - CSMA, CSMA/CD, CSMA/CA (wireless, covered later)

Where it all started: AlohaNet



- Norm Abramson left Stanford in 1970 (*so he could surf!*)
- Setup the first data communication system for Hawaiian islands
- Central hub at U. Hawaii, Oahu


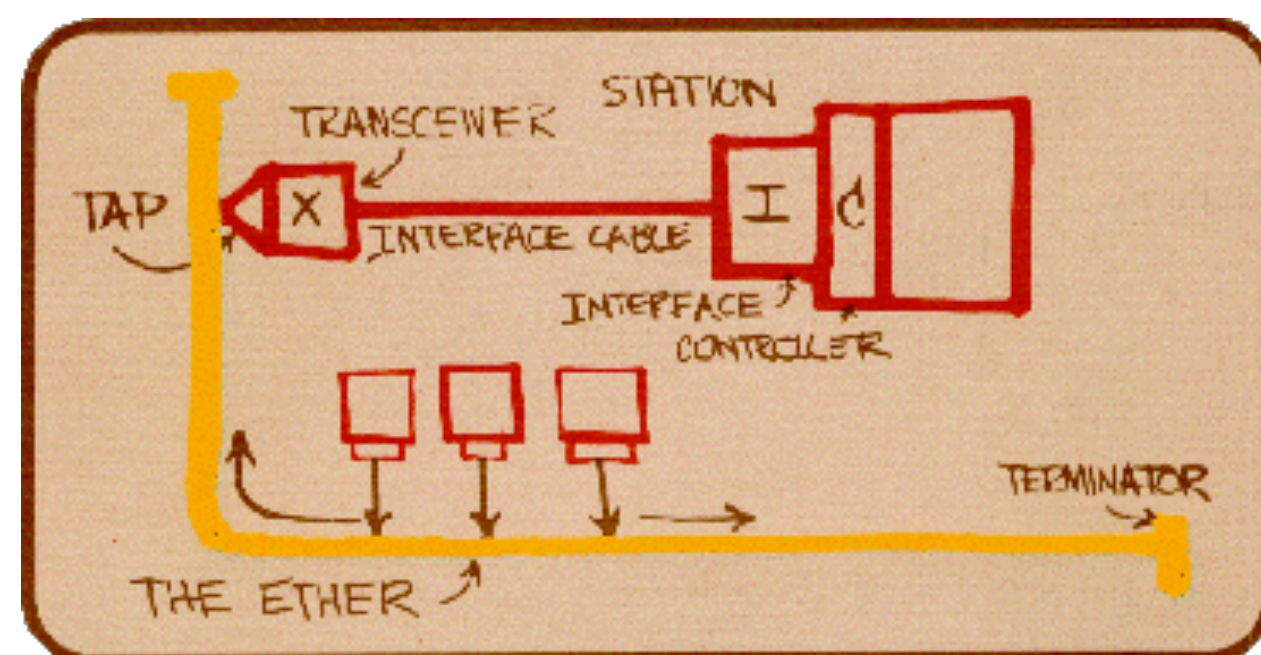
ALOHA Signaling

- Two channels: random access, broadcast
- Sites send packet to hub (random-access channel)
 - If not received (due to collision), site resends
- Hub sends packets to all sites (broadcast channel)
 - Sites can receive even if they are also sending

Ethernet



- Bob Metcalfe, Xerox PARC, visits Hawaii and gets an idea!
 - (Maybe we all need to go to Hawaii for good ideas...)
- Shared wired medium
 - coax cable



Evolution

- **Ethernet was invented as a broadcast technology**
 - Hosts share channel
 - Each packet received by all attached hosts
 - CSMA/CD for media access control
- **Current Ethernets are “switched” (later)**
 - Point-to-point links between switches; between a host and switch
 - No sharing, no CSMA/CD
 - Uses “self learning” and “spanning tree” algorithms for routing


Questions?

CSMA (Carrier Sense Multiple Access)


- CSMA: **Listen** before transmit
 - If channel sensed idle: transmit entire frame
 - If channel sensed busy: defer transmission
- Human analogy: don't interrupt others!
- Does this eliminate all collisions?
 - *No*, because of non-zero propagation delay

CSMA Collisions

- Propagation delay: two nodes may not each other before sending
- CSMA reduces but does not eliminate collisions
- Biggest remaining problem?
- Collisions still take full transmission slot!



CSMA Collisions

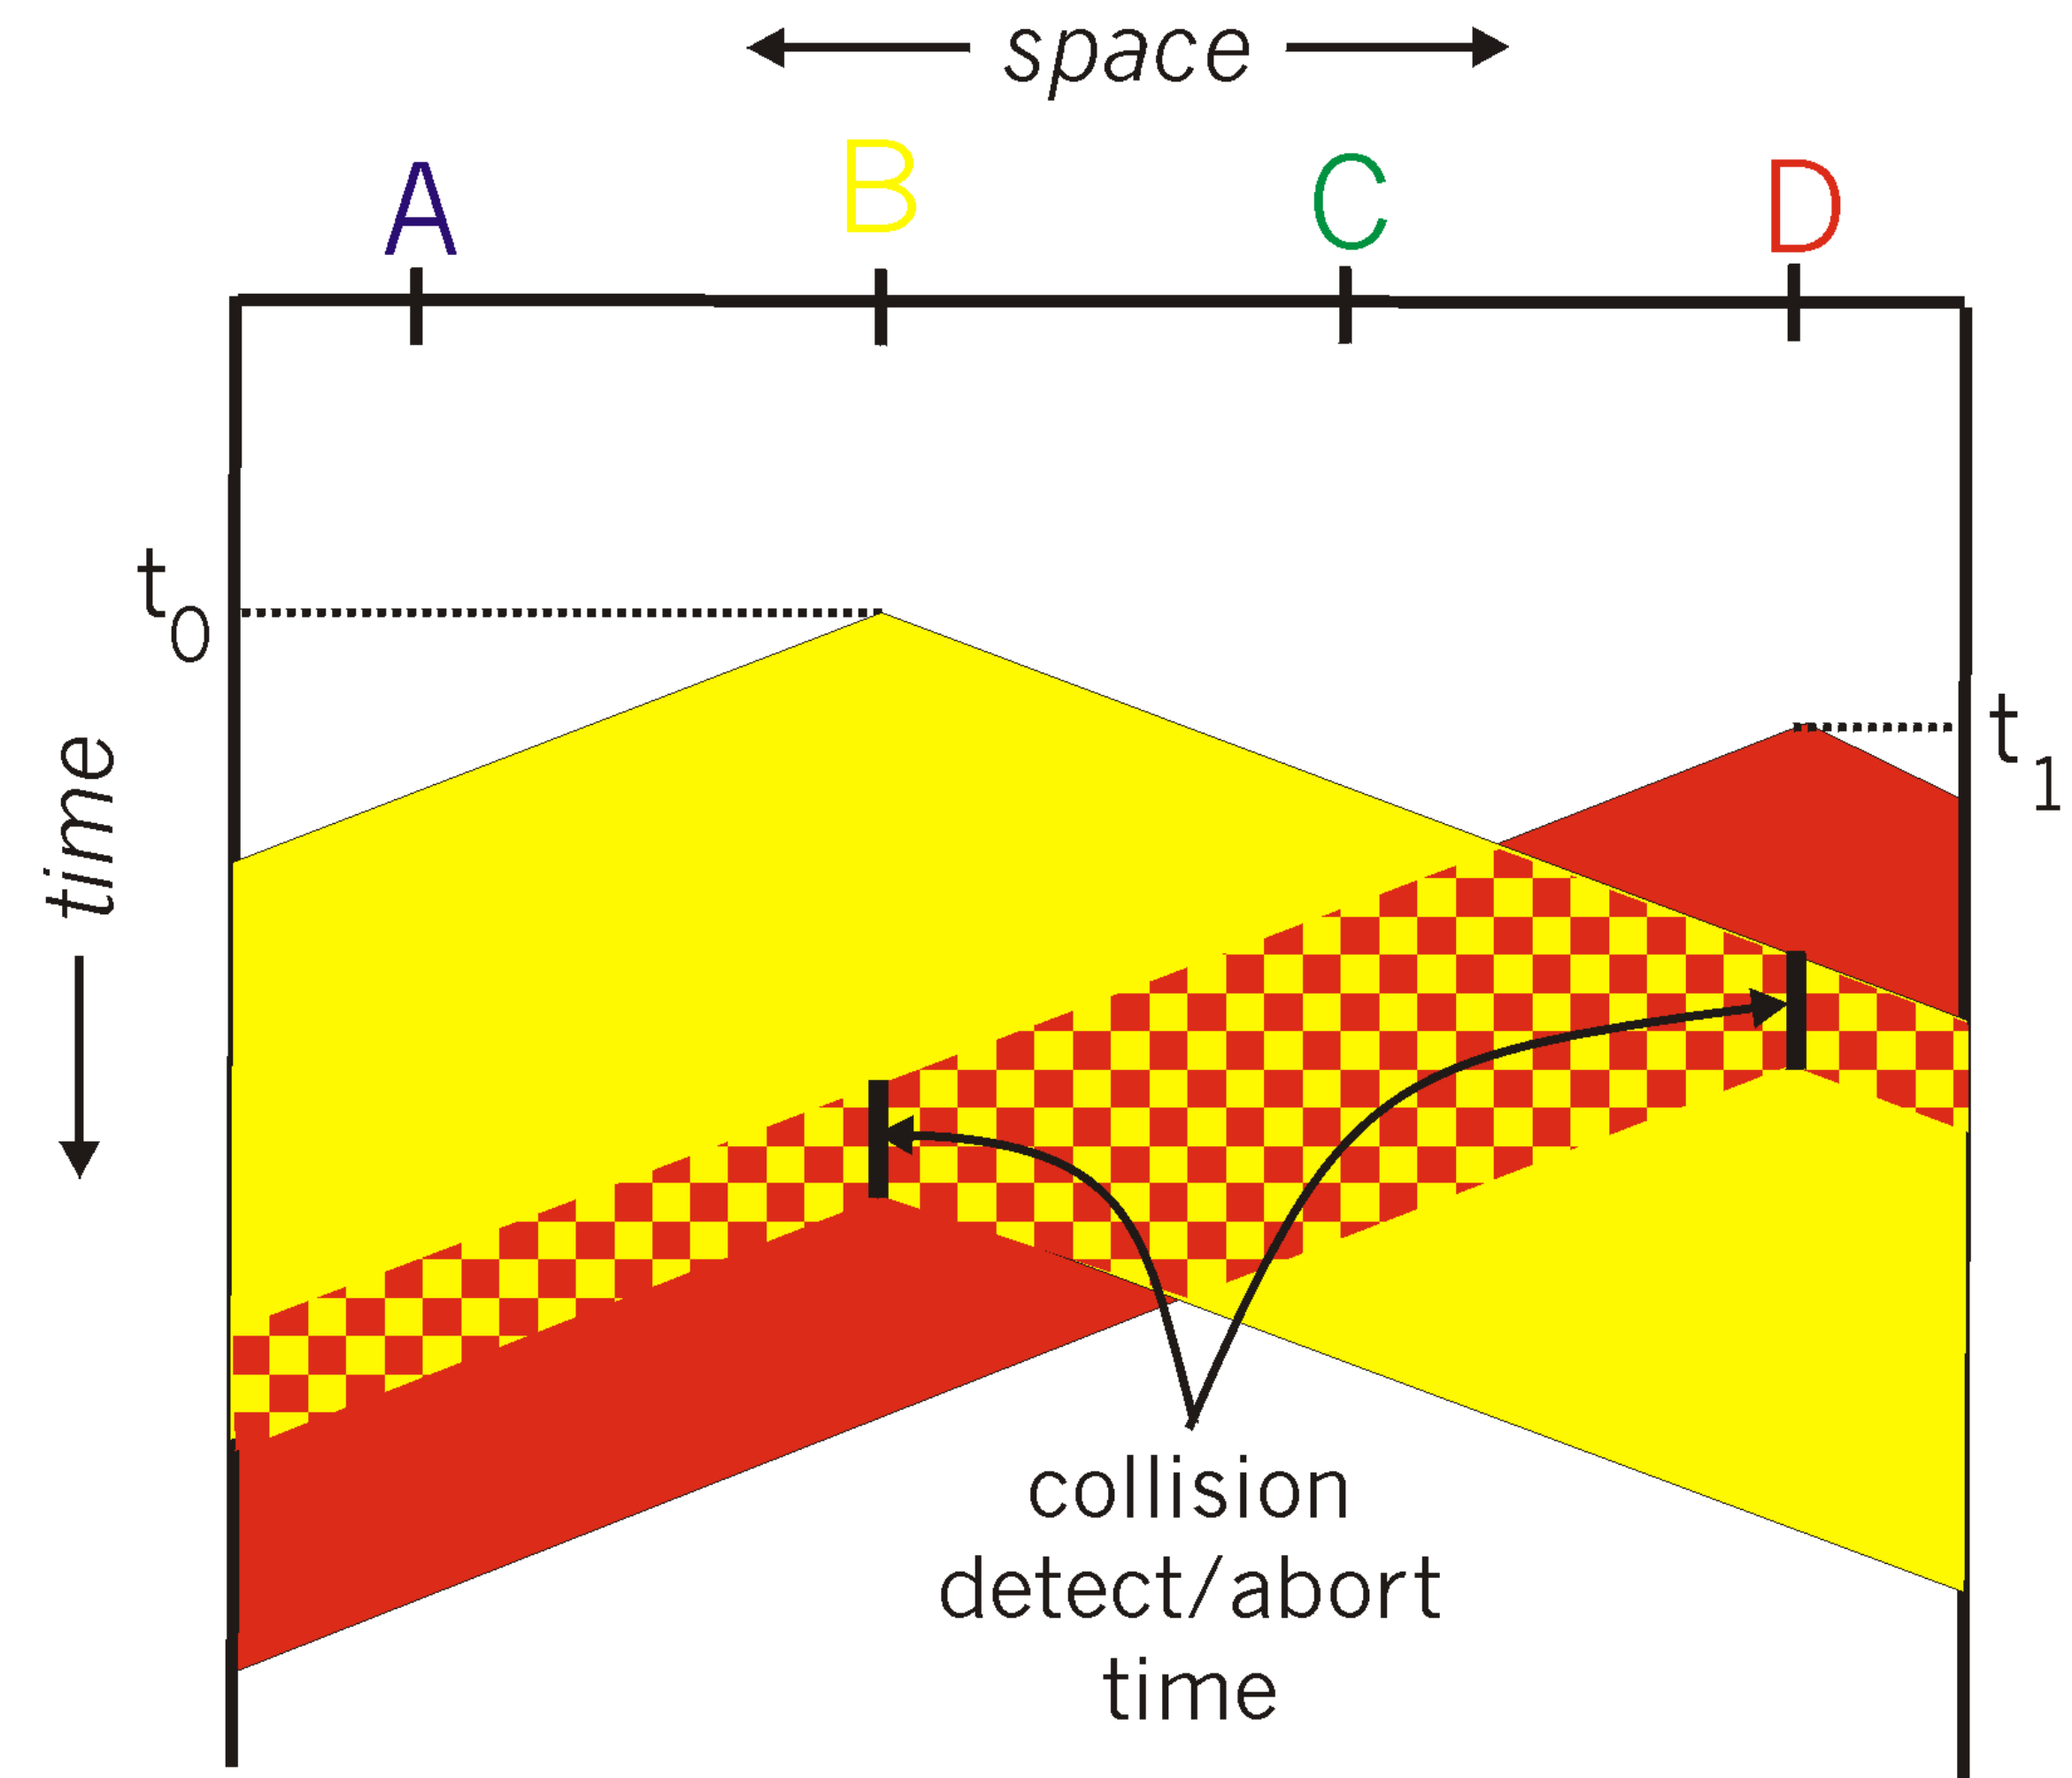


CSMA/CD (Collision Detection)

- **CSMA/CD: Carrier sensing, deferral as in CSMA**
 - *Collisions detected within short time*
 - Colliding transmissions aborted, reducing wastage
- **Collision detection easy in wired (broadcast) LANs**
 - Compare transmitted, received signals
- **Collision detection difficult in wireless LANs**
 - *Will learn more in Lecture on Wireless*

CSMA Collisions

- **B** and **D** can tell the collision occurred
- They will try again...
 - *When?*



How long should you wait?

- After collision, when should you resend?
- Should it be immediate?
- Should it be a random number with a fixed distribution?

Ethernet: CSMA/CD Protocol


- Carrier sense: wait for link to be idle
- Collision Detection: listen while transmitting
 - No collision: transmission is complete
 - Collision: abort transmission & send *jam* signal
- Random access: **binary exponential backoff**
 - After collision, wait a random time before trying again
 - After m^{th} collision, chose K randomly from $\{0, \dots, 2^{m-1}\}$
 - ...and wait for K units of time before trying again
 - If transmission occurring when ready to send, wait until end of transmission (CSMA)

Questions?

Broadcast vs. Switched Ethernet

- Ethernet was invented as a broadcast technology
 - Each packet received by all attached hosts
 - CSMA/CD for media access control
- Current Ethernets are “switched”
 - Point-to-point links between switches and to hosts
 - No sharing, no CSMA/CD

Why Switched Ethernet?



- Enables concurrent communication
 - **Host A can talk to C, while B talks to D**
 - No collisions → no need for CSMA, CD
 - No constraints on link lengths, etc.

The evolution of Ethernet


- From the shared media coax cables to dedicated links
- From 3 Mbit/s experimental ethernet to 100 Gbit/s
- From electrical signaling to optical
- **Changed almost everything except the frame format**
- Lesson: the **right interface** can accommodate many changes
 - Evolve the implementation while maintaining the interface (backward compatibility)

Topics

- Frames and framing
- Addressing
- Routing
- Forwarding
- Discovery: Bootstrapping end-to-end communication

Ethernet “Frames”

- Encapsulates IP datagram




- Preamble
 - 7 bytes for clock synchronization
 - 1 byte to indicate start of frame
- **Addresses:** 6 bytes
- **Type:** 2 bytes, indicating higher-layer protocol (e.g., IP, Appletalk)
- **Data payload:** max 1500 bytes, minimum 46 bytes
- **CRC:** 4 bytes for error correction

Framing frames


- Physical layer puts bits on a link
- But, two hosts connected on the same physical medium need to be able to exchange *frames*
 - Service provided by the link layer
 - Implemented by the network adapter
- Framing problem: how does the link layer determine where each frame begins and ends?

Simple approach: count bytes

- Sender includes number of bytes in header



- Receiver extracts this number of bytes of body
- But what if the count field is corrupted?



- L2 will frame the wrong bytes → a framing error
- CRC tells you to discard this frame, but what about the next one?

Desynchronization

- Once framing on a link is desynchronized, *it can stay that way*
- Need a method to *resynchronize*

Ethernet “Frames”

- Delineate frame with special “sentinel” bit pattern
 - E.g., 01111110 → start, 01111111 → end



- Problem: what if the sentinel bits occurs in the frame?
- Solution: bit stuffing
 - Sender always inserts a 0 after five 1s in the frame contents
 - Receiver always removes a 0 appearing after five 1s

Ethernet “Frames”



- **If next bit 0, remove it, and begin counting again**
 - Because this must be a stuffed bit; we can't be at beginning/end of frame (those had six or seven 1s)
- **If next bit 1 (i.e., we have seen six 1s) then:**
 - If following bit is 0, this is start of frame
 - Because the receiver has seen 01111110
 - If following bit is 1, this is end of frame
 - Because the receiver has seen 01111111

Example: Sentinel bits

- Original data, including start/end of frame:
 - 0|1|1|1|1|1|0|0|1|1|1|1|1|0|1|1|1|1|1|0|0|1|0|1|1|1|1|1|1|1
- Sender rule: five 1s → insert a 0
 - After bit stuffing at the sender:
 - 0|1|1|1|1|1|0|0|1|1|1|1|1|0|0|1|1|1|1|1|0|0|1|1|1|1|1|00|0|0|1|0|1|1|1|1|1|1|1|1|1
- Receiver rule: five 1s and next bit 0 → remove 0
 - After bit de-stuffing at the receiver:
 - 0|1|1|1|1|1|0|0|1|1|1|1|1|1|0|1|1|1|1|1|1|0|1|1|1|1|1|10|0|1|0|1|1|1|1|1|1|1|1|1|1

Questions?